
CS2030 Lecture 8

Declarative Programming with Integer Streams

Henry Chia (hchia@comp.nus.edu.sg)

Semester 2 2018 / 2019

Lecture Outline

- ❑ Declarative versus imperative programming
- ❑ Internal versus external iteration
- ❑ Stream concepts using `java.util.stream.IntStream`
 - Stream elements
 - Stream pipelines
 - Intermediate and terminal operations
 - Lazy and eager evaluations
 - Lambda expressions
 - Mapping
 - Reduction
 - Method references
 - Infinite streams

External Iteration

- An imperative loop that specifies **how** to loop and sum

```
int sum = 0;
for (int x = 1; x <= 10; x++) {
    sum += x;
}
```

- Realize the variables `i` and `sum` *mutates* at each iteration
- Errors could be introduced when
 - `sum` is initialized wrongly before the loop
 - looping variable `x` is initialized wrongly
 - loop condition is wrong
 - increment of `x` is wrong
 - aggregation of `sum` is wrong

Internal Iteration

- A *declarative* approach that specifies **what** to do

```
int sum = IntStream  
    .rangeClosed(1, 10)  
    .sum();
```

- sum is assigned with the result of a **stream pipeline**
- Literal meaning “for the range 1 through 10, sum them”
- A **stream** is a sequence of elements on which tasks are performed; the stream pipeline moves the stream’s elements through a sequence of tasks
- No need to specify how to iterate through elements or use any *mutable* variables — no variable state, no problem 😊
- IntStream handles all the iteration details
- A key aspect of functional programming

Streams and Pipelines

- A stream pipeline starts with a **data source**
- Static method `IntStream.rangeClosed(1, 10)` creates an `IntStream` containing the ordered sequence $1, 2, \dots, 9, 10$
 - `range(1, 10)` produces the ordered sequence $1, 2, \dots, 8, 9$
- Instance method `sum` is the processing step, or **reduction**
 - it reduces the stream of values into a single value
 - Other reductions include `count`, `min`, `max`, `average`

```
long count = IntStream  
    .rangeClosed(1, 10)  
    .count();
```

```
OptionalInt max = IntStream  
    .rangeClosed(1, 10)  
    .max();  
System.out.println(max.getAsInt()
```

- Reductions are **terminal operations** that initiate a stream pipeline's processing so as to produce a result

Mapping

- Most stream pipelines contain **intermediate operations** that specify tasks to perform on a stream's elements before a terminal operation produces a result
- **Mapping** is a common intermediate operation which
 - transforms a stream's elements to new values
 - resulting stream has the same number of elements
 - type of the mapped elements can be different from that of the original stream's elements
- Example, given the following *external iteration*

```
int sum = 0;  
for (int x = 1; x <= 10; x++) {  
    sum += (2 * x);  
}
```

Mapping

- Using *internal iteration*

```
int sum = IntStream
    .rangeClosed(1, 10)
    .map(/* operation that maps x to 2x */)
    .sum();
```

- map is the processing step that would map each element in the stream to that multiplied by 2, giving a stream of even integers
- From Java 9 API,

```
IntStream map(IntUnaryOperator mapper)
```

Returns a stream consisting of the results of applying the given function to the elements of this stream.

This is an intermediate operation.

IntStreams's map

- map operation takes in an instance of a `IntUnaryOperator` as argument
- `IntUnaryOperator` is a **functional interface** with a *single abstract method*

```
int applyAsInt(int operand)
```

Applies this operator to the given operand.

- The familiar `Comparator` is also a **functional interface** with a single abstract method

```
int compare(T o1, T o2)
```

- How did we pass a `Comparator` object to, say `ArrayList.sort`? Can we do something similar for `map`?

IntStreams's map

- The usual way is to create a class that implements the IntUnaryOperator interface and override the map method

```
import java.util.function.IntUnaryOperator;

class MultiplyByTwo implements IntUnaryOperator {
    @Override
    public int applyAsInt(int x) {
        return 2 * x;
    }
}
```

```
int sum = IntStream
    .rangeClosed(1, 10)
    .map(new MultiplyByTwo())
    .sum();
```

Anonymous Inner Class

- Rather than creating another class and pass an instance of the class to map, we can replace the argument with an anonymous inner class definition instead

```
int sum = IntStream
    .rangeClosed(1, 10)
    .map(new IntUnaryOperator() {
        @Override
        public int applyAsInt(int x) {
            return 2 * x;
        }
    })
    .sum();
```

- Which part of the anonymous inner class is *really* the useful bit? Can we simplify it?

Anonymous Method: Lambda Expression

- Class and method names (IntUnaryOperator and applyAsInt) do not add value
- Use an *anonymous method* without a name

```
int sum = IntStream
    .rangeClosed(1, 10)
    .map((int x) -> { return 2 * x; })
    .sum();
```

- Lambda expression (Lambda): `(int x) -> {return 2 * x;}`
 - receives an integer parameter x and returns that value multiplied by two, much like

```
int applyAsInt(int x) {
    return 2 * x;
}
```

Lambda Expression

- Lambda syntax: $(parameterList) \rightarrow \{statements\}$
- Lambda does not require a method name, and the compiler infers the return type
- Other lambda variants: no need parameter type
 - $(x) \rightarrow \{\text{return } 2 * x;\}$: compiler infers parameter type
 - $(x) \rightarrow 2 * x$: body contains a single expression
 - $x \rightarrow 2 * x$: only one parameter
 - $() \rightarrow \text{System.out.println("Lambdas!!!")}$
- Methods can now be treated as data! 😎
 - pass lambdas as arguments to other methods (like map)
 - assign lambdas to variables for later use
 - return lambdas from methods

Intermediate and Terminal Operations

- Intermediate operations (like map) use **lazy evaluation**
- Does not perform any operations on stream's elements until a terminal operation is called. Using filtering as an example,

- Select elements that match a condition, or **predicate**

```
int sum = 0;
for (int x = 1; x <= 10; x++) {
    if (x % 2 == 0) {
        sum += (2 * x);
    }
}
```

allows true to pass

```
int sum = IntStream
    .rangeClosed(1, 10)
    .filter(x -> x % 2 == 0)
    .map(x -> 2 * x)
    .sum();
```

- `filter` receives a method that takes one parameter and returns a **boolean** result; if it is true the element is included in the resulting stream
- Terminal operation use **eager evaluation**, i.e. perform the requested operation when they are called

Stream Elements

- Each intermediate operation results in a new stream
- Each new stream is an object representing the processing steps that have been specified up to that point in the pipeline
 - Chaining intermediate operations adds to the set of processing steps to perform on each stream element
 - The last stream object contains all processing steps to perform on each stream element
- When initiating a stream pipeline with a terminal operation, the intermediate operations' processing steps are applied one stream element after another
- Stream elements within a stream can only be consumed once
 - Cannot iterate through a stream multiple times

Stream Elements

- The following illustrates the movement of stream elements

```
int sum = IntStream
    .rangeClosed(1, 10)
    .filter(
        x -> {
            System.out.println("filter: " + x);
            return x % 2 == 0;
        })
    .map(
        x -> {
            System.out.println("map: " + x);
            return 2 * x;
        })
    .sum();
System.out.println(sum);
```

```
filter: 1
filter: 2
map: 2
filter: 3
filter: 4
map: 4
filter: 5
filter: 6
map: 6
filter: 7
filter: 8
map: 8
filter: 9
filter: 10
map: 10
sum is 60
```

Method References

- A lambda that simply calls another method can be replaced with just that method's name, e.g. in the `forEach` terminal

```
IntStream
    .rangeClosed(1, 10)
    .forEach(x -> System.out.println(x));
```

- Using method reference

```
IntStream
    .rangeClosed(1, 10)
    .forEach(System.out::println);
```

- Types of method references:
 - reference to a static method
 - reference to an instance method
 - reference to a constructor

IntStream Operations for Arrays

- Consider the typical array operations below

```
int[] values = {7, 9, 5, 2, 8, 4, 1, 6, 10, 3};

int count = 0;
int min = values[0];
int max = values[0];
int sum = 0;
for (int x : values) {
    count++;
    if (x < min) {
        min = x;
    }
    if (x > max) {
        max = x;
    }
    sum += x;
}
double average = 1.0 * sum / values.length;
System.out.println("count: " + count);
System.out.println("sum: " + sum);
System.out.println("min: " + min);
System.out.println("max: " + max);
System.out.println("average: " + average);
```

IntStream Operations for Arrays

- Using IntStream operations

```
int[] values = {7, 9, 5, 2, 8, 4, 1, 6, 10, 3};

System.out.println("count: " +
    IntStream.of(values).count());
System.out.println("sum: " +
    IntStream.of(values).sum());
System.out.println("min: " +
    IntStream.of(values).min().getAsInt());
System.out.println("max: " +
    IntStream.of(values).max().getAsInt());
System.out.println("average: " +
    IntStream.of(values).average().getAsDouble());
```

- `IntStream.of(int... values)` creates an `IntStream` from the array `values`; `T...` is the *varargs* construct for variable-length arguments
- `min`, `max` returns `OptionalInt`; `average` returns `OptionalDouble`
- Use `getAsInt()` and `getAsDouble()` correspondingly since we know there are elements in the stream

User-defined Reductions

- Using IntStream's reduce method
- Terminal operations are specific implementations of reduce
- For example, using reduce in place of sum

```
IntStream  
    .of(values)  
    .reduce(0, (x, y) -> x + y)
```

- First argument to reduce is the operation's identity value
- Second argument is the lambda that receives two int values, adds them and returns the result; in the above
 - ▷ First calculation uses identity value 0 as left operand
 - ▷ Subsequent calculations uses the result of the prior calculation as the left operand
 - ▷ If stream is empty, the identity value is returned

Boolean Terminal Operations

- Useful terminal operations that return a **boolean** result
 - `noneMatch` returns **true** if none of the elements pass the given predicate
 - `allMatch` returns **true** if every element passes the given predicate
 - `anyMatch` returns **true** if at least one element passes the given predicate
- Example: primality checking using external iteration

```
static boolean isPrime(int n) {  
    for (x = 2; x < n; x++) {  
        if if (n % x == 0) {  
            return false;  
        }  
    }  
    return true;  
}
```

Boolean Terminal Operations

- Using streams

```
static boolean isPrime(int n) {  
    return IntStream  
        .range(2, n)  
        .noneMatch(x -> n % x == 0);  
}
```

- How about finding the first 500 prime numbers?

```
static void fiveHundredPrime() {  
    int count = 0;  
    int i = 2;  
    while (count < 500) {  
        if (isPrime(i)) {  
            System.out.println(i);  
            count++;  
        }  
        i++;  
    }  
}
```

Infinite Stream

- Lazy evaluation allows us to work with infinite streams that represent an infinite number of elements
- Since streams are lazy until a terminal operation is performed, intermediate operations can be used to restrict the total number of elements in the stream
- `iterate` generates an ordered sequence starting using the first argument as a seed value

```
static void fiveHundredPrime() {  
    IntStream  
        .iterate(2, x -> x + 1)  
        .filter(x -> isPrime(x))  
        .limit(500)  
        .forEach(System.out::println);  
}
```

Infinite Stream to Finite Stream

- `limit` takes in an **int** `n` and returns a stream containing the first `n` elements of the stream
- `takeWhile` takes in a predicate and returns a stream containing the elements of the stream, until the predicate becomes false;
 - the resulting stream might still be infinite if the predicate never becomes false

```
static void primesLessThanFiveHundred() {  
    IntStream  
        .iterate(2, x -> x + 1)  
        .filter(x -> isPrime(x))  
        .takeWhile(x -> x <= 500)  
        .forEach(System.out::println);  
}
```

Lecture Summary

- Appreciate the declarative style of programming using `IntStream`
- Understand how Java Functional Interface with a single abstract method can be used in stream operations
- Familiarity with writing lambda expressions as anonymous methods/functions
- Appreciate how lazy evaluations are used for intermediate operations, eager evaluation for terminal operations
- Know how to define reductions for use in a stream pipeline
- Appreciate how lazy evaluations support infinite streams