

---

# CS2030 Lecture 11

## Parallel and Concurrent Programming

Henry Chia (hchia@comp.nus.edu.sg)

Semester 2 2018 / 2019

# Lecture Outline

---


- Concurrency versus parallelism
- Parallel streams
  - Debugging parallel streams
  - Comparing sequential and parallel streams
- Correctness of parallel streams
  - reduce operator
  - Accumulator and combiner
- Fork and join in parallel streams
  - Overhead of parallelization

# Concurrency vs Parallelism

---

- A single core processor executes one instruction at a time
  - Only one process can run at any one time
  - Context-switching allows multi-tasking on a single processor
- Concurrent programs run concurrently via threads
  - OS switches between threads
  - Separate unrelated tasks into separate threads
  - Improves processor utilization
- Parallel computing involves multiple subtasks running at the same time on multiple (possibly multi-core) processors
- Parallel programs are concurrent, but not all concurrent programs are parallel

# Concurrency vs Parallelism



# Parallel Streams

---

- ❑ Execute Streams in parallel to increase runtime performance
  - ❑ Parallel streams use a common `ForkJoinPool` via the static `ForkJoinPool.commonPool()` method  

```
System.out.println(ForkJoinPool.commonPool().getParallelism());
```
  - ❑ The level of parallelism can be controlled by setting the following system property, either using  

```
System.setProperty("java.util.concurrent.ForkJoinPool.common.parallelism", "4")
```

or including the following flag when you running the program  

```
-Djava.util.concurrent.ForkJoinPool.common.parallelism=4
```
  - ❑ Collections support the method `parallelStream()` to create a parallel stream of elements
  - ❑ Alternatively, the intermediate operation `parallel()` can be invoked on a given stream to parallelize a sequential stream
-


# Parallel Streams

- Using prime number testing as an example

```
static boolean isPrime(int n) {  
    return IntStream  
        .rangeClosed(2, (int) Math.sqrt(n))  
        .noneMatch(x -> n % x == 0);  
}
```

- Count number of primes between 2,000,000 and 3,000,000

```
long count = IntStream.range(2_000_000, 3_000_000)  
    .filter(x -> isPrime(x))  
    .count();
```




# Parallel Streams

- Parallelizing the search for primes

```
long count = IntStream.range(2_000_000, 3_000_000)  
    .parallel()  
    .filter(x -> isPrime(x))  
    .count();
```

- The `parallel()` intermediate operation turns on a boolean flag that switches the stream pipeline to be parallel
  - Invoked anywhere between the data source and terminal
  - The counter operation is `sequential()`



# Parallel Streams

---

```
public static void main(String[] args) {  
    if (args.length != 0) {  
        System.setProperty(  
            "java.util.concurrent.ForkJoinPool.common.parallelism",  
            args[0]);  
    }  
    System.out.println("Number of worker threads: " +  
        ForkJoinPool.commonPool().getParallelism());  
  
    Instant start = Instant.now();  
    long howMany = IntStream.range(2_000_000, 3_000_000)  
        .parallel()  
        .filter(x -> isPrime(x))  
        .count();  
    Instant stop = Instant.now();  
  
    System.out.println(howMany + " : " +  
        Duration.between(start, stop).toMillis() + "ms");  
}
```



# Debugging Parallel Streams

---

- To time the execution of a process,
  - `java.time.Instant`'s `now()` method returns the current `Instant` from the system clock
  - `java.time.Duration`'s `between()` returns the `Duration` of two `Instances` (an implementation of `Temporal`)
  - `Duration`'s `toMillis()/toNanos()/...` extracts the desired representation of the duration

```
java.util.Instant;  
java.util.Duration;
```

```
Instant start, stop;  
start = Instant.now();  
/* perform some task */  
stop = Instant.now();
```

```
long timeInMillis = Duration.between(start, stop).toMillis();
```

# Debugging Parallel Streams

- To debug and manage each execution thread
  - `Thread.currentThread()` (or `Thread.currentThread().getName()`) to retrieve the identity of the thread
  - `Thread.sleep(long millis)` causes the currently executing thread to sleep (i.e. temporarily cease execution) for the specified number of milliseconds
    - ▷ Used within a **try.. catch** block
    - ▷ Example, letting a thread sleep for one second

```
try {  
    ...  
    Thread.sleep(1000);  
    ...  
} catch (InterruptedException e) { }
```

# Debugging Parallel Streams

- Effect of parallelizing a stream

```
int sum = IntStream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter: " + x + " "
            + Thread.currentThread().getName());
        return true;
    })
    .map(x -> {
        System.out.println("map: " + x + " "
            + Thread.currentThread().getName());
        return x;
    })
    .reduce(0, (x, y) -> {
        System.out.println("reduce: " + x + " + " + y + " "
            + Thread.currentThread().getName());
        return x + y;
    });
System.out.println(sum);
```

# Comparing Sequential and Parallel Streams

- Suppose given the following task unit

```
class OneSecondTask {
    int ID;

    public OneSecondTask(int ID) {
        this.ID = ID;
    }

    public int compute() {
        System.out.println(Thread.currentThread().getName());
        try {
            Thread.sleep(1000);
        } catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        return ID;
    }
}
```

# Comparing Sequential and Parallel Streams

```
public static void sequentialRun(List<OneSecondTask> tasks) {  
    Instant start = Instant.now();  
    List<Integer> result = tasks.stream()  
        .map(x -> x.compute())  
        .collect(Collectors.toList());  
    Instant stop = Instant.now();  
    System.out.print(result + " ");  
    System.out.println(Duration.between(start, stop).toMillis() + "ms")  
}
```

□ Sequential stream on 4 worker threads:

```
main  
main  
main  
main  
main  
main  
main  
main  
main  
main  
main  
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 10003ms
```

# Comparing Sequential and Parallel Streams

```
public static void parallelStreamRun(List<OneSecondTask> tasks) {  
    Instant start = Instant.now();  
    List<Integer> result = tasks.parallelStream()  
        .map(x -> x.compute())  
        .collect(Collectors.toList());  
    Instant stop = Instant.now();  
    System.out.print(result + " ");  
    System.out.println(Duration.between(start, stop).toMillis() + "ms")  
}
```

- Parallel stream on 4 worker threads:

```
main  
ForkJoinPool.commonPool-worker-1  
ForkJoinPool.commonPool-worker-3  
ForkJoinPool.commonPool-worker-2  
ForkJoinPool.commonPool-worker-4  
ForkJoinPool.commonPool-worker-3  
ForkJoinPool.commonPool-worker-2  
ForkJoinPool.commonPool-worker-4  
ForkJoinPool.commonPool-worker-1  
ForkJoinPool.commonPool-worker-3  
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 3006ms
```

# Correctness of Parallel Streams

- To ensure correct parallel execution, stream operations must not interfere with stream data, preferably stateless and have no side effects

- Example of interference:

```
List<String> list = new ArrayList<>(
    List.of("abc", "def", "xyz"));

list.stream()
    .peek(str -> {
        if (str.equals("xyz")) {
            list.add("pqr");
        }
    })
    .forEach(x -> {});
```

- Interference is not allowed in both sequential and parallel streams

# Correctness of Parallel Streams

- Another example:

```
List<Integer> list = new ArrayList<>(
    Arrays.asList(1, 3, 5, 7, 9, 11, 13, 15, 17, 19));
List<Integer> result = new ArrayList<>();
```

- The following is erroneous due to side effects

```
list.parallelStream() // list.stream().parallel()
    .filter(x -> isPrime(x))
    .forEach(x -> result.add(x));
```

- Use `.collect` instead

```
result = list.parallelStream()
    .filter(x -> isPrime(x))
    .collect(Collectors.toList());
```

- Side effects are a problem in parallel streams
- Consider using a thread-safe list, e.g. `CopyOnWriteArrayList`



# Inherently Parallelizable reduce

- Consider Stream's three-argument reduce method:

```
<U> U reduce(U identity,  
            BiFunction<U,? super T,U> accumulator,  
            BinaryOperator<U> combiner)
```

- Rules to follow when parallelizing

- `combiner.apply(identity, i)` must be equal to `i`
- combiner and accumulator must be associative, i.e. order of application does not matter
- combiner and accumulator must be compatible, i.e. `combiner.apply(u, accumulator.apply(identity, t))` must be equal to `accumulator.apply(u, t)`
- The following example compiles with the above rules:

```
Stream.of(1,2,3,4)  
    .parallel()  
    .reduce(1, (x,y) -> x * y, (x,y) -> x * y)
```


# Accumulator and Combiner

- Accumulator and combiner functions are executed in parallel

```
String s = Stream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter: " + x + " "
            + Thread.currentThread().getName());
        return true;
    })
    .map(x -> {
        System.out.println("map: " + x + " "
            + Thread.currentThread().getName());
        return x;
    })
    .reduce("",
        (x, y) -> {
            System.out.println("accumulate: " + x + " + " + y + " "
                + Thread.currentThread().getName());
            return x + y;
        },
        (x, y) -> {
            System.out.println("combine: " + x + " + " + y + " "
                + Thread.currentThread().getName());
            return x + y;
        }
    );
```

NEVER System.out.println 2 times when using thread

# Accumulator and Combiner



# Accumulator and Combiner


- Erroneous examples where rules are violated
  - `combiner.apply(identity, i)` not equal to `i`  
**double** `result = Stream`  
    `.parallel()`  
    `.of(1, 2, 3, 4)`  
    `.reduce(1.0,`  
        `(x, y) -> x + y,`  
        `(x, y) -> x + y);`
  - for division, the order of application **does** matter

```
double result = Stream
    .of(1, 2, 3, 4)
    .parallel()
    .reduce(24.0,
        (x, y) -> 1.0 * x / y,
        (x, y) -> 1.0 * x / y);
```

# Fork and Join

- Given the following program fragment and *computation* graph

```
b = f(a);  
c = g(b);  
d = h(b);  
e = n(c,d);
```



- $f(a)$  invoked before  $g(b)$  and  $h(b)$ ;  $n(c,d)$  invoked after
- How about the order of  $g(b)$  and  $h(b)$ ?
  - If  $g$  and  $h$  does not produce side effects, then parallelize
  - **Fork** task  $g$  to execute at the same time as  $h$ , and **join** back task  $g$  later

# Fork and Join in Parallel Streams

- `parallel()` runs `fork` to create sub-tasks running the same chain of operations on sub-streams
  - Processes for sub-tasks are run in multiple threads when appropriate
  - Threads are shared from a common **Fork Join Pool**
- `combiner` in `reduce` runs `join` to combine the results
- Should we exploit parallelism to the fullest?
  - Parallelizing primality testing

```
return IntStream
    .rangeClosed(2, (int) Math.sqrt(n))
    .parallel()
    .noneMatch(x -> n % x == 0);
```

# Fork and Join in Parallel Streams

---

- Parallelizing a trivial task actually creates more work in terms of parallelizing overhead
- Parallelization is worthwhile only if the task is complex enough that the benefit of parallelization outweighs the overhead
  - In primality testing, checking  $(n \% x == 0)$  is trivial;
  - Parallelizing it induces more overhead in terms of processing the forks and joins
- Holds true for all parallel and concurrent programs, either in the context of parallel streams or otherwise

# Lecture Summary

---

- Familiarity with the use of sequential and parallel streams
- Able to compare performances between sequential and parallel streams
- Able to debug parallel streams
- Adherence to rules for parallelizing streams
- Appreciate fork and join in parallel streams
- Appreciate fork/join overhead