

# LECTURE 20: SUMMARY AND CLOSING LECTURE

---

Harold Soh  
[harold@comp.nus.edu.sg](mailto:harold@comp.nus.edu.sg)

# ADMINISTRATIVE ISSUES

TAs are grading your assignments

- All scores (except late submissions) will be online by next Monday (18<sup>th</sup>).
- **Please verify** (1 participation point!)

Problem Set 5 was due yesterday!

- All Problem Sets should be submitted by midnight Friday (Nov 15<sup>th</sup>).
- **Email me and your DG TA** if you make a late submission!
- Late Submissions
  - 24 hours: 20% penalty
  - 2 weeks: 40% penalty
  - Last Second of Last Day of class: 60% penalty

No lecture tomorrow. DGs and Tutorials are still on.

# **FINAL EXAM**

**4 DEC 2018 (Morning, 9-11am)**

**2 hours**

**44 Questions (120 points)**

**All multiple choice**

- Bring pencils to shade in the form

**Open-book exam**



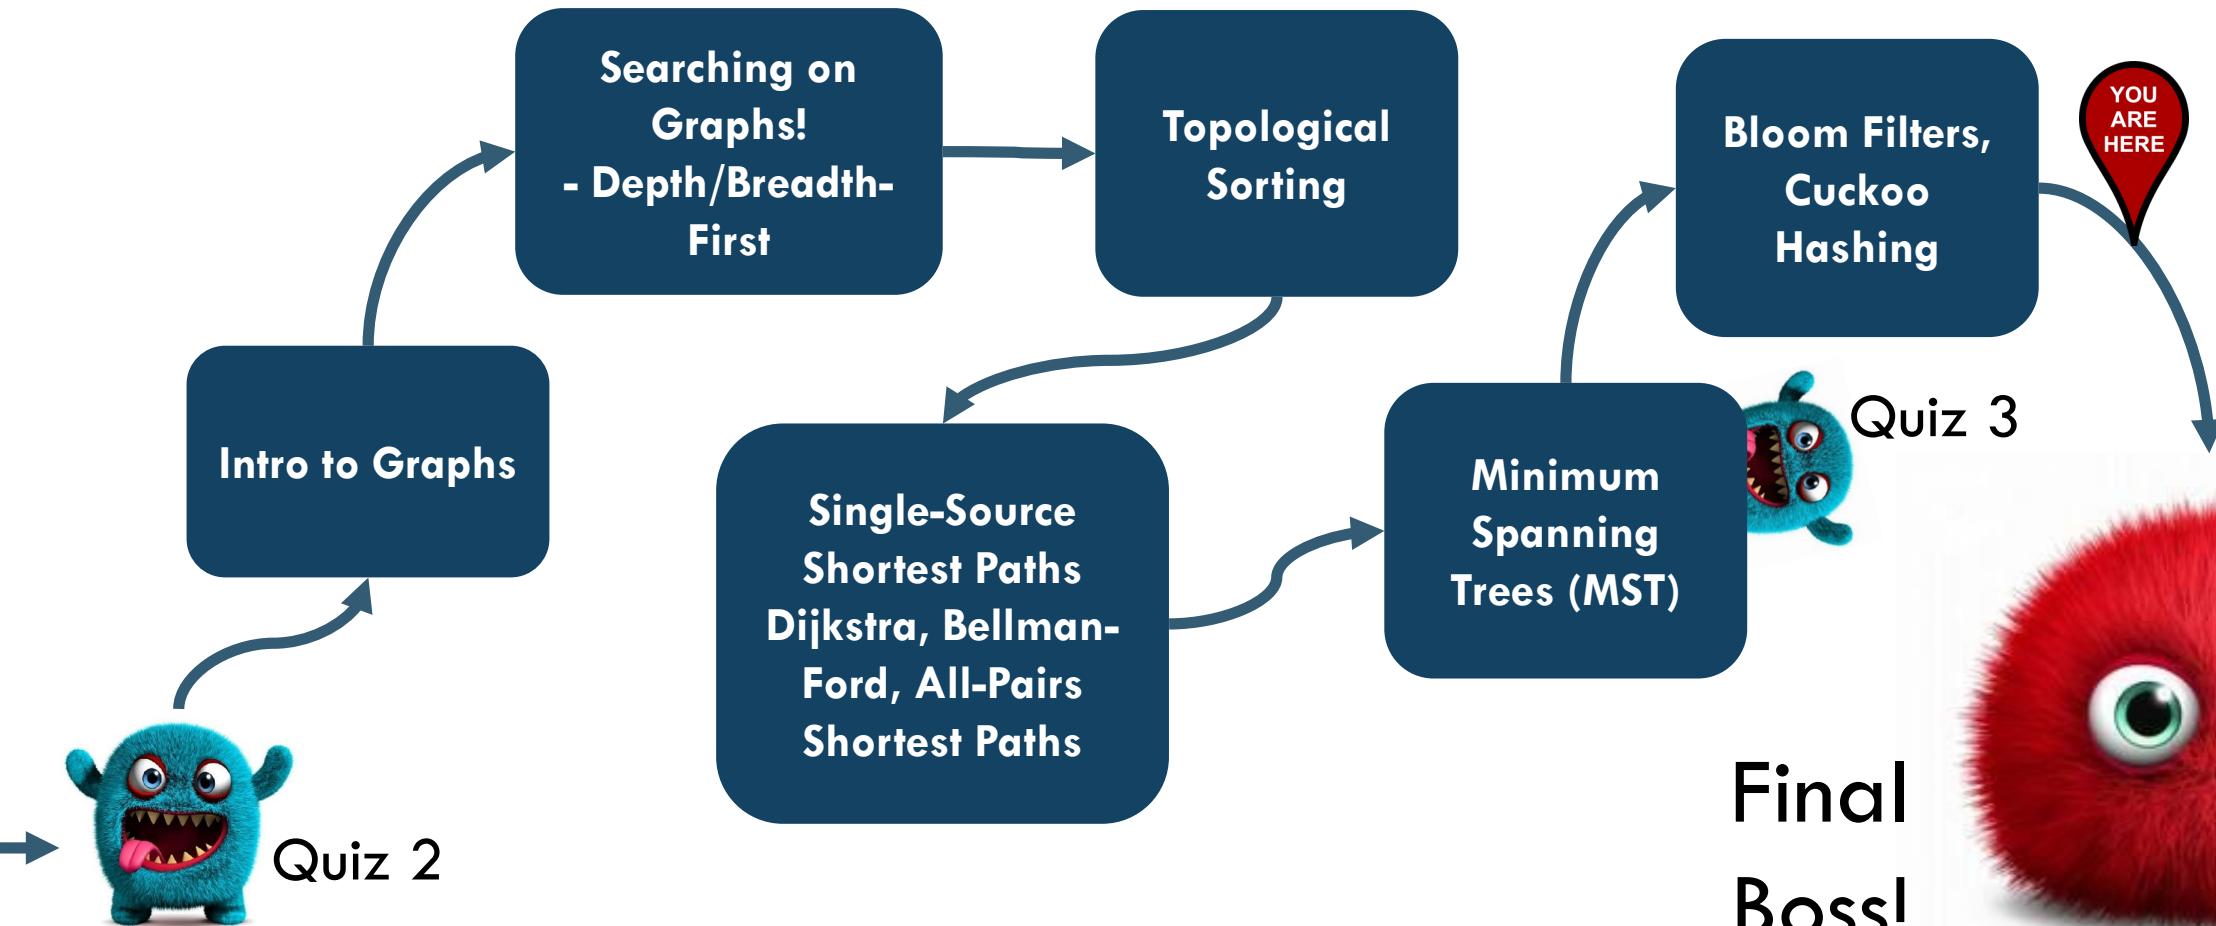


# QUESTIONS?

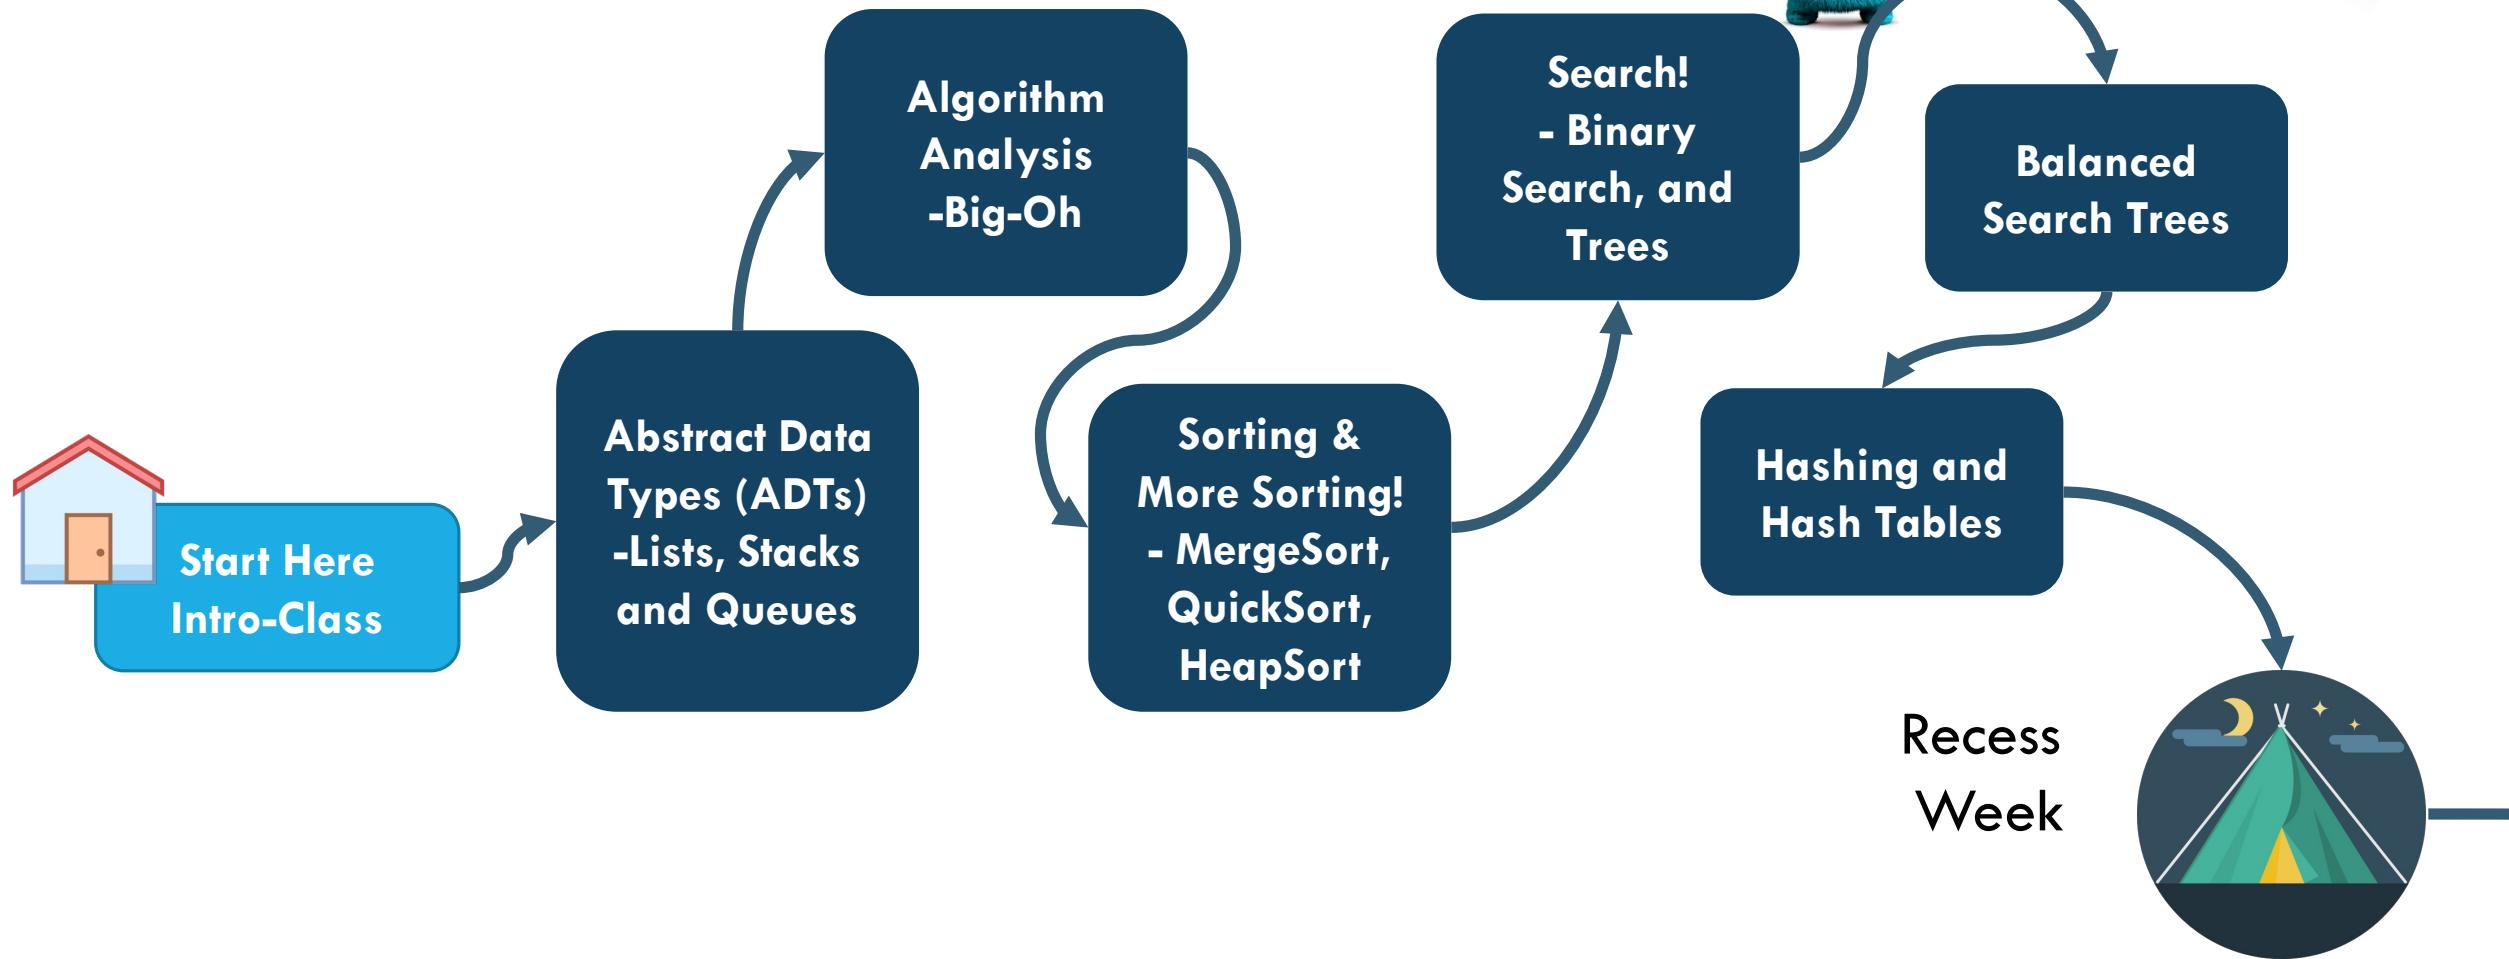


# PATH TO MASTERY / COURSE STRUCTURE

DRAFT



# PATH TO MASTERY / COURSE STRUCTURE



# CS2040S GAME SHOW!!

4 teams

Answer on polleverywhere

Let's go there now!

# QUESTION

CS2040 is designed to develop which of these skills?

- A. Monkeying around.
- B. Problem solving with algorithmic thinking.
- C. Ultra-Competitive programming.
- D. Eating cake.

# QUESTION

CS2040 is designed to develop which of these skills?

- A. Monkeying around.
- B. **Problem solving with algorithmic thinking.**
- C. Ultra-Competitive programming.
- D. Eating cake.

# COMPUTER SCIENTISTS HAVE CHANGED THE WORLD...



# ...BY SOLVING PROBLEMS

CS2040S is about **solving problems**.

It prepares you for what lies ahead...

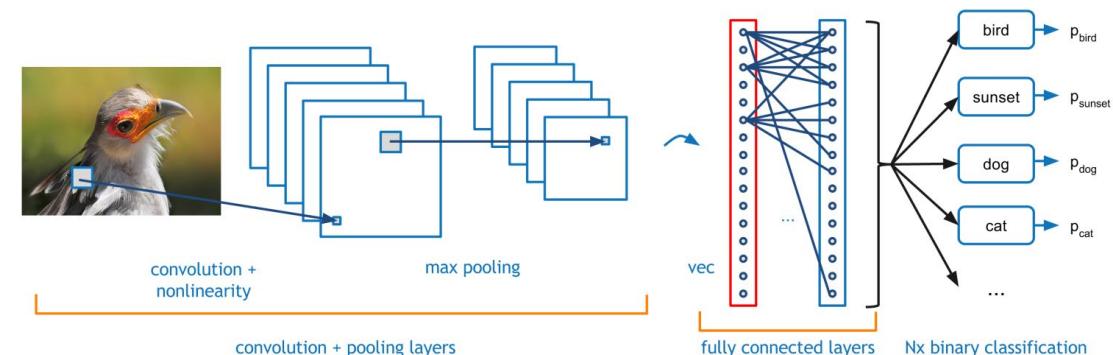
## Search



Google Search

I'm Feeling Lucky

## Image Recognition



# QUESTION

People are not perfect. Code is complex.

How can we prevent them from making stupid mistakes when working with your code?

- A. Use ADTs.
- B. Specify the invariants.
- C. Hire smarter people (or monkeys).
- D. Develop AI that does all the work for us.

# QUESTION

People are not perfect. Code is complex.

How can we prevent them from making stupid mistakes when working with your code?

- A. Use ADTs.
- B. Specify the invariants.
- C. Hire smarter people (or monkeys).
- D. Develop AI that does all the work for us.

# ABSTRACT DATA TYPES (ADTs)

Define **behavior**, not internal operation.

Interface SimpleStore  
add(x)

- **Description:** adds x to the stored items.
- **Pre:** x is the item to be added
- **Post:** x is added to the store
- **Computational cost:**
  - takes constant time
  - adds constant space (memory).

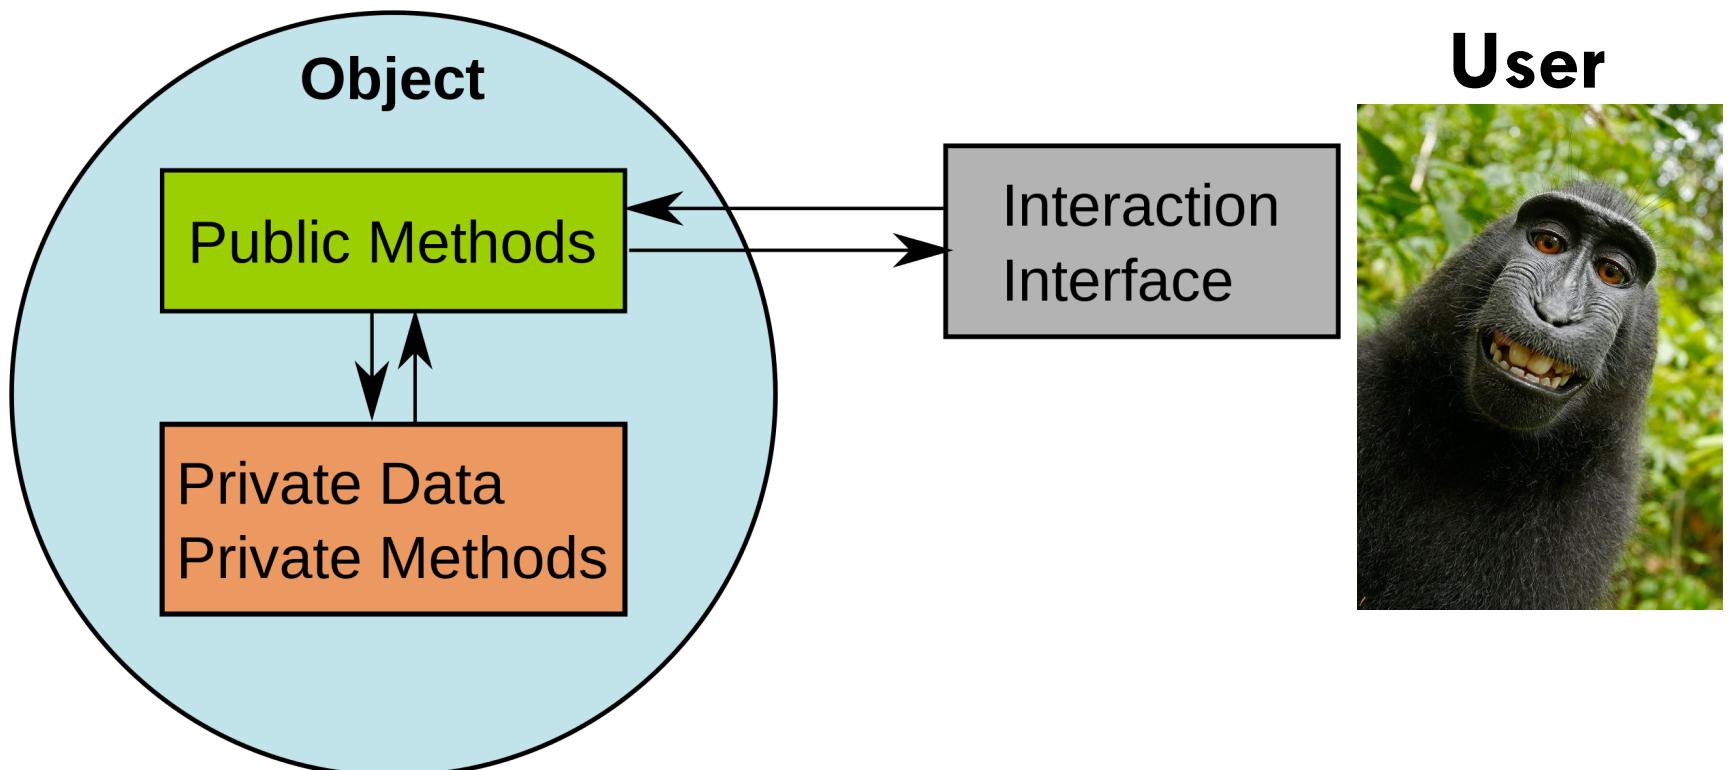
V.S.

Class SimpleStore  
add(x)

- adds x to the stored items.
- if arr has already reached its current maximum size,
  - create a new temporary array
  - copy the original array over
  - add the x
  - update the end marker
- and so on ....

# ENCAPSULATION AND INFORMATION HIDING

**Implementer**



**User**



[Image from Wikipedia commons]

# QUESTION

Which ADT can I use to transform a recursive function into an iterative function?

- A. Stack
- B. Queue
- C. Priority Queue
- D. Hash Table

# QUESTION

Which ADT can I use to transform a recursive function into an iterative function?

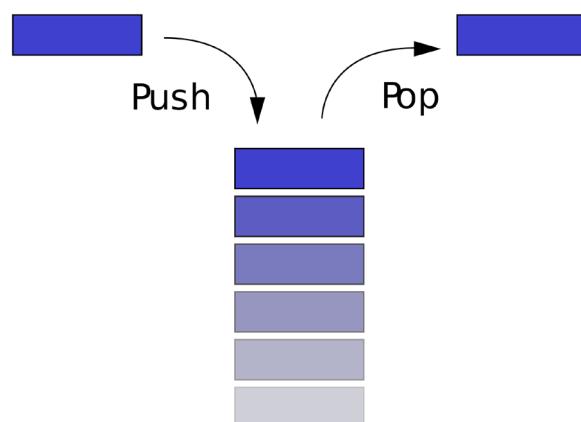
- A. Stack
- B. Queue
- C. Priority Queue
- D. Hash Table

# STACK

## Stack Operations:

- push(x)
- peek()
- pop()

## Last-in-First-Out (LIFO)



# STACKS ARE VERY USEFUL!

Implement Recursion

Expression Evaluation:  $(x+2)*(2/45) + 2$

Depth-First-Search (later in the semester!)

and others...

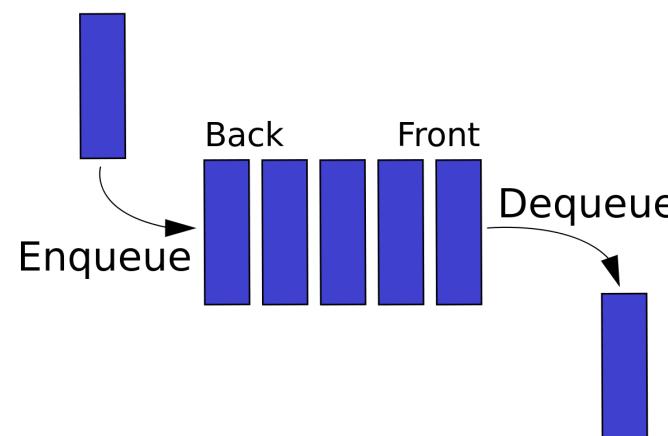


# QUEUE

## Queue Operations:

- enqueue(x)
- peek()
- dequeue()

## First-in-First-Out (FIFO)



# QUESTION

I have a word with  $n$  characters stored in a singly linked list. If I can only use  $O(1)$  space, detecting if the word is a palindrome requires:

- A.  $O(1)$  time
- B.  $O(n)$  time
- C.  $O(n \log n)$  time
- D.  $O(n^2)$  time

# QUESTION

I have a word with  $n$  characters stored in a singly linked list. If I can only use  $O(1)$  space, detecting if the word is a palindrome requires:

- A.  $O(1)$  time
- B.  $O(n)$  time
- C.  $O(n \log n)$  time
- D.  $O(n^2)$  time

# DETECTING IF A LINKED LIST IS A PALINDROME

**Key idea:**

Iterate to middle

Reverse second half of linked list

Iterate from beginning and middle, comparing the characters.

**Time Complexity:**  $O(n)$

# REVERSING A SINGLY LINKED LIST PSEUDOCODE

```
Node reversed_list = null
Node current = head
while current is not null
    Node next = current.next
    current.next = reversed_list
    reversed_list = current
    current = next
head = reversed_list
```

# QUESTION

What's the time complexity of the following piece of code?

```
int a = 0;  
for (int i=0; i<n; i++) {  
    a = a + 2;  
    A[i] = a;  
}
```

# RUNNING TIME

*Remember when I had to explain this step by step?*

Number of primitive instructions / “steps” executed.

Code	Cost	Times
int a = 0;	$c_1$	1
for (int i=0; i<n; i++) {	$c_2$	$n$
a = a + 2;	$c_3$	$n$
A[i] = a;	$c_4$	$n$
}		

$$T(n) = c_1 + nc_2 + nc_3 + nc_4 = cn + c_1$$

# QUESTION

The following functions requires:

```
function testme(n)
    if (n <= 2)
        return n
    return 1 + testme(n/2)
```

- A.  $O(1)$  time
- B.  $O(\log n)$  time
- C.  $O(n)$  time
- D.  $O(n \log n)$  time

# QUESTION

The following functions requires:

```
function testme(n)
    if (n <= 2)
        return n
    return 1 + testme(n/2)
```

- A.  $O(1)$  time
- B.  $O(\log n)$  time
- C.  $O(n)$  time
- D.  $O(n \log n)$  time

# QUESTION

**True or False:** two functions take  $O(n^2)$  time to run. I can choose either in practice.

- A. True
- B. False

# QUESTION

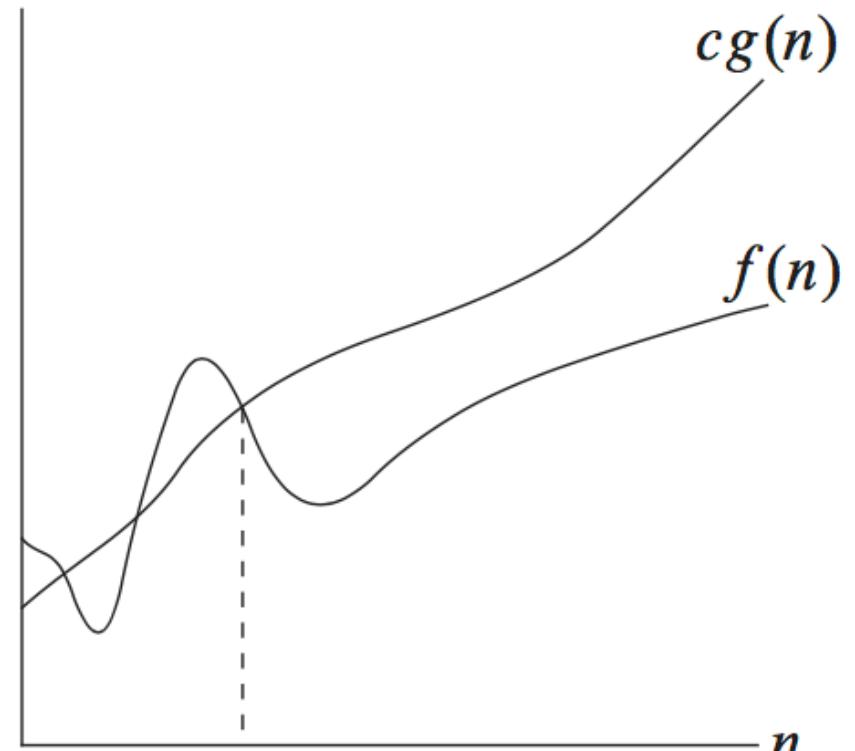
**True or False:** two functions take  $O(n^2)$  time to run. I can choose either in practice.

A. True

B. False

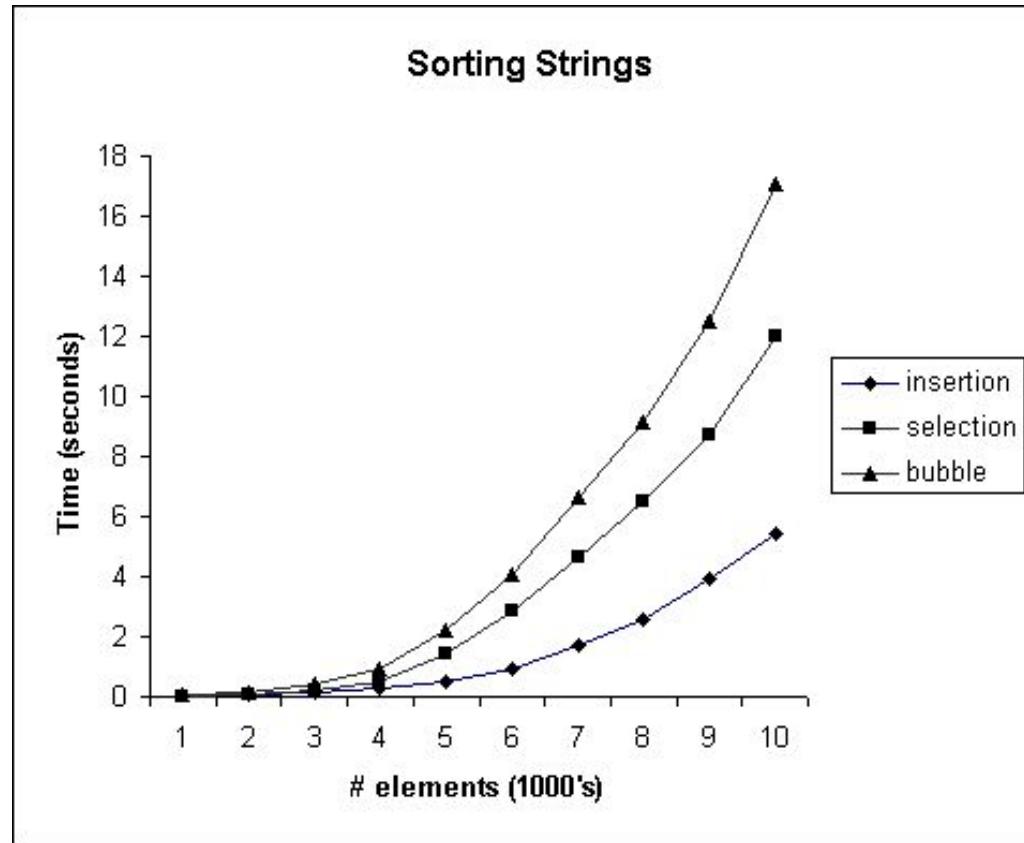
## BIG-OH $O(g(n))$

$O(g(n)) = \{f(n): \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0\}$



$$n_0 \quad f(n) = O(g(n))$$

# SOME PRACTICE WITH BIG-OH: LIMITATIONS



[Bubble Sort: An Archaeological Algorithmic Analysis, Owen Astrachan, cs.duke.edu]

# QUESTION

$T(n) = T\left(n - \frac{n}{4}\right) + T\left(\frac{n}{4}\right) + O(n)$  is bounded asymptotically by:

- A.  $O(1)$
- B.  $O(n)$
- C.  $O(n^2)$
- D.  $O(n \log n)$

# QUESTION

$T(n) = T\left(n - \frac{n}{4}\right) + T\left(\frac{n}{4}\right) + O(n)$  is bounded asymptotically by:

- A.  $O(1)$
- B.  $O(n)$
- C.  $O(n^2)$
- D.  $O(n \log n)$

# QUESTION

Which of these quotes is not by a famous computer scientist?

- A. Simplicity is a prerequisite for reliability
- B. The use of COBOL cripples the mind
- C. There should be no such thing as boring mathematics
- D. Make things as simple as possible, but not simpler
- E. It is practically impossible to teach good programming to students that have prior experience to BASIC: as potential programmers they are mentally mutilated beyond hope of regeneration.

# EDSGER W. DIJKSTRA

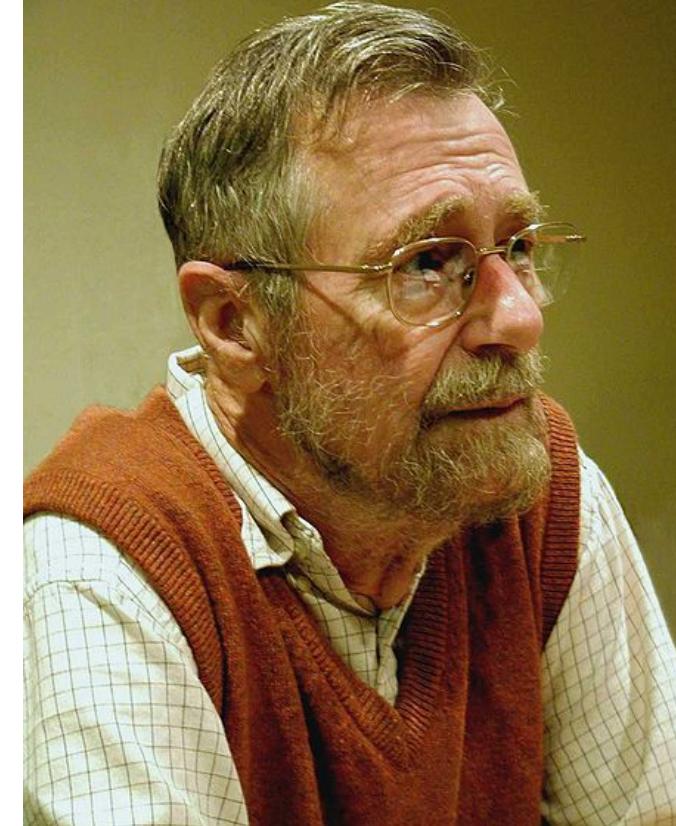
*“Computer science is no more about computers than astronomy is about telescopes.”*

*“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*

*“There should be no such thing as boring mathematics.”*

*“Elegance is not a dispensable luxury but a factor that decides between success and failure.”*

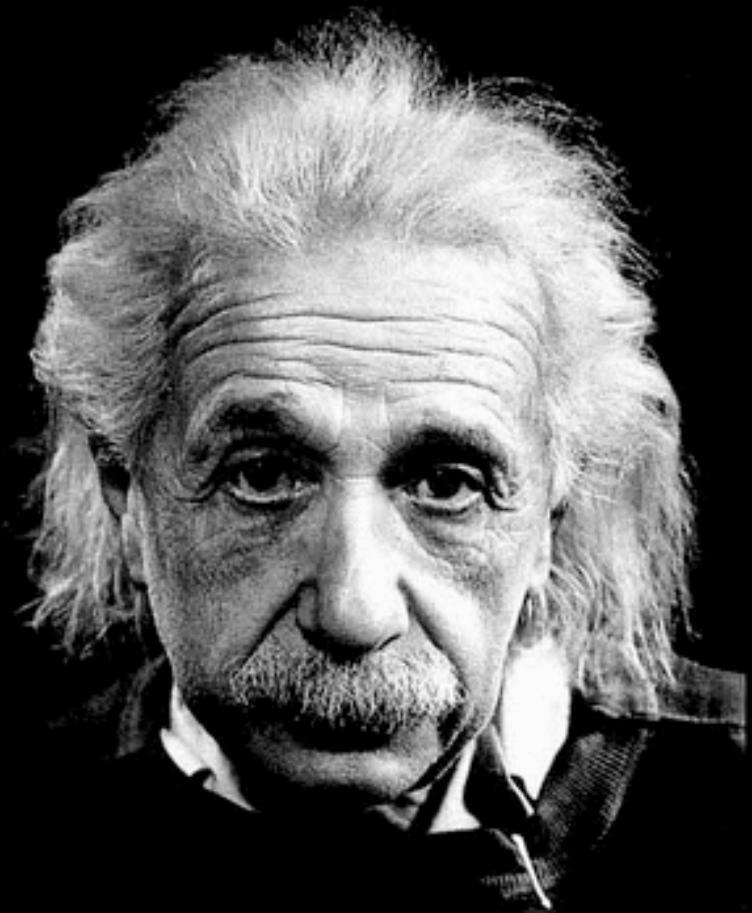
*“Simplicity is prerequisite for reliability.”*



1930-2002

“Everything should be made  
as simple as possible,  
but not simpler.”

Albert Einstein



# QUESTION

Naruto has a large tribe:  $2^{20}$  members. After sorting all of them by height, he wants find a family member that is exactly 3.14159 feet tall. What algorithm should he use?

- A. Mergesort
- B. k-way Merge
- C. Binary Search
- D. 3-way partition
- E. Quickselect

# QUESTION

Naruto has a large tribe:  $2^{20}$  members. After sorting all of them by height, he wants find a family member that is exactly 3.14159 feet tall. What algorithm should he use?

- A. Mergesort
- B. k-way Merge
- C. **Binary Search**
- D. 3-way partition
- E. Quickselect

# PSEUDOCODE FOR BINARY SEARCH

```
function binarySearch (A, key, n)
    low = 0
    high = n - 1
    while low <= high
        mid = (low + high )/2
        if key < A[mid] then
            high = mid - 1
        else if key > A[mid]
            low = mid + 1
        else
            return mid
    return not_found
```

2	5	7	11	13
---	---	---	----	----

*Which line can cause problems when used on real systems?*

# QUESTION

I have to sort a large array but on a embedded system with very little memory. I can only use Quicksort or Mergesort: which should I use?

- A. Mergesort
- B. Quicksort

# QUESTION

I have to sort a large array but on a embedded system with very little memory. I can only use Quicksort or Mergesort: which should I use?

- A. Mergesort
- B. **Quicksort**

From Sedgewick and Wayne's slides

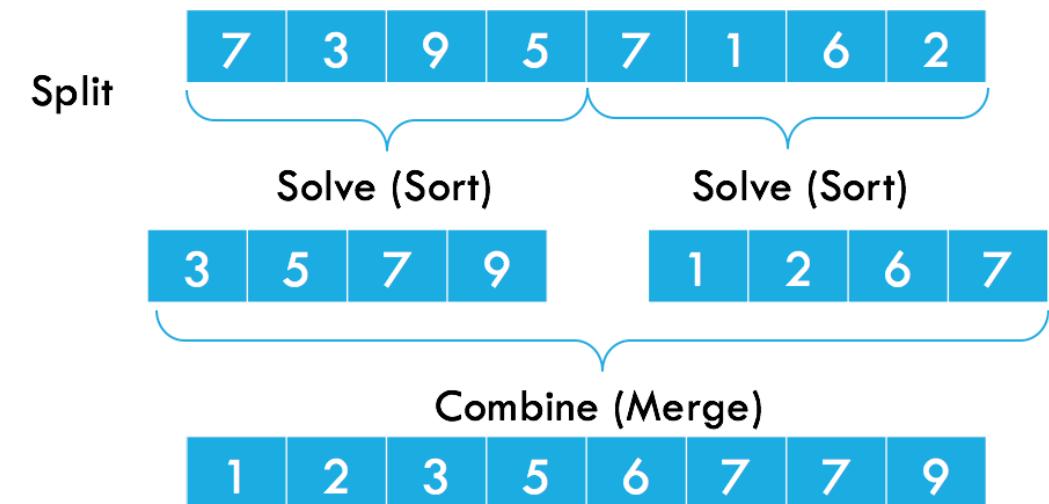
<http://algs4.cs.princeton.edu/lectures/24PriorityQueues.pdf>

## Sorting algorithms: summary

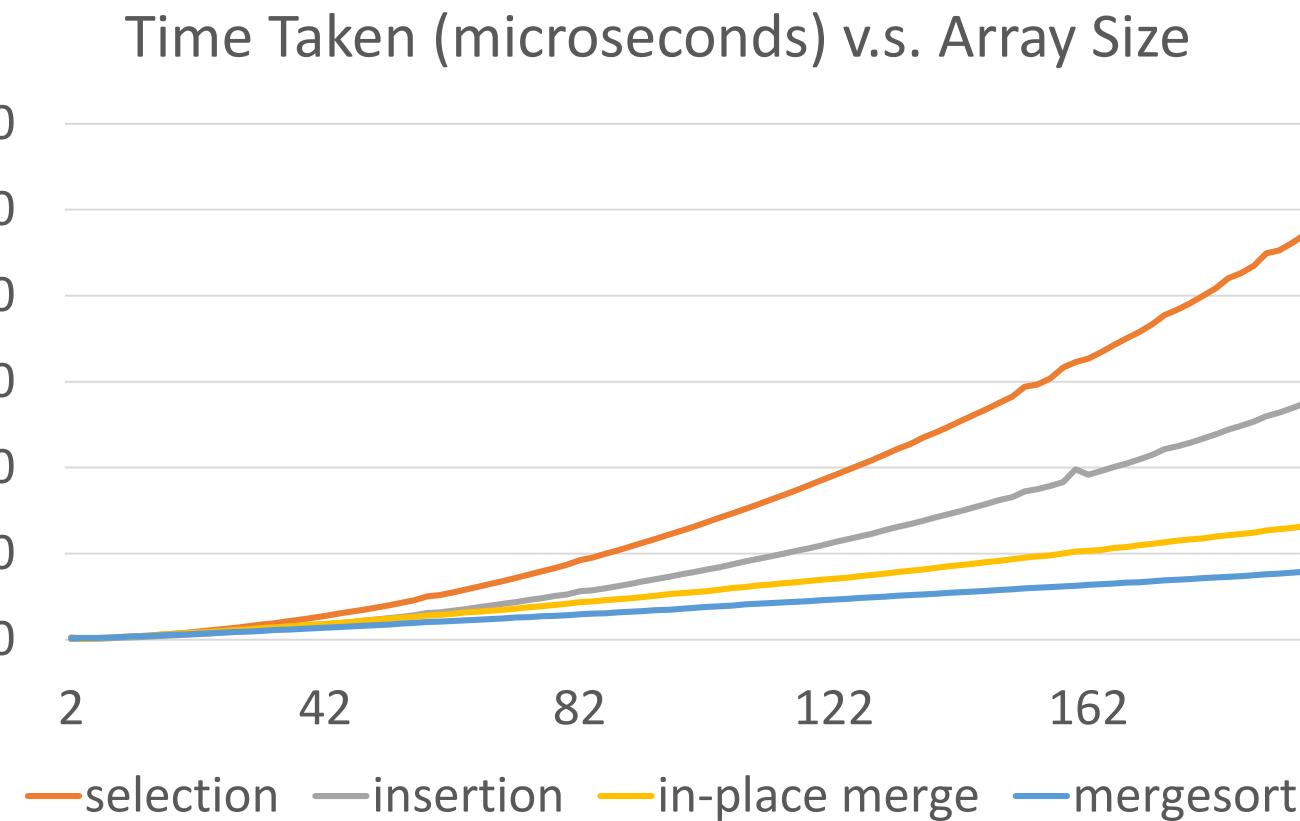
	inplace?	stable?	best	average	worst	remarks
selection	✓		$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	$n$ exchanges
insertion	✓	✓	$n$	$\frac{1}{4} n^2$	$\frac{1}{2} n^2$	use for small $n$ or partially ordered
shell	✓		$n \log_3 n$	?	$c n^{3/2}$	tight code; subquadratic
merge		✓	$\frac{1}{2} n \lg n$	$n \lg n$	$n \lg n$	$n \log n$ guarantee; stable
timsort		✓	$n$	$n \lg n$	$n \lg n$	improves mergesort when preexisting order
quick	✓		$n \lg n$	$2 n \ln n$	$\frac{1}{2} n^2$	$n \log n$ probabilistic guarantee; fastest in practice
3-way quick	✓		$n$	$2 n \ln n$	$\frac{1}{2} n^2$	improves quicksort when duplicate keys
heap	✓		$3 n$	$2 n \lg n$	$2 n \lg n$	$n \log n$ guarantee; in-place
?	✓	✓	$n$	$n \lg n$	$n \lg n$	holy sorting grail

# MERGESORT: THE ALGORITHM

```
function mergeSort(A, low, high)
    if low < high
        mid = (high + low)/2
        mergeSort(A, low, mid)
        mergeSort(A, mid+1, high)
        merge(A, low, mid, high)
```



# IN PLACE MERGE SORT?



# QUESTION

Which algorithm is this snippet of code from?

- A. Mergesort
- B. Quicksort
- C. Heapsort
- D. Selectionsort
- E. Dijkstra's SSSP

```
1 int i, j;
char x, y;

i = left; j = right;
x = items[(left + right) / 2];

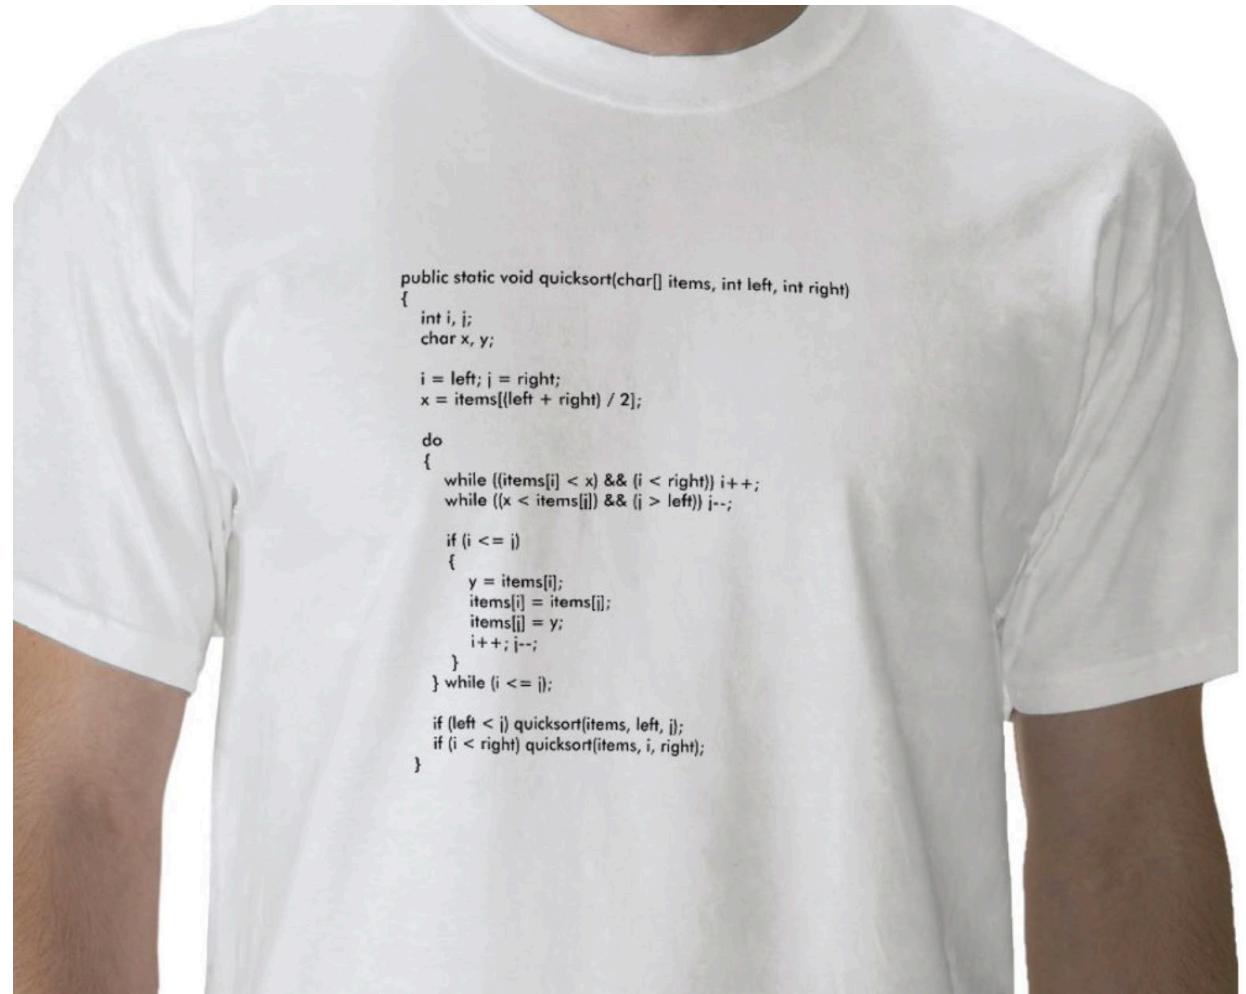
do
{
    while ((items[i] < x) && (i < right)) i++;
    while ((x < items[j]) && (j > left)) j--;

    if (i <= j)
    {
        y = items[i];
        items[i] = items[j];
        items[j] = y;
        i++; j--;
    }
} while (i <= j);
```

# QUESTION

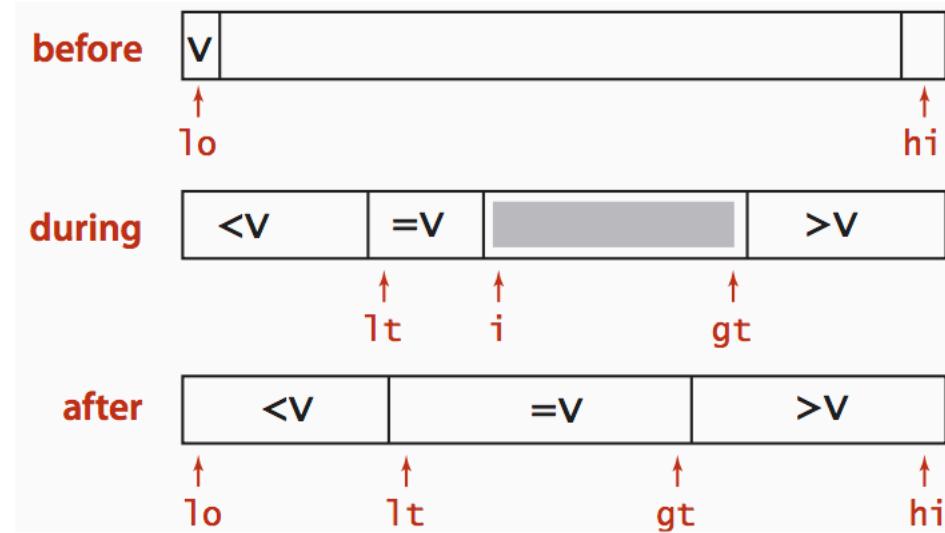
Which algorithm is this snippet of code from?

- A. Mergesort
- B. Quicksort
- C. Heapsort
- D. Selectionsort
- E. Dijkstra's SSSP



# 3 WAY PARTITION + QUICKSORT

```
quicksort(int[] a, int lo, int hi) {  
    if (hi <= lo) return;  
    int lt = lo, gt = hi;  
    int v = a[lo];  
    int i = lo + 1;  
    while (i <= gt) {  
        if (a[i] < v) swap(a, lt++, i++);  
        else if (a[i] > v) swap(a, i, gt--);  
        else i++;  
    }  
    quicksort(a, lo, lt-1);  
    quicksort(a, gt+1, hi);  
}
```



# QUESTION

True or False: Quicksort will preserve the ordering of elements that have the same key.

- A. True
- B. False

# QUESTION

True or False: Quicksort will preserve the ordering of elements that have the same key.

- A. True
- B. False

# QUICKSORT: STABLE?

A sorting algorithm is stable if two objects with equal keys appear in the same order in the output as in input.

- No changing places if you have the same key!

Why? Sort by lastname only:

soh, aaron

soh, harold

chan, hazel

input

chan, hazel

soh, aaron

soh, harold

stable

chan, hazel

soh, harold

soh, aaron

unstable

sort

Is quicksort stable?

- A. Yessssss.
- B. **Noooooo**
- C. Who knows?
- D. Stable as an ikea table.

Consider [a,b,c] where c < a = b

# QUESTION

I have an unsorted array  $A$  of size  $n$  that has *distinct* elements from 1 to  $n$ . e.g.,  $A = [2, 3, 1]$ , or  $A = [4, 2, 1, 3, 5]$ .

The time complexity required to output  $A$  in sorted order is:

- A.  $O(1)$
- B.  $O(n)$
- C.  $O(n \log n)$
- D.  $O(n^2)$
- E. Larger than the above

# QUESTION

I have an unsorted array  $A$  of size  $n$  that has *distinct* elements from 1 to  $n$ . e.g.,  $A = [2, 3, 1]$ , or  $A = [4, 2, 1, 3, 5]$ .

The time complexity required to print  $A$  in sorted order is:

- A.  $O(1)$
- B.  $O(n)$
- C.  $O(n \log n)$
- D.  $O(n^2)$
- E. Larger than the above

# QUESTION

The depth of any 2 leaves in a heap differs at most by 1.

True or False?

- A. True
- B. False

# QUESTION

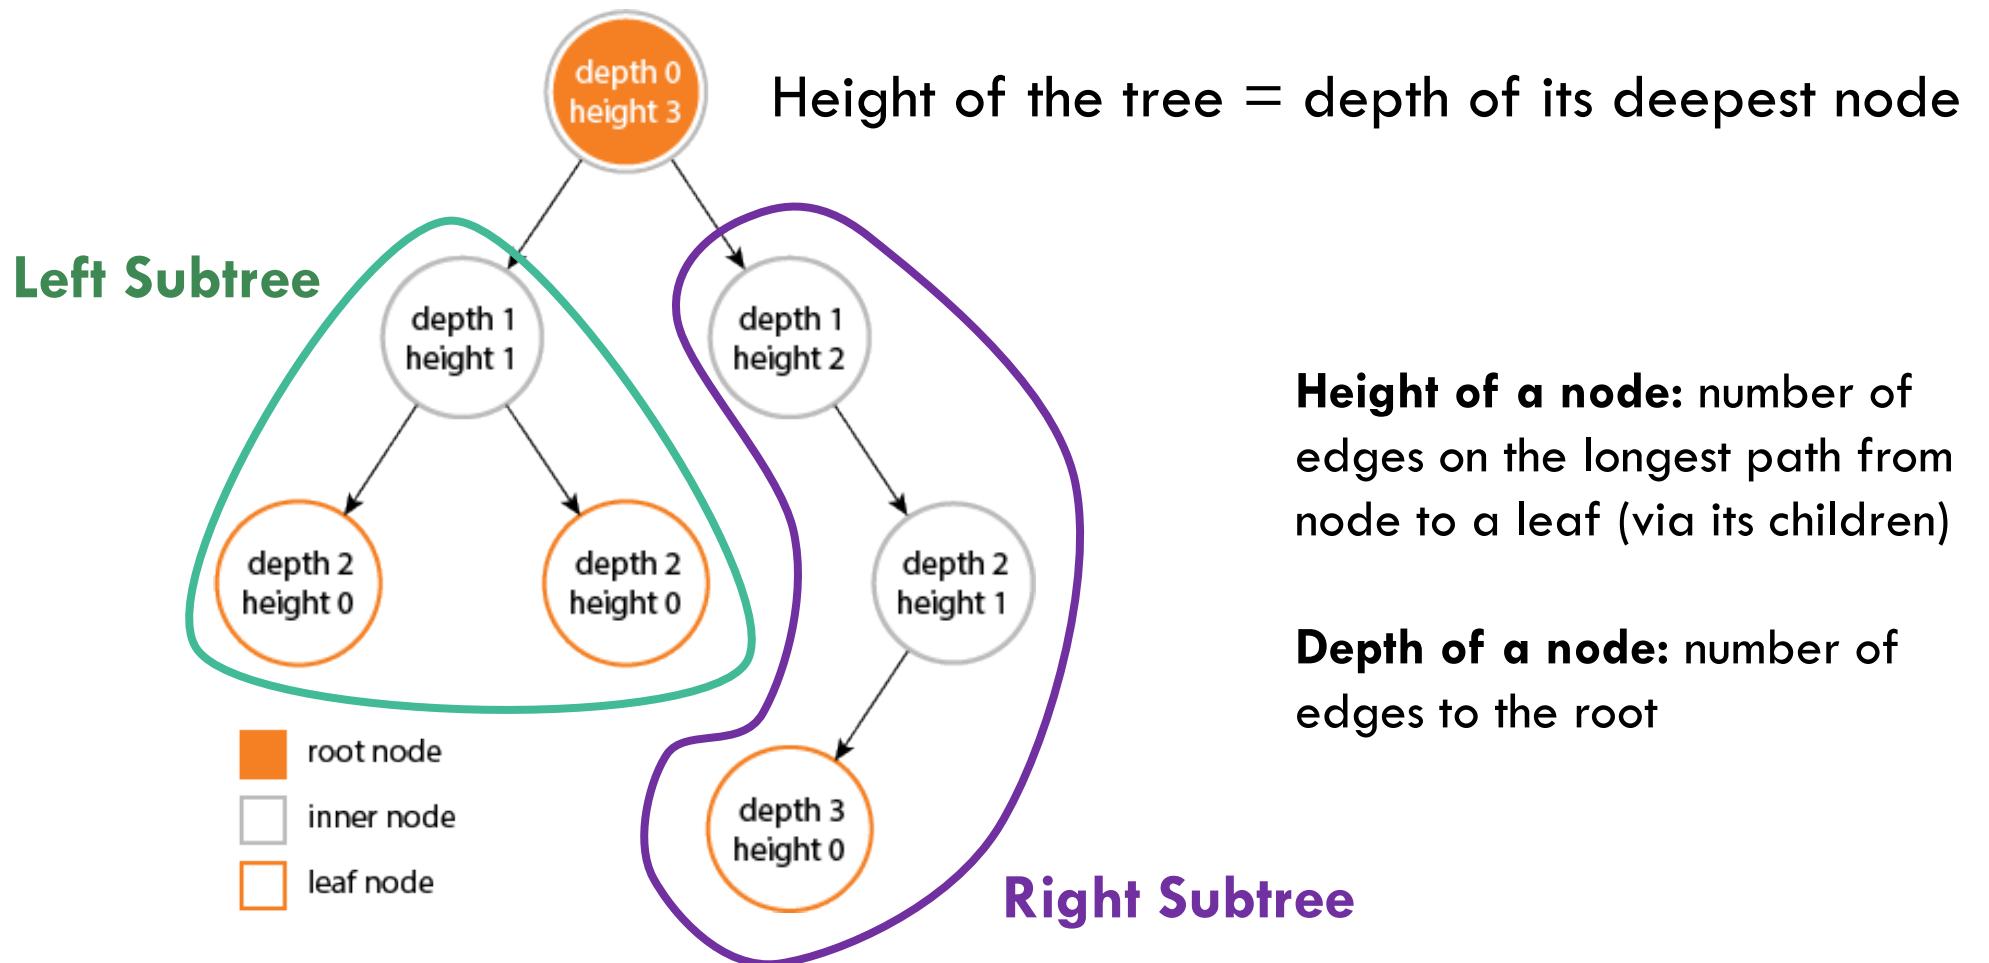
The depth of any 2 leaves in a heap differs at most by 1.

True or False?

A. True

B. False

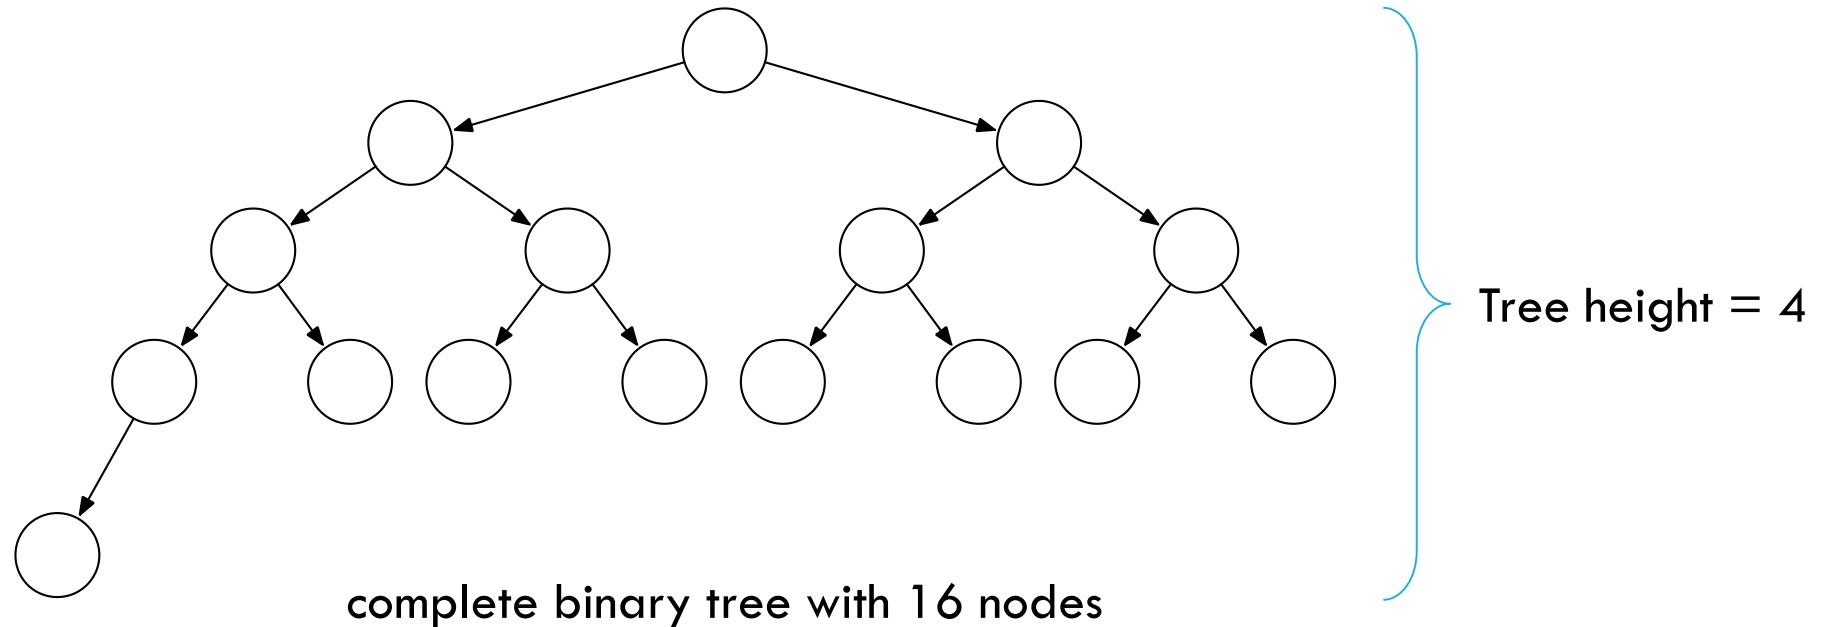
# A BINARY TREE



# COMPLETE BINARY TREE

Binary Tree: nodes with links to left and right binary trees (or empty)

Complete tree: Perfectly balanced, except for bottom level.



# QUESTION

Which data structure is space and time efficient for storing a heap?

- A. Linked List
- B. Doubly-Linked List
- C. Array
- D. Tree
- E. None of the above

# QUESTION

Which data structure is space and time efficient for storing a heap?

- A. Linked List
- B. Doubly-Linked List
- C. **Array**
- D. Tree
- E. None of the above

# TREE REPRESENTATION

**Class Node**

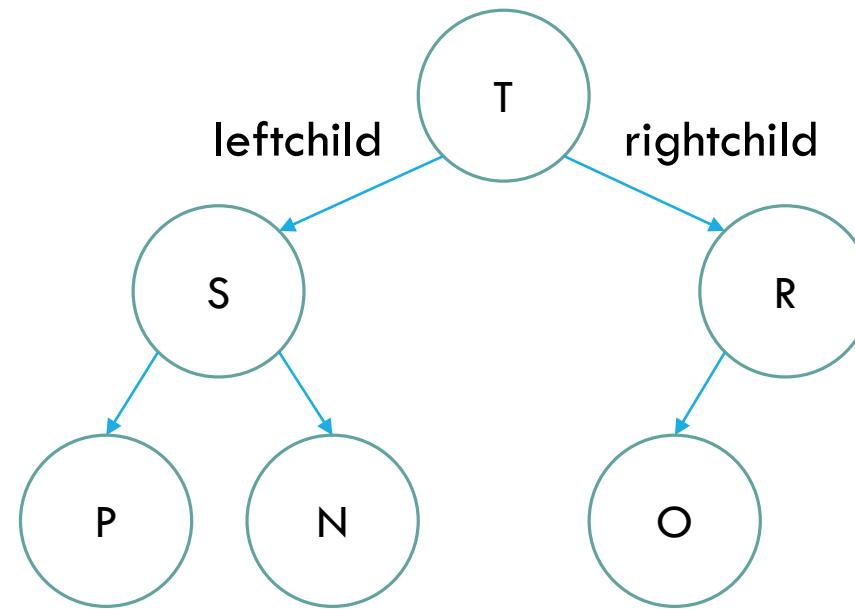
Object key

Object data

Node leftchild

Node rightchild

Node parent





# ARRAY REPRESENTATION

Indices start at 1

Take nodes in level order

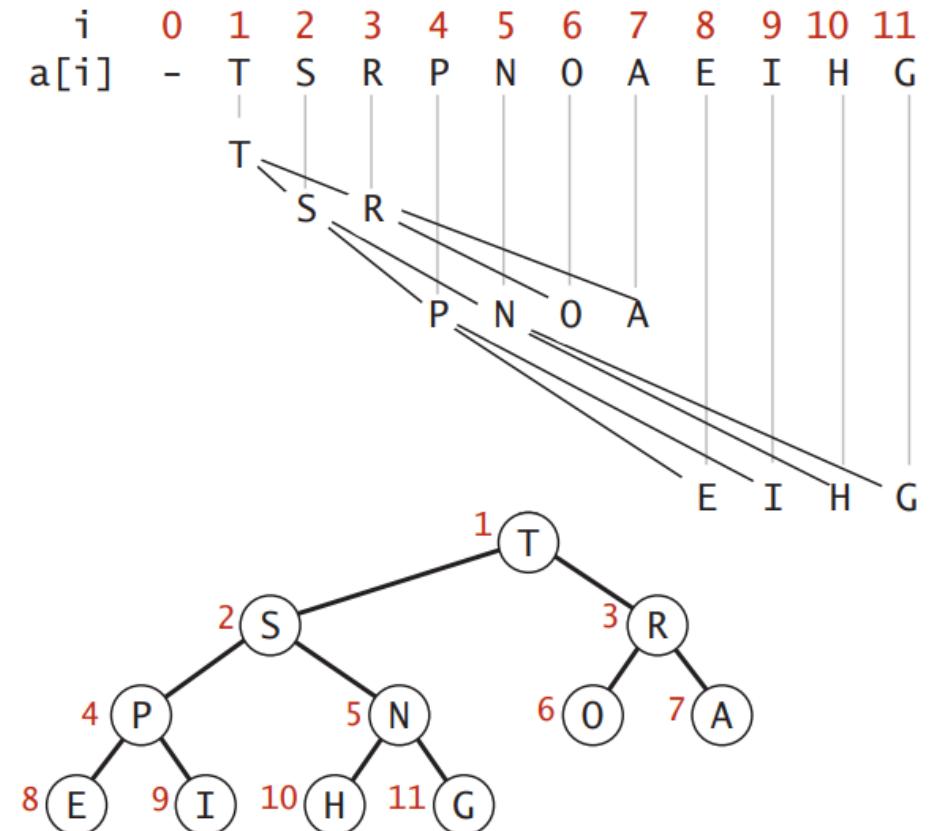
No links needed

Parent of node  $k$  is at  $k/2$

Children of node  $k$  are at:

Left Child:  $2k$

Right Child:  $2k + 1$



Heap representations

# QUESTION

If a key in a max-heap is decreased, what operation is needed to restore the max-heap property?

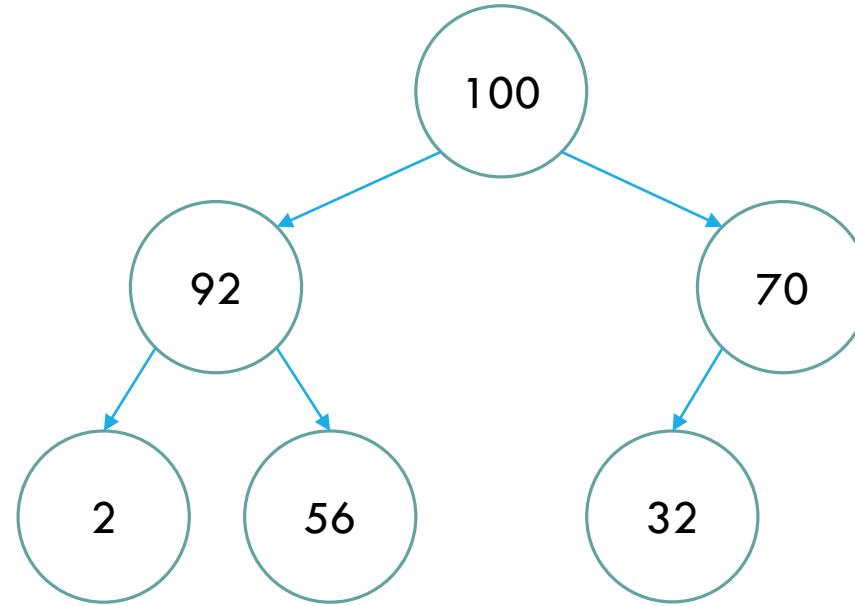
- A. Swim
- B. Sink
- C. Merge
- D. None of the above

# QUESTION

If a key in a max-heap is decreased, what operation is needed to restore the max-heap property?

- A. Swim
- B. **Sink**
- C. Merge
- D. None of the above

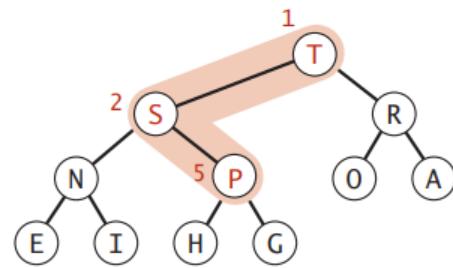
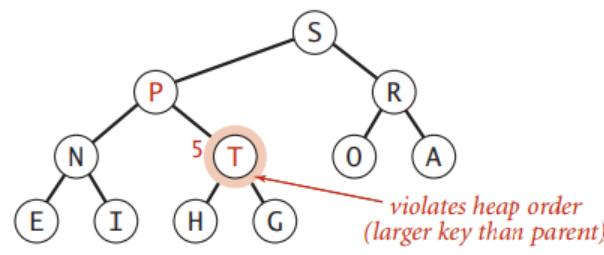
# BINARY HEAP



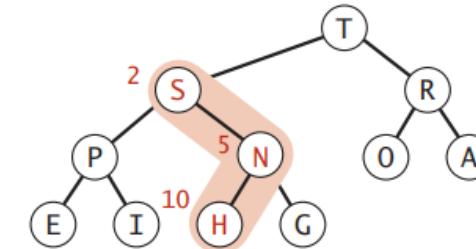
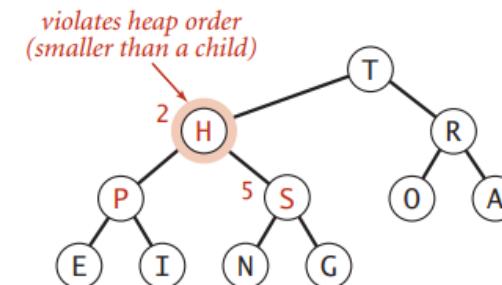
## **Max heap property:**

- the key of each node is larger than or equal to the keys in its children

# TWO IMPORTANT OPERATIONS: SWIM & SINK



**Swim**



**Sink**

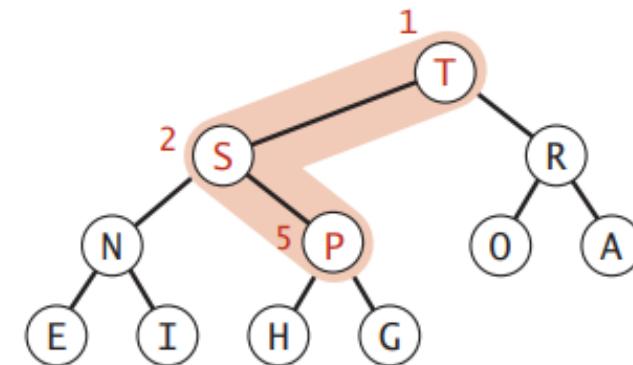
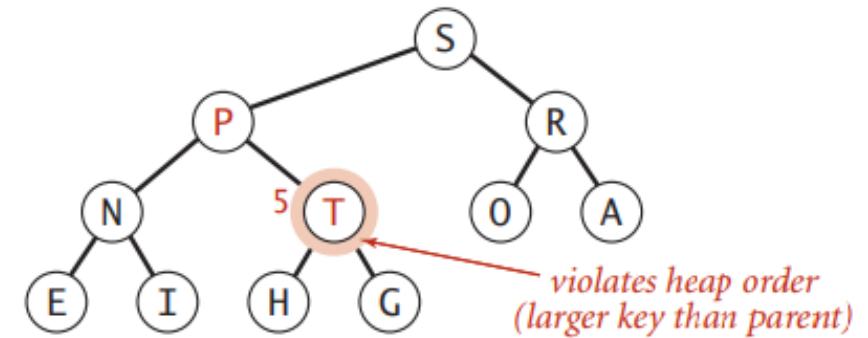
We use these operations to correct binary trees to satisfy the max heap property

# SWIM or (SHIFTUP, BUBBLEUP, INCREASEKEY)

**Problem:** A key becomes larger than its parent's key.

**Solution:**

- Swap child with parent
- Repeat until heap order restored



# SWIM or (SHIFTUP, BUBBLEUP, INCREASEKEY)

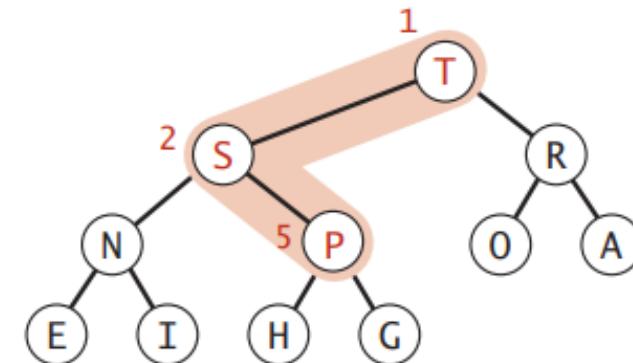
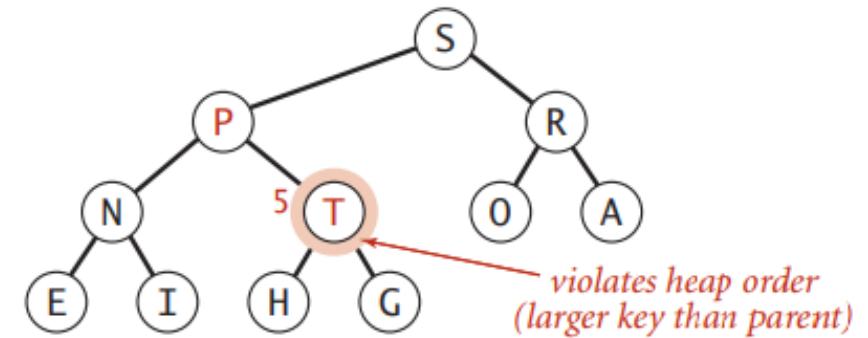
**Problem:** A key becomes larger than its parent's key.

**Solution:**

- Swap child with parent
- Repeat until heap order restored

Pseudocode for array implementation

```
function swim(A, k)
    while (k > 1) and A[k/2].key < A[k].key
        swap(A[k], A[k/2])
        k = k/2
```

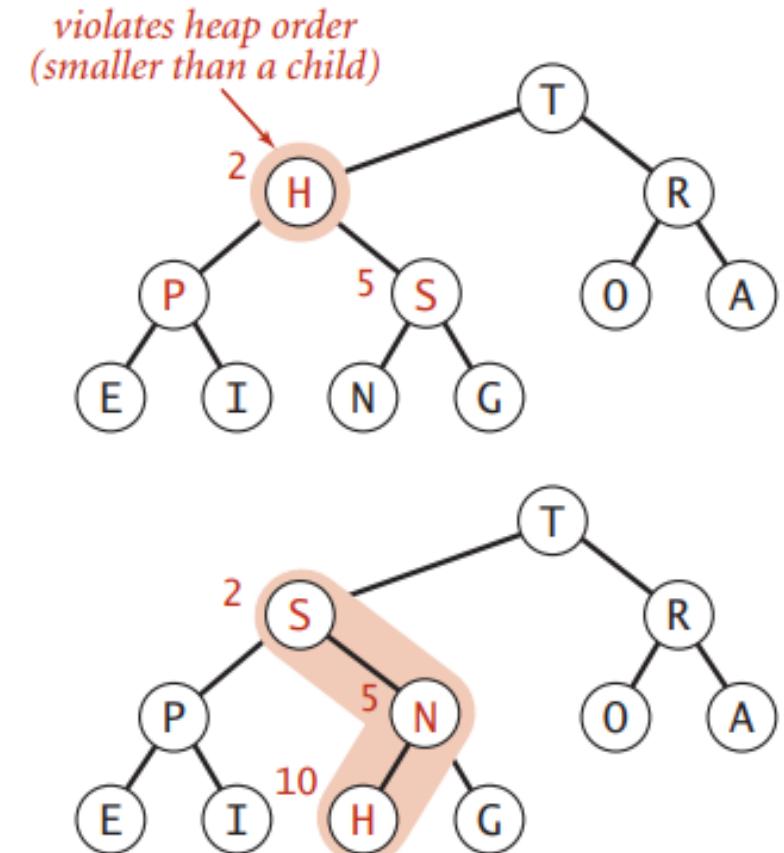




# SINK or (SHIFTDOWN, BUBBLEDOWN, HEAPIFY)

**Problem:** A key becomes smaller than one (or both) of its children's keys.

**Solution:**



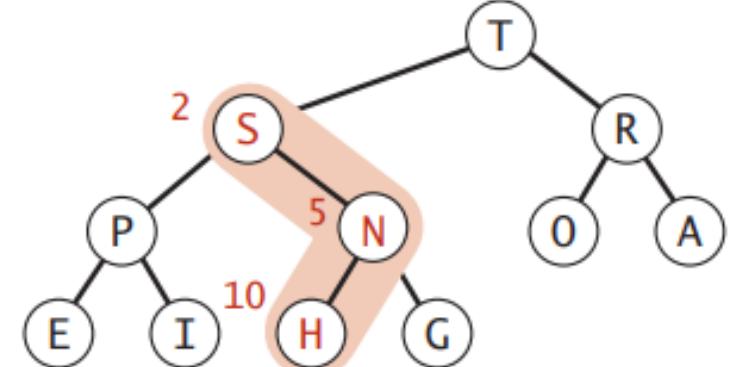
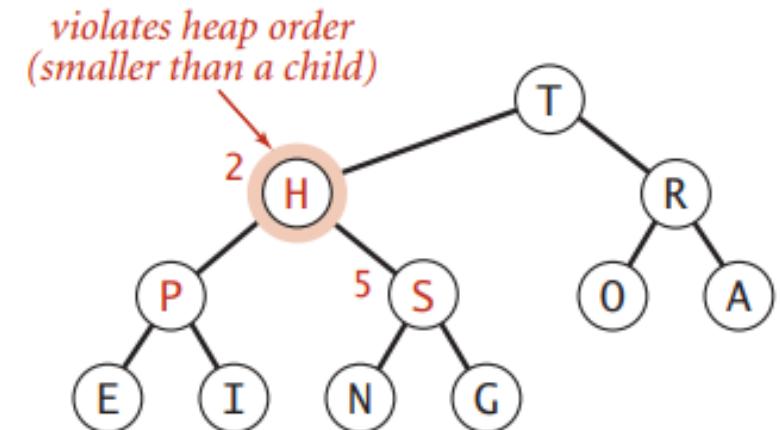


# SINK or (SHIFTDOWN, BUBBLEDOWN, HEAPIFY)

**Problem:** A key becomes smaller than one (or both) of its children's keys.

**Solution:**

- Swap larger child with parent
- Repeat until heap order restored





# SINK or (SHIFTDOWN, BUBBLEDOWN, HEAPIFY)

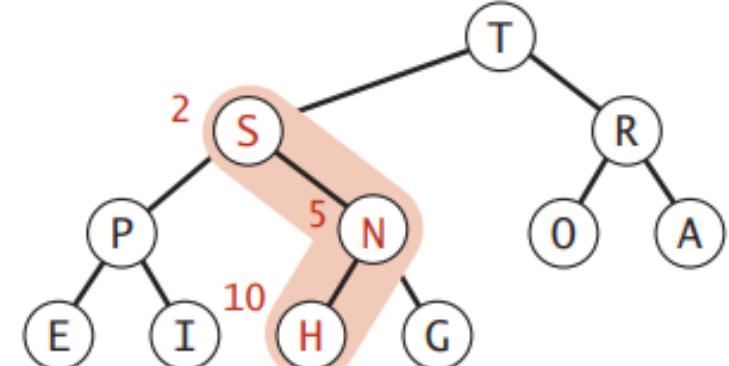
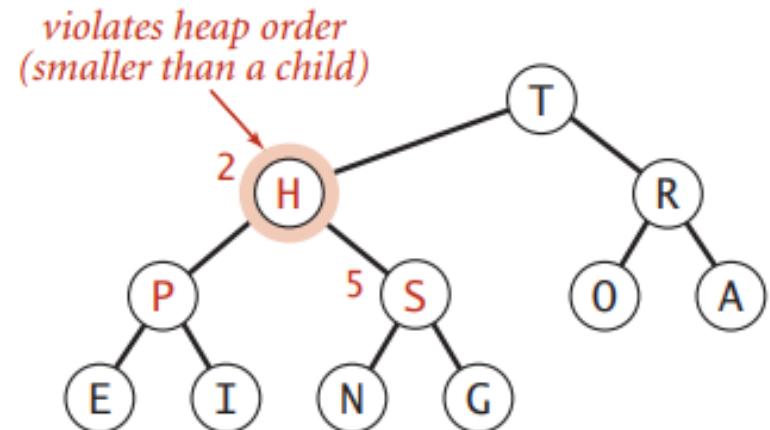
**Problem:** A key becomes smaller than one (or both) of its children's keys.

**Solution:**

- Swap larger child with parent
- Repeat until heap order restored

Pseudocode for array implementation

```
function sink(A, k)
    while (2k <= n)
        j = 2k
        if (j < n and (A[j].key < A[j+1].key)) j++
        if not (A[k].key < A[j].key) break
        swap(A[k], A[j])
        k = j
```



# QUESTION

Who invented the algorithm for creating a heap in linear time?

- A. Donald Knuth
- B. Robert Floyd
- C. Edsger Dijkstra
- D. Dlorah Hos
- E. None of the above

# QUESTION

Who invented the algorithm for creating a heap in linear time?

- A. Donald Knuth
- B. **Robert Floyd**
- C. Edsger Dijkstra
- D. Dlorah Hos
- E. None of the above

# CLEVER CREATION IN $O(n)$ TIME



Invented by Robert Floyd in 1964

- invented invariants (among other things)
- we'll hear about him again in when we meet graphs!

The idea:

- View the input array as a binary tree
- “Bottom up” fixing of the tree to satisfy MaxHeap property

# **CODE FOR CREATION IN $O(n)$ TIME**

```
function buildHeap(A)
  for i from A.length/2 to 1
    sink(i)
```



# QUESTION

I need to sort an array of  $n$  elements, but with two constraints:

- only  $O(1)$  additional space, and
- $O(n \log n)$  worst case time guarantee

What algorithm should I use?

- A. Insertion sort
- B. Mergesort
- C. Quicksort
- D. Heapsort
- E. Bubble sort

# QUESTION

I need to sort an array of  $n$  elements, but with two constraints:

- only  $O(1)$  additional space, and
- $O(n \log n)$  worst case time guarantee

What algorithm should I use?

- A. Insertion sort
- B. Mergesort
- C. Quicksort
- D. **Heapsort**
- E. Bubble sort

# HEAPSORT CODE

```
function heapsort(A)
    n = A.length
    // create the binary heap
    for i = n/2 to 1
        sink(A, k, n)
    // swap and sink
    while (n > 1)
        swap(A, 1, n);
        sink(A, 1, --n);
```

Creating a binary heap in place takes  $O(n)$

Swapping and Sinking takes \_\_\_\_\_

Note: uses modified sink and swim with array length

# HEAPSORT CODE

```
function heapsort(A)
    n = A.length
    // create the binary heap
    for i = n/2 to 1
        sink(A, k, n)
    // swap and sink
    while (n > 1)
        swap(A, 1, n);
        sink(A, 1, --n);
```

Creating a binary heap in place takes  $O(n)$

Swapping and Sinking takes  $O(n \log n)$

Note: uses modified sink and swim with array length

# QUESTION

True/False:

In a Binary Search Tree, a node's key will always be smaller than its left child and larger than its right child.

- A. True
- B. False

# QUESTION

True/False:

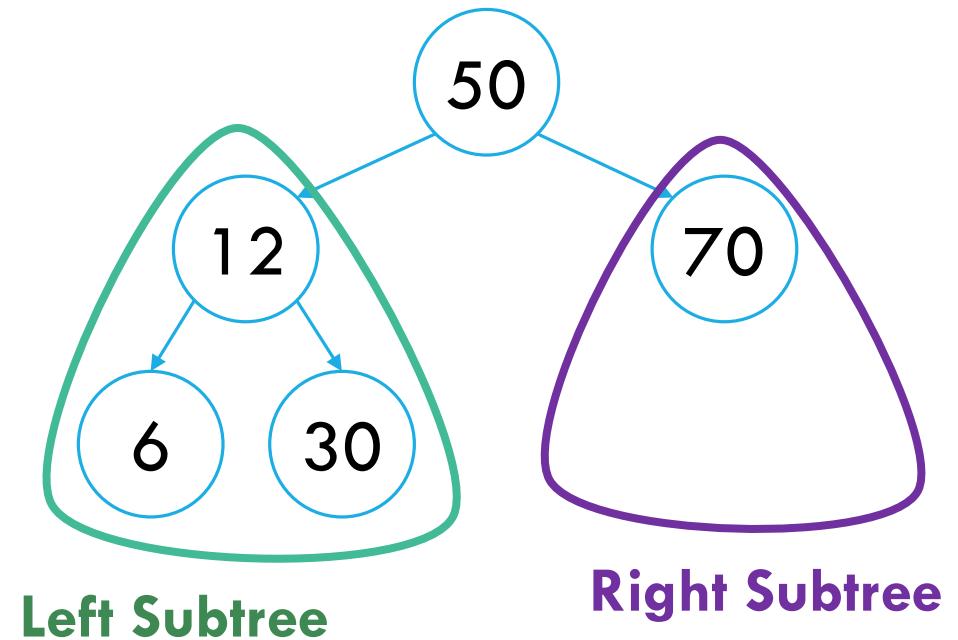
In a Binary Search Tree, a node's key will always be smaller than its left child and larger than its right child.

- A. True
- B. False

# BINARY SEARCH TREES

## Binary Search Tree Properties:

- A node's left subtree only contains nodes with keys strictly less than the node's key
- A node's right subtree only contains nodes with keys strictly larger than the node's key
- The left and right subtrees are binary trees.
- All keys belong to a total order\*.



\*  $\forall u, v \in K$  when  $u \neq v$  either  $u < v$  or  $v > u$  (no two different keys can be considered equal)

# QUESTION

Naruto's uncle, Uzumaki, was killed in an unfortunate skiing accident. Everyone in Naruto's family is stored in a Balanced BST by height, and the next tallest family member inherits Uzumaki's estate. How can we efficiently find this member?

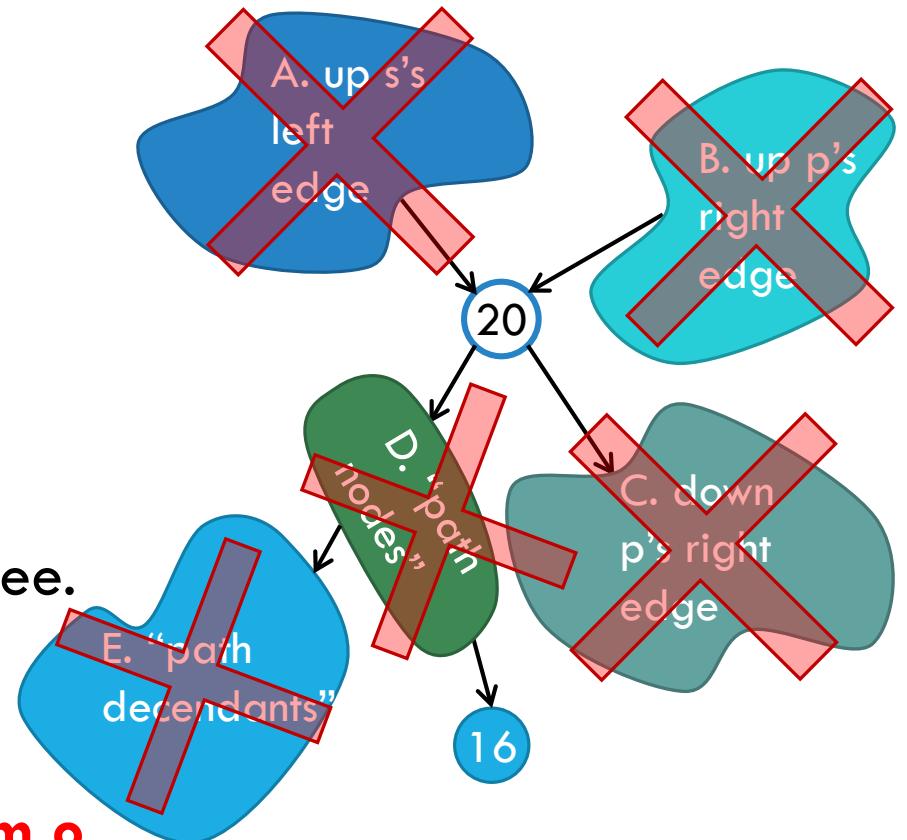
- A. Binary Search
- B. Successor
- C. Predecessor
- D. Array Halving

# FINDING THE SUCCESSOR NODE

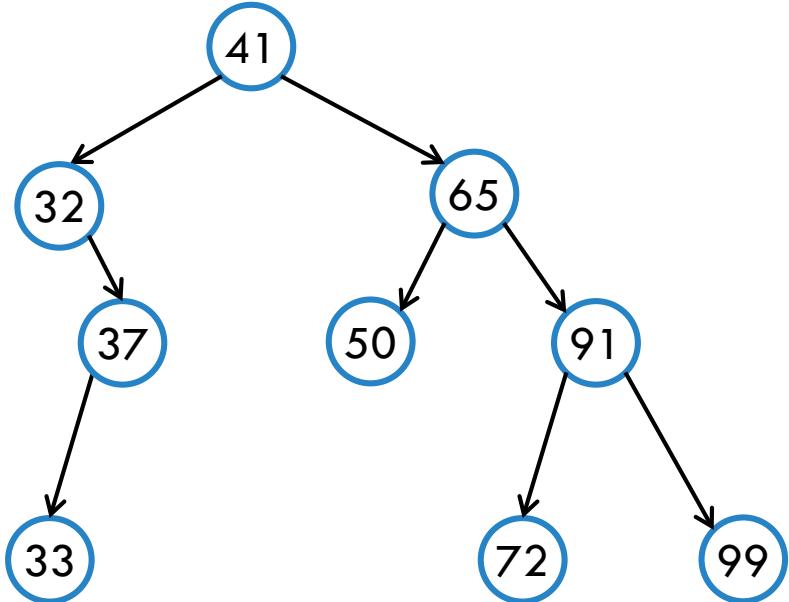
Given a node  $o$ , there are 2 cases:

- it has a right subtree
  - return `searchMin(o.m_rightTree)`
- it doesn't have a right subtree
  - **Claim:** successor will be higher up in the tree.
  - 2 cases:
    - $o$  is the left child: successor is  $o$ 's parent
    - $o$  is the right child: **first ancestor up from  $o$  that has key > k. ("first right ancestor")**

Examine each  
in turn



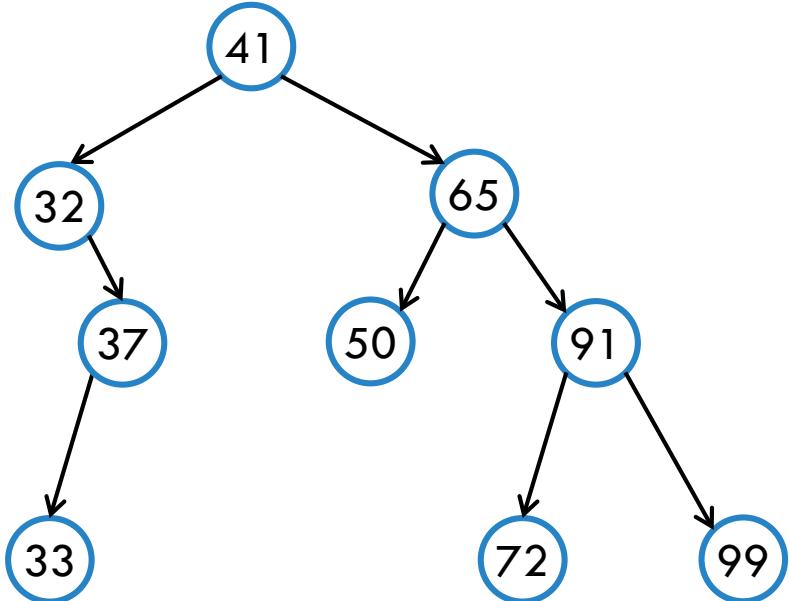
# QUESTION



To make this tree height balanced, what should I do and to which nodes?

- A. Left rotate 32
- B. Left rotate 37
- C. Right rotate 37
- D. Right rotate 37, left rotate 32
- E. Left rotate 32, right rotate 33

# QUESTION

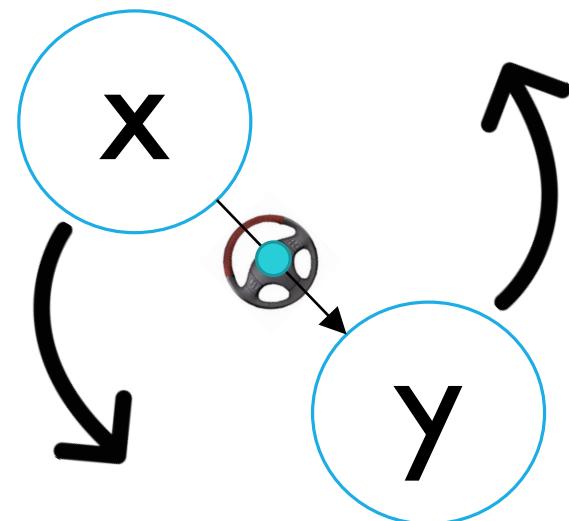


To make this tree height balanced, what should I do and to which nodes?

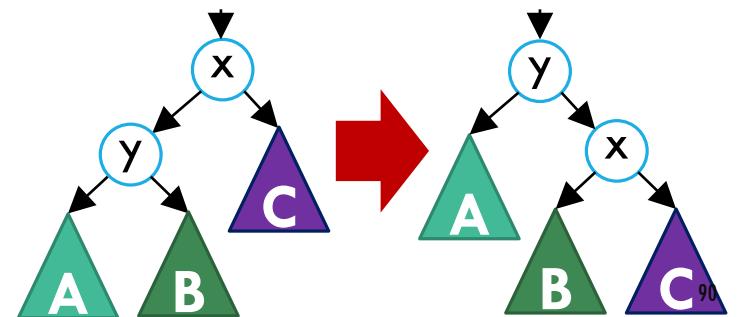
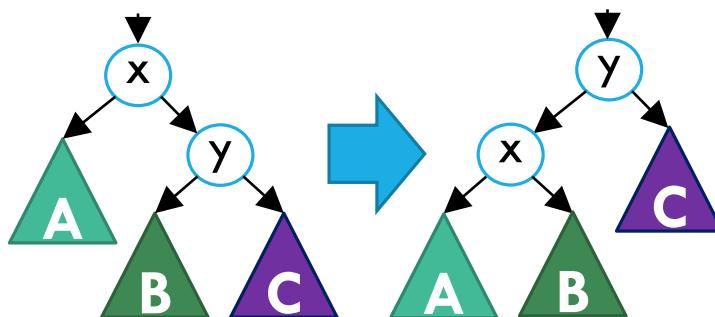
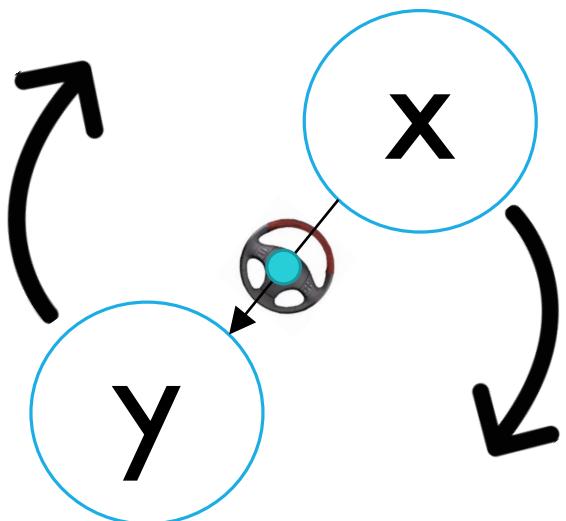
- A. Left rotate 32
- B. Left rotate 37
- C. Right rotate 37
- D. **Right rotate 37, left rotate 32**
- E. Left rotate 32, right rotate 33

# HOW I REMEMBER IT

left / anti-clockwise



right / clockwise



# ALGORITHM TO FIX IMBALANCES

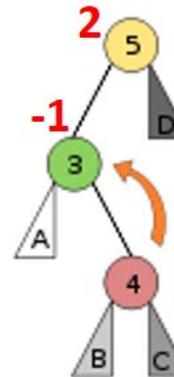
For both insertion and deletions:

- walk up the tree (to the root) from the inserted/deleted node & update balance factors.
- if we find a height-balance violation, fix it via a rotation:

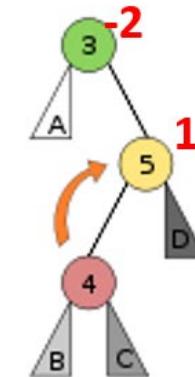
```

if tree is right heavy
    if tree's right subtree is left heavy
        right rotate, left rotate
    else
        left rotate
else if tree is left heavy
    if tree's left subtree is right heavy
        left rotate, right rotate
    else
        right rotate
  
```

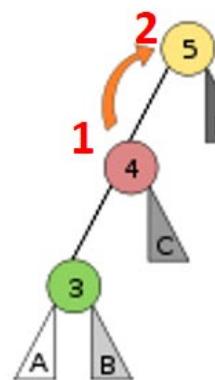
Left Right Case



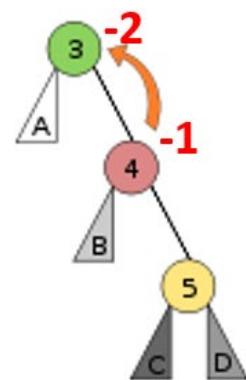
Right Left Case



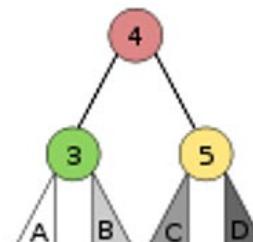
Left Left Case



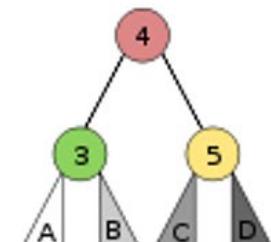
Right Right Case



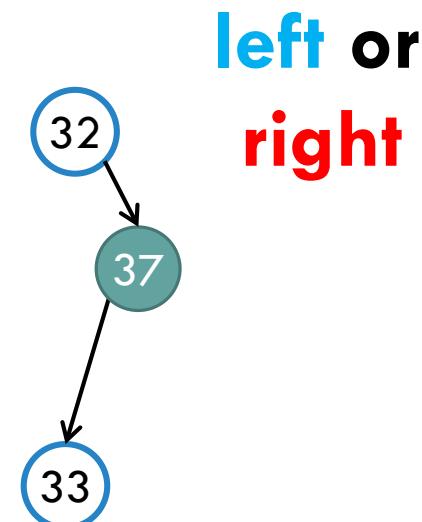
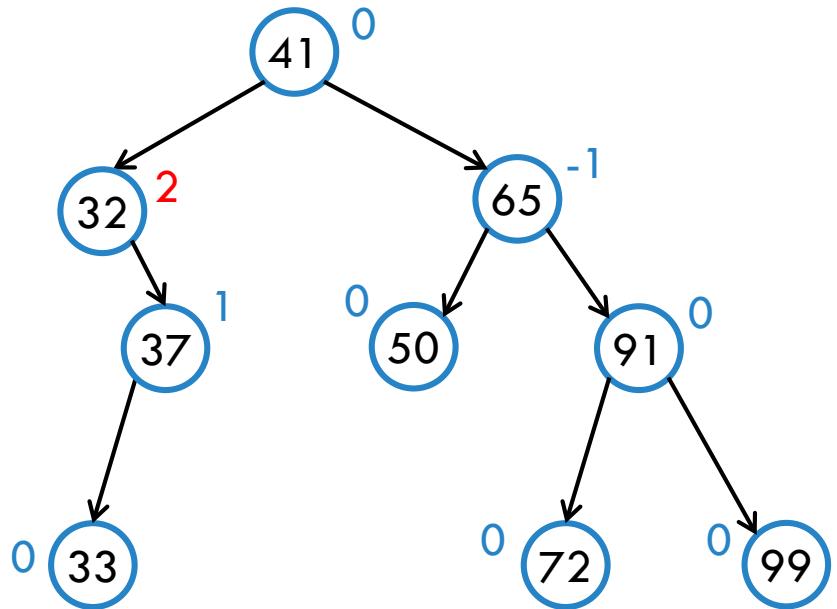
Balanced



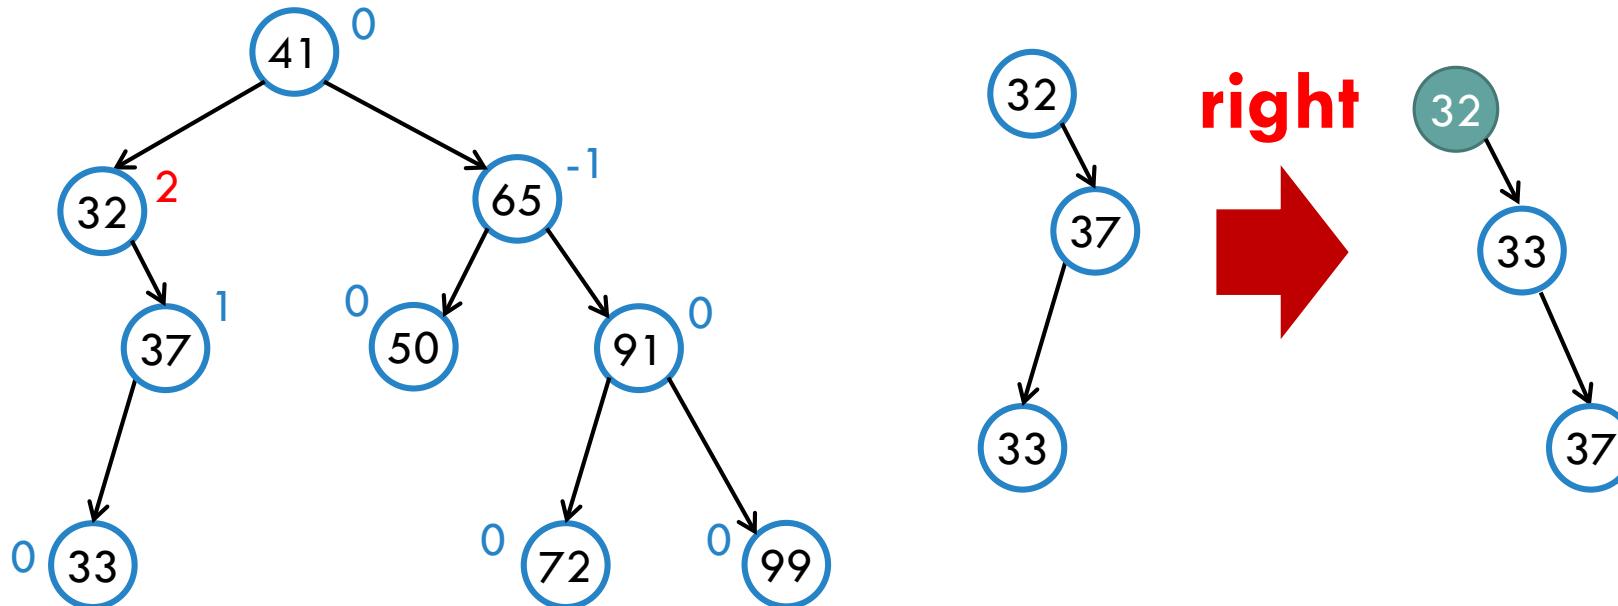
Balanced



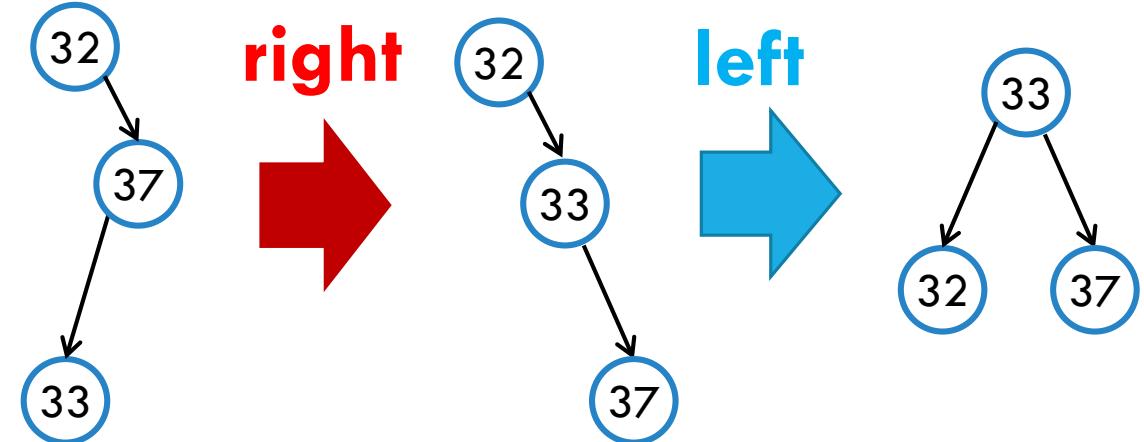
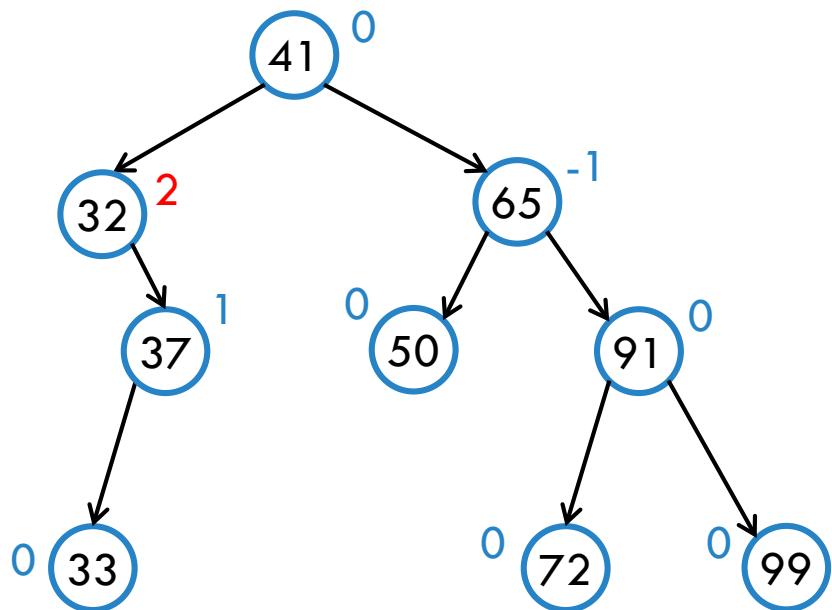
# LET'S TRY THE DOUBLE ROTATION!



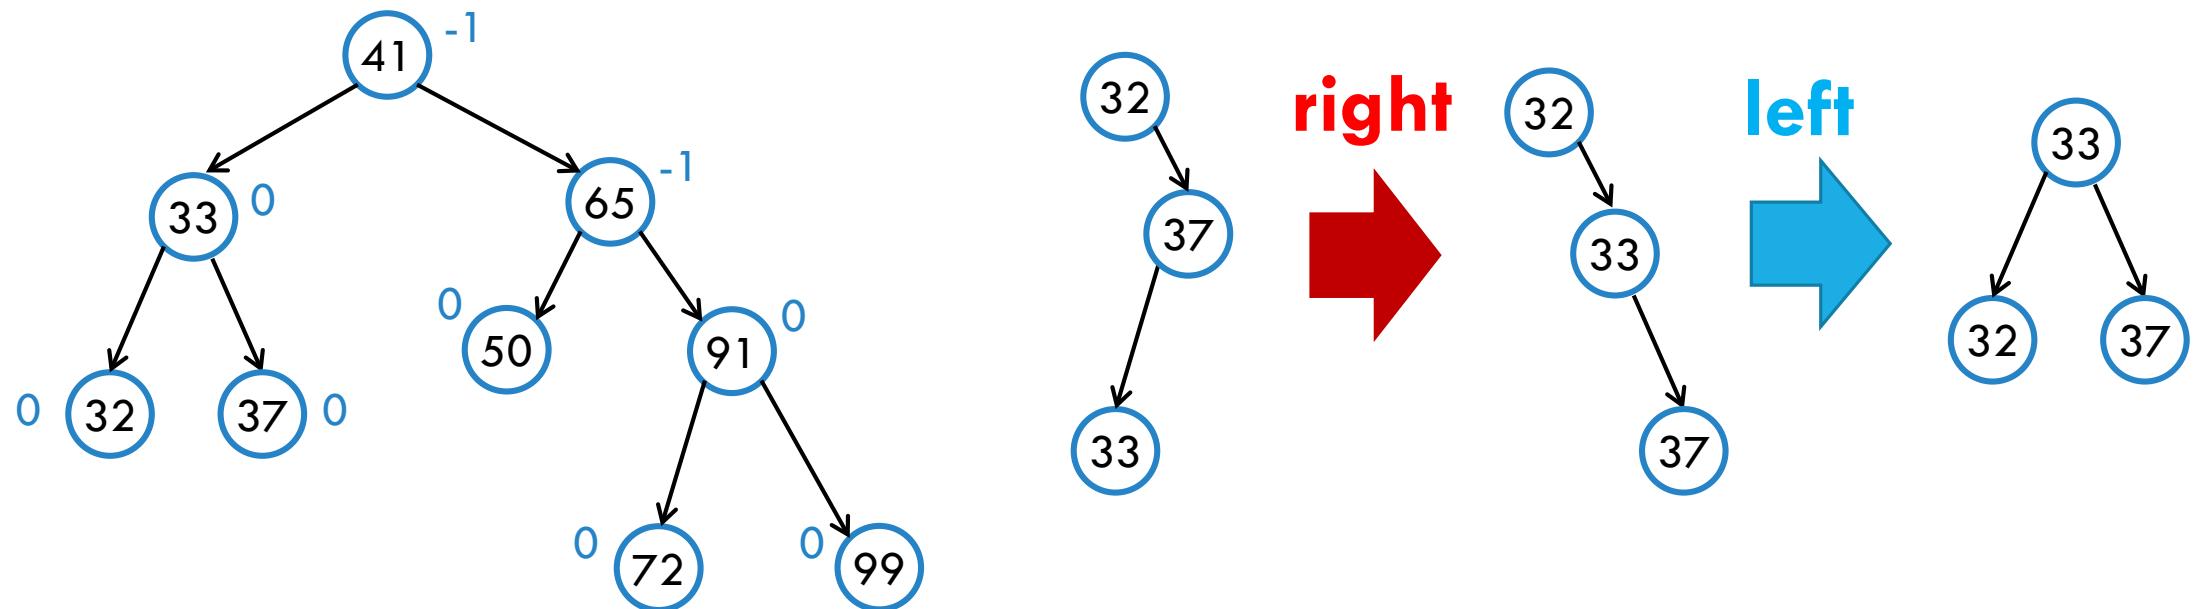
# LET'S TRY THE DOUBLE ROTATION!



# LET'S TRY THE DOUBLE ROTATION!



# LET'S TRY THE DOUBLE ROTATION!



**It worked!**

# QUESTION

In a hash table, which of the following operations are fast (e.g., constant or log time)?

- A. Floor
- B. Insert
- C. Successor
- D. Max

# QUESTION

In a hash table, which of the following operations are fast (e.g., constant or log time)?

- A. Floor
- B. Insert
- C. Successor
- D. Max

# WHAT DO WE GIVE UP?



**no free  
lunch!**

Data Structure	Avg. Insert Time	Avg. Search Time	Avg. Max/Min	Avg. Floor/Ceiling
Unordered Array / Linked List	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Ordered Array / Linked List	$O(n)$	$O(\log n)$	$O(1)$	$O(\log n)$
Balanced Binary Search Tree (AVL)	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$
Hash Table	$O(1)$	$O(1)$	$O(n)$	$O(n)$

# SIMPLE UNIFORM HASHING ASSUMPTION

An optimistic assumption:

Every key is **equally likely** to map to every bucket

Keys are mapped independently.

Intuition:

- Each key is put in a random bucket.
- As long as enough buckets, not too many keys in any one bucket.



OPTIMISM

It doesn't matter if the glass is half full or half empty if you have a lot of glasses

# OPEN ADDRESSING: LINEAR PROBING

**Idea:** On collision, probe until you find an empty slot.

**Question:** How to probe?

**Linear Probing:** keep checking next bucket until you find an empty slot.

$$\text{index } i = (h(k) + \text{step} \times 1) \bmod m$$

$$h(k_5) = 2$$

insert( $k_5$ , E)

0	null
1	null
2	( $k_1$ , A)
3	( $k_3$ , C)
4	( $k_4$ , D)
5	( $k_5$ , E)
6	null
7	null
8	( $k_2$ , B)
9	null

Collision!

Success!

100

# A PROBLEM: PRIMARY CLUSTERS

cluster = collection of consecutive occupied slots

In a hash table of size 10, consider 2 clusters:

- A: size 5
- B: size 2

Probability that a new inserted key k has a bucket in

- cluster A?  $5/10$
- cluster B?  $2/10$

What happens after k is added?

Cluster grows by 1 element

0	null
1	
2	
3	A
4	
5	
6	
7	null
8	B
9	

# WHY DO CLUSTERS MATTER?

With table size  $m$  and  $n = \alpha m$  keys, the average number of linear probes is:

$$\sim \frac{1}{2} \left( 1 + \frac{1}{1-\alpha} \right) \text{ for search hits}$$

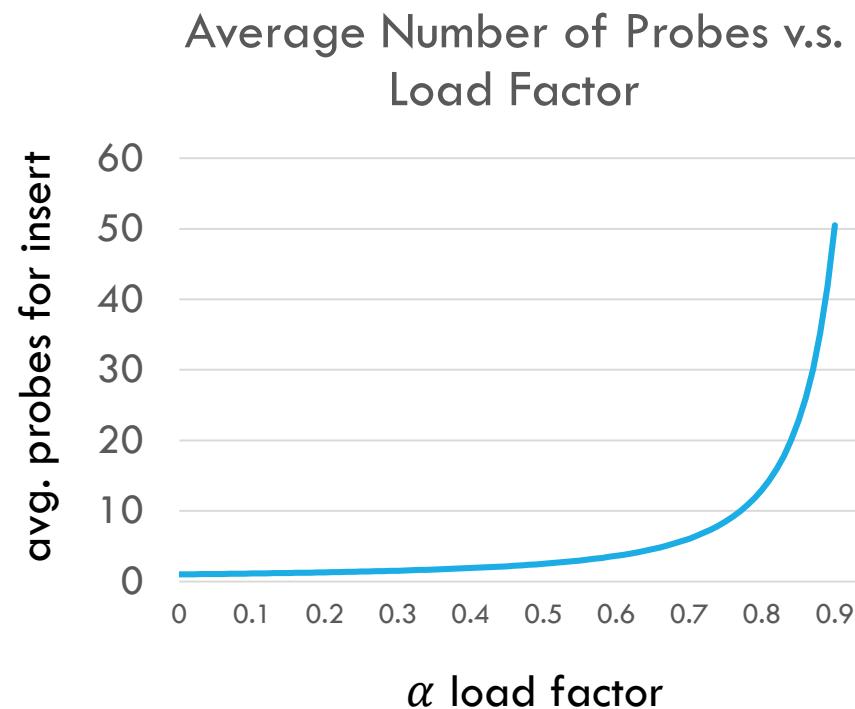
$$\sim \frac{1}{2} \left( 1 + \frac{1}{(1-\alpha)^2} \right) \text{ for search misses}$$

Recall  $\alpha \leq 1$  is the “load factor”

(Proof beyond scope of CS2040S)

$\lambda$	1/4	1/2	2/3	3/4	9/10
<b>successful</b>	1.2	1.5	2.0	3.0	5.5
<b>unsuccessful</b>	1.4	2.5	5.0	8.5	50.5

<http://www.cs.cmu.edu/afs/cs/academic/class/15210-f13/www/lectures/lecture24.pdf>



# OPEN ADDRESSING: QUADRATIC PROBING

**Linear probing:** index  $i = (h(k) + \text{step} \times 1) \bmod m$



**Quadratic probing:** index  $i = (h(k) + \text{step}^2) \bmod m$

**Example:**

- $h(k) = 3, m = 7$
- Step 0:  $i = h(k) = 3$
- Step 1:  $i = (h(k) + 1) \bmod 7 = 4$
- Step 2:  $i = (h(k) + 4) \bmod 7 = 0$

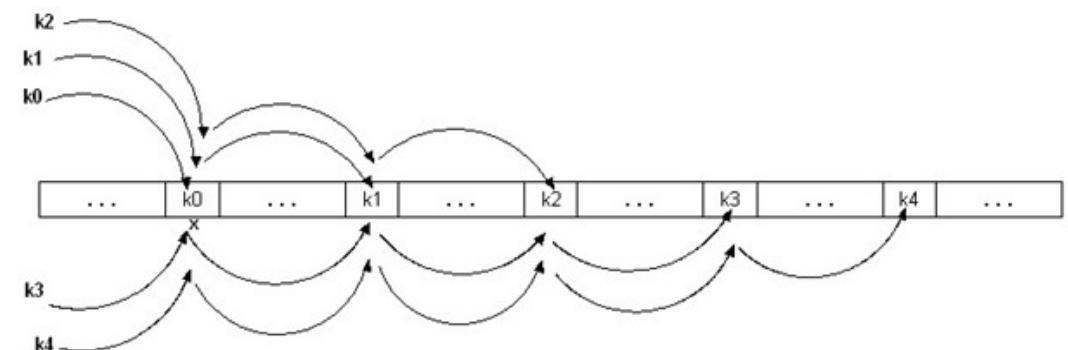
Is this a good probing method?

# ONE PROBLEM: SECONDARY CLUSTERING

Milder form of the clustering problem.

Because: if two keys have the same probe position, their probe sequences are the same.

Clustering around different points  
(rather than the primary probe point)



**How many probe sequences can there be?  $m$**

# DOUBLE HASHING

Use a second hashing:

$$\text{index } i = (h_1(k) + \text{step} \times h_2(k)) \bmod m$$

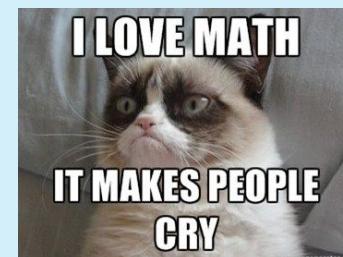
Avoids secondary clustering by providing more unique probing sequences.

## Intuition:

- $h_1(k)$  provides good “random” base address
- $h_2(k)$  provides good “random” sequence

Up to how many unique indexing sequences does double hashing provide?

- A.  $m$
- B.  $2m$
- C.  $m^2$
- D.  $2^m$
- E.



# QUESTION

I have a hash table of size  $2^r$  where  $r = 8$ . Should I use the division or multiplication hashing method?

- A. Division
- B. Multiplication
- C. Either would work just fine.

# QUESTION

I have a hash table of size  $2^r$  where  $r = 8$ . Should I use the division or multiplication hashing method?

- A. Division
- B. **Multiplication**
- C. Either would work just fine.

# DIVISION METHOD

$$h(k) = k \bmod m$$

- $m = 7, h(17) = 3$
- $m = 20, h(100) = 0$
- $m = 20, h(97) = 17$
- $m = 13, h(102) = 12$

Two keys  $k_1$  and  $k_2$  **collide** when:

$$k_1 = k_2 \pmod{m}$$

Want:

- A. Consistent
- B. Fast to compute,  $O(1)$
- C. as uniform as possible  $\in [0..m - 1]$

**How to choose  $m$ ?**

# CHOICE OF $m = 2^x$

Does it satisfy the desirability C?

**Problem:** Regularity of keys

- Input keys are often not uniformly distributed

If all input keys are even:

- $h(k) = k \bmod m$  must be even.

**No longer uniform:** will waste  $\frac{1}{2}$  the space and increase collisions!

Want:

- A. Consistent
- B. Fast to compute,  $O(1)$
- C. as uniform as possible  $\in [0..m - 1]$

**Fast to compute, but does not use table well.  
Also, will only use lowest order bits of k (Why?)**

# DIVISION METHOD

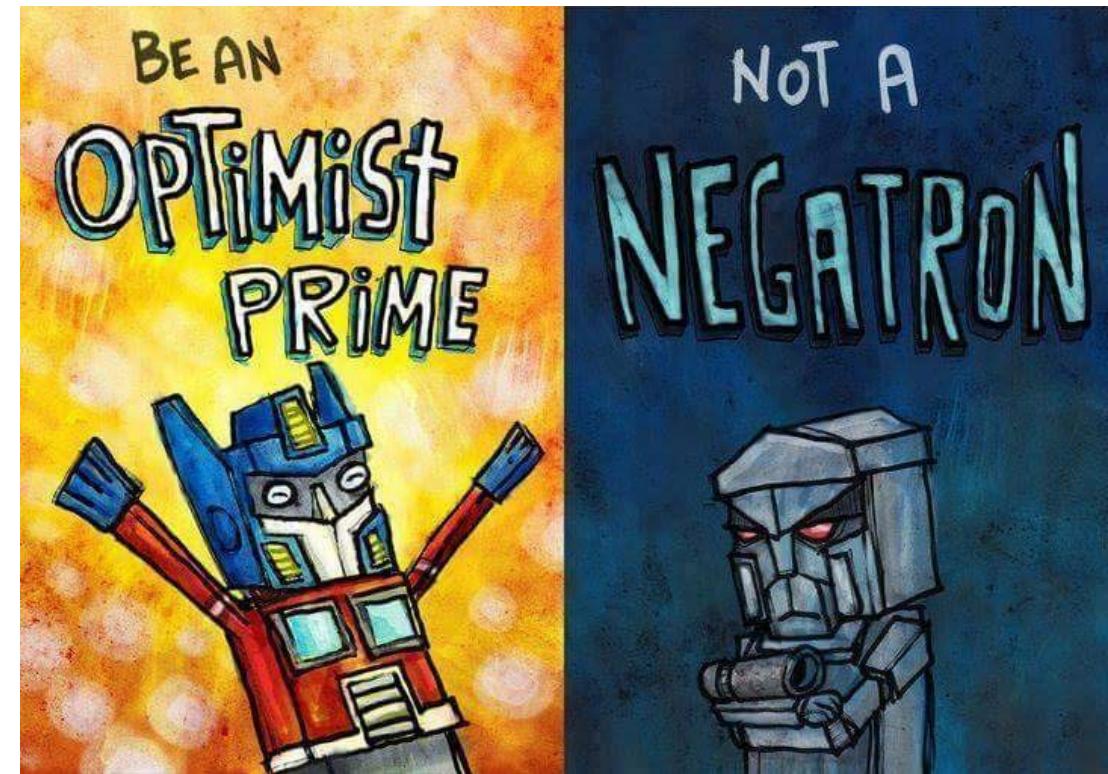
Choose  $m$  to be **prime**

- Avoid powers of 2 and powers of 10

In practice: popular and easy

But not always the most effective.

Slow (no more shifts)



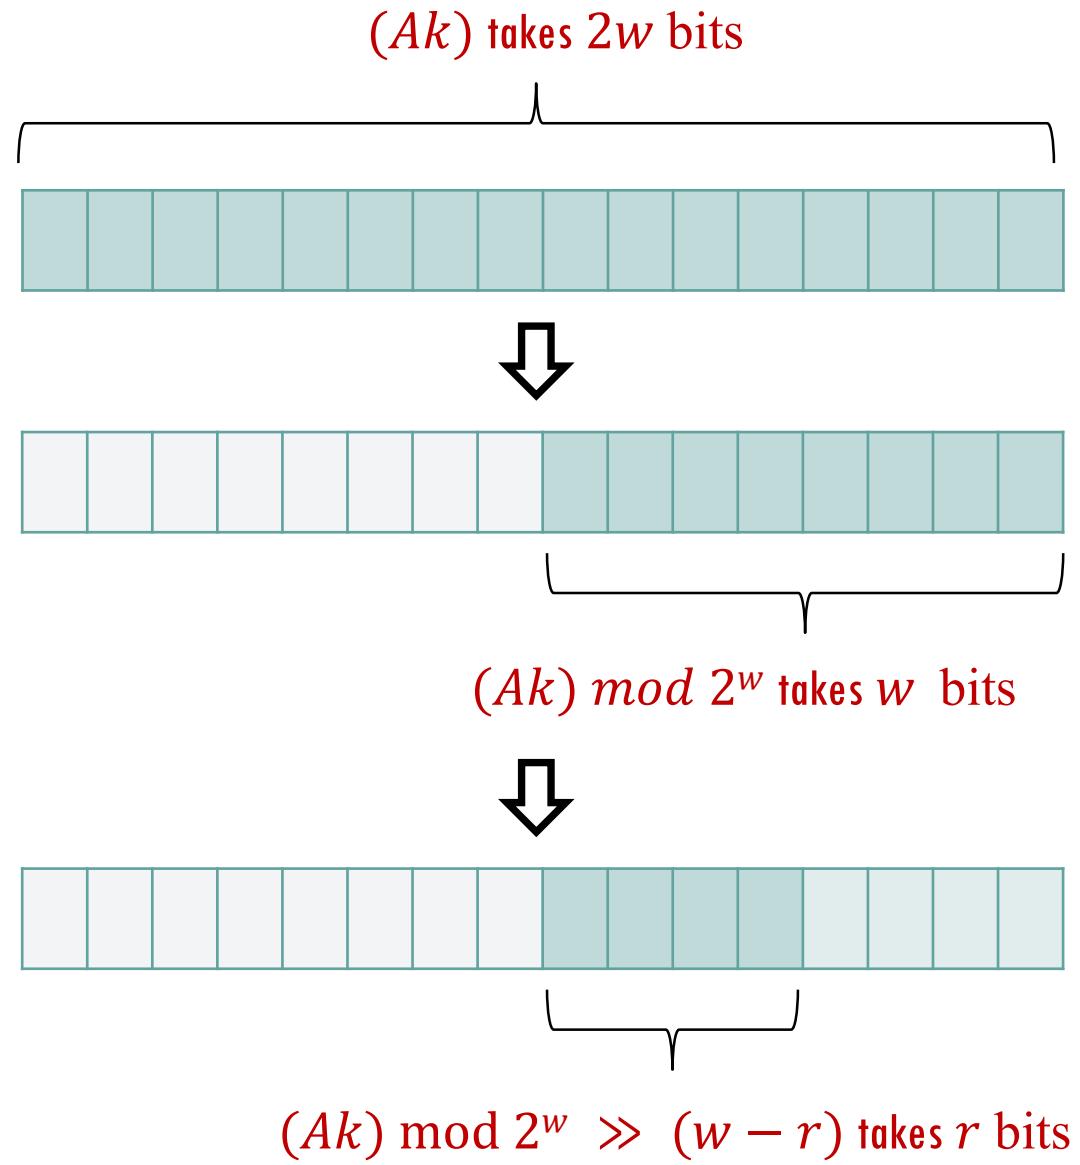
# MULTIPLICATION METHOD

**Fix**

- table size:  $m = 2^r$
- word size:  $w$  (size of a key in bits)
- constant:  $2^{w-1} < A < 2^w$

**Then:**

$$h(k) = (Ak) \bmod 2^w \gg (w - r)$$

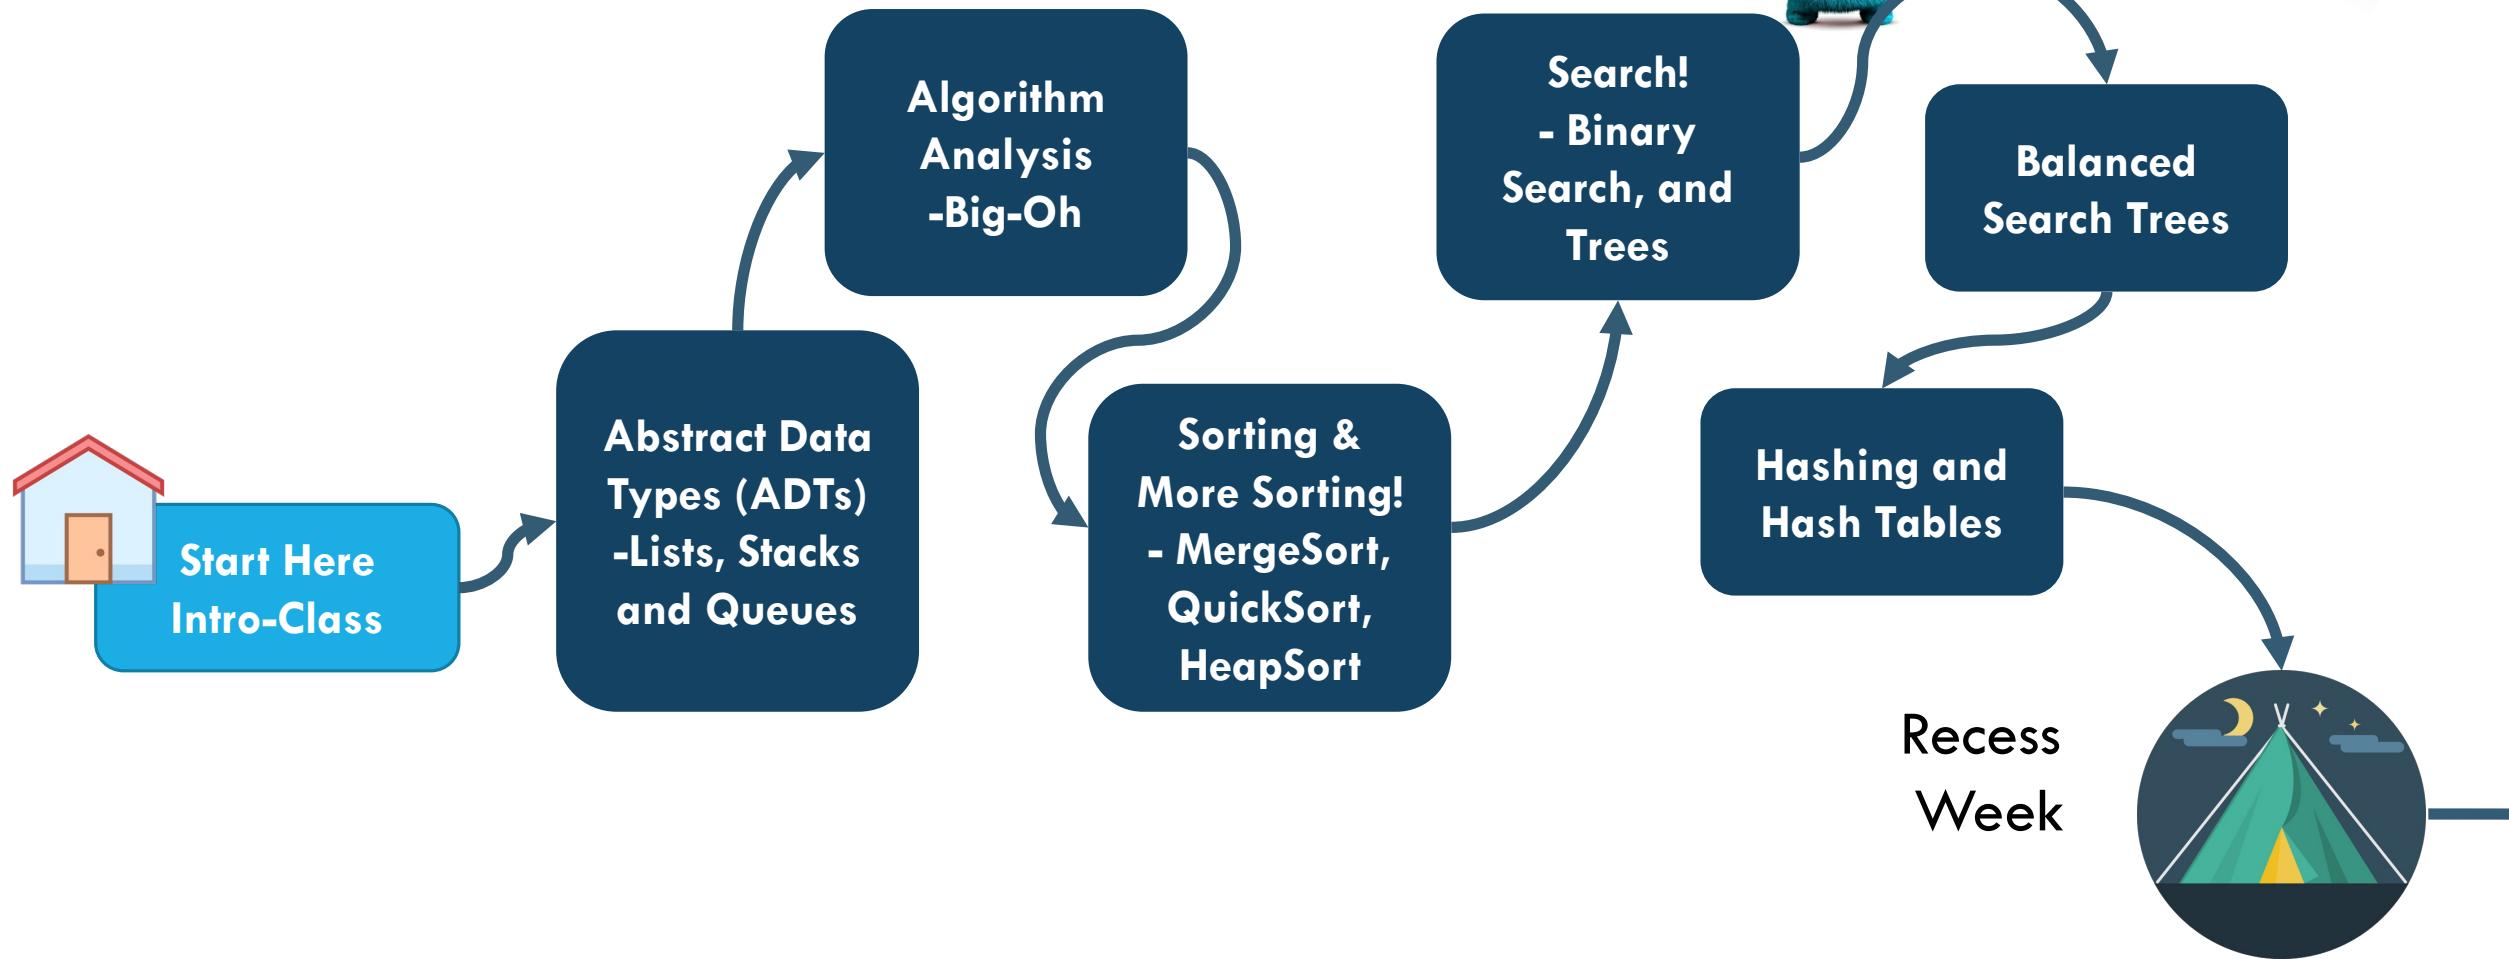




# QUESTIONS?

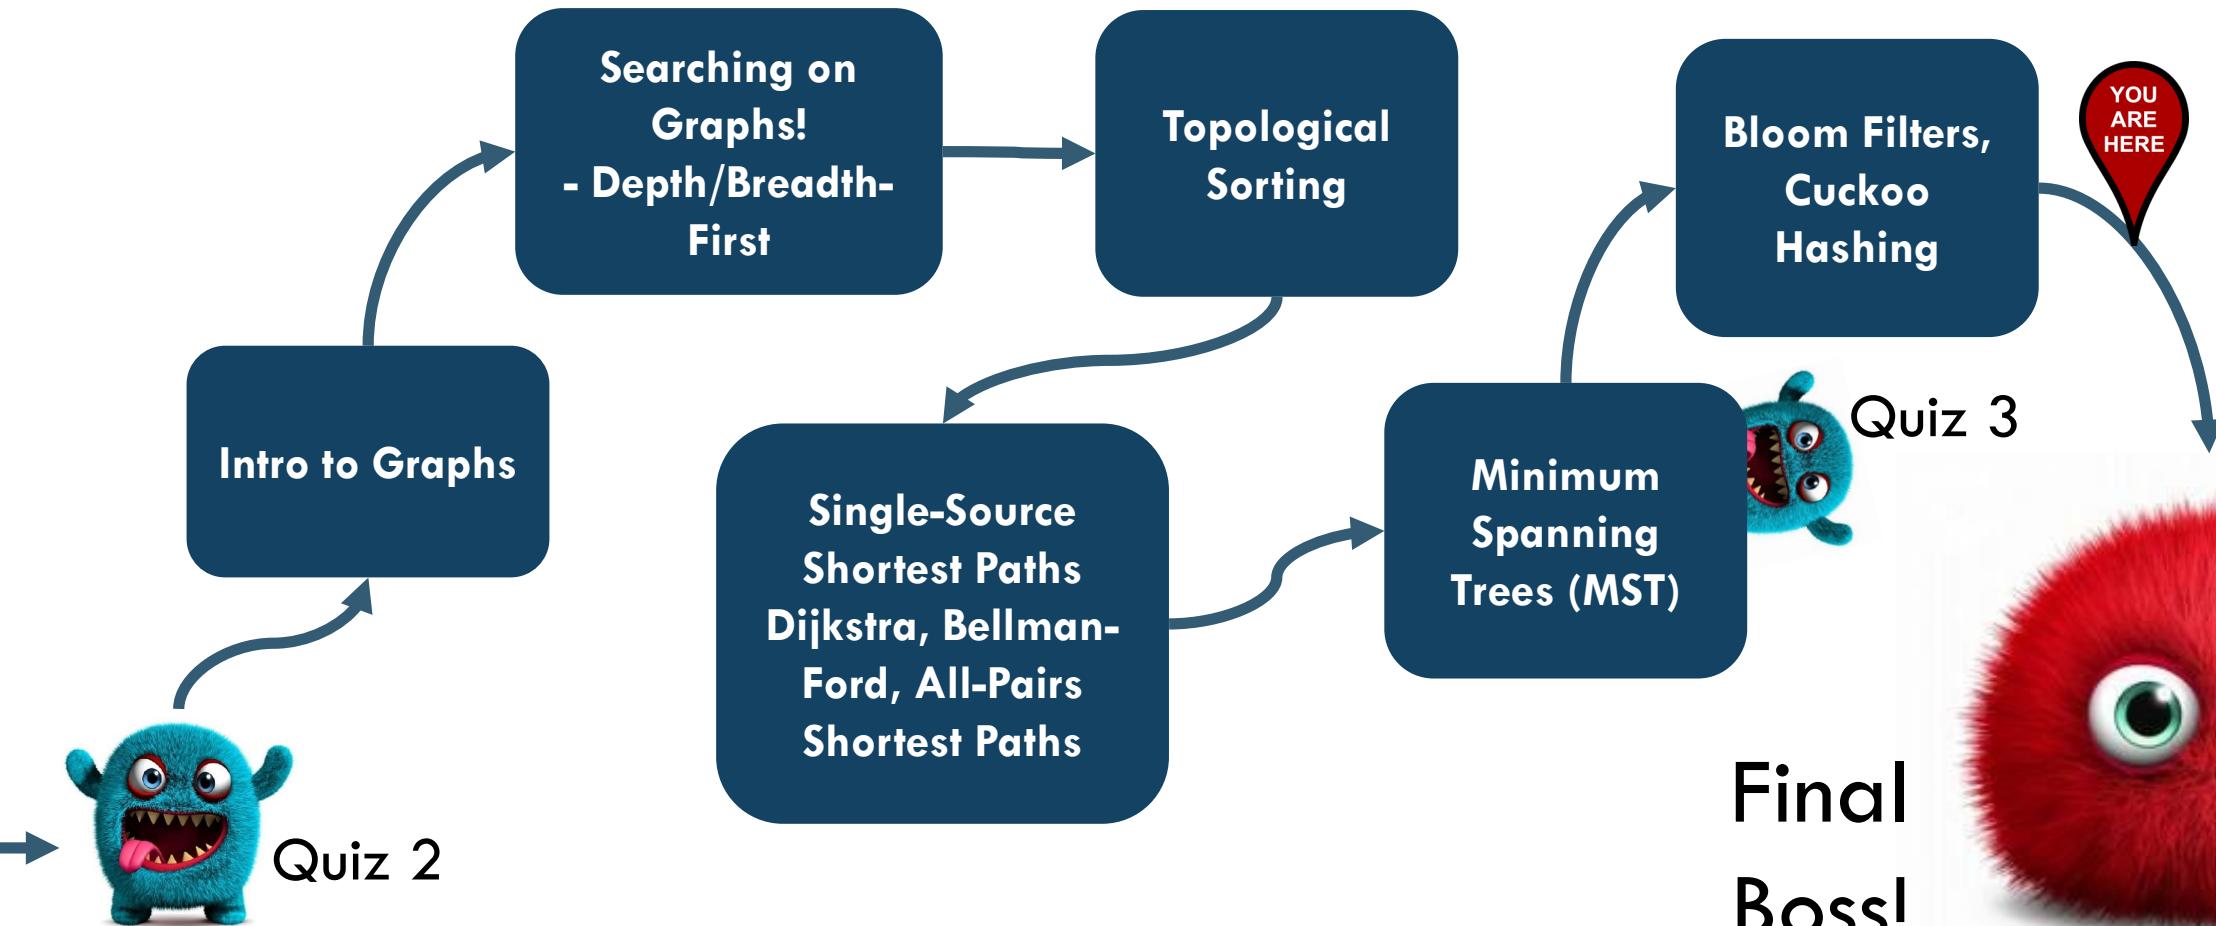


# PATH TO MASTERY / COURSE STRUCTURE



# PATH TO MASTERY / COURSE STRUCTURE

DRAFT



# QUESTION

I have to store a clique. Which takes more space (in terms of big-O): an adjacency matrix or an adjacency list?

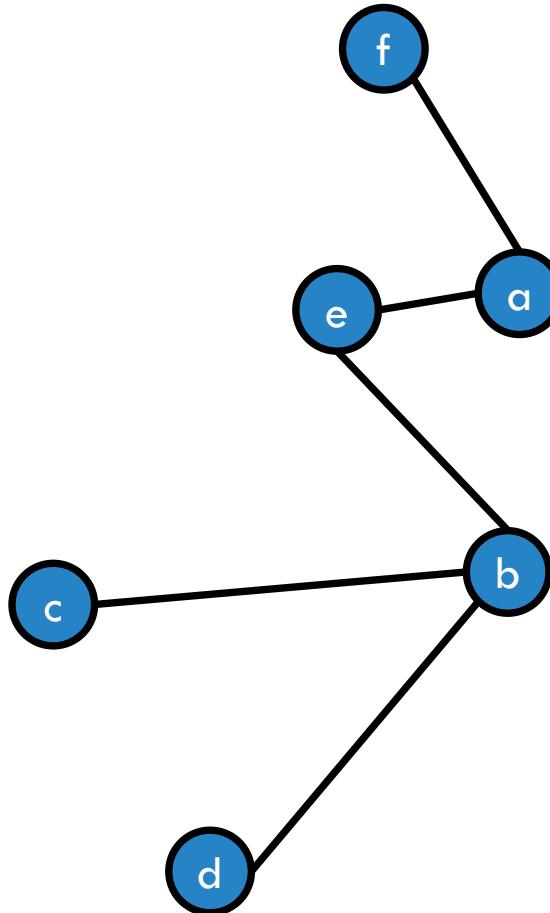
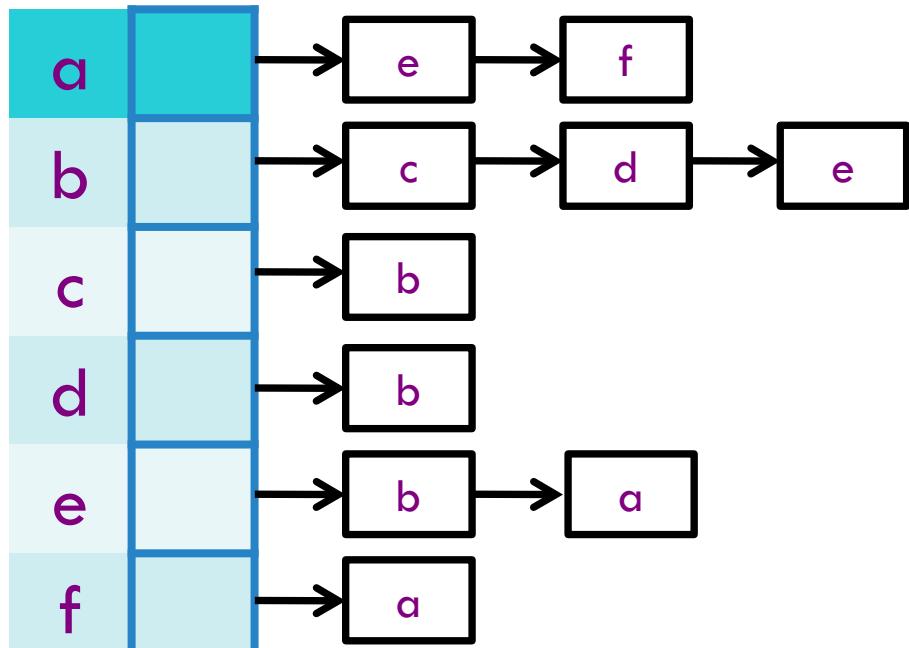
- A. Adjacency List
- B. Adjacency Matrix
- C. They both take the same space

# QUESTION

I have to store a clique. Which takes more space (in terms of big-O): an adjacency matrix or an adjacency list?

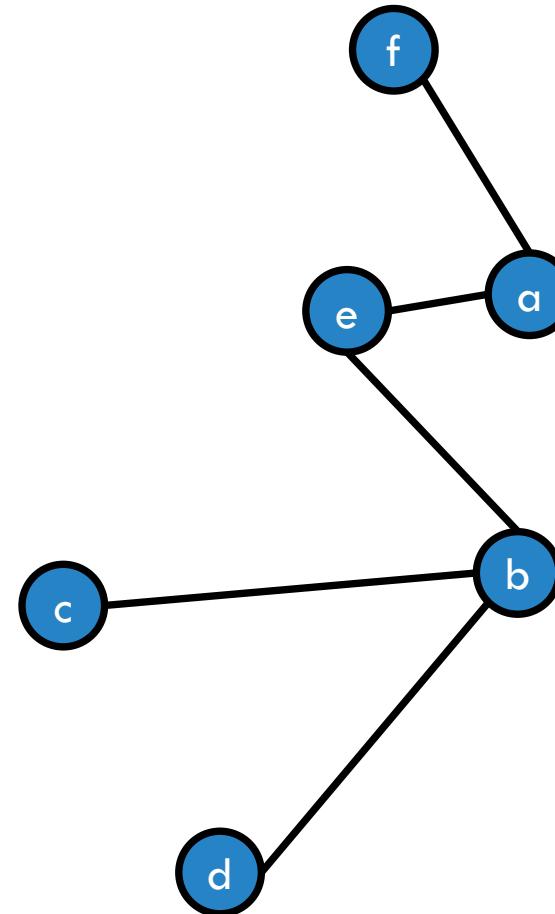
- A. Adjacency List
- B. Adjacency Matrix
- C. **They both take the same space**

# ADJACENCY LIST



# ADJACENCY MATRIX

	a	b	c	d	e	f
a	0	0	0	0	1	1
b	0	0	1	1	1	0
c	0	1	0	0	0	0
d	0	1	0	0	0	0
e	1	1	0	0	0	0
f	1	0	0	0	0	0



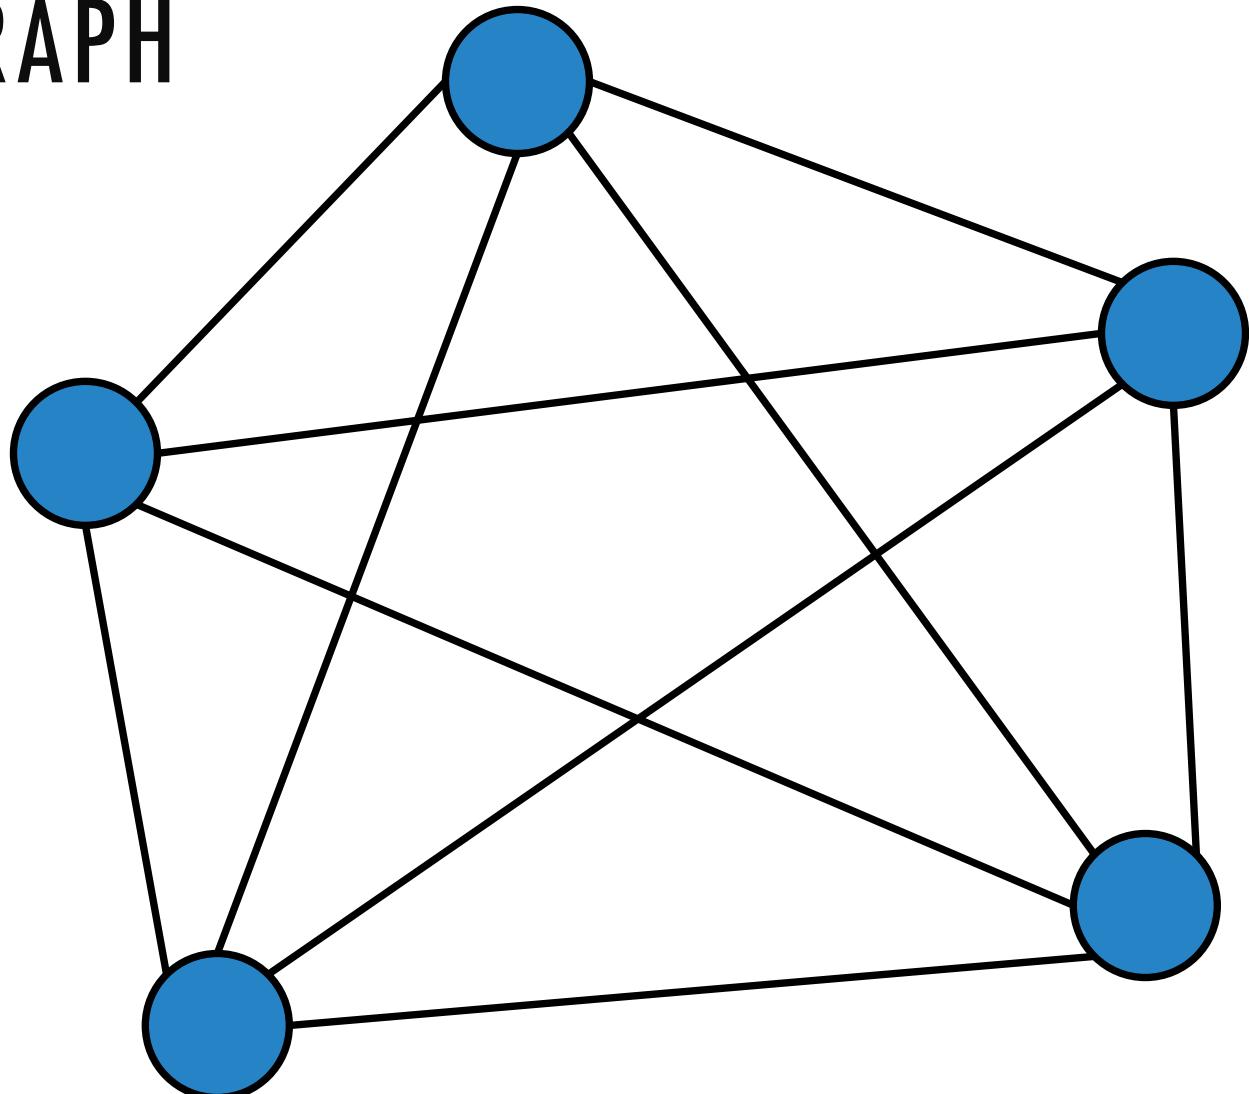
# CLIQUE: A COMPLETE GRAPH

Fully connected

All pairs connected by edges

**Question:** In a clique, what is the longest path from one node to another?

**One.** Every node is only 1 hop from any other node. The diameter of a clique = 1



# ADJACENCY LIST

Memory usage for graph  $G$ :  
array of size  $|V|$

linked lists of size  $|E|$

Total:  $O(V + E)$

For a cycle:  $O(V)$

For a clique:  $O(V^2)$

# ADJACENCY MATRIX

Memory usage for graph  $G$ :  
array of size  $|V| \times |V|$

Total:  $O(V^2)$

For a cycle:  $O(V^2)$

For a clique:  $O(V^2)$

# QUESTION

I want to construct a graph of countries and flights connecting the countries. A very common operation I will have to execute is:

“Find me all flights departing from country X”

Which representation is preferred: adjacency list or adjacency matrix?

- A. Adjacency List
- B. Adjacency Matrix

# QUESTION

I want to construct a graph of countries and flights connecting the countries. A very common operation I will have to execute is:

“Find me all flights departing from country X”

Which representation is preferred: adjacency list or adjacency matrix?

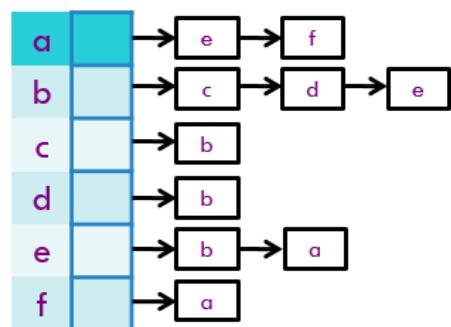
- A. **Adjacency List**
- B. Adjacency Matrix

# TRADE-OFFS

	a	b	c	d	e	f
a	0	0	0	0	1	1
b	0	0	1	1	1	0
c	0	1	0	0	0	0
d	0	1	0	0	0	0
e	1	1	0	0	0	0
f	1	0	0	0	0	0

## Adjacency Matrix:

- Fast query: are v and w neighbors?
- Slow query: find me any neighbor of v.
- Slow query: enumerate all neighbors.



## Adjacency List:

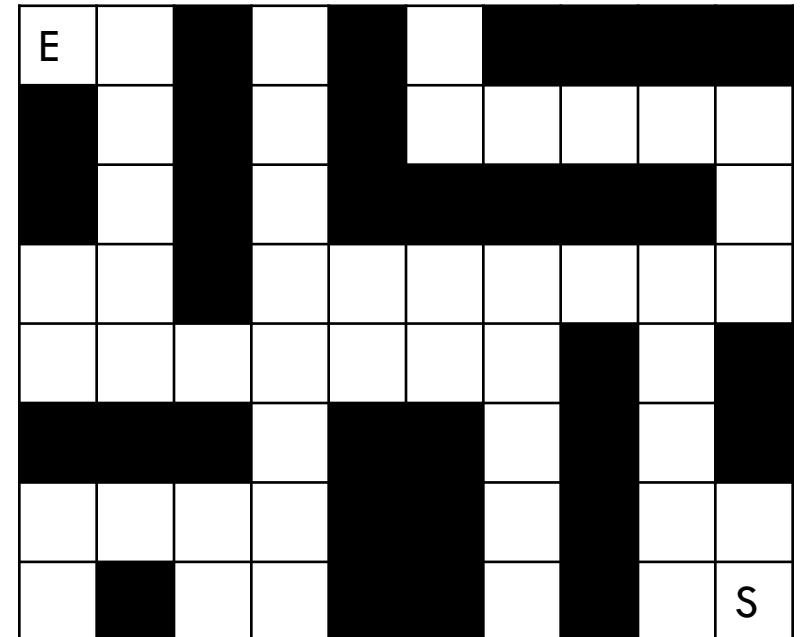
- Fast query: find me any neighbor.
- Fast query: enumerate all neighbors.
- Slower query: are v and w neighbors?

# QUESTION

I'm in a grid-maze where I need to find the quickest route to the exit. I can only move left, right, up or down. Taking each step requires 2 seconds.

Which of the following algorithms is most relevant for solving this problem?

- A. Breadth First Search
- B. Depth First Search
- C. Bellman-Ford
- D. Floyd-Warshall

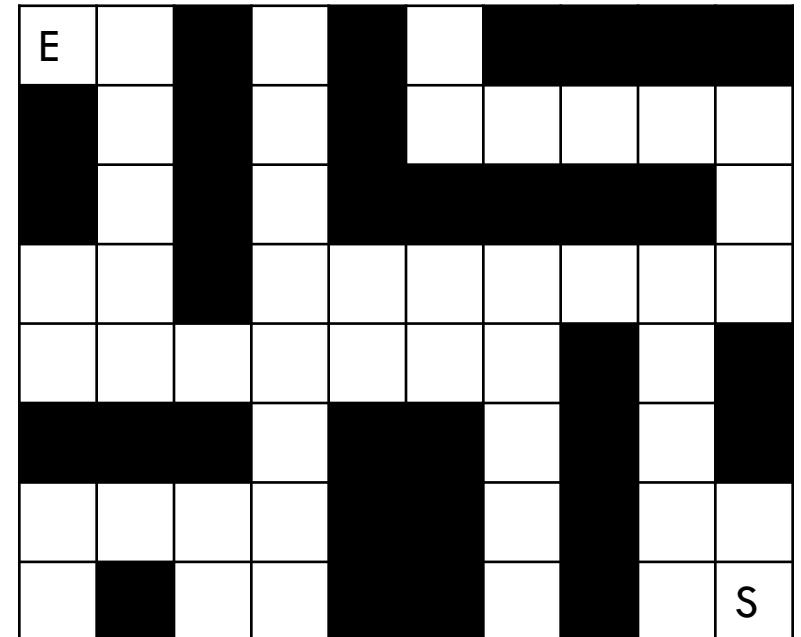


# QUESTION

I'm in a grid-maze where I need to find the quickest route to the exit. I can only move left, right, up or down. Taking each step requires 2 seconds.

Which of the following algorithms is most relevant for solving this problem?

- A. Breadth First Search
  - B. Depth First Search
  - C. Bellman-Ford
  - D. Floyd-Warshall



# BREADTH-FIRST SEARCH

Explore level by level

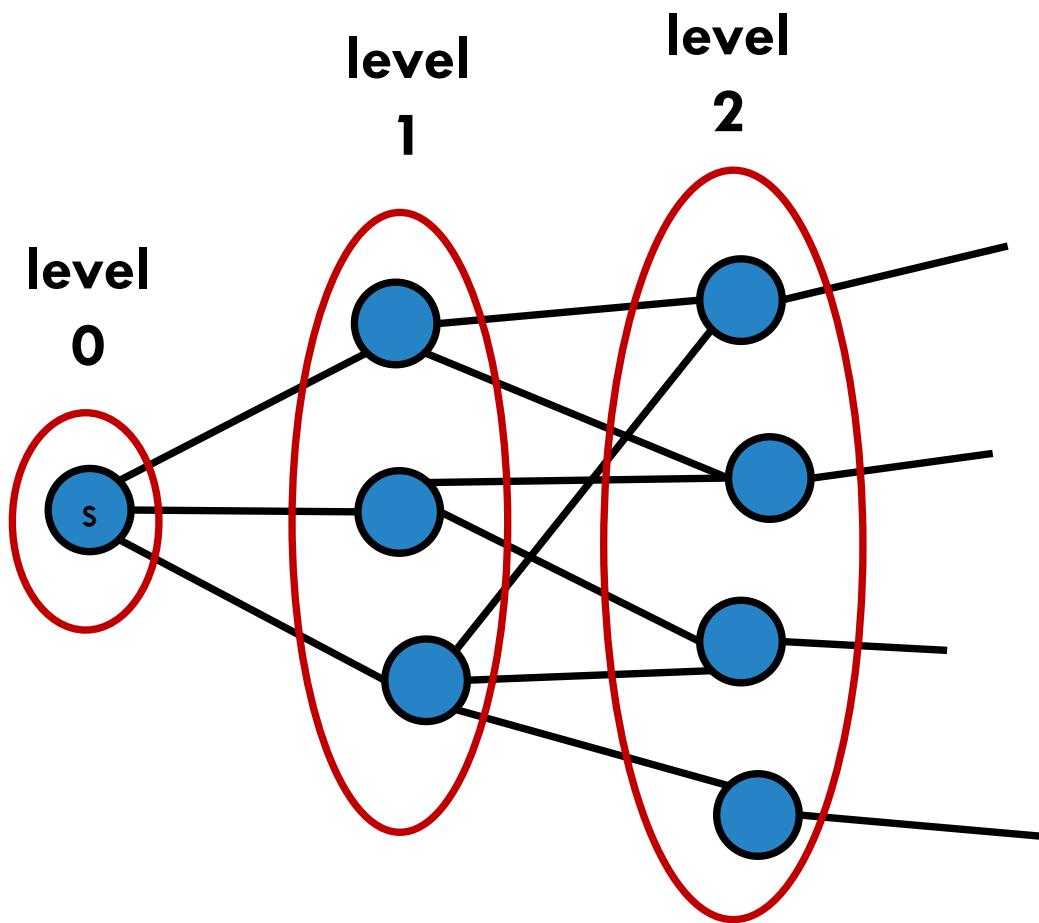
Frontier: current level

Initial frontier:  $\{s\}$

Advance frontier.

Don't go backward!

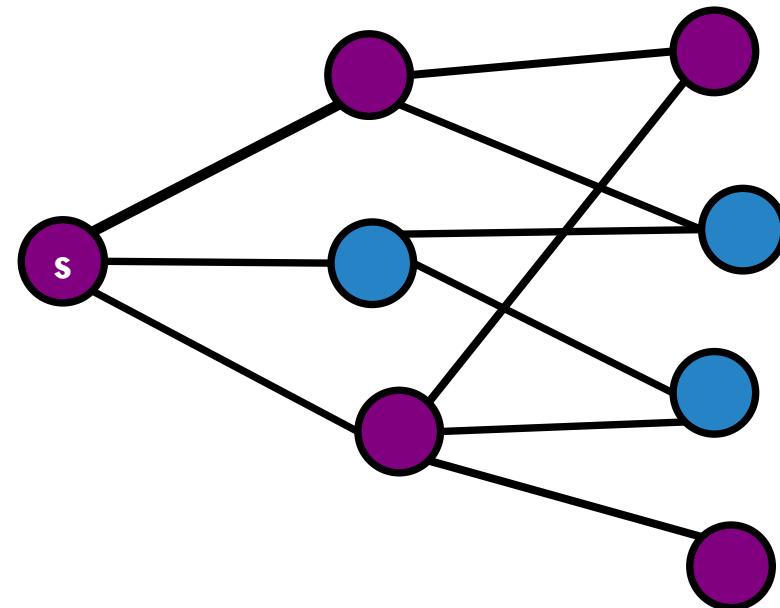
Finds shortest paths.



# DEPTH-FIRST SEARCH (DFS)

## Strategy

- Follow path until you get stuck
- Backtrack until you find a new edge
- Recursively explore it
- Don't repeat a vertex.



# QUESTION

I have a pretty complicated list of things to study before the CS2040 exam. Some topics have pre-requisite topics (e.g., I should study algorithm analysis before I study graphs algorithms).

How can order the topics so that I can optimize my studying?

- A. Breadth-First Search
- B. Topological Sort
- C. Dijkstra's Algorithm
- D. None of the above.

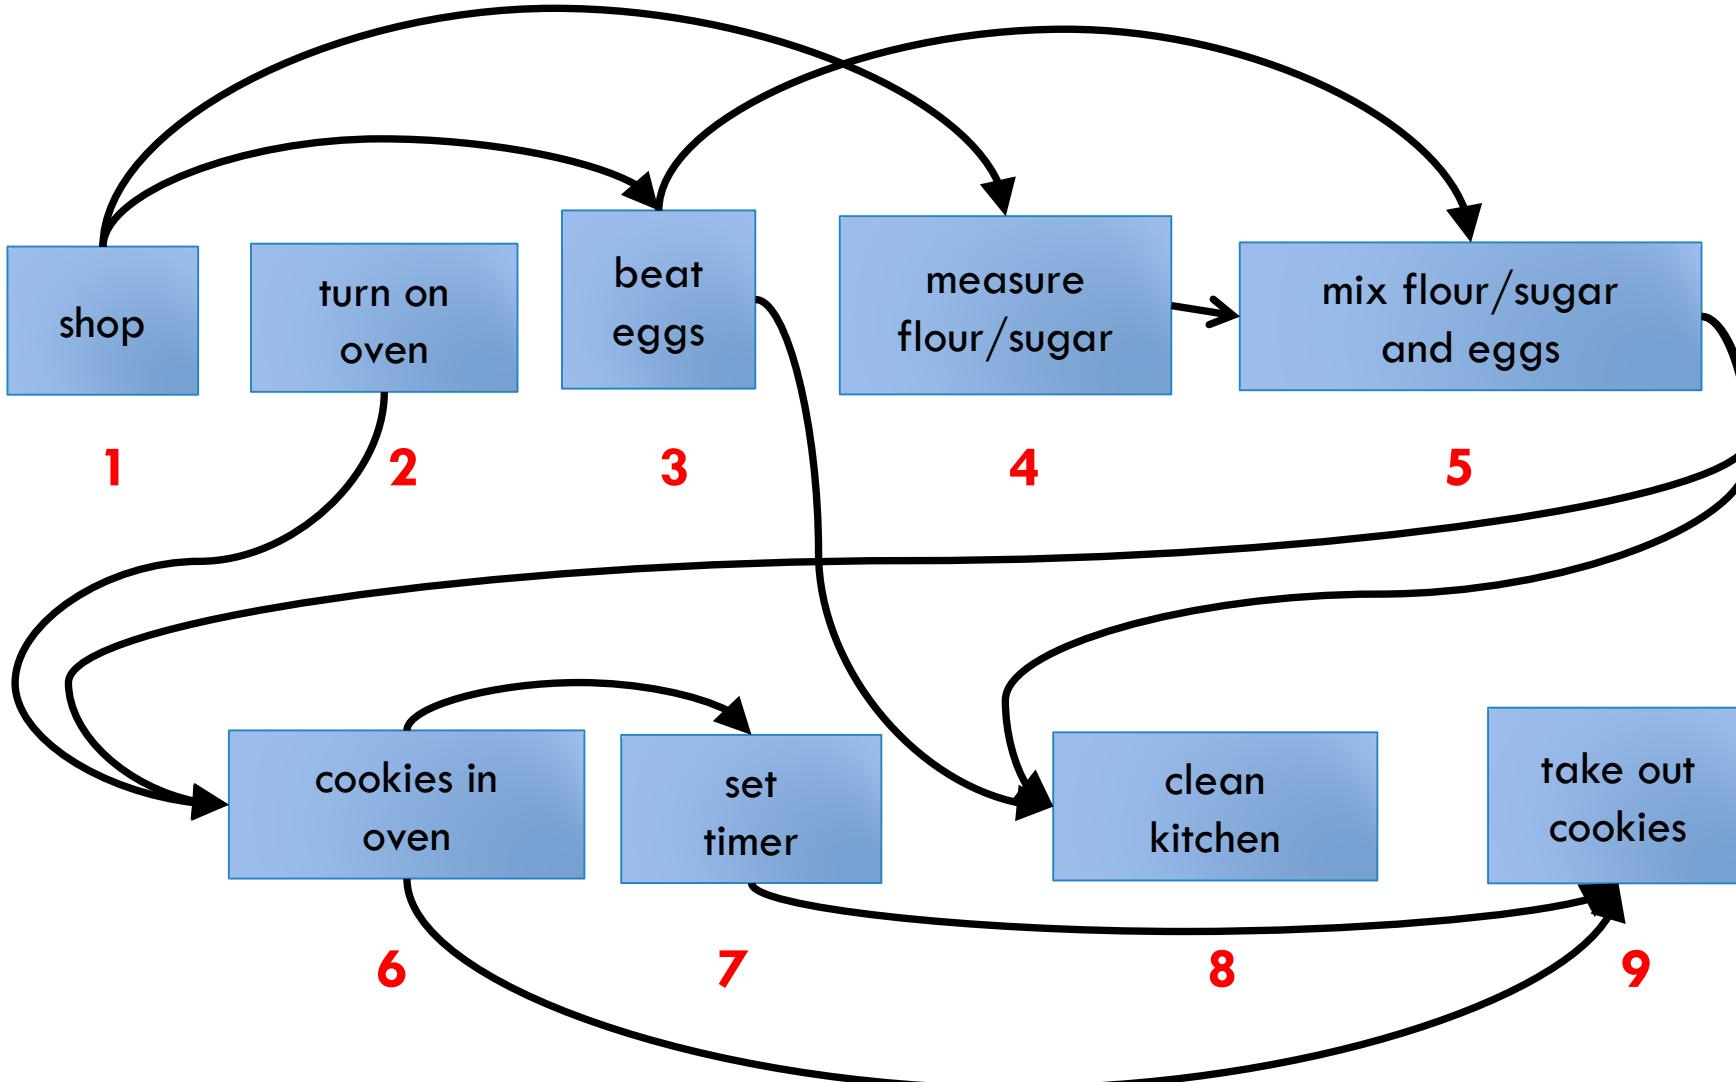
# QUESTION

I have a pretty complicated list of things to study before the CS2040 exam. Some topics have pre-requisite topics (e.g., I should study algorithm analysis before I study graphs algorithms).

How can order the topics so that I can optimize my studying?

- A. Breadth-First Search
- B. **Topological Sort**
- C. Dijkstra's Algorithm
- D. None of the above.

# TOPOLOGICAL ORDERING



## Assume adjacency list

# KAHN'S ALGORITHM

Start at any node  $v$  with no incoming edges.

Add  $v$  to our list

Remove  $v$  and all its outgoing edges.

Repeat

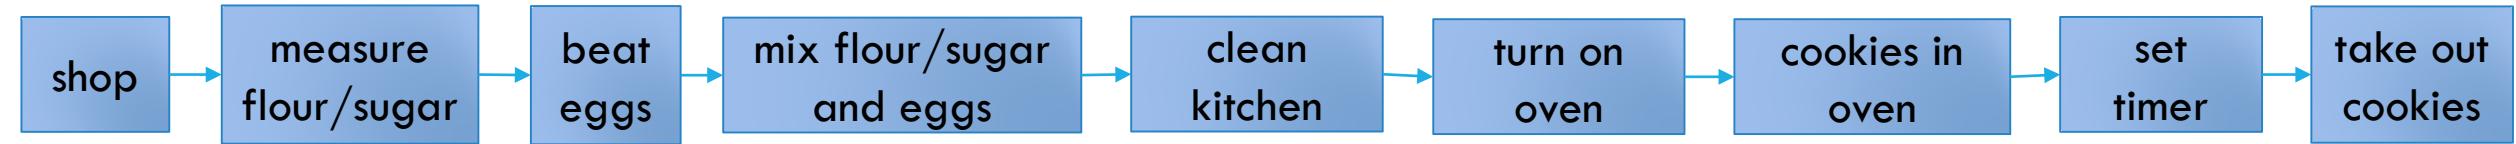
Pseudocode:

```
L = list()
S = list()

O(V) add all nodes with no incoming edge to S
while S is not empty:
    remove node v from S
    add v to tail of L
    for each of v's neighbors u
        remove edge e where source is v
        if u has no other incoming edges
            add u to S
```

$O(E)$

What is the time complexity?  $O(V + E)$



# TOPOLOGICAL SORT USING DFS (ASSUME DAG)

Idea: Process node when it is “last” visited.

```
L = list()  
while there are unvisited nodes  
    v = select unvisited node  
    DFS(G, v, L)
```

```
DFS(G,v,L)  
    if v is visited  
        return  
    else  
        for each of v's neighbor u  
            DFS(G, u, L)  
        visit(v)  
        L.pushFront(v)
```

# QUESTION

The code for relaxation is:

```
relax(int u, int v) {  
    if (dist[v] < dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```

- A. True
- B. False

# QUESTION

The code for relaxation is:

```
relax(int u, int v) {  
    if (dist[v] < dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```

- A. True
- B. False

# SHORTEST PATHS

```
relax(int u, int v) {  
    if (dist[v] > dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```

Maintain estimate for each distance:

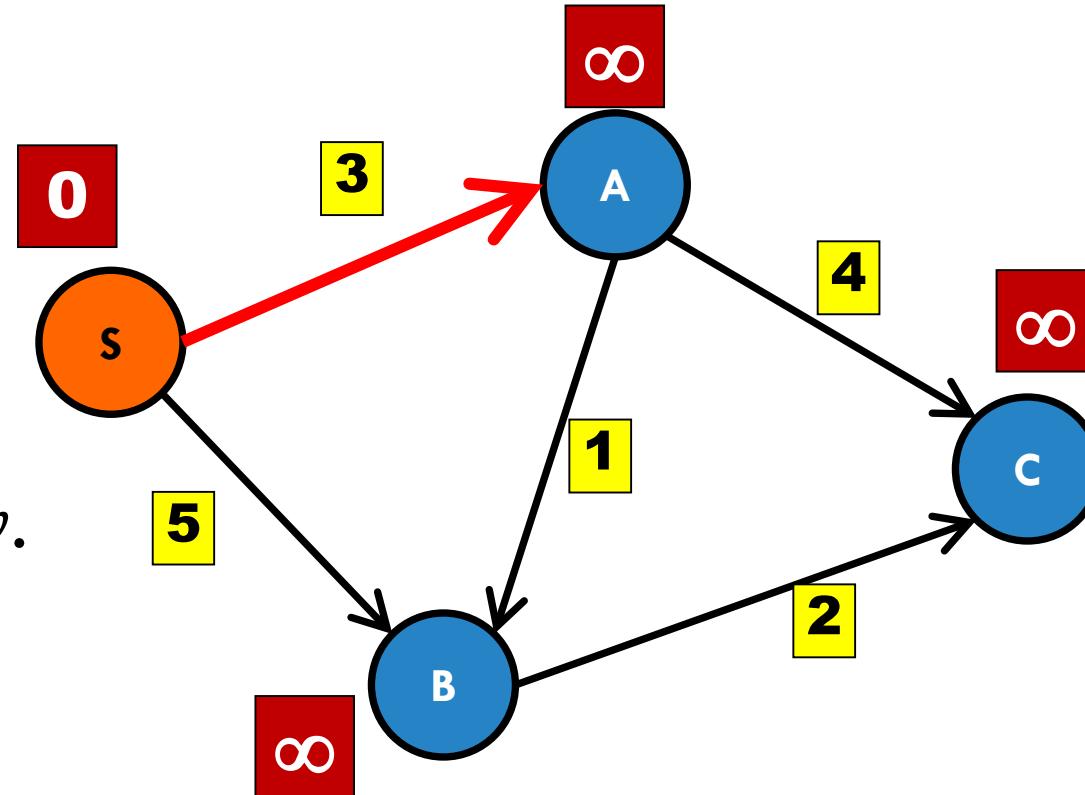
`relax(S, A)`

**The idea:**

`relax(w,v):`

Test if the best way to  
get from  $s \rightarrow v$  is to  
go from  $s \rightarrow w$ , then  $w \rightarrow v$ .

Update dist



# SHORTEST PATHS

```
relax(int u, int v) {  
    if (dist[v] > dist[u] + weight(u,v))  
        dist[v] = dist[u] + weight(u,v);  
}
```

Maintain estimate for each distance:

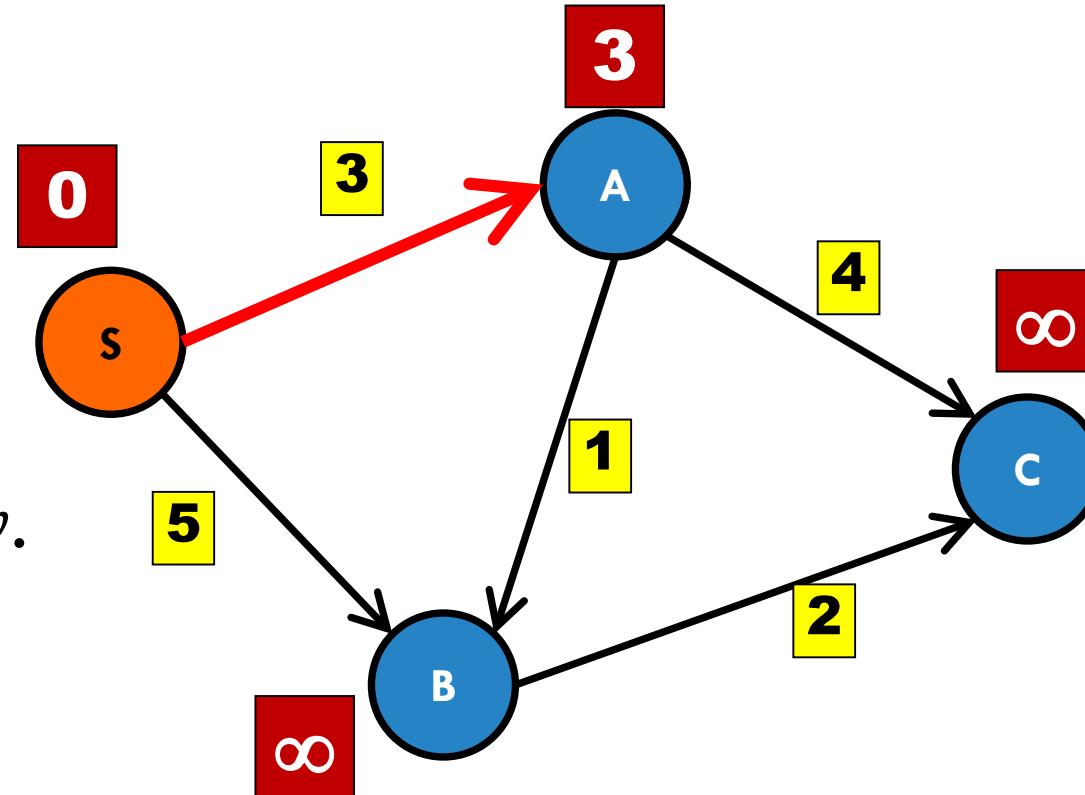
`relax(S, A)`

**The idea:**

`relax(w,v):`

Test if the best way to  
get from  $s \rightarrow v$  is to  
go from  $s \rightarrow w$ , then  $w \rightarrow v$ .

Update dist



# QUESTION

Who first invented the Bellman-Ford algorithm?

- A. Alfonso Shimbel
- B. Richard Bellman
- C. Lester Ford
- D. Edward Moore

# QUESTION

Who first invented the Bellman-Ford algorithm?

- A. **Alfonso Shimbel**
- B. Richard Bellman
- C. Lester Ford
- D. Edward Moore

# BELLMAN-FORD

## Bellman–Ford algorithm

---

From Wikipedia, the free encyclopedia

The Bellman–Ford algorithm is an algorithm that computes shortest paths from a single source vertex to all of the other vertices in a weighted digraph.<sup>[1]</sup> It is slower than Dijkstra's algorithm for the same problem, but more versatile, as it is capable of handling graphs in which some of the edge weights are negative numbers. The algorithm was first proposed by Alfonso Shimbel in 1955, but is instead named after Richard Bellman and Lester Ford, Jr., who published it in 1958 and 1956, respectively.<sup>[2]</sup> Edward F. Moore also published the same algorithm in 1957, and for this reason it is also sometimes called the Bellman–Ford–Moore algorithm.<sup>[1]</sup>

# QUESTION

Poor Shimbler. But he's in good company: many things aren't named after their original discoverers. What "law" describes this phenomena?

- A. Merton's Law
- B. Soh's Law
- C. Stigler's Law
- D. Moore's Law

# QUESTION

Poor Shimbler. But he's in good company: many things aren't named after their original discoverers. What "law" describes this phenomena?

- A. Merton's Law
- B. Soh's Law
- C. **Stigler's Law**
- D. Moore's Law

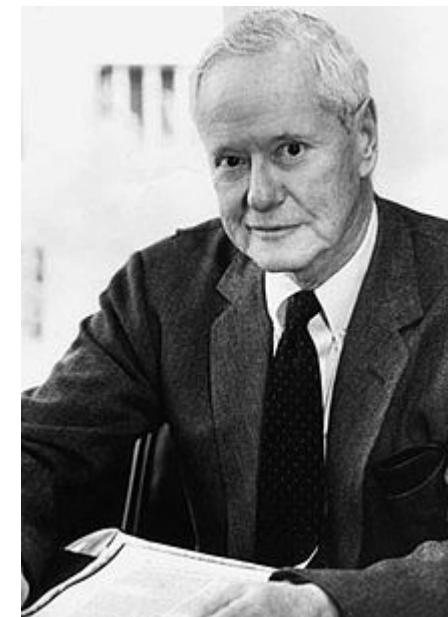
# STIGLER'S LAW OF EPONYMY

*No scientific discovery is named after its original discoverer.*

proposed by University of Chicago statistics professor Stephen Stigler.

Stigler named Columbia University sociology professor **Robert K. Merton** as the original discoverer of “Stigler's law”

Stigler's law follows Stigler's law.



# QUESTION

True/False: Bellman-Ford involves relaxing edges and requires relaxing all the edges for  $|V|-1$  iterations.

- A. True
- B. False

# QUESTION

True/False: Bellman-Ford involves relaxing edges and requires relaxing all the edges for  $|V|-1$  iterations.

- A. True
- B. False

# EARLY TERMINATION?

```
n = V.length  
for i = 1 to n-1  
    for Edge e in Graph  
        relax(e)
```

When can we terminate early?

- A. When a relax operation has no effect.
- B. When two consecutive relax operations have no effect.
- C. When an entire sequence of  $|E|$  relax operations have no effect.
- D. Never. Only after  $|V|$  complete iterations.

# QUESTION

I want to get from SoC to ION Orchard so that I can get some LeTAO cheesecakes.

Which algorithm should I use to find the shortest path (distance)?

- A. Bellman-Ford
- B. Dijkstra's Algorithm
- C. Floyd Warshall
- D. Breadth-First Search

# QUESTION

I want to get from SoC to ION Orchard so that I can get some LeTAO cheesecakes.

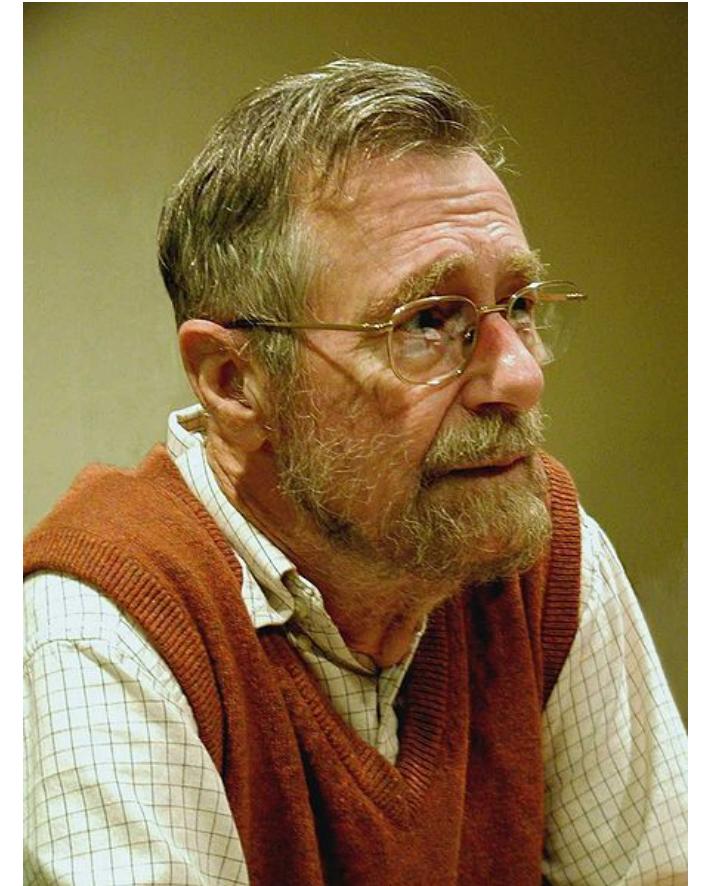
Which algorithm should I use to find the shortest path (distance)?

- A. Bellman-Ford
- B. **Dijkstra's Algorithm**
- C. Floyd Warshall
- D. Breadth-First Search

# DIJKSTRA'S ALGORITHM

## Basic idea:

- Maintain distance **estimate** for every node.
- Begin with empty shortest-path-tree.
- Repeat:
  - Consider vertex with minimum **estimate**.
  - Add vertex to **shortest-path-tree**.
  - Relax all outgoing edges.



1930-2002

# QUESTION

Which ADT and data structure should I use to store the vertices for use in the standard Dijkstra's algorithm?

- A. Priority Queue (Heap)
- B. Priority Queue (AVL Tree)
- C. Priority Queue (Linked List)
- D. Stack (Array)
- E. Stack (AVL Tree)

# QUESTION

Which ADT and data structure should I use to store the vertices for use in the standard Dijkstra's algorithm?

- A. Priority Queue (Heap)
- B. **Priority Queue (AVL Tree)**
- C. Priority Queue (Linked List)
- D. Stack (Array)
- E. Stack (AVL Tree)

# DIJKSTRA'S ALG: PERFORMANCE

PQ Implementation	insert	deleteMin	decreaseKey	Total
Array	1	$V$	1	$O(V^2)$
AVL Tree	$\log V$	$\log V$	$\log V$	$O((V + E)\log V)$
d-way Heap	$d\log_d V$	$d\log_d V$	$\log_d V$	$O(E \log_{\frac{V}{d}} V)$
Fibonacci Heap	1	$\log V$	1	$O(E + V \log V)$

# QUESTION

I give you two identical graphs  $G$  and  $G'$  that differ only in the following respect: for every edge  $e \in G$ , the corresponding graph  $e' \in G'$  has weight  $w(e') = w(e)^2$

Then every shortest path in  $G$  is also a shortest path in  $G'$ .

True or False?

- A. True
- B. False

# QUESTION

I give you two identical graphs  $G$  and  $G'$  that differ only in the following respect: for every edge  $e \in G$ , the corresponding graph  $e' \in G'$  has weight  $w(e') = w(e)^2$

Then every shortest path in  $G$  is also a shortest path in  $G'$ .

True or False?

- A. True
- B. False

# QUESTION

Given a graph  $G$ , if I partition the nodes into two subsets, the maximum edge across the cut is never in the MST.

- A. True
- B. False

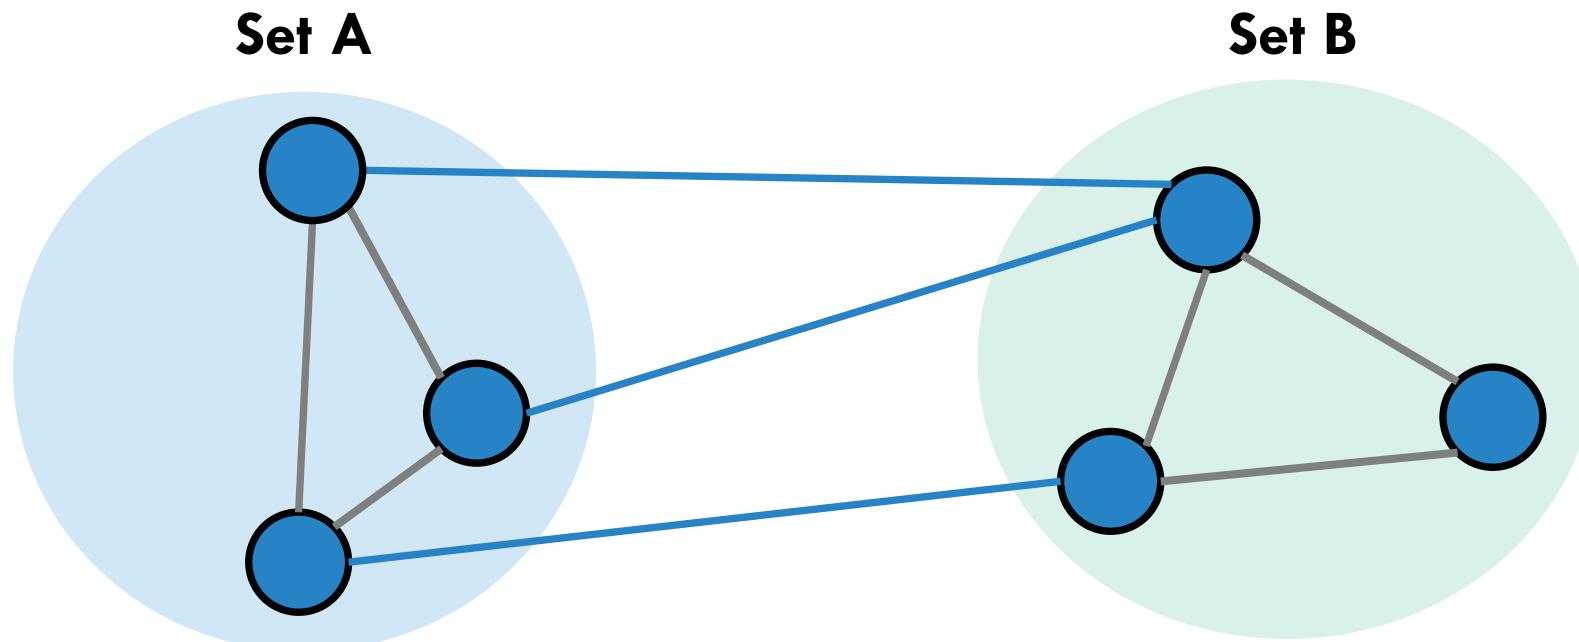
# QUESTION

Given a graph  $G$ , if I partition the nodes into two subsets, the maximum edge across the cut is never in the MST.

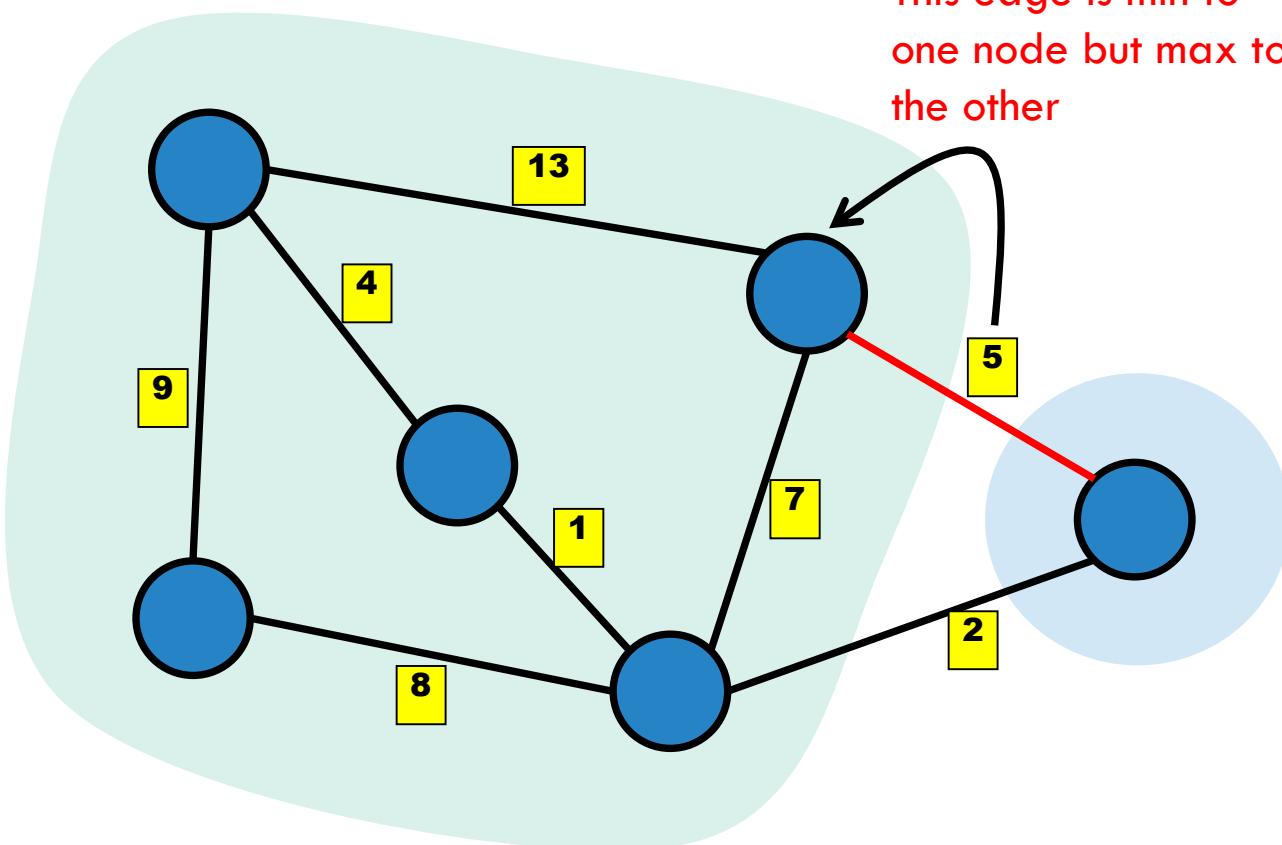
- A. True
- B. False

# WHAT IS A CUT?

**Definition:** A ***cut*** of a graph  $G = (V, E)$  is a partition of the vertices  $V$  into two disjoint subsets.



# TRUE OR FALSE?



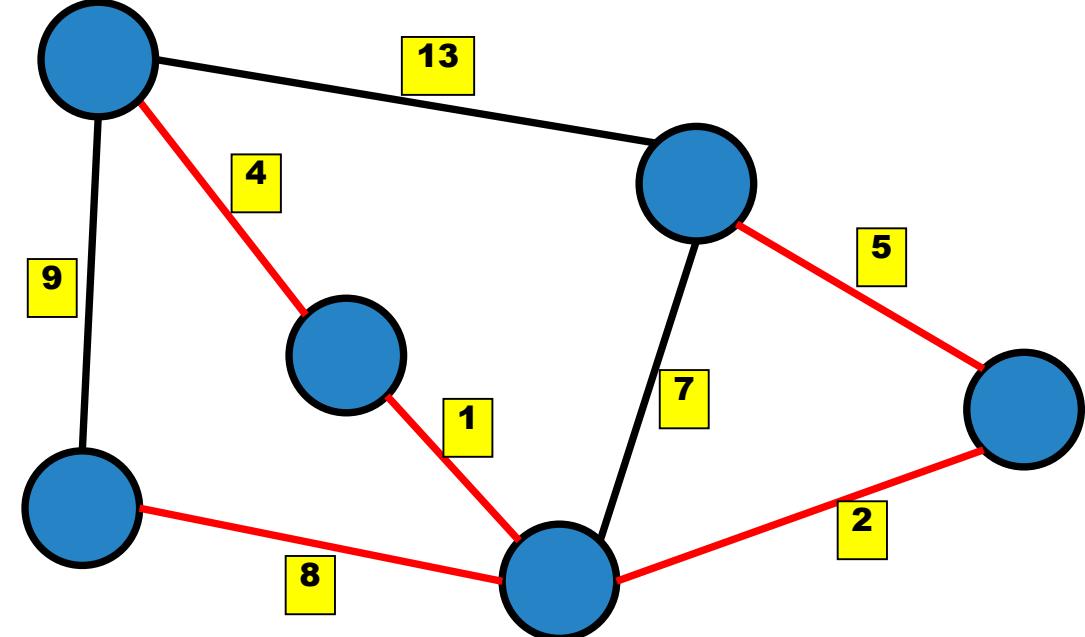
For every vertex, the maximum incident/outgoing edge is **never** part of the MST.

- A. True! You're not tricking me again!
- B. **False.**
- C. THE STATEMENT BELOW IS FALSE



# MST PROPERTIES: SUMMARY

1. No cycles
2. If you cut an MST, the two pieces are both MSTs.
3. Cycle property
  - For every cycle, the maximum weight edge is not in the MST.
4. Cut property
  - For every cut D, the minimum weight edge that crosses the cut is in the MST.



# QUESTION

I have a graph with edges already sorted by their weights. To construct a MST, which algorithm is most relevant?

- A. Kruskal's algorithm
- B. Prim's algorithm
- C. Floyd-Warshall
- D. Dijkstra's Shortest Path Algorithm

# QUESTION

I have a graph with edges already sorted by their weights. To construct a MST, which algorithm is most relevant?

- A. Kruskal's algorithm
- B. Prim's algorithm
- C. Floyd-Warshall
- D. Dijkstra's Shortest Path Algorithm

# KRUSKAL'S ALGORITHM

## Basic idea:

- Graph  $F$  : a set of trees (initially each vertex is a separate tree)
- Set of edges  $S = \{e \in E\}$
- While  $S$  is nonempty and  $F$  is not spanning:
  - Remove minimum weight edge from  $S$
  - If removed edge connects two trees
    - add it to the  $F$  (combine the trees)

Sort takes  $O(E \log E)$

$$O(E \log E) = O(E \log V^2) = O(E \log V)$$

For  $E$  edges, find/union take  $O(E\alpha)$

**Total:**  $O(E \log V)$

# REVIEW: PRIM'S ALGORITHM

## Basic idea:

- $S$  : set of nodes connected by blue edges.
- Initially:  $S = \{A\}$
- Repeat:
  - Identify cut:  $\{S, V - S\}$
  - Find minimum weight edge on cut.
  - Add new node to  $S$ .

Prim's “grows” the tree one node at a time.

## Analysis:

Each vertex added/removed once from the priority queue:  
 $O(V \log V)$

Each edge  $\rightarrow$  one decreaseKey:  
 $O(E \log V)$

# QUESTION

Let  $G = (V, E)$  be a weighted undirected graph and  $T$  be a MST of  $G$ . For any pair of vertices  $u$  and  $v$ , the path between  $u$  and  $v$  in  $T$  must be the shortest path between  $u$  and  $v$  in  $G$ .

Is the above statement True or False?

- A. True
- B. False

# QUESTION

Let  $G = (V, E)$  be a weighted undirected graph and  $T$  be a MST of  $G$ . For any pair of vertices  $u$  and  $v$ , the path between  $u$  and  $v$  in  $T$  must be the shortest path between  $u$  and  $v$  in  $G$ .

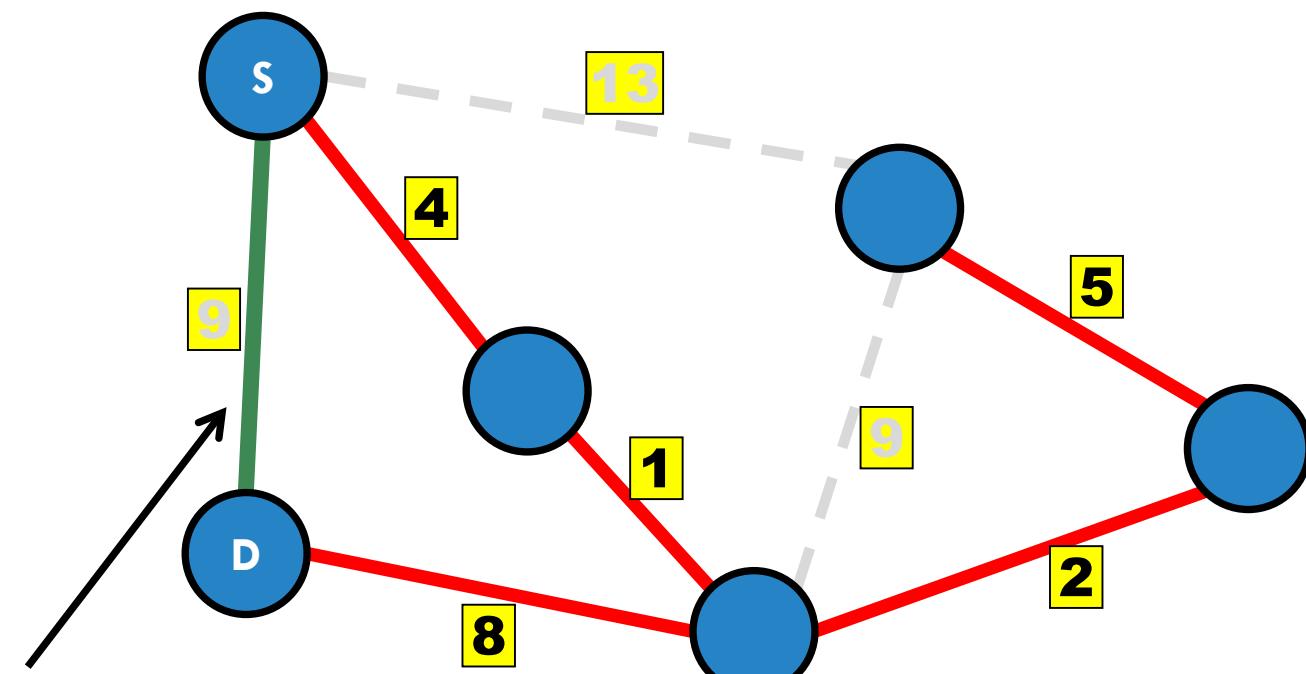
Is the above statement True or False?

- A. True
- B. False



# MINIMUM SPANNING TREE

MST is not the same as shortest paths



Shortest path  
from S to D

Can we use the MST to find the shortest paths?

- A. Yes
- B. Only on connected graphs
- C. Only on dense graphs
- D. **Nope.**
- E.

# QUESTION

The depth of a BFS tree on an undirected graph from an arbitrary vertex  $v$  is the diameter of the graph.

True or False?

- A. True
- B. False

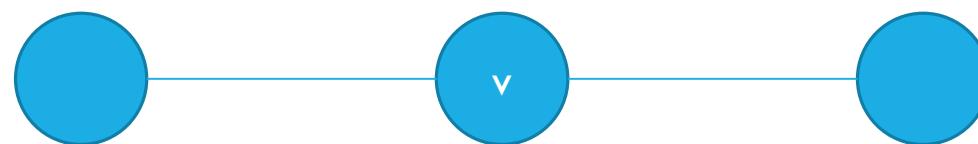
# QUESTION

The depth of a BFS tree on an undirected graph from an arbitrary vertex  $v$  is the diameter of the graph.

True or False?

- A. True
- B. False

# COUNTER-EXAMPLE



# QUESTION

You work for Foogle Maps. A very frequent query on Foogle Maps is “*How long does it take to get from A to B?*” where A and B are two arbitrary locations on the map. What algorithm is most relevant for this problem?

- A. Dijkstra's Algorithm
- B. Breadth-First Search
- C. Floyd-Warshall
- D. Bellman-Ford
- E. Kruskal's Algorithm

# QUESTION

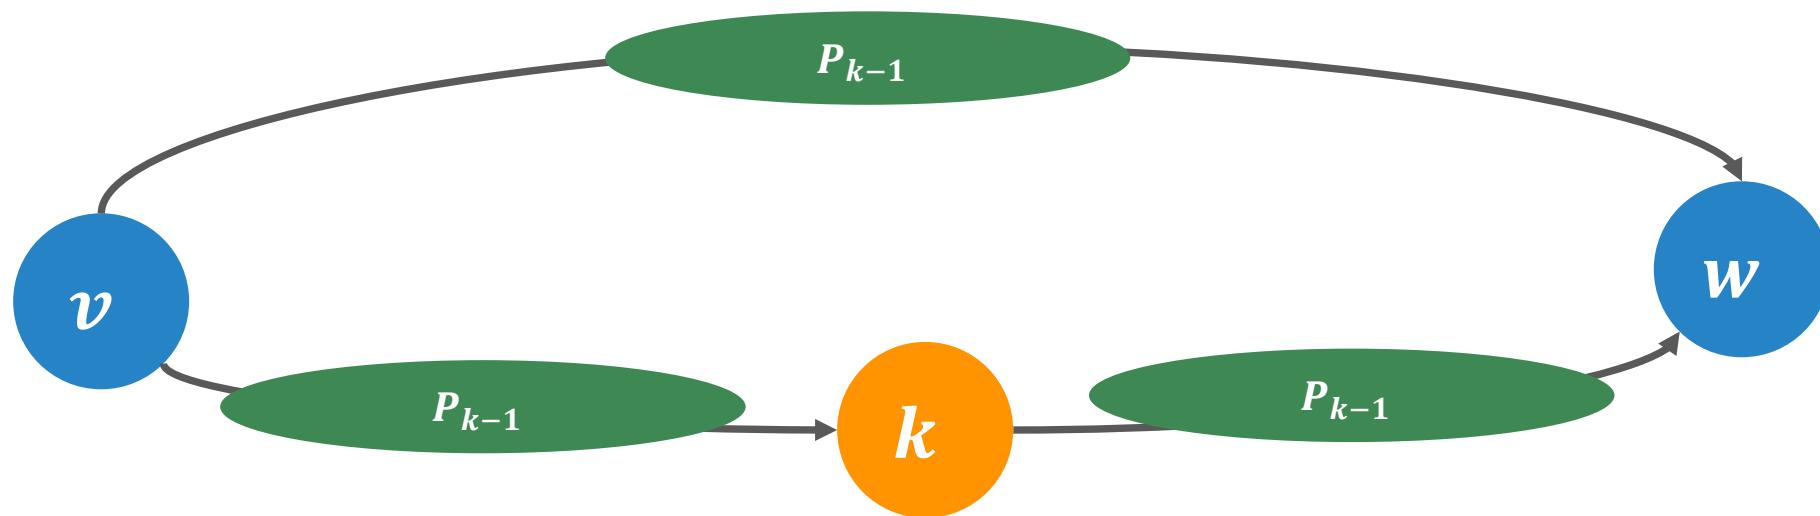
You work for Foogle Maps. A very frequent query on Foogle Maps is “*How long does it take to get from A to B?*” where A and B are two arbitrary locations on the map. What algorithm is most relevant for this problem?

- A. Dijkstra's Algorithm
- B. Breadth-First Search
- C. **Floyd-Warshall**
- D. Bellman-Ford
- E. Kruskal's Algorithm

# FLOYD-WARSHALL

Use the **precalculated** subproblems:

$$S[v, w, P_k] = \min( S[v, w, P_{k-1}], S[v, k, P_{k-1}] + S[k, w, P_{k-1}] )$$



# FLOYD-WARSHALL: PSEUDOCODE

```
Function FloydWarshall(G)
```

```
    S = Array of size |V|×|V| //memoization table S has |V| rows and |V| columns
```

```
// Initialize every pair of nodes
```

```
    for v = 0 to |V|-1
```

```
        for w = 0 to |V|-1
```

```
            S[v,w] = E[v,w]
```

```
// For sets P0, P1, P2, P3, ..., for every pair (v,w)
```

```
    for k = 0 to |V|-1
```

```
        for v = 0 to |V|-1
```

```
            for w = 0 to |V|-1
```

```
                S[v,w] = min(S[v,w], S[v,k]+S[k,w])
```

```
return S
```

**What is the running time?  $O(V^3)$**

# QUESTION

When is the CS2040 exam?

- A. 5-Dec-2019 (9am)
- B. 4-Dec-2019 (9am)
- C. 3-Dec-2019 (9am)
- D. 2-Dec 2019 (12pm)

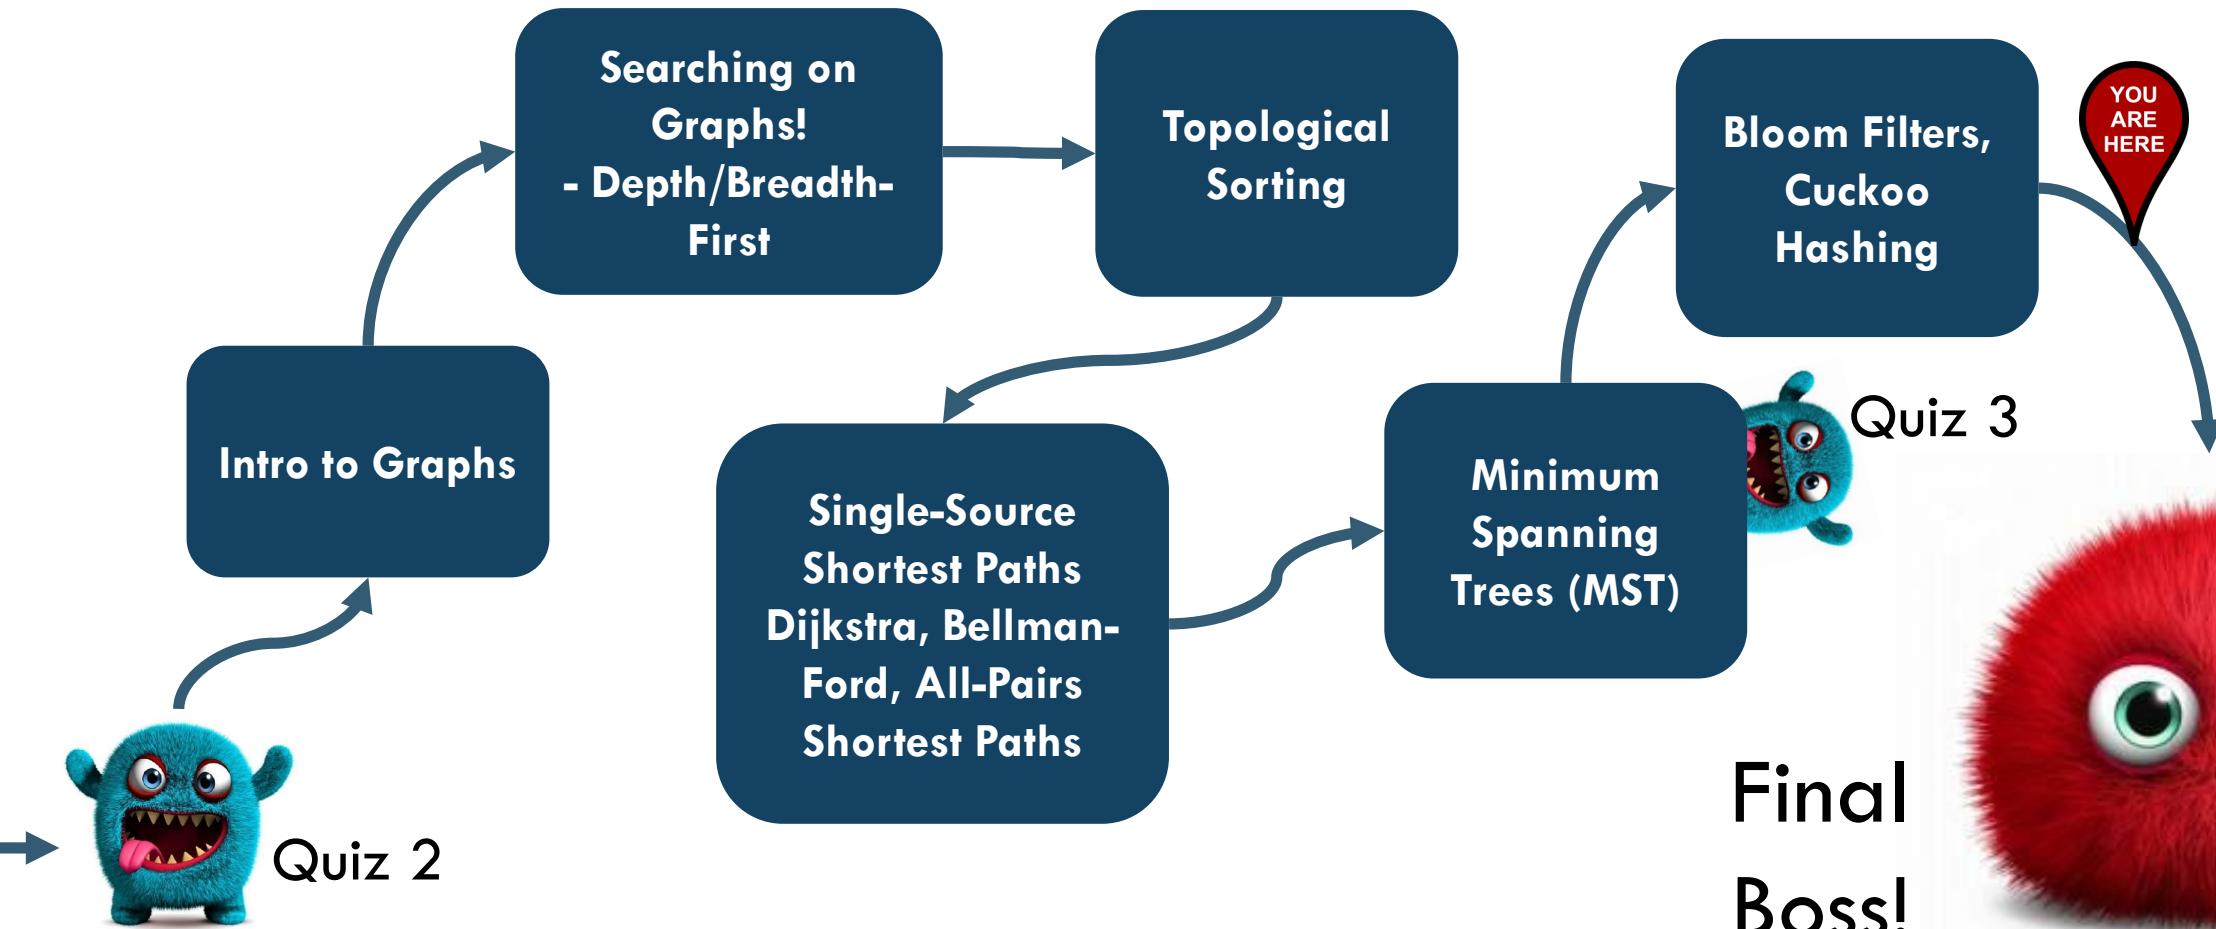
# QUESTION

When is the CS2040 exam?

- A. 5-Dec-2019 (9am)
- B. **4-Dec-2019 (9am)**
- C. 3-Dec-2019 (9am)
- D. 2-Dec 2019 (12pm)

# PATH TO MASTERY / COURSE STRUCTURE

DRAFT



# ~~QUIZ~~ FINAL ADVICE

Don't panic!

Go through Visualgo and the lecture notes.

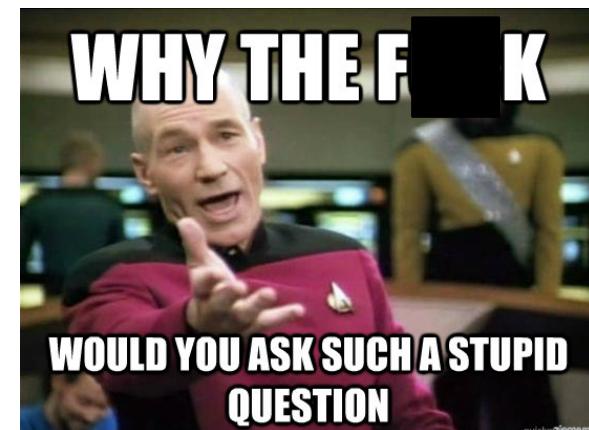
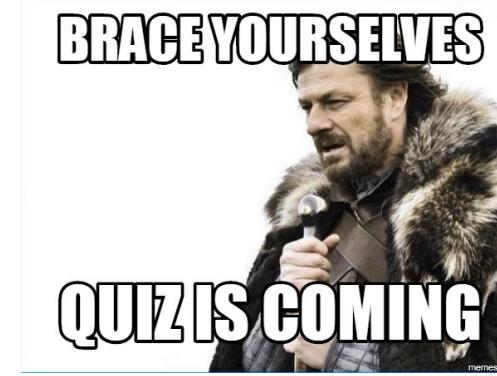
Sleep well the night before.

Sometimes, questions just look hard.

If you are unsure about something, raise your hand and ask.

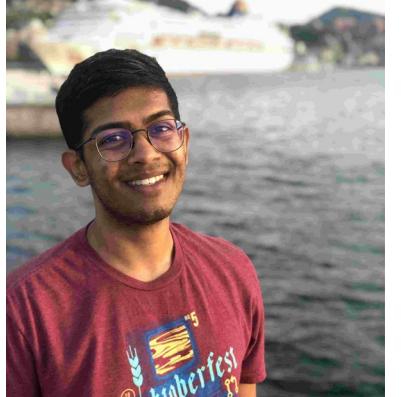
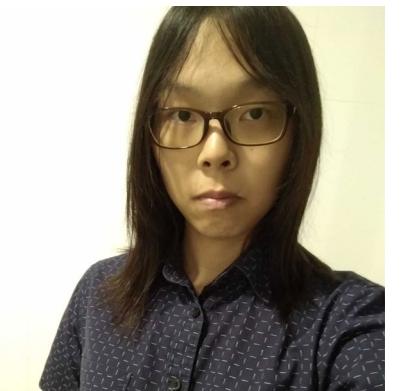
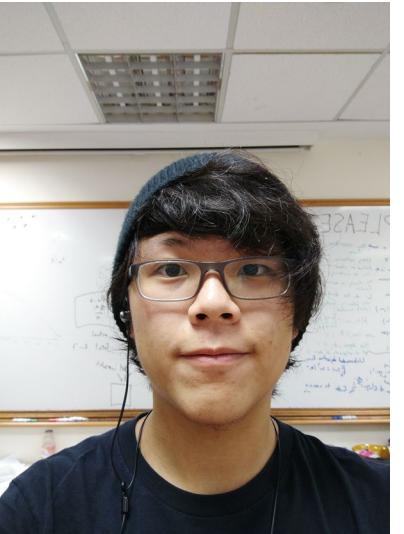
Many questions will require some on-the-spot thinking.

Pick the best answer.



**QUESTIONS?**





Thank you!!!

# On Piazza

## Class At A Glance

# BEHIND THE SCENES On Slack

### Active members

See how many people are active — meaning they viewed at least one public channel.

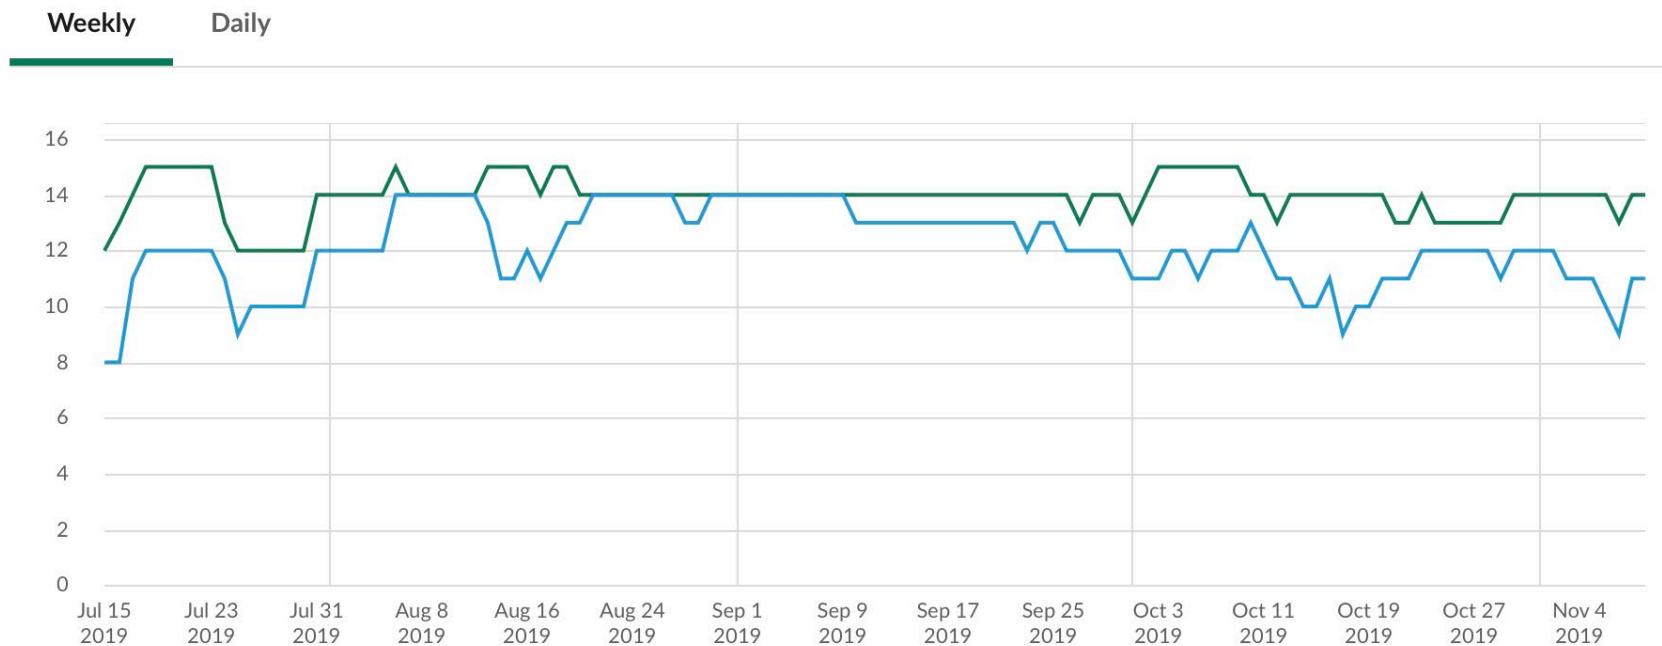
**119** total posts\*

**941** total contributions\*\*

**134** instructors' responses

**112** students' responses

**7 min** avg. response time



# TOP STUDENT CONTRIBUTORS & ANSWERERS!

## Top Student Contributors

57 contributions; 94 days online	<b>Christian James Welly</b>
37 contributions; 69 days online	<b>Nelson Tan</b>
33 contributions; 84 days online	<b>Chen Shenghao</b>
32 contributions; 77 days online	<b>Nathanael Seen Zhong Qi</b>
25 contributions; 86 days online	<b>Mario Lorenzo</b>



# THANK YOU!

Good luck for the finals!!!

Please provide feedback!

Please verify your scores on  
Luminus!

