

# **CS2102 Lecture 1**

## **Introduction**

# Database Management System (DBMS)

- What is a DBMS?
  - Software for managing large persistent data
- Advantages of using DBMS
  - Data Independence
  - Efficient Data Access
  - Data Integrity & Security
  - Data Administration
  - Transaction Management
    - Concurrent Access & Crash Recovery
  - Query language

# Traditional Data Processing: File Processing Techniques

```
initialize some book-keeping information I
open data file F
while (F is not empty)
    read next record r from F
    if (r satisfies some condition) then
        do something with r
    update I if necessary
do something with I if necessary
close file F
```

# Study of DBMS

- Database design
  - How to model the data requirements of applications
  - How to organize data using a DBMS
  - Topics: relational model, ER model, schema refinement
- Database programming
  - How to create, query, and update a database
  - How to specify data constraints
  - How to use SQL in applications
  - Topics: SQL, relational algebra/calculus
- DBMS implementation
  - How to build a DBMS (Covered in CS3223 & CS4224)

# Describing Data in a DBMS

- A DBMS allows users to define and query data in terms of a *data model*
- A **data model** is a collection of concepts for describing data
- A **schema** is a description of the structure of a database using a data model
- A **schema instance** is the content of the database at a particular time

# Data Models

- Network Model (e.g., General Electric's IDS (1964))
- Hierarchical Model (e.g., IBM's IMS (1966))
- **Relational Model**
  - **Commercial RDBMS**: IBM DB2, Microsoft SQL Server, Oracle, SAP ASE, etc.
  - **Open-source RDBMS**: MariaDB, MySQL, SQLite, etc.
- Object-oriented Model (e.g., ObjectStore 1988)
- Object-relational Model (e.g., Postgres 1986)
- etc.

# Relational Database Systems



(Image: Software Engineering Daily)

# Relational Data Model

- Introduced by **Edgar Codd** of IBM Research Laboratory in 1970
- Data is modeled using **relations** (tables with rows & columns)

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	4.0

**Degree/Arity** = Number of columns

**Cardinality** = Number of rows

- Each relation has a definition called a **relation schema**
  - Schema specifies **attributes** and **data constraints**
  - Data constraints include **domain constraints**

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

- Each row in a relation is called a **tuple/record**; it has one **component** for each attribute of relation

(1423, 'Bob', 2000-05-27, 4.3)



# Relational Data Model (cont.)

- **Domain** - a set of atomic values (e.g., integer, numeric, text)
- Let  $\text{domain}(A_i)$  denote the domain of an attribute  $A_i$  (set of possible values for  $A_i$ )
- Each value of attribute  $A_i$  is either a value in  $\text{domain}(A_i)$  or **null**
- **null** is a special value used to indicate that the value is either not applicable or unknown
- A relation is a **set of tuples**
  - Consider a relation schema  $R(A_1, A_2, \dots, A_n)$  with  $n$  attributes  $A_1, \dots, A_n$
  - Each **instance of schema**  $R$  is a relation which is a subset of  $\{(a_1, a_2, \dots, a_n) \mid a_i \in \text{domain}(A_i) \cup \{\text{null}\}\}$

# Relational Data Model (cont.)

- Consider the relation schema **Lectures(course, day, hour)**

- $\text{domain}(\text{course}) = \{\text{'cs101'}, \text{'cs203'}, \text{'cs305'}\}$
- $\text{domain}(\text{day}) = \{1, 2, 3, 4, 5\}$
- $\text{domain}(\text{hour}) = \{8, 10, 12, 14, 16\}$

- Each instance of Lectures is a subset of

$$\{\text{'cs101'}, \text{'cs203'}, \text{'cs305'}, \text{null}\} \times \{1, 2, 3, 4, 5, \text{null}\} \times \{8, 10, 12, 14, 16, \text{null}\}$$

<i>course</i>	<i>day</i>	<i>hour</i>
cs101	1	8
cs101	1	10
cs101	1	12
cs101	1	14
cs101	1	16
cs101	2	8
⋮	⋮	⋮
null	null	16
null	null	null

# Relational Data Model (cont.)

- A **relational database schema** consists of a set of relation schemas

Students    (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

Courses    (*courseId*: **integer**, *name*: **text**, *credits*: **integer**)

Enrols      (*sid*: **integer**, *cid*: **integer**, *grade*: **numeric**)

- A **relational database** is a collection of tables

Students				Courses		
<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>	<i>courseId</i>	<i>name</i>	<i>credits</i>
3118	Alice	1999-12-25	3.8	101	Programming in C	5
1423	Bob	2000-05-27	4.3	112	Discrete Mathematics	4
5609	Carol	1999-06-11	4.0	204	Analysis of Algorithms	4
				311	Database Systems	5

Enrols		
<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	204	3.0
1423	112	4.5
.....	.....	.....

- Relational database schema = relational schemas + data constraints

# Relation/Database Schema/Instance

- Relation schema

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

- Database schema

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

Courses (*courseId*: **integer**, *name*: **text**, *credits*: **integer**)

Enrols (*sid*: **integer**, *cid*: **integer**, *grade*: **numeric**)

- Relation (or relation instance)

- Database (or database instance)

Students			
<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	4.0

Courses		
<i>courseId</i>	<i>name</i>	<i>credits</i>
101	Programming in C	5
112	Discrete Mathematics	4
204	Analysis of Algorithms	4
311	Database Systems	5

Enrols		
<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	204	3.0
1423	112	4.5
.....	.....	.....

# Integrity Constraints (ICs)

- **Integrity constraint**: a condition that restricts the data that can be stored in database instance
  - Specified when schema is defined
  - ICs are checked when relations are updated
- A **legal relation instance** is a relation that satisfies all specified ICs.
- A DBMS enforces ICs - allows only legal instances to be stored

# Integrity Constraints (ICs) (cont.)

- Without any additional integrity constraints, each instance of  $R(A_1, \dots, A_n) \subseteq \{(a_1, a_2, \dots, a_n) \mid a_i \in \text{domain}(A_i) \cup \{\text{null}\}\}$

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	6.5
1423	Dave	2000-10-05	3.7

Courses

<i>courseId</i>	<i>name</i>	<i>credits</i>
101	Programming in C	5
112	Discrete Mathematics	4
204	Analysis of Algorithms	4
null	Compiler Design	4
311	Database Systems	5

Enrols

<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	202	3.0
1423	112	3.7
5609	101	4.5

# Types of Integrity Constraints

- Domain constraints restrict attribute values of relations
- Key constraints
- Foreign key constraints
- Other general constraints

# Key Constraints

- A **superkey** is a subset of attributes in a relation that uniquely identifies its tuples
  - No two distinct tuples of a relation have the same values in all attributes of superkey
- **Example:** Which of the following is a superkey for the relation **Students (studentId, name, birthDate, cap)?**
  - {studentId}
  - {name}
  - {birthDate}
  - {cap}
  - {studentId, name}
  - {studentId, birthDate}
  - {studentId, cap}
  - {name, birthDate}
  - {name, cap}
  - {birthDate, cap}
  - {studentId, name, birthDate}
  - {studentId, name, cap}
  - {studentId, birthDate, cap}
  - {name, birthDate, cap}
  - {studentId, name, birthDate, cap}



# Key Constraints (cont.)

- A **key** is a superkey that satisfies the additional property:
  - No *proper subset* of the key is a superkey
- Thus, a **key** is a minimal subset of attributes in a relation that uniquely identifies its tuples
- **Example:** Which of the following is a key for the relation Students (studentId, name, birthDate, cap)?
  - {studentId}
  - {name}
  - {birthDate}
  - {cap}
  - {studentId, name}
  - {studentId, birthDate}
  - {studentId, cap}
  - {name, birthDate}
  - {name, cap}
  - {birthDate, cap}
  - {studentId, name, birthDate}
  - {studentId, name, cap}
  - {studentId, birthDate, cap}
  - {name, birthDate, cap}
  - {studentId, name, birthDate, cap}

# Key Constraints (cont.)

- Key attribute values cannot be *null*
- A relation could have multiple keys called **candidate keys**
- One of the candidate keys is selected as the **primary key**
- **Example:**
  - Accounts (acctId, name, email, address)
  - Accounts has two candidate keys: {*acctId*} and {email}
  - Any one of them could be selected as the primary key

# Key Constraints (cont.)

- Consider the relation schema

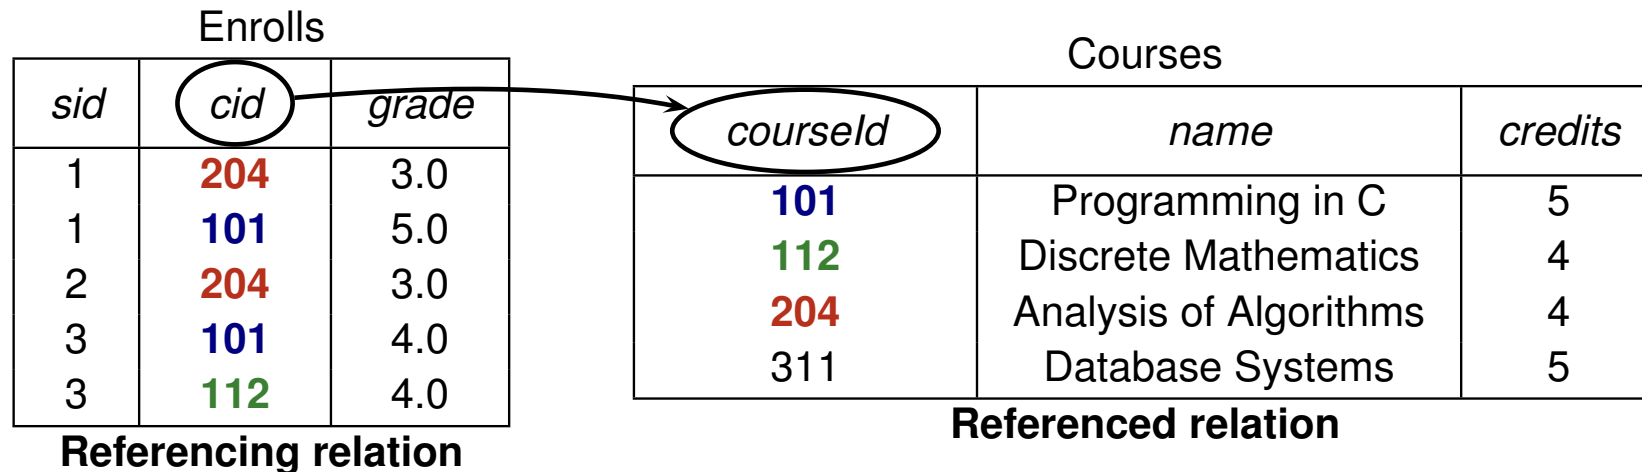
## **Lectures (cname, pname, day, hour)**

- cname** is a course taught by professor **pname** at time given by **day** & **hour**
- Lectures satisfies these constraints:
  - $\text{day} \in \{1, 2, 3, 4, 5\}$
  - $\text{hour} \in \{8, 10, 12, 14, 16\}$
  - At any time, each professor is teaching at most one course
  - Each course is taught by exactly one professor
  - Each course could have multiple lectures

<b>cname</b>	<b>pname</b>	<b>day</b>	<b>hour</b>
cs101	alice	1	10
cs101	alice	3	14
cs200	bob	2	8
ma300	bob	1	10

# Foreign Key Constraints

- A subset of attributes in a relation is a **foreign key** if it refers to the primary key of a second relation



- cid* is a foreign key in **Enrolls** that refers to the primary key *courseid* in **Courses**
- Foreign key constraint:** each **foreign key** value in **referencing relation** must either (1) appear as **primary key** value in **referenced relation** or (2) be a null value

# Foreign Key Constraints (cont.)

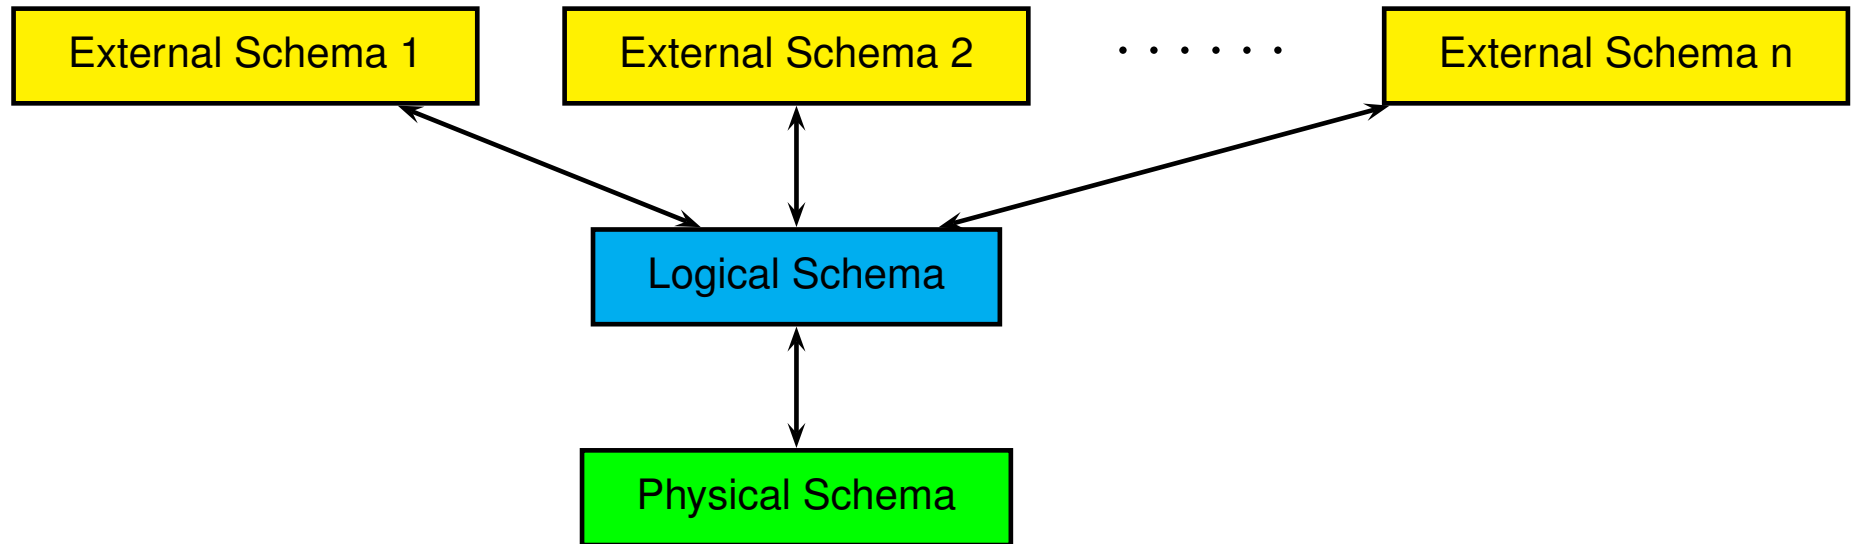
- The referencing & referenced relations could be the same relation
- **Example:** Each employee has at most one manager

Employees

<i>eid</i>	<i>ename</i>	<i>managerid</i>
001	Alice	<i>null</i>
002	Bob	001
003	Carol	001
007	Dave	<i>null</i>
008	Eve	007

- Constraints on Employees table:
  - *eid* is the primary key
  - *managerid* is a foreign key that refers to *eid*
- Foreign key constraints are also known as **referential integrity constraints**

# Levels of Data Abstraction



- Data in DBMS is described at three levels of abstractions
- **Logical Schema** - logical structure of data in DBMS
- **Physical Schema** - how the data described by logical schema is physically organized in DBMS
- **External Schema** - A customized view of logical schema for a group of users or an individual user

# External Schema Example

- Consider the following logical database schema:

Students    (*studentId, sname, birthDate, cap*)

Profs        (*profId, pname, email, office*)

Courses     (*courseId, cname, credits, profId, lectureTime*)

Enrols       (*sid, cid, grade*)

- External schema for Alice:

CourseInfo   (*cname, pname, lectureTime, totalEnrollment*)

- External schema for Bob:

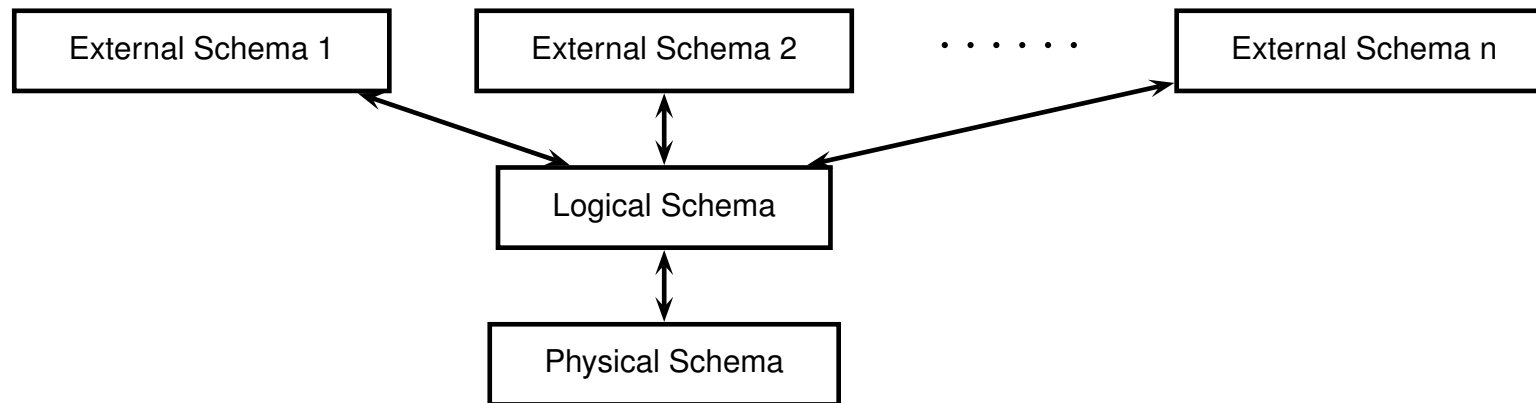
StudentInfo   (*studentId, sname*)

CourseInfo    (*courseId, cname, credits, profId, lectureTime*)

EnrollInfo     (*sid, cid, cname, pname, lectureTime*)

# Data Independence

- Insulate users/applications from changes in how data is structured and stored
- Data independence is achieved via the three levels of abstraction



- **Physical data independence** - protection from changes in physical schema
- **Logical data independence** - protection from changes in logical schema
- Data independence is an important advantage of using DBMS!



# Transactions

- Abstraction for representing a logical unit of work
- **ACID Properties**
  - **Atomicity**: Either all the effects of a transaction are reflected in the database or none are
  - **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
  - **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
  - **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

# Transaction Example

## **Transfer(X, Y, amount)**

```
fromBal := read balance from X's account
if fromBal ≥ amount then
    toBal := read balance from Y's account
    update Y's balance to toBal + amount
    update X's balance to fromBal - amount
end if
```

# Serial Transaction Executions

Two possible serial executions of `Transfer(1,2,100)` & `Transfer(2,1,100)`

(1): `fromBal := read 1's balance`  
`toBal := read 2's balance`  
`Update 2's balance to toBal + 100`  
`Update 1's balance to fromBal - 100`  
`fromBal := read 2's balance`  
`toBal := read 1's balance`  
`Update 1's balance to toBal + 100`  
`Update 2's balance to fromBal - 100`

(2): `fromBal := read 2's balance`  
`toBal := read 1's balance`  
`Update 1's balance to toBal + 100`  
`Update 2's balance to fromBal - 100`  
`fromBal := read 1's balance`  
`toBal := read 2's balance`  
`Update 2's balance to toBal + 100`  
`Update 1's balance to fromBal - 100`

# Concurrent Transaction Executions

A concurrent execution of  $\text{Transfer}(1,2,100)$  &  $\text{Transfer}(2,1,100)$

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal:
toBal:

**Accounts**

(1, \$400)
(2, \$2000)

**Transfer(2,1,100)**

fromBal:
toBal:

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400

toBal:

**Accounts**

(1, \$400)

(2, \$2000)

**Transfer(2,1,100)**

fromBal:

toBal:

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400

toBal: 2000

**Accounts**

(1, \$400)

(2, \$2000)

**Transfer(2,1,100)**

fromBal:

toBal:

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400

toBal: 2000

**Accounts**

(1, \$400)

(2, \$2100)

**Transfer(2,1,100)**

fromBal:

toBal:



# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal:
-------------------------

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 400
-----------------------------

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$300) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 400
-----------------------------

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$500) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 400
-----------------------------

# Example 1

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance  
toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100  
Update 2's balance to fromBal + 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$500) (2, \$2000)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 400
-----------------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal:
toBal:

**Accounts**

(1, \$400)
(2, \$2000)

**Transfer(2,1,100)**

fromBal:
toBal:

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400

toBal:

**Accounts**

(1, \$400)

(2, \$2000)

**Transfer(2,1,100)**

fromBal:

toBal:

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2000)
---------------------------

**Transfer(2,1,100)**

fromBal: toBal:
--------------------



# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: toBal:
--------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal:
-------------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$300) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal:
-------------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$300) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 300
-----------------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal - 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400) (2, \$2100)
---------------------------

**Transfer(2,1,100)**

fromBal: 2100 toBal: 300
-----------------------------

# Example 2

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

**Initial Database State:** (1, \$400), (2, \$2000)

fromBal := read 1's balance  
toBal := read 2's balance  
Update 2's balance to toBal + 100

fromBal := read 2's balance

Update 1's balance to fromBal - 100

toBal := read 1's balance  
Update 1's balance to toBal + 100  
Update 2's balance to fromBal + 100

**Transfer(1,2,100)**

fromBal: 400 toBal: 2000
-----------------------------

**Accounts**

(1, \$400)
(2, \$2000)

**Transfer(2,1,100)**

fromBal: 2100 toBal: 300
-----------------------------

# Serializable Transaction Executions

- A concurrent execution of a set of transactions is **serializable** if this execution is equivalent to some serial execution of the same set of transactions
- A concurrent execution is **equivalent** to a serial execution (of the same set of transactions) if
  - both executions produce the same final database state, and
  - every read operation in one execution returns the same value as the corresponding read operation in the other execution
- A serial execution is trivially serializable

# Revisiting Examples

## Serial Execution 1

```
fromBal := read 1's balance
toBal := read 2's balance
Update 2's balance to toBal + 100
Update 1's balance to fromBal - 100
fromBal := read 2's balance
toBal := read 1's balance
Update 1's balance to toBal + 100
Update 2's balance to fromBal - 100
```

## Serial Execution 2

```
fromBal := read 2's balance
toBal := read 1's balance
Update 1's balance to toBal + 100
Update 2's balance to fromBal - 100
fromBal := read 1's balance
toBal := read 2's balance
Update 2's balance to toBal + 100
Update 1's balance to fromBal - 100
```

## Concurrent Execution 1

```
fromBal := read 1's balance
toBal := read 2's balance
Update 2's balance to toBal + 100
fromBal := read 2's balance
toBal := read 1's balance
Update 1's balance to fromBal - 100
Update 1's balance to toBal + 100
Update 2's balance to fromBal - 100
```

## Concurrent Execution 2

```
fromBal := read 1's balance
toBal := read 2's balance
Update 2's balance to toBal + 100
fromBal := read 2's balance
Update 1's balance to fromBal - 100
toBal := read 1's balance
Update 1's balance to toBal + 100
Update 2's balance to fromBal - 100
```



# Summary

- DBMS used to store, update, and query data
- **Relational data model**
  - Tabular representation of data
  - Integrity constraints specify restrictions on data based on application semantics
- Levels of data abstraction provide data independence
- Transactions simplify application development