

CS2102 Lecture 5

SQL (Part 2)

Example Database

Pizzas

pizza
Diavola
Funghi
Hawaiian
Margherita
Marinara
Siciliana

Customers

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

Restaurants

rname	area
Corleone Corner	North
Gambino Oven	Central
Lorenzo Tavern	Central
Mamma's Place	South
Pizza King	East

Contains

pizza	ingredient
Diavola	cheese
Diavola	chilli
Diavola	salami
Funghi	ham
Funghi	mushroom
Hawaiian	ham
Hawaiian	pineapple
Margherita	cheese
Margherita	tomato
Marinara	seafood
Siciliana	anchovies
Siciliana	capers
Siciliana	cheese

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Likes

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Sciliana
Ralph	Diavola

Simple Queries

- Basic form of SQL query consists of three clauses:

select [**distinct**] select-list
from from-list
[**where** qualification]

- **from-list**: Specifies list of relations
- **qualification**: Specifies conditions on relations
- **select-list**: Specifies columns to be included in output table
- Output relation could contain duplicate records if **distinct** is not used in the select clause

Simple Queries (cont.)

select distinct a_1, a_2, \dots, a_m
from r_1, r_2, \dots, r_n
where c

$$\pi_{a_1, a_2, \dots, a_m} \left(\sigma_c \left(r_1 \times r_2 \times \dots \times r_n \right) \right)$$

Simple Queries (cont.)

Find the names of restaurants, the pizzas that they sell and their prices, where the price is under \$20

```
select rname, pizza, price
from Sells
where price < 20;
```

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	pizza	price
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Pizza King	Diavola	17

Simple Queries (cont.)

Find the names of restaurants, the pizzas that they sell and their prices, where the price is under \$20

```
select rname, pizza, price  
from Sells  
where price < 20;
```

```
select *  
from Sells  
where price < 20;
```

Simple Queries (cont.)

Find all tuples from `Sells` relation such that (1) the price is under \$20 and the restaurant name is not “Pizza King” or (2) the restaurant name is “Corleone Corner”.

```
select * from Sells
where ((price < 20) and (rname <> 'Pizza King'))
or      rname = 'Corleone Corner';
```

`Sells`

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16

Removing Duplicate Records

Q1: **select** A, C **from** R;

Q2: **select** **distinct** A, C **from** R;

R			Q1		Q2	
A	B	C	A	C	A	C
10	1	2	10	2	10	2
10	7	2	10	2	20	null
20	3	null	20	null	30	2
20	9	null	20	null	30	9
30	3	2	30	2		
30	5	9	30	9		

Two tuples (a_1, c_1) and (a_2, c_2) are considered to be distinct if the following evaluates to *true*:

“(a_1 IS DISTINCT FROM a_2) or (c_1 IS DISTINCT FROM c_2)”

Expressions in SELECT Clause

select item **as** product, price * qty **as** cost
from Orders;

Orders

item	price	qty
A	2.50	100
B	4.00	100
C	7.50	100

product	cost
A	250.00
B	400.00
C	750.00

Expressions in SELECT Clause (cont.)

select 'Price of ' || pizza || ' is ' || **round**(price / 1.3) || ' USD' **as** menu
from Sells
where rname = 'Corleone Corner';

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

menu
Price of Diavola is 11 USD
Price of Hawaiian is 20 USD
Price of Margherita is 15 USD

|| is the string concatenation operator

round() = function that rounds off to nearest integer value

Set Operations

- Let Q_1 & Q_2 denote SQL queries that output union-compatible relations
- Q_1 **union** $Q_2 = Q_1 \cup Q_2$
- Q_1 **intersect** $Q_2 = Q_1 \cap Q_2$
- Q_1 **except** $Q_2 = Q_1 - Q_2$
- **union**, **intersect**, and **except** eliminate duplicate records
- **union all**, **intersect all**, and **except all** preserves duplicate records

Set Operations (cont.)

Example 1: Find all customer/restaurant names

select cname **from** Customers

union

select rname **from** Restaurants;

Example 2: Find pizzas that contain both cheese and chilli

select pizza **from** Contains **where** ingredient = 'cheese'

intersect

select pizza **from** Contains **where** ingredient = 'chilli';

Example 3: Find pizzas that contain cheese but not chilli

select pizza **from** Contains **where** ingredient = 'cheese'

except

select pizza **from** Contains **where** ingredient = 'chilli';

Set Operations (cont.)

Q1: select B from R **except select B from S;**

Q2: select B from R **except all select B from S;**

R		S		Q1	Q2
A	B	A	B	B	B
10	1	10	1	4	1
20	1	20	5		1
30	1	30	2		4
40	2	40	2		4
50	3	50	3		
60	4				
70	4				

Multi-relation Queries: Example 1

Find distinct pairs of customers and restaurants that are located in the same area

```
select  cname, rname
from    Customers, Restaurants
where   Customers.area = Restaurants.area;
```

Multi-relation Queries: Example 1 (cont.)

Find distinct pairs of customers and restaurants that are located in the same area

```
select  cname, rname
from    Customers as C, Restaurants as R
where    C.area = R.area;
```

```
select  cname, rname
from    Customers C, Restaurants R
where    C.area = R.area;
```

Multi-relation Queries: Example 1 (cont.)

Find distinct pairs of customers and restaurants that are located in the same area

```
select  cname, rname
from    Customers C inner join Restaurants R
on C.area = R.area;
```

```
select  cname, rname
from    Customers C join Restaurants R
on C.area = R.area;
```


Multi-relation Queries: Example 1 (cont.)

Find distinct pairs of customers and restaurants that are located in the same area

```
select  cname, rname  
from    Customers natural join Restaurants;
```

Multi-relation Queries: Example 2

Find distinct restaurant pairs (R1,R2) where $R1 < R2$ and they sell some common pizza

```
select distinct S1.rname, S2.rname  
from Sells S1, Sells S2  
where S1.rname < S2.rname  
and S1.pizza = S2.pizza;
```

```
select distinct S1.rname, S2.rname  
from Sells S1 inner join Sells S2  
  on (S1.rname < S2.rname)  
  and (S1.pizza = S2.pizza);
```

Multi-relation Queries: Example 3

Find customers and the pizzas they like; include also customers who don't like any pizza

```
select  C.cname, L.pizza
from    Customers C left outer join Likes L
on C.cname = L.cname;
```

```
select  C.cname, L.pizza
from    Customers C left join Likes L
on C.cname = L.cname;
```

Multi-relation Queries: Example 3 (cont.)

Find customers and the pizzas they like; include also customers who don't like any pizza

```
select  cname, pizza
from    Customers natural left outer join Likes;
```

```
select  cname, pizza
from    Customers natural left join Likes;
```

Multi-relation Queries: Example 4

Find all customer-pizza pairs (C,P) where the pizza is sold by some restaurant that is located in the same area as that of the customer. Include customers whose associated set of pizzas is empty

```
select  C.cname, S.pizza
from    Customers C left join
          (Restaurants R join Sells S on R.rname = S.rname)
on C.area = R.area;
```



```
select  C.cname, RS.pizza
from    Customers C left join
          (Restaurants R join Sells S on R.rname = S.rname) RS
on C.area = RS.area;
```

Subquery Expressions

- EXISTS subqueries
- IN subqueries
- ANY/SOME subqueries
- ALL subqueries

EXISTS Subqueries

- EXISTS (subquery)
- Returns *true* if the result subquery is non-empty; otherwise, *false*

EXISTS Subqueries (cont.)

Find distinct customers who like some pizza sold by “Corleone Corner”

Likes

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Sciliana
Ralph	Diavola

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma’s Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

cname
Homer
Ralph

EXISTS Subqueries (cont.)

Find distinct customers who like some pizza sold by
“Corleone Corner”

```
select  distinct cname
from    Likes L
where   exists (
        select  1
        from    Sells S
        where   S.rname = 'Corleone Corner'
        and     S.pizza = L.pizza
        );
```

EXISTS Subqueries (cont.)

Find distinct customers who like some pizza sold by “Corleone Corner”

```
select distinct cname
from Likes L
where exists (
    select 1
    from Sells S
    where S.rname = 'Corleone Corner'
    and S.pizza = L.pizza
);
```

Likes

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Sciliana
Ralph	Diavola

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

NOT EXISTS Subqueries

Find distinct customers who does not like any pizza sold by “Corleone Corner”

Customers

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

Likes

cname	pizza
Homer	Hawaiian
Homer	Margherita
Lisa	Funghi
Maggie	Funghi
Moe	Funghi
Moe	Sciliana
Ralph	Diavola

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

cname
Lisa
Maggie
Moe
Willie

NOT EXISTS Subqueries (cont.)

Find distinct customers who does not like any pizza sold by “Corleone Corner”

```
select  cname
from    Customers C
where   not exists (
          select  1
          from    Likes L, Sells S
          where   L.pizza = S.pizza
          and     S.rname = 'Corleone Corner'
          and     L.cname = C.cname
          );
```

IN Subqueries

- **expression IN (subquery)**
- Subquery must return exactly one column
- Returns *false* if result of subquery is empty; otherwise return the result of the boolean expression

$$((v = v_1) \text{ or } (v = v_2) \text{ or } \cdots \text{ or } (v = v_n))$$

where

- v denote the result of expression
- $\{v_1, v_2, \cdots, v_n\}$ denote the result of subquery

IN Subqueries (cont.)

Find distinct customers who like some pizza sold by “Corleone Corner”

```
select  distinct cname
from    Likes
where   pizza in (
           select  pizza
           from    Sells
           where   rname = 'Corleone Corner'
        );
```

Another Form of IN Predicate

- **expression IN (value1, value2, ..., valuen)**

Example: Find pizzas that contain ham or seafood

```
select distinct pizza from Contains  
where ingredient in ('ham', 'seafood');
```

```
select distinct pizza from Contains  
where ingredient = 'ham' or ingredient = 'seafood';
```

```
select pizza from Contains where ingredient = 'ham'  
union  
select pizza from Contains where ingredient = 'seafood';
```

ANY/SOME Subqueries

- expression operator **ANY** (subquery)
- Subquery must return exactly one column
- Returns *false* if result of subquery is empty; otherwise return the result of the boolean expression

$$((v \text{ op } v_1) \text{ or } (v \text{ op } v_2) \text{ or } \dots \text{ or } (v \text{ op } v_n))$$

where

- v denote the result of expression
- $\{v_1, v_2, \dots, v_n\}$ denote the result of subquery
- op denote operator

ANY/SOME Subqueries (cont.)

Find distinct restaurants that sell some pizza P1 that is more expensive than some pizza P2 sold by “Corleone Corner”. P1 and P2 are not necessarily the same pizza. Exclude “Corleone Corner” from the query result.

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname
Lorenzo Tavern
Mamma's Place
Pizza King

ANY/SOME Subqueries (cont.)

Find distinct restaurants that sell some pizza P1 that is more expensive than some pizza P2 sold by “Corleone Corner”. P1 and P2 are not necessarily the same pizza. Exclude “Corleone Corner” from the query result.

```
select  distinct rname
from    Sells
where   rname <> 'Corleone Corner'
and     price > any (
           select  price
           from    Sells
           where   rname = 'Corleone Corner'
           );
```

ALL Subqueries

- expression operator **ALL** (subquery)
- Subquery must return exactly one column
- Returns *true* if result of subquery is empty; otherwise return

$((v \text{ op } v_1) \text{ and } (v \text{ op } v_2) \text{ and } \dots \text{ and } (v \text{ op } v_n))$

where

- v denote the result of expression
- $\{v_1, v_2, \dots, v_n\}$ denote the result of subquery
- op denote operator

ALL Subqueries (cont.)

For each restaurant, find the name and price of its most expensive pizzas. Exclude restaurants that do not sell any pizza. Assume that all prices are non-null values.

Sells

rname	pizza	price
Corleone Corner	Diavola	25
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	pizza	price
Corleone Corner	Diavola	25
Corleone Corner	Hawaiian	25
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Hawaiian	21

ALL Subqueries (cont.)

For each restaurant, find the name and price of its most expensive pizzas. Exclude restaurants that do not sell any pizza. Assume that all prices are non-null values.

```
select  rname, pizza, price
from    Sells S1
where   price >= all (
           select  S2.price
           from    Sells S2
           where   S2.rname = S1.rname
        );
```

Subqueries & Row Constructors

- So far, **IN/ANY/ALL** subqueries must return exactly one column
 - expression IN (subquery)
 - expression operator ANY (subquery)
 - expression operator ALL (subquery)
- Possible to use subqueries that return more than one column
 - rowConstructor IN (subquery)
 - rowConstructor operator ANY (subquery)
 - rowConstructor operator ALL (subquery)

Row Constructors

Find all courses with lectures scheduled after Wednesday 11am.

Lectures

cname	pname	day	hour
CS101	Alice	1	10
CS123	Alice	3	10
CS200	Bob	4	8
MA300	Bob	3	14

cname
CS200
MA300

```
select  cname
from    Lectures
where   day > 3
or      ((day = 3) and (hour > 11));
```

Row Constructors (cont.)

Find all courses with lectures scheduled after Wednesday 11am.

Lectures

cname	pname	day	hour
CS101	Alice	1	10
CS123	Alice	3	10
CS200	Bob	4	8
MA300	Bob	3	14

cname
CS200
MA300

```
select  cname
from    Lectures
where   row(day, hour) > row(3, 11);
```


Subqueries & Row Constructors

For each lecturer, find the time of his/her earliest lecture during the week.

Lectures

cname	pname	day	hour
CS101	Alice	1	10
CS123	Alice	3	10
CS200	Bob	4	8
MA300	Bob	3	14

pname	day	hour
Alice	1	10
Bob	3	14

```
select  pname, day, hour
from    Lectures L
where    (day, hour) <= all (
            select day, hour
            from    Lectures L2
            where L2.pname = L.pname
        );
```

Scalar Subqueries

- A **scalar subquery** is a subquery that returns at most one tuple with one column
 - If the subquery's result is empty, its return value is `null`
- A scalar subquery can be used as a scalar expression

Scalar Subqueries (cont.)

For each restaurant that sells Funghi, find its name, area, and selling price.

```
select  S.rname, R.area, S.price
from    Sells S, Restaurants R
where   S.rname = R.rname
and     S.pizza = 'Funghi';
```

```
select  rname,
         (select R.area from Restaurants R
         where R.rname = S.rname), price
from    Sells S
where   pizza = 'Funghi';
```

Subquery Expressions: Scoping Rules

- Queries with subquery expressions are also called **nested queries**
- A subquery expression is referred to as an **inner query** that is nested within an **outer query**
- **Scoping rules for table alias (a.k.a. tuple variable):**
 - A tuple variable declared in a subquery/query Q can be used only in Q and any subquery nested within Q
 - If a tuple variable is declared both locally in a subquery Q as well as in an outer query, the local declaration applies in Q

Correlated Nested Queries

- A nested query with a subquery that references a tuple variable declared in an outer query is called a **correlated nested query**
- Example of **correlated nested query**

```
select  distinct cname from Likes L where exists (  
    select    1 from Sells S  
    where    (S.rname = 'Corleone Corner') and (S.pizza = L.pizza));
```

- Example of **non-correlated nested query**

```
select  distinct cname from Likes  
where   pizza in (select pizza from Sells where rname = 'Corleone Corner');
```

Usage of Subqueries

- Non-scalar subquery expressions can be used in different parts of SQL queries:
 - WHERE clause
 - FROM clause
 - HAVING clause (to be discussed in the next lecture)

Example: Subqueries in From Clause

```
select      *  
from      (  
            values (('Alice', 24), ('Bob', 22), ('Carol', 23))  
            ) as persons(name, age);
```

name	age
Alice	24
Bob	22
Carol	23

- Subqueries in from clause must be enclosed in parentheses & assigned a table alias. Column aliases are optional
- More examples of subqueries in from clause will be illustrated in the next lecture

Database Modifications with Subqueries

```
create table Students (  
    studentId    integer,  
    name         varchar(100),  
    birthDate    date,  
    year         integer,  
    primary key (studentId));
```

```
create table Enrolls (  
    sid    integer  
           references Students,  
    cid    integer  
           references Courses,  
    grade  char(2),  
    primary key (sid, cid));
```

```
-- Enroll all first-year students in the course 101  
insert into Enrolls (sid, cid)  
    select studentId, 101  
    from    Students  
    where year = 1;
```


ORDER BY Clause

For each restaurant that sells some pizza, find its name, area, and the pizzas it sells together with their prices. Show the output in ascending order of the area, followed by in descending order of the price.

Restaurants

rname	area
Corleone Corner	North
Gambino Oven	Central
Lorenzo Tavern	Central
Mamma's Place	South
Pizza King	East

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	area	pizza	price
Lorenzo Tavern	Central	Funghi	23
Gambino Oven	Central	Siciliana	16
Pizza King	East	Hawaiian	19
Pizza King	East	Diavola	17
Corleone Corner	North	Hawaiian	25
Corleone Corner	North	Diavola	24
Corleone Corner	North	Margherita	19
Mamma's Place	South	Marinara	22

ORDER BY Clause (cont.)

For each restaurant that sells some pizza, find its name, area, and the pizzas it sells together with their prices. Show the output in ascending order of the area, followed by in descending order of the price.

```
select    *  
from      Restaurants, Sells  
where     Restaurants.rname = Sells.rname  
order by  area asc, price desc;
```

```
select    *  
from      Restaurants, Sells  
where     Restaurants.rname = Sells.rname  
order by  area, price desc;
```

LIMIT Clause

Find the top three most expensive pizzas. Show the pizza name, the name of the restaurant that sells it, and its selling price for each output tuple; and show the output in descending order of price.

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

pizza	rname	price
Hawaiian	Corleone Corner	25
Diavola	Corleone Corner	24
Funghi	Lorenzo Tavern	23

```
select    pizza, rname, price
from      Sells
order by  price desc
limit     3;
```

OFFSET Clause

For each pizza that is sold by some restaurant, find the pizza name, the restaurant name and its selling price; show the output in descending order of price and exclude the top three pizzas.

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

pizza	rname	price
Mamma's Place	Marinara	22
Pizza King	Hawaiian	21
Corleone Corner	Margherita	19
Pizza King	Diavola	17
Gambino Oven	Siciliana	16

```
select    pizza, rname, price
from      Sells
order by  price desc
offset    3;
```

OFFSET Clause

Find the 4th and 5th most expensive pizzas. Show the pizza name, the name of the restaurant that sells it, and its selling price for each output tuple; and show the output in descending order of price.

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

pizza	rname	price
Mamma's Place	Marinara	22
Pizza King	Hawaiian	21

```
select    pizza, rname, price
from      Sells
order by  price desc
offset    3
limit     2;
```

SQL:2008 Syntax for LIMIT & OFFSET

select	pizza, rname, price
from	Sells
order by	price desc
offset	3
limit	2;

select	pizza, rname, price
from	Sells
order by	price desc
offset	3 rows
fetch next	2 rows only ;

Summary

- Basic SQL queries

select	distinct select-list
from	from-list
where	where-condition
order by	orderby-list
offset	offset-specification
limit	limit-specification

- Non-scalar subqueries can be used in FROM, WHERE, and HAVING clauses
- SQL Reference: <https://www.postgresql.org/docs/current/index.html>