
CS2102 Database Systems

Intro

- Xiaokui Xiao (Cedric)
- xkxiao@nus.edu.sg
- COM1, #03-25

Announcements (March 12)

- Submission of revised answers for midterm test's questions 8, 9 & 10
 - Please refer to LumiNUS announcement on submission of revised answers to fix minor typo errors
 - The deadline for this is **today (March 12, 11pm)**
 - Assignment 2 on SQL is due on March 19 (Friday, 6pm)
 - Late submission penalty: One mark will be deducted for each late day up to two late days; submissions after the second late day will receive zero marks and will not be graded.
-

Programming with SQL


- The previous lectures have discussed how to use SQL to access a database
 - But what if we are asked to write a C/Java/... program that accesses a database?
 - We have several options:
 - Mixing C/Java/... and SQL
 - Write a C/Java/... program that calls some API to run SQL queries
-

Mixing a host language with SQL

- Basic idea:

- Write a program that mixes a host language with SQL
- Pass it to a preprocessor, which returns a program in the host language
- Compile the program to executable code

- In PostgreSQL:



Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

Example


```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part declares the **shared** variables that can be used by both the host language statements and SQL statements

Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This specifies the database to connect and the username to use



Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part is in the
host language


Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part inserts a record into the Sells table.

Note that the each shared variable must be preceded by a colon

Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;   
    return 0;  
}
```

This closes the
database
connection.

Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char supplName[30], prodName[20];  
        float price;  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that assigns values to  
    // supplName, prodName, and price;  
  
    EXEC SQL INSERT INTO  
        Sells(supplName, prodName, price)  
        VALUES(:supplName, :prodName, :price);  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

In this example, the SQL statement is "fixed" beforehand.

This is referred to as **static SQL**.

The alternative is **dynamic SQL**.

Dynamic SQL

- Basic idea:
 - The SQL statement to be executed is not fixed in advance
 - Instead, the statement is stored in a string variable, and it is compiled and executed at runtime
 - i.e., syntax checking happens at runtime
-

Dynamic SQL Example

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char stmt[500];  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    // some code that writes a SQL statement  
    // in stmt;  
  
    EXEC SQL EXECUTE IMMEDIATE :stmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part executes the SQL statement stored in **stmt**.

Example:
"SELECT * FROM
Sells;"

Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```


Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```



This part initializes a SQL statement with two parameters

Dynamic SQL Example 2


```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
    EXEC SQL PREPARE mystmt FROM :stmt;  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

← This part connects to the database

Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```


This part parses the SQL statement, and gives it a name, mystmt




Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part executes
the SQL statement,
and using 42 and
'foobar' as
parameters



Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;   
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

This part releases the resources allocated to mystmt

Every prepared statement should be deallocated when it is no longer needed

Dynamic SQL Example 2

```
int main() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        const char *stmt = "INSERT INTO test VALUES(?, ?);";  
    EXEC SQL END DECLARE SECTION;  
  
    EXEC SQL CONNECT TO @localhost USER john;  
  
    EXEC SQL PREPARE mystmt FROM :stmt;  
  
    EXEC SQL EXECUTE mystmt USING 42, 'foobar';  
  
    EXEC SQL DEALLOCATE PREPARE mystmt;  
  
    EXEC SQL DISCONNECT;  
    return 0;  
}
```

Static and Dynamic SQL: Summary

- Static and dynamic SQL allow us to execute SQL statements in a host language
 - Using EXEC SQL ...
 - They are referred to as **statement-level interfaces**
 - What if we don't want to mix the host language with SQL?
 - We can use a **call-level interface** (CLI)
 - i.e., instead of using EXEC SQL, we call some API to run SQL queries
 - Everything is in the host language
-

Call-Level Interface (CLI)

- Examples:

- Java Database Connectivity (JDBC)

- <https://jdbc.postgresql.org/>

- Open Database Connectivity (ODBC)

- <https://odbc.postgresql.org/>

ODBC Example (in C++)

```
void main() {
    char * sql;
    connection C("dbname = testdb user = postgres \
        password = cohondob hostaddr = 127.0.0.1 port = 5432");

    // Some code here

    // Create SQL statement
    sql = "CREATE TABLE COMPANY(" \
        "ID INT PRIMARY KEY   NOT NULL," \
        "NAME          TEXT   NOT NULL)";

    work W(C);    // Create a transactional object

    W.exec( sql ); // Execute SQL query
    W.commit();
    C.disconnect ();
}
```

ODBC Example (in C++)

```
void main() {
    char * sql;
    connection C("dbname = testdb user = postgres \
        password = cohondob hostaddr = 127.0.0.1 port = 5432");


    // Some code here

    // Create SQL statement
    sql = "CREATE TABLE COMPANY(" \
        "ID INT PRIMARY KEY   NOT NULL," \
        "NAME          TEXT   NOT NULL)";

    work W(C);    // Create a transactional object

    W.exec( sql ); // Execute SQL query
    W.commit();
    C.disconnect ();
}
```


This part creates a connection to the database



ODBC Example (in C++)

```
void main() {  
    char * sql;  
    connection C("dbname = testdb user = postgres \  
        password = cohondob hostaddr = 127.0.0.1 port = 5432");  
  
    // Some code here  
  
    // Create SQL statement  
    sql = "CREATE TABLE COMPANY(" \  
        "ID INT PRIMARY KEY   NOT NULL," \  
        "NAME          TEXT   NOT NULL);" ;  
  
    work W(C);    // Create a transactional object  
  
    W.exec( sql ); // Execute SQL query  
    W.commit();  
    C.disconnect ();  
}
```

This part assigns
the SQL statement
as a string to the sql
variable



ODBC Example (in C++)


```
void main() {  
    char * sql;  
    connection C("dbname = testdb user = postgres \  
        password = cohondob hostaddr = 127.0.0.1 port = 5432");  
  
    // Some code here  
  
    // Create SQL statement  
    sql = "CREATE TABLE COMPANY(" \  
        "ID INT PRIMARY KEY   NOT NULL," \  
        "NAME          TEXT   NOT NULL);";  
  
    work W(C);    // Create a transactional object  
  
    W.exec( sql ); // Execute SQL query  
    W.commit();  
    C.disconnect ();  
}
```

This part creates an
objective W for
executing the SQL
statement

ODBC Example (in C++)

```
void main() {  
    char * sql;  
    connection C("dbname = testdb user = postgres \  
        password = cohondob hostaddr = 127.0.0.1 port = 5432");  
  
    // Some code here  
  
    // Create SQL statement  
    sql = "CREATE TABLE COMPANY(" \  
        "ID INT PRIMARY KEY   NOT NULL," \  
        "NAME          TEXT   NOT NULL);";  
  
    work W(C);    // Create a transactional object  
  
    W.exec( sql ); // Execute SQL query  
    W.commit();  
    C.disconnect ();  
}
```


This part executes
the SQL statement
recorded in sql



ODBC Example (in C++)


```
void main() {  
    char * sql;  
    connection C("dbname = testdb user = postgres \  
        password = cohondob hostaddr = 127.0.0.1 port = 5432");  
  
    // Some code here  
  
    // Create SQL statement  
    sql = "CREATE TABLE COMPANY(" \  
        "ID INT PRIMARY KEY   NOT NULL," \  
        "NAME          TEXT   NOT NULL);";  
  
    work W(C);    // Create a transactional object  
  
    W.exec( sql ); // Execute SQL query  
    W.commit();  
    C.disconnect ();  
}
```

This part closes the
database
connection



Coming Next

- Statement-level interface
 - Code in a mix of a host language and SQL
- Call-level interface
 - Code in the host language only
- What if we want to code in SQL only?
 - "Why would you do that?"
 - To create functions/procedures in the database



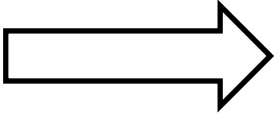
Functions/procedures in SQL

- Advantages:
 - Code reuse
 - Ease of maintenance
 - Performance
 - Security
-

SQL-based Procedure Language

- ISO standard: SQL/PSM (Persistent Stored Modules)
 - Different vendors' implementations:
 - Oracle: PL/SQL
 - PostgreSQL: PL/pgSQL
 - SQL Server: TransactSQL
 - ...
 - Our discussions will be based on PL/pgSQL
-

PL/pgSQL Example

Scores	Name	Mark			Name	Grade
	Alice	92			Alice	A
	Bob	63			Bob	B
	Cathy	58			Cathy	C
	David	47			David	D

- Suppose that we want to create a function that converts numeric marks to letter grades as follows:
 - ❑ [70, 100] -> A
 - ❑ [60, 70) -> B
 - ❑ [50, 60) -> C
 - ❑ [0, 50) -> D

PL/pgSQL Example

Scores

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION convert (mark INT)
RETURNS char(1) AS $$
    SELECT CASE
        WHEN mark >= 70 THEN 'A'
        WHEN mark >= 60 THEN 'B'
        WHEN mark >= 50 THEN 'C'
        ELSE 'D'
    END;
$$ LANGUAGE sql;
```

PL/pgSQL Example

Scores

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47

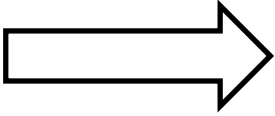
```
CREATE OR REPLACE FUNCTION convert (mark INT)
RETURNS char(1) AS $$
```

```
    SELECT CASE
        WHEN mark >= 70 THEN 'A'
        WHEN mark >= 60 THEN 'B'
        WHEN mark >= 50 THEN 'C'
        ELSE 'D'
    END;
```

```
$$ LANGUAGE sql;
```

- Query: `SELECT convert(90);`
- Result: A
- Query: `SELECT convert(40);`
- Result: D


PL/pgSQL Example

Scores	<u>Name</u>	Mark			<u>Name</u>	Grade
	Alice	92			Alice	A
	Bob	63			Bob	B
	Cathy	58			Cathy	C
	David	47			David	D

```
CREATE OR REPLACE FUNCTION convert (mark INT)
RETURNS char(1) AS $$
```

```
    SELECT CASE
        WHEN mark >= 70 THEN 'A'
        WHEN mark >= 60 THEN 'B'
        WHEN mark >= 50 THEN 'C'
        ELSE 'D'
    END;
```

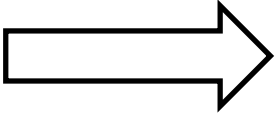
```
$$ LANGUAGE sql;
```

SELECT Name, convert(Mark)
FROM Scores;

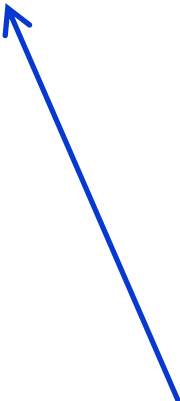
This generates the results above

PL/pgSQL Example

Scores	<u>Name</u>	<u>Mark</u>
	Alice	92
	Bob	63
	Cathy	58
	David	47



<u>Name</u>
Bob



```
CREATE OR REPLACE FUNCTION convert (mark INT)
RETURNS char(1) AS $$
```

```
    SELECT CASE
        WHEN mark >= 70 THEN 'A'
        WHEN mark >= 60 THEN 'B'
        WHEN mark >= 50 THEN 'C'
        ELSE 'D'
    END;
```

```
$$ LANGUAGE sql;
```

```
SELECT Name
FROM   Scores
WHERE  convert( Mark ) = 'B';

This generates the results above
```

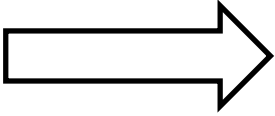
PL/pgSQL Example

- The function here returns one character
- There are other options:
 - Return an existing tuple or tuple set
 - Return a new tuple or tuple set

```
CREATE OR REPLACE FUNCTION convert (mark INT)
RETURNS char(1) AS $$
    SELECT CASE
        WHEN mark >= 70 THEN 'A'
        WHEN mark >= 60 THEN 'B'
        WHEN mark >= 50 THEN 'C'
        ELSE 'D'
    END;
$$ LANGUAGE sql;
```

Returning an existing tuple

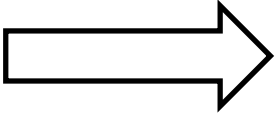
Scores	Name	Mark
	Alice	92
	Bob	63
	Cathy	58
	David	47



Name	Mark
Alice	92

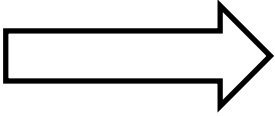
- Create a function `topStudent()` that returns any one tuple in Scores with the highest mark
- That is, `SELECT topStudent();` should return Alice's record

Returning an existing tuple

Scores	Name	Mark			Name	Mark
	Alice	92			Alice	92
	Bob	63				
	Cathy	58				
	David	47				

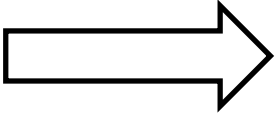
```
CREATE OR REPLACE FUNCTION topStudent ()  
RETURNS Scores AS $$  
    SELECT      *  
    FROM        Scores  
    ORDER BY    Mark DESC LIMIT 1;  
$$ LANGUAGE sql;
```

Returning an existing set of tuples

Scores	Name	Mark		Name	Mark
	Alice	92		Alice	92
	Bob	63		David	92
	Cathy	58			
	David	92			

- Create a function topStudents() that returns **all** tuples in Scores with the highest mark
- That is, SELECT * FROM topStudents(); should return Alice and David's records

Returning an existing set of tuples

Scores	Name	Mark		Name	Mark
	Alice	92		Alice	92
	Bob	63		David	92
	Cathy	58			
	David	92			

```
CREATE OR REPLACE FUNCTION topStudents ()  
RETURNS SETOF Scores AS $$  
    SELECT *  
    FROM Scores  
    WHERE Mark =  
           (SELECT MAX(Mark) FROM Scores);  
$$ LANGUAGE sql;
```

Returning existing tuples

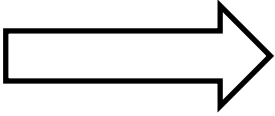
```
CREATE OR REPLACE FUNCTION topStudent ()  
RETURNS Scores AS $$  
    SELECT *  
    FROM Scores  
    ORDER BY Mark DESC LIMIT 1;  
$$ LANGUAGE sql;
```

```
CREATE OR REPLACE FUNCTION topStudents ()  
RETURNS SETOF Scores AS $$  
    SELECT *  
    FROM Scores  
    WHERE Mark =  
        (SELECT MAX(Mark) FROM Scores);  
$$ LANGUAGE sql;
```

- The above two examples return tuples from the Scores table
- What if we want to return some other values?

Returning a new tuple

Scores	Name	Mark
	Alice	92
	Bob	63
	Cathy	58
	David	92

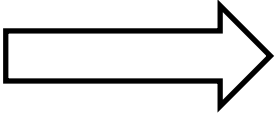


TopMark	Cnt
92	2

- Create a function topMarkCnt that returns the highest mark and its number of occurrences
- That is, `SELECT topMarkCnt();` should return the results above

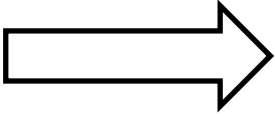
Returning a new tuple

Scores	Name	Mark		TopMark	Cnt
	Alice	92		92	2
	Bob	63			
	Cathy	58			
	David	92			



```
CREATE OR REPLACE FUNCTION topMarkCnt (OUT TopMark INT,  
                                         OUT Cnt INT)  
RETURNS RECORD AS $$  
    SELECT Mark, COUNT(*)  
    FROM   Scores  
    WHERE Mark = (SELECT MAX(Mark) FROM Scores)  
    GROUP BY      Mark ;  
$$ LANGUAGE sql;
```

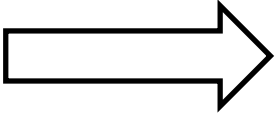
Returning a new set of tuples

Scores	Name	Mark		Mark	Cnt
	Alice	92		92	2
	Bob	63		63	1
	Cathy	58		58	1
	David	92			

- Create a function MarkCnt that returns each distinct Mark and its number of occurrences
- That is, SELECT MarkCnt();
should return the results above

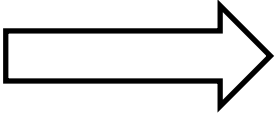
Returning a new set of tuples

Scores	Name	Mark		Mark	Cnt
	Alice	92		92	2
	Bob	63		63	1
	Cathy	58		58	1
	David	92			



```
CREATE OR REPLACE FUNCTION MarkCnt (OUT Mark INT,  
                                     OUT Cnt INT)  
RETURNS SETOF RECORD AS $$  
    SELECT Mark, COUNT(*)  
    FROM   Scores  
    GROUP BY    Mark ;  
$$ LANGUAGE sql;
```


Returning a new set of tuples 2

Scores	Name	Mark		Mark	Cnt
	Alice	92		92	2
	Bob	63		63	1
	Cathy	58		58	1
	David	92			

```
CREATE OR REPLACE FUNCTION MarkCnt ()  
RETURNS TABLE(Mark INT, Cnt INT) AS $$  
    SELECT Mark, COUNT(*)  
    FROM Scores  
    GROUP BY Mark ;  
$$ LANGUAGE sql;
```

SQL Functions vs. Procedures

- The above examples illustrate SQL functions
 - Each of which returns some values/tuples
 - If no return value/tuple is needed, we may use **SQL procedures** instead
-

SQL Procedure Example

```
CREATE OR REPLACE PROCEDURE transfer (fromAcct TEXT,  
                                       toAcct TEXT, toAcct TEXT, amount INT)  
AS $$  
    UPDATE Accounts  
    SET balance = balance - amount  
    WHERE name = fromAcct;  
  
    UPDATE Accounts  
    SET balance = balance + amount  
    WHERE name = toAcct;  
$$ LANGUAGE SQL
```

- To execute this procedure:
- **CALL** transfer('Alice', 'Bob', 100);

SQL with control structures

- The previous SQL functions/procedures are relatively simple
 - Each of them just sequentially executes some SQL statements
 - But in general, SQL functions/procedures can be complex, having variables and **control structures**:
 - IF ... THEN ... ELSE ... ENF IF
 - LOOP ... END LOOP
 - EXIT ... WHEN ...
 - WHILE ... LOOP ... END LOOP
 - FOR ... IN ... LOOP ... END LOOP
 - ...
-

Example: Using variables

- The following function swaps two input integers

```
CREATE OR REPLACE FUNCTION swap (INOUT val1 INT, INOUT val2 INT)
RETURNS RECORD AS $$
DECLARE
    temp_val INTEGER;
BEGIN
    temp_val := val1;
    val1 := val2;
    val2 := temp_val;
END;
$$ LANGUAGE plpgsql;
```

- Query: `SELECT swap(11, 22);`
- Result: (22, 11)

Example: IF ... THEN ... END IF

- The following function sorts two integers

```
CREATE OR REPLACE FUNCTION swap (INOUT val1 INT, INOUT val2 INT)
RETURNS RECORD AS $$
DECLARE
    temp_val INTEGER;
BEGIN
    IF val1 > val2 THEN
        temp_val := val1;
        val1 := val2;
        val2 := temp_val;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

- SELECT swap(99, 11); → (11, 99)
- SELECT swap(11, 99); → (11, 99)

Example: LOOP ... END LOOP

- The following function computes $1+2+\dots+x$

```
CREATE OR REPLACE FUNCTION sum_to_x (IN x INT, OUT s INT)
RETURNS INTEGER AS $$
DECLARE
    temp_val INTEGER;
BEGIN
    s := 0;
    temp_val := 1;
    LOOP
        EXIT WHEN temp_val > x;
        s := s + temp_val;
        temp_val := temp_val + 1;
    END LOOP;
END;
$$ LANGUAGE plpgsql;
```

Question

Scores	Name	Mark			
	Alice	92	Name	Mark	Gap
	Bob	63	Alice	92	NULL
	Cathy	58	Bob	63	29
	David	47	Cathy	58	5
			David	47	11

- Write a function for the following task:
 - Sort the students in Scores in descending order of their Marks
 - For each student, compute the different between his/her mark and the previous student

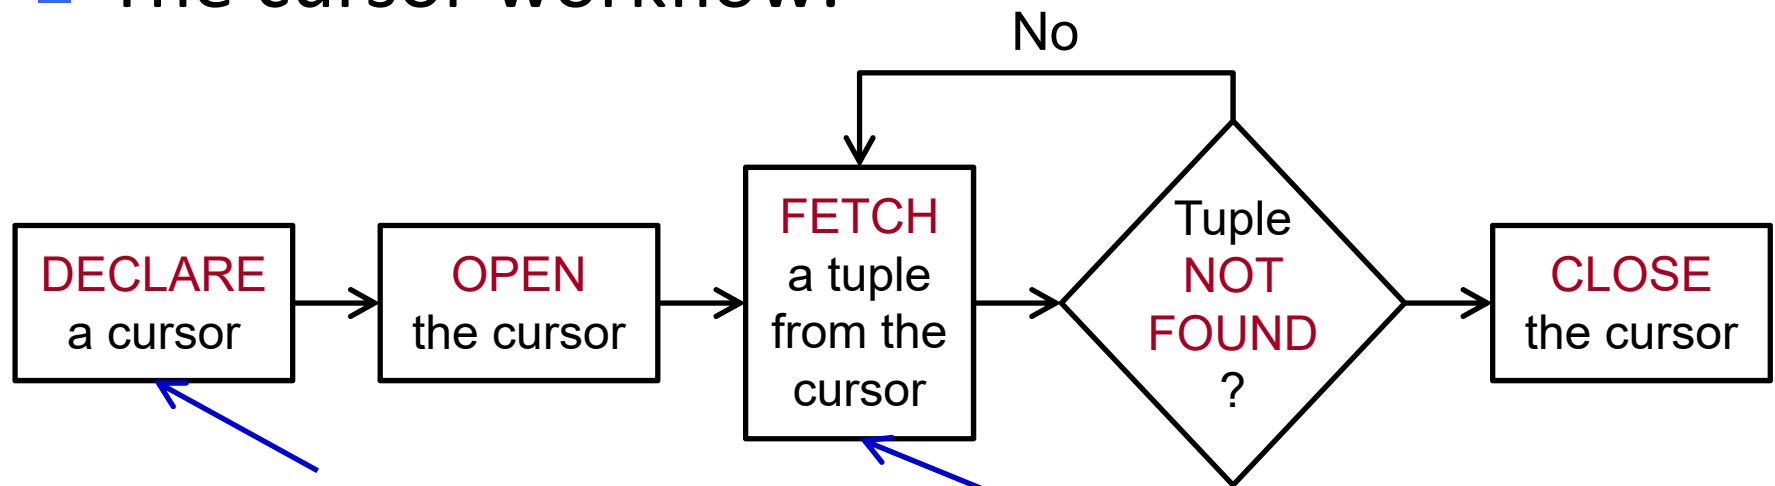
Question

Scores	Name	Mark			
	Alice	92	Name	Mark	Gap
	Bob	63	Alice	92	NULL
	Cathy	58	Bob	63	29
	David	47	Cathy	58	5
			David	47	11

- Idea:
 - `SELECT * FROM Scores ORDER BY Mark DESC;`
 - Loop over the sorted sequence of students
 - For each student, computes the score difference with the previous student
- Problem: How we do loop over the sorted sequence?
- Answer: Use a **cursor**

Cursor

- A cursor enables us to access each individual row returned by a SELECT statement
- The cursor workflow:

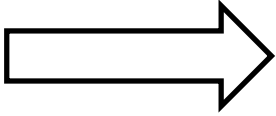


The cursor is associated with a SELECT statement at declaration

Once a tuple is fetched, we can do some operations based on it

Cursor Example

Scores	Name	Mark			
	Alice	92	Name	Mark	Gap
	Bob	63	Alice	92	NULL
	Cathy	58	Bob	63	29
	David	47	Cathy	58	5
			David	47	11



- Write a function for the following task:
 - Sort the students in Scores in descending order of their Marks
 - For each student, compute the different between his/her mark and the previous student

Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;
BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;

BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

This part declares
a cursor variable



Cursor Example


Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;

BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

This executes the SQL statement, and let curs point to the beginning of the result



Cursor Example


Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;

BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

Read the next
tuple from curs,
and put it into r.




Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;
BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

If the FETCH operation did not get any tuple, terminate the loop




Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;
BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

If the FETCH operation got a tuple, then we decide the values of name, mark, and gap.




Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;
BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

This inserts a tuple (name, mark gap) to the output of the function.



Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;

BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

This updates the
prv_mark
variable.


Cursor Example

Scores

Name	Mark
Alice	92
Bob	63
Cathy	58
David	47

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prv_mark INT;
BEGIN
    prv_mark := -1;
    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name;
        mark := r.Mark;
        IF prv_mark >= 0 THEN gap := prv_mark - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT;
        prv_mark := r.Mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

This releases the
resources
allocated to curs.



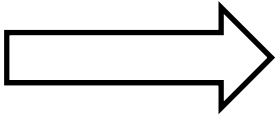
Cursor Movement

- In the previous example, we have seen
FETCH curs INTO r
 - There are other variants of FETCH
 - FETCH **PRIOR FROM** cur INTO r
 - Fetch the prior row
 - FETCH **FIRST FROM** cur INTO r
 - FETCH **LAST FROM** cur INTO r
 - FETCH **ABSOLUTE 3 FROM** cur INTO r
 - Fetch the 3-rd tuple
 - ...
-

SQL/PSM Summary

- An extension of SQL with a procedural language
 - Enables us to create complex functions/procedures in databases
 - as well as **triggers** (will be covered next week)
 - For additional details about PL/pgSQL, please refer to
 - <https://www.postgresql.org/docs/current/plpgsql.html>
-

Exercise

Scores	Name	Mark		Name	Mark
	Alice	92		Bob	63
	Bob	63		Cathy	58
	Cathy	58			
	David	47			

- Write a function that retrieves the student(s) with the **median** mark(s):
- What we mean by "median"
 - Let n be the total number of students
 - If n is odd, then the median is the $((n+1)/2)$ -th
 - If n is even, then the medians are the $(n/2)$ -th and the $(n/2+1)$ -th
 - Ties are broken arbitrarily

```
CREATE OR REPLACE FUNCTION median_stu()
RETURNS TABLE ( name TEXT, mark INT ) AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    n INT;
BEGIN
    OPEN curs;
    SELECT COUNT(*) INTO n FROM Scores;
    IF n % 2 = 1 THEN
        FETCH ABSOLUTE (n+1)/2 FROM curs INTO r;
        name := r.name;
        mark := r.Mark;
        RETURN NEXT;
    ELSE
        FETCH ABSOLUTE n/2 FROM curs INTO r;
        name := r.name;
        mark := r.Mark;
        RETURN NEXT;
        FETCH ABSOLUTE n/2 + 1 FROM curs INTO r;
        name := r.name;
        mark := r.Mark;
        RETURN NEXT;
    END IF;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE OR REPLACE FUNCTION median_stu()
RETURNS SETOF Scores AS $$
DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    n INT;
BEGIN
    OPEN curs;
    SELECT COUNT(*) INTO n FROM Scores;
    IF n % 2 = 1 THEN
        FETCH ABSOLUTE (n+1)/2 FROM curs INTO r;
        RETURN NEXT r;
    ELSE
        FETCH ABSOLUTE n/2 FROM curs INTO r;
        RETURN NEXT r;
        FETCH ABSOLUTE n/2 + 1 FROM curs INTO r;
        RETURN NEXT r;
    END IF;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;
```

Coming Next

- SQL Injection Attacks

SQL Injection Attacks

- A type of attacks on dynamic SQL
 - Example:
 - Suppose that we have a user-provided parameter *input_name*
 - We construct a string *stmt* by concatenating three sub-strings:
 - SELECT * FROM T WHERE Name = '
 - *input_name*
 - ';
 - And we have the following dynamic SQL:
 - EXEC SQL EXECUTE IMMEDIATE :*stmt*;
 - For instance, if *input_name* = "Alice", then the executed query is:
 - SELECT * FROM T WHERE Name = 'Alice';
-

SQL Injection Attacks

- `stmt` is the concatenation of three sub-strings:
 - `SELECT * FROM T WHERE Name = '`
 - *`input_name`*
 - `';`
 - `EXEC SQL EXECUTE IMMEDIATE :stmt;`
 - A malicious user sets *`input_name`* to the following:
`a'; DROP TABLE T; SELECT 'a`
 - What would be the SQL statement in `stmt`?
 - `SELECT * FROM T WHERE Name = 'a'; DROP TABLE T; SELECT 'a';`
 - T will be deleted once `stmt` is executed
-

SQL Injection Attacks

- A malicious user sets *input_name* to the following:
a'; DROP TABLE T; SELECT 'a
 - What would be the SQL statement in *stmt*?
 - SELECT * FROM T WHERE Name = 'a'; DROP TABLE T;
SELECT 'a';
 - Main issue:
 - *input_name* is supposed to be a name
 - But the malicious user hides a DROP TABLE query in it
 - The database executes it as it is
 - Solution: Let the database know what it is supposed to execute
-

Solution: Use function/procedures

```
CREATE OR REPLACE FUNCTION queryT (IN  
in_name TEXT)  
RETURNS SETOF T AS $$  
    SELECT * FROM T WHERE Name = in_name;  
$$ LANGUAGE sql;
```

■ Rationale:

- ❑ The function is compiled and stored in the database
 - ❑ At runtime, anything in in_name would be treated as a string constant
 - ❑ Even if in_name is set to **a'; DROP TABLE T; SELECT 'a**
-

Functions/procedures in SQL

- Advantages:
 - Code reuse
 - Ease of maintenance
 - Performance
 - Security
-

Solution: use prepared statements

```
EXEC SQL BEGIN DECLARE SECTION;  
    const char *stmt = "SELECT * FROM T WHERE Name = ?;";  
    char in_name[100];  
EXEC SQL END DECLARE SECTION;  
  
// get the user-provided name and put it into in_name;  
  
EXEC SQL PREPARE mystmt FROM :stmt;  
  
EXEC SQL EXECUTE mystmt USING :in_name;
```

■ Rationale:

- ❑ The SQL statement is compiled when it is prepared
- ❑ Anything in `in_name` is treated a string constant