

1. Solution:

| <table><tr><th>a</th><th>b</th></tr><tr><td>2</td><td>20</td></tr><tr><td>3</td><td>30</td></tr></table> | a | b | 2 | 20 | 3 | 30 | Invalid query | <table><tr><th>a</th><th>b</th></tr><tr><td>2</td><td>20</td></tr><tr><td>3</td><td>30</td></tr></table> | a | b | 2 | 20 | 3 | 30 | <table><tr><th>a</th><th>b</th></tr><tr><td>1</td><td>10</td></tr><tr><td>2</td><td>20</td></tr><tr><td>3</td><td>30</td></tr><tr><td>4</td><td>40</td></tr></table> | a | b | 1 | 10 | 2 | 20 | 3 | 30 | 4 | 40 | <table><tr><th>a</th><th>b</th></tr><tr><td>3</td><td>30</td></tr></table> | a | b | 3 | 30 |
|--|-----|-----|-----|-----|---|----|---------------|--|---|---|---|----|---|----|--|---|---|---|----|---|----|---|----|---|----|--|---|---|---|----|
| a | b | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 2 | 20 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 | 30 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| a | b | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 2 | 20 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 | 30 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| a | b | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 1 | 10 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 2 | 20 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 | 30 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 4 | 40 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| a | b | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 3 | 30 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| (a) | (b) | (c) | (d) | (e) | | | | | | | | | | | | | | | | | | | | | | | | | | |

Query (b) is invalid because the alias *f* for table foo declared in the inner query is not visible to the outer query.

In query (c), “a = b.a” is equivalent to “f.a = b.a”. In query (d), “f > 100” is equivalent to “f.f > 100”, and “a = a” is equivalent to “f.a = f.a”. In query (e), “b > 20” is equivalent to “b.b > 20”.

It is a good defensive programming practice to use distinct aliases and use explicitly qualified column names.

2. Solution:

- (a) `select distinct cname
from Likes L, Sells S
where L.pizza = S.pizza
and S.rname = 'Corleone Corner';`
- (b) `select cname from Customers
except
select cname
from Likes L, Sells S
where L.pizza = S.pizza
and S.rname = 'Corleone Corner';`
- (c) `select distinct S.rname
from Sells S, Sells S2
where S.rname <> 'Corleone Corner'
and S2.rname = 'Corleone Corner'
and S.price > S2.price;`
- (d) `select rname, pizza, price
from Sells
where price is not null
except
select S.rname, S.pizza, S.price
from Sells S, Sells S2
where S.rname = S2.rname
and S.price < S2.price;`

3. Solution:**(a) Solution 1:**

```
select pizza
from Likes
where cname = 'Alice'
and pizza not in
    (select pizza
     from Likes
     where cname = 'Bob');
```

Solution 2:

```
select pizza
from Likes L1
where cname = 'Alice'
and not exists (
    select 1
    from    Likes L2
    where   L2.cname = 'Bob'
    and     L2.pizza = L1.pizza
);
```

Solution 3:

```
select pizza
from Likes
where cname = 'Alice'
and not pizza = any (
    select pizza
    from Likes
    where cname = 'Bob'
);
```

Note that if Bob doesn't like any pizza, then the ANY subquery will evaluate to *false*, and “not *false*” will evaluate to *true*; thus, the query will return all the pizzas that Alice likes.

Solution 4:

```
select pizza from Likes where cname = 'Alice' except
select pizza from Likes where cname = 'Bob';
```

Wrong answer: The following answer is incorrect.

```
select pizza
from Likes
where cname = 'Alice'
and pizza <> any (
    select pizza
    from    Likes
```

```

    where    cname = 'Bob'
  );

```

This answer looks similar to Solution 3 but it is incorrect. If Bob doesn't like any pizza, then the ANY subquery will evaluate to *false* and the query will return an empty set, which is incorrect if Alice likes some pizza.

- (b) A pizza is the output if there does not exist two distinct restaurants that are located in the same area selling that pizza.

```

select distinct pizza
from Sells S3
where not exists (
    select 1
    from Sells S, Restaurants R, Sells S2, Restaurants R2
    where S.rname = R.rname
    and S2.rname = R2.rname
    and S.pizza = S2.pizza
    and R.area = R2.area
    and R.rname <> R2.rname
    and S.pizza = S3.pizza
);

```

Wrong answer: The following answer is incorrect.

```

select distinct pizza
from Sells S, Restaurants R
where S.rname = R.rname
and not exists (
    select 1
    from Sells S2, Restaurants R2
    where S2.rname = R2.rname
    and S2.pizza = S.pizza
    and R.area = R2.area
    and R2.rname <> R.rname
);

```

This answer computes the pizzas that are sold by at most one restaurant in **some** area, which is a weaker condition than what is required by the question.

- (c)

```

select distinct R.area, S.pizza, S.price
from Restaurants R, Sells S
where R.rname = S.rname
and S.price <= all (
    select S2.price
    from Restaurants R2, Sells S2
    where R2.rname = S2.rname
    and R2.area = R.area
    and S2.pizza = S.pizza
);

```

- (d) You should recognize that this query is simply an extension of the previous query requiring an additional information (highest selling price) for each area-pizza pair. For a given area-pizza pair (A, P) , the following query will compute the highest price of pizza P in area A :

```
select distinct S2.price
from Restaurants R2, Sells S2
where R2.rname = S2.rname
and R2.area = A
and S2.pizza = P
and R2.price >= all (
    select price
    from Restaurants R3, Sells S3
    where R3.rname = S3.rname
    and R3.area = A
    and S3.pizza = P
);
```

Since the above query will return a single one-column tuple, it can be used as a scalar subquery to extend the previous question's solution as follows.

```
select distinct R.area, S.pizza, S.price as minPrice, (
    select distinct S2.price
    from Restaurants R2, Sells S2
    where R2.rname = S2.rname
    and R2.area = R.area
    and S2.pizza = S.pizza
    and S2.price >= all (
        select S3.price
        from Restaurants R3, Sells S3
        where R3.rname = S3.rname
        and R3.area = R.area
        and S3.pizza = S.pizza
    )
) as maxPrice
from Restaurants R, Sells S
where R.rname = S.rname
and S.price <= all (
    select S2.price
    from Restaurants R2, Sells S2
    where R2.rname = S2.rname
    and R2.area = R.area
    and S2.pizza = S.pizza
);
```

We will learn about other (simpler and more elegant) solutions for such queries later in class.

4. **Solution:** Queries $Q1$ and $Q2$ are equivalent. Whether the selection predicate "area = 'East'" is evaluated before or after the join/cross product operation does not change the semantics of the query.
5. **Solution:** Queries $Q1$ and $Q2$ are not equivalent. Observe that if the **Likes** relation is empty, then query $Q1$ simplifies to $\pi_{rname}(Sells \bowtie \sigma_{\text{area}='East'}(Restaurants))$, but the result of $Q2$ is always an empty set due to $Sells \times Restaurants \times \emptyset$.

6. **Solution:**

```
update Employees
set officeId =
    (select officeId from Offices
     where building = 'Tower1'
     and level = 5
     and roomNumber = 11)
where officeId = 123;
```

7. **Solution:**

(a) **select * from R natural join S;**

| A | B | X | Y | Z | C | D |
|---|---|----|---|---|----|-----|
| 8 | 5 | 30 | 0 | 1 | 60 | 100 |
| 4 | 3 | 60 | 1 | 3 | 30 | 100 |

(b) **select * from R inner join S on R.A = S.A;**

| X | A | Y | B | Z | A | B | C | D |
|----|---|---|---|---|---|---|----|-----|
| 30 | 8 | 0 | 5 | 1 | 8 | 5 | 60 | 500 |
| 60 | 4 | 1 | 3 | 3 | 4 | 2 | 40 | 200 |
| 60 | 4 | 1 | 3 | 3 | 4 | 3 | 30 | 100 |

(c) **select * from R left outer join S on R.A = S.A;**

| X | A | Y | B | Z | A | B | C | D |
|----|----|---|---|---|------|------|------|------|
| 0 | 10 | 0 | 9 | 2 | null | null | null | null |
| 30 | 8 | 0 | 5 | 1 | 8 | 5 | 60 | 500 |
| 60 | 4 | 1 | 3 | 3 | 4 | 2 | 40 | 200 |
| 60 | 4 | 1 | 3 | 3 | 4 | 3 | 30 | 100 |
| 90 | 0 | 0 | 4 | 5 | null | null | null | null |

(d) **select * from R right outer join S on R.A = S.A;**

| X | A | Y | B | Z | A | B | C | D |
|------|------|------|------|------|----|---|----|-----|
| 30 | 8 | 0 | 5 | 1 | 8 | 5 | 60 | 500 |
| 60 | 4 | 1 | 3 | 3 | 4 | 2 | 40 | 200 |
| 60 | 4 | 1 | 3 | 3 | 4 | 3 | 30 | 100 |
| null | null | null | null | null | 17 | 1 | 20 | 100 |

(e) **select * from R full outer join S on R.A = S.A;**

| X | A | Y | B | Z | A | B | C | D |
|------|------|------|------|------|------|------|------|------|
| 0 | 10 | 0 | 9 | 2 | null | null | null | null |
| 30 | 8 | 0 | 5 | 1 | 8 | 5 | 60 | 500 |
| 60 | 4 | 1 | 3 | 3 | 4 | 2 | 40 | 200 |
| 60 | 4 | 1 | 3 | 3 | 4 | 3 | 30 | 100 |
| 90 | 0 | 0 | 4 | 5 | null | null | null | null |
| null | null | null | null | null | 17 | 1 | 20 | 100 |