
CS2106

Process Management

Inter-Process Communication

Lecture 4

Overview

- Inter-process Communication
 - Motivation
 - Common communication mechanisms:
 - Shared memory
 - Message passing
 - Pipes (Unix-specific)
 - Signal (Unix-specific)

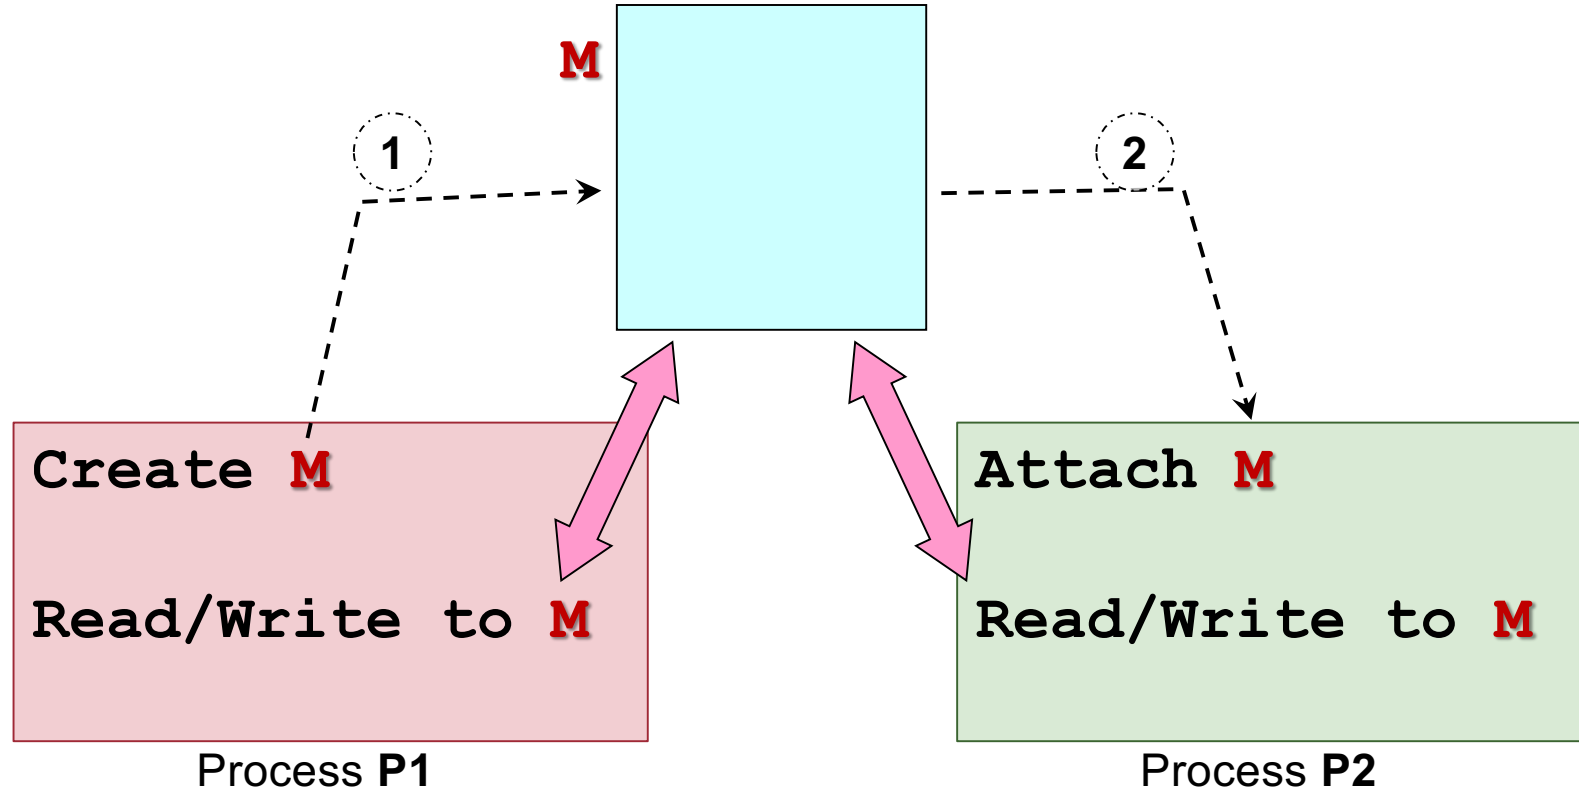
Inter-Process Communication (**IPC**)

- It is hard for cooperating processes to share information
 - Memory space is independent!
 - Inter-Process Communication mechanisms (IPC) is needed
- Two common IPC mechanisms:
 - Shared-Memory and Message Passing
- Two Unix-specific IPC mechanisms:
 - Pipe and Signal

Shared-Memory

- Communication happens through reads/writes to shared variables
- General Idea:
 - Process P_1 creates a shared memory region M
 - Process P_2 attaches memory region M to its own memory space
 - P_1 and P_2 can now communicate using memory region M
 - M behaves very similar to normal memory region
 - Any writes to the region can be seen by all other parties
- The same model is applicable to multiple processes sharing the same memory region

Shared-Memory: Illustration



- OS is involved in step 1 and 2 only

Shared-Memory: Pros and Cons

■ Advantages:

- ❑ Efficient:
 - Only the initial steps (e.g., **Create** and **Attach** shared memory region) involves OS
- ❑ Ease of use:
 - Shared memory region behaves the same as normal memory space
 - i.e., information of any type or size can be written easily

■ Disadvantages:

- ❑ Requires Synchronization:
 - Shared resource → Need to synchronize accesses (more later)
 - Without synchronization: data races (a.k.a. race conditions) → incorrect behavior
- ❑ Implementation is usually harder

Example

```
int counter=5;  
//shared variable  
  
void incCounter() {  
    counter++;  
}  
  
void decCounter() {  
    counter--;  
}
```

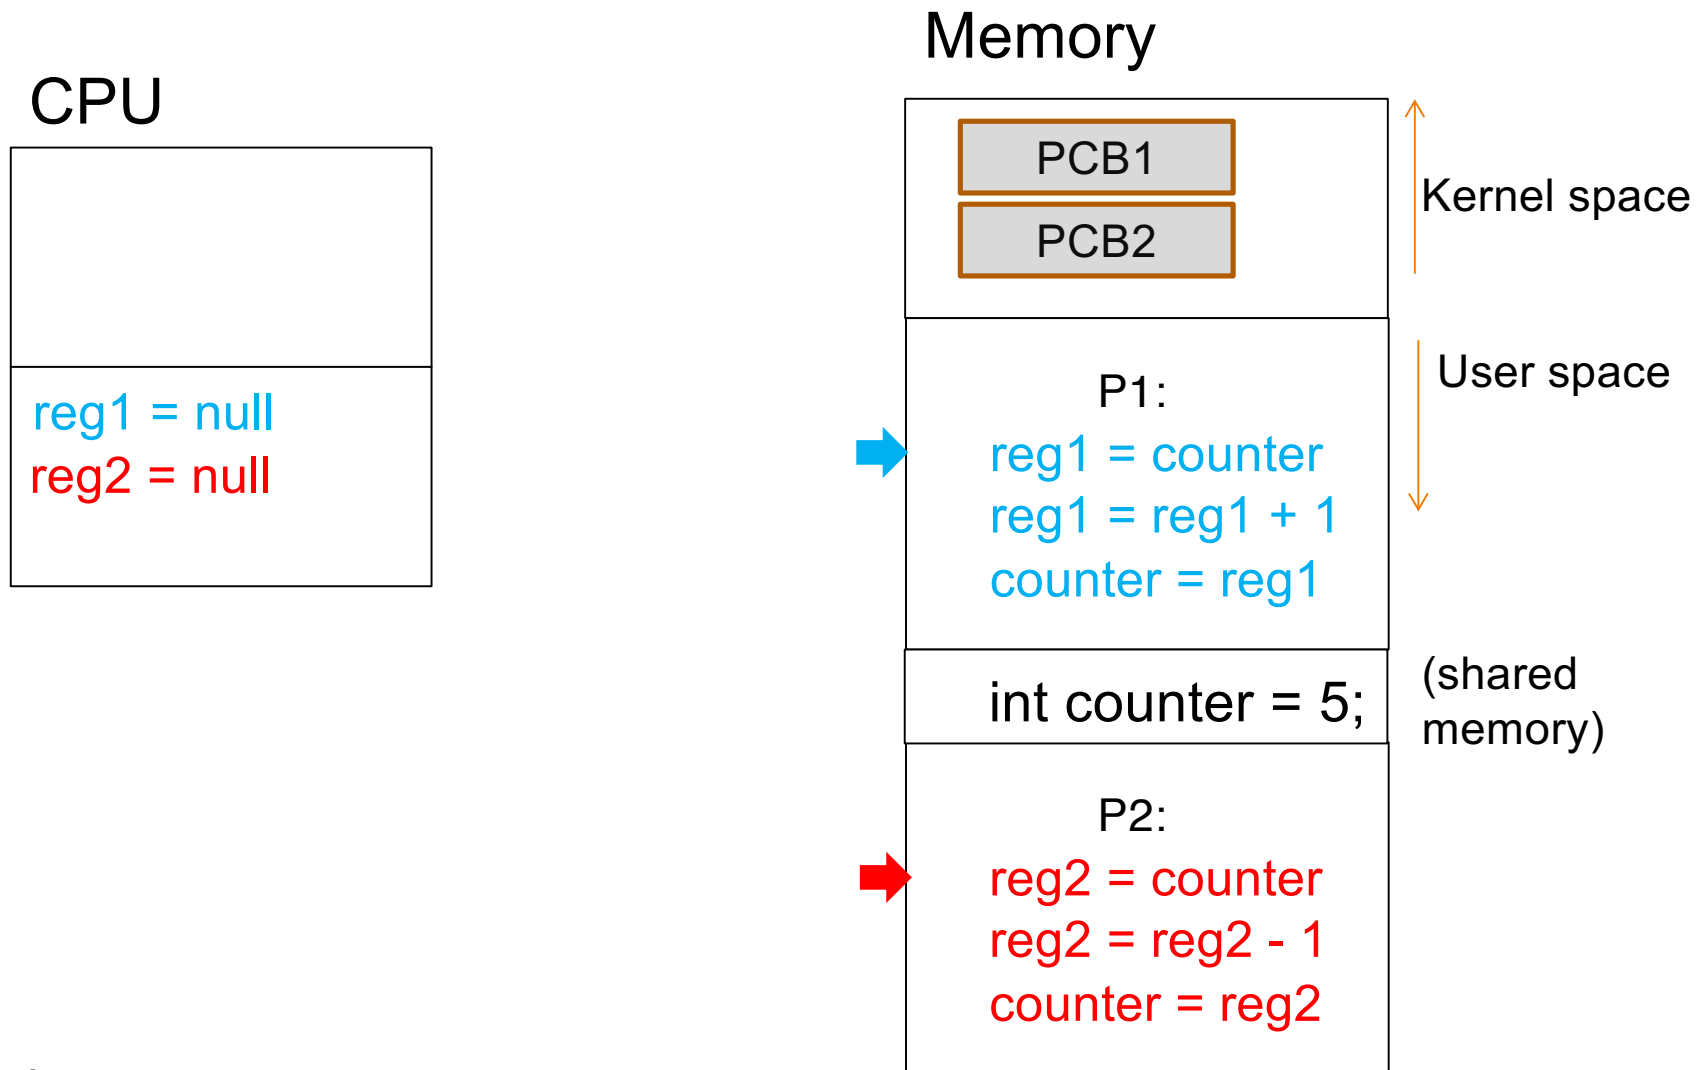
```
Process P0: incCounter();  
Process P1: decCounter();
```

compiled:

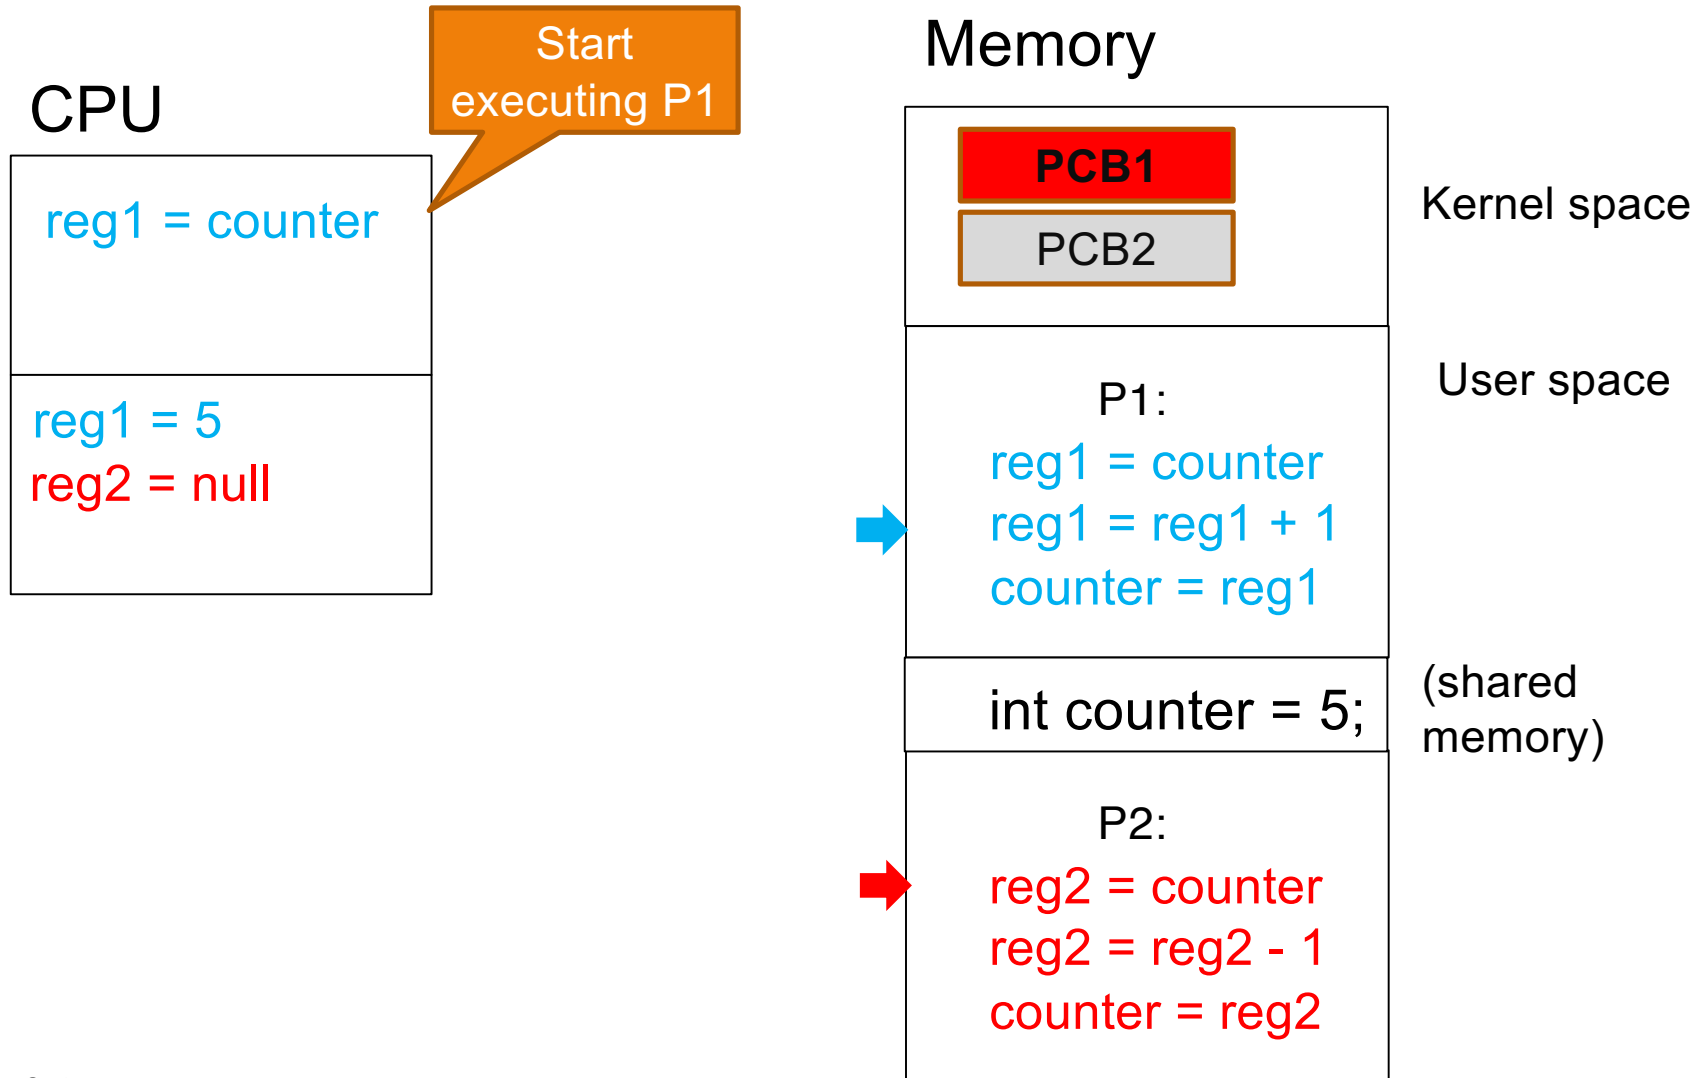
```
incCounter()  
    load    reg1, counter  
    add     reg1, reg1, 1  
    store   reg1, counter
```

```
decCounter()  
    load    reg2, counter  
    sub     reg2, reg2, 1  
    store   reg2, counter
```

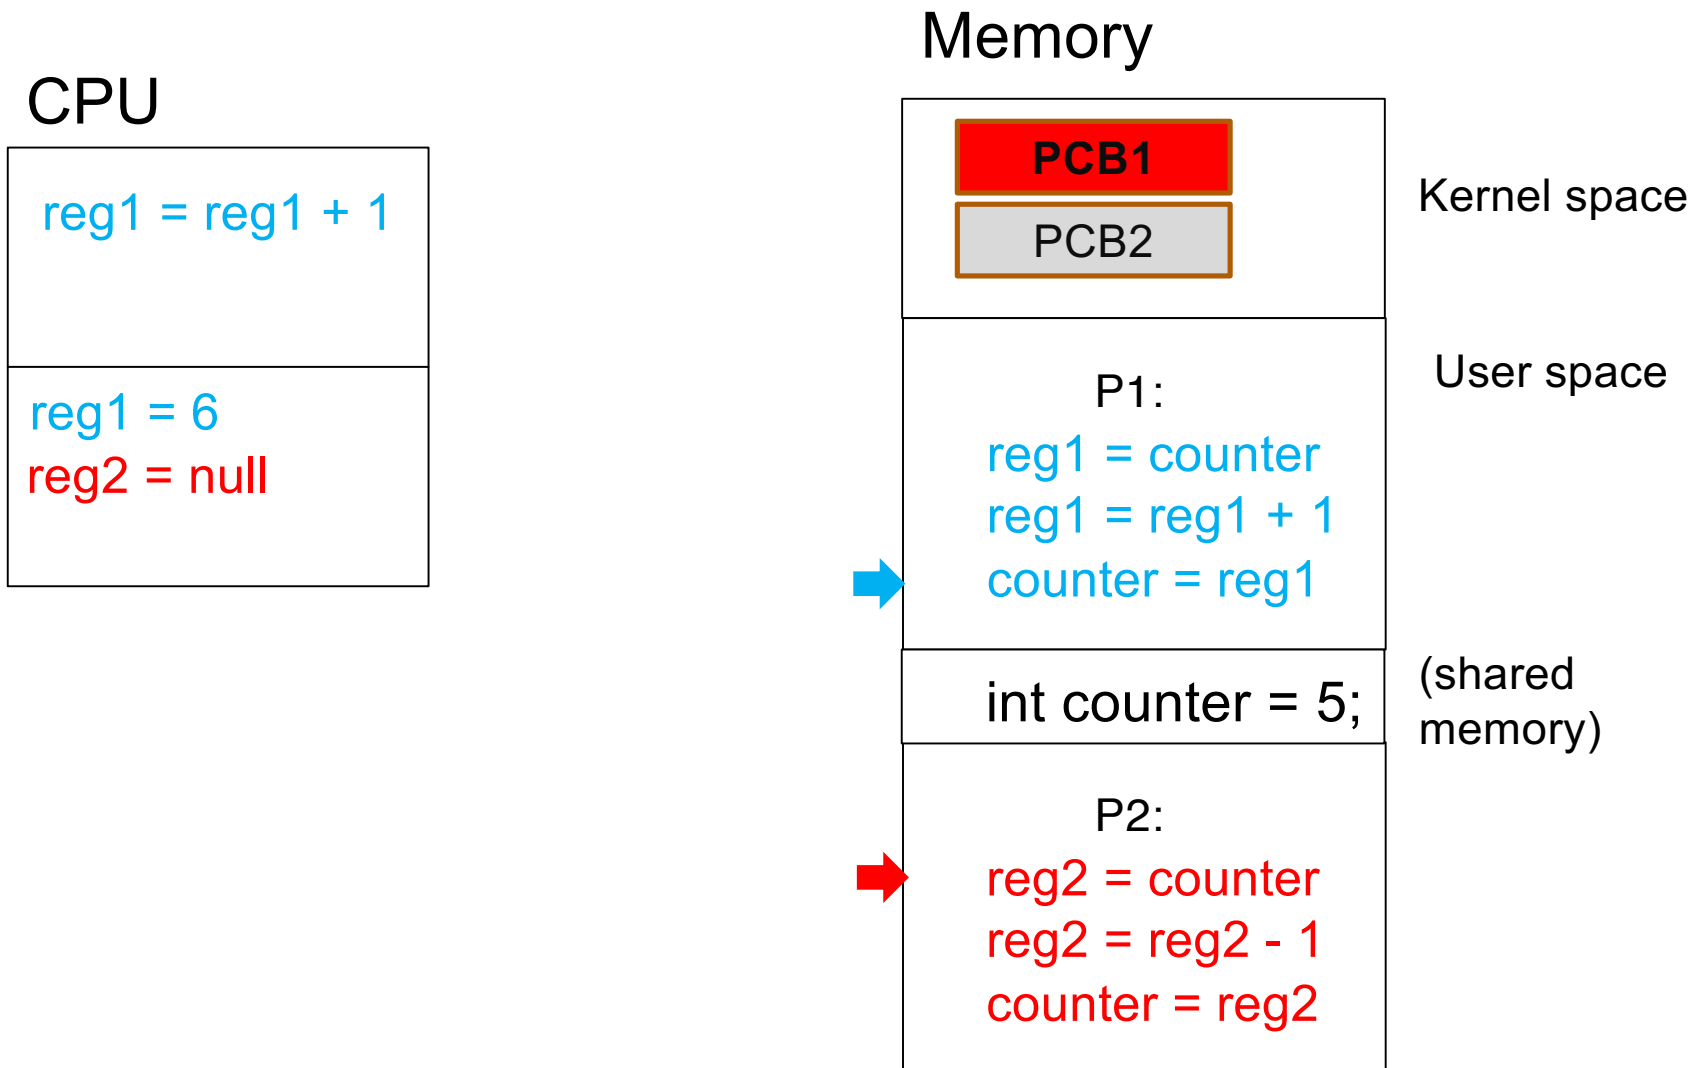
Example



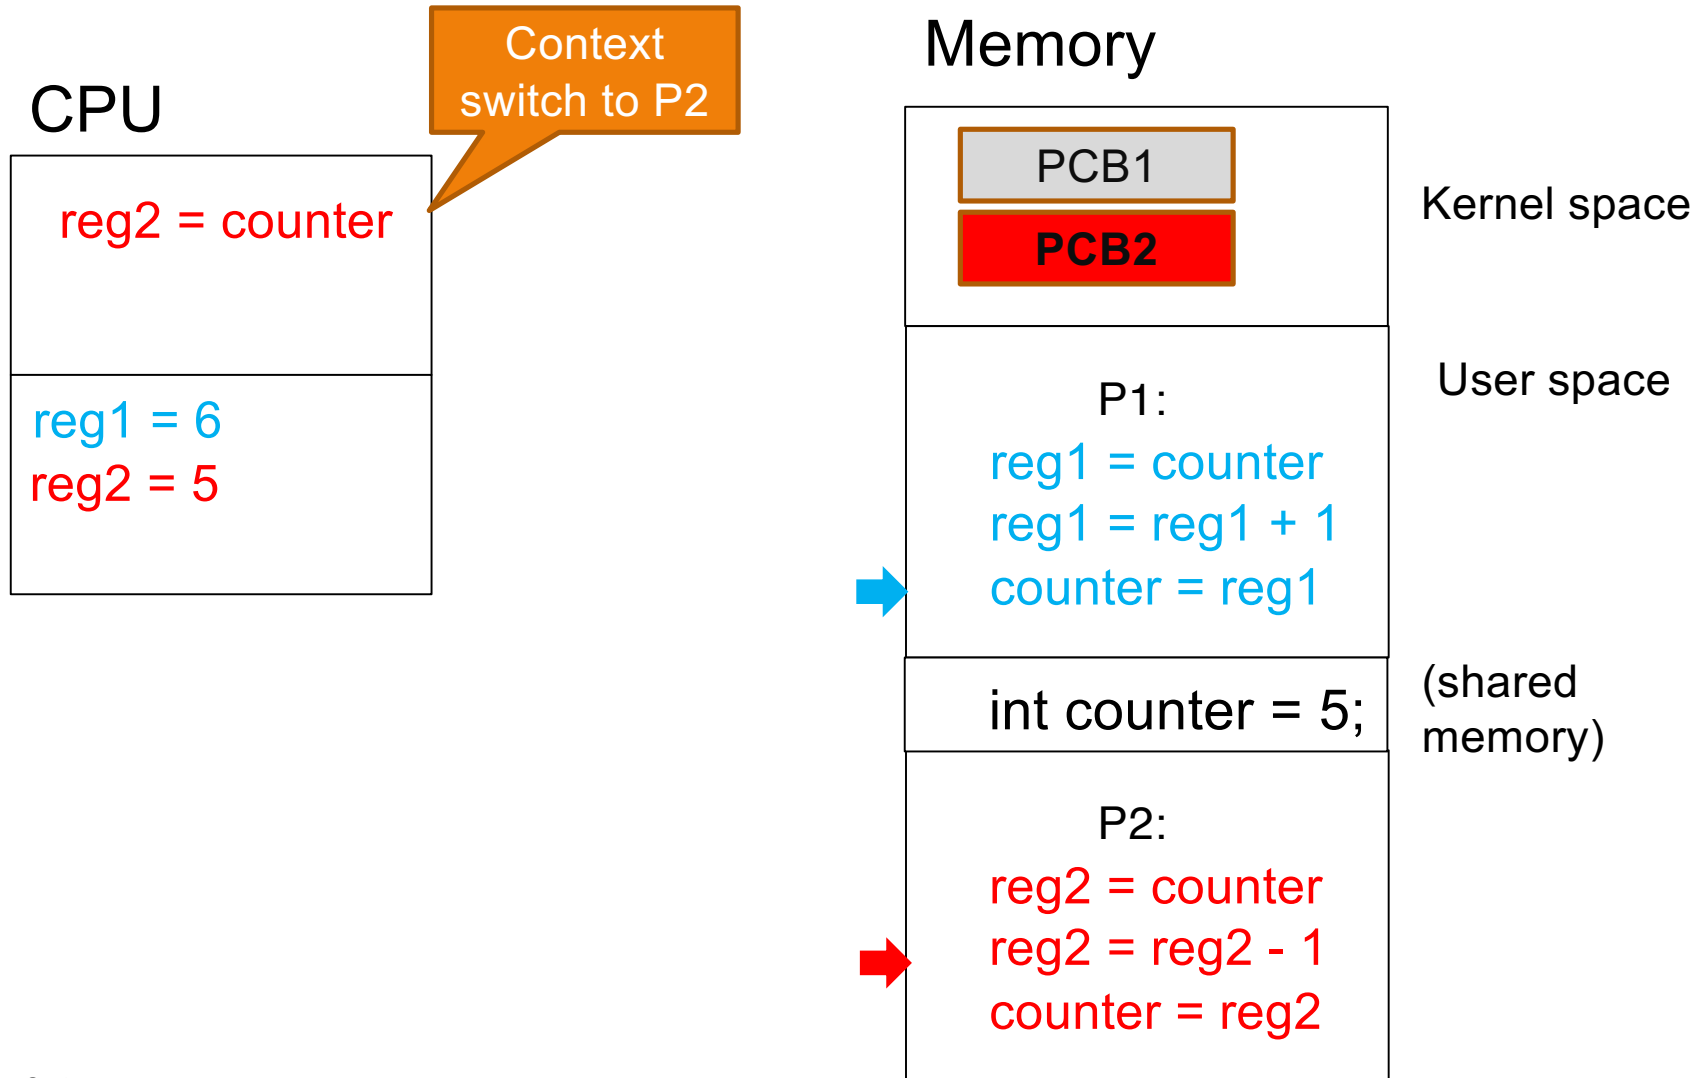
Example



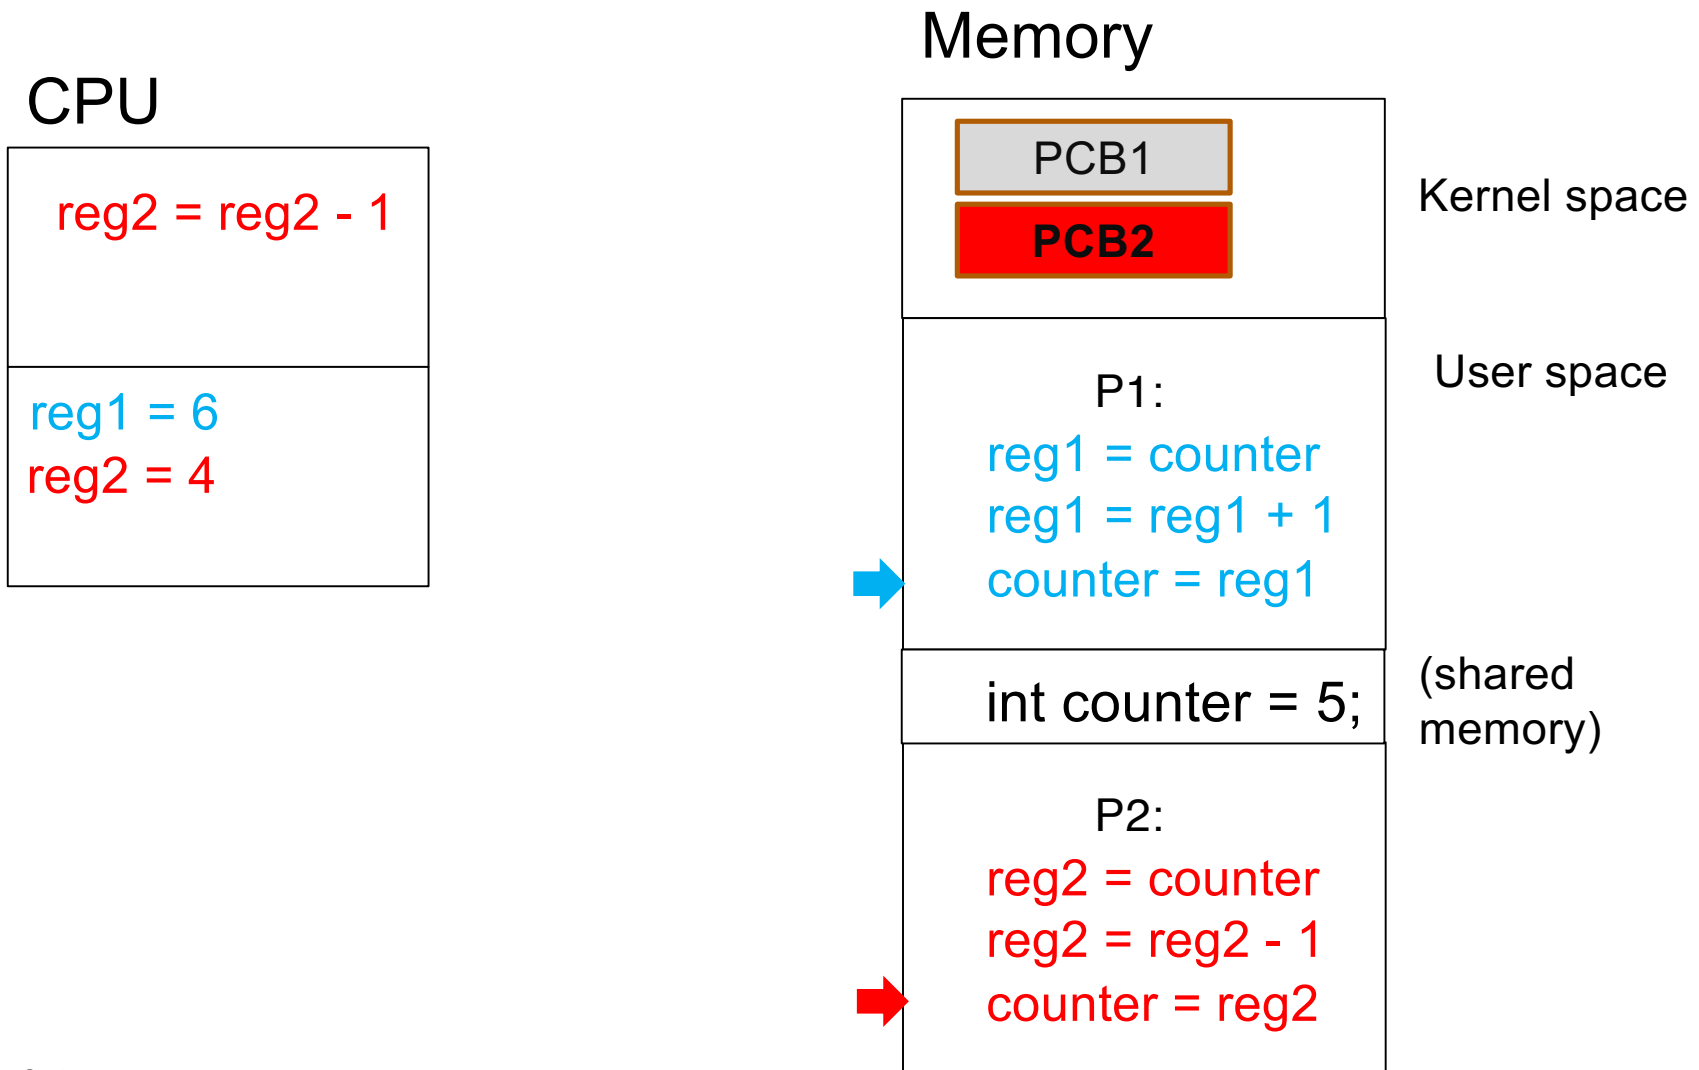
Example



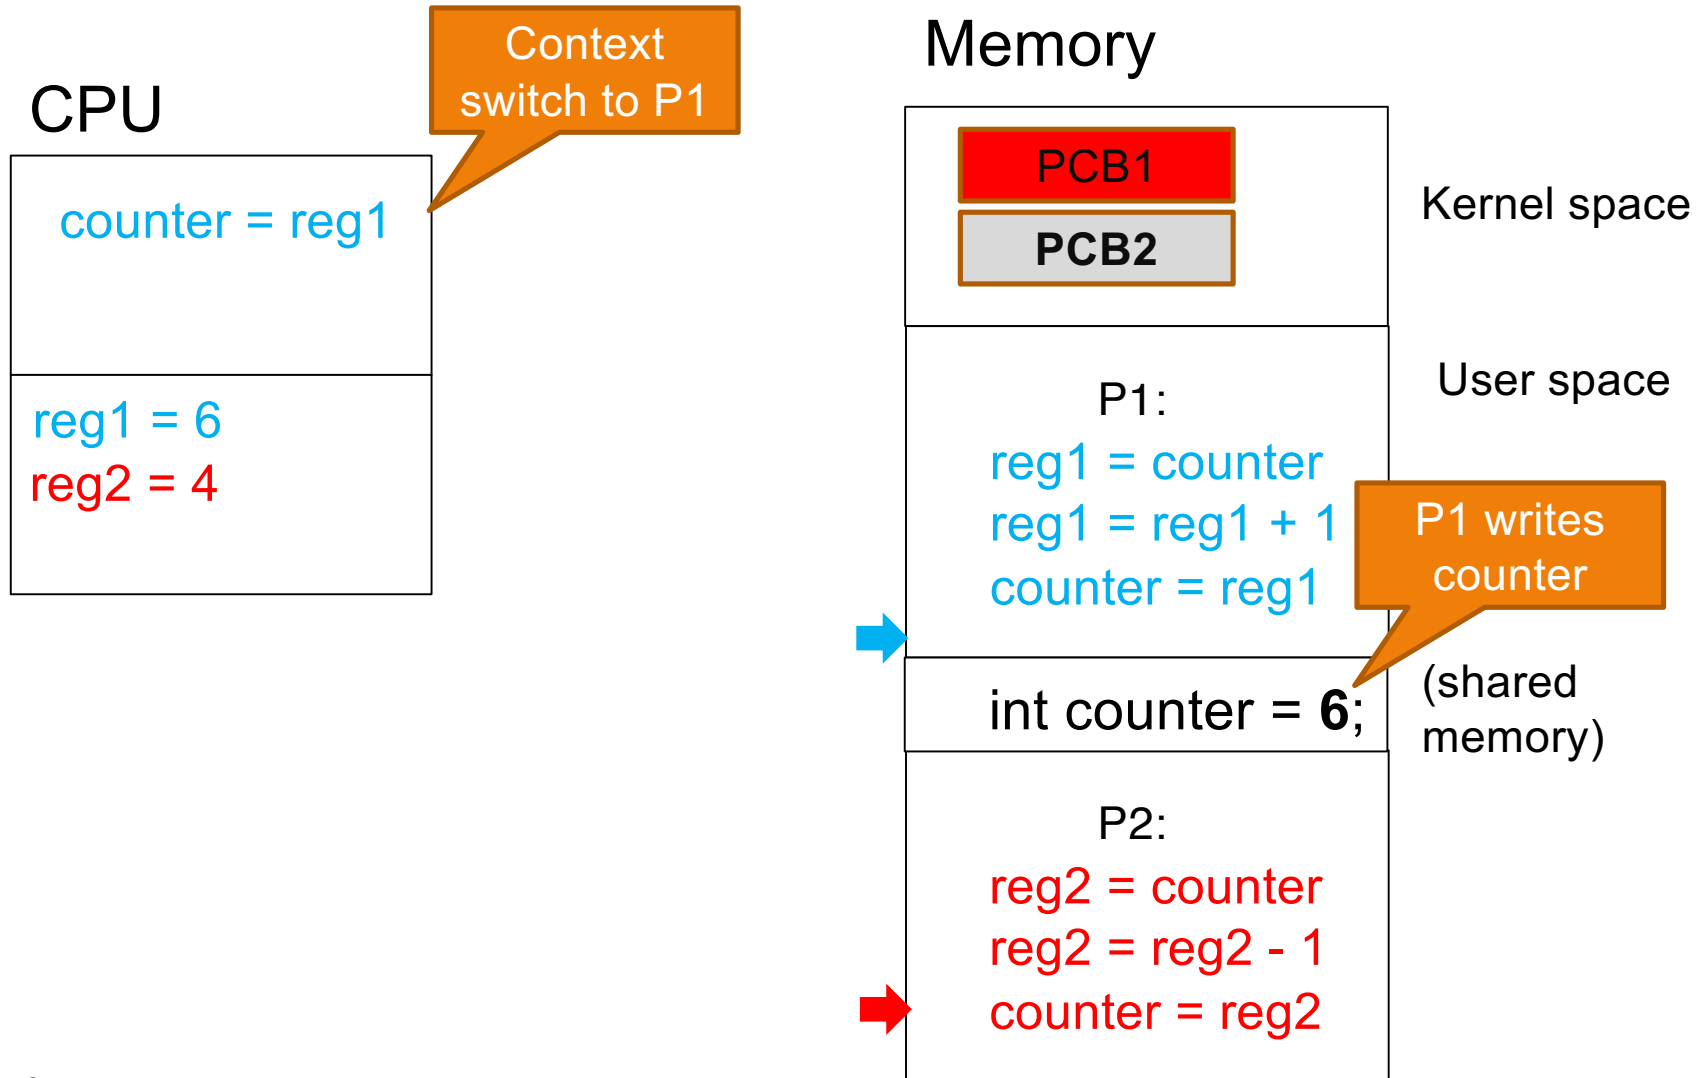
Example



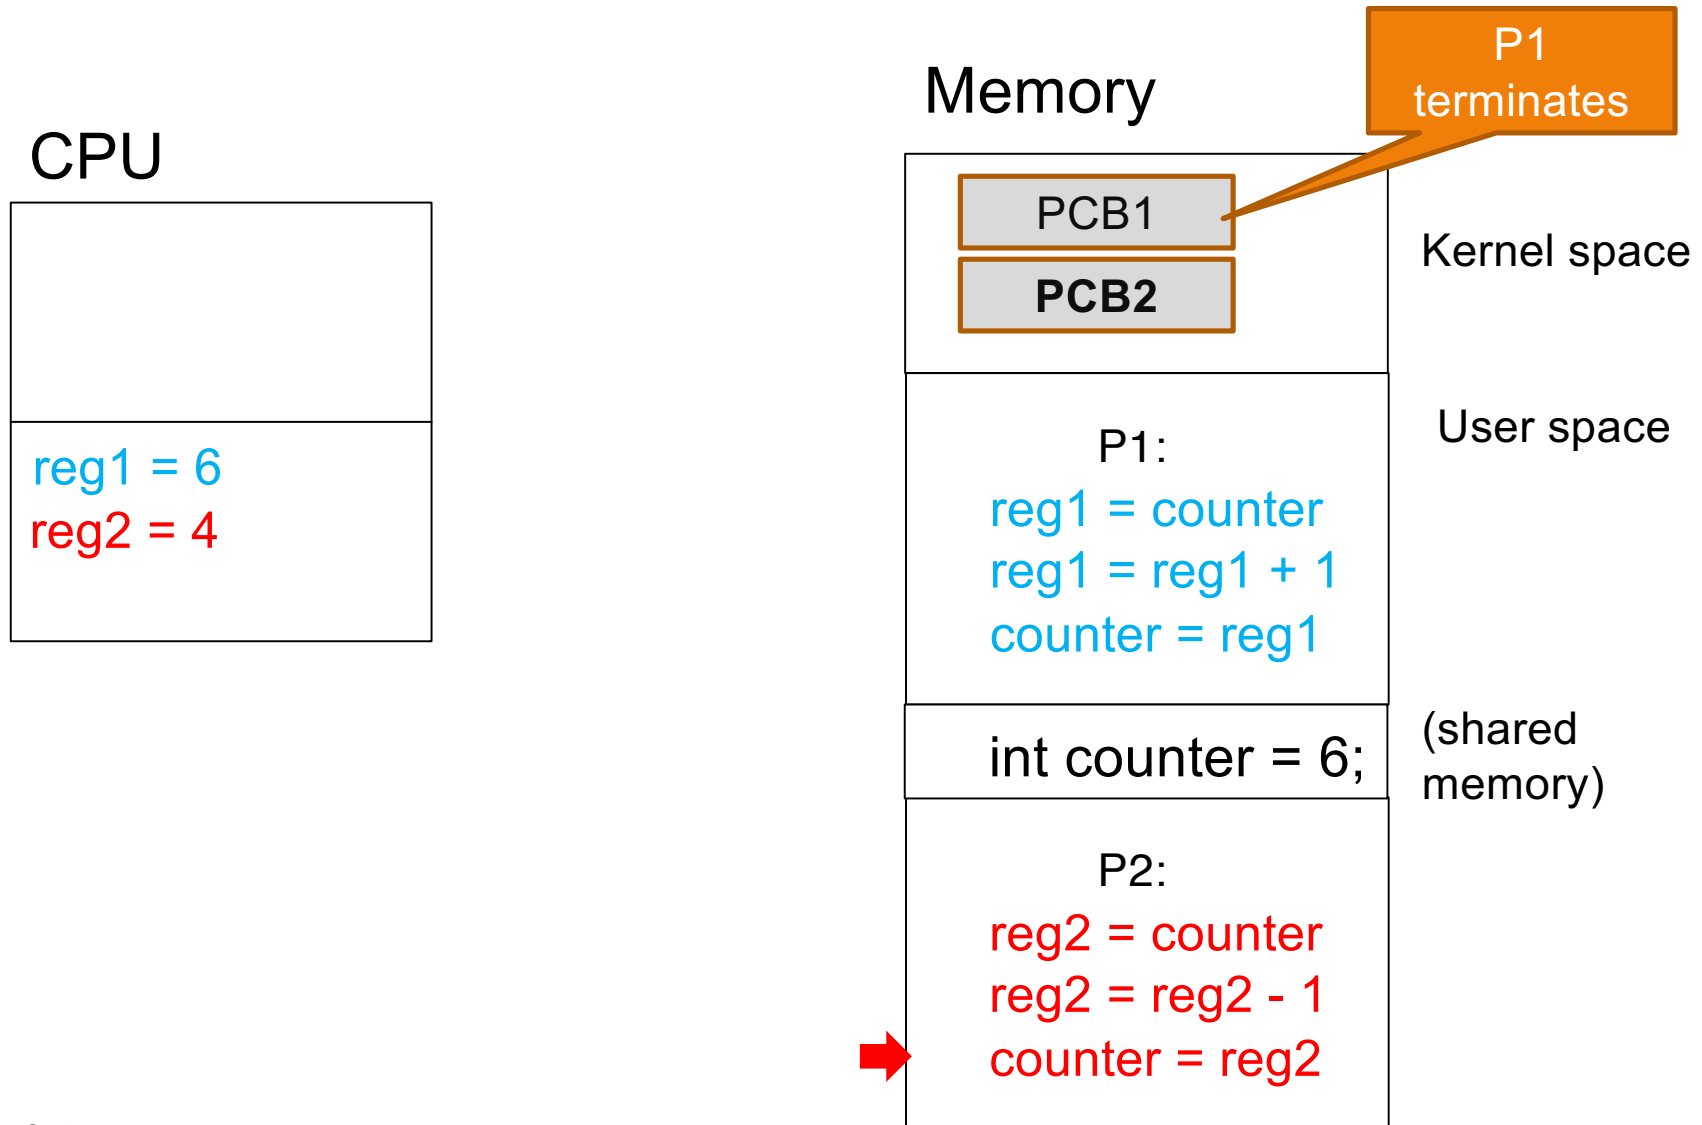
Example



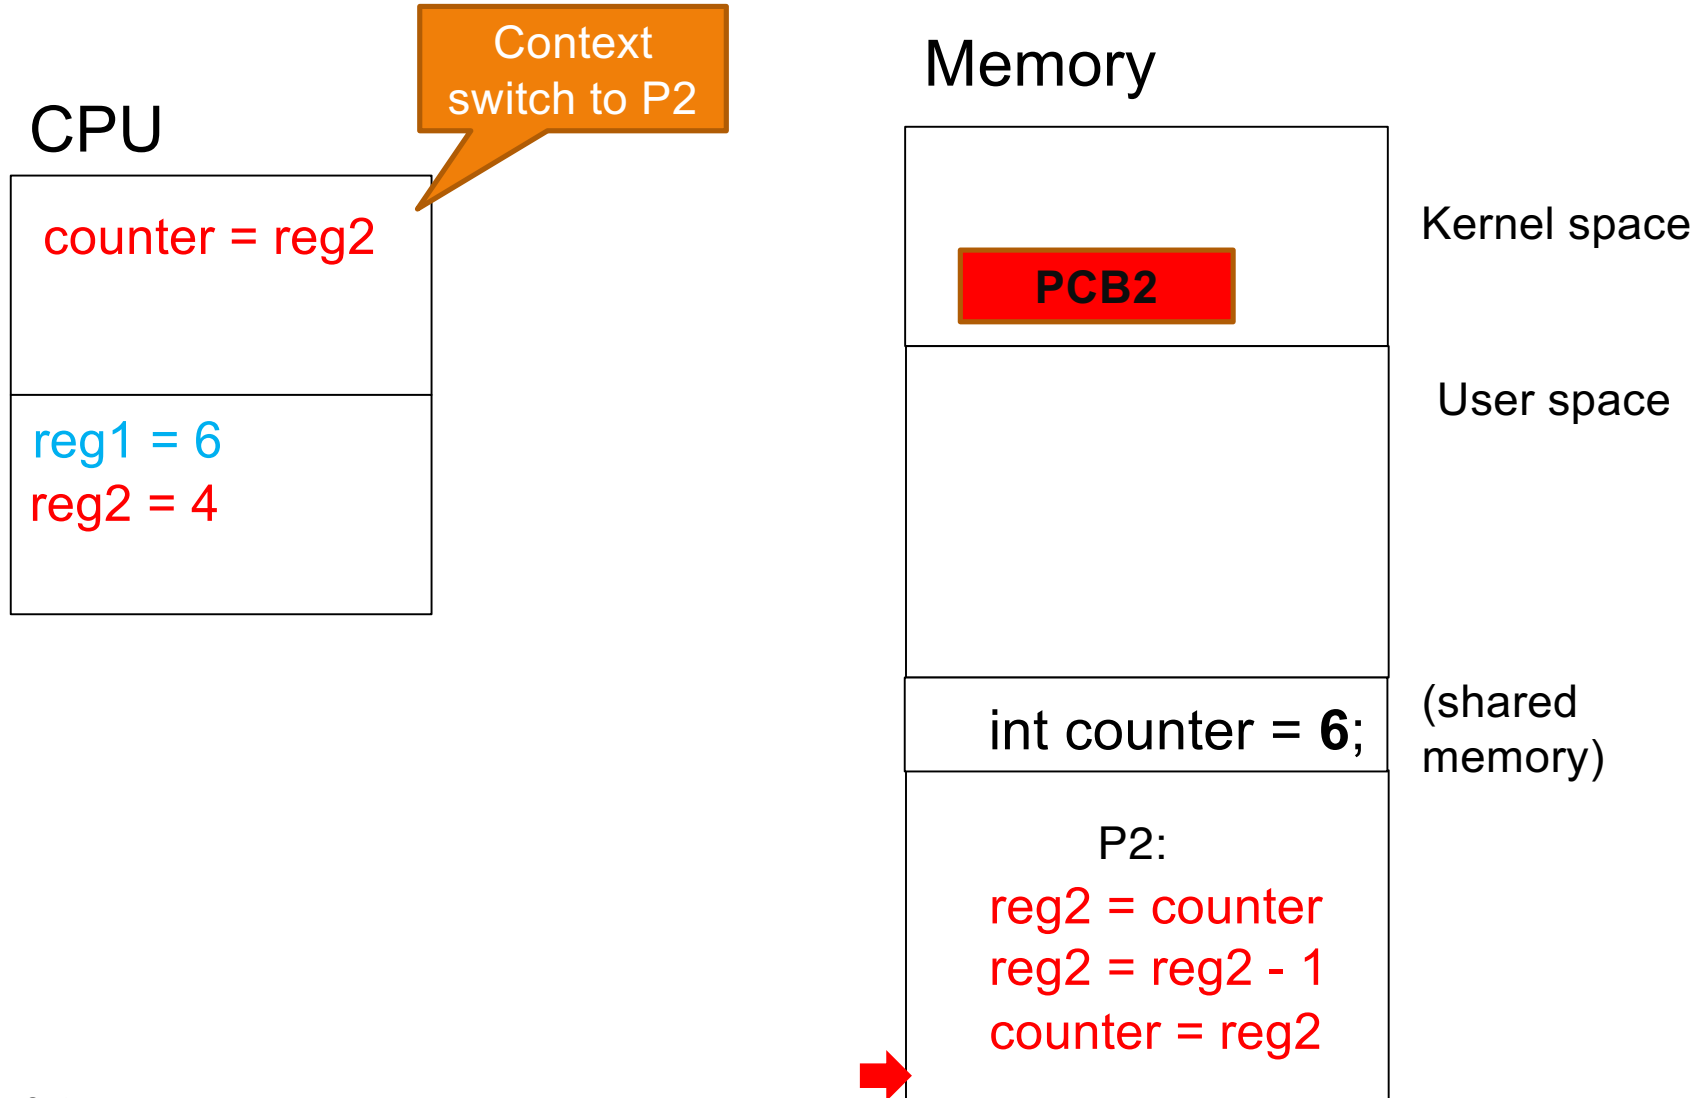
Example



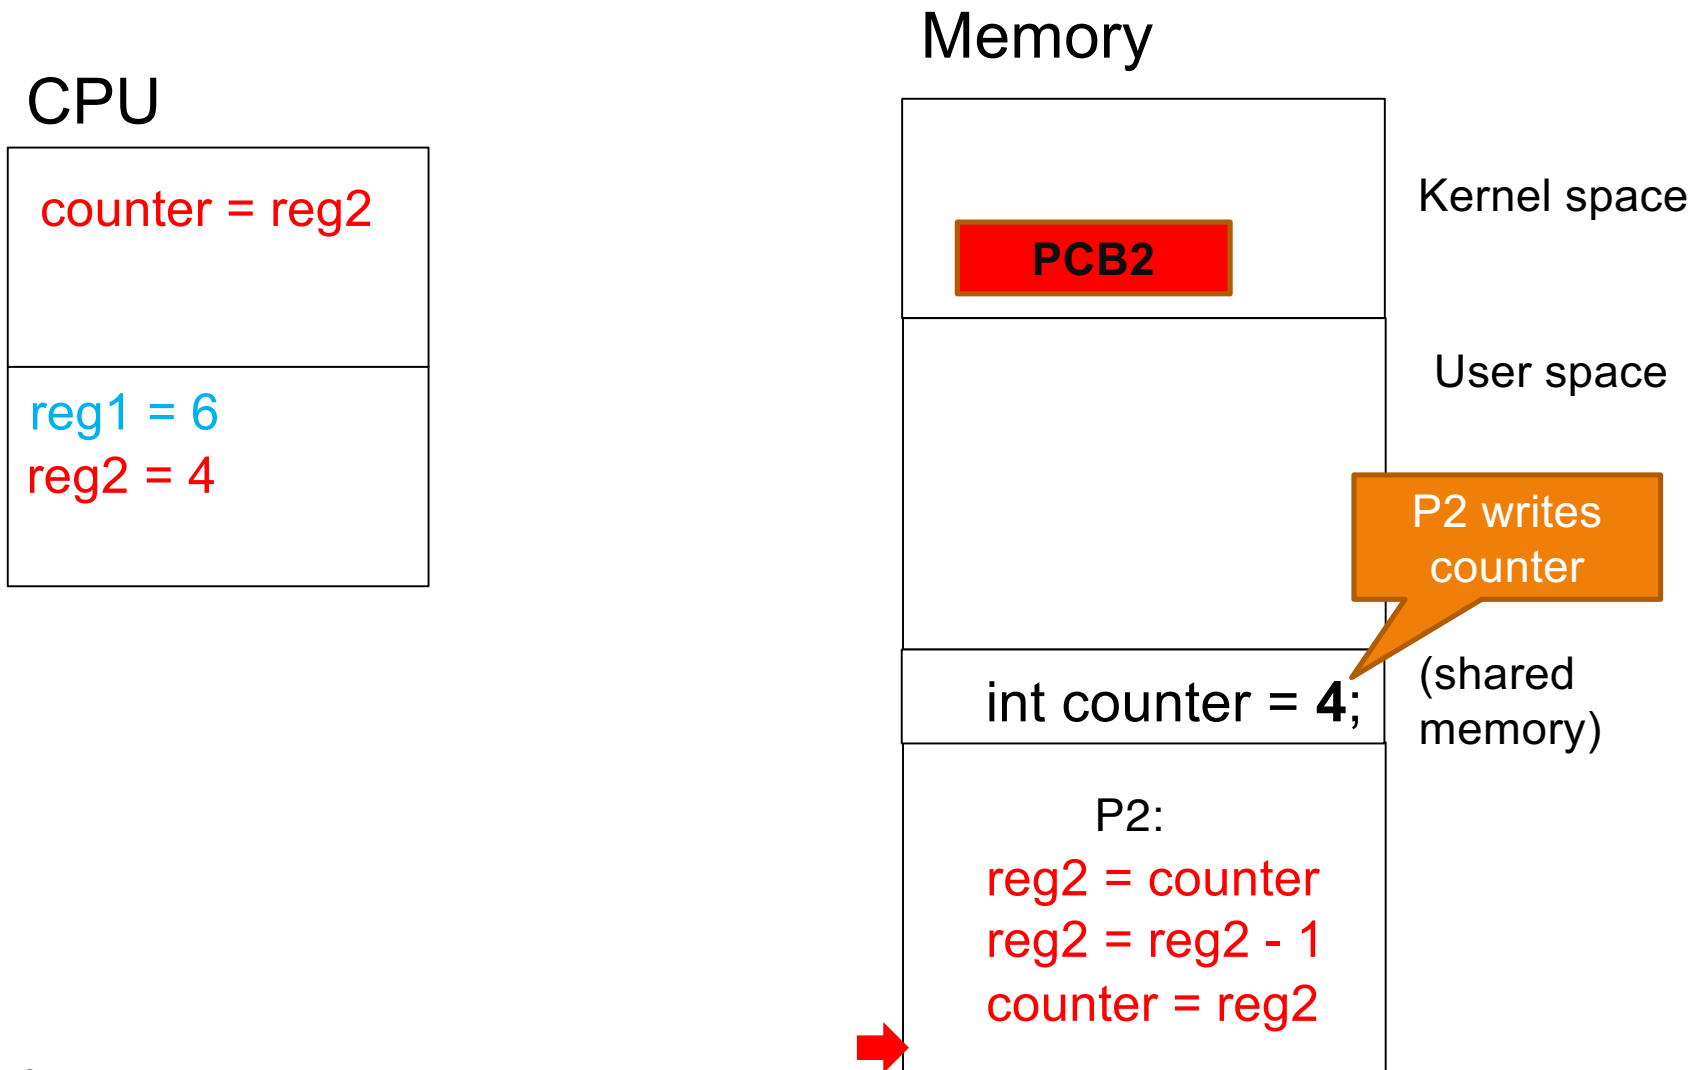
Example



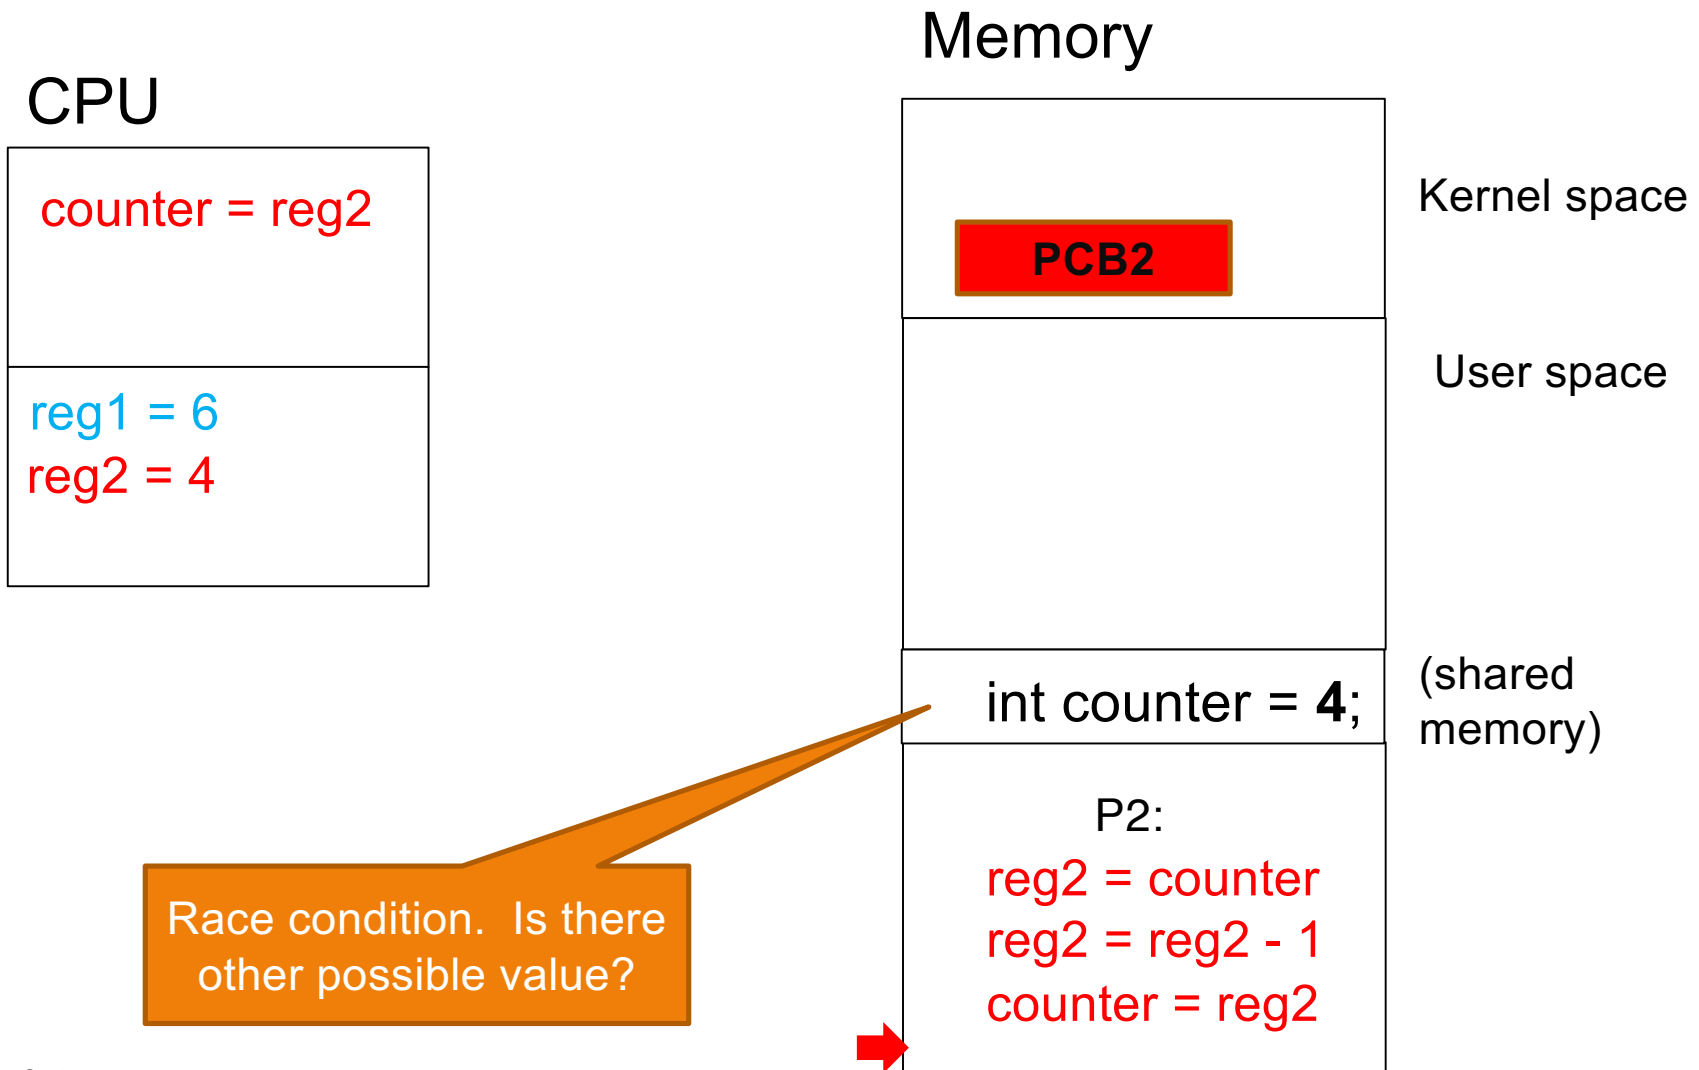
Example



Example



Example



Problems in the example

incCounter()

```
load  reg1, counter
add   reg1, reg1, 1
store reg1, counter
```

decCounter()

```
load  reg2, counter
sub   reg1, reg1, 1
store reg1, counter
```

- The final outcome depends on the interleaving of read and write operations to shared variable counter
- How many interleavings are possible between two processes? $4!/2*2 = 6$
 - inc-LD, inc-ST, dec-LD, dec-ST \rightarrow counter = 5
 - inc-LD, dec-LD, inc-ST, dec-ST \rightarrow counter = 4
 - inc-LD, dec-LD, dec-ST, inc-ST \rightarrow counter = 6
 - dec-LD, dec-ST, inc-LD, inc-ST \rightarrow counter = 5
 - dec-LD, inc-LD, dec-ST, inc-ST \rightarrow counter = 6
 - dec-LD, inc-LD, inc-ST, dec-ST \rightarrow counter = 4

Are they all
OK? No!

Race Conditions

- The system behavior depends on the exact interleaving
 - Systems that have race conditions are incorrect
 - The impact is context-dependent (could be catastrophic)
- Possibly huge number of interleaving scenarios!
 - Grows exponentially w.r.t. the number of processes, shared variables, and operations on shared variables
- Some interleavings are OK, some are not.
 - **How to avoid bad interleavings, while allowing as many good ones as possible? → Synchronization! (next week)**

POSIX Shared Memory in *nix

■ Basic steps of usage:

1. Create/locate a shared memory region **M**
2. Attach **M** to process memory space
3. Read from/Write to **M**
 - Written values visible to all process that share **M**
4. Detach **M** from memory space after use
5. Destroy **M**
 - Only one process need to do this
 - Allowed only if **M** is not attached to any process

Example: Master program (1/2)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int main()
{
    int shmid, i, *shm;

    shmid = shmget( IPC_PRIVATE, 40, IPC_CREAT | 0600 );

    if (shmid == -1){
        printf("Cannot create shared memory!\n");
        exit(1);
    } else
        printf("Shared Memory Id = %d\n", shmid);

    shm = (int*) shmat( shmid, NULL, 0 );
    if (shm == (int*) -1){
        printf("Cannot attach shared memory!\n");
        exit(1);
    }
}
```

The master program create the shared memory region and wait for the “worker” program to produce values before proceeding.

Step 1. Create Shared Memory region.

Step 2. Attach Shared Memory region.

Example: Master program (2/2)

```
shm[0] = 0;

while(shm[0] == 0){
    sleep(3);
}

for (i = 0; i < 3; i++){
    printf("Read %d from shared memory.\n", shm[i+1]);
}

shmdt( (char*) shm);
shmctl( shm, IPC_RMID, 0);

return 0;
}
```

The first element in the shared memory region is used as “control” value in this example (0: values not ready, 1: values ready).

The next 3 elements are values produced by the worker program.

Step 4+5. Detach and destroy Shared Memory region.

Example: Worker program

```
//similar header files
int main()
{   int shmid, i, input, *shm;

    printf("Shared memory id for attachment: ");
    scanf("%d", &shmid);

    shm = (int*)shmat( shmid, NULL, 0);
    if (shm == (int*)-1) {
        printf("Error: Cannot attach!\n");
        exit(1);
    }

    for (i = 0; i < 3; i++){
        scanf("%d", &input);
        shm[i+1] = input;
    }
    shm[0] = 1;

    shmdt( (char*)shm );
    return 0;
}
```

Step 1. By using the shared memory region id directly, we skip **shmget()** in this case.

Step 2. Attach to shared memory region.

Write 3 values into shm[1 to 3]

Let master program know we are done!

Step 4. Detach Shared Memory region.

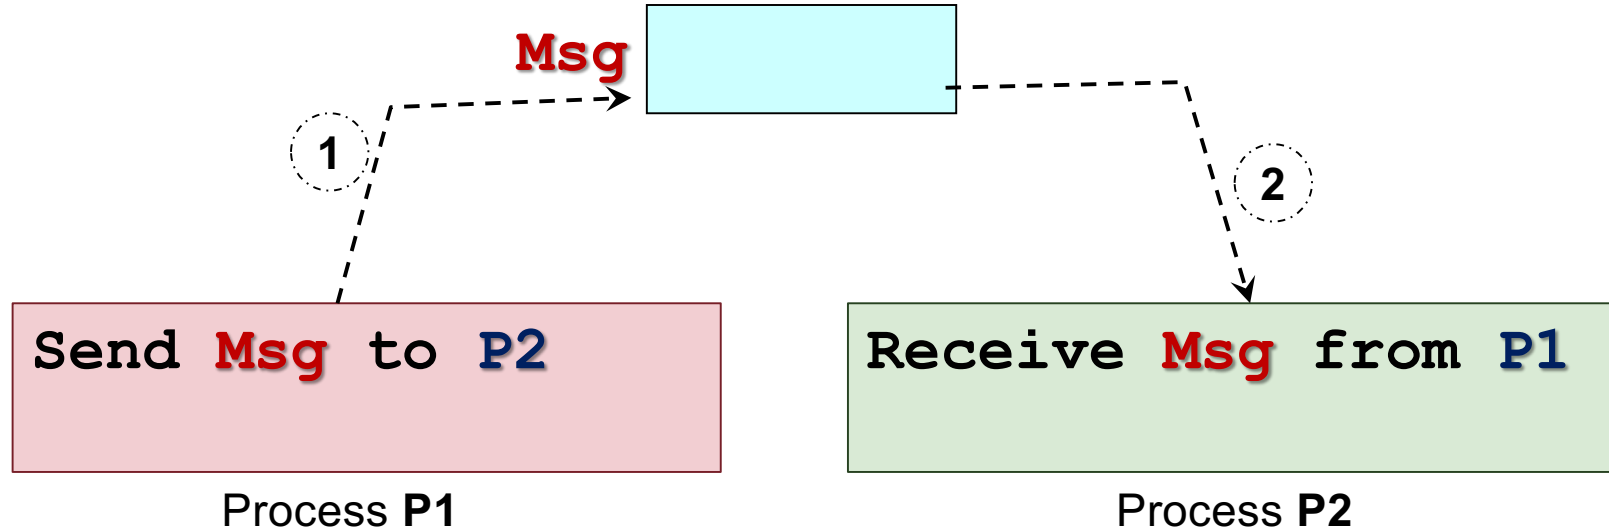
You have 1,023,428 messages waiting...

MESSAGE PASSING

Message Passing

- Explicit Communication through exchange of messages
- General Idea:
 - Process P_1 prepares a message M and sends it to Process P_2
 - Process P_2 receives the message M
 - Message sending and receiving are usually provided as system calls
- Additional properties:
 - **Naming**
 - How to identify the other party in the communication
 - **Synchronization**
 - Implicit: through the blocking behavior of the sending/receiving operations

Message Passing: Illustration



- The **Msg** have to be stored in kernel memory space
- All send/receive operations must go through OS (i.e., a system call)

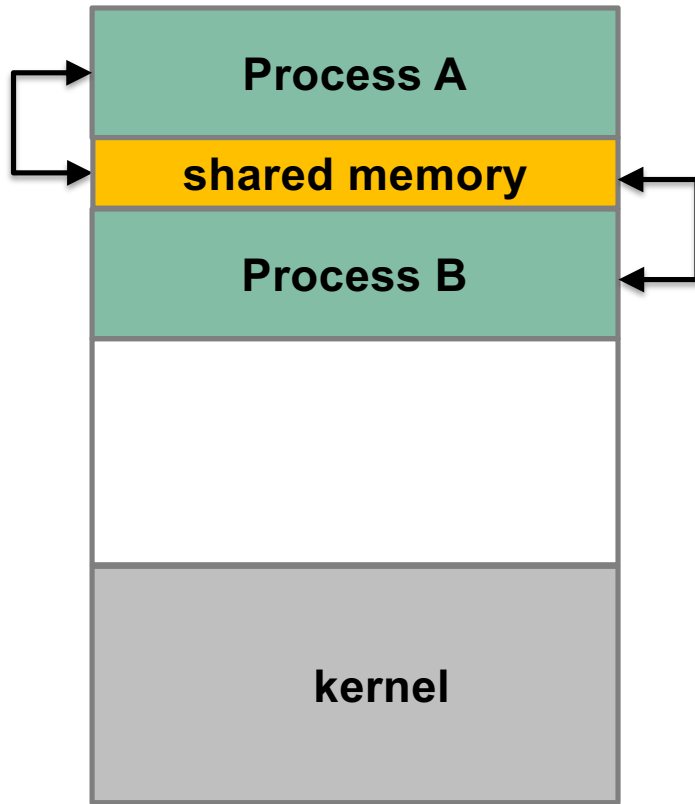
Naming Scheme: **Direct Communication**

- Sender/Receiver of message **explicitly name the other party**
- Example:
 - **Send** (P_2 , **Msg**) : Send Message **Msg** to Process P_2
 - **Receive** (P_1 , **Msg**) : Receive Message **Msg** from Process P_1
- Characteristics:
 - One link per pair of communicating processes
 - Need to know the identity of the other party

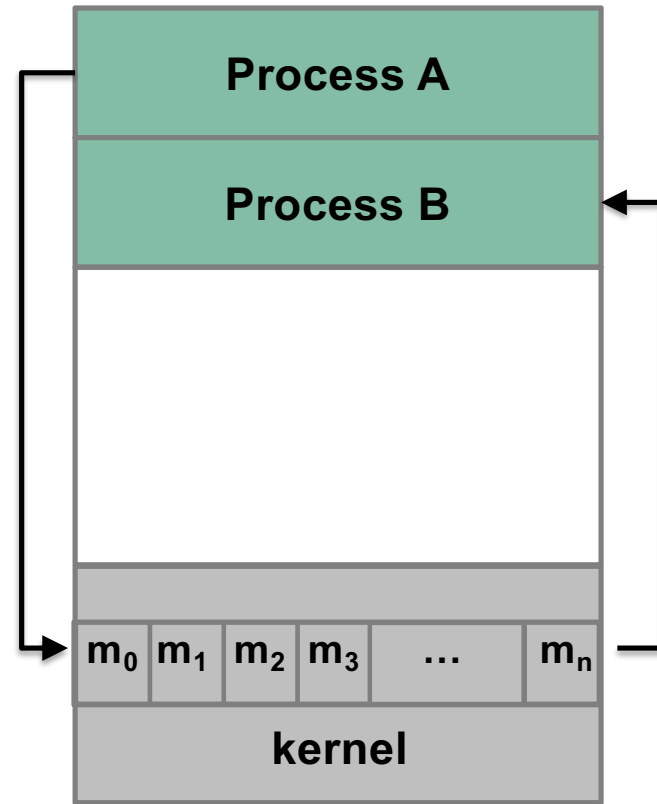
Naming Scheme: Indirect Communication

- Message are sent to / received from message storage:
 - Usually known as *mailbox* or *port*
- Example:
 - `Send(MB, Msg)` : Send Message `Msg` to Mailbox `MB`
 - `Receive(MB, Msg)` : Receive Message `Msg` from Mailbox `MB`
- Characteristics:
 - One mailbox can be shared among a number of processes

Shared Memory versus Message Passing



Shared memory



Message passing

Two Synchronization Behaviors

■ **Blocking Primitives** (Synchronous):

- ❑ Send(): Sender is blocked until the message is received
- ❑ Receive(): Receiver is blocked until a message has arrived

■ **Non-Blocking Primitives** (asynchronous):

- ❑ Send(): Sender resume operation immediately
- ❑ Receive(): Receiver either receive the message if available or some indication that message is not ready yet

Synchronization Model: Receiver

■ Blocking receive

- ❑ Most common
- ❑ Receiver must wait for a message it is not already available

■ Non-blocking receive

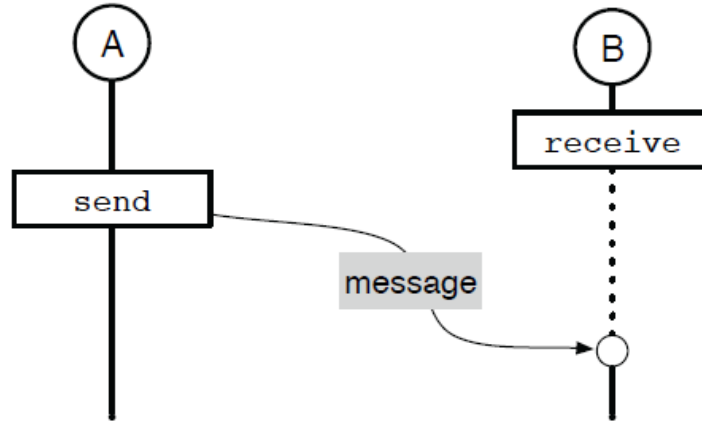
- ❑ Checks whether a message is available
- ❑ If message is available, retrieves it and moves on
- ❑ If message not available, continues without a message

■ We will assume blocking receive from now on

Non-blocking Send = Asynchronous Message Passing

- **Sender is never blocked** even if the receiver has not yet executed the matching `receive()`
- **System buffers the message** (up to certain capacity).
- `receive()` performed by the receiver later will be completed immediately.
- Asynchronous MP is generally a good choice, BUT:
 - ❑ Too much freedom for the programmer; program is complex to understand.
 - ❑ **Finite buffer size means system is not truly asynchronous** (sender waits when the buffer is full or returns with error).

Asynchronous Message Passing




Scenario 1:

Message sent after `receive()` call.

Sender always continues.

Receiver is blocked waiting for message



Scenario 2:

Message sent before `receive()` call.

Sender always continues.

`receive()` gets data immediately

— Executing
..... Blocked

Message Buffers

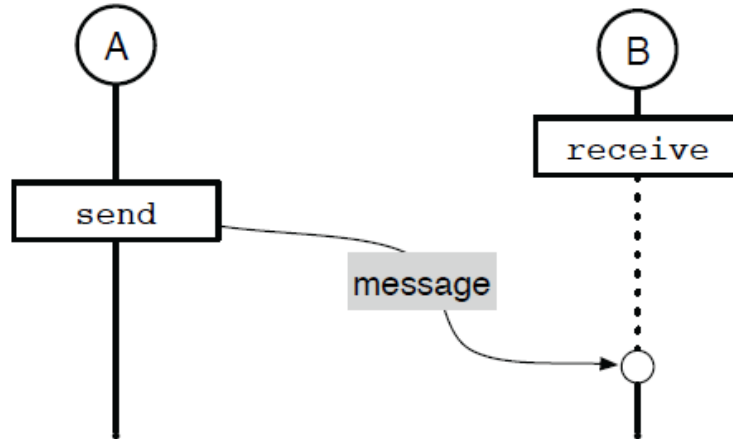
- Intermediate buffers between sender and receiver in asynchronous communication
- Under OS control → no synchronization necessary
- Large buffer decouples sender and receiver making them less sensitive to variations in execution; they do not wait for each other unnecessarily
- No amount of buffering helps when the sender is always faster than the receiver
- User needs to declare in advance the capacity of the mailbox mailbox is created
 - ❑ Sender blocks if the buffer is completely full OR
 - ❑ Sender returns immediately with an error

Synchronous Message Passing

- Also known as **rendezvous**
- **Sender** invoking `send()` **is blocked** till the receiver performs the matching `receive()`
- Sender forced to wait till the receiver is ready → **no intermediate buffering** required
 - Message can be kept by the sender till the receiver is ready and then directly copied into the receiver address space



Synchronous Message Passing




Scenario 1:

Message sent after receive() call

Sender does not have to block.

Receiver is blocked waiting for message



Scenario 2:

Message sent before receive() call

Sender is blocked.

Sender is blocked waiting for receiving.

— Executing
..... Blocked

Message Passing: Pros and Cons

■ Advantages:

- ❑ Portable:
 - Can easily be implemented on different processing environment, e.g. distributed system, wide area network etc
- ❑ Easier synchronization:
 - E.g. when synchronous primitive is used, sender and receiver are implicitly synchronized

■ Disadvantages:


- ❑ Inefficient:
 - Usually require OS intervention
- ❑ Harder to use:
 - Message are usually limited in size and/or format

Plumber needed! Leaking pipes all around!

UNIX PIPES

Pipes: Communication channels


- In Unix, a process has 3 default communication channels:



- Example:
 - In a typical C program, `printf()` uses `stdout`, `scanf()` uses `stdin`.

Piping in Shell

- Unix shell provides the “|” symbol to link the input/output channels of one process to another, this is known as **piping**
- For example (“ A | B ”):



- The output of A (instead of going to screen) directly goes into B as input (as if it come from keyboard)


Unix Pipes

- One of the earliest IPC mechanism
- General Idea:
 - A communication channel is created with 2 ends:
 - 1 end for reading, the other for writing
 - Just like a water pipe in the real world



- The piping “|” in shell is achieved using this mechanism internally

Unix Pipes: as a IPC Mechanism



- A pipe can be shared between two processes
- A form of Producer-Consumer relationship
 - ❑ P produces (writes) **n** bytes
 - ❑ Q consumes (reads) **m** bytes
- Behavior:
 - ❑ Like an anonymous file
 - ❑ FIFO → must access data in order

Unix Pipes: Semantic


- Pipe functions as **circular bounded byte buffer** with **implicit synchronization**:
 - Writers **wait** when buffer is **full**
 - Readers **wait** when buffer is **empty**
- Variants:
 - Can have multiple readers/writers
 - The normal shell pipe has 1 writer and 1 reader
 - Depends on Unix version, pipes may be **half-duplex**
 - **unidirectional**: with one write end and one read end
 - Or **full-duplex**
 - **bidirectional**: any end for read/write

Unix Pipe: System Calls

Header File	<code>#include <unistd.h></code>
Syntax	<code>int <i>pipe</i>(int fd[])</code>

■ Returns:

- ❑ 0 to indicate success; !0 for errors
- ❑ An array of file descriptors is returned:
 - `fd[0]` == reading end
 - `fd[1]` == writing end



Unix Pipes: Example Code

```
#define READ_END 0
#define WRITE_END 1

int main()
{
    int pipeFd[2], pid, len;
    char buf[100], *str = "Hello There!";

    pipe( pipeFd );
    if ((pid = fork()) > 0) { /* parent */
        close(pipeFd[READ_END]);
        write(pipeFd[WRITE_END], str, strlen(str)+1);
        close(pipeFd[WRITE_END]);
    } else { /* child */
        close(pipeFd[WRITE_END]);
        len = read(pipeFd[READ_END], buf, sizeof buf);
        printf("Proc %d read: %s\n", pid, buf);
        close(pipeFd[READ_END]);
    }
}
```

Unix Pipes: More to explore

- It is possible to:
 - ❑ Attach/change the standard communication channels (**`stdin`**, **`stdout`**, **`stderr`**) to one of the pipes
 - ➔ Redirect the input/output from one program to another!
- Unix system calls to explore:
 - ❑ **`dup()`**
 - ❑ **`dup2()`**
- Wikipedia article on **`dup()`** system call has a great program example

pssst! pssst!

UNIX SIGNAL

Unix Signal: Quick Overview

- A form of inter-process communication
 - An asynchronous notification regarding an event
 - Sent to a process/thread
- The recipient of the signal must handle the signal by:
 - A default set of handlers OR
 - User supplied handler (only applicable to some signals)
- Common signals in Unix:
 - Kill, Stop, Continue, Segmentation Fault Signal....

Example: Custom Signal Handler

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>

void myOwnHandler( int signo )
{
    if (signo == SIGSEGV){
        printf("Memory access blows up!\n");
        exit(1);
    }
}

int main(){
    int *ip = NULL;

    if (signal(SIGSEGV, myOwnHandler) == SIG_ERR)
        printf("Failed to register handler\n");

    *ip = 123;

    return 0;
}
```

User defined function to handle signal. In this example, we handle the "SIGSEGV" signal, i.e. the memory segmentation fault signal.

Register our own code to replace the default handler.

This statement will cause a segmentation fault.

Summary

- Common Inter Process Communication mechanisms:
 - ❑ Shared Memory
 - POSIX example
 - ❑ Message Passing
 - ❑ Unix Pipes
 - ❑ Unix Signals

References

- Modern Operating Systems (4th Edition)
 - Chapters 2.3.1 and 2.3.8
- Operating System Concepts (9th Edition)
 - Chapters 3.4 and 3.5
- Operating Systems: Three Easy Pieces
 - None!