

Process Management

Process Alternative - Threads

Lecture 5

Overview

■ Threads

- ❑ Motivation
- ❑ Basic Idea

■ Threads models

- ❑ Kernel vs. User Thread
- ❑ Hybrid model

■ Thread in Unix

- ❑ POSIX thread:
 - Create, Exit, Synchronization
 - Exploration

Motivation for Threads

- Process is expensive:
 - ❑ Process creation under the `fork()` model:
 - Duplicate memory space
 - Duplicate most of the process context etc
 - ❑ Context switch:
 - Requires saving/restoration of process Information
- It is hard for independent processes to communicate with each other:
 - ❑ Independent memory space → No easy way to pass information
 - ❑ Requires Inter-Process Communication (IPC)

Motivation for Thread (cont)

- Thread is invented to overcome the problems with process model
 - ❑ Started out as a "quick hack" and eventually matured to be very popular mechanism
- **Basic Idea:**
 - ❑ A traditional process has a single thread of control:
 - Only one instruction of the whole program is executing at any time
 - ❑ We "simply" add more threads of control to the same process:
 - Multiple parts of the programs is executing at the same time conceptually

Threads of control: Illustration



- Suppose we are preparing lunch, which consists of the following tasks:
 - ❑ Steam rice
 - ❑ Fry fish
 - ❑ Cook soup
- A pseudo-C program:

```
int main()  
{  
    steamRice( twoBowls );  
    fryFish ( bigFish );  
    cookSoup( cornSoup );  
  
    printf( "Lunch READY!!\n" );  
    return 0;  
}
```

Threads of control: Illustration (cont)

- Process with a single thread will go through the functions sequentially


```
int main()  
{  
    steamRice( twoBowls );  
    fryFish( bigFish );  
    cookSoup( cornSoup );  
  
    printf( "Lunch Ready\n" );  
    return 0;  
}
```




- Suppose the tasks are independent:
 - Lets try to have multiple threads of control

Multiple "threads of control" with **fork()**

```
int result;  
result = fork();  
if (result != 0){  
    result = fork();  
    if (result != 0){  
        steamRice( twoBowls );  
    } else {  
        fryFish( bigFish );  
    }  
} else {  
    cookSoup( cornSoup );  
}  
... ..
```



```
int result;  
result = fork();  
if (result != 0){  
    result = fork();  
    if (result != 0){  
        steamRice( twoBowls );  
    } else {  
        fryFish( bigFish );  
    }  
} else {  
    cookSoup( cornSoup );  
}  
... ..
```




We effectively
have two
"threads" of
control after
fork()

New approach: Multi-Threading

```
int main()
{
    Thread 1 do { steamRice( twoBowls ); }
    Thread 2 do { fryFish( bigFish ); }
    Thread 3 do { cookSoup( cornSoup ); }

    Wait for all threads to finish;

    printf( "Lunch Ready\n" );
    return 0;
}
```



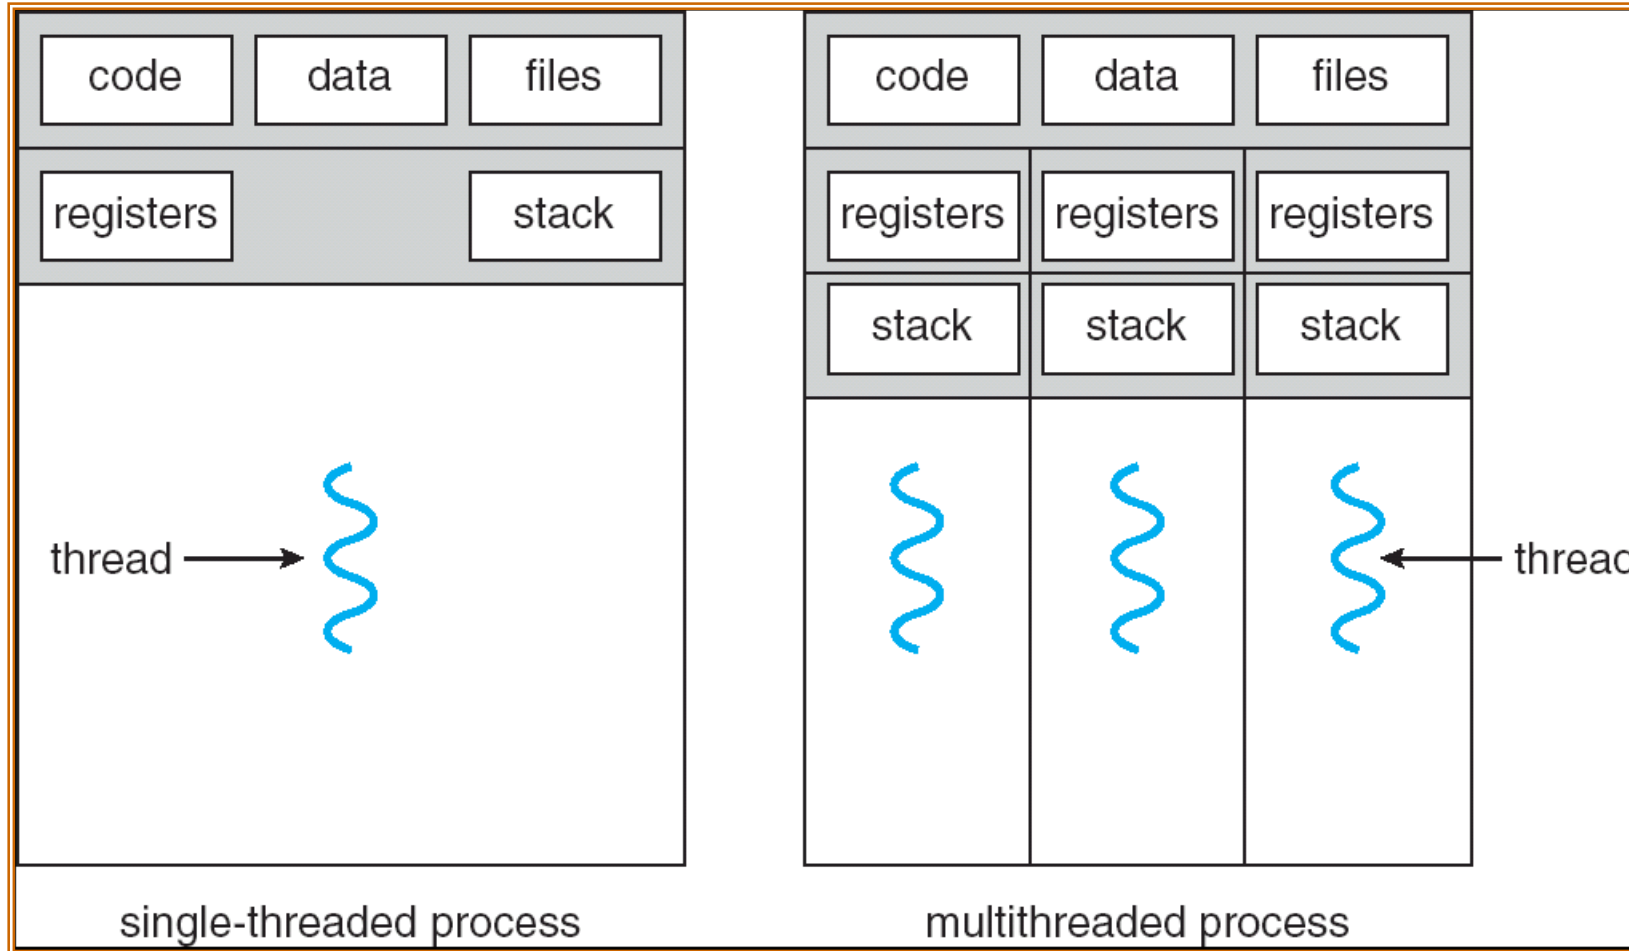
■ Note:

- ❑ Pseudocode 😊
- ❑ The three threads are concurrent

Process and Thread

- A single process can have multiple threads:
 - known as **multithreaded process**
- Threads in the same process shares:
 - **Memory Context**: Text, Data, Heap
 - **OS Context**: Process id, other resources like files, etc
- Unique information needed for each thread:
 - Identification (usually **thread id**)
 - Registers (General purpose and special)
 - “Stack” (more about this later)

Process and thread: Illustration



- Taken from Operating System Concepts (7th Edition) by Silberschatz, Galvin & Gagne, published by Wiley

Process Context Switch VS Thread Switch

- Process context switch involves:
 - ❑ OS Context
 - ❑ Hardware Context
 - ❑ Memory Context
- Thread switch within the same process involves:
 - ❑ Hardware context:
 - Registers
 - "Stack" (actually just changing FP and SP registers)
- Thread is much “lighter” than process:
 - ❑ a.k.a. ***lightweight process***

Threads: Benefits

Economy:

- Multiple threads in the same process requires much less resources to manage compared to multiple processes

Resource sharing:

- Threads share most of the resources of a process
- No need for additional mechanism for passing information around

Responsiveness:

- Multithreaded programs can appear much more responsive

Scalability:

- Multithreaded program can take advantage of multiple CPUs

Thread: **Problems**

■ **System Call Concurrency**

- Parallel execution of multiple threads → parallel system call possible
 - Have to guarantee correctness and determine the correct behavior

■ **Process Behavior:**

- Impact on process operations
- Example:
 - `fork()` duplicate process, how about threads?
 - If a single thread executes `exit()`, how about the whole process?
 - If a single thread calls `exec()`, how about other threads?

Thread Models

What are the ways to support threads?

Two Major Thread Implementations


■ User Thread:

- ❑ Thread is implemented as a **user library**
 - A runtime system (in the process) will handle thread related operation
- ❑ Kernel is **not aware** of the threads in the process

■ Kernel Thread:

- ❑ Thread is implemented in the OS
 - Thread operation is handled as system calls
- ❑ Thread-level scheduling is possible:
 - Kernel schedule by threads, instead of by process
- ❑ Kernel may make use of threads for its own execution

User Thread: Illustration



- Adapted from Modern Operating System Concepts (3rd Edition) by Andrew Tanenbaum, published by Pearson

User Thread: Pros and Cons


■ Advantages:

- ❑ Can have multithreaded program on ANY OS
- ❑ Thread operations are just library calls
- ❑ Generally more configurable and flexible
 - e.g. Customized thread scheduling policy

■ Disadvantages:

- ❑ OS is not aware of threads, scheduling is performed at process level
 - One thread blocked → Process blocked → All threads blocked
 - Cannot exploit multiple CPUs!

Kernel Thread: Illustration



- Adapted from Modern Operating System Concepts (3rd Edition) by Andrew Tanenbaum, published by Pearson

Kernel Threads: Pros and Cons

■ **Advantages:**


- ❑ Kernel can schedule on thread levels:
 - More than 1 thread in the same process can run simultaneously on multiple CPUs

■ **Disadvantages:**


- ❑ Thread operations is now a system call!
 - Slower and more resource intensive
- ❑ Generally less flexible:
 - Used by all multithreaded programs
 - If implemented with many features:
 - ❑ Expensive, overkill for simple program
 - If implemented with few features:
 - ❑ Not flexible enough for some programs

Hybrid Thread Model

- Have both Kernel and User threads
 - ❑ OS schedule on kernel threads only
 - ❑ User thread can bind to a kernel thread
- Offer great flexibility
 - ❑ Can limit the concurrency of any process / user



Hybrid Model Example: Solaris



Threads on Modern Processor

- Threads started off as a software mechanism
 - User space library → OS aware mechanism
- There are hardware support on modern processors:
 - Essentially supply multiple sets of registers (GPRs, and special registers) to allow threads to run natively and in parallel on the same core
 - known as **simultaneous multi-threading (SMT)**
- Example:
 - **Hyperthreading** on Intel processor

POSIX Threads

A popular thread API

POSIX Threads: **pthread**

- Standard defined by the IEEE
 - Supported by most Unix variants
- Defines the API as well as the behavior
 - However, **implementation** is not specified
 - So, **pthread** can be implemented as user thread (Windows) or kernel thread (Linux)
- Will show a few example to highlight the differences between process and thread only

Basics of pthread

- Header File:

- ❑ `#include <pthread.h>`

- Compilation (flag is system dependent):

- ❑ `gcc XXXX.c -lpthread`

- Useful datatypes:

- ❑ `pthread_t` : Data type to represent a thread id (TID)

- ❑ `pthread_attr`: Data type to represents attributes of a thread

pthread Creation Syntax

```
int pthread_create(  
    pthread_t* tidCreated,  
    const pthread_attr_t* threadAttributes,  
    void* (*startRoutine) (void*),  
    void* argForStartRoutine );
```

- Returns (0 = success; != 0 = errors)
- Parameters:
 - ❑ **tidCreated**: Thread Id for the created thread
 - ❑ **threadAttributes**: Control the behavior of the new thread
 - ❑ **startRoutine**: Function pointer to the function to be executed by thread
 - ❑ **argForStartRoutine**: Arguments for the `startRoutine` function

pthread Termination Syntax

```
int pthread_exit( void* exitValue );
```

- Parameters:
 - **exitValue**: Value to be returned to whoever synchronizes with this thread (more later)
- If *pthread_exit*() is not used, a *pthread* will terminate automatically at the end of the **startRoutine**
 - If a “return XYZ;” statement is used, then “XYZ” is captured as the exitValue
 - Otherwise, the exitValue is not well defined

pthread Creation & Termination: Example

```
//header files not shown
```

```
void* sayHello( void* arg )  
{  
    printf("Just to say hello!\n");  
    pthread_exit( NULL );  
}
```

Function to be executed
by a pthread

Pthread Termination

```
int main()  
{
```

```
    pthread_t tid;
```

Pthread Creation

```
    pthread_create( &tid, NULL, sayHello, NULL );  
    printf("Thread created with tid %i\n", tid);
```

```
    return 0;
```

```
}
```

pthread: Sharing of memory space

```
//header files not shown
```

```
int globalVar;
```

Variable shared between pthreads

```
void* doSum( void* arg)
```

```
{    int i;
```

```
    for (i = 0; i < 1000; i++)
```

```
        globalVar++;
```

```
}
```

Using a shared variables

```
int main()
```

```
    pthread_t tid[5];    //5 threads id
```

```
    int i;
```

```
    for (i = 0; i < 5; i++)
```

```
        pthread_create( &tid[i], NULL, doSum, NULL );
```

```
    printf("Global variable is %i\n", globalVar);
```

```
    return 0;
```

```
}
```

- What is the sum that we get (or should get)?

pthread Simple Synchronization - Join

```
int pthread_join( pthread_t threadID,  
                  void **status );
```

- To wait for the termination of another *pthread*:
- Returns(0 = success; !=0 = errors)
- Parameters:
 - **threadID**: TID of the *pthread* to wait for
 - **status**: Exit value returned by the target *pthread*

Pthread: Sharing of memory space V2.0

```
//header files not shown
int globalVar;

void* doSum( void* arg)
{ //same as before }

int main()
    pthread_t tid[5];    //5 threads id
    int i;
    for (i = 0; i < 5; i++)
        pthread_create( &tid[i], NULL, doSum, NULL );

    //Wait for all threads to finish
    for (i = 0; i < 5; i++)
        pthread_join( tid[i], NULL );

    printf("Global variable is %i\n", globalVar);
    return 0;
}
```

**Pthread
Synchronization**

Pthread: A lot more!

- There are more interesting stuff about ***pthread***:
 - ❑ Yielding (giving up CPU voluntarily)
 - ❑ Advanced synchronization
 - ❑ Scheduling policies
 - ❑ Binding to kernel threads
 - ❑ Etc.
- As we cover new topics, you can explore the ***pthread*** library to see the application!

Summary

- Thread:
 - Motivation
 - Difference between thread and process
- Thread Models:
 - User vs Kernel thread
- Simple introduction to ***pthread*** library

References

- Modern Operating System (4th Edition)
 - Chapter 2.2
- Operating System Concepts (9th Edition)
 - Chapter 4
- Operating Systems: Three Easy Pieces
 - Chapters 26 and 27
- Linux Pthread Man Pages
 - “man pthread.....” for more