

File System Management

# File System Implementations

Lecture 12

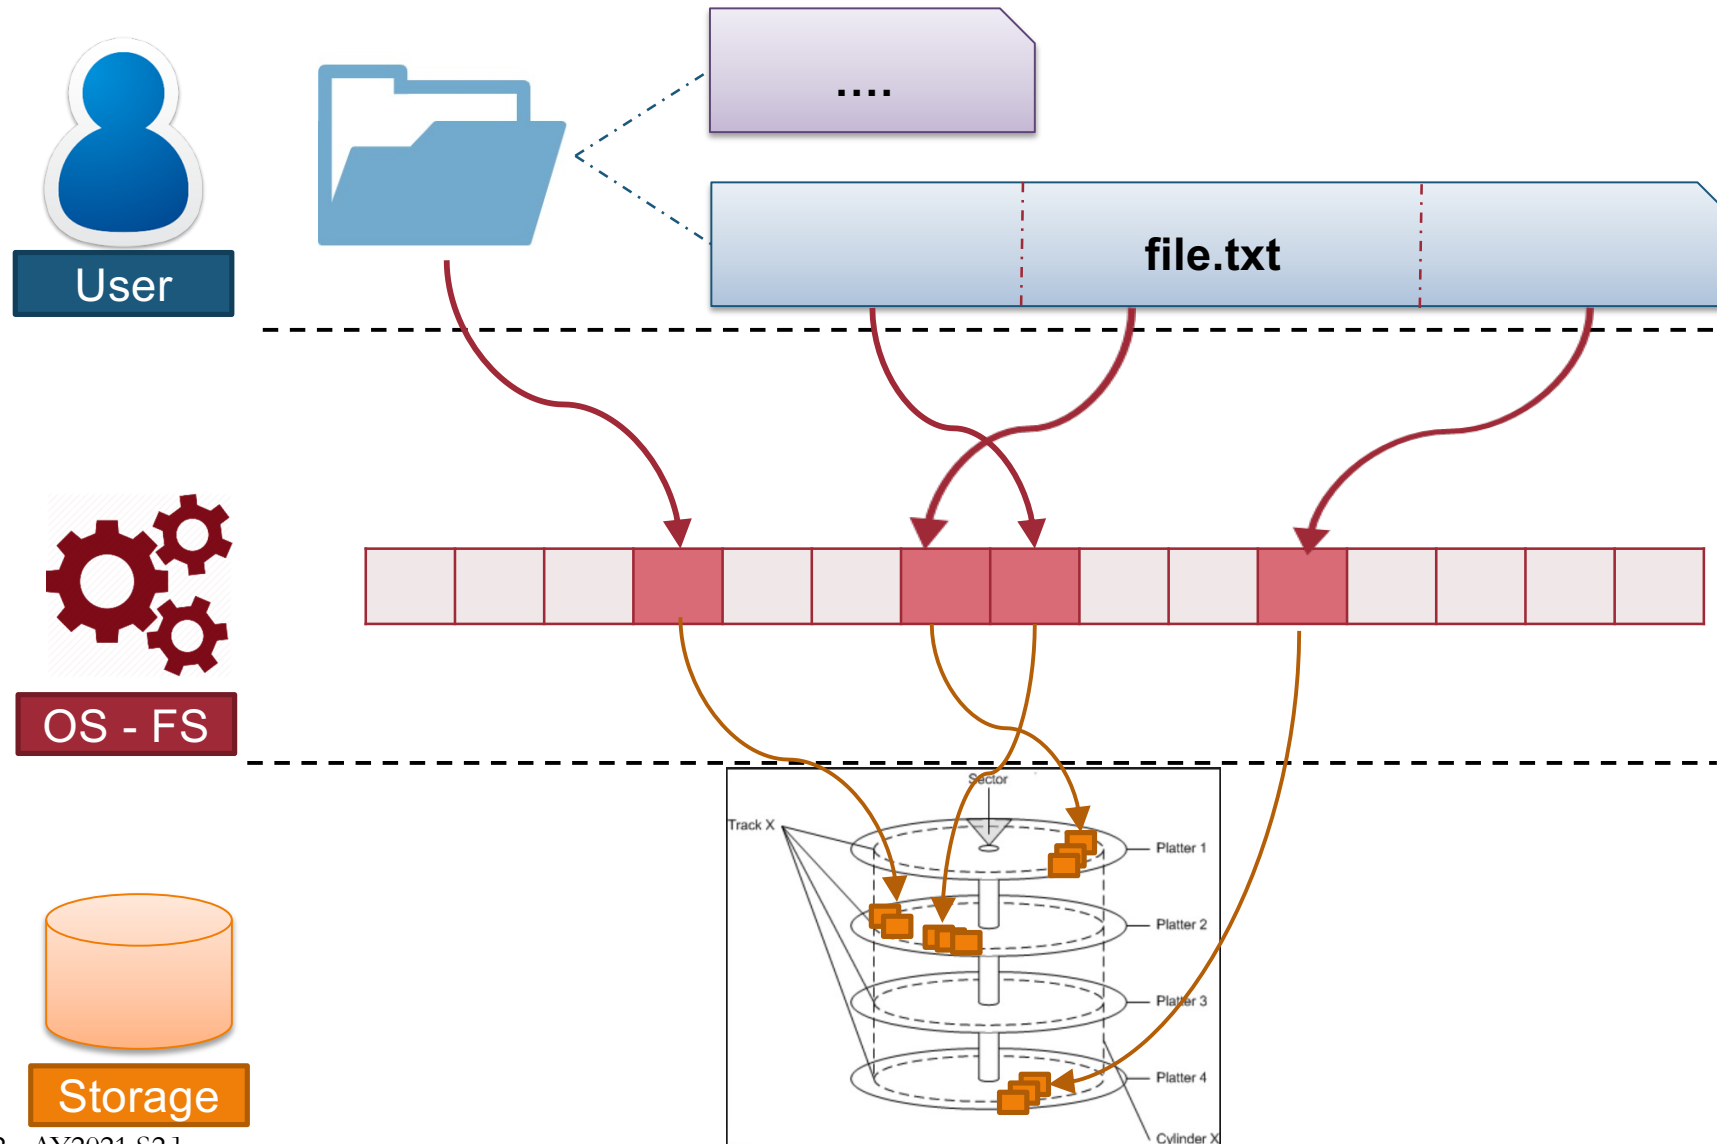
# Overview

- File System Implementation:
  - File system layout
  - Disk organization
- Implementation details for:
  - File Information
  - Free Space Management
  - Directory Structure
- File System in Action
- Disk I/O Scheduling

# File System Implementation: Overview

- File systems are stored on storage media:
  - e.g., Hard disk, CD/DVD, SRAM etc
- Concentrate on hard disk in this lecture
  - Though the ideas are generally applicable
- General **Disk Structure**:
  - Can be treated as a 1-D array of **logical blocks**
  - Logical block:
    - Smallest accessible unit (Usually 512-bytes to 4KB)
  - Logical block is mapped into **disk sector(s)**
    - Layout of disk sector is **hardware dependent**

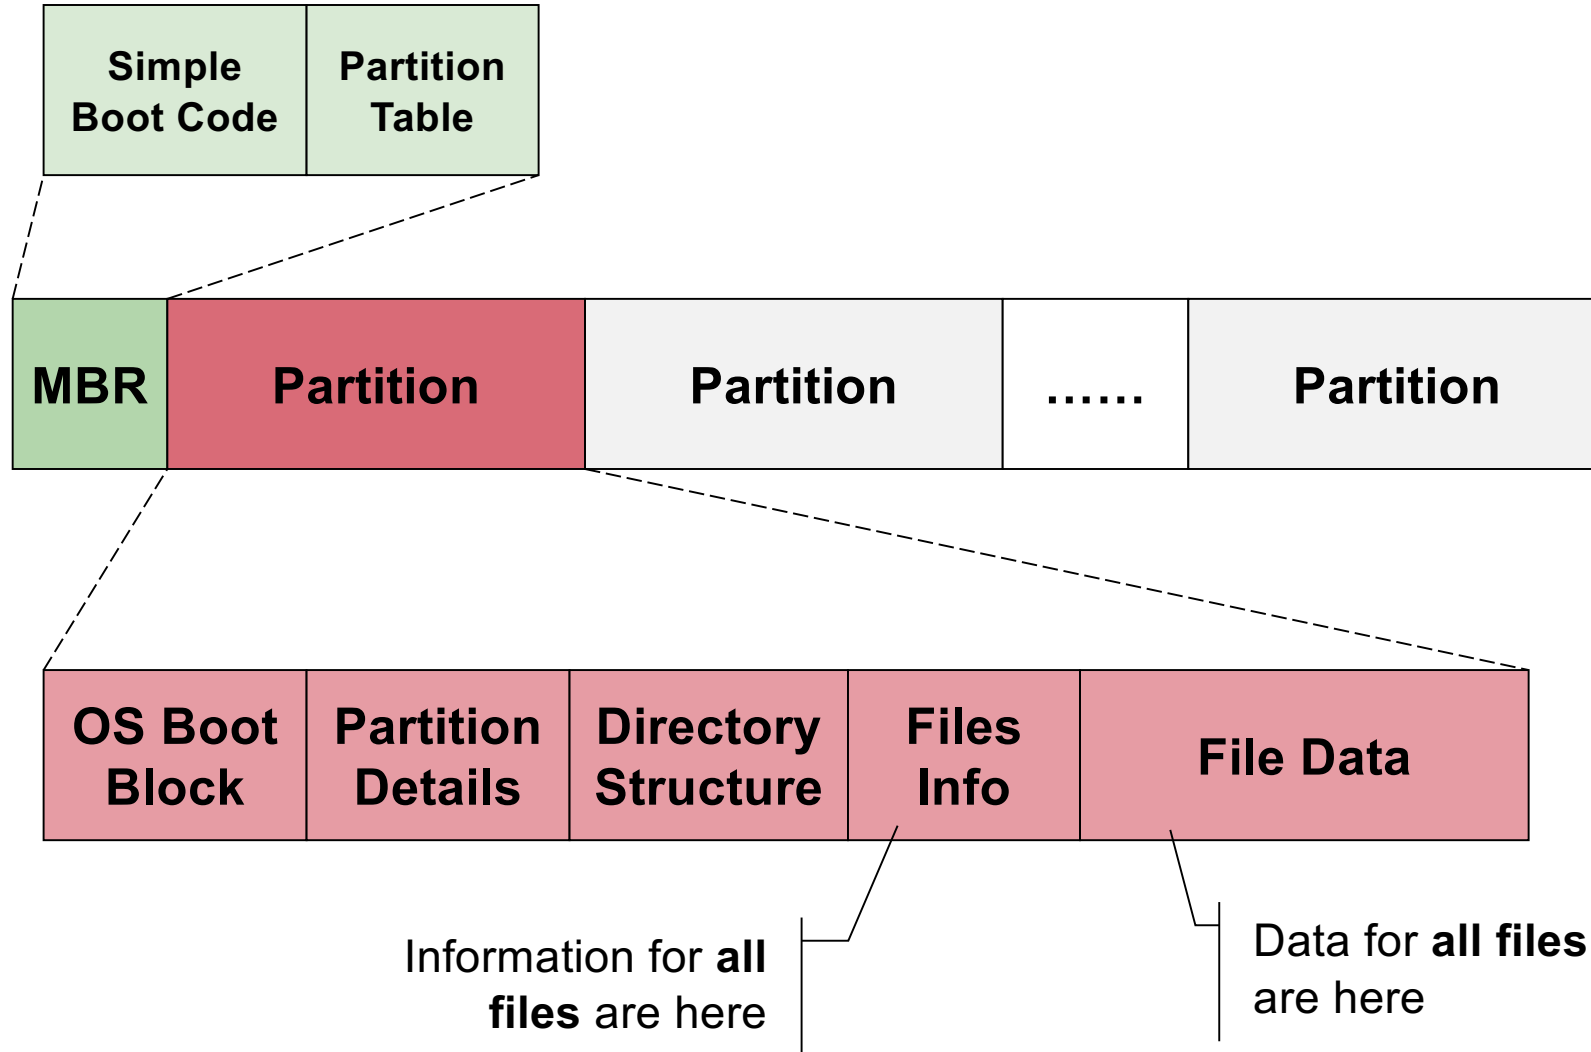
# User $\leftrightarrow$ OS $\leftrightarrow$ Hardware: Views



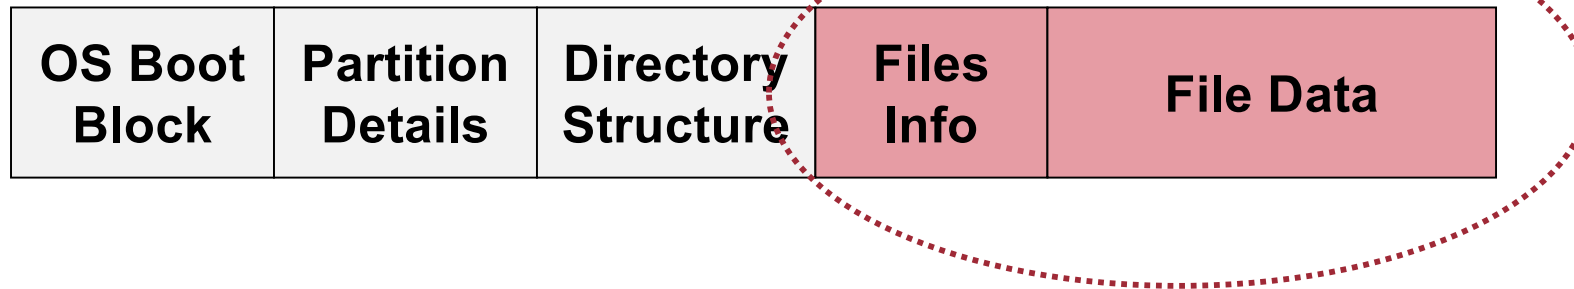
# Disk Organization: Overview

- Disk organization:
  - ❑ **Master Boot Record (MBR)** at sector 0 with partition table
  - ❑ Followed by one or more **partitions**
    - Each partition can contains an independent **file system**
- A file system generally contains:
  - ❑ OS Boot-Up information
  - ❑ Partition details:
    - Total Number of blocks
    - Number and location of free disk blocks
  - ❑ Directory Structure
  - ❑ Files Information
  - ❑ Actual File Data

# Generic Disk Organization: Illustration



# Implementing File



# File Implementation: Overview

- Logical view of a file:
  - A collection of logical blocks
- When file size  $\neq$  multiple of logical blocks
  - Last block may contain wasted space
  - i.e. **internal fragmentation**
- A good file implementation must:
  - Keep track of the logical blocks
  - Allow efficient access
  - Disk space is utilized effectively
- Basically focuses on **how to allocate** file data on disk



# File Block Allocation 1: **Contiguous**

## ■ **General Idea:**

- ❑ Allocate consecutive disk blocks to a file

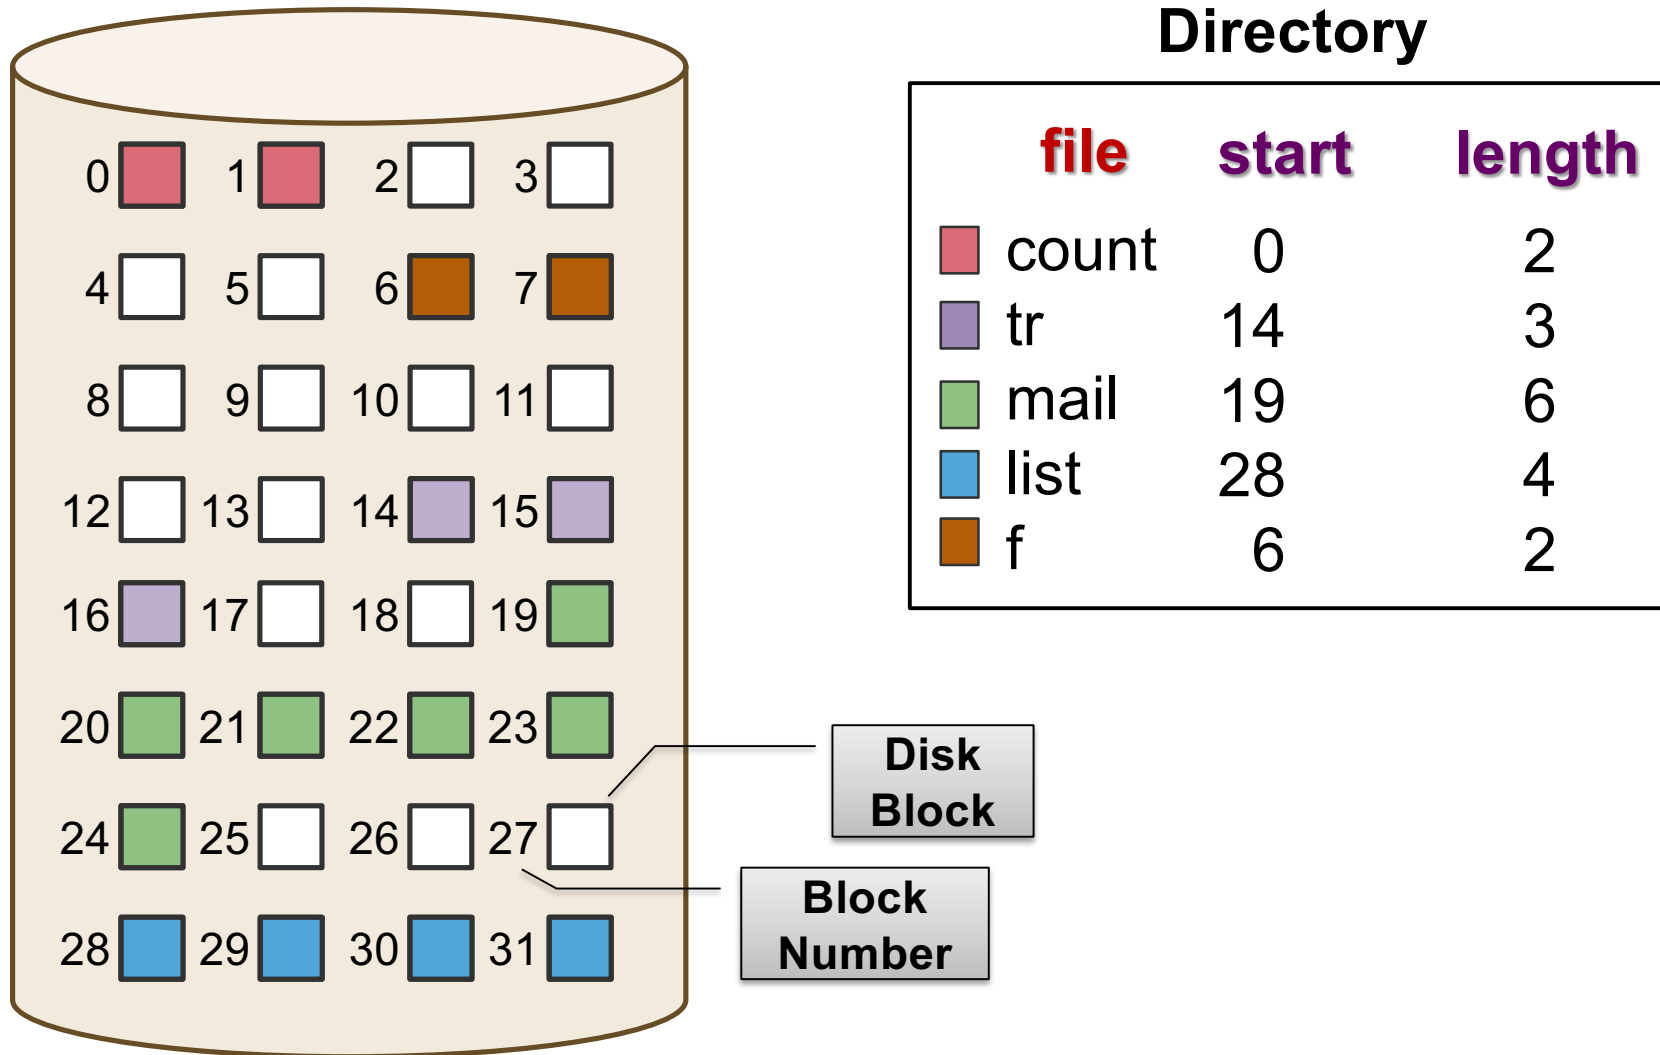
## ■ **Pros:**

- ❑ Simple to keep track:
  - Each file only needs: Starting block number + Length
- ❑ Fast access (only need to seek to first block)

## ■ **Cons:**

- ❑ **External Fragmentation**
  - Think of each file as a variable-size "partition"
  - Over time, with file creation/deletion, disk can have many small "holes"
- ❑ File size need to be specified in advance

# Contiguous Block Allocation



# File Block Allocation 2: **Linked List**

## ■ **General Idea:**

- ❑ Keep a linked list of disk blocks
- ❑ Each disk block stores:
  - The next disk block number (i.e. act as **pointer**)
  - Actual file data
- ❑ File information stores:
  - First and last disk block number

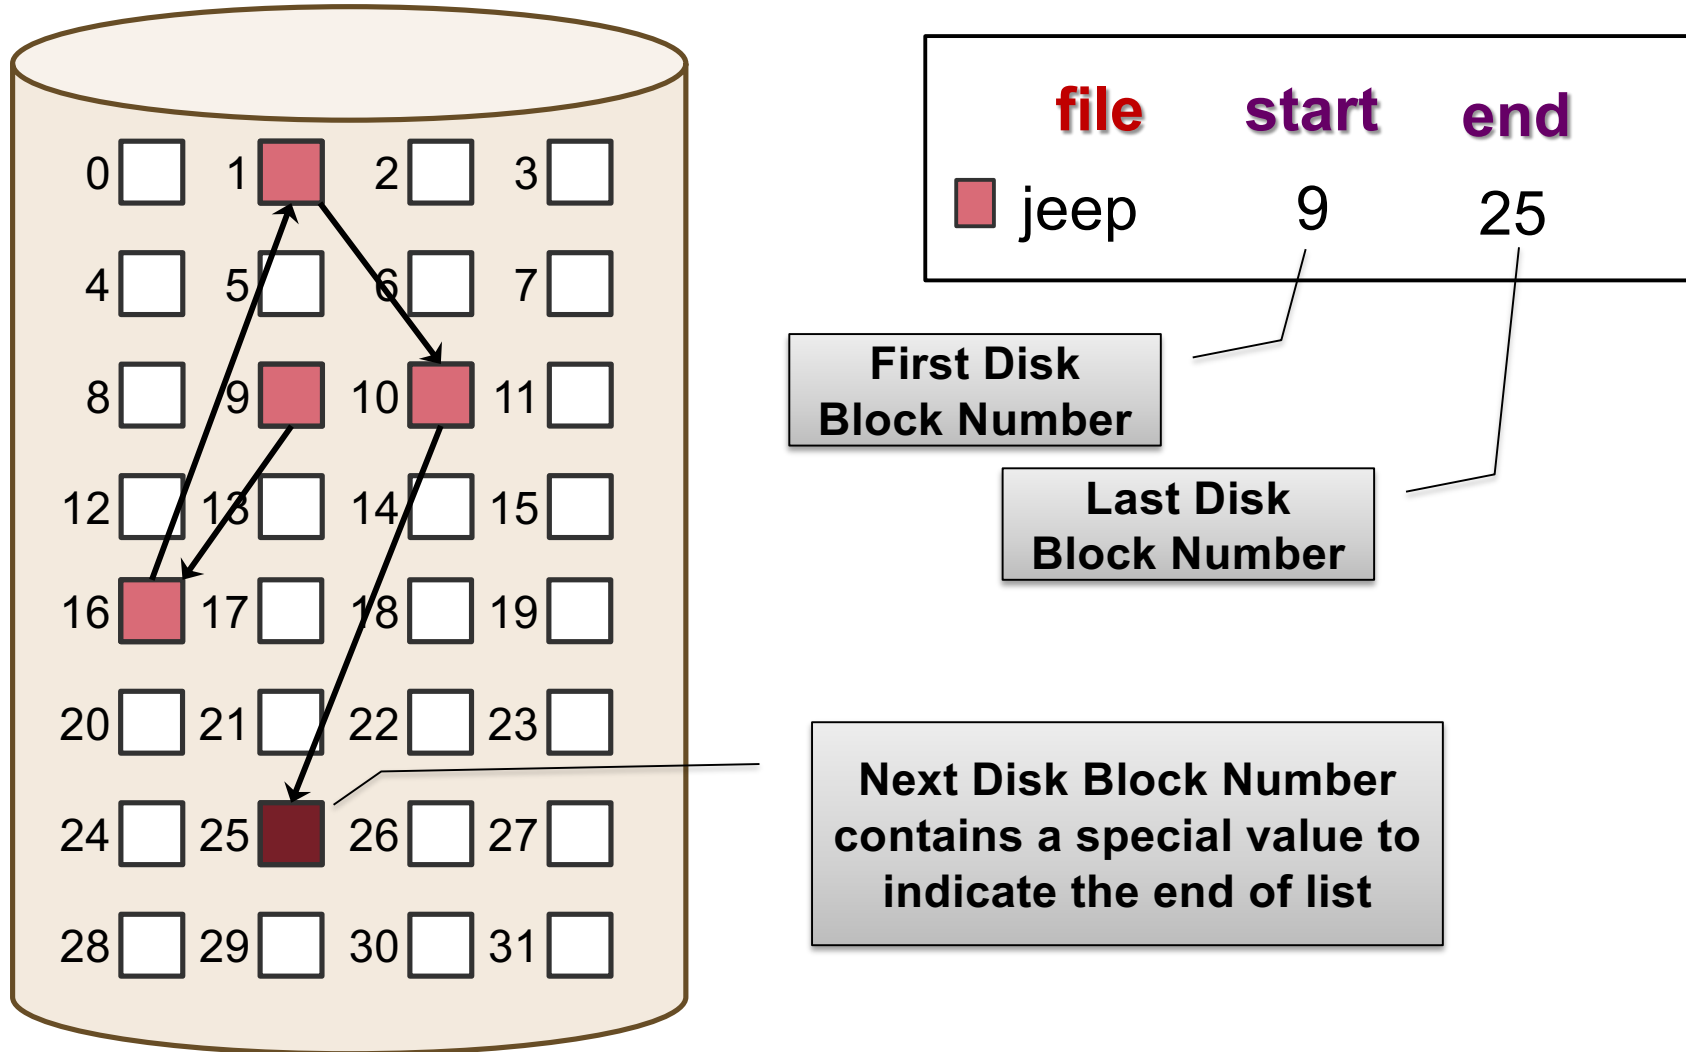
## ■ **Pros:**

- ❑ Solve fragmentation problem

## ■ **Cons:**

- ❑ Random access in a file is very slow
- ❑ Part of disk block is used for pointer
- ❑ Less reliable (what if one of the pointers is incorrect?)

# Linked List Allocation



# File Block Allocation 2: **Linked List V2.0**

## ■ **General Idea:**

- ❑ Move all the block pointers into a single table
  - known as **File Allocation Table (FAT)**
  - FAT is in memory at all time
- ❑ Simple yet efficient
  - Used by MS-DOS

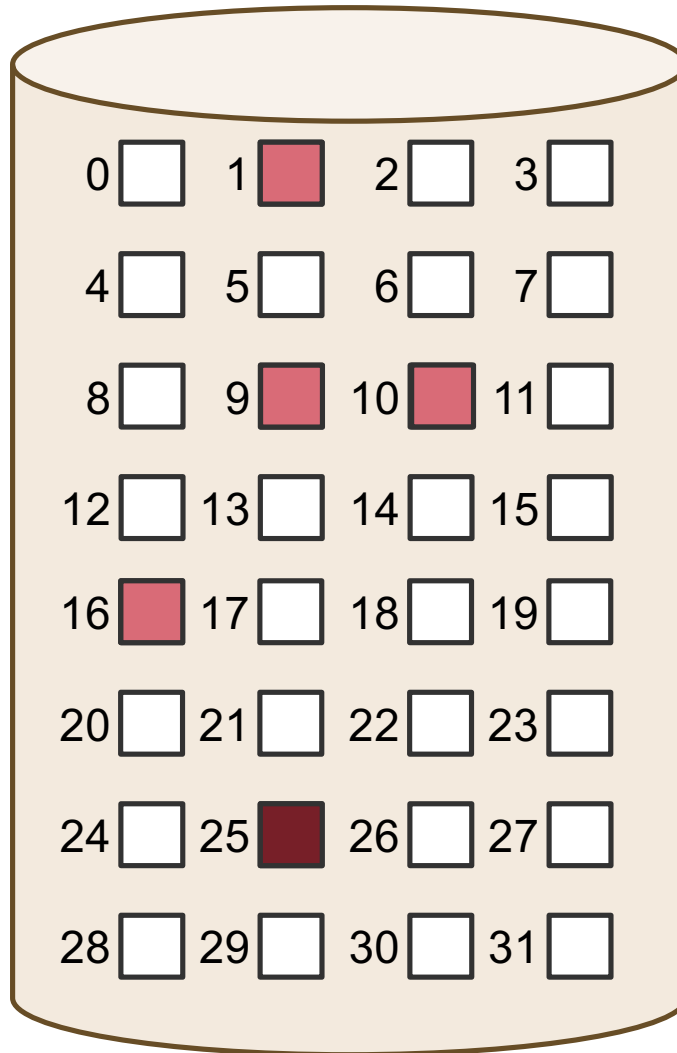
## ■ **Pros:**

- ❑ Faster Random Access
  - The linked list traversal now takes place in memory

## ■ **Cons:**

- ❑ FAT keep tracks of **all disk blocks** in a partition
  - Can be huge when disk is large
  - Consume valuable memory space

# FAT Allocation



file	start
jeep	9

Disk Block  
Number used  
to index table

0	
1	10
9	16
10	25
16	1
25	-1
n-1	

## File Block Allocation 3: **Indexed Allocation**

### ■ **General Idea:**

- ❑ Each file has an **index block**
  - An **array of** disk block addresses
  - $\text{IndexBlock}[N] == N^{\text{th}} \text{ Block address}$

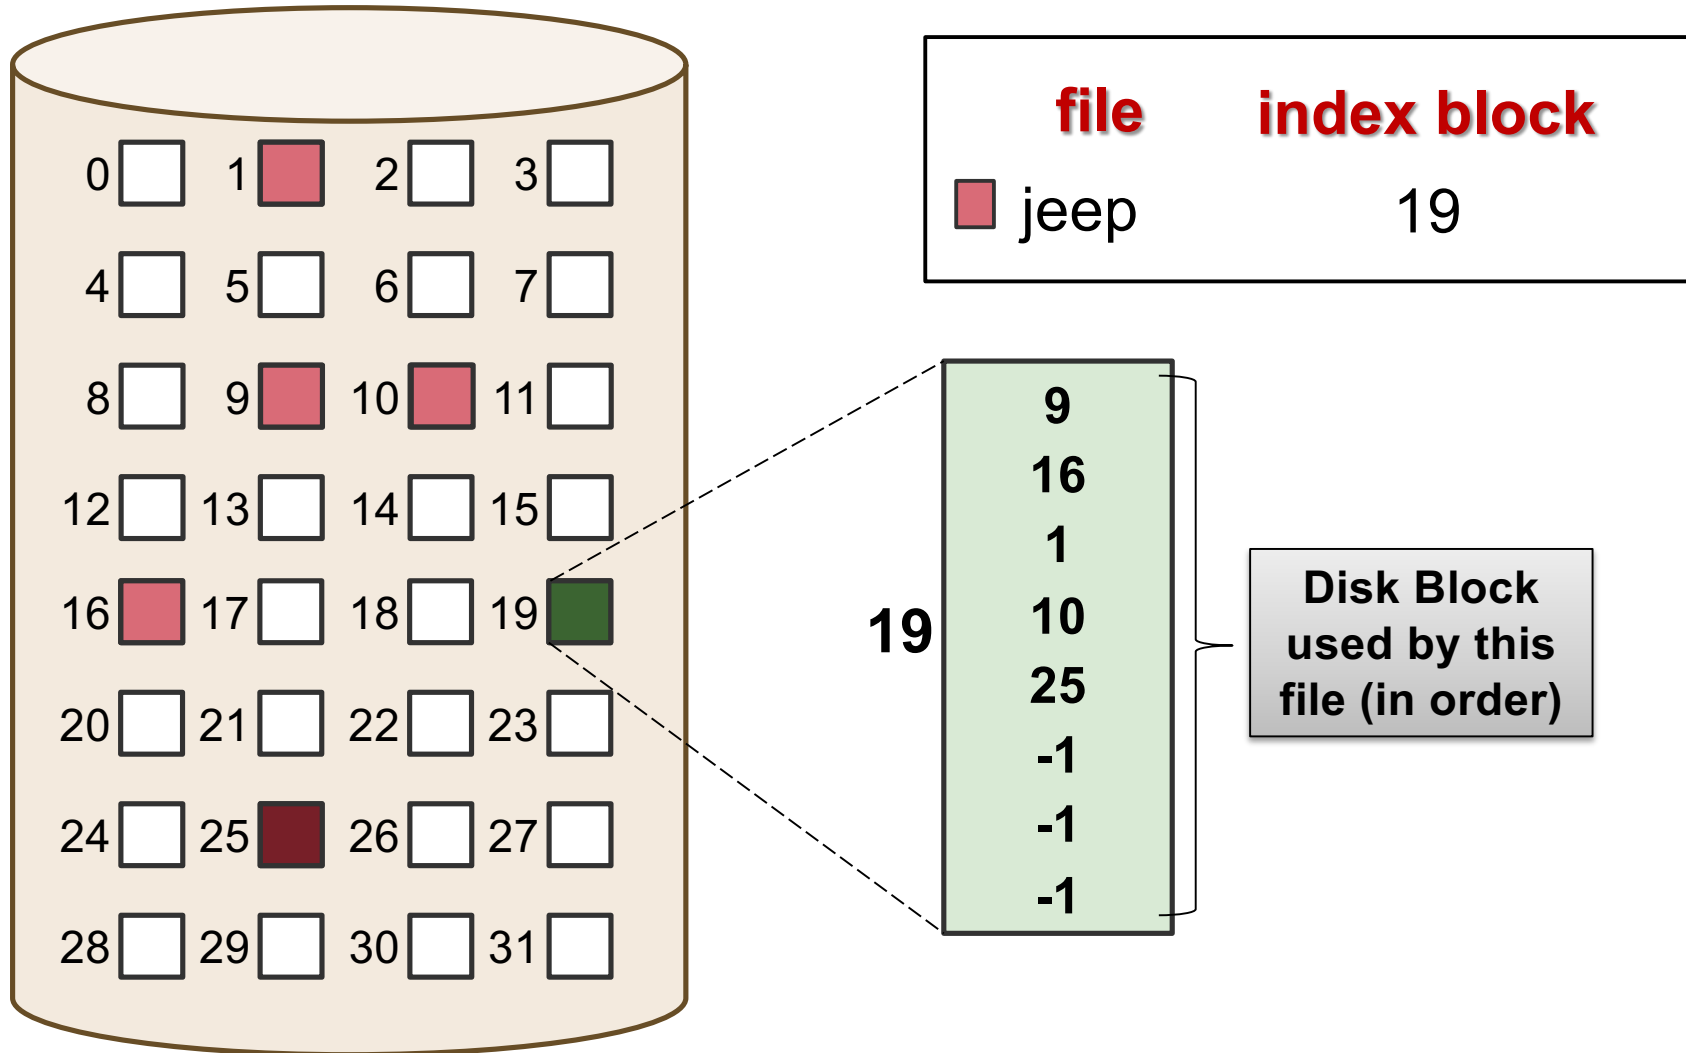
### ■ **Pros:**

- ❑ Lesser memory overhead
  - Only index block of opened file needs to be in memory
- ❑ Fast direct access

### ■ **Cons:**

- ❑ Limited maximum file size
  - Max number of blocks == Number of index block entries
- ❑ Index block overhead

# Indexed Allocation





# Indexed Block Allocation: **Variation**

- Several schemes to:
  - ❑ Allow larger file size
- Linked scheme:
  - ❑ Keep a **linked list** of index blocks
  - ❑ Each index block contains the pointer to next index block
- Multilevel index:
  - ❑ Similar idea **as multi-level paging**
  - ❑ First level index block points to a number of **second level index blocks**
    - Each second level index blocks point to actual disk block
  - ❑ Can be generalized to any number of levels

## Indexed Block Allocation: **Variation** (cont)

### ■ Combined scheme:

- ❑ Combination of direct indexing and multi-level index scheme
- ❑ Example: Unix I-node has:
  - 12 **direct pointers** that point to disk block directly
  - 1 **single indirect block**
    - ❑ which contains a number of direct pointers
  - 1 **double indirect block**
    - ❑ which points to a number of **single indirect blocks**
  - 1 **triple indirect block**
    - ❑ which points to a number of **double indirect blocks**
  - A combination of efficiency (for small file) and flexibility (still allow large file)

# Unix Indexed Node (I-Node): Illustration

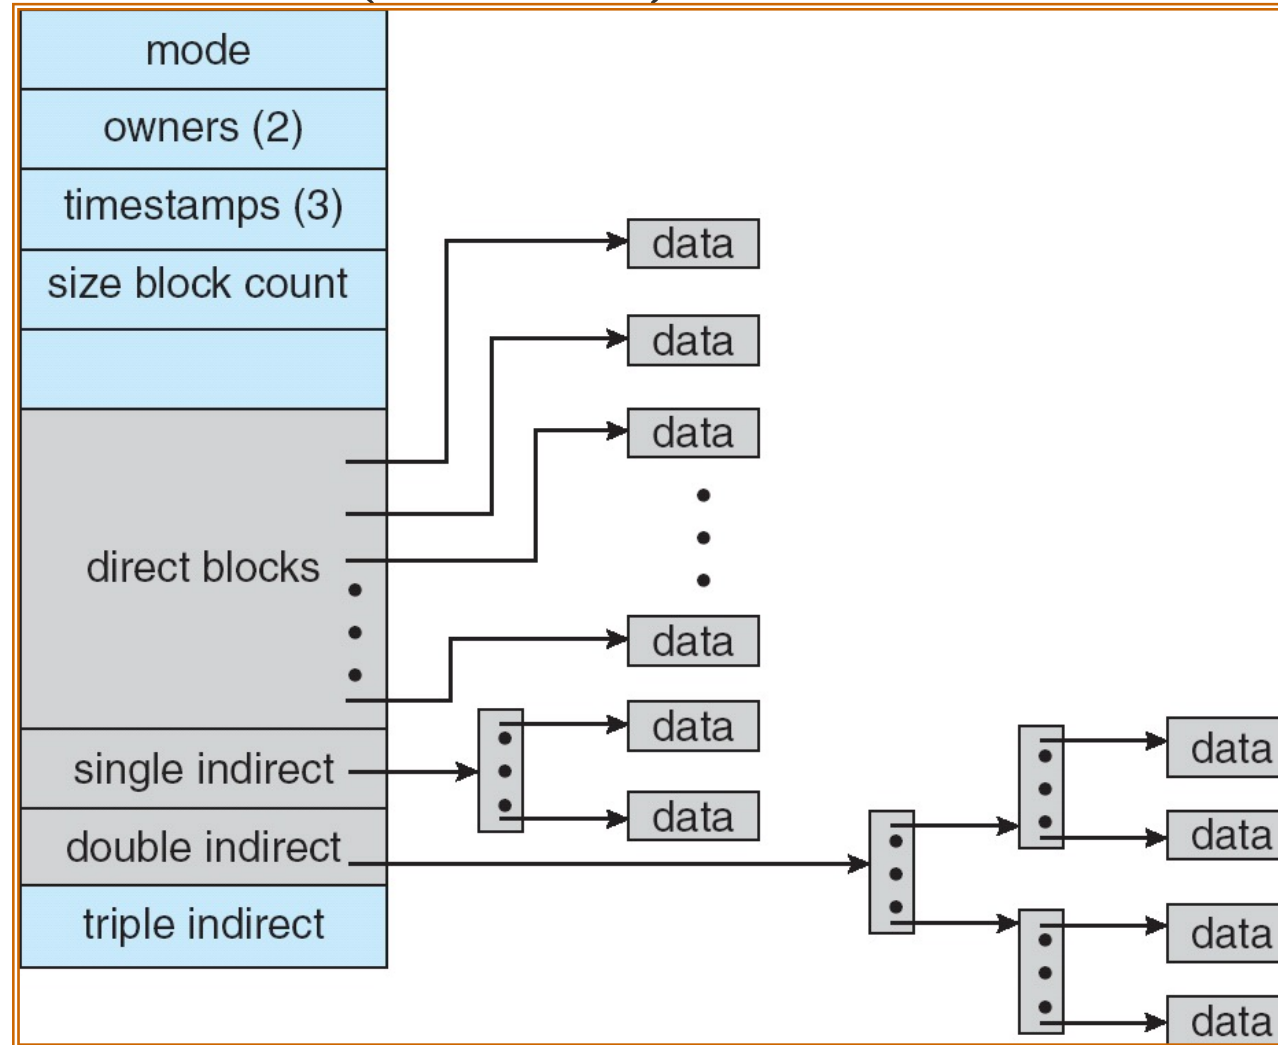



Image taken from "Operating System Concepts" 7<sup>th</sup> Edition by Silberschatz, Galvin and Gagne

# I-Node Unrolled




- 12 **direct pointers** that point to disk block directly
- 1 **single indirect block**
  - contains a number of **direct pointers**
- 1 **double indirect block**
  - points to a number of **single indirect blocks**
- 1 **triple indirect block**
  - points to a number of **double indirect blocks**
- The combination ensures efficiency for small files and flexibility (still support larger files)

---

# Free Space Management

---



# Free Space Management: Overview

- To perform file allocation:
  - Need to know which disk block is free
  - i.e. maintain a **free space list**
- Free space management:
  - Maintain free space information
  - **Allocate:**
    - Remove free disk block from free space list
    - Needed when file is created or enlarged (appended)
  - **Free:**
    - Add free disk block to free space list
    - Needed when file is deleted or truncated

# Free Space Management: **Bitmap**

- Each disk block is represented by 1 bit
  - E.g. 1 == free, 0 == occupied

- Example:

0	1	0	1	1	1	0	0	1	0	1	1	...
---	---	---	---	---	---	---	---	---	---	---	---	-----

- Occupied Blocks = **0, 2, 6, 7, 9, ...**
- Free Blocks = **1, 3, 4, 5, 8, 10, 11, ...**

- **Pros:**

- Provide a good set of manipulations
  - E.g. can find the first free block, n-consecutive free blocks easily by bit level operation

- **Cons:**

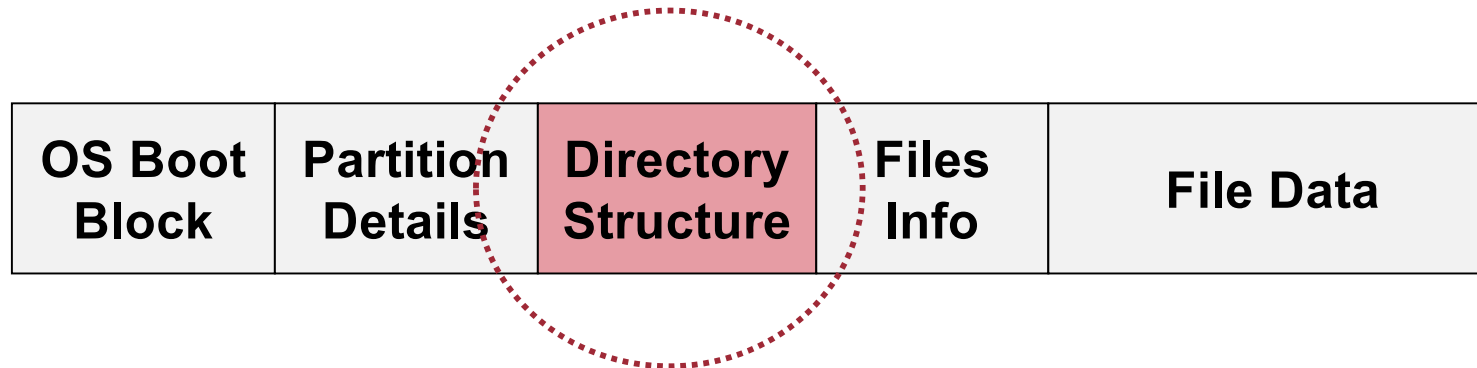
- Need to keep in memory for efficiency reason

# Free Space Management: **Linked List**

- Use a linked list of disk blocks:
  - Each disk block contains:
    - A number of free disk block numbers
    - A pointer to the next free space disk block
- **Pros:**
  - Easy to locate free block
  - Only the first pointer is needed in memory
    - Though other blocks can be cached for efficiency
- **Cons:**
  - High overhead
    - Can be mitigated by storing the free block list in free blocks!



# Implementing Directory



# Directory Structure: Overview

- The main tasks of a directory structure:
  1. Keep tracks of the files in a directory
    - Possibly with the file metadata
  2. Map the file name to the file information
- Remember:
  - ❑ File must be opened before use
    - Something like `open( "data.txt" );`
  - ❑ The purpose of the open operation:
    - Locate the file information using pathname + file name
- Path name
  - ❑ List of directory names traversed from root
  - ❑ E.g. `/dir2/dir3/data.txt`

# Directory Structure: Overview (cont)

- Given a full path name:
  - Need to recursively search the directories along the path to arrive at the file information
- Example:
  - Full path name: `/dir2/dir3/data.txt`
    1. Find "`dir2`" in directory "/"
      - Stop if not found (or incorrect type)
    2. Find "`dir3`" in directory "`dir2`"
      - Stop if not found (or incorrect type)
    3. Find "`data.txt`" in directory "`dir3`"
      - Stop if not found (or incorrect type)
- Sub-directory is usually stored as file entry with special type in a directory

# Directory Implementation: **Linear List**

- Directory consists of a list:
  - Each entry represents a file:
    - Store file name (minimum) and possibly other metadata
    - Store file information or pointer to file information
- Locate a file using list:
  - Requires a linear search
    - Inefficient for large directories and/or deep tree traversal
  - Common solution:
    - Use cache to remember the latest few searches
      - User usually move up/down a path

# Directory Implementation: **Hash Table**

- Each directory contains a
  - Hash table of size  $N$
- To locate a file by file name:
  - File name is hashed into index  $K$  from 0 to  $N-1$
  - **HashTable**[ $K$ ] is inspected to match file name
    - Usually chained collision resolution is used
    - i.e., file names with same hash value is chained together
      - to form a linked list with list head at **HashTable**[  $K$  ]
- **Pros:**
  - Fast lookup
- **Cons:**
  - Hash table has limited size
  - Depends on good hash function

# Directory Implementation: **File Information**

- File information consists of:
  - File name and other metadata
  - Disk blocks information
    - As discussed in the file allocation schemes earlier
- Two common approaches:
  1. Store everything in directory entry
    - A simple scheme is to have a fixed size entry
      - All files have the same amount of space for information
  2. Store only file name and points to some data structure for other info

File locked and loaded!

# FILE SYSTEM IN ACTION

# File System in Action: Overview

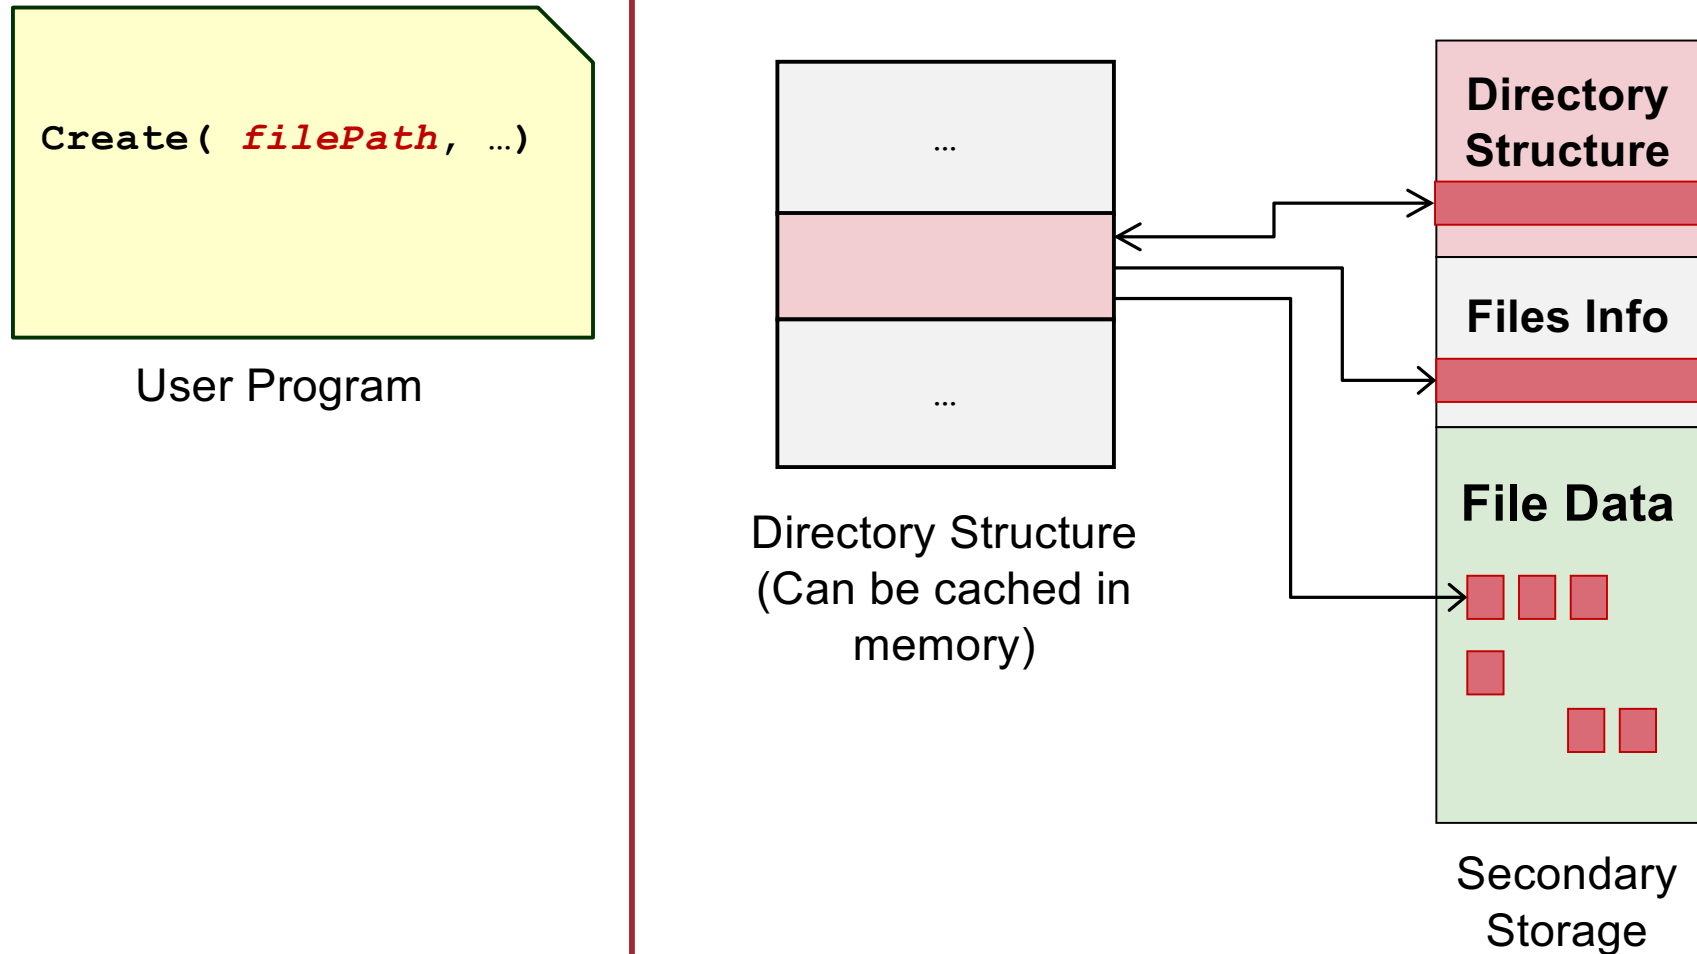
- Previous sections are on **static information** for a FS stored on media
- At runtime, when user interacts with file:
  - **Run-time information** is needed
  - Maintained by OS in memory
- [Recap] Common in-memory information:
  - System-wide open-file table:
    - Contain a copy of file information for each open file + other info
  - Per-process open-file table:
    - Contains pointer to system-wide table + other info
  - Buffers for disk blocks read from/written to disk



# Walkthrough on file operation: **Create**

- Let us relook at the file operation
  - With the newly covered details
- To create a file `/.../.../parent/F`:
  - Use full pathname to locate the **parent** directory
    - Search for filename **F** to avoid duplicates
      - If found, file creation terminates with error
    - Search could be on the cached directory structure
  - Use free space list to find free disk block(s)
    - Depends on allocation scheme
  - Add an entry to **parent** directory
    - With relevant file information
    - File name, disk block information etc

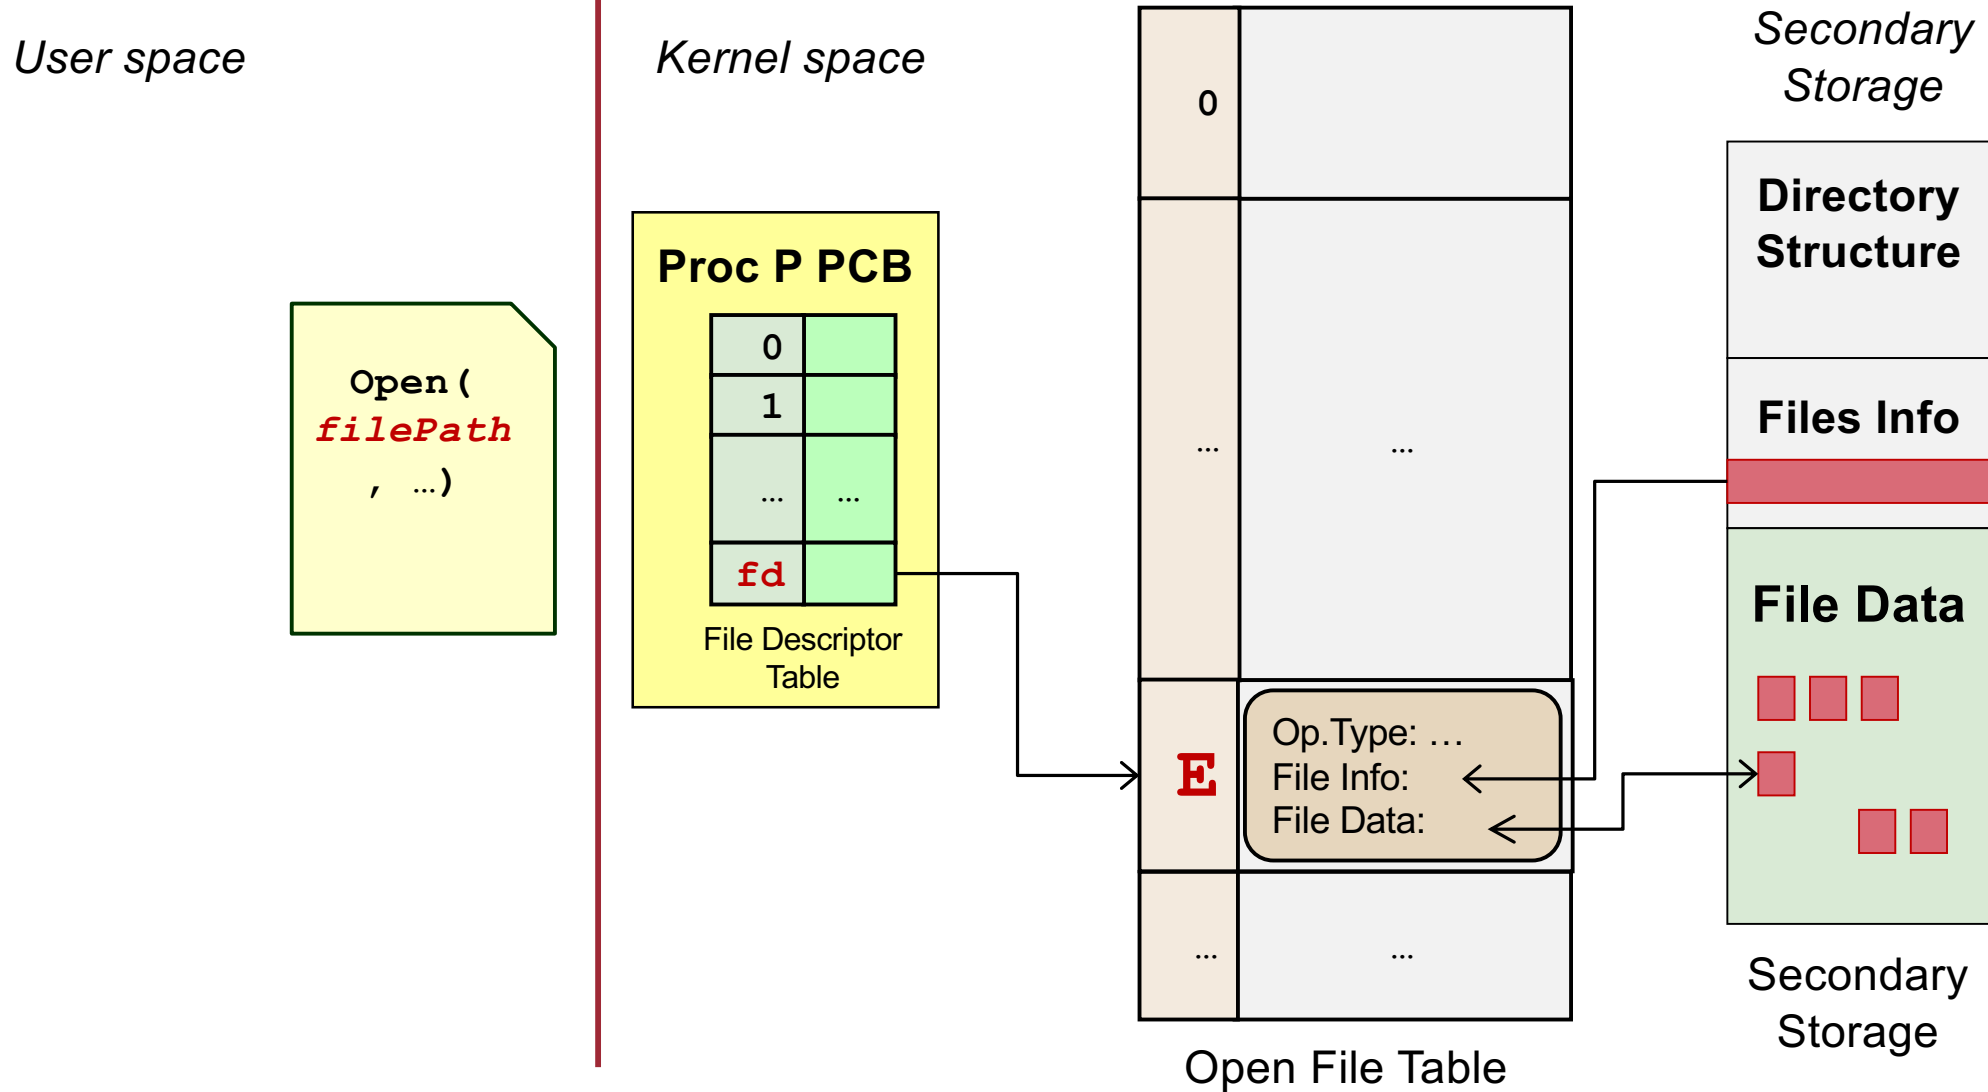
# File Creation: Illustration



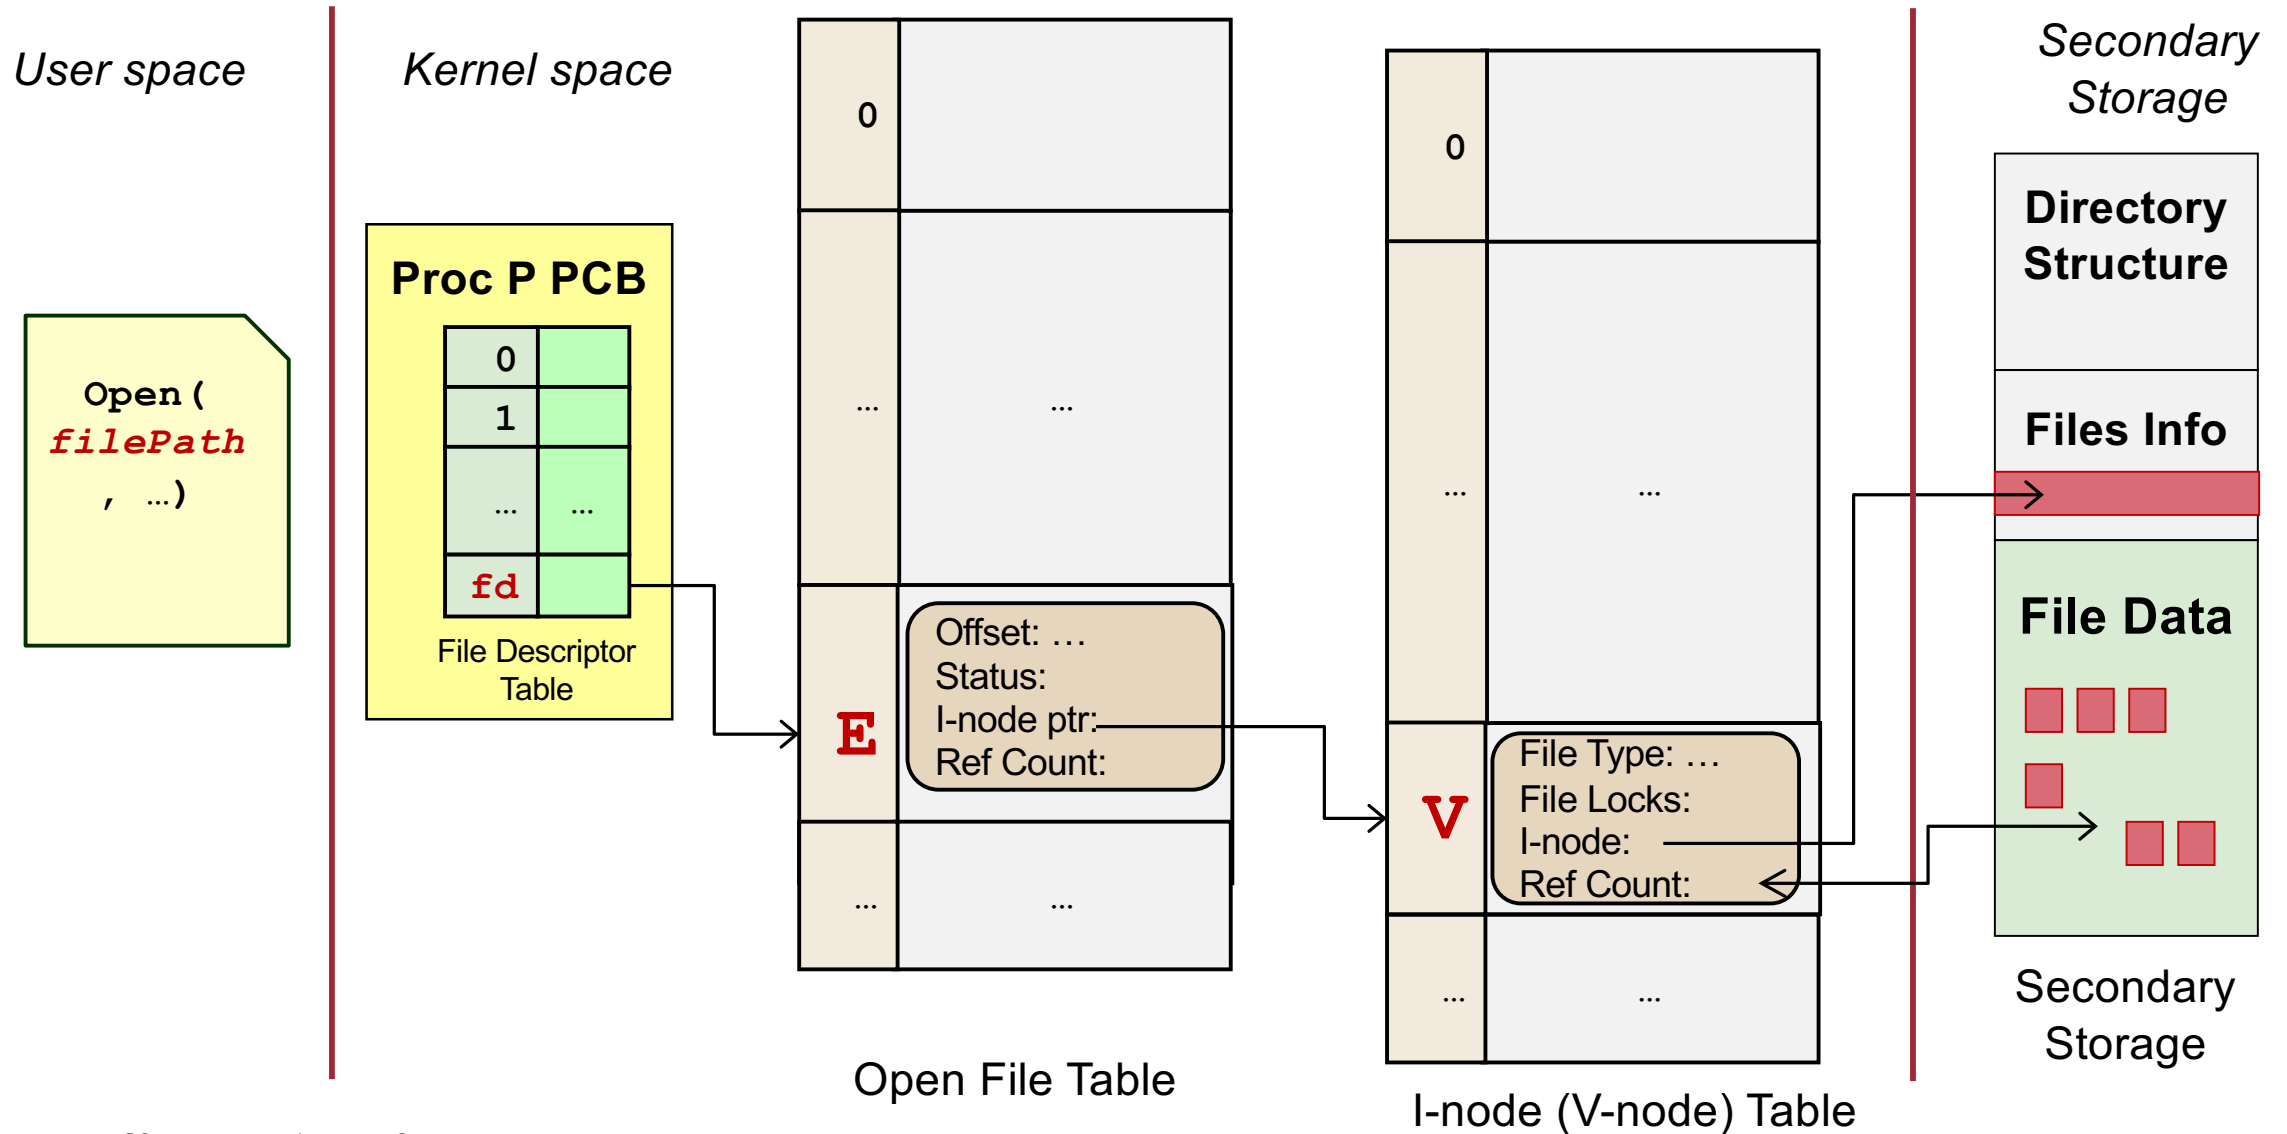
# Walkthrough on file operation: **Open**

- Process **P** open file **/.../.../.../F**:
  - Search system-wide table for existing entry **E**
    - If found:
      - Creates an entry in **P**'s table to point to **E**
      - Return a pointer to this entry
    - If not found, continue to next step
  - Use full pathname to locate file **F**
    - If not found, open operation terminates with error
    - When **F** is located, its file information is loaded into a new entry **E** in system-wide table
    - Creates an entry in **P**'s table to point to **E**
    - Return a pointer to this entry
- The returned pointer is used for further read/write operation

# File Open: Improved Understanding




# File Open in Linux (non-examinable)



I'm afraid you have to wait.....

# DISK I/O SCHEDULING

# Magnetic Disk in One Glance



# Disk Scheduling: The Problem

- Due to the significant seek and rotational latency, OS should schedule the disk I/O requests
- I/O (disk) scheduling:
  - ❑ Intention of reducing **overall waiting time**
  - ❑ As rotational latency is hard to mitigate, we focus on reducing the **seeking time**
  - ❑ Balance the need for high throughput while trying to fairly share I/O requests amongst processes




# Disk Scheduling: **Algorithms**

- Consider the following disk I/O requests indicated by only the **track number (magnetic disks)**:
  - ❑ 13, 14, 2, 18, 17, 21, 15
- A few obvious candidates:
  - ❑ **FCFS**
  - ❑ **SSF (Shortest Seek First)**
    - "SJF" modified for the disk context
  - ❑ The **SCAN** family (aka **Elevator**):
    - Bi-Direction [Innermost  $\leftrightarrow$  Outermost] (SCAN)
    - 1-Direction [Outermost  $\rightarrow$  Innermost, then seek back and start again from the outermost] (C-SCAN)
    - Very intuitive: Imagine the tracks are floors in a building, and the disk head is the elevator servicing the floors (Figure out the algorithm before lecture 😊)

# SCAN: Disk Head Movement

- disk I/O requests indicated by only the **track number** :  
**[13, 14, 2, 10, 17, 21, 7]**



# I/O Scheduling: **Newer Algorithms**

- **Deadline** - 3 queues for I/O requests:
  - Sorted
  - Read FIFO - read requests stored chronologically
  - Write FIFO - write requests stored chronologically
- **noop (No-operation)** - no sorting
- **cfq (Completely Fair Queueing)** - time slice and per-process sorted queues
- **bfq (Budget Fair Queuing) (Multiqueue)** - fair sharing based on the number of sectors requested

# Summary

- Covered implementation details for file system
  - File Information
    - Allocation schemes
  - Free Space management
  - Directory Structure
- Relook at file operations from the OS viewpoint
- Discussed OS responsibility in I/O scheduling
  - for hard disks

# References

- **OS Concepts**, 9<sup>th</sup> Edition
  - ❑ FS abstraction: 11.1 – 11.3
  - ❑ FS implementation: 12.1 – 12.5
  - ❑ Disk Scheduling: 10.4
- **Modern Operating Systems**, 4<sup>th</sup> Edition
  - ❑ FS abstraction: 4.1, 4.2
  - ❑ FS implementation: 4.3
  - ❑ Disk Scheduling: 5.4.3.1
- **Three Easy Pieces:**
  - ❑ Chapters 39, 40