

## CS2106 Operating Systems

### Semester 2 2020/2021

#### Week of 15<sup>th</sup> February 2021

#### Tutorial 4 Suggested Solutions

#### **More on Process Scheduling and IPC**

1. [MLFQ] As discussed in the lecture, the simple MLFQ has a few shortcomings. Describe the scheduling behavior for the following two cases.
  - a. (Change of heart) A process with a lengthy CPU-intensive phase followed by I/O-intensive phase.
  - b. (Gaming the system) A process repeatedly gives up CPU just before the time quantum lapses.

The following are two simple tweaks. For each of the rules, identify which case (a or b above) it is designed to solve, then briefly describe the new scheduling behavior.

- i. (Rule – Accounting matters) The CPU usage of a process is now accumulated across time quanta. Once the CPU usage exceeds a single time quantum, the priority of the task will be decremented.
- ii. (Rule – Timely boost) All processes in the system will be moved to the highest priority level periodically.

ANS:

- a. The process can sink to the lowest priority during the CPU intensive phase. With the low priority, the process may not receive CPU time in a timely fashion during the I/O phase, which degrades the responsiveness. The general shape of the timing chart is the same as the example 1 shown in lecture.
- b. If a process gives up / blocks before the time quantum lapses, it will retain its priority. Since all processes enter the system with the highest priority, a process can keep its high priority indefinitely by using this trick and receive disproportionately more CPU time than other processes.

#### Amendments to MLFQ

- i. This tweak fixes case (b). The trick in (b) works because the scheduler is “memory-less”, i.e. the CPU usage is counted from fresh every time a process receives a time quantum. If the CPU usage is accumulated, then a CPU intensive process will still

exhaust the allowed time quantum and be demoted in priority. This will prevent the process from hogging the CPU.

- ii. This tweak is for case (a). By periodically boosting the priority of all processes (essentially treat all process as “new” and hence have highest priority), a process with different behavior phases may get a chance to be treated correctly even after it has sank to the lowest priority.
2. [Adapted from AY1920S1 Midterm – Evaluating scheduling algorithms] Briefly answer each of the following questions regarding process scheduling, stating your assumptions, if any.
- a. Under what conditions does FCFS (FIFO) scheduling result in the shortest possible average response time?
  - b. Under what conditions does round-robin (RR) scheduling behave identically to FIFO?
  - c. Under what conditions does RR scheduling perform poorly compared to FIFO?
  - d. Does reducing the time quantum for RR scheduling help or hurt its performance relative to FIFO, and why?
  - e. Do you think a CPU-bound (CPU-intensive) process should be given a higher priority for I/O than an I/O-bound process? Why?

ANS:

- a. FIFO minimises the average response time if the jobs arrive in the ready queue in order of increasing job lengths. This avoids short jobs arriving later from waiting substantially for an earlier longer job. A special case also exists: when all jobs have the same completion time.
- b. RR behaves identically to FIFO if the job lengths are shorter than the time quantum, since it is essentially a pre-emptive variant of FIFO.
- c. There are multiple accepted responses, depending on the criteria of evaluation for RR scheduling. Some possible responses include:
  - When the job lengths are all the same and much greater than the time quantum, RR performs poorly in average turnaround time
  - When there are many jobs and the job lengths exceed the time quantum, RR results in reduced throughput due to greater overhead from the OS incurred due to context-switches when jobs are pre-empted

- d. Both *yes* and *no* responses are accepted, if the justification provided references the criteria by which RR is being evaluated by.
    - *Yes*: reducing the time quantum increases the amount of CPU time spent by the OS performing context-switching, reducing throughput and increasing the average turnaround time
    - *No*: reducing the time quantum reduces the waiting time for a task to first receive CPU time, improving the initial response time (*responsiveness*)
  - e. Two main lines of reasoning were accepted.
    - *Yes*: CPU processes can be given higher priority for I/O so they may return to waiting for the CPU, decreasing overall turnaround time at the expense of response time of I/O-bound processes
    - *No*: maximise responsiveness of I/O-bound processes
3. [Shared Memory] In this question, we are going to analyze the pitfalls of having multiple processes accessing and modifying data at the same time. Compile and run the code in **shm.c**. The executable accepts two command line arguments, **n** and **nChild**; if no command line arguments are given, the default values will be used: **n** is initialized to 100 and **nChild** is initialized to 1.

The code does a simple job: the parent creates **nChild** processes and then each process, the parent included, increases a shared memory location by 1 for **n** number of times. After all processes have finished incrementing the shared value, the parent prints the shared result and exits.

- a. The value of the shared memory is initialized to 0, what is the expected final printed value given **n** and **nChild**?
- b. Run the program multiple times with various values for **n** and **nChild**. Observe the result when **n** is high (e.g. **n** = 10000 and **nChild** = 10). Explain the results.
- c. Do you think (b) is caused by multi-core processors since processes can run in parallel? Would the issue persist if you use a single core processor? [You can run experiment to find out. Use the Linux command **taskset** to bind a process (and all its child) to a single processor core, e.g. "**taskset -c 5 ./a.out 10000 100**" will bind the executable to run on core no. 5.]

ANS:

- a. We expect the printed value to be equal to  $(1 + \text{nChild}) * n$  at the end.
- b. We see final value smaller than expected in some test cases when **n** is sufficiently high. When incrementing a variable, the process must perform the following steps:
  - I. Read the value from the memory
  - II. Modify it, and

### III. Write the value back to the memory.

If two or more processes have their executions interleaved, it is very likely that one process will overwrite the value written by another process. In this case, it is impossible to have a deterministic value.

There are two possible scenarios for the interleaving:

1. [Possible on multi-core processors] Multiple processes executes in parallel and happen to interleave the instructions.
2. [Possible on single core processors] Process A swapped out (e.g. due to time quantum) before step (3), the next process B will read the memory value and continue to update. When A swapped back in, its previously read value (stored in register) will be outdated. Upon writing back to memory, A has wipe out work done by B.

This also explains why we need higher **n** to increase the likelihood of this interleaving.

c. As explained in (b) scenario 2, you still can get wrong result on a single core. You need to use higher **n** to increase the chance that a process swapped out at the "right" place for the error to show up. On Compute Cluster Node, a **n** value of 10,000,000 with **nChild** > 3 should give you good chance to see the mistake.

4. [Protecting the Shared Memory] Allowing processes to access and modify shared data whenever they please can be problematic! Therefore, we would like to modify the code in **shm.c** such that the output is deterministic regardless of how large **n** and **nChild** are. More precisely, we want the processes to take turns when modifying the result value that resides in the shared memory.

To this end, we will add another field in the shared memory, called the **order** value, that specifies which process' turn it is to increment the shared result. If the **order** value is 0, then the parent should increase the shared result, if the **order** value is 1, then the first child should increase it, if the **order** value is 2, then the second child and so on. Each process has an associated **pOrder** and checks whether the value **order** is equal to its **pOrder**; if it is, then it proceeds to increment the shared result. Otherwise, it waits until the order value is equal to its **pOrder**.

Your task is to modify the code in **shm\_protected.c** to achieve the desired result. You will have to:

- Create a shared memory region with two locations, one for the shared result, and the other one for the order value;

- Write the logic that allows a process to modify the shared result only if the order value is equal to its **pOrder** (the skeleton already takes care of assigning the right **pOrder** to each process);
- Print the result value and cleanup the shared memory.

Why is **shm** faster than **shm\_protected**? Why is running **shm\_protected** with large values for **nChild** particularly slow? (Hint: You may want to take a look at the output of the Linux command *htop*)

ANS:

To guarantee that two processes won't try to modify the shared memory location at the same time, we must **serialize the access** to it, i.e., make it impossible for two processes to access it at the same time. Serializing the code will, of course, make the execution time longer.

The technique we use to prevent simultaneous access is called **busy waiting**. Each process is continuously checking whether its time to modify the variable has come. This requires the process to run on the CPU and consume CPU cycles (thus the name). Note that at the moment when it's Child X's time to increment the value, there are **nChild - X** processes doing busy waiting – each of these processes will get scheduled to run on the CPU but the execution of the program will not make any progress. Because of this, your program will take a long time to complete for large values of **nChild**.

---

### Interesting topics for your own exploration

Unix Pipes ( [https://en.wikipedia.org/wiki/Dup\\_\(system\\_call\)](https://en.wikipedia.org/wiki/Dup_(system_call)) ) – pretty good article with code(!) example).

Linux Message Queue ( <https://www.geeksforgeeks.org/ipc-using-message-queues/> - Simple code example to highlight the communication).

Shared Memory in Parallel Programming ( [https://en.wikipedia.org/wiki/Shared\\_memory](https://en.wikipedia.org/wiki/Shared_memory) - Generalized shared memory model that can work *across network*)