

# CS2106 Operating Systems

## Semester 1 2020/2021

Week of 5<sup>th</sup> April 2021

Tutorial 10

### File Abstraction

1. [Working Set Model] Working set model allows OS to make better decision for frame allocation among processes. Let us consider several aspects of the idea in this question.
  - a. [Basic Idea] Given the following memory reference string (from lecture 9 – page replacement section):

Time	1	2	3	4	5	6	7	8	9	10	11	12
Page	2	3	2	1	5	2	4	5	3	2	5	2

Given that  $W(T, \Delta)$  gives the working set at time  $T$  with a window size of length  $\Delta$ , give the following:

**$W(9, 3)$ ,  $W(11, 3)$ ,  $W(9, 4)$ ,  $W(11, 4)$**

- b. [Application] Suppose we know the  $W(T, \Delta)$  value for all processes, how can OS use this information? Hint: think about virtual memory + demand paging. Demand paging refers to the idea that OS only loads a page into RAM upon page fault instead of trying to predict and pre-load pages for processes.

(Exploration Question) Using the same idea, is there any easy way to dynamically adjust the  $\Delta$  value?

- c. [Implementation] Consider the following "mysterious" algorithm:

1. Every Page Table Entry has an additional K-bit mysterious value,  $M_0, M_1, \dots, M_{K-1}$ . All K-bit are initialized to '0'.
2. Whenever a page P is referenced, its corresponding Mysterious value are updated as follows:
  - a.  $M_{K-1} \leftarrow M_{K-2}$
  - b. ...
  - c.  $M_2 \leftarrow M_1$
  - d.  $M_1 \leftarrow M_0$
  - e.  $M_0 \leftarrow 1$
3. All other Non-P pages' Mysterious value are updated as follows:
  - a.  $M_{K-1} \leftarrow M_{K-2}$
  - b. ...
  - c.  $M_2 \leftarrow M_1$
  - d.  $M_1 \leftarrow M_0$
  - e.  $M_0 \leftarrow 0$

What does the mysterious value represents? If the above algorithm is handled by the hardware, what should be done every K memory accesses?

2. [Wrapping File Operation] File operations are very expensive in terms of time. There are several reasons: a) As we learned in the lecture, each file operation is a system call, which requires an execution mode change (user → kernel); b) Secondary storage has high latency access time.

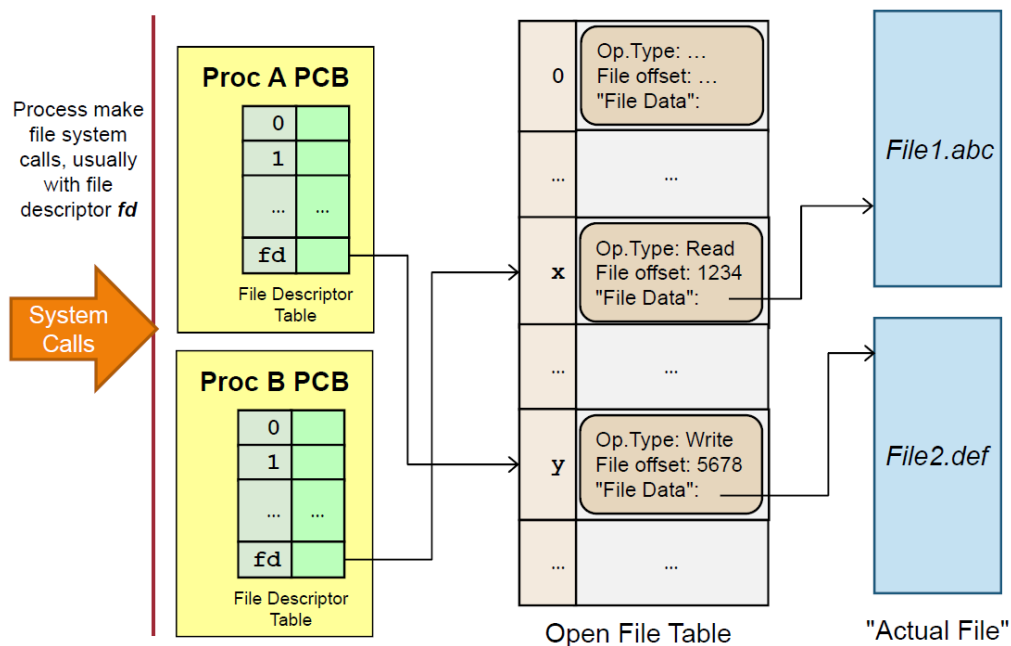
This leads to a strange phenomenon: it is generally true that the total time to perform 100 file operations for 1 item each is **much longer** than performing a single file operation for 100 items instead. e.g. write one byte 100 times takes longer than writing 100 bytes in one go.

So, most high level programming languages provide **buffered file operations** that wraps around the primitive file operations. The buffered version essentially maintains an internal intermediate storage in memory (i.e. buffer) to store user read/write values from/to the file. For example, a **buffered write operation** will wait until the internal memory buffer is full before doing a large one time file write operation to *flush* the buffer content into file.

- a. [Generalization] Give one or two examples of buffered file operations found in your favorite programming language(s). Other than the "chunky" read/write benefit, are there any other additional features provided by those high level buffered file operations?
- b. [Application] Take a look at the given "**weird.c**" source code. Compiles and performs the following experiments: Change the trigger value from 100, 200, ... until you see values printed on screen **before the program crashes**. Can you explain both the behavior and the significance of the "trigger" value? If you add a new line character "\n" to the the **printf()** statement, how does the output pattern changes? How can this information be useful?
- c. [Design] Give a **high level pseudo code** to provide buffered file read operation. Use the following "function header" as a starting point:

```
BufferedFileRead( File, outputArray, arraySize )  
//Read "arraySize" items from "File" and place in  
// the "outputArray"
```

3. [Open File Table] Below is the illustration taken from lecture 10:



Discuss how this organization helps OS to handle the following scenarios. Your answer should refer to the relevant structure(s) if possible.

- Process A tries to open a file that is currently being written by Process B.
  - Process A tries to use a bogus file descriptor in a file-related system call.
  - Process A can never "accidentally" access files opened by Process B.
  - Process A and Process B reads from the same file. However, their reading should not affect each other.
  - Redirect Process A's standard input / output, e.g. "**a.out < test.in > test.out**". (Hint: \*nix processes has 3 default file descriptors: 0 = stdin, 1 = stdout, 2 = stderr)
4. [Cover if time permits - Understanding directory permission] In \*nix system, a directory has the same set of permission settings as a file. For example:

```
sooyj@sunfire [13:22:52] ~/tmp/Parent $ ls -l
total 8
drwx--x--x 2 sooyj  compsc  4096 Nov  8 13:22 Directory
sooyj@sunfire [13:22:53] ~/tmp/Parent $
```

You can see that directory **Directory** has the read, write, execute permission for owner, but only execution permission for group and others. It is easy to understand the permission bits for a regular file (read = can only access, write = can modify, execute = can execute this file). However, the same cannot be said for the directory permission bits.

Let's perform a few experiments to understand the permission bits for a directory.

**Setup:**

- Unzip **DirExp.zip** on any \*nix platform (Solaris, Mac OS X included).
- Change directory to the **DirExp/** directory, there are 4 subdirectories with the same set of files. Let's set their permission as follows:

<code>chmod 700 NormDir</code>	NormDir is a normal directory with read, write and execute permissions.
<code>chmod 500 ReadExeDir</code>	ReadExeDir has read and execute permission.
<code>chmod 300 WriteExeDir</code>	WriteExeDir has write and execute permission.
<code>chmod 100 ExeOnlyDir</code>	ExeOnlyDir has only execute permission.

Perform the following operations on each of the directory and note down the result. Make sure you are at the **DirExp/** directory at the beginning. DDDD is one of the subdirectories.

- Perform "**ls -l DDDD**".
- Change into the directory using "**cd DDDD**".
- Perform "**ls -l**".
- Perform "**cat file.txt**" to read the file content.
- Perform "**touch file.txt**" to modify the file.
- Perform "**touch newfile.txt**" to create a new file.

Can you deduce the meaning of the permission bits for directory after the above? Can you use the "directory entry" idea to explain the behavior?

---

**Questions for your own exploration**

1. Explain the following concepts in your own words clearly; the shorter, the better! Your explanation should be easily understandable for non-CS2106 students.
  - a. What is a file?
  - b. Name and describe the two basic classifications of files.
  - c. Distinguish between a file type and a file extension.
  - d. What does it mean to open and close a file?
  - e. What does it mean to truncate a file?
2. (Adapted from [SGG]) Why do many operating systems have a system call to "open" a file, rather than just passing a path to the read or write system calls each time?