

# Access Methods II: Hash-based Indexes

"If you don't find it in the index, look very carefully through the entire catalogue."


-- Sears, Roebuck, and Co.,  
Consumer's Guide, 1897

# *Hash-based Index*

- *Hash-based indexes*
  - Data (or data entries) are stored in buckets
  - Key  $k$ , hash function  $h$ ,  $h(k)$  returns the bucket/page ID that stores record with key  $k$
  - (Ideally) best for *equality selections*
  - Performance degenerate for *skewed* data distributions
  - Inefficient for range searches
    - Depends on hash function used
- Static and dynamic hashing techniques exist

# Static Hashing


- Data is stored in M **buckets**:  $B_0, B_1, \dots, B_{M-1}$ 
  - M is *fixed* at creation time
  - Each bucket consists of one **primary data page** and a chain of zero or more **overflow data pages**
  - Primary pages are allocated sequentially, never de-allocated



- Hashing function  $h(\cdot)$  is used to determine the bucket to store a record
  - A record with search key  $k$  is inserted into bucket  $B_j$ , where  $j = h(k) \bmod M$
  - Example: hash function of the form  $h(\text{key}) = a \cdot \text{key} + b$ 
    - $a$  and  $b$  are constants
    - $h$  has to be tuned for different applications
  - Long overflow chains can develop and degrade performance


# Static Hashing: Example

- $k^*$  denotes a data entry  $e$  with key  $k$
- Suppose we have the following keys, and we have a hash index with  $M = 4$  buckets, each of which can hold 3 records (hashing function:  $j = h(k) = k \bmod 4$ )
  - 9, 14, 18, 32, 25, 5, 10, 44, 37, 31, 35, 29, 36, 7, 43, 30, 22, 6, 38, 11




# Static Hashing: Example

- What if we have  $M = 6$  buckets (hashing function:  $j = h(k) = k \bmod 6$ )?
  - 9, 14, 18, 32, 25, 5, 10, 44, 37, 31, 35, 29, 36, 7, 43, 30, 22, 6, 38, 11




# Static Hashing (Cont.)

- Buckets may contain *data records (Format 1)* or *pointers (Format 2)*
  - Unless otherwise stated, we assume the **former (Format 1)**



Format 1: Buckets contain the data records

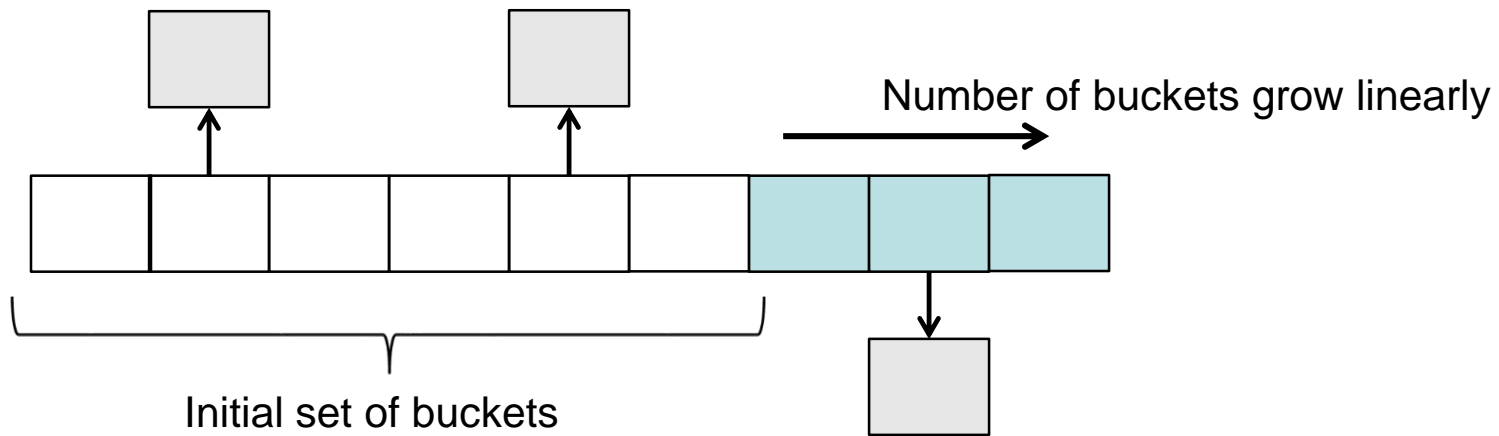


Format 2: Buckets contain the (key, pointer)-pairs

- How to cope with growth?
  - Overflows and reorganization (may need to change the hash fn)
  - Dynamic hashing: *Linear Hashing* and *Extendible Hashing*

# Dynamic Hashing: Linear Hashing

- Linear Hashing (LH) is a *dynamic* scheme
  - Hash file grows *linearly (one bucket at a time)* by systematic splitting of buckets
- LH handles the problem of long overflow chains
  - Overflow pages are still needed as an overflowed bucket might not be split immediately




# Challenges

- How to split a bucket  $B_i$ ?
  - Obtain a new bucket  $B_j$  (known as **split image** of  $B_i$ )
  - Redistribute entries in  $B_i$  between  $B_i$  and  $B_j$
  - Issues
    - How to determine  $j$ ?
    - How to redistribute entries?
    - How to modify hash function?
- When to split a bucket?



# How to determine split image?

- Suppose the initial file size has  $N$  buckets:  $B_0, \dots, B_{N-1}$
- Imagine a corresponding set of **virtual** buckets (not created until needed):  $B_N, \dots, B_{2N-1}$



- To split bucket  $B_i$ , choose  $B_{N+i}$  as its split image (and create  $B_{N+i}$ )
- File size doubles after one round of splitting (i.e., every bucket has been split)

# File doubles after each round ...

Initial




After 1 round



After 2 rounds



After 3 rounds



NOTE:

(a) File grows one bucket at a time

(b) Hash function has to be changed after each round!

# How to redistribute entries?

- Let the initial file size be  $N_0 = 2^m$  ( $m > 1$ )
  - Initial file has buckets  $B_0, \dots, B_{N_0-1}$
- Define the initial hash function  $h_0$  as follows:
$$h_0(k) = \text{last } m \text{ bits of } h(k)$$
- $h_0(k) \in [0, N_0 - 1]$
- An entry  $e$  belongs to bucket  $B_i$  if  $h_0(e.\text{key}) = i$
- Example:  $m = 2$ ,  $N_0 = 4$ , initial hash file has buckets  $B_0, \dots, B_3$

k	$h(k)$	$h_0(k)$
Alice	$\dots 010$	10
Bob	$\dots 110$	10
Carol	$\dots 111$	11
Dave	$\dots 110$	10

Alice belongs to bucket  $B_2$   
Bob belongs to bucket  $B_2$   
Carol belongs to bucket  $B_3$   
Dave belongs to bucket  $B_2$

# How to redistribute entries?

k	h(k)	h <sub>0</sub> (k)
Alice	...010	10
Bob	...110	10
Carol	...111	11
Dave	...110	10

- B<sub>0</sub>
- B<sub>1</sub>
- B<sub>2</sub>
- B<sub>3</sub>

↑  
bucket

B<sub>4</sub> (B<sub>000</sub> B<sub>100</sub>)

B<sub>5</sub> (B<sub>001</sub> B<sub>101</sub>)

B<sub>6</sub> (B<sub>010</sub> B<sub>110</sub>)

B<sub>7</sub> (B<sub>011</sub> B<sub>111</sub>)

↑  
split image


m (2) bits needed  
to identify bucket

m+1 (3) bits needed to identify  
split image;  
(m+1)th bit needed to distinguish  
a bucket from its split image

# How to redistribute entries?

k	h(k)	h <sub>0</sub> (k)
Alice	...010	10
Bob	...110	10
Carol	...111	11
Dave	...110	10

- $B_0 \rightarrow B_0$  or  $B_4$  ( $B_{000}$  or  $B_{100}$ )
- $B_1 \rightarrow B_1$  or  $B_5$  ( $B_{001}$  or  $B_{101}$ )
- $B_2 \rightarrow B_2$  or  $B_6$  ( $B_{010}$  or  $B_{110}$ )
- $B_3 \rightarrow B_3$  or  $B_7$  ( $B_{011}$  or  $B_{111}$ )



How to redistribute  $\{Alice, Bob, Dave\}$  in  $B_2$ ?

Split image of  $B_2$  is  $B_6$

Alice remains in  $B_2$

Bob & Dave are moved to  $B_6$

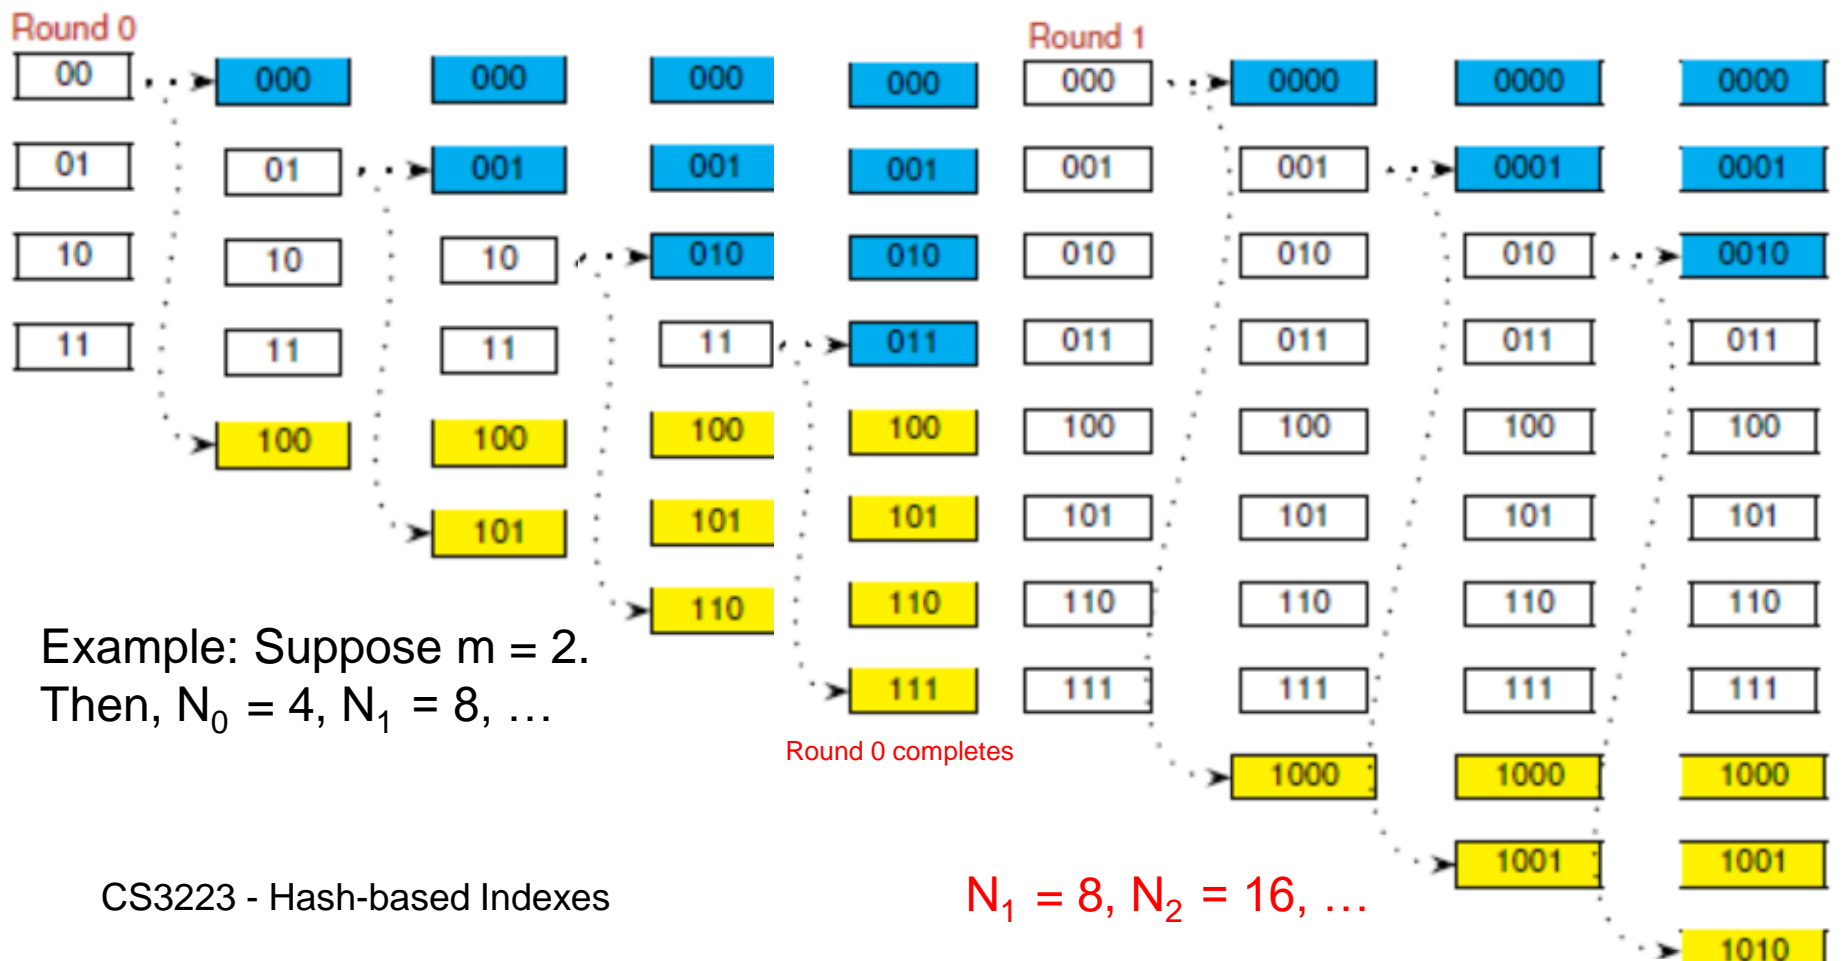
Think of it as re-hashing using  
3 bits instead of 2

# How to redistribute entries?

- To distribute entries in  $B_i$  between  $B_i, \dots, B_{N_0+i}$ 
  - Use the last  $(m+1)^{\text{th}}$  bit of  $h(e.\text{key})$  to redistribute entry  $e \in B_i$ 
    - If the last  $(m+1)^{\text{th}}$  bit of  $h(e.\text{key})$  is 0,  $e$  remains in  $B_i$
    - Otherwise, move  $e$  to  $B_{N_0+i}$
  - If  $e$  is moved to  $B_{N_0+i}$ , the last  $(m+1)$  bits of  $h(e.\text{key})$  must be  $(=) N_0 + i$

# Linear Hashing Example

- To grow file linearly, buckets are split *sequentially*
  - Bucket  $B_i$  must be *split before* bucket  $B_{i+1}$
- File size doubles after each round of splitting




# How to modify the hash function?

- Splitting progresses in rounds represented by level = 0, 1, 2, ...
- Let the initial file size be  $N = 2^m$  buckets
- At round  $i$ 
  - Number of buckets at the beginning is  $N_i = 2^i N = 2^{m+i}$
  - Number of buckets at the end of round  $i$  is  $2 N_i$
  - The bucket to be split is chosen in **round robin fashion**
    - **next** = a pointer to the next bucket to be split
    - $\text{next} \in \{0, 1, \dots, N_i - 1\}$
    - When bucket  $B_{\text{next}}$  is split, its split image is  $B_{\text{next} + N_i}$
  - Uses a pair of hash functions:  $h_i$  and  $h_{i+1}$ 
    - $h_{i+1}(k) = \text{last } (m+i) \text{ bits of } h(k)$
    - $h_i(k) = h(k) \bmod N_i$  (**this is essentially the last  $m$  bits**)



# How to modify the hash function?

- Buckets can be classified into three regions:



# When to split a bucket?


- The time to split the next bucket can be decided with various criteria:
  - Split whenever **some** bucket/page overflows
  - Split whenever space utilization of file is above some threshold
  - etc
- Overflow pages are still needed since an overflowed bucket might not be split immediately
- Some split pages may not be overloaded
- We shall assume that a bucket is triggered whenever **some bucket overflows**
  - Bucket  $B_j$  overflows if an entry is to be inserted into  $B_j$  and all the pages in  $B_j$  (i.e., primary and overflow pages) are full

# So far, ...

- Initially
  - We have  $2^m$  buckets
  - We have 2 hash functions  $h_i$  and  $h_{i+1}$  (based on last  $m+i$  bits and last  $m+i+1$  bits)
  - We have a next (*next-bucket-to-split*) pointer; this is pointing at bucket 0
- File grows one bucket at a time
  - Whenever a bucket is split, its records are redistributed using the last  $m+i+1$  bits
  - Increment the *next* pointer

# Inserting a data entry with search key $k$

- Find bucket by applying hash function  $h_i$ 
  - Use  $h_i(k)$  if  $h_i(k) \geq \text{next}$  (why?)
  - Otherwise, apply  $h_{i+1}$



- If bucket to insert into is full
  - Add overflow page and insert data entry
  - Split  $next$  bucket and increment  $next$ 
    - Redistribute entries in  $B_{next}$  to  $B_{next+N_i}$  using  $h_{i+1}$
    - $next = next + 1$
    - Reset  $next$  to bucket 0 when a round completes
- Since buckets are split round-robin, long overflow chains not expected to develop!

# Example of Linear Hashing

- Recall: On **split**,  $h_{i+1}$  is used to **re-distribute** entries
- Insert  $43^*$  (1010**11**)
  - Bucket  $B_{11}$  overflows (but we are not splitting it!)
  - Split bucket  $B_{00}$  ( $32 = 100$ **0**00,  $36 = 100$ **1**00,  $44 = 101**1**00)$
  - Increment next to 1
  - Insert  $43^*$  into overflow page

$i=0, m=2, N_0=4$

**next=0**

**PRIMARY  
PAGES**

00

32*	44*	36*	
-----	-----	-----	--

01


9*	25*	5*	
----	-----	----	--

10

14*	18*	10*	30*
-----	-----	-----	-----

11

31*	35*	7*	11*
-----	-----	----	-----



$i=0$

**PRIMARY  
PAGES**

**OVERFLOW  
PAGES**

000

32*			
-----	--	--	--

**Next=1**

01

9*	25*	5*	
----	-----	----	--

10

14*	18*	10*	30*
-----	-----	-----	-----

11

31*	35*	7*	11*
-----	-----	----	-----

43*			
-----	--	--	--

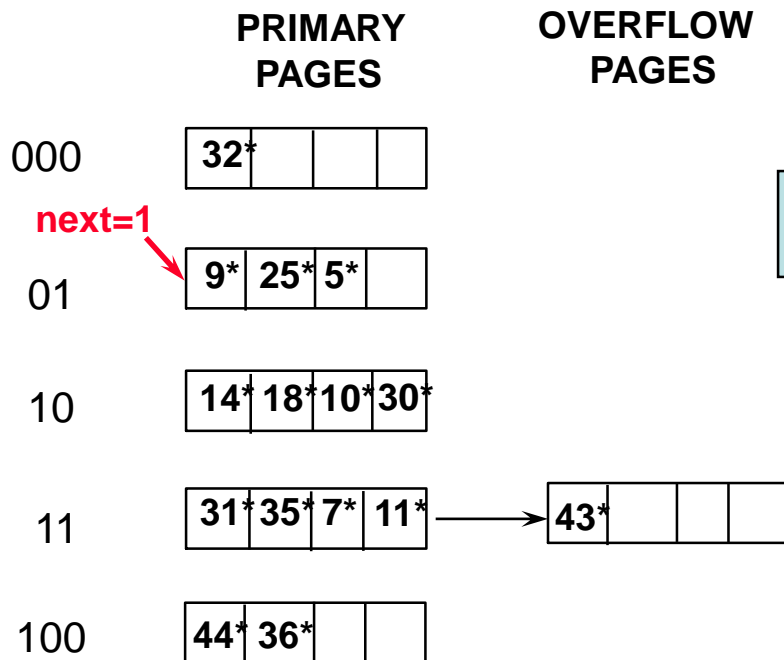
100

44*	36*		
-----	-----	--	--

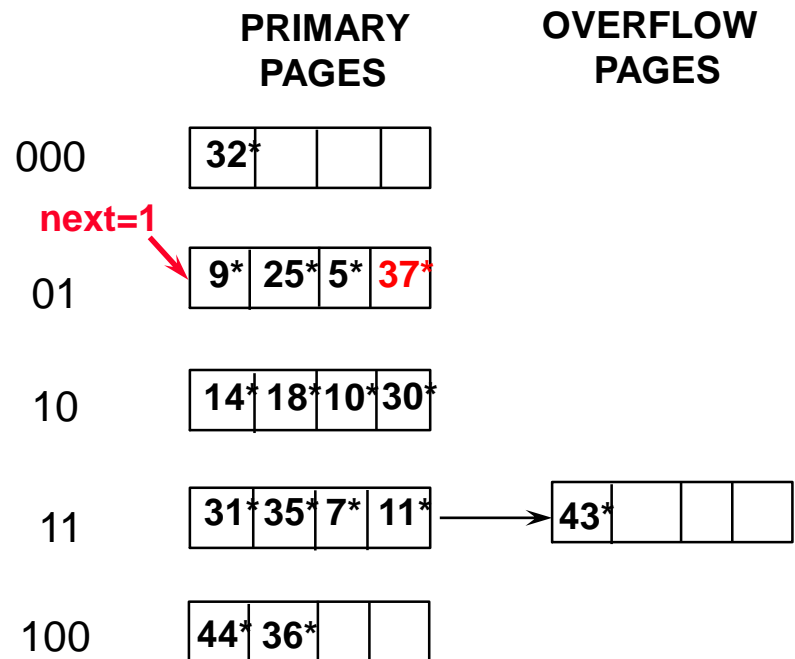
# Example of Linear Hashing

- Insert  $37^*$  (100101)
  - No overflow, no splitting
  - Insert  $37^*$  into bucket  $B_{01}$

$i=0, N_0=4$



$i=0, N_0=4$




# Example of Linear Hashing

- Insert  $29^*$  (0111**01**)
  - Bucket  $B_{01}$  overflows
  - Split bucket  $B_{01}$  ( $5 = 101$ ,  $9 = 1001$ ,  $25 = 11001$ ,  $37 = 100101$ )
  - Increment next to 2
  - Insert  $29^*$  into bucket  $B_{101}$

$i=0, N_0=4$

$i=0, N_0=4$



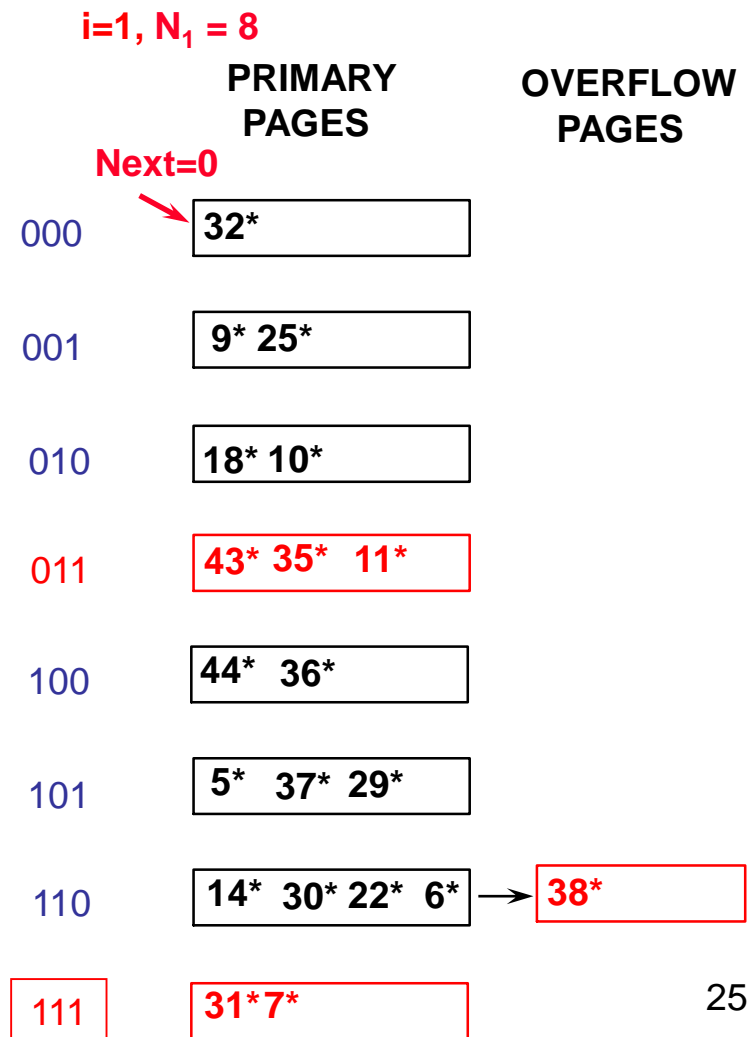
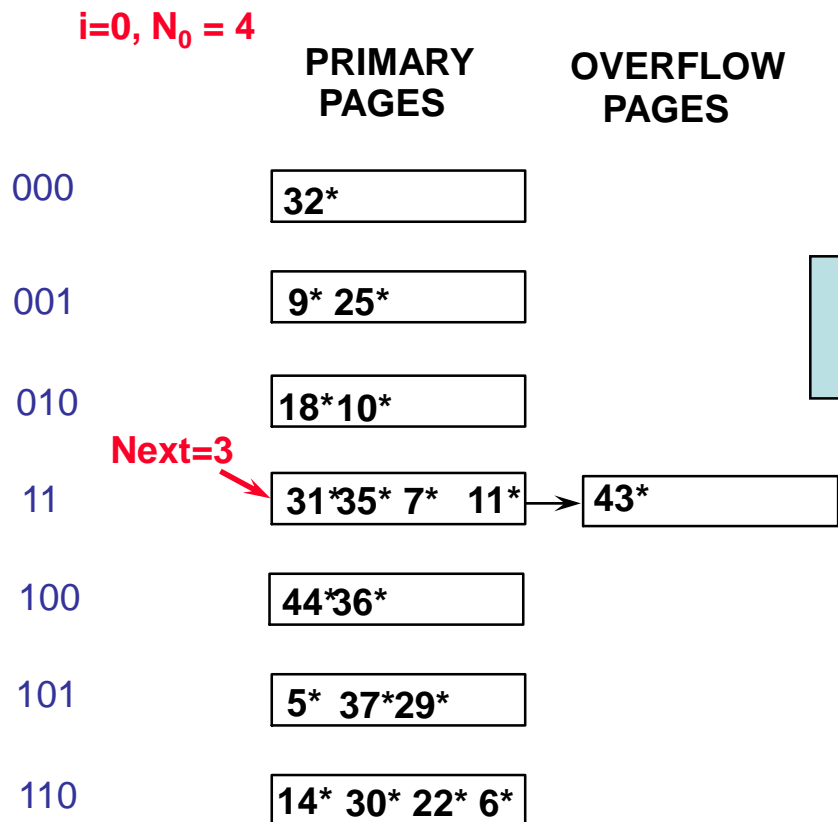
# More insertions ...

- Insert  $22^*$  (010110)
  - Bucket  $B_{10}$  overflows
  - Split bucket  $B_{10}$  ( $10 = 1\textcolor{red}{0}10$ ,  $14 = 1\textcolor{red}{1}10$ ,  $18 = 10\textcolor{red}{0}10$ ,  $30 = 11\textcolor{red}{1}10$ )
  - Increment next to 3
  - Insert  $22^*$  into bucket  $B_{110}$
- Insert  $6^*$  (000110)
  - No overflow, no splitting
  - Insert  $6^*$  into bucket  $B_{110}$





# Example: Insert 38\* (End of a Round)

- Insert 38\* (100**110**)
  - Bucket  $B_{110}$  overflows
  - Split bucket  $B_{11}$  ( $7 = 111$ ,  $11 = 1011$ ,  $31 = 11111$ ,  $35 = 100011$ ,  $43 = 101011$ )
  - Increment level to 1; reset next to 0
  - Insert 38\* into overflow page



# Linear Hashing: Deletion


- Locate bucket and delete entry
- If the last bucket  $B_{N_i + \text{next} - 1}$  becomes empty, it can be removed
- Case 1: If  $\text{next} > 0$ 
  - Decrement next




# Linear Hashing: Deletion

- Case 1: If ( $\text{next} = 0$ ) and ( $i > 0$ )
  - Update next to point to the last bucket in previous level  $B_{ni/2-1}$
  - Decrement level

$i = 1, N_0 = 2, \text{next} = 0$



$i = 0, N_0 = 2, \text{next} = 1$



# Linear Hashing: Performance

- One disk I/O unless the bucket has overflow pages
  - On average 1.2 disk I/O for uniform or lowerly skewed data distribution
  - Worst case: I/O cost is linear in the number of data entries
- Poor space utilization with skewed data distribution

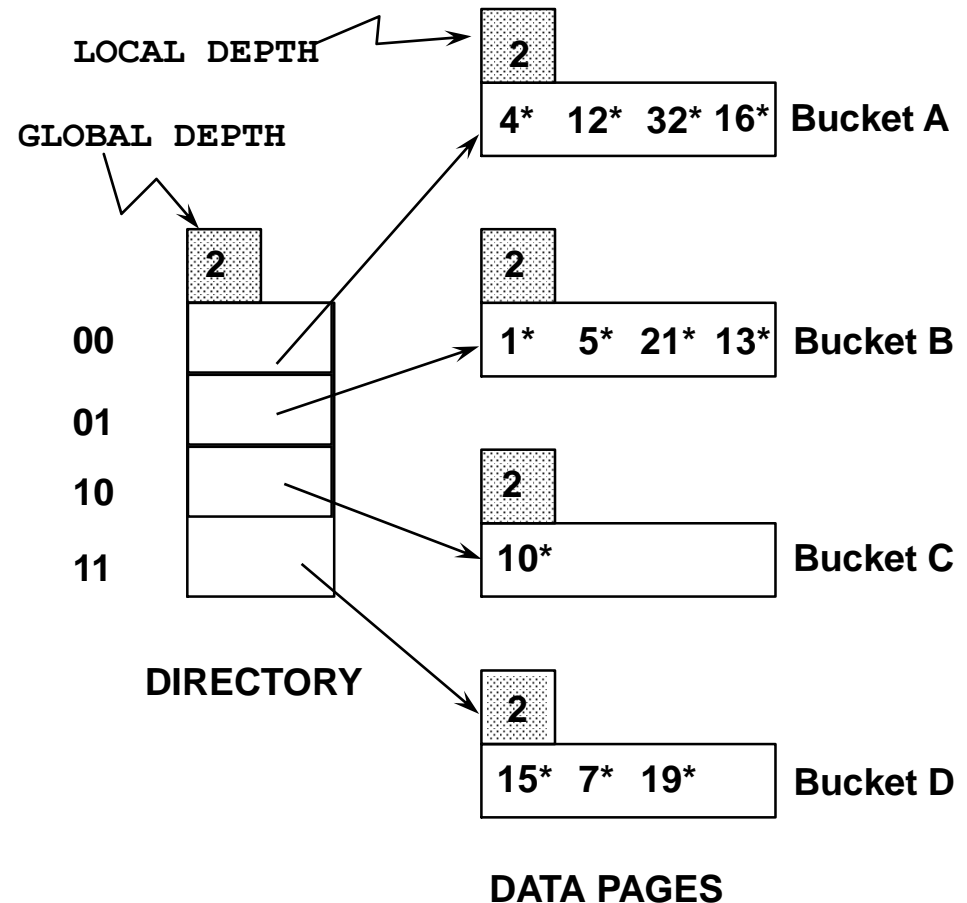
Why not simply double the file  
each time?

# Dynamic Hashing: Extendible Hashing

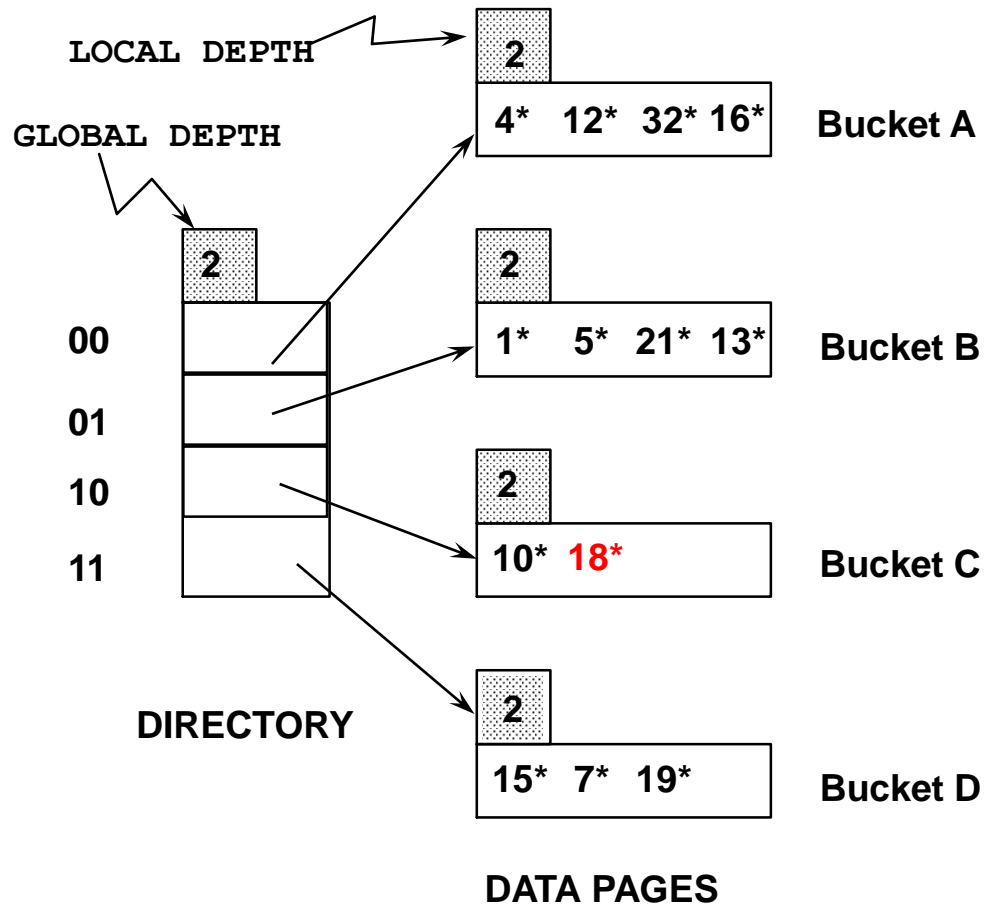
- Handle file growth by *doubling* # of buckets!
  - Idea: Use *directory of pointers to buckets*, double # of buckets by *doubling the directory* (not the actual buckets), *splitting just the bucket that overflows*!
  - Directory much smaller than file, so doubling it is much cheaper. Only one page of data entries is split.  
*No (?) overflow page!*
  - Trick lies in how hash function is adjusted!

# Example

- Directory is array of size 4
- To find bucket for  $r$ , take last '*global depth*' # bits of  $h(r)$ 
  - If  $h(r) = 5 = \text{binary } 101$ , it is in bucket pointed to by 01

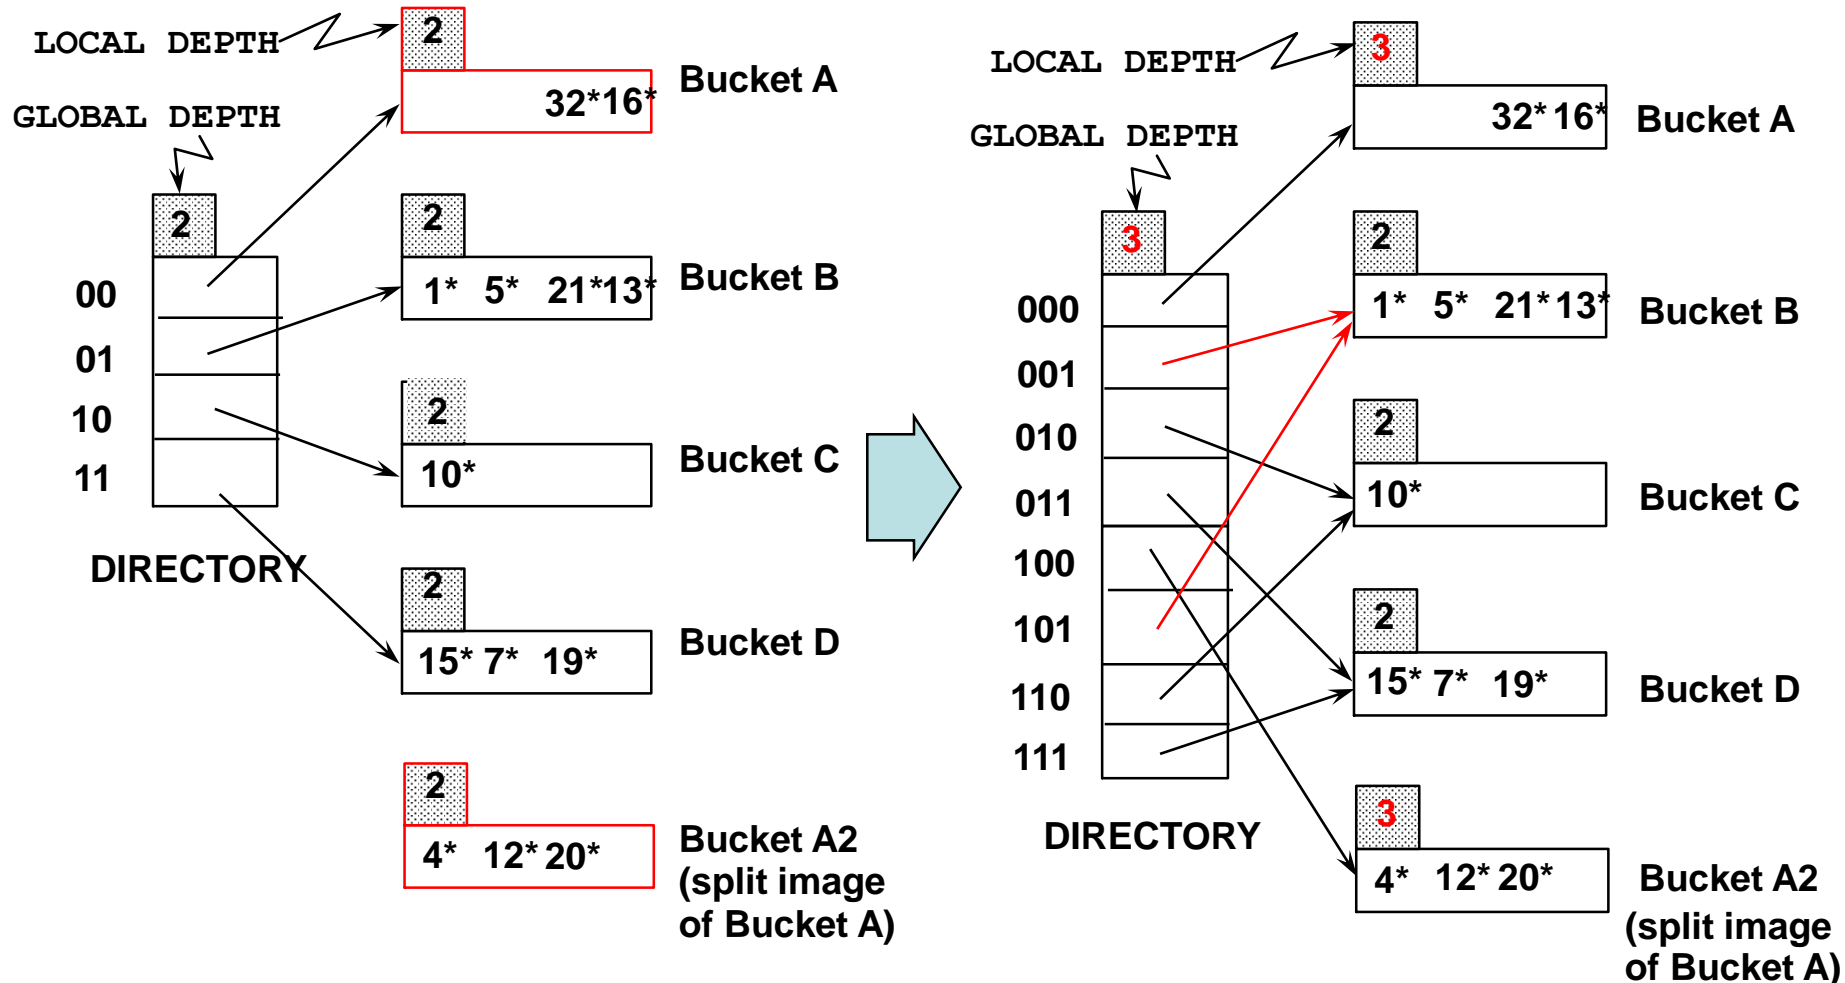


# Example: Insert 18 (10010)







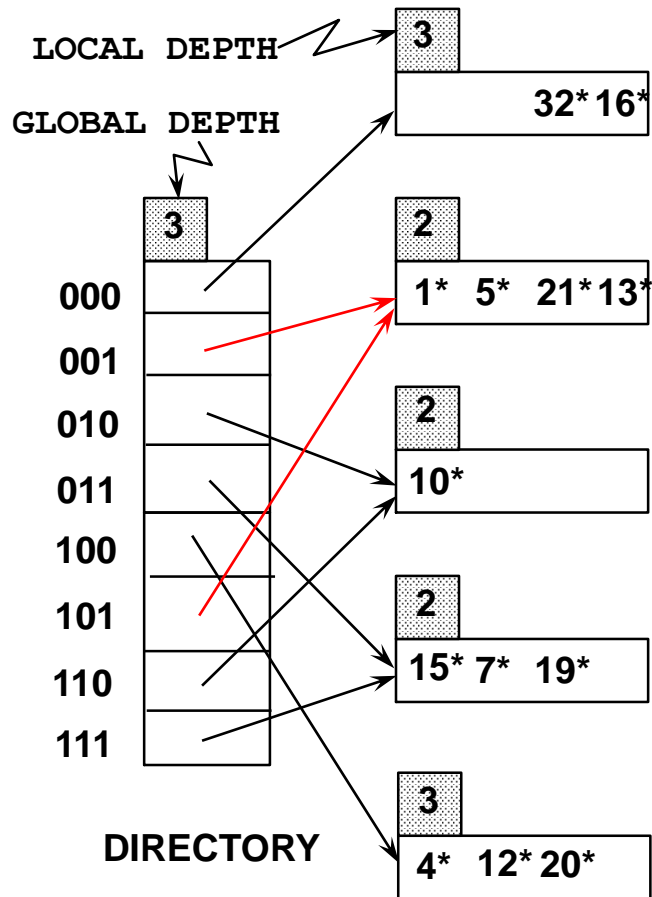
# Insert $h(r)=20$ (10100) – Causes *Doubling*



# Points to Note

- When does bucket split cause directory doubling?
  - Before insert, local depth of bucket to split = global depth
- 20 = binary 10100. Last **2** bits (00) tell us  $r$  belongs in  $A$  or  $A2$ . Last **3** bits needed to tell which.
  - *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
  - *Local depth of a bucket*: # of bits used to determine if an entry belongs to this bucket.

# Example: Insert 9 (01001)



# *Insertion*

- If bucket is full, split it (allocate new page, re-distribute)
- If the split bucket has a local depth that is smaller than the global depth
  - we need to **adjust** the directory entries to point to the right buckets
- If the split bucket's local depth is equal to the global depth
  - we need to **double** the directory as well

# Comments on Extendible Hashing

- If directory fits in memory, equality search answered with one disk access; else two.
  - 100MB file, 100 bytes/rec, 1,000,000 records, 4K pages.
    - 25,000 directory elements; chances are high that directory will fit in memory
- Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large
- **Delete:** If removal of data entry makes bucket empty, can be merged with 'split image'. If each directory element points to same bucket as its split image, can halve directory

# Summary

- Hash-based indexes: best for equality searches, cannot (?) support range searches
- Static Hashing can lead to long overflow chains
- Dynamic hashing methods such as Linear Hashing and Extendible Hashing avoid overflow pages (and hence speed up equality search cost) by splitting full buckets when new data entries are added