

# Query Optimization in Relational Database Systems

It is safer to accept any chance  
that offers itself, and extemporize  
a procedure to fit it, than to get a  
good plan matured, and wait  
for a chance of using it.



Thomas Hardy (1874)  
in *Far from the Madding Crowd*

# Query Optimization

- Since each relational op returns a relation, ops can be *composed*!
- Queries that require multiple ops to be composed may be composed in different ways - thus *optimization* is necessary for good performance, e.g.  $A \bowtie B \bowtie C \bowtie D$  can be evaluated as follows:
  - $((A \bowtie B) \bowtie C) \bowtie D$
  - $((A \bowtie B) \bowtie (C \bowtie D))$
  - $((B \bowtie A) \bowtie (D \bowtie C))$
  - ...

# Query Optimization

- Each strategy can be represented as a **query evaluation plan (QEP)** - Tree of R.A. ops, with **choice of algorithms** for each op.



- Goal of optimization: To find the “best” plan that compute the same answer (to **avoid “bad”** plans)

# ***Motivating Examples***

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)



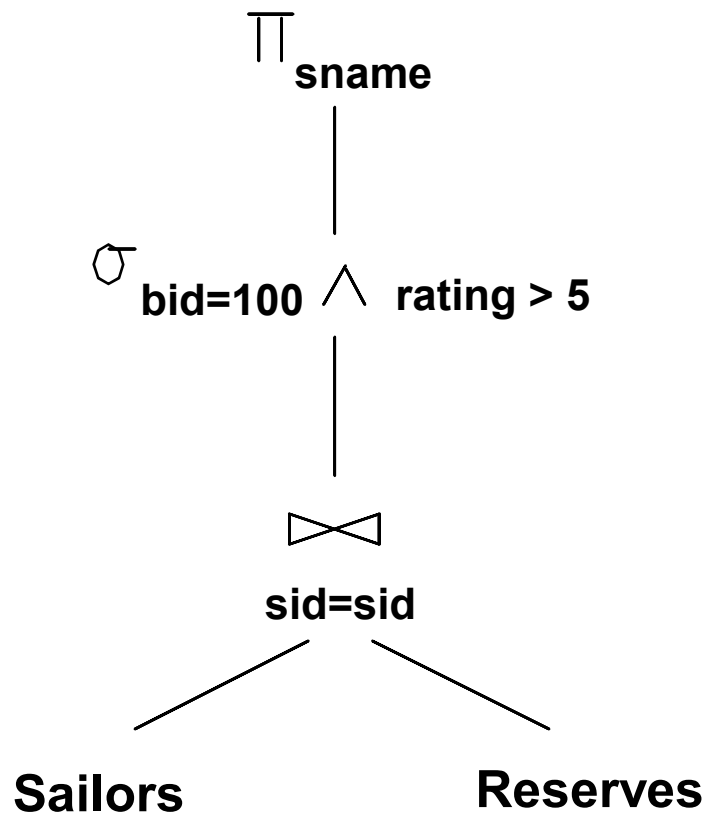
- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

```

SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5

```

# Example



Logical plans

```

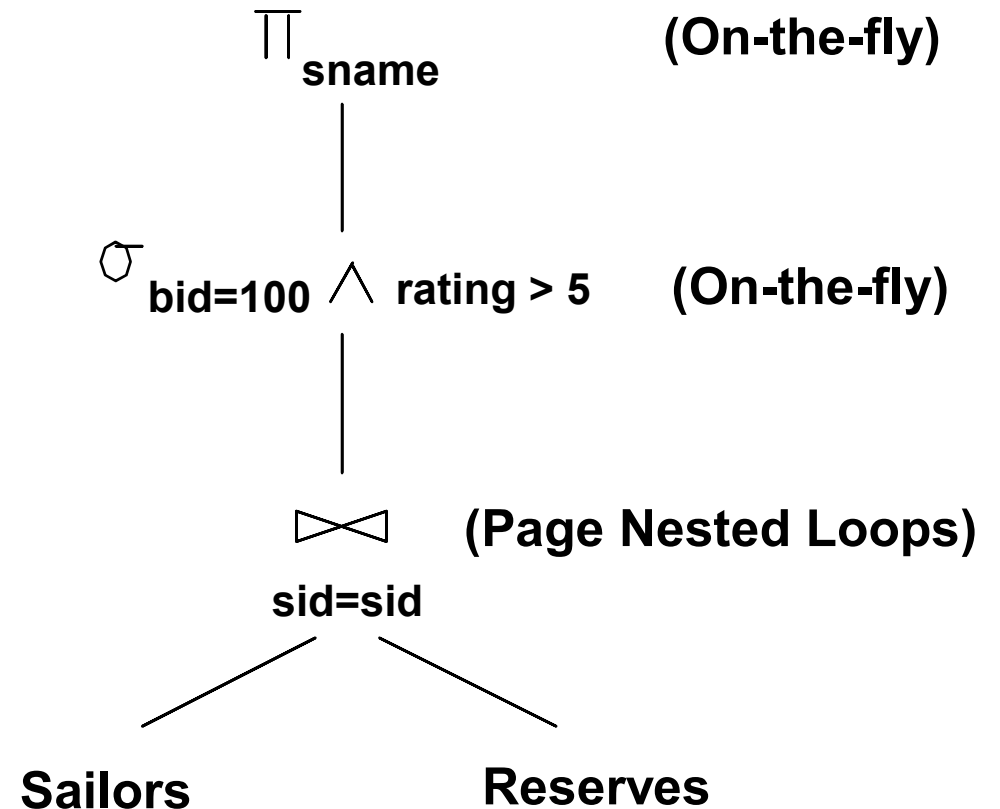
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5

```

- **Cost:**  $500 + 500 * 1000$  IO
- **Memory:** 3 pages

## Example (Cont)


### Query Evaluation Plan:



Physical plan


# Alternative Plan 1 (No Indexes)

- Main difference: push selections down
- Assume 5 buffers, T1 = 10 pages (100 boats, uniform distribution), T2 = 250 pages (10 ratings, uniform distribution)
- Cost of plan:
  - Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution)
  - Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings)
  - Sort T1 ( $2 \times 2 \times 10$ ), sort T2 ( $2 \times 4 \times 250$ ), merge (10+250)
  - Total: 4060 page I/Os




# Alternative Plan 2 (With Indexes)

- **Clustered index** on bid of Reserves
  - $100,000/100 = 1000$  tuples on  $1000/100 = 10$  pages
- Hash index on sid (format 2). Join column sid is a key for Sailors
- INL with pipelining (outer is not materialized)
  - Project out unnecessary fields from outer doesn't help
- At most one matching tuple, unclustered index on sid OK
- Did not push "rating>5" before the join. Why?
- Cost?
  - Selection of Reserves tuples (10 I/Os); for each, must get matching Sailors tuple ( $1000 \times 2.2$ ); total **2210 I/Os**





# *Query Optimizaton: Find Optimal Plan From A Set of QEPs*



# Relational Algebra Equivalences

What about  $\sigma_{p_1 \vee p_2 \dots \vee p_n}(R)$ ?

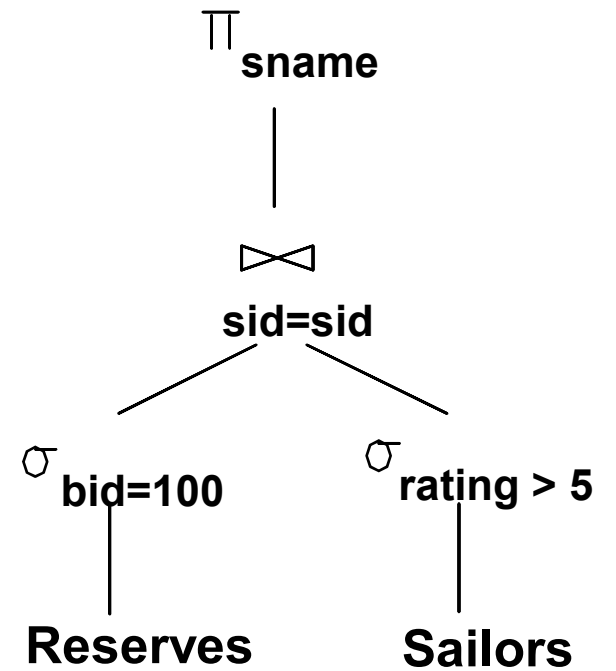

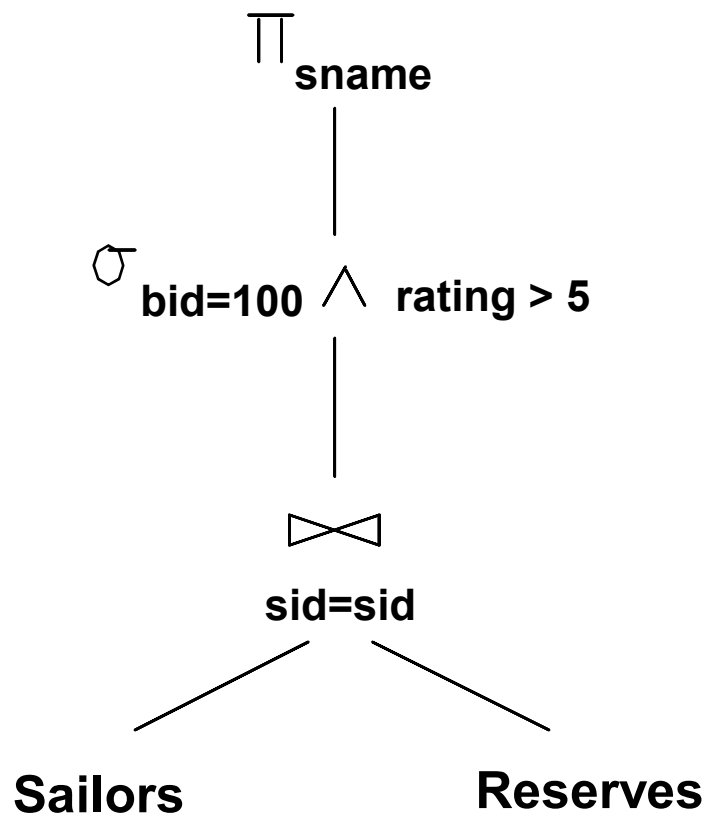
- ▶ **Cascading of selections:**  $\sigma_{p_1 \wedge p_2 \wedge \dots \wedge p_n}(R) \equiv \sigma_{p_1}(\sigma_{p_2}(\dots(\sigma_{p_n}(R))\dots))$
- ▶ **Commutativity of selections:**  $\sigma_{p_1}(\sigma_{p_2}(R)) \equiv \sigma_{p_2}(\sigma_{p_1}(R))$
- ▶ **Cascading of projections:**  $\pi_{L_1}(R) \equiv \pi_{L_1}(\pi_{L_2}(\dots(\pi_{L_n}(R))\dots))$ , where  $L_i \subseteq L_{i+1}$  for  $i \in [1, n)$
- ▶ **Commutativity of cross-products:**  $R \times S \equiv S \times R$
- ▶ **Associativity of cross-products:**  $R \times (S \times T) \equiv (R \times S) \times T$
- ▶ **Commutativity of joins:**  $R \bowtie S \equiv S \bowtie R$
- ▶ **Associativity of joins:**  $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$
- ▶ **Others:**  $R \cup S = S \cup R$ ,  $R \cap S = S \cap R$ ,  $R \cup (S \cup T) = (R \cup S) \cup T$ ,  $R \cap (S \cap T) = (R \cap S) \cap T$ , etc.

# Relational Algebra Equivalences

- ▶  $\pi_L(\sigma_p(R)) \equiv \sigma_p(\pi_L(R))$  if  $\sigma$  involves only attributes retained by  $\pi$
- ▶  $R \bowtie_p S \equiv \sigma_p(R \times S)$
- ▶  $\sigma_p(R \times S) \equiv \sigma_p(R) \times S$  if  $\sigma$  refers to attributes only in  $R$  but not in  $S$
- ▶  $\sigma_p(R \bowtie S) \equiv \sigma_p(R) \bowtie S$  if  $\sigma$  refers to attributes only in  $R$  but not in  $S$
- ▶  $\pi_L(R \times S) \equiv \pi_{L_1}(R) \times \pi_{L_2}(S)$  if  $L_1 = L \cap \text{attr}(R)$  &  $L_2 = L \cap \text{attr}(S)$
- ▶  $\pi_L(R \bowtie_p S) \equiv \pi_{L_1}(R) \bowtie_p \pi_{L_2}(S)$  if  $L_1 = L \cap \text{attr}(R)$ ,  $L_2 = L \cap \text{attr}(S)$ , & every attribute in  $p$  also appears in  $L$
- ▶ Others:  $\sigma_p(R \cup S) = \sigma_p(S) \cup \sigma_p(R)$ , etc.

# Example

$\sigma_p(R \bowtie S) \equiv \sigma_p(R) \bowtie S$  if  $\sigma$  refers to attributes only in  $R$  but not in  $S$



# *Bags vs. Sets*

$$R = \{a,a,b,b,b,c\}$$

$$S = \{b,b,c,c,d\}$$

$$R \cup S = ?$$

- SUM is implemented:  $R \cup S = \{a,a,b,b,b,b,c,c,c,d\}$
- Some rules cannot be used for bags
  - e.g.  $A \cap_s (B \cup_s C) = (A \cap_s B) \cup_s (A \cap_s C)$

Let A, B and C be  $\{x\}$




$$B \cup_B C = \{x, x\} \quad \mathbf{A \cap_B (B \cup_B C) = \{x\}}$$

$$A \cap_B B = \{x\} \quad A \cap_B C = \{x\} \quad \mathbf{(A \cap_B B) \cup_B (A \cap_B C) = \{x, x\}}$$

# Query Optimizer


- Find the “best” plan (more often avoid the bad plans)
- Comprises the following
  - Plan space
    - huge number of alternative, **semantically** equivalent plans
    - computationally expensive to examine all
    - Conventional wisdom: **avoid bad plans**
      - need to include plans that have **low cost**
  - Enumeration algorithm (**Search space**)
    - search strategy (optimization algorithm) that *searches through the plan space*
    - has to be efficient (low optimization overhead)
  - Cost model
    - facilitate comparisons of alternative plans
    - has to be “**accurate**”

# *Join Plan Notation*




# Plan Space


- Left-deep trees: right child has to be a **base** table
- Right-deep trees: left child has to be a base table
- Deep trees: one of the two children is a base table
- Bushy tree: unrestricted



Bushy tree




Left-deep tree



Deep tree



# Query Plan Space for $R \bowtie S \bowtie T$



This has not accounted for the algorithms!

# Search Algorithms (and Search Space)


- Exhaustive (*Complete space*)
  - enumerate each possible plan, and pick the best
- Greedy Techniques (*Very small – polynomial*)
  - smallest relation next
  - smallest result next
  - typically polynomial time complexity
- Randomized/Transformation Techniques (*Large space – can be complete if you run the algorithms indefinitely*)
- System R approach (*Almost complete*)
  - Dynamic Programming with Pruning

# *Multi-Join Queries*


- Focus on **multi-join queries** first
  - Join is the most expensive operations
  - Selections and projections can be **pushed down** as early as possible
- Query
  - A query graph whose nodes are relations and edges represent a join condition between the two nodes

# Greedy Algorithm (Example)


- Heuristic 1: Smallest relation next
  - Suppose  $R_i < R_k$  for  $i < k$



Join Graph



All plans must begin with  $R_1$




$R_3$   $R_4$   $R_5$

All plans beginning with  $R_2$ - $R_5$  have been pruned!

# Greedy Algorithm (Example)

- Smallest relation next
  - What if  $R_1 < R_5 < R_3 < R_2 < R_4$ ???



Another heuristic:  
Smallest result next?

no join condition, cross product

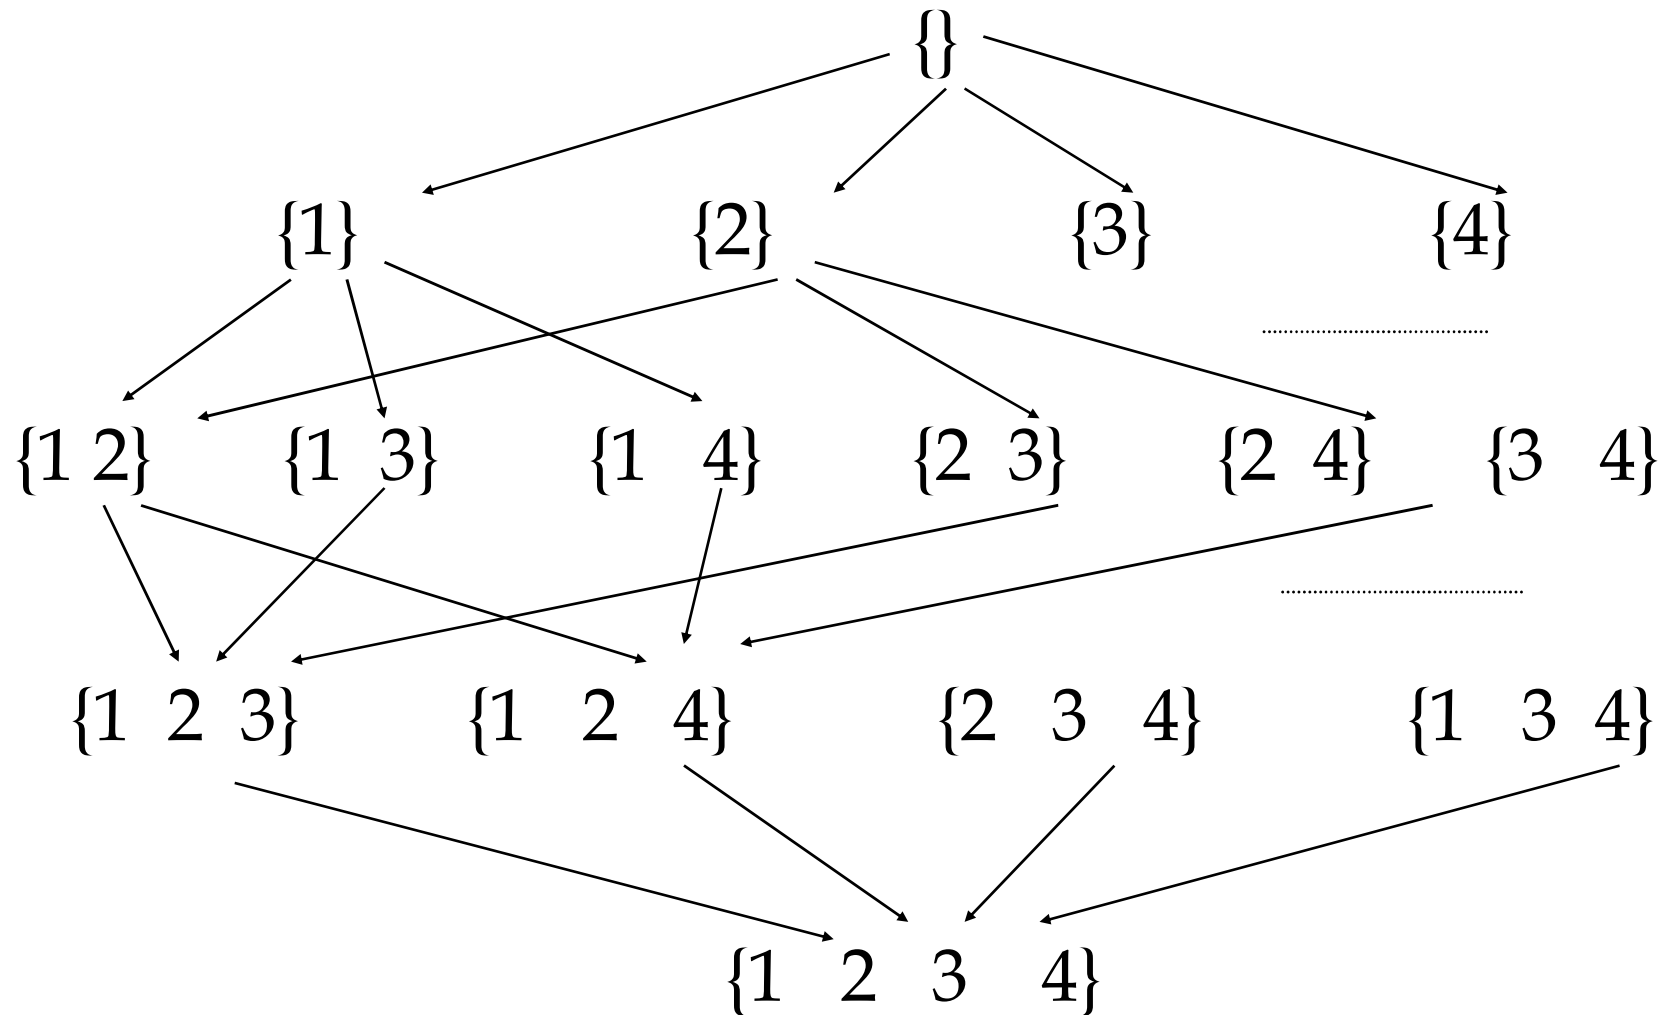
# *Dynamic Programming (**Left-Deep Trees**)*

## *(System R)*

- The algorithm proceeds by considering increasingly larger subsets of the set of all relations
  - Builds a plan bottom-up (beginning from 1 table, then to 2, and so on)
  - Plans for a set of cardinality  $i$  are constructed as extensions of the **best** plan for a set of cardinality  $i-1$ 
    - For each set of cardinality  $i$ , we only keep ONE best plan

# Dynamic Programming (Cont)

find best plan for tbl 1, index on diff attribute / table scan




# Dynamic Programming (*Left-Deep Trees*) (*System R*)

- The algorithm proceeds by considering increasingly larger subsets of the set of all relations
  - Builds a plan bottom-up (beginning from 1 table, then to 2, and so on)
  - Plans for a set of cardinality  $i$  are constructed as extensions of the **best** plan for a set of cardinality  $i-1$ 
    - Keep only **ONE** best plan for each set of cardinality  $i$
- Search space can be pruned based on the **principal of optimality**
  - if two plans differ only in a subplan, then the plan with the better subplan is also the better plan



# *Principle of Optimality*



# Dynamic Programming (**Left-Deep Trees**) (**System R**)

- The algorithm proceeds by considering increasingly larger subsets of the set of all relations
  - Builds a plan bottom-up (beginning from 1 table, then to 2, and so on)
  - Plans for a set of cardinality  $i$  are constructed as extensions of the **best** plan for a set of cardinality  $i-1$ 
    - Keep only **ONE** best plan for each set of cardinality  $i$
- Search space can be pruned based on the **principal of optimality**
  - if two plans differ only in a subplan, then the plan with the better subplan is also the better plan
- Computation overhead reduced due to **overlapping subproblems**
  - Multiple sets of cardinality  $i$  uses same set at cardinality  $(i-1)$


# *Dynamic Programming (Left-Deep Trees)*

- `accessPlan(R)` produces the best plan for relation (single table) `R`
- `joinPlan(p1,R)` extends the (partial) **join plan p1** into another plan `p2` in which the result of `p1` is joined with `R` in the best possible way
  - `p1 = R1 JOIN R2 JOIN R3`
  - `p2 = joinPlan(p1, R) = (R1 JOIN R2 JOIN R3) JOIN R4`
- Optimal plans for subsets are stored in `optplan()` array and are reused rather than recomputed

# Dynamic Programming (Cont)

```

for i = 1 to N
  optPlan({Ri}) = accessPlan(Ri)
for i = 2 to N {
  for all subset of table size i
  forall S subset of {R1, R2, ... Rn} such that |S|=i {
    bestPlan = dummy plan with infinite cost
    forall Rj, Sj, |Sj| = i-1 such that S = {Rj} U Sj {
      p = joinPlan(optPlan(Sj), Rj)
      if cost(p) < cost(bestPlan)
        bestPlan = p
    }
    optPlan(S) = bestPlan
  }
}
Popt = optPlan{R1, R2, ... Rn}
  
```



# Dynamic Programming: A Concrete Example

► Schema: **R**(A,B,C,D), **S**(X,Y), **T**(E,F,G)

► Query:

```
select *  
from   R join S on R.A = S.X join T on R.D = T.F  
where  R.B > 10  
and    R.C = 20  
and    T.E < 100
```

► Available B<sup>+</sup>-tree indexes:  $I_B$ ,  $I_C$ ,  $I_E$

► Assumptions on database system

- Supports only one join algorithm: hash join
- Avoids cartesian products

$$\sigma_p(R \bowtie_{R.A=S.X} S \bowtie_{R.D=T.F} T), p = (R.B > 10) \wedge (R.C = 20) \wedge (T.E < 100)$$

# Enumeration of Single-relation Plans

$\sigma_p(R \bowtie_{R.A=S.X} S \bowtie_{R.D=T.F} T), p = (R.B > 10) \wedge (R.C = 20) \wedge (T.E < 100)$

## ► Plans for {R}

- **Plan P1**: Table scan with “ $(B > 10) \wedge (C = 20)$ ”
- **Plan P2**: Index seek with  $I_B$  & RID-lookups with “ $C = 20$ ”
- **Plan P3**: Index seek with  $I_C$  & RID-lookups with “ $B > 10$ ”
- **Plan P4**: Index intersection with  $I_B$  &  $I_C$ , and RID-lookups
- Assume  $cost(P3) < cost(P4) < cost(P2) < cost(P1)$
- $optPlan(\{R\}) = P3$

## ► Plans for {S}

- **Plan P5**: Table scan of S
- $optPlan(\{S\}) = P5$


## ► Plans for {T}

- **Plan P6**: Table scan of T with “ $(E < 100)$ ”
- **Plan P7**: Index seek with  $I_E$  & RID-lookups
- Assume  $cost(P7) < cost(P6)$
- $optPlan(\{T\}) = P7$

# Enumeration of Two-Relation Plans


$\sigma_p(R \bowtie_{R.A=S.X} S \bowtie_{R.D=T.F} T), p = (R.B > 10) \wedge (R.C = 20) \wedge (T.E < 100)$

## ► Plans for {R, S}



- Assume  $cost(P8) < cost(P9)$
- $optPlan(\{R, S\}) = P8$



## ► Plans for {R, T}



- Assume  $cost(P11) < cost(P10)$
- $optPlan(\{R, T\}) = P11$

# Enumeration of Three-Relation Plans

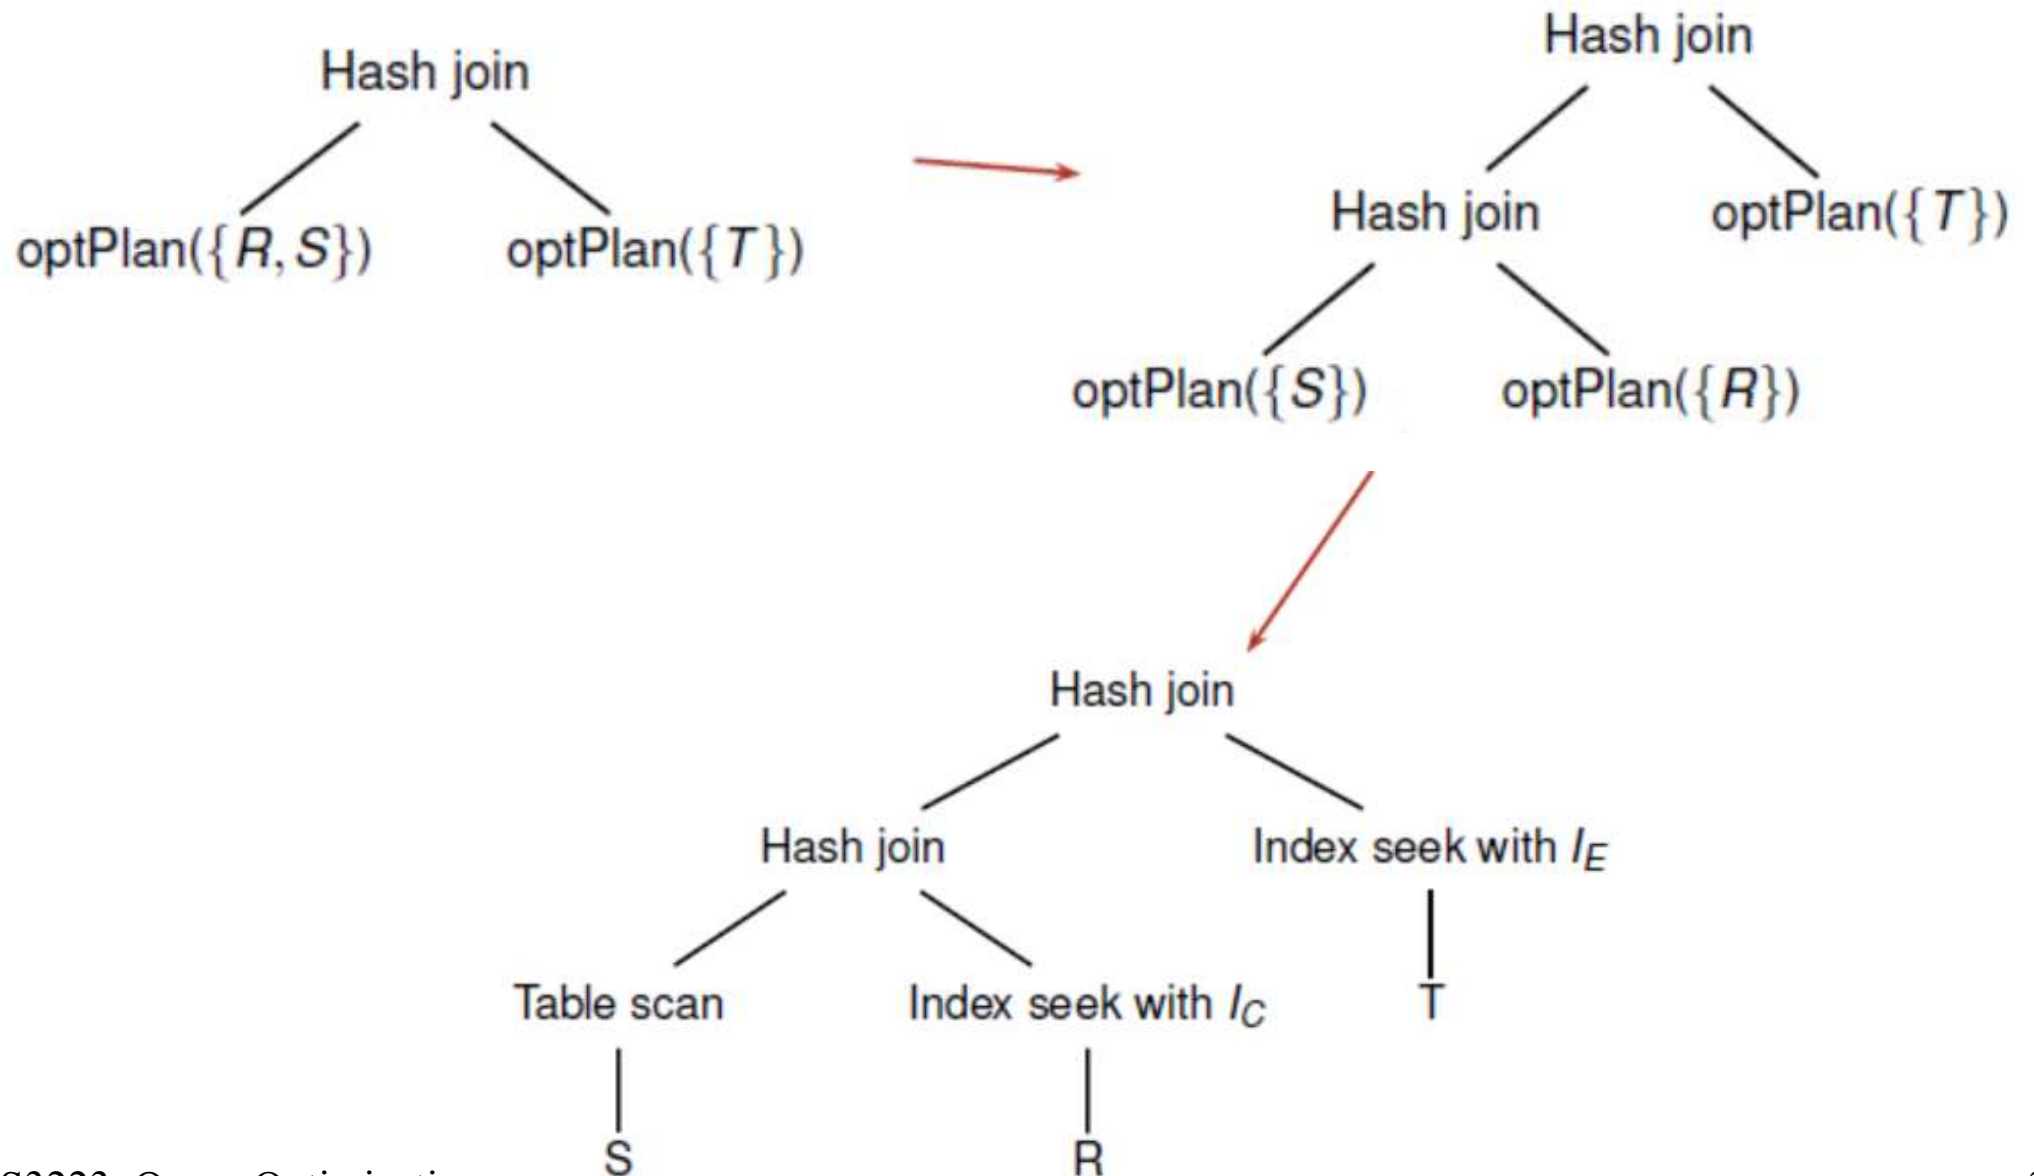
$\sigma_p(R \bowtie_{R.A=S.X} S \bowtie_{R.D=T.F} T), p = (R.B > 10) \wedge (R.C = 20) \wedge (T.E < 100)$



- ▶ Assume  $cost(P12) < cost(P13)$
- ▶  $optPlan(\{R, S, T\}) = P12$



# Optimal Plan



# ***Dynamic Programming (Cont)***

- Time & Space complexity
  - For  $k$  relations, for left-deep trees,  $2^k - 1$  entries!
  - For bushy trees,  $O(3^k)$
- Is DP (as presented) optimal?

# *Dynamic Programming (Cont)*

- Time & Space complexity
  - For  $k$  relations, for left-deep trees,  $2^k - 1$  entries!
  - For bushy trees,  $O(3^k)$
- Is DP optimal?
- DP may maintain **multiple plans** per subset of relations
  - **Interesting orders**




# *Dynamic Programming (Cont)*

- Time & Space complexity
  - For  $k$  relations, for left-deep trees,  $2^k - 1$  entries!
  - For bushy trees,  $O(3^k)$
- Is DP optimal?
- DP may maintain **multiple plans** per subset of relations
  - **Interesting orders**
- **Is DP with interesting orders optimal?**

# *Randomized Techniques*

- Employ randomized/transformation techniques for query optimization
- **State space** -- space of plans, **State** -- plan
- Each state has a **cost** associated with it
  - determined by some cost model
- A **move** is a perturbation applied to a state to get to another state
  - a move set is the set of moves available to go from one state to another
  - any one move is chosen from this move set **randomly**
  - each move set has a probability associated to indicate the probability of selecting the move
- Two states are **neighboring states** if one move suffices to go from one state to the other

# Randomized Algorithm (Example)



# *More on Randomized Techniques*

- A **local minimum** in the state space is a state such that its cost is lower than that of all neighboring states
- A **global minimum** is a state which has the lowest cost among all local minima
  - at most one global minimum
- A move that takes one state to another state with a lower cost is called a **downward move**; otherwise it is an **upward move**
  - in a local/global minimum, all moves are upward moves



# Local Optimization

Repeat until  
a near-optimal  
minimum  
is reached

By doing so  
repeatedly,  
a local minimum  
can  
be reached

```
S = initialize() // initial plan
minS = S // cost of plan S – currently the best
repeat {
  repeat {
    newS = move(S) // move to a new plan
    if (cost(newS) < cost(S))
      S = newS
  } until ("local minimum reached")
  if (cost(S) < cost(minS))
    minS = S
  newStart(S); // iterate with a different initial plan
} until ("stopping condition satisfied")
return (minS);
```

A move is accepted if it is a downward move, i.e., has a lower cost

# *Issues on Local Optimization*

- How is the start state obtained?
  - The state in which we start a run
  - The start state of the first run is the initial state
  - **All start states should be different**
  - Should be obtained quickly
    - Random
    - greedy heuristics
    - making a number of moves from the local minimum, except that this time each move is accepted irrespective of whether it increases or decreases the cost
- How is the local minimum detected?
- How is the stopping criterion detected?

*Run*: sequence of moves to a local minimum from the start state

# *Issues on Local Optimization (Cont)*

- How is the local minimum detected?
  - Not practical to examine all neighbors to verify that one has reached a local minimum
  - Based on random sampling
    - examine a sufficiently large number of neighbors
      - if any one is lower, we move to that state, and repeat the process
      - if no tested neighbor is of lower cost, the current state can be considered a local minimum
    - the number of neighbors to examine can be specified as a parameter, and is called the **sequence length**
      - Can also be time-based

# *Issues in Local Optimization (Cont)*

- How is the stopping criterion detected?
  - Determines the number of times that the outer loop is executed
  - Can be fixed and is given by  $\text{sizeFactor} * N$ , where  $\text{sizeFactor}$  is a parameter,  $N$  is the number of relations
    - Why  $N$ ? Can it be a constant?

# ***What about the MOVEs?***

## ***Transformation Rules***

- Restricted to left-deep trees
  - all possible permutations of the N relations
  - let S be the current state,  $S = (... i ... j ... k ...)$
  - swap
    - select two relations, say i and j at random. Swapped i and j to get the new state  $newS = (... j ... i ... k ...)$
  - 3Cycle
    - select three relations, say i and j and k at random. The move consists of cycling i, j and k: i is moved to the position of j, j is moved to the position of k and k is moved to the position of i. The resultant new state  $newS = (... k ... i ... j ...)$
- Other methods (e.g., join methods)? Bushy trees?

# Comparison between Exhaustive, Greedy and Randomized Algorithms


- Search space
- Plan quality
- Optimization overhead

# Cost Models

- Typically, a combination of CPU and I/O costs
- Objective is to be able to rank plans
  - exact value is not necessary
- Relies on
  - statistics on relations and indexes
  - formulas to estimate CPU and I/O cost
  - formulas to estimate **selectivities** of operators and intermediate results

# Cost Estimation

- For each plan considered, must estimate cost:
  - Must **estimate cost** of each operation in plan tree
    - Depends on input cardinalities
    - Depends on buffer size, availability of indexes, algorithms used, etc.
      - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
  - Must **estimate size of result** for each operation in tree!
    - Use information about the input relations
    - Typical assumptions like *uniform distribution* of data and *independence* of predicates can simplify size estimation but is *error prone*





# ***Statistics and Catalogs***

- Need information about the relations and indexes involved  
**Catalogs** typically contain at least:
  - # tuples of  $R$  ( $||R||$ ), #bytes in each  $R$  tuple ( $S(R)$ )
  - # blocks/pages to hold all  $R$  tuples ( $|R|$ )
  - # distinct values in  $R$  for attribute  $A$  ( $V(R,A)$ )
  - NPages for each index
  - Index height, low/high key values (Low/High) for each tree index
- Catalogs updated periodically
  - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok

# *Estimation Assumptions*

- Uniformity assumption
  - Uniform distribution of attribute values
- Independence assumption
  - Independent distribution of values in different attributes
- Inclusion assumption
  - For  $R \bowtie_{R.A=S.B} S$ , if  $V(R, A) \leq V(R, B)$ , then  $\pi_A(R) \subseteq \pi_B(S)$
  - $V(R, A)$  is the number of distinct values of  $R.A$

# Example

R

A	B	C	D
cat	1	10	a
cat	1	20	b
dog	1	30	a
dog	1	40	c
bat	1	50	d

A: 20 byte string

B: 4 byte integer

C: 8 byte string

D: 5 byte string

$$\|R\| = 5 \quad S(R) = 37$$

$$V(R,A) = 3 \quad V(R,C) = 5$$

$$V(R,B) = 1 \quad V(R,D) = 4$$

# Size estimate for $W = \sigma_{Z=\text{val}}(R)$

R	A	B	C	D
	cat	1	10	a
	cat	1	20	b
	dog	1	30	a
	dog	1	40	c
	bat	1	50	d

$$V(R,A)=3$$

$$V(R,B)=1$$

$$V(R,C)=5$$

$$V(R,D)=4$$

$$||W|| = \frac{||R||}{V(R,Z)}$$

$$S(W) = S(R)$$

Assumption:

Values in select expression  $Z = \text{val}$  are *uniformly distributed* over possible  $V(R,Z)$  values

Alternative assumption: use  $\text{DOM}(R,Z)$

***What about  $W = \sigma_{z \geq \text{val}} (R)$ ?***

Solution: Estimate values in range

R	Z		
		Min=1	$V(R,Z)=10$
		↕	$W = \sigma_{z \geq 15} (R)$
		Max=20	

$$f \text{ (fraction of range)} = \frac{20-15+1}{20-1+1} = \frac{6}{20} \quad ||W|| = f \times ||R||$$

Alternative:  $(\text{Max}(Z) - \text{value}) / (\text{Max}(Z) - \text{Min}(Z))$

$$W = R1 \bowtie R2$$

R1	A	B	C

R2	A	D


Assumption:

$V(R1,A) \leq V(R2,A) \Rightarrow$  Every A value in R1 is in R2

$V(R2,A) \leq V(R1,A) \Rightarrow$  Every A value in R2 is in R1

“containment of value sets”

# Computing $T(W)$ when $V(R1,A) \leq V(R2,A)$



1 tuple of R1 matches with  $\frac{||R2||}{V(R2,A)}$  tuples of R2

$$\text{so } ||W|| = ||R1|| \times \frac{||R2||}{V(R2,A)}$$

$$\text{If } V(R2,A) \leq V(R1,A) \quad ||W|| = \frac{||R2|| \times ||R1||}{V(R1,A)}$$

***For complex expressions, need  
intermediate T,S,V results***

$$\text{E.g. } W = \underbrace{[\sigma_{A=a}(R1)]}_{\text{Treat as relation U}} \bowtie R2$$

Treat as relation U

$$||U|| = ||R1||/V(R1,A) \quad S(U) = S(R1)$$

Also need  $V(U, *)$  !!



# Example

R1

A	B	C	D
cat	1	10	10
cat	1	20	20
dog	1	30	10
dog	1	40	30
bat	1	50	10

$$V(R1, A) = 3$$

$$V(R1, B) = 1$$

$$V(R1, C) = 5$$

$$V(R1, D) = 3$$

$$U = \sigma_{A=a}(R1)$$

$$V(U, A) = ? \quad V(U, B) = ?$$

$$V(U, C) = ?$$

$V(D, U)$  ... somewhere in between  $V(U, B)$  and  $V(U, C)$

***For Joins*     $U = R1(A,B) \bowtie R2(A,C)$**

$$V(U,A) = \min \{ V(R1, A), V(R2, A) \}$$

$$V(U,B) = V(R1, B)$$

$$V(U,C) = V(R2, C)$$

(Assumption: Preservation of value sets)

## *Example*

$$Z = R1(A,B) \bowtie R2(B,C) \bowtie R3(C,D)$$

R1	$\ R1\  = 1000$	$V(R1,A)=50$	$V(R1,B)=100$
----	-----------------	--------------	---------------

R2	$\ R2\  = 2000$	$V(R2,B)=200$	$V(R2,C)=300$
----	-----------------	---------------	---------------

R3	$\ R3\  = 3000$	$V(R3,C)=90$	$V(R3,D)=500$
----	-----------------	--------------	---------------

$$Z = R1(A,B) \bowtie R2(B,C) \bowtie R3(C,D)$$

$$\|R1\| = 1000 \quad V(R1,A)=50 \quad V(R1,B)=100$$

$$\|R2\| = 2000 \quad V(R2,B)=200 \quad V(R2,C)=300$$

$$\|R3\| = 3000 \quad V(R3,C)=90 \quad V(R3,D)=500$$

**Partial Result:  $U = R1 \bowtie R2$**

$$\|U\| = \frac{1000 \times 2000}{200}$$

$$V(U,A) = 50$$

$$V(U,B) = 100$$

$$V(U,C) = 300$$

**$Z = U \bowtie R3$**

$$\|Z\| = \frac{1000 \times 2000 \times 3000}{200 \times 300}$$

$$V(Z,A) = 50$$

$$V(Z,B) = 100$$

$$V(Z,C) = 90$$

$$V(Z,D) = 500$$


# ***Errors in Estimating Size of Plan***

- Errors

- source include uniformity assumption, and inability to capture correlation, accuracy of cost model, statistics, etc.
- *propagated to other operators at the higher level of the plan tree*

- Dealing with errors

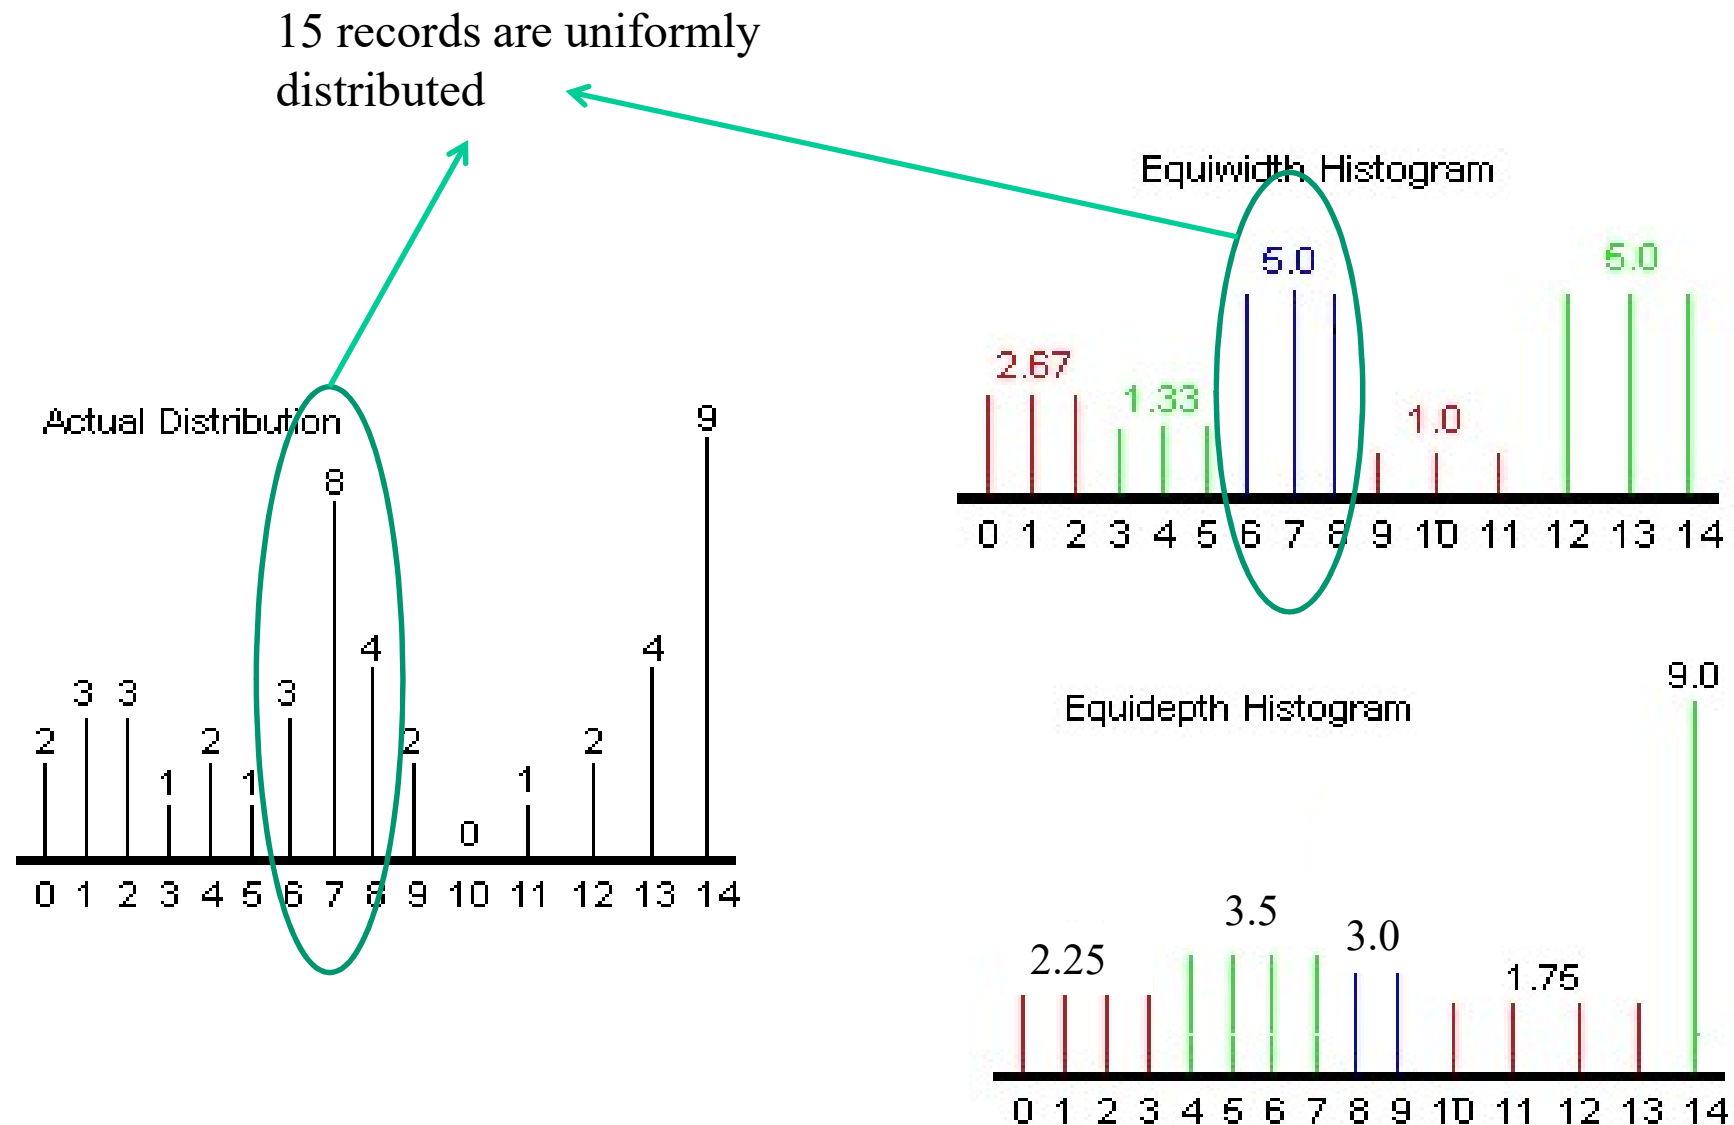
- Maintain more detailed statistics (at finer granularity)
- During runtime, may need to sample the actual intermediate results
  - *dynamic* query optimization




# Statistical Summaries of Data

- More detailed information are sometimes stored e.g., **histograms** of the values in some attributes
  - a histogram divides the values on a column into  $k$  buckets
    - $k$  is predetermined or computed based on space allocation
  - several choices for “bucketization” of values
    - If a table has  $n$  records, an **equi-depth** histogram divides the set of values on a column into  $k$  ranges such that **each range** has *approximately* the **same number of records**, i.e.,  $n/k$
    - **Equi-width** histogram – **each bucket** has (almost) **equal number of values**
    - Within each bucket, records are uniformly distributed across the range of the bucket
    - Frequently occurring values may be placed in singleton buckets

# Histograms



# *Estimations with Histograms*



Query Q:  $\sigma_{A=6} (R)$

Actual value,  $\|Q\| = 3$

Without histogram,  $\|Q\| = 45/15 = 3$

Equiwidth histogram,  $\|Q\| = 15/3 = 5$

Equidepth histogram,  $\|Q\| = 14/4 = 3.5$

Query Q:  $\sigma_{A=10} (R)$

Actual value,  $\|Q\| = 0$

Without histogram,  $\|Q\| = 45/15 = 3$

Equiwidth histogram,  $\|Q\| = 1$

Equidepth histogram,  $\|Q\| = 1.75$



# *Summary*

- Query optimization is NP-hard
- Instead of finding the best, the objective is largely to avoid the bad plans
- Many different optimization strategies have been proposed
  - greedy heuristics are fast but may generate plans that are far from optimal
  - dynamic programming is effective at the expense of high optimization overhead