

Admin

- Lab 2 due today!
- Tutorial participation
- Telegram Group Link: <u>https://t.me/+pnE_EW3wsl0yMmM1</u>





Quick Review
5-10 min

Tutorial Questions
30 – 40 min

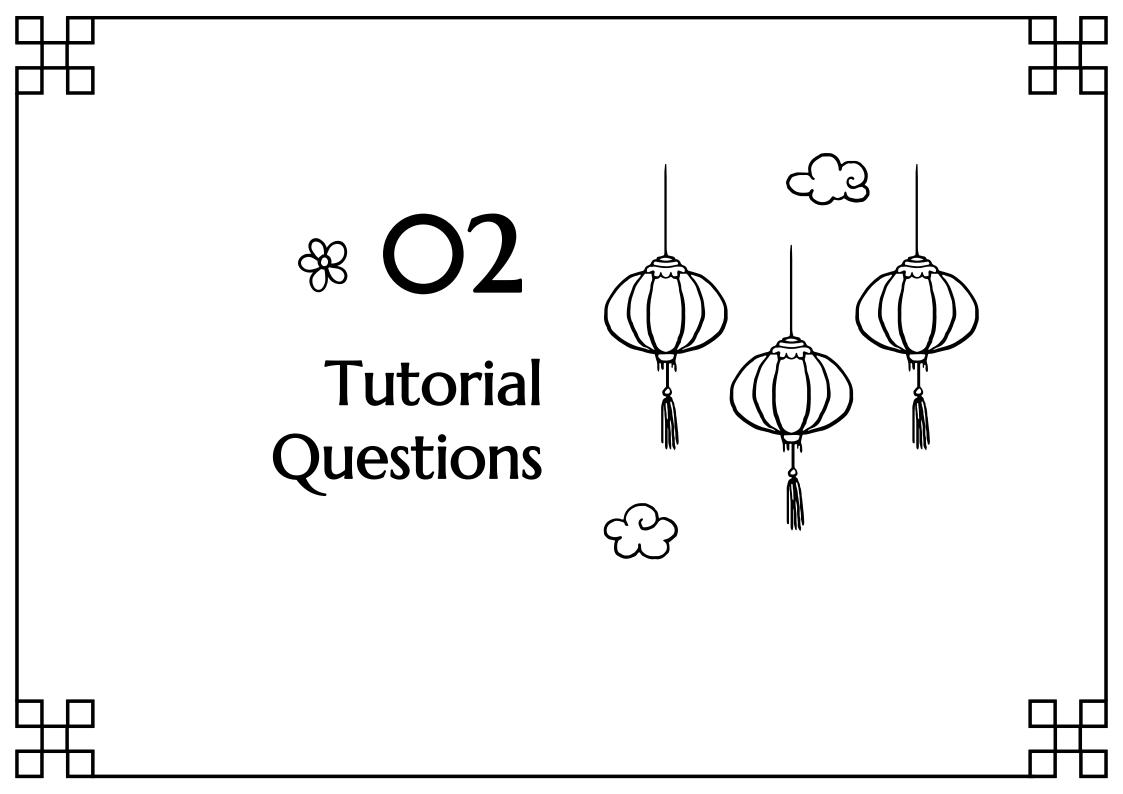




Q&A/Extra Questions
5-10 min



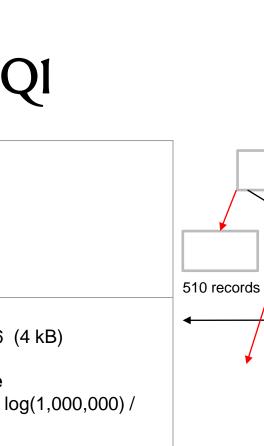
Hash Index

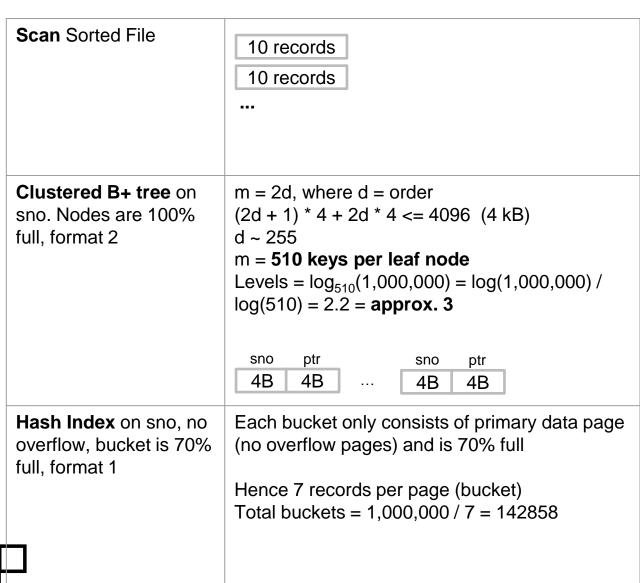


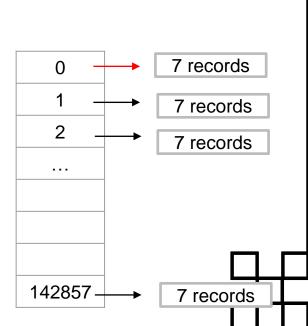
- 1. Consider a relation EMP(eno, sno, name, salary) containing 1,000,000 employee records. Suppose each page of the relation contains 10 records, and relation EMP is organized as a sorted file. Suppose sno is a candidate key, with values lying in the range 0-999,999, and that EMP is stored in sno order. Assume that each page is 4K bytes, eno and sno are each 4 bytes, and a pointer to a page is also 4 bytes. For each of the following queries, state which of the following three approaches is most likely to be the cheapest (in terms of I/O count). Justify your answer. For simplicity, assume that the data are uniformly distributed in the domain.
 - Scan the sorted file for EMP.
 - Use a (clustered) B+-tree index on sno. Assume that the nodes are 100% full, and that format 2 is used for the index.
 - Use a hash index on sno. You may assume ideal situation of no overflows, and each bucket is 70% full (and a minimum number of buckets is used). You may assume that format 1 is used for the index.
 - a) Find all employees whose sno < 100,000.
 - b) Find the employee whose sno = 10.
 - c) Find the employees whose sno is in the range 10,000-10,010.
 - d) Find the employees whose sno is not 10.

Scan Sorted File	10 records 10 records
Clustered B+ tree on sno. Nodes are 100% full, format 2	m = 2d, where d = order $(2d + 1) * 4 + 2d * 4 <= 4096 (4 kB)$ d ~ 255 m = 510 keys per leaf node Levels = $log_{510}(1,000,000) = log(1,000,000) / log(510) = 2.2 = approx. 3$
Hash Index on sno, no overflow, bucket is 70% full, format 1	Each bucket only consists of primary data page (no overflow pages) and is 70% full Hence 7 records per page (bucket) Total buckets = 1,000,000 / 7 = 142858

- **1,000,000** employee records.
- **10 records** per page of relation (data page)
- sorted on sno
- sno is a candidate key, with values lying in the range 0-999,999
- 4K bytes per page (B+ tree node / Hash Index primary/overflow page)
- eno and sno are each 4 bytes
- apointer to a page is 4 bytes
- data are uniformly distributed in the domain.







1 root node

510 records

1,000,000 records

~1961 leaf nodes

510 records

~ 4 internal

510 records

nodes

- **1,000,000** employee records.
- 10 records per page of relation (data page)
- sorted on sno
- sno is a candidate key, with values lying in the range 0-999,999
- **4K bytes** per page (B+ tree node / Hash Index bucket)
- eno and sno are each 4 bytes, and a
- pointer to a page is 4 bytes
- data are uniformly distributed in the domain.
- a) Find all employees whose sno < 100,000.

Scan Sorted File	10 records per page, 100,000 records in total 10,000 I/O
Clustered B+ tree on sno. Nodes are 100% full, format 2	10,000 I/O + additional I/O to access B+ tree nodes (each node is a page)
Hash Index on sno, no overflow, bucket is 70% full, format 1	142,858 I/O by scanning all

- **1,000,000** employee records.
- 10 records per page
- sorted on sno
- sno is a candidate key, with values lying in the range 0-999,999
- 4K bytes per page
- eno and sno are each 4 bytes, and a
- pointer to a page is 4 bytes
- data are uniformly distributed in the domain.
- b) Find the employee whose sno = 10.

Scan Sorted File	If we have access to the first page of the file, and it is chained, then we need 2 pages at most. But consider the general case where sno can be anything, e.g. sno = 429580
Clustered B+ tree on sno. Nodes are 100% full, format 2	4 I/O. One I/O to retrieve the index, two I/O to reach leaf node, one I/O to retrieve data page
Hash Index on sno, no overflow, bucket is 70% full, format 1	2 I/O. One I/O to retrieve the index, the other to access the primary page

Ql

- **1,000,000** employee records.
- 10 records per page
- sorted on sno
- sno is a candidate key, with values lying in the range 0-999,999
- 4K bytes per page
- eno and sno are each 4 bytes, and a
- pointer to a page is 4 bytes
- data are uniformly distributed in the domain.
- c) Find the employees whose sno is in the range 10,000-10,010.

Scan Sorted File	In the best case, we scan till sno = 10,000 1001 I/O					
Clustered B+ tree on sno. Nodes are 100% full, format 2	4 + 1 = 5 I/O. Four I/O to get to the data page consisting 10,000 to 10,009, then use the pointer to the next page to retrieve 10,010					
Hash Index on sno, no overflow, bucket is 70% full, format 1	11 I/O. 11 for each of the data entries, assuming they are hashed to different buckets					

QI

- **1,000,000** employee records.
- 10 records per page
- sorted on sno
- sno is a candidate key, with values lying in the range 0-999,999
- 4K bytes per page
- eno and sno are each 4 bytes, and a
- pointer to a page is 4 bytes
- data are uniformly distributed in the domain.
- d) Find the employees whose sno is not 10.

Scan Sorted File	100,000 I/O
Clustered B+ tree on sno. Nodes are 100% full, format 2	100,000 I/O + additional 3 I/O to access B+ tree nodes
Hash Index on sno, no overflow, bucket is 70% full, format 1	142,858 I/O by scanning all



Consider a relation Book(Bookid: integer, Author: string, ...) where **Bookid** is the key, and its values are assigned in **increasing order** starting from 1. Suppose Book has 6 million records of 200 bytes each. Suppose we have to perform 10,000 single-record accesses, and 100 range queries of 0.005% of the file, with Bookid as the search key.

300 records per range query

Which of the above two methods is better for the application? Under what circumstance will the "loser" outperform the "winner"?

Hash Index

Load factor of 70%, bucket size of 4096 bytes, Format 1, no overflow buckets

Records per bucket = bucket size / record size * 0.7 = floor(4096 / 200) * 0.7 = 14

Number of buckets = total records / records per

bucket = 6000000 / 14 = 428,572 buckets

10,000 single-record accesses = 10,000 I/O **100** range queries = 100 * 300 = 30,000 I/O **Total** = 40,000 I/O

Size per page = **4096 bytes**Records per page = 20 * 0.7 = **14**Pages required for 300 records = **22**22 entries can span across **2-3 leaf nodes**

B+ Tree Index

Nodes are 70% full, node has a size of 4096 bytes, key is 8 bytes each, address is 4 bytes each

m = 2d

(2d + 1) * 4 + 2d * 8 <= 4096

d = 170 (round down)

m = **340** records per leaf node

If 70% full, 238 records per leaf node

 $\log_{238}(6000000) = 3$ levels

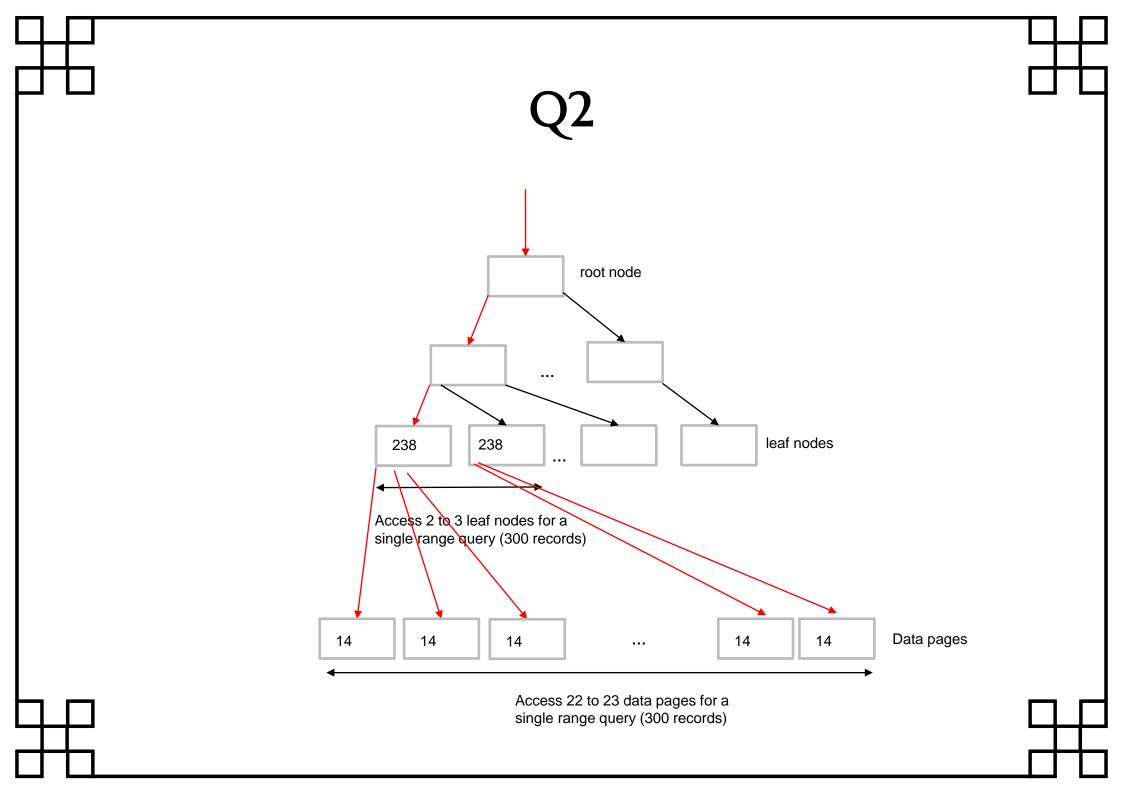
+1 to access data page

10,000 single-record accesses = 10,000 * (3+1) =

40,000 I/O

100 range queries = 100 * (3+1+22) = 2600 I/O

Total = 42,700 I/O



Discussion

Winner is hash index (based on our assumptions)

When will B+ tree index (loser) outperform hash index (winner)?

 number of range queries increases or the range increases, or the number of single-record queries reduces

When will hash index (winner) outperform B+ tree index (winner)?

• The other way round (more single record retrieval)



a) What can you say about the last entry that was inserted into the index if you know that there have been no deletions from this index so far?

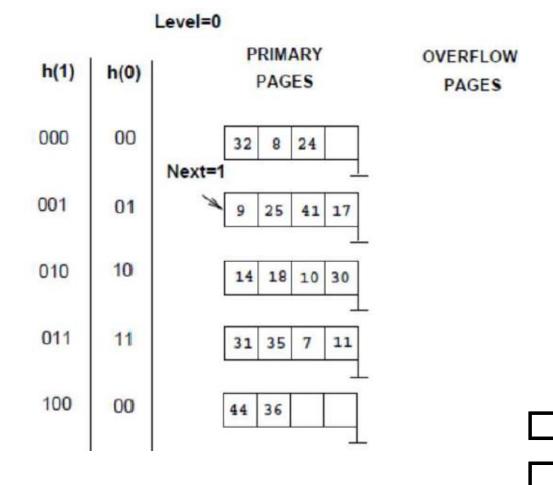
If the last item that was inserted had a hashcode h0(keyvalue) = 00 then it caused a split (as there are no overflow pages for any of the other buckets), otherwise, any value could have been inserted.

		Level=0	-				
h(1)	h(0)		PRIMARY				OVERFLOW PAGES
000	00		32	8	24		
		Next=1		69 0		1	
001	01	A	9	25	41	17	
010	10		14	18	10	30	
				20			
011	11		31	35	7	11	
100	00		44	36			



b) Show the index after inserting an entry with hash value 4, and 15.

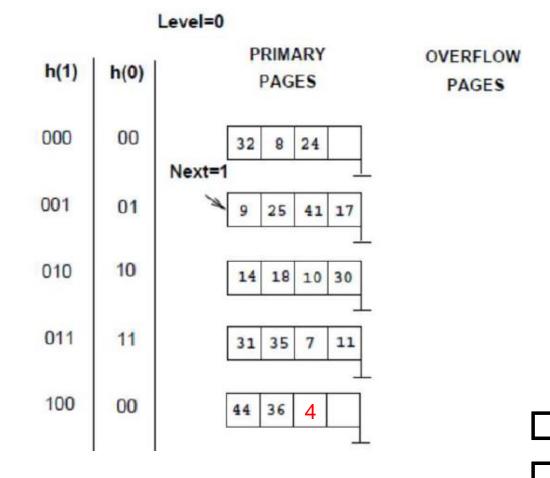
Binary representation 4-100 15-1111





b) Show the index after inserting an entry with hash value 4, and 15.


Binary representation **4 – 100** 15 – 1111





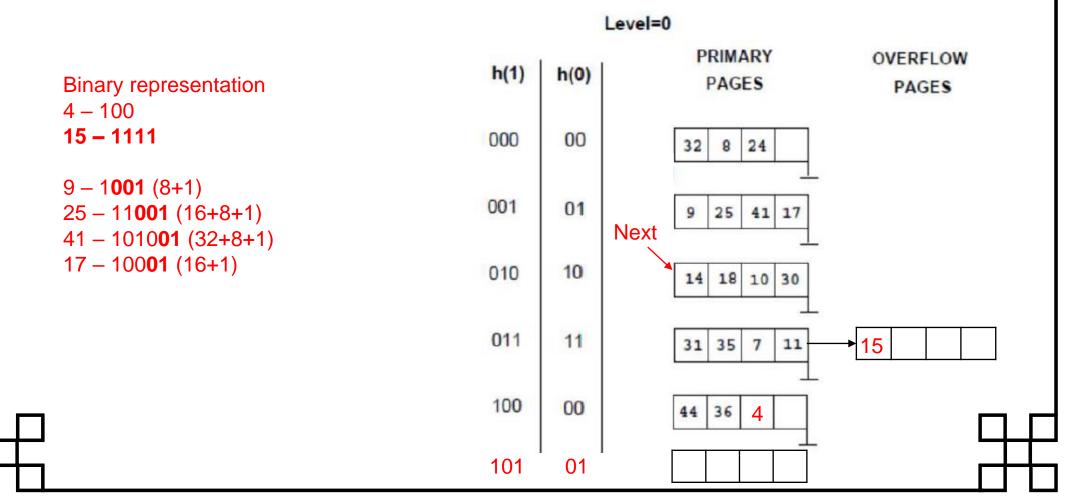
b) Show the index after inserting an entry with hash value 4, and 15.

Binary representation 4 – 100 **15 – 1111**





b) Show the index after inserting an entry with hash value 4, and 15.



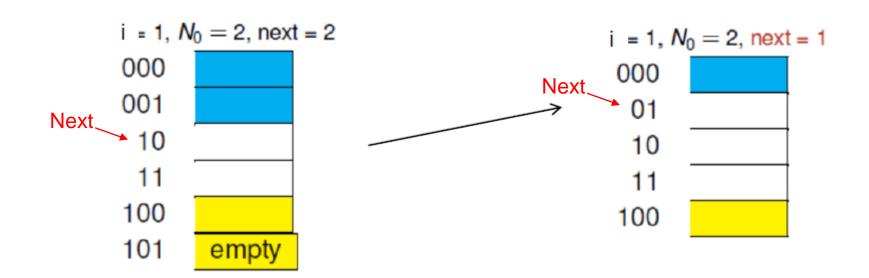
Consider the Linear Hashing index shown in the figure below. Assume that we split whenever an overflow page is created. Answer the following questions about this index:

c) Show the index after deleting the entries with hash values 36 and 44 into the original tree. Assume that the full deletion algorithm is used.

Level=0 PRIMARY **OVERFLOW** h(1) h(0) Delete 36 and 44 from the last bucket PAGES PAGES 00 000 32 24 Next=1 001 01 25 41 10 010 18 10 011 11 35 100 00 44 36

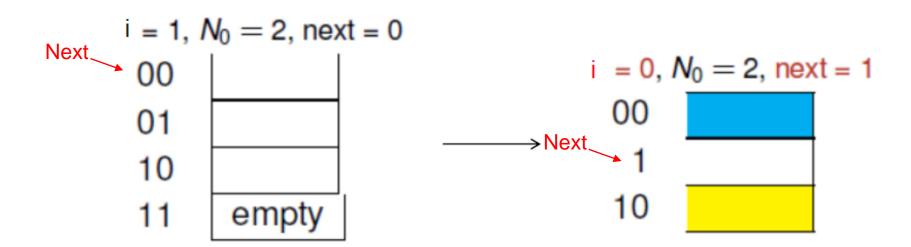
Linear Hashing: Deletion

- Locate bucket and delete entry
- If the last bucket B_{Ni + next 1} becomes empty, it can be removed
- Case 1: If next > 0
 - Decrement next



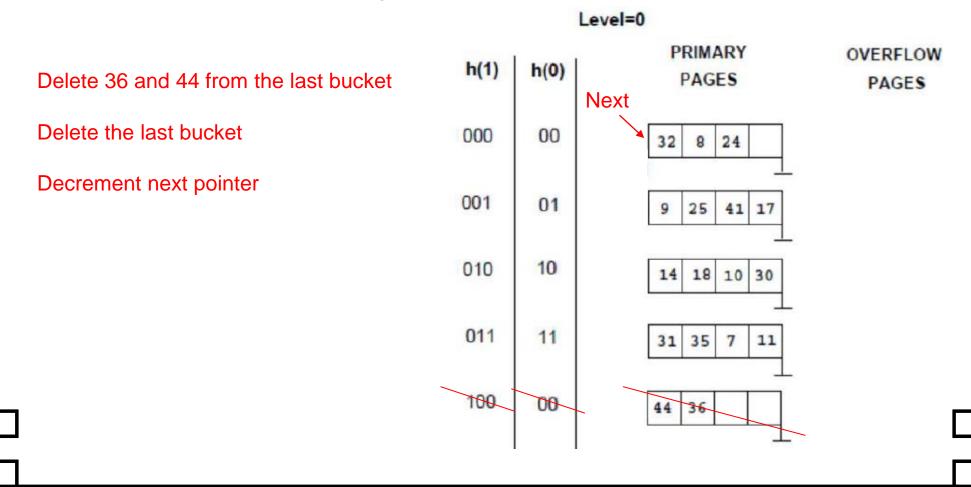
Linear Hashing: Deletion

- Case 2: If (next = 0) and (i > 0)
 - Update next to point to the last bucket in previous level B_{ni/2-1}
 - Decrement level





c) Show the index after deleting the entries with hash values 36 and 44 into the original tree. Assume that the full deletion algorithm is used.



Consider the Linear Hashing index shown in the figure below. Assume that we split whenever an overflow page is created. Answer the following questions about this index:

d) Find a list of entries whose insertion into the original index would lead to a bucket with two overflow pages. Use as few entries as possible to accomplish this.

h(1)	h(0)		PRIMARY PAGES				OVERFLOW PAGES
000	00		32	8	24		
		Next=1		100			
001	01	A	9	25	41	17	
010	10		14	18	10	30	
011	11		31	35	7	11	
100	00		44	36			
						一	

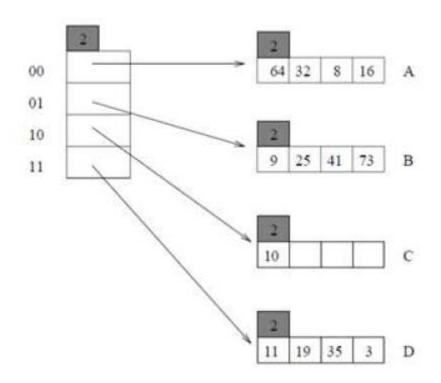
I accel-0

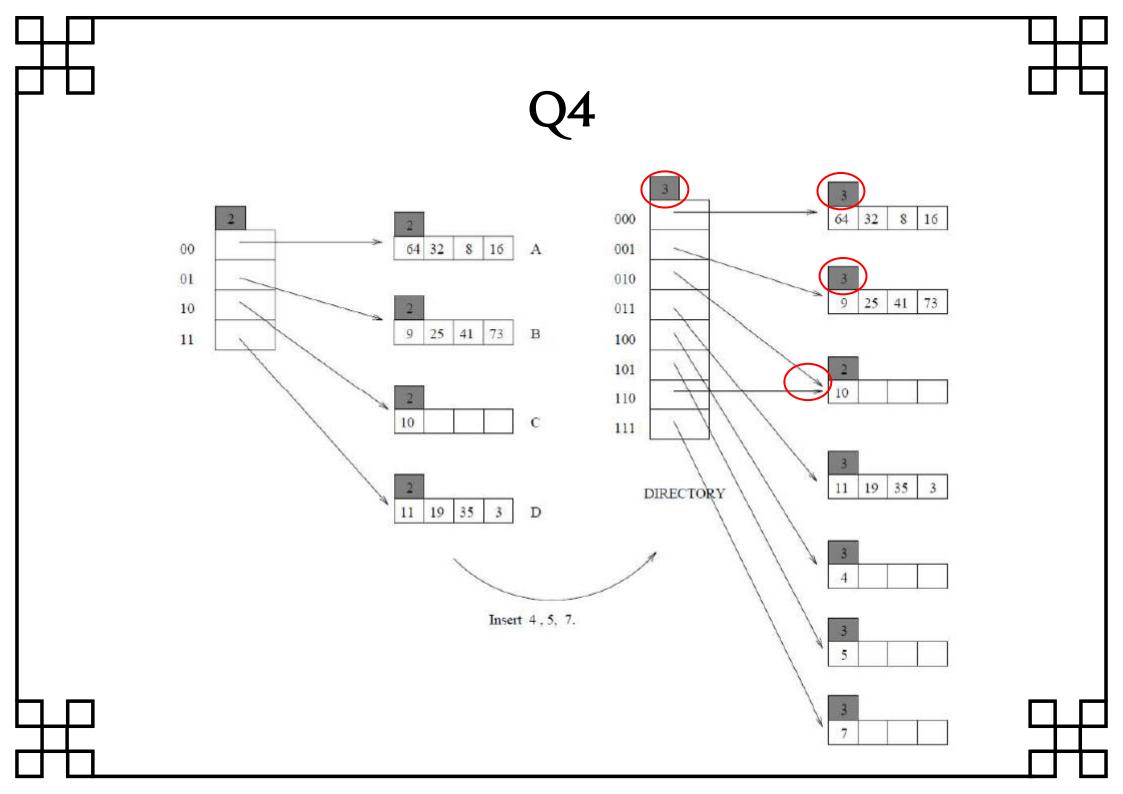
Consider the Linear Hashing index shown in the figure below. Assume that we split whenever an overflow page is created. Answer the following questions about this index:

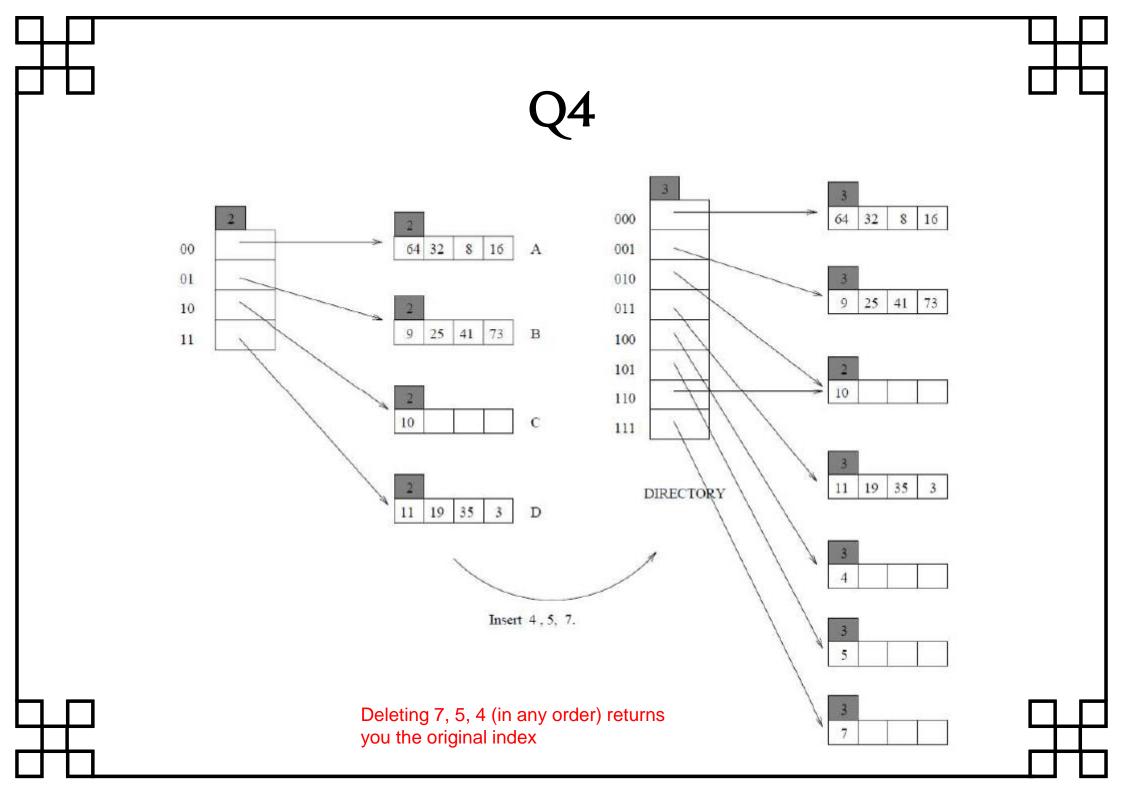
d) Find a list of entries whose insertion into the original index would lead to a bucket with two overflow pages. Use as few entries as possible to accomplish this.

	Level=0							
All ending with 111 63, 127, 255, 511, 1023 The first insertion causes a split and causes an update of Next to 2. The	h(1)	h(0)	PRIMARY PAGES	OVERFLOW PAGES				
insertion of 1023 causes a subsequent split and Next is updated to 3 which points to this bucket. This overflow	000	00	32 8 24 Next=1					
chain will not be redistributed until three more insertions (a total of 8 entries) are made.	001	01	9 25 41 17					
OR Keep inserting -001 entries will cause	010	10	14 18 10 30					
bucker 01 to overflow into 2 pages OR	011	11	31 35 7 11					
any other valid answers	100	00	44 36	[
Keep inserting -001 entries will cause bucker 01 to overflow into 2 pages OR								

Consider the extendible hash table below. Insert 4, 5 and 7. Show the resultant hash table? Now, delete 7, 5, and 4. What is the resultant hash table?







O3 % Extra Qs

Extra

1. Is Linear Hashing order independent?

Ans:

- An index is order dependent; inserting the same set of keys in a different ordering result in different structure.
- Here, we look at the structure of the hash index. It is order dependent. Counter example: 4 buckets initially. Insert a1-a8 into the buckets. Case 1: insert a1-a3 into bucket 1; then insert a4-a8 to bucket 4. Case 2: insert a4-a8 to bucket 4 first; then insert a1-a3. We will have two different structures (due to different occurrence of overflow pages and bucket splitting)

Extra 2. Suppose we use a linear hashing scheme, but there are pointers to records from outside. These pointers prevent us from moving records between blocks (which is what LH does redistribute records whenever there is a split). Suggest ways that we could modify the structure to allow pointers from outside?

