

The background of the image consists of numerous large, 3D-rendered blue numbers of various sizes and orientations, creating a dense, textured pattern.

CS3223 Tutorial 5

Gary Lim

Welcome Back!

- ❖ **Midterms** is on the **4th of March** (Friday) **10 to 11:30am**
 - ❖ Venue: MPSH2
 - ❖ See course webpage for Seat Allocation
- ❖ If tested covid positive
 - ❖ Online arrangement: Email prof to attempt the **same paper** online at the **same time** on luminus
 - ❖ Make-up test: If unable to sit for the paper (online or onsite), please email Prof to have a separate makeup test (a different paper) arranged for you

Chapter Review

- ❖ Access Paths
- ❖ Duplicate Removal
 - ❖ Sort-based
 - ❖ Hash-based

Access Paths

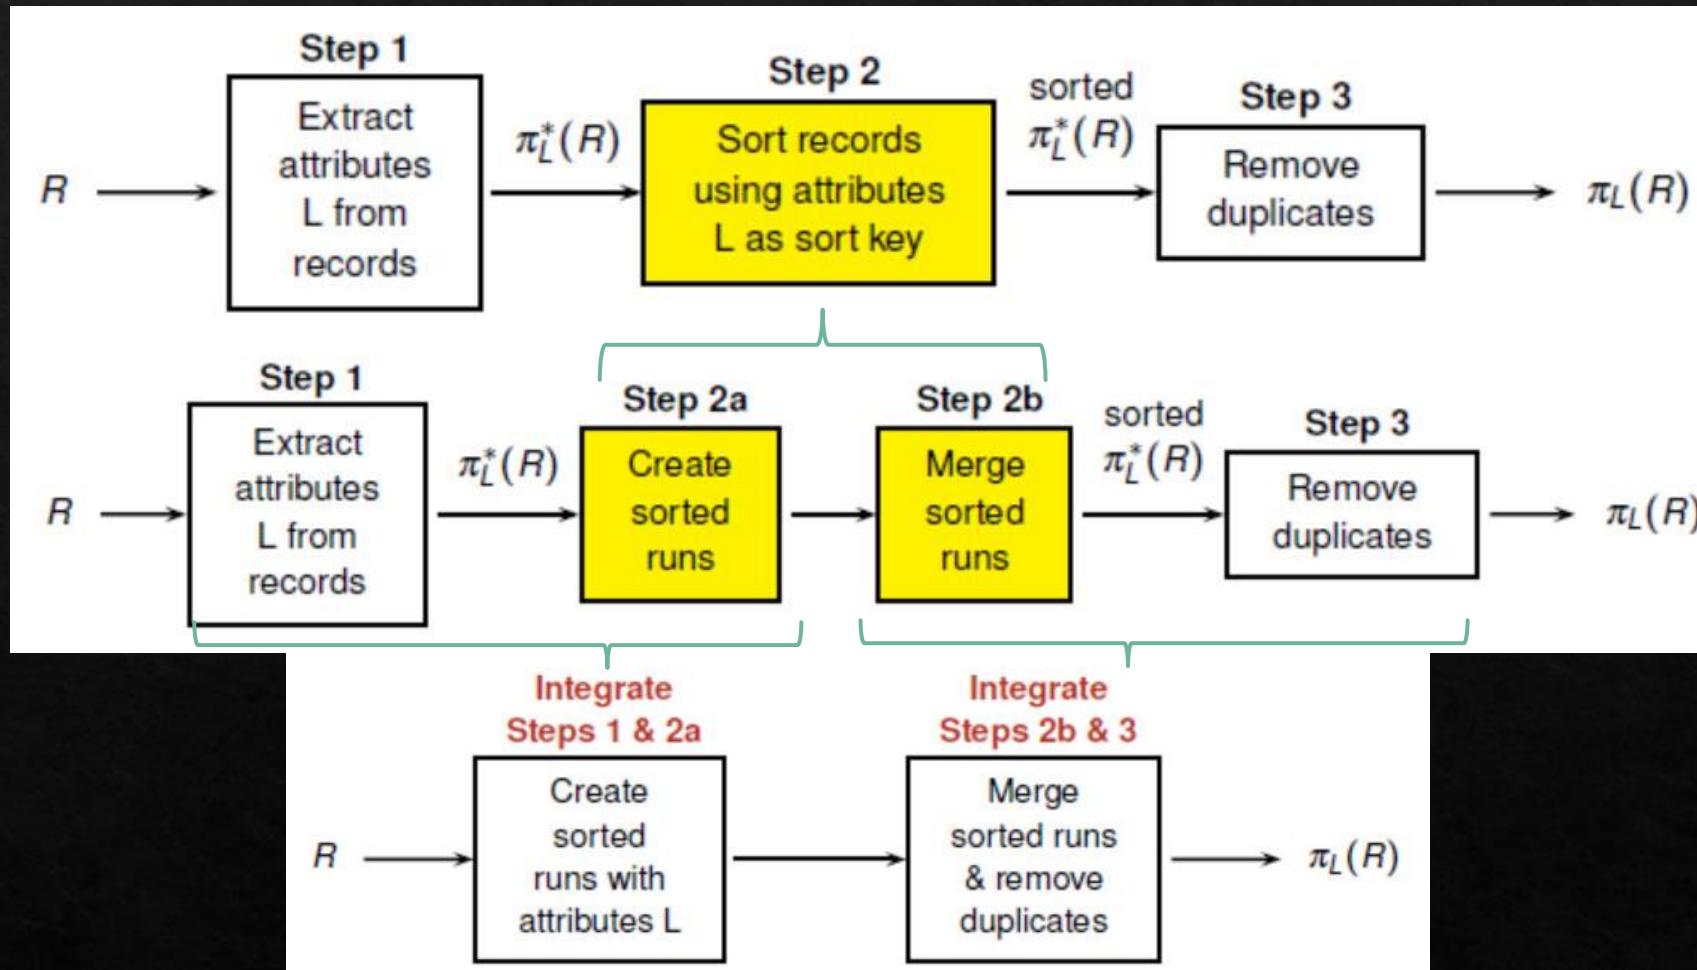
- ❖ Table Scan – Scan all pages
 - ❖ Projection
 - ❖ Selection (no index)
- ❖ Index-Only Scan
 - ❖ Scan the index and its information (e.g. format 1)
- ❖ Index Scan
 - ❖ Scan the index and retrieve data pages

Note that cost of reading and writing are usually different as the result would usually be different size (after a projection, selection or a join)

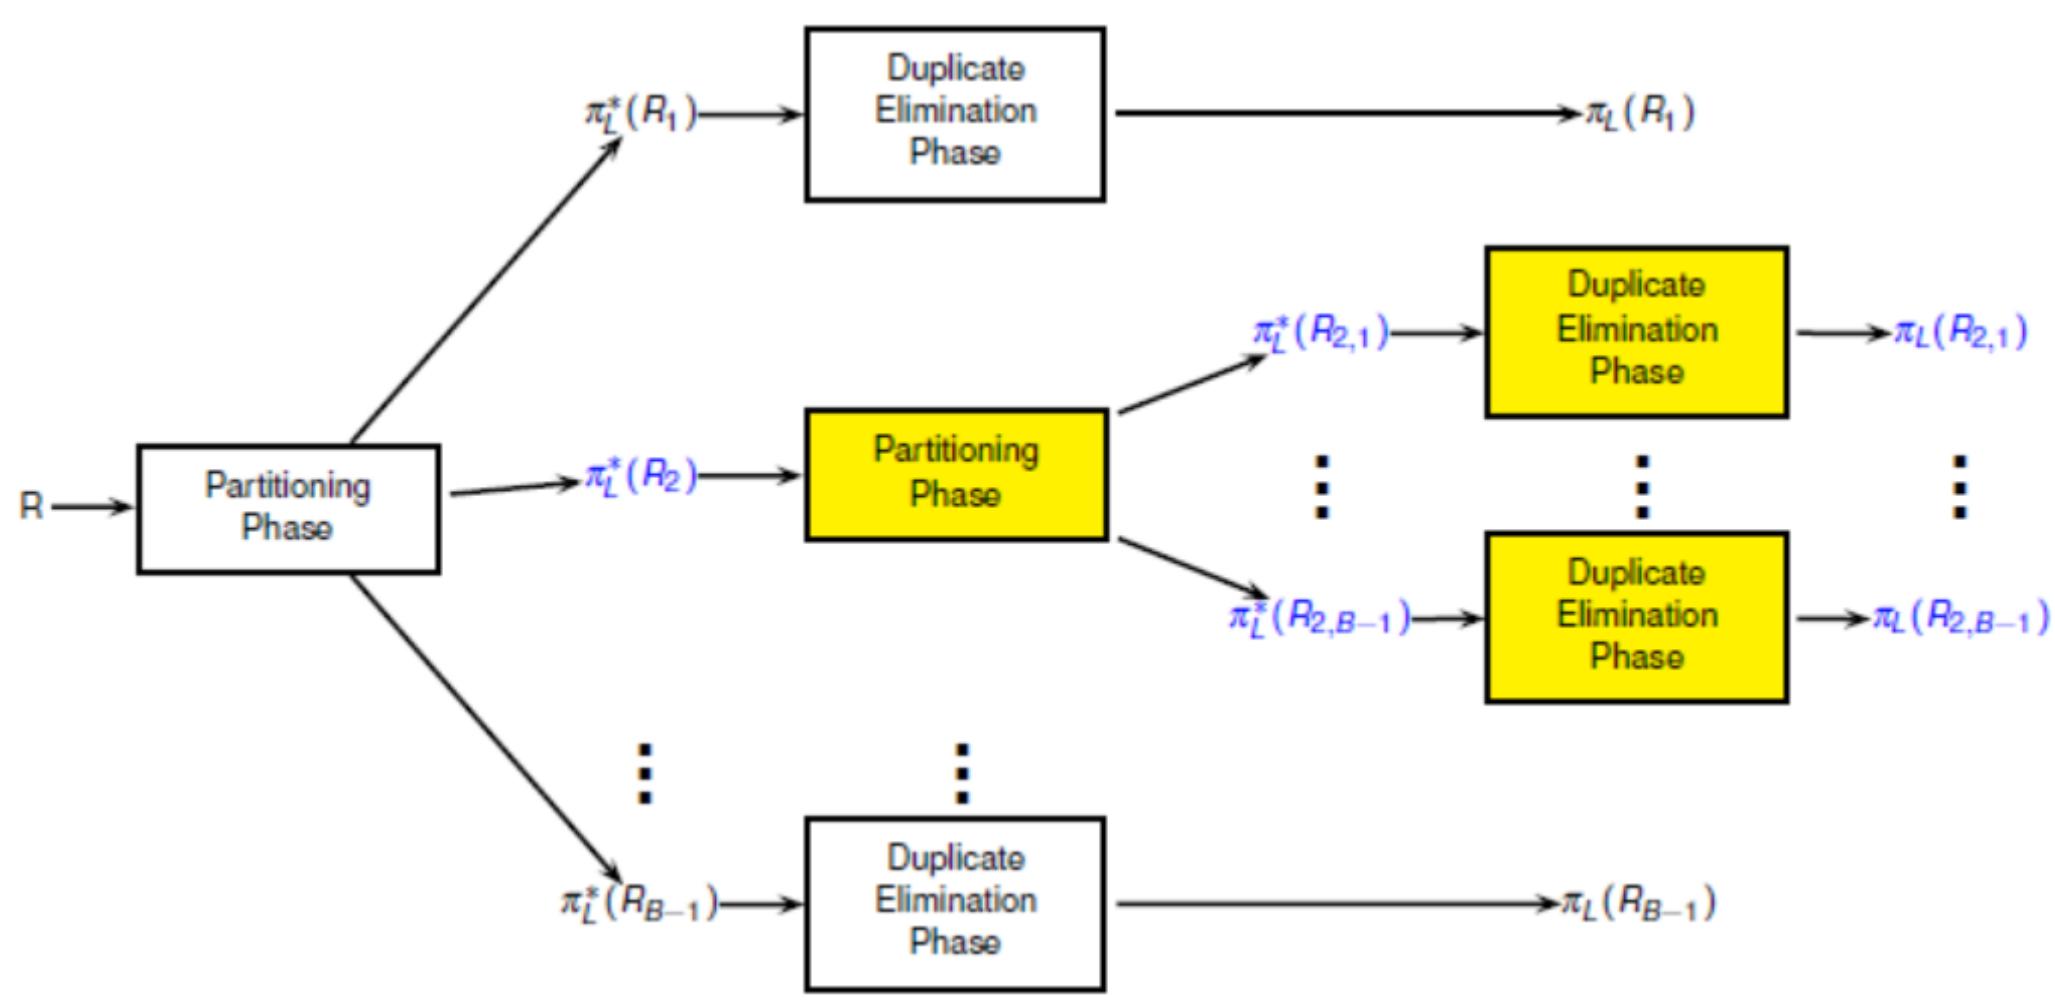
Projection (with duplicate removal)

- ❖ Sort-based approach
- ❖ Hash-based approach

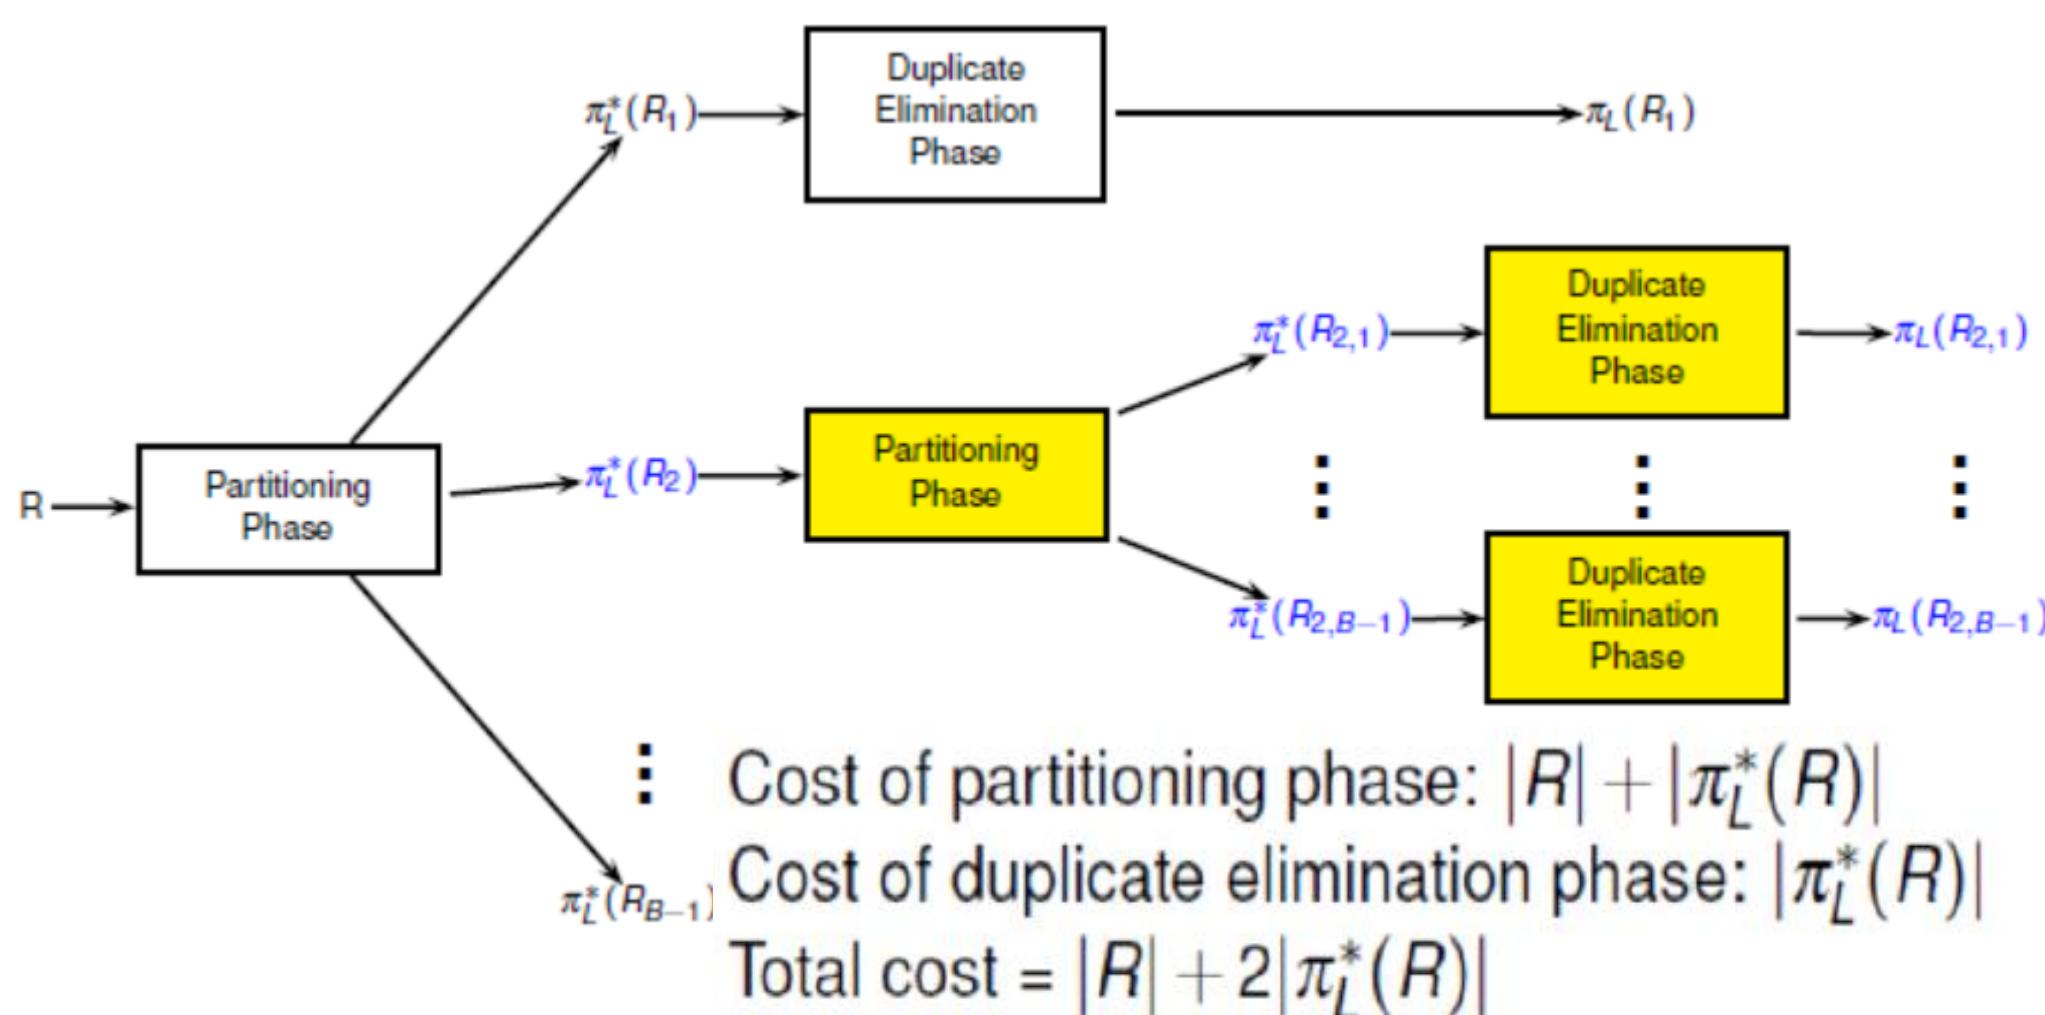
Optimized Sort-Based Approach



Hash-Based Approach



Hash-Based Approach



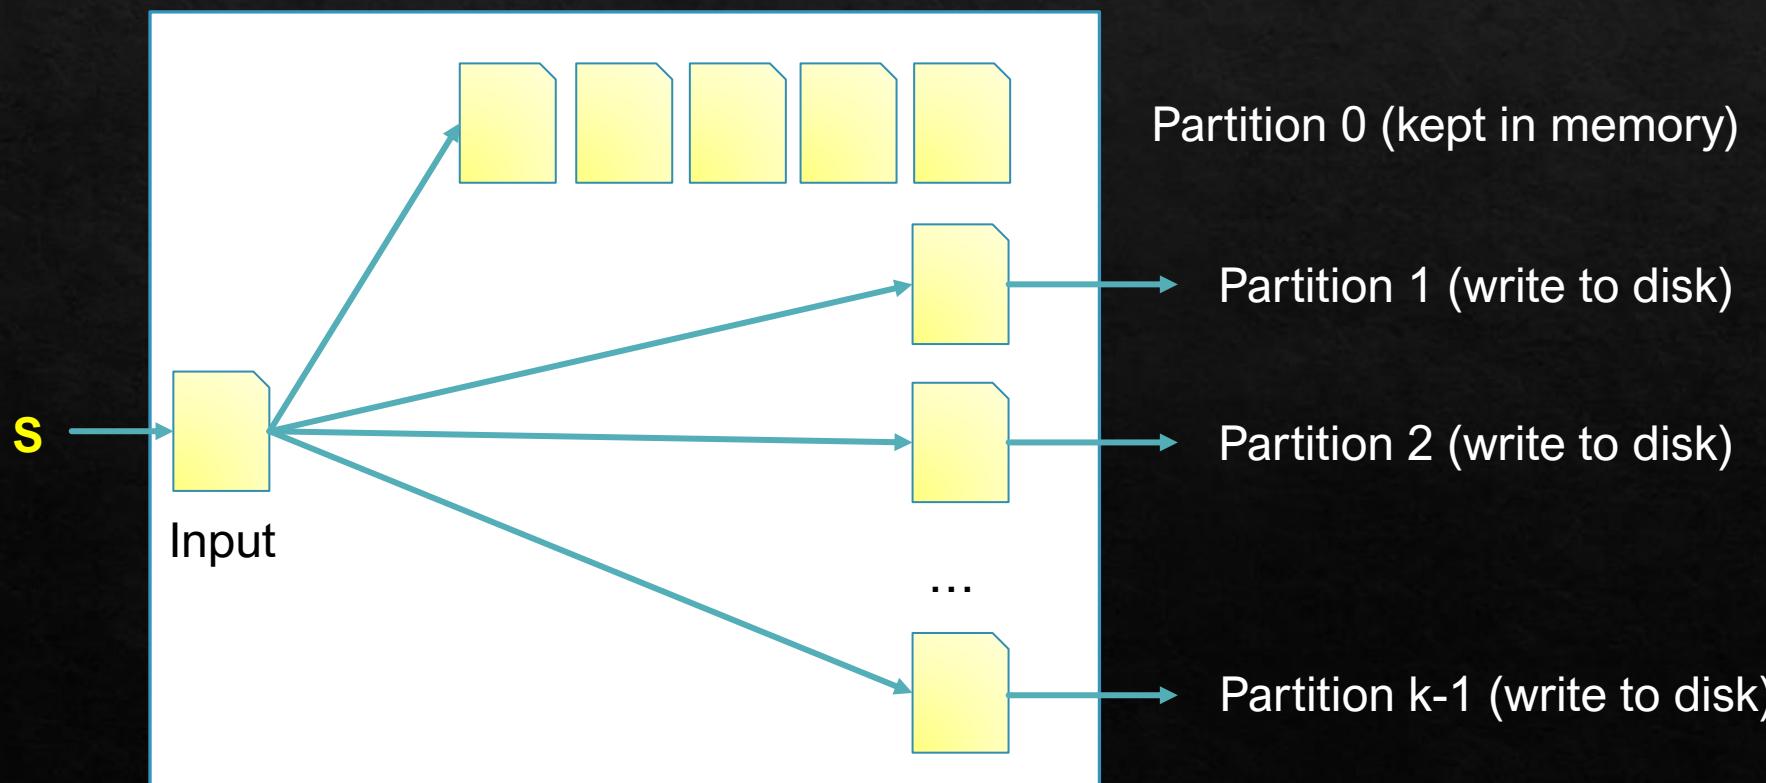
Tutorial Questions

Q1

- ❖ Hybrid Hash Join
 - ❖ k buckets
 - ❖ Keep bucket 0 of relation **S** in memory and join it on the fly as bucket 0 elements of relation **R** are being generated
- ❖ What is the cost of the algorithm?

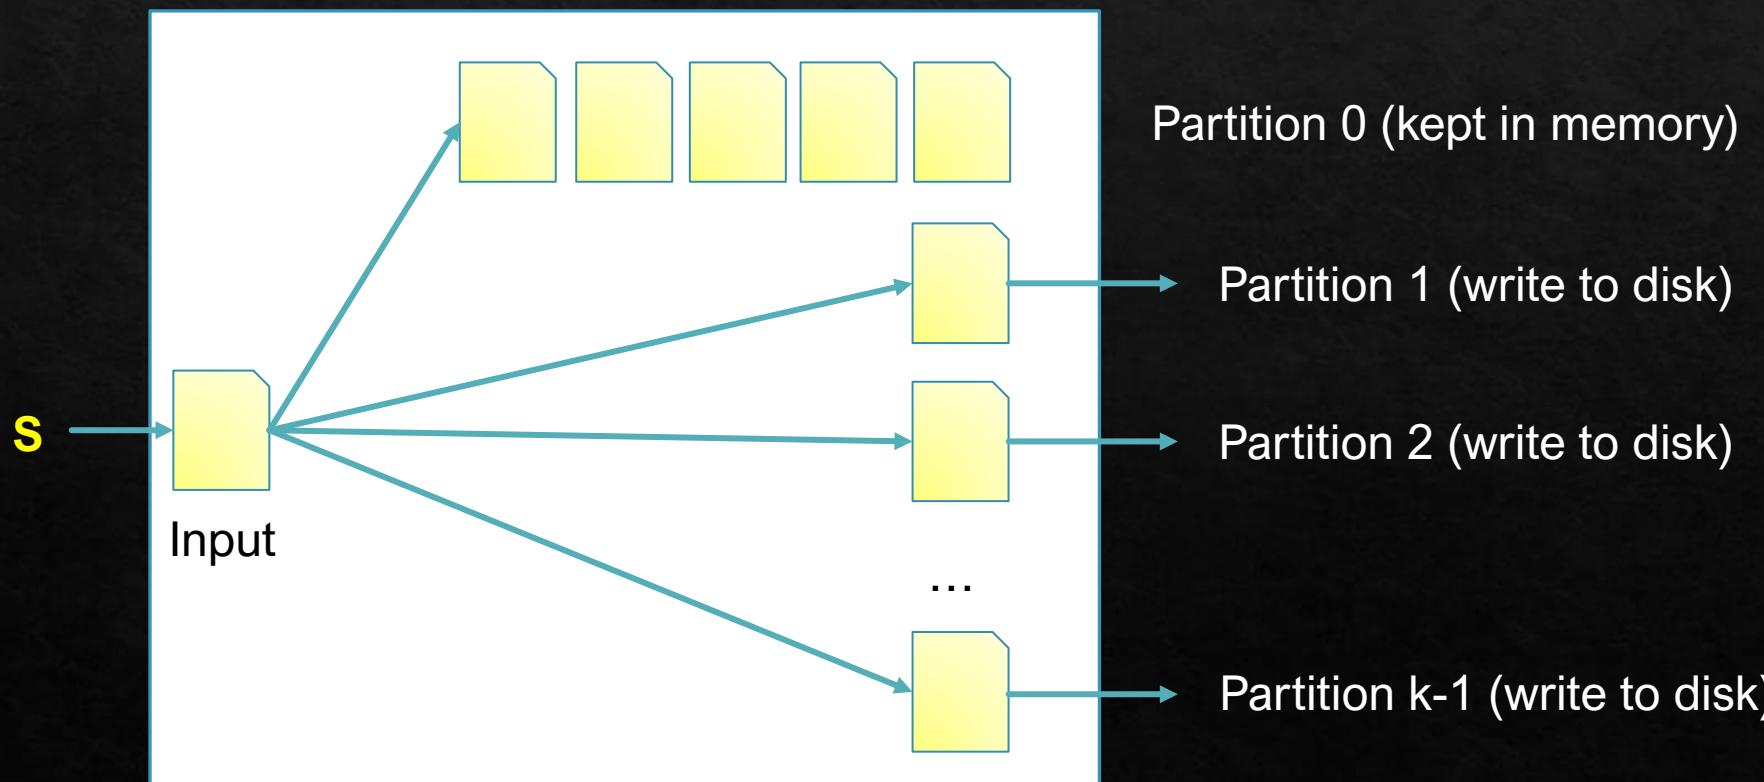
Q1

❖ Partition Phase: **Partition S**



Q1

❖ Partition Phase: Partition **S**

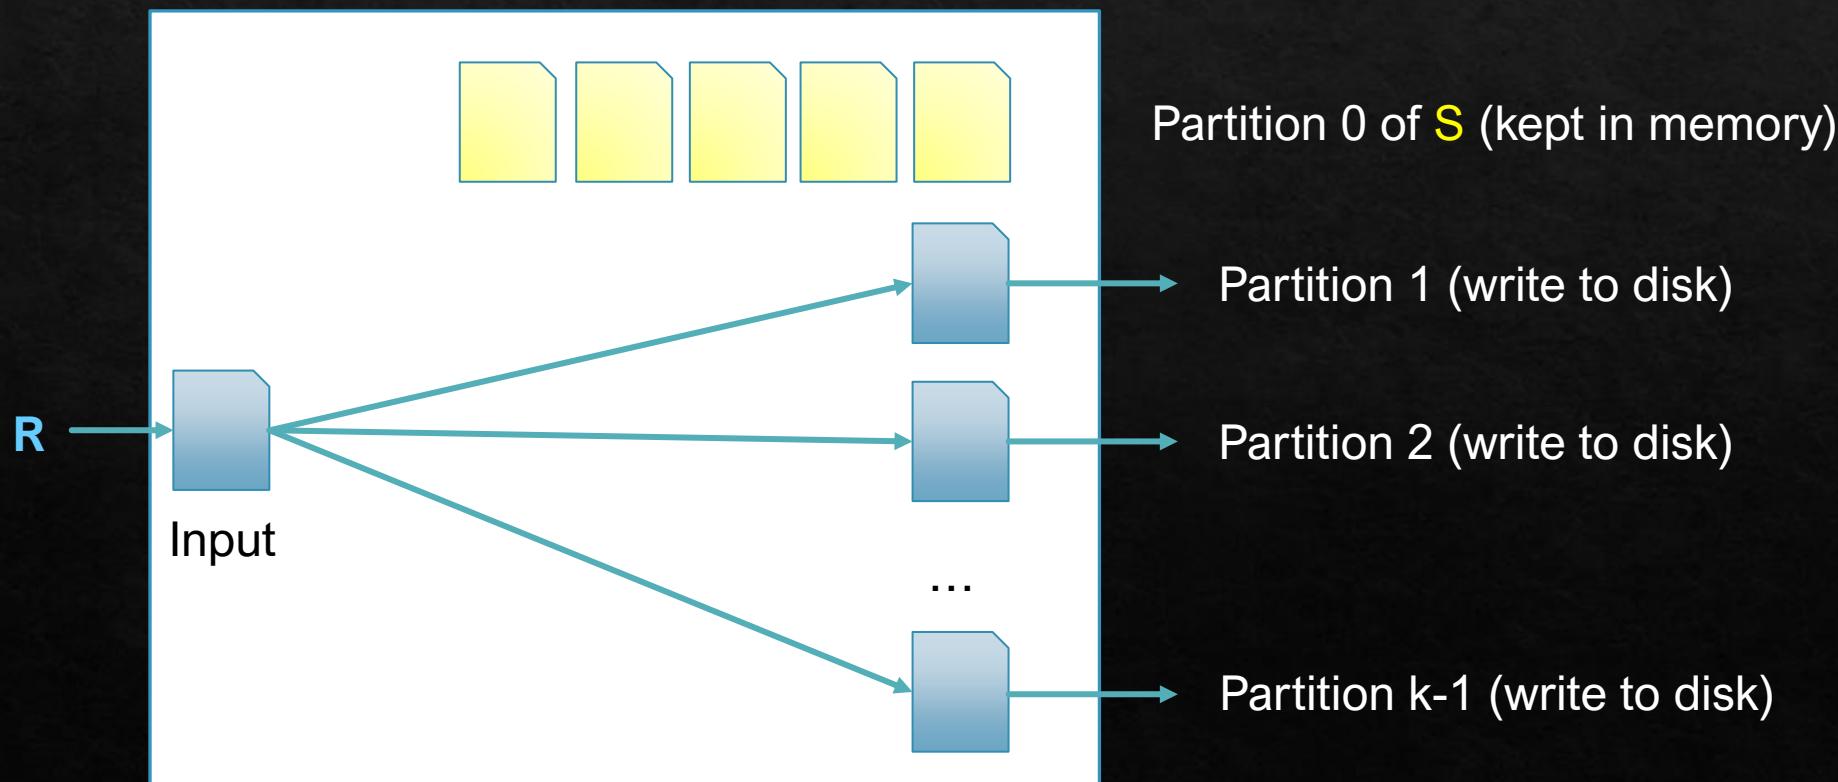


Size of Partition 0:
 $|S| / k$

Cost:
Read + Write
 $= |S| + |S| - |S|/k$
 $= (2-1/k)|S|$

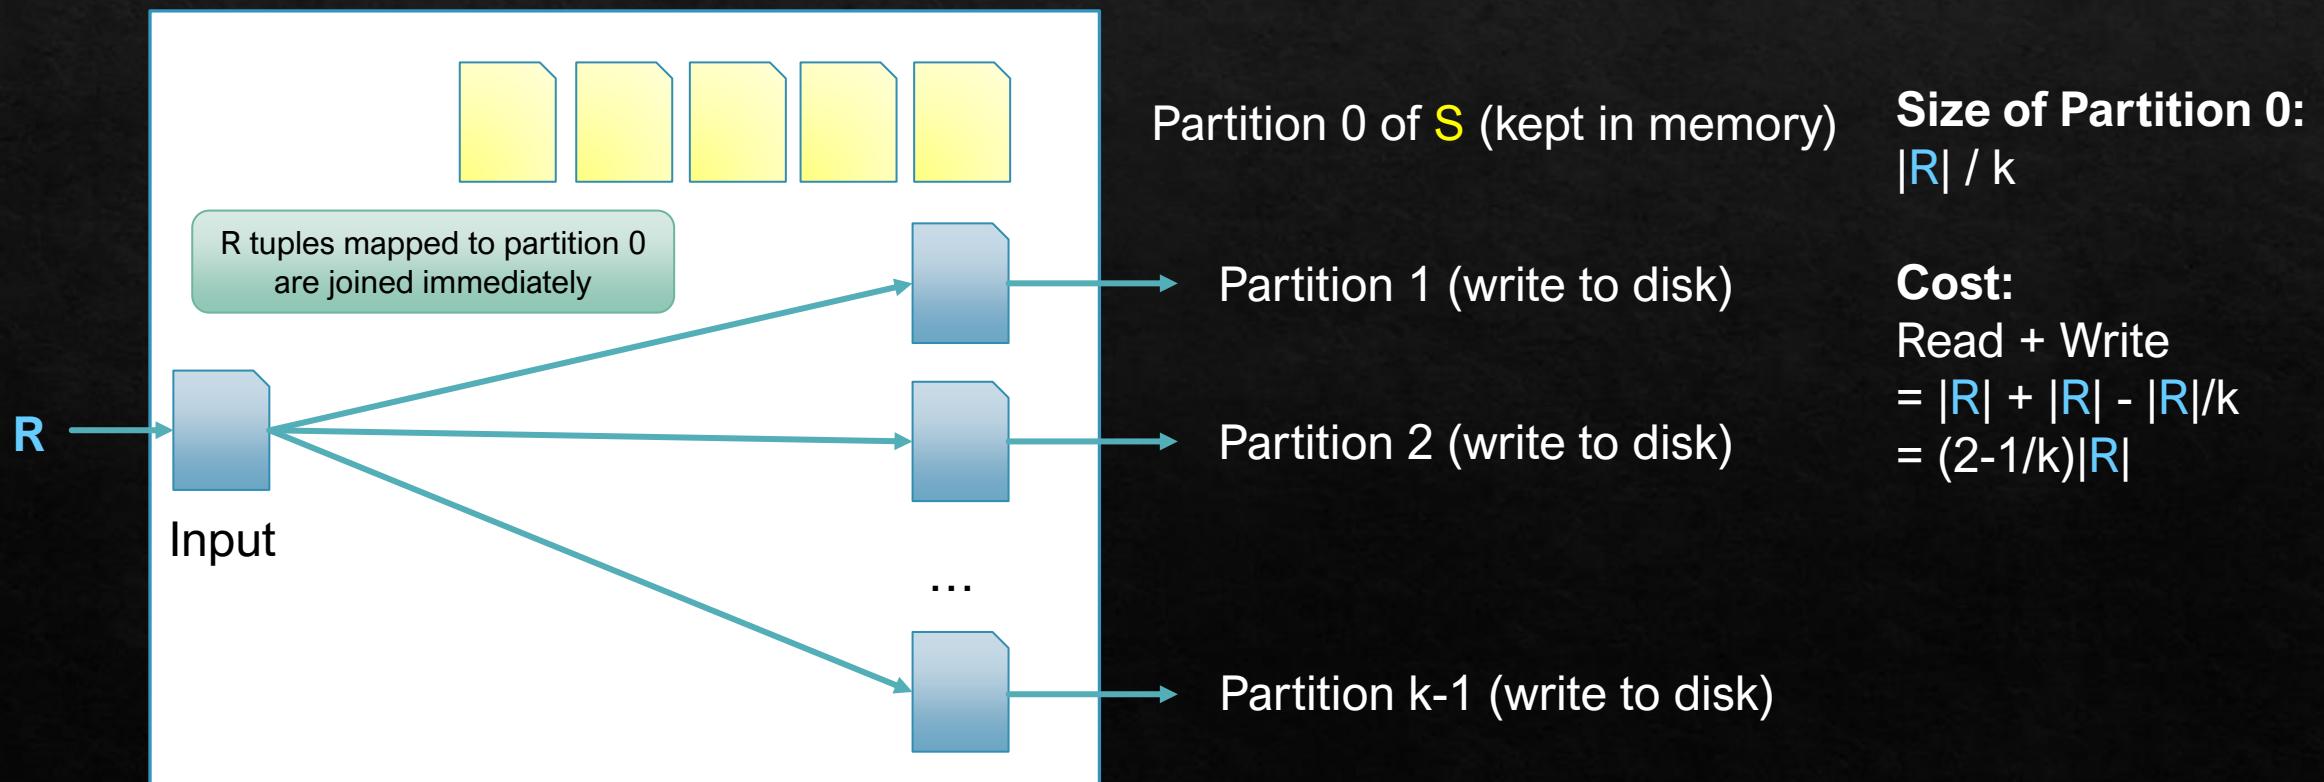
Q1

❖ Partition Phase: **Partition R**



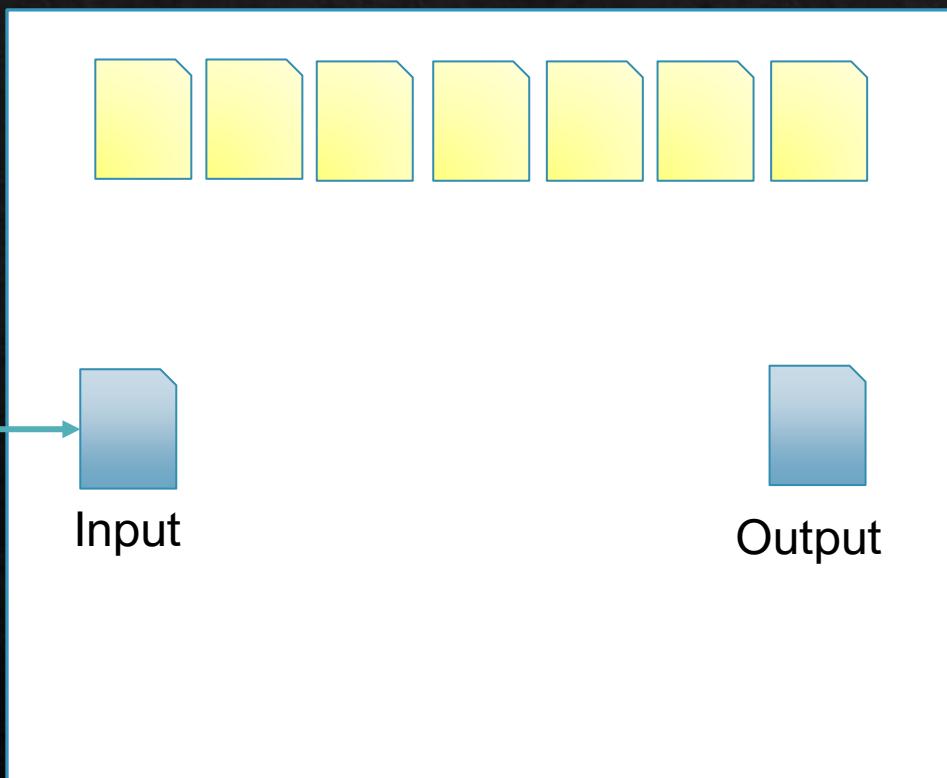
Q1

❖ Partition Phase: Partition R



Q1

- Join Phase: Join bucket 1 to $(k-1)$ or **S** and **R**



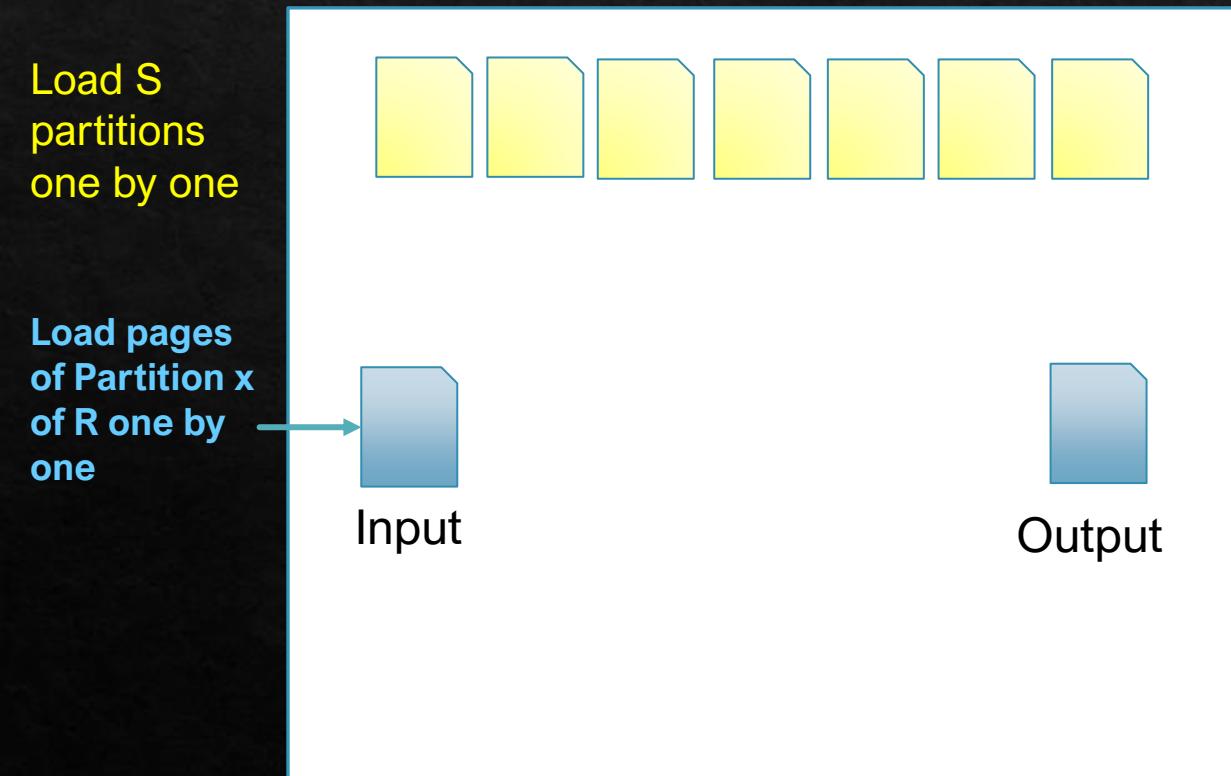
Same as original

Partition x of **S** (kept in memory)
where $x = \{1, 2, \dots, k-1\}$

Cost:
Read + Write
 $= |S| - |S|/k + |R| - |R|/k$
 $= (1-1/k)(|S| + |R|)$

Q1

- ◇ Join Phase: Join bucket 1 to $(k-1)$ or **S** and **R**



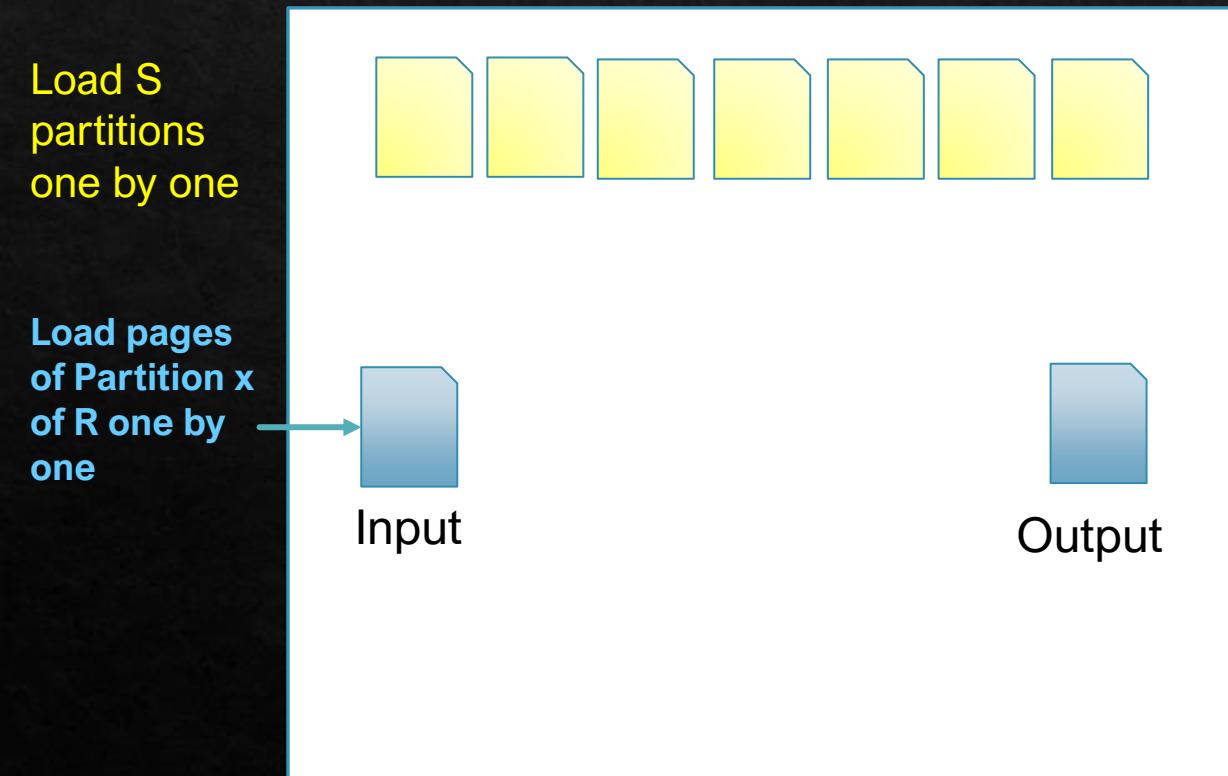
Partition x of **S** (kept in memory)
where $x = \{1, 2, \dots, k-1\}$

Total Cost:

$$\begin{aligned} &= (2-1/k)|S| + (2-1/k)|R| + (1-1/k)(|S|+|R|) \\ &= (3-2/k)(|S|+|R|) \end{aligned}$$

Q1

❖ Join Phase: Join bucket 1 to $(k-1)$ or **S** and **R**



Partition x of **S** (kept in memory)
where $x = \{1, 2, \dots, k-1\}$

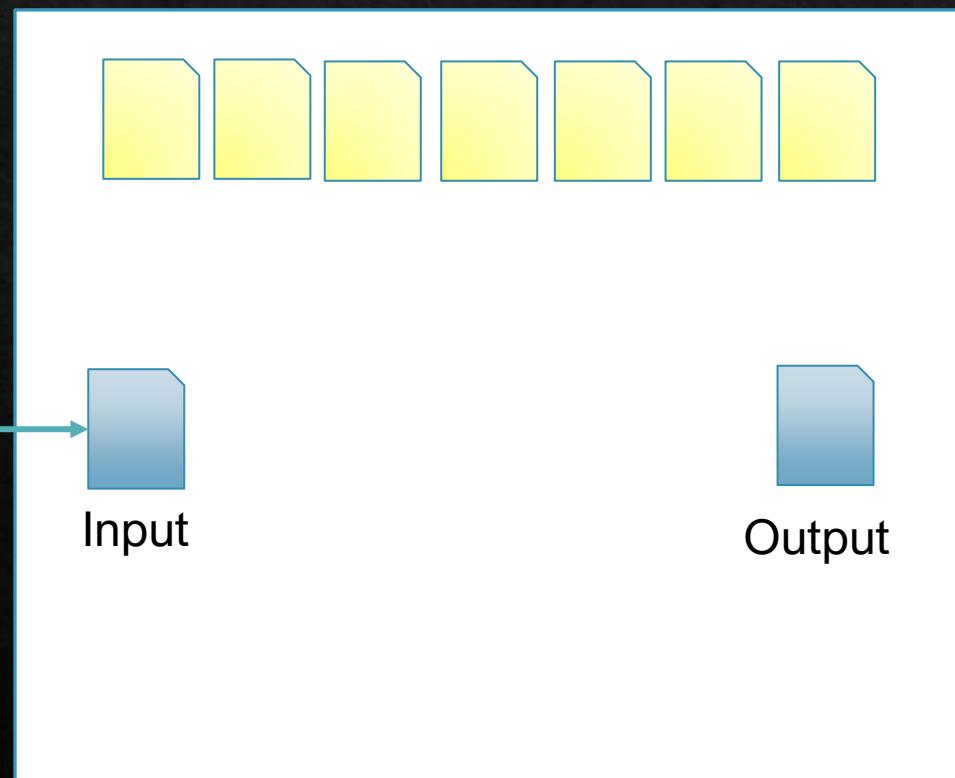
Total Cost:

$$\begin{aligned} &= (2-1/k)|S| + (2-1/k)|R| + (1-1/k)(|S|+|R|) \\ &= (3-2/k)(|S|+|R|) < 3(|S|+|R|) \end{aligned}$$

- Always better than original GRACE Hash join
- As k gets smaller, cost savings increase!
- Able to return initial answers faster (faster initial response time)

Q1

❖ Join Phase: Join bucket 1 to (k-1) or **S** and **R**



Partition x of **S** (kept in memory)
where $x = \{1, 2, \dots, k-1\}$

Total Cost:

$$\begin{aligned} &= (2-1/k)|S| + (2-1/k)|R| + (1-1/k)(|S|+|R|) \\ &= (3-2/k)(|S|+|R|) < 3(|S|+|R|) \end{aligned}$$

- Always better than original GRACE Hash join
- As k gets smaller, cost savings increase!
- Able to return initial answers faster (faster initial response time)

- Assumes uniformly distributed buckets (which is not always the case)
- For partition phase of **R**, we need 1 input buffer, $k-1$ output buffer for partitions of **R**, and 1 output buffer for join results of partition 0. Hence in memory partition must not be larger than **B-2-(k-1)**

Q2

- ❖ Advantages of join index (rid, sid)
- ❖ Disadvantages

Q2

- ❖ Advantages of join index (rid, sid)
 - ❖ No need to perform join processing, simply scan index and retrieve relevant data pages
 - ❖ Initial response time is good
 - ❖ If it is B-tree, results are sorted by join columns
- ❖ Disadvantages
 - ❖ Updates are costly to maintain. E.g. every insert into relation R require join tuples of S to be computed and added to index. Same applies to deletion
 - ❖ Can it perform worse than standard join processing?

Q2

- ❖ Advantages of join index (rid, sid)
 - ❖ No need to perform join processing, simply scan index and retrieve relevant data pages
 - ❖ Initial response time is good
 - ❖ If it is B-tree, results are sorted by join columns
- ❖ Disadvantages
 - ❖ Updates are costly to maintain. E.g. every insert into relation R require join tuples of S to be computed and added to index. Same applies to deletion
 - ❖ Can it perform worse than standard join processing? Yes, if the search result is extremely huge wrt to the original relations, probing the index might cost significant overhead

Q3a

- ❖ applicant(**pid**, cityid, income, gmat)
- ❖ location(**cityid**, country)
- ❖ Find applicants that earn a salary greater than 60,000 and have GMAT scores higher than the average score of applicants from USA

Q3a

- ❖ applicant(**pid**, cityid, income, gmat)
- ❖ location(**cityid**, country)
- ❖ Find applicants that earn a salary greater than 60,000 and have GMAT scores higher than the average score of applicants from USA
- ❖

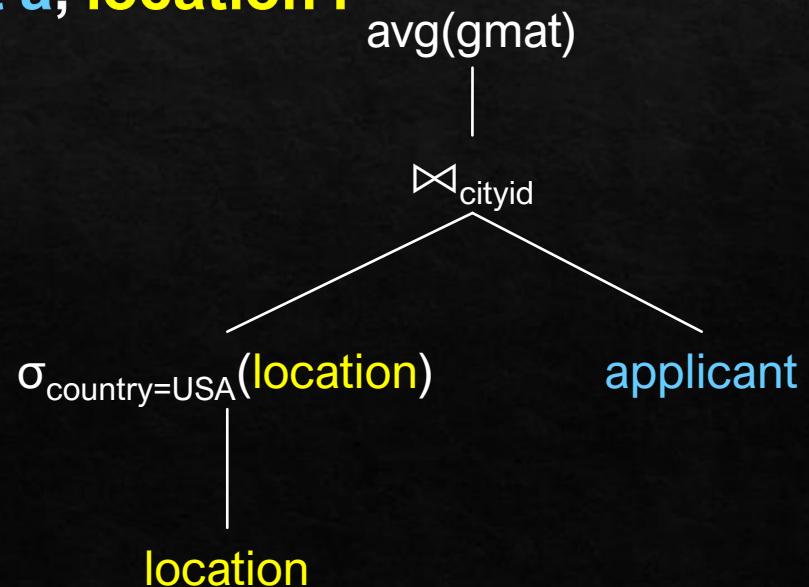
```
SELECT pid FROM applicant, location  
WHERE income > 60000  
AND gmat >  
(SELECT AVG(gmat) FROM applicant a, location l  
WHERE a.cityid = l.cityid  
AND l.country = "USA")
```

Q3a

```
◊ SELECT pid FROM applicant  
WHERE income > 60000  
AND gmat >  
(SELECT AVG(a.gmat) FROM applicant a, location l  
WHERE a.cityid = l.cityid  
AND l.country = "USA")
```

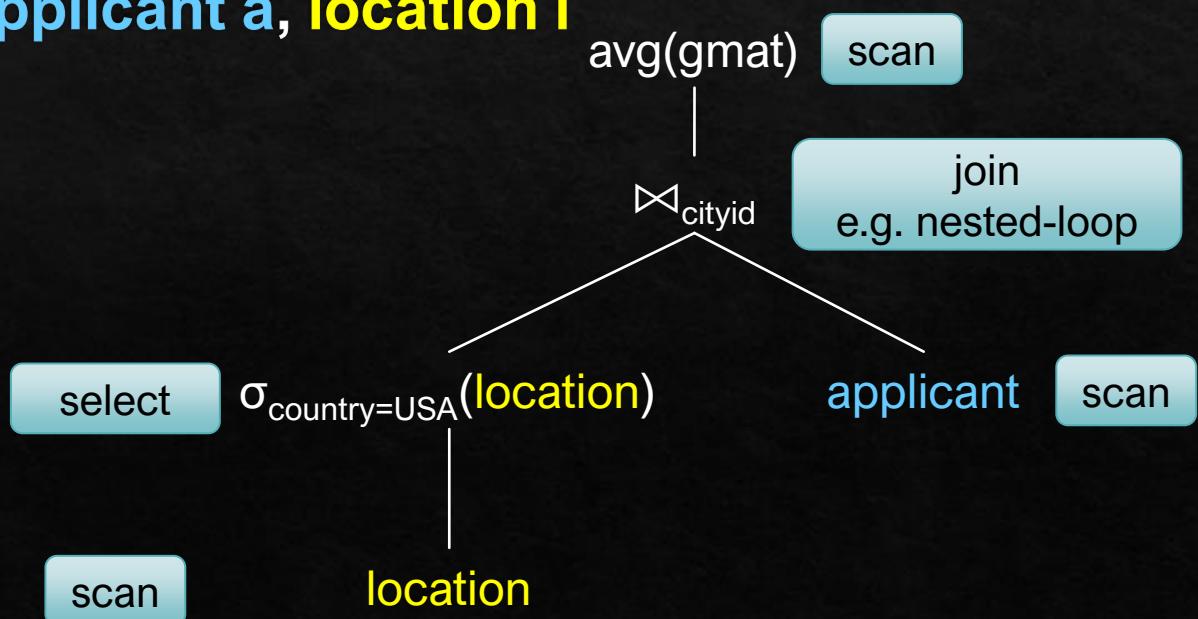
Q3a

❖ SELECT pid FROM applicant
WHERE income > 60000
AND gmat >
(SELECT AVG(a.gmat) FROM applicant a, location l
WHERE a.cityid = l.cityid
AND l.country = “USA”)



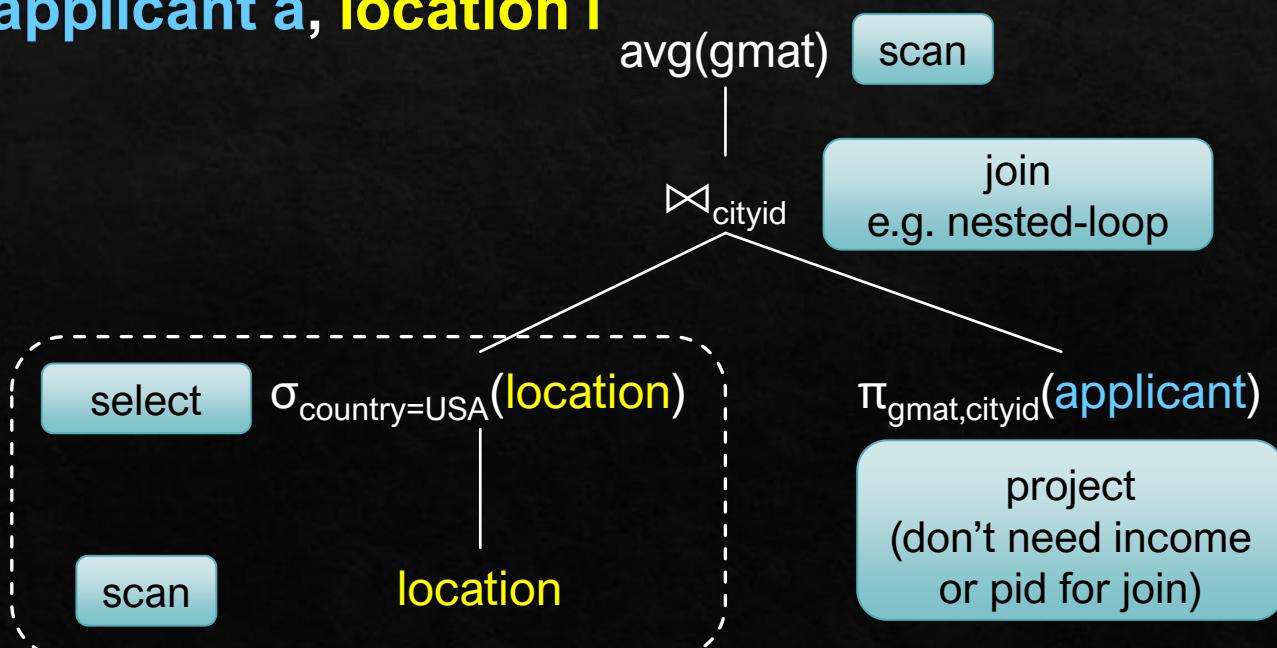
Q3

◊ SELECT pid FROM applicant
WHERE income > 60000
AND gmat >
(SELECT AVG(a.gmat) FROM applicant a, location l
WHERE a.cityid = l.cityid
AND l.country = “USA”)



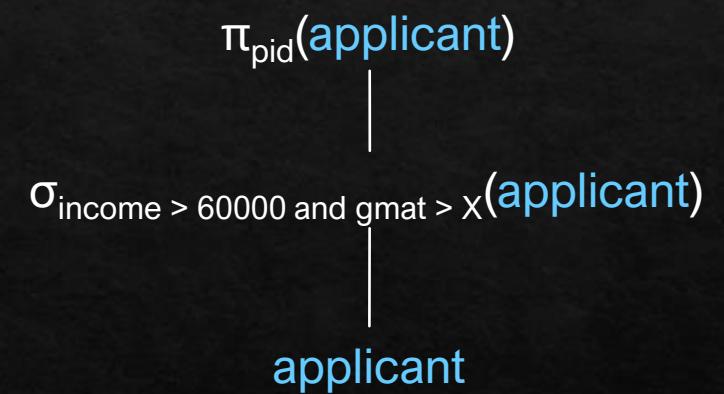
Q3a

◊ SELECT pid FROM applicant
WHERE income > 60000
AND gmat >
**(SELECT AVG(a.gmat) FROM applicant a, location l
WHERE a.cityid = l.cityid
AND l.country = “USA”)**



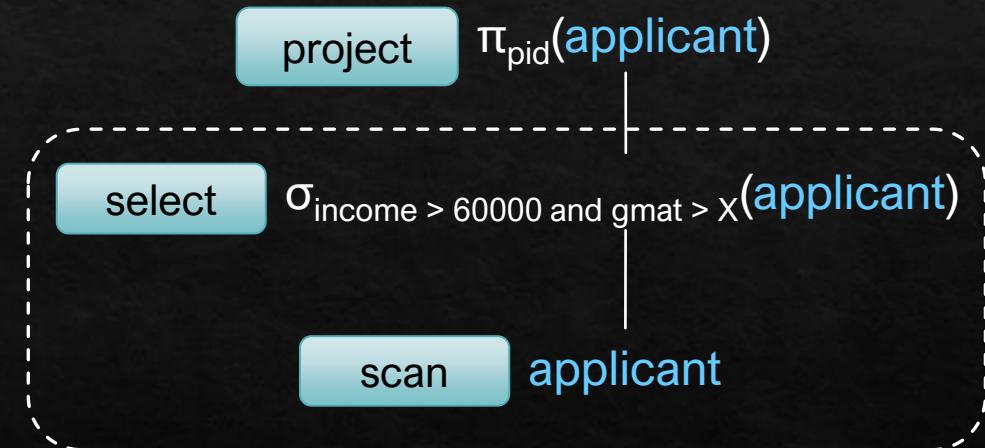
Q3a

◊ **SELECT pid FROM applicant,
WHERE income > 60000
AND gmat > X**



Q3a

◊ **SELECT pid FROM applicant,
WHERE income > 60000
AND gmat > X**



Q3b

- ❖ applicant(**pid**, cityid, income, gmat)
- ❖ location(**cityid**, country)
- ❖ Find applicants that earn a salary greater than 60,000 and have GMAT scores higher than the average score of applicants from the **same US city**"

Q3b

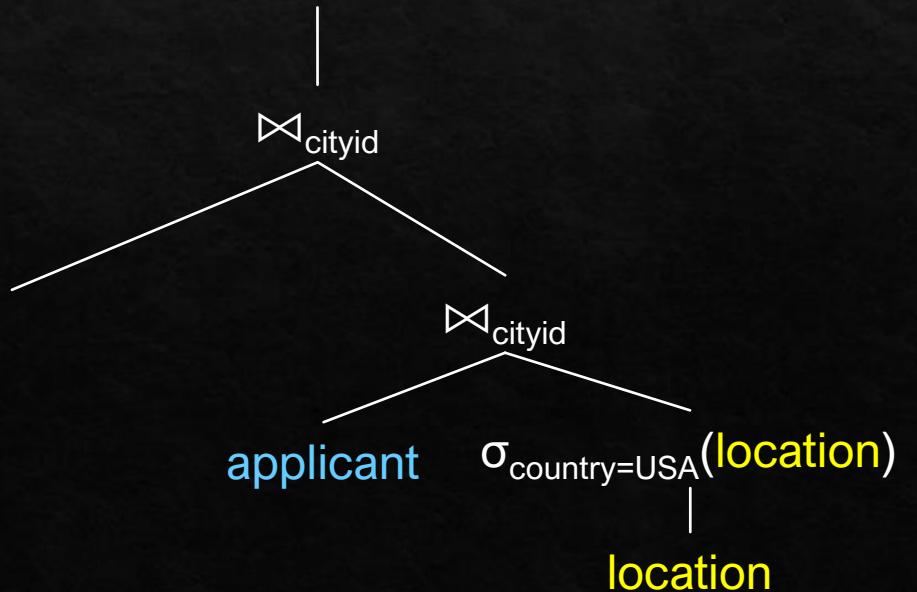
- ❖ applicant(**pid**, cityid, income, gmat)
- ❖ location(**cityid**, country)
- ❖ Find applicants that earn a salary greater than 60,000 and have GMAT scores higher than the average score of applicants from the **same US city**"
- ❖

```
SELECT a1.pid, a1.income FROM applicant a1
WHERE a1.income > 60000
AND a1.gmat >
(SELECT avg(a2.gmat)
FROM applicant a2, location l2
WHERE a1.cityid = a2.cityid
AND a2.cityid = l2.cityid
AND l2.country = "USA");
```

Q3b

Naive

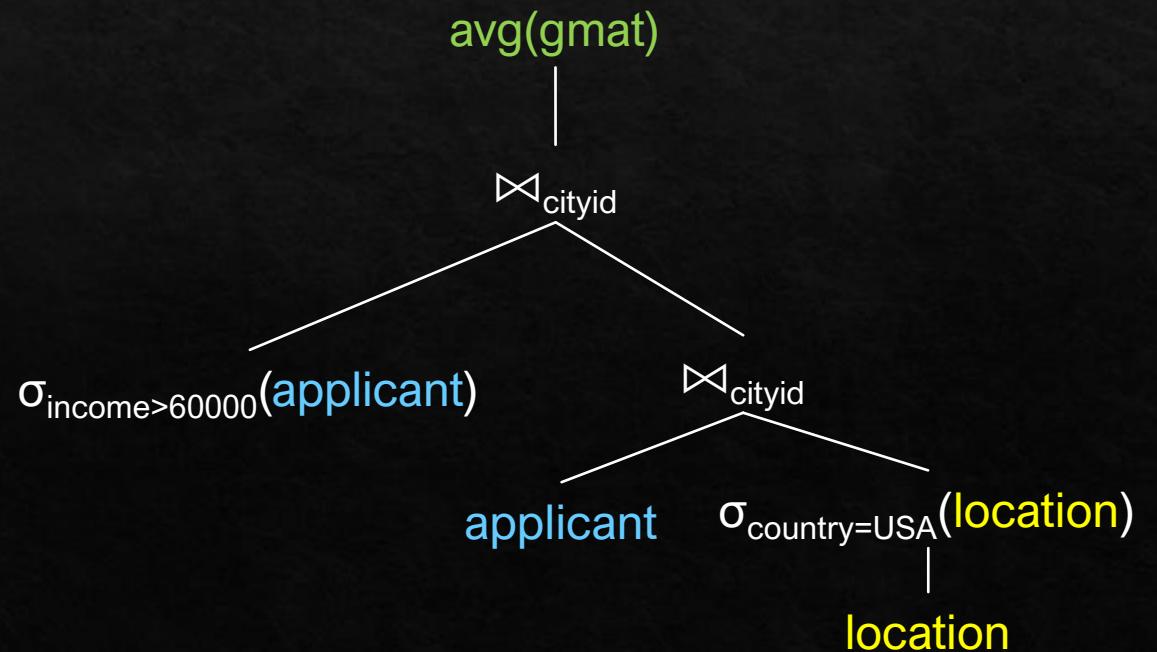
```
❖ SELECT a1.pid, a1.income FROM applicant a1  
WHERE a1.income > 60000  
AND a1.gmat >  
(SELECT avg(a2.gmat)  
FROM applicant a2, location l2  
WHERE a1.cityid = a2.cityid  
AND a2.cityid = l2.cityid  
AND l2.country = "USA");
```



Q3b

Naive

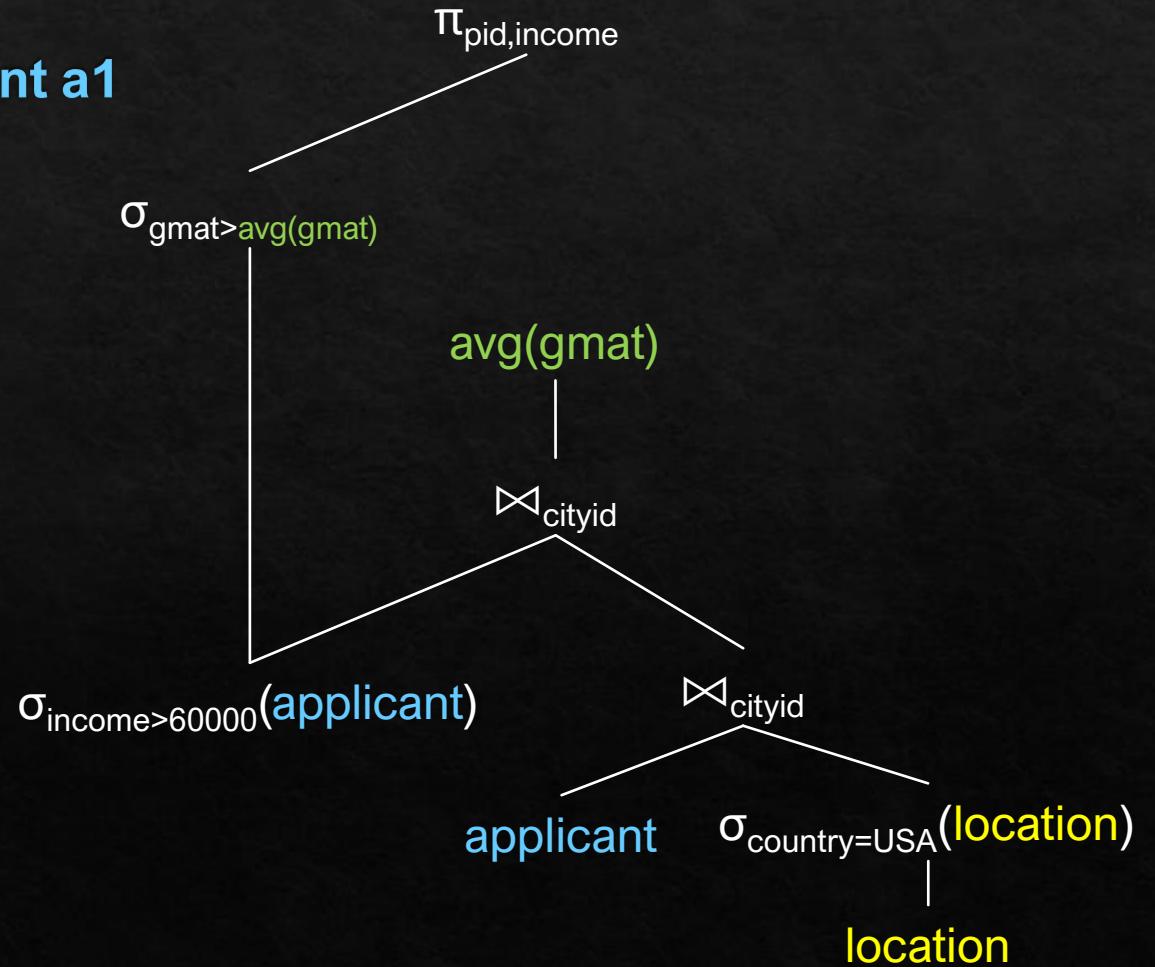
```
❖ SELECT a1.pid, a1.income FROM applicant a1  
WHERE a1.income > 60000  
AND a1.gmat >  
(SELECT avg(a2.gmat)  
FROM applicant a2, location l2  
WHERE a1.cityid = a2.cityid  
AND a2.cityid = l2.cityid  
AND l2.country = "USA");
```



Q3b

Naive

```
❖ SELECT a1.pid, a1.income FROM applicant a1  
WHERE a1.income > 60000  
AND a1.gmat >  
(SELECT avg(a2.gmat)  
FROM applicant a2, location l2  
WHERE a1.cityid = a2.cityid  
AND a2.cityid = l2.cityid  
AND l2.country = "USA");
```

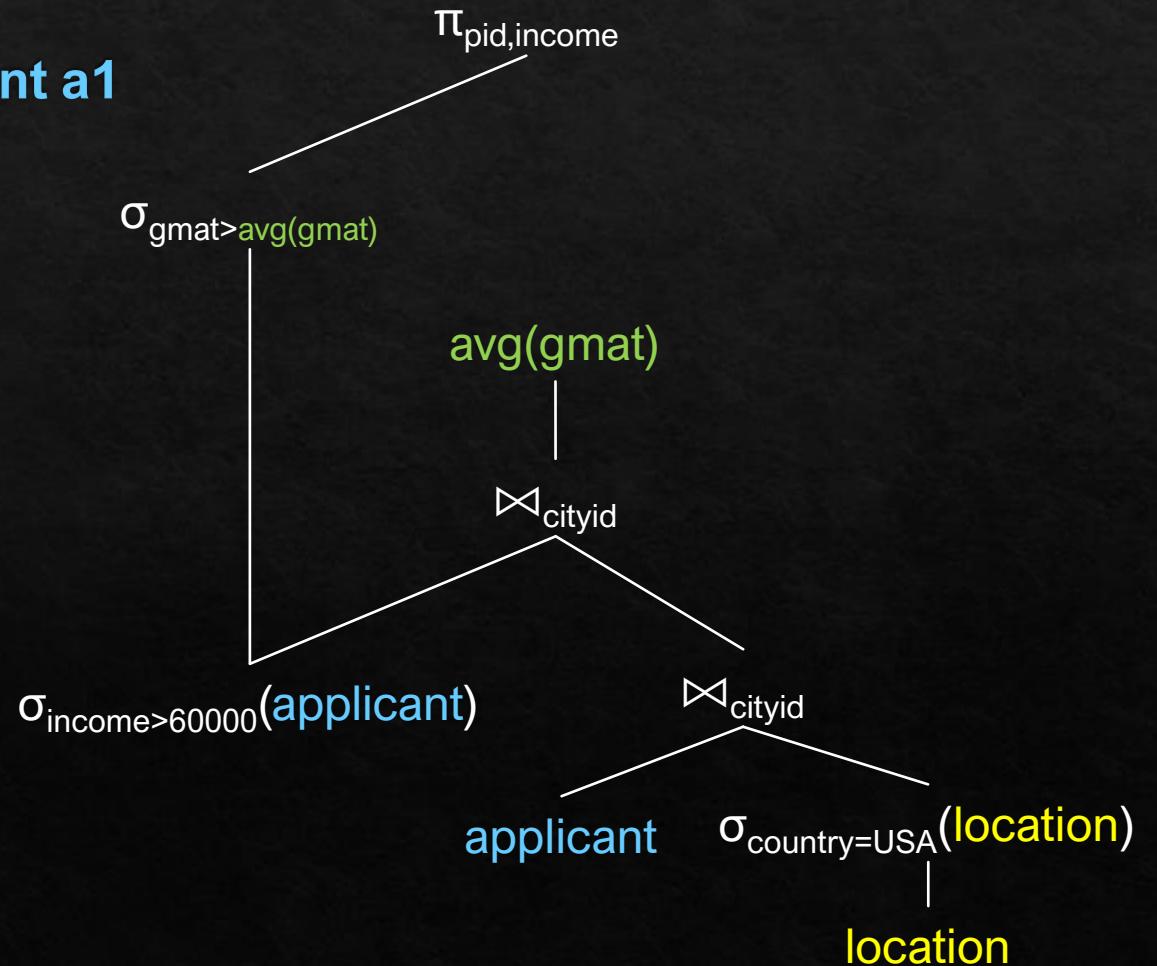


Q3b

Naive

```
❖ SELECT a1.pid, a1.income FROM applicant a1  
WHERE a1.income > 60000  
AND a1.gmat >  
(SELECT avg(a2.gmat)  
FROM applicant a2, location l2  
WHERE a1.cityid = a2.cityid  
AND a2.cityid = l2.cityid  
AND l2.country = "USA");
```

Expensive, computing the inner query
(`SELECT avg...`) for every a1 tuple



Q3b

Improved

❖ Inner query:

```
(SELECT cityid, avg(a.gmat)
FROM applicant a, location l
WHERE a.cityid = l.cityid
AND l.country = "USA"
GROUP BY cityid) as t;
```

❖ Outer query:

```
SELECT a.pid, a.income
FROM applicant a, tmpAgg t
WHERE a.income > 60000
AND a.cityid = t.c1
AND a.gmat > t.c2
```

cityid	avg(gmat)
0	87
1	85
...	...

Q3b

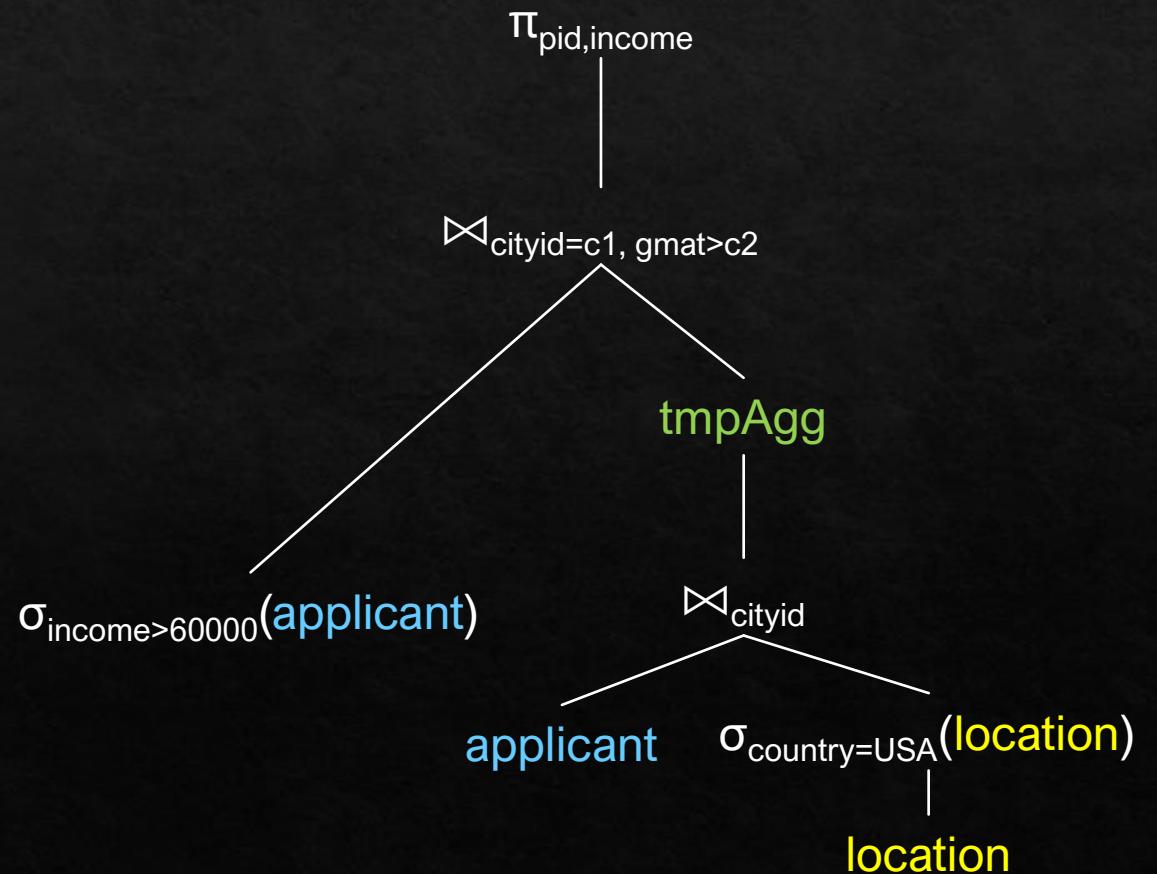
Improved

❖ **Inner query:**

```
(SELECT cityid, avg(a.gmat)
  FROM applicant a, location l
 WHERE a.cityid = l.cityid
   AND l.country = "USA"
 GROUP BY cityid) as t;
```

❖ **Outer query:**

```
SELECT a.pid, a.income
  FROM applicant a, tmpAgg t
 WHERE a.income > 60000
   AND a.cityid = t.c1
   AND a.gmat > t.c2
```



End