

CS3223: Database Management Systems

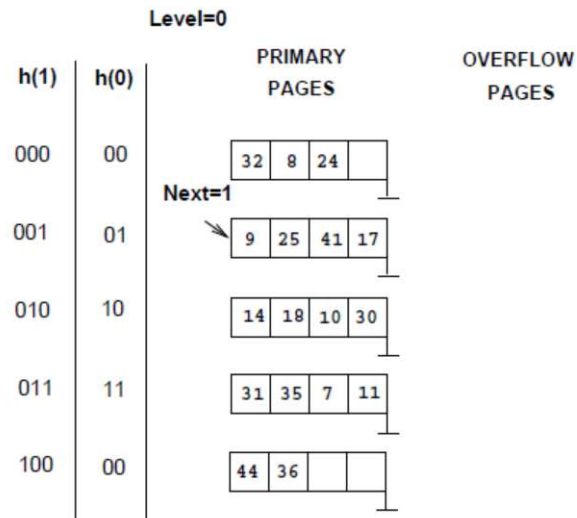
Tutorial 2 (Week 4, Feb 2022)

1. Consider a relation EMP(eno, sno, name, salary) containing 1,000,000 employee records. Suppose each page of the relation contains 10 records, and relation EMP is organized as a sorted file. Suppose sno is a candidate key, with values lying in the range 0-999,999, and that EMP is stored in sno order. Assume that each page is 4K bytes, eno and sno are each 4 bytes, and a pointer to a page is also 4 bytes. For each of the following queries, state which of the following three approaches is most likely to be the cheapest (in terms of I/O count). Justify your answer. For simplicity, assume that the data are uniformly distributed in the domain.
 - Scan the sorted file for EMP.
 - Use a (clustered) B+-tree index on sno. Assume that the nodes are 100% full, and that format 2 is used for the index.
 - Use a hash index on sno. You may assume ideal situation of no overflows, and each bucket is 70% full (and a minimum number of buckets is used). You may assume that format 1 is used for the index.
 - a) Find all employees whose sno < 100,000.
 - b) Find the employee whose sno = 10.
 - c) Find the employees whose sno is in the range 10,000-10,010.
 - d) Find the employees whose sno is not 10.
2. Consider a relation Book(Bookid: integer, Author: string, ...) where Bookid is the key, and its values are assigned in increasing order starting from 1. Suppose Book has 6 million records of 200 bytes each. Suppose we have to perform 10,000 single-record accesses, and 100 range queries of 0.005% of the file, with Bookid as the search key.
 - Use hashing (with key-to-address transformation of the form $x \bmod y$). Suppose the hash table has a load factor of 70% and the bucket size is 4096 bytes. Moreover, assume that records are stored in the bucket (format 1), and there is no overflow of buckets.
 - Use B+-tree. Suppose each node is 70% full, and the sizes of a node, key and address are 4096, 8 and 4 bytes respectively.

Which of the above two methods is better for the application? Under what circumstance will the “loser” outperform the “winner”?

3. (Question taken from the main reference text – Ex 11.2) Consider the Linear Hashing index shown in the figure below. Assume that we split whenever an overflow page is created. Answer the following questions about this index:

- What can you say about the last entry that was inserted into the index if you know that there have been no deletions from this index so far?
- Show the index after inserting an entry with hash value 4, and 15.
- Show the index after deleting the entries with hash values 36 and 44 into the original tree. Assume that the full deletion algorithm is used.
- Find a list of entries whose insertion into the original index would lead to a bucket with two overflow pages. Use as few entries as possible to accomplish this.



- (Adapted from question in the main reference text – Ex 11.11.4) Consider the extendible hash table below. Insert 4, 5 and 7. Show the resultant hash table? Now, delete 7, 5, and 4. What is the resultant hash table?

