

## Sample Solutions

### Question 1

Number of blocks of R =  $500,000/500 = 1000 = |R|$

Number of blocks of S =  $100,000/100 = 1000 = |S|$

Time for one random I/O =  $T_r$

Time for one sequential I/O =  $T_s$

We read R one block at a time.

Cost =  $|R|$  random I/Os (because read head needs to access S relation when checking a tuple from R)

For every tuple of R, we retrieve all the tuples of S that have a matching B value. Assuming that all R.B values are also found in column S.B of relation S, there will be (number of S tuples/number of distinct B values in S) matches for each tuple of R. Since the index on S.B is not clustered, there will be one IO for each tuple.

Cost =  $500,000 * (\text{number of S tuples/number of distinct B values})$  random I/Os

$IO_{plan1} = |R| + \text{number of R tuples} * (\text{number of S tuples/number of distinct B values}) = 1000 + 500,000 * 100,000/40,000 = 1,251,000$  I/Os.

All these I/Os are random I/Os. Thus,  $Time_{plan1} = 1,251,000 * T_r$ .

### Question 2

We read R one block at a time, so  $|R|$  I/Os, so 1000 random I/Os

For each tuple of R, we retrieve all the tuples of S that have a matching C value. Assuming that all the R.C values are also found in the S.C column of relation S, there will be (number of S tuples)/(number of distinct C values in S) matches for each tuple of R. However, since there is clustering on attribute C, this will take only  $(|S|/\text{number of distinct C values})$  IO's. So,

$IO_{plan2} = |R| + \text{number of R tuples} * 1$  will be random whereas the remaining  $(|S|/\text{number of distinct C values} - 1)$  will be sequential. Hence, we have:

$Time_{plan2} = (|R| + \text{number of R tuples} * 1) T_r + (\text{number of R tuples} * [|S|/\text{#distinct C values} - 1]) * T_s$   
 $= 501,000 T_r + 12,000,000 T_s$

Comparison:

In terms of IOs, clearly  $IO_{plan2} > IO_{plan1}$

But what is more important (most of the time) is access time. Now plan 1 is better than plan 2 if

$Time_{plan1} < Time_{plan2}$

$1,251,000 T_r < 501,000 T_r + 12,000,000 T_s$

$$750,000 T_r < 12,000,000 T_s$$

$$T_r < 16 T_s$$

Hence, as long as random access is no more than 16 times more time consuming than sequential IO, plan 1 wins - otherwise plan 2 wins.

### Question 3

$$R = 1000000/20 = 50000 \text{ pages}$$

$$S = 2000000/40 = 50000 \text{ pages}$$

a)

For Block-nested loops join, cost =  $50000 + \lceil 50000/98 \rceil * 50000 = 50000 + 511 * 50000 = 25,600,000$

For Sort-merge join, sort R = sort S =  $2 * 50000 * (1 + \lceil \log_{99} \lceil 50000/100 \rceil \rceil) = 300,000$

So, cost of sort-merge join =  $300,000 + 300,000 + 50,000 + 50,000 = 700,000$

If tables are already sorted, then it costs only 100,000.

For hash join, cost =  $5 * (|R| + |S|) = 5 * 100,000 = 500,000$ :

Since  $\sqrt{50k} > 100$ , assuming uniform distribution of data, we need to recursively partition the file.

So, partitioning phase involves: read R, write partitions of R; read partitions of R, write more partitions of R.

Total cost =  $4 * 50000 = 200,000$

Similar cost for S in partitioning = 200,000

Joining phase only reads partitions of R and the corresponding partitions of S. So, its one read of R and S. So, another 100,000.

b)

Essentially, for each R-tuple, it will probe the index for a match.

Cost to read R = 50000

In the worst case, R is not sorted and we are not able to buffer any record of S.

The probing cost =  $1000000 * (3+1) = 4,000,000$ . This would also be the result if the index is unclustered.

Even if we assume we can buffer the top 2 or 3 levels, we still need at least 1,000,000 I/Os.

The best case happens when R.a has only a single value. The probing cost =  $3 + 1 = 4$ .

### Question 4

Basic idea:

Generate sorted runs of R. Instead of merging the runs into a single sorted run, stop the merging when the number of runs is  $< \lfloor (B-1)/2 \rfloor$  where B is the number of buffer pages.

Repeat the above step for S.

Allocate 1 buffer page for each run of R and S. Allocate 1 buffer page for join output. As tuples of the sorted R and S are produced, they can be checked whether the join predicate is satisfied.

Savings: Cost of reading and writing a single sorted run of R and S