

Analysis and Design of Algorithms



CS3230
C23530

Week 9

Greedy Algorithms

Diptarka Chakraborty

Ken Sung

Dynamic Programming (Recap)

- Expressing the solution recursively
- Overall there are only polynomial number of subproblems
- But there is a huge overlap among the subproblems. So the recursive algorithm takes exponential time (solving same subproblem multiple times)
- So we compute the recursive solution iteratively in a bottom-up fashion (like in case of Fibonacci numbers). This avoids wastage of computation and leads to an efficient implementation

Question 4 (last lecture)

We have n cents and need to get change in terms of denominations d_1, d_2, \dots, d_k . Goal is to use the fewest total number of coins.

Example: If denominations are 25c, 10c, and 1c, then solution for $n = 30c$ should be 10c+10c+10c.

Let $M[j]$ be the fewest number of coins needed to change j cents. Write a recursive formula for $M[j]$ in terms of $M[i]$ with $i < j$.

Question 4: Solution

Optimal substructure: Suppose $M[j] = t$, meaning that

$$j = d_{i_1} + d_{i_2} + \cdots + d_{i_t}$$

for some $i_1, \dots, i_t \in \{1, \dots, k\}$. Then, if $j' = d_{i_1} + d_{i_2} + \cdots + d_{i_{t-1}}$, $M[j'] = t - 1$, because otherwise if $M[j'] < t - 1$, by **cut-and-paste** argument, $M[j] < t$.

$$M[j] = \begin{cases} 1 + \min_{i \in [k]} M[j - d_i], & j > 0 \\ 0, & j = 0 \\ \infty, & j < 0 \end{cases}$$

Question 5 (last lecture)

Using the above, derive a DP algorithm to compute the minimum number of coins of denomination d_1, \dots, d_k needed to change n cents.

Question 5: Solution

NUM-COINTS-DP(n, d):

for $j = 0, \dots, n$:

$M[j] \leftarrow \infty$

$M[0] \leftarrow 0$

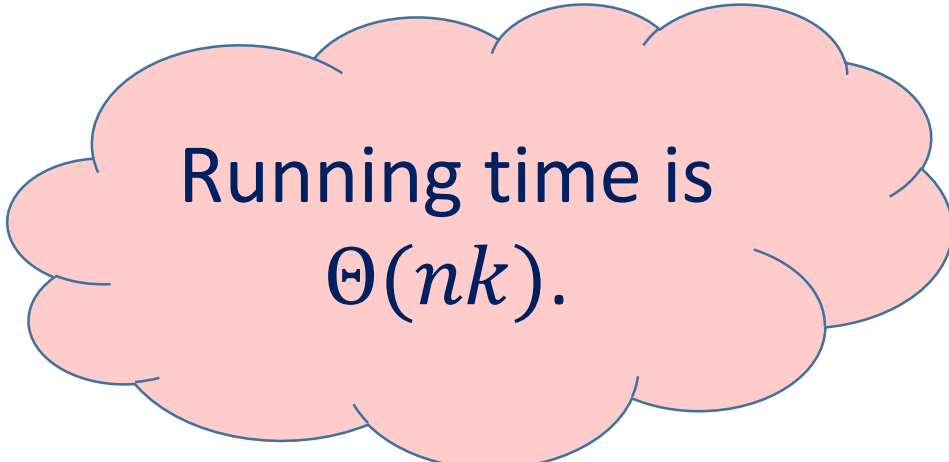
for $j = 1, \dots, n$:

for $i = 1, \dots, k$:

if $(j - d_i \geq 0) \wedge (M[j - d_i] + 1 < M[j])$:

$M[j] \leftarrow M[j - d_i] + 1$

return $M[n]$



Running time is
 $\Theta(nk)$.

Today: Greedy Algorithms

A very general technique, like divide-and-conquer and dynamic programming

Technique is to recast the problem so that only one subproblem needs to be solved at each step. Beats divide-and-conquer and dynamic programming, **when it works.**

Longest Increasing Subsequence

Applications in genome comparisons, ...

What is a subsequence?

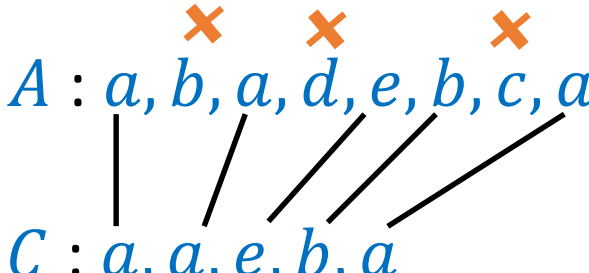
Sequence $A : a_1, a_2, \dots, a_n$

Can be stored in an array $A[1..n]$, $A[j] : a_j$

$A[1..k] : a_1, a_2, \dots, a_k$

Definition: C is said to be a subsequence of A if we can obtain C by removing 0 or more elements from A .

Example: $A : a, b, a, d, e, b, c, a$
 $C : a, a, e, b, a$



A more formal definition:

C is a subsequence of A if there exists k integers: $1 \leq i_1 < \dots < i_k \leq n$ s.t.
for all $1 \leq j \leq k$ $C[j] = A[i_j]$

Longest Increasing Subsequence (LIS)

Problem: Given a sequence of n numbers find a longest increasing subsequence (**LIS**)

10, 3, 7, 2, 4, 6, 11, 8, 5, 9

Question 1

Can you design an $O(n^2)$ time algorithm for the LIS problem using the LCS algorithm discussed in the last lecture?

Longest Increasing Subsequence (LIS)

Given a sequence of n numbers

10, 3, 7, 2, 4, 6, 11, 8, 5, 9

$O(n \log n)$

1. Sort n numbers

2, 3, 4, 5, 6, 7, 8, 9, 10, 11

2. Use **dynamic programming** to find **longest common subsequence** between **sorted** and the **original** sequence

$O(n^2)$

Can we do better?


Patience Solitaire

Given n cards deal into piles according to the following two rules

1. Not allowed to place a higher-valued card over a lower-valued card
2. Can form a new pile and place card


Objective: Form as few piles as possible.

Patience Solitaire



First card to deal

Patience Solitaire



Question 2


What strategy will you use while playing Patience Solitaire game so that you can always win?

1. Place card on the leftmost allowed pile; if not possible then create a new pile
2. Place card on the rightmost allowed pile; if not possible then create a new pile
3. Place card on any arbitrary allowed pile; if not possible then create a new pile
4. Since we cannot see the future cards we have to rely on luck to win. We could only ensure number of piles upto a constant times the minimum one.


Patience Solitaire



Greedy strategy: Place on the **leftmost** “allowed” pile; else create new pile




Patience Solitaire



Observations:


- ❖ Cards in each pile are in decreasing order
- ❖ Any **increasing subsequence** contains at most one card from each pile



Patience Solitaire and LIS

Weak Duality:

Length of any increasing subsequence \leq # of piles in any valid game



Patience Solitaire and LIS




Length of LIS \leq Min. number of piles

\leq # of piles in
greedy strategy

Why do not we take the top cards from each pile?

May not be a valid
subsequence




Patience Solitaire and LIS

Length of LIS \leq Min. number of piles

\leq # of piles in
greedy strategy


Can we take some card from each pile to form
an increasing subsequence?



Patience Solitaire and LIS

Strong Duality:

Length of LIS = Min. number of piles




Patience Solitaire and LIS

Length of LIS \leq Min. number of piles \leq # of piles in greedy strategy

Can we take some card from each pile to form an increasing subsequence?


YES!!!
why?



Patience Solitaire



Greedy strategy: Place on **leftmost** “allowed” pile; else create new pile




Patience Solitaire and LIS

Length of LIS \leq Min. number of piles

\leq # of piles in greedy strategy

Can we take some card from each pile to form an increasing subsequence?

YES!!!
why?




Patience Solitaire and LIS

Strong Duality:

Found by our greedy strategy

Length of LIS = Min. number of piles



Running Time

Greedy strategy: Place on the **leftmost** “allowed” pile; else create new pile

At most **n** stacks

- ❖ Use stack to implement each pile
- ❖ How do you find the leftmost “allowed” pile?

Recall, top cards are in sorted order

Binary search

Time = $O(n \log n)$

A few important points

- How do you use this greedy algorithm to sort?
- Each pile is sorted in decreasing order
- **Simple Exercise:** Use idea of merge sort to merge all the piles in total time $O(n \log n)$
- What is the best case in this sorting algorithm?
- If numbers are $1, 2, \dots, n$ then the above algorithm can be implemented in time $O(n \log \log n)$


Lets leave it for advanced
algorithm course !!!

Paradigm for greedy algorithms

1. Cast the problem where we have to **make a choice and are left with one subproblem** to solve.
2. Prove that there is always an **optimal solution to the original problem that makes the greedy choice**, so the greedy choice is safe.
3. Use **optimal substructure** to show that we can combine an optimal solution to the subproblem with the greedy choice to get an optimal solution to the original problem.

Question 3

Who is the Master of Algorithms pictured below?



- ☐ Edsger Dijkstra
- ☐ Robert Prim
- ☐ Joseph Kruskal
- ☐ Jack Edmonds

Question 3

Edsger Dijkstra

Turing Award winner

Known for Dijkstra's algorithm
(greedy algorithm!), structured
programming (GOTO statement
considered harmful),
semaphores, mutex, deadlock
(Dining Philosophers problem),
...




Huffman Code

Applications in data compression, ...



Binary coding

Alphabet set $A : \{a_1, a_2, \dots, a_n\}$

A text File: a sequence of alphabets



A text file F



Binary coding of F

Question: How many bits needed to encode a text file with m characters?

Answer: $m \lceil \log_2 n \rceil$ bits.

Fixed length coding

Alphabet set $A : \{a_1, a_2, \dots, a_n\}$

Question: What is a binary coding of A ?

Answer: $\gamma: A \rightarrow \text{binary strings}$

Question: What is a **fixed length** coding of A ?

Answer: each alphabet \leftarrow a unique binary string of length $\lceil \log_2 n \rceil$.

Question: How to decode a **fixed length binary** coding?

Answer: easy 😊

0100|1010|0000|1011|...

Fixed length coding

Alphabet set A : $\{a_1, a_2, \dots, a_n\}$


Question: Can we use fewer bits to store alphabet set A ?

Answer: No.

Question: Can we use fewer bits to store a file ?

Answer: Yes

Huge variation in the frequency of alphabets in a text



ENGLISH LETTER FREQUENCIES			
Per One-Thousand Letters			
Sorted By Letter		Sorted By Frequency	
A	73	E	130
B	9	T	93
C	30	N	78
D	44	R	77
E	130	I	74
F	28	O	74
G	16	A	73
H	35	S	63
I	74	D	44
J	2	H	35
K	3	L	35
L	35	C	30
M	25	F	28
N	78	P	27
O	74	U	27
P	27	M	25
Q	3	Y	19
R	77	G	16
S	63	W	16
T	93	V	13
U	27	B	9
V	13	X	5
W	16	K	3
X	5	Q	3
Y	19	J	2
Z	1	Z	1

http://en.wikipedia.org/wiki/Letter_frequency

Huge variation in the frequency of alphabets in a text

Question: How to exploit variation in the frequencies of alphabets ?

Answer:

More frequent alphabets ← coding with **shorter bit string**

Less frequent alphabets ← coding with **longer bit string**

Variable length encoding

Average bit length per symbol using γ :

$$\text{ABL}(\gamma) = \sum_{x \in A} f(x) \cdot |\gamma(x)|$$

$$= 0.45 \times 1 + 0.18 \times 2 + (0.15 + 0.12 + 0.10) \times 3$$

$$= 1.92$$

Alphabets	Frequency f	Encoding γ
a	0.45	0
b	0.18	10
c	0.15	110
d	0.12	101
e	0.10	111

There is a serious problem with the γ encoding. Can you see?

Variable length encoding

Can you
fix it?

Symbols	Frequency f	Encoding γ
a	0.45	0
b	0.18	10
c	0.15	110
d	0.12	101
e	0.10	111

Average bit length per symbol using γ :

$$\text{ABL}(\gamma) = \sum_{x \in A} f(x) \cdot |\gamma(x)|$$

$$= 0.45 \times 1 + 0.18 \times 2 + (0.15 + 0.12 + 0.10) \times 3$$

$$= 1.92$$

Question: How will you decode 01010111 ?

Answer: *abbe* or *adae* ☹

Question: What is the source of this ambiguity ?

Answer: $\gamma(b)$ is a prefix of $\gamma(d)$.

Variable length Coding

Alphabets	Frequency f	Encoding γ
a	0.45	0
b	0.18	100
c	0.15	110
d	0.12	101
e	0.10	111

Average bit length per symbol using γ :

$$\begin{aligned} \text{ABL}(\gamma) &= \sum_{x \in A} f(x) \cdot |\gamma(x)| \\ &= 0.45 \times 1 + 0.18 \times 3 + (0.15 + 0.12 + 0.10) \times 3 \\ &= 2.1 \end{aligned}$$

Prefix Coding

Definition:

A coding $\gamma(A)$ is called **prefix coding** if there does not exist $x, y \in A$ such that

$$\gamma(x) \text{ is prefix of } \gamma(y)$$

Algorithmic Problem: Given a set A of n alphabets and their frequencies, compute coding γ such that

- γ is prefix coding
- $ABL(\gamma)$ is **minimum**.

The challenge of the problem





Among all possible binary coding of A , how to find the **optimal prefix coding** ?

The novel idea of Huffman




Binary coding



Binary tree


A labeled binary tree



Leaf nodes \rightarrow alphabets

Code of an alphabet = **Label of path from root**

A labeled binary tree



Leaf nodes → alphabets

Code of an alphabet = **Label of path from root**

01,
001,
0000,
11,
100,
10110,
10111


Variable length Coding

Alphabets	Frequency f	Encoding γ
a	0.45	0
b	0.18	100
c	0.15	110
d	0.12	101
e	0.10	111

Question:

How to build the labeled tree for a prefix code ?

$\{0, 100, 101, 110, 111\}$



Prefix code and Labelled Binary tree

Theorem:

For each prefix code of a set A of n alphabets, there exists a binary tree T on n leaves s.t.


- There is a bijjective mapping between the alphabets and the leaves.
- The label of a path from root to a leaf node corresponds to the prefix code of the corresponding alphabet.

Question: Can you express **Average bit length** of γ in terms of its binary tree T ?

$$\begin{aligned}\text{ABL}(\gamma) &= \sum_{x \in A} f(x) \cdot |\gamma(x)| \\ &= \sum_{x \in A} f(x) \cdot |\text{depth}_T(x)|\end{aligned}$$

Finding the **labeled binary tree** for the optimal
prefix codes

Is the following prefix coding **optimal** ?



NO

Observations on the binary tree of the optimal prefix code

Lemma:

The binary tree corresponding to optimal prefix coding must be a **full binary tree**:

Every internal node has degree exactly 2.

Question: What next ?

We need to see the influence of frequencies on the optimal binary tree.

Let a_1, a_2, \dots, a_n be the alphabets of A in non-decreasing order of their frequencies.

Observations on the binary tree of the optimal prefix code

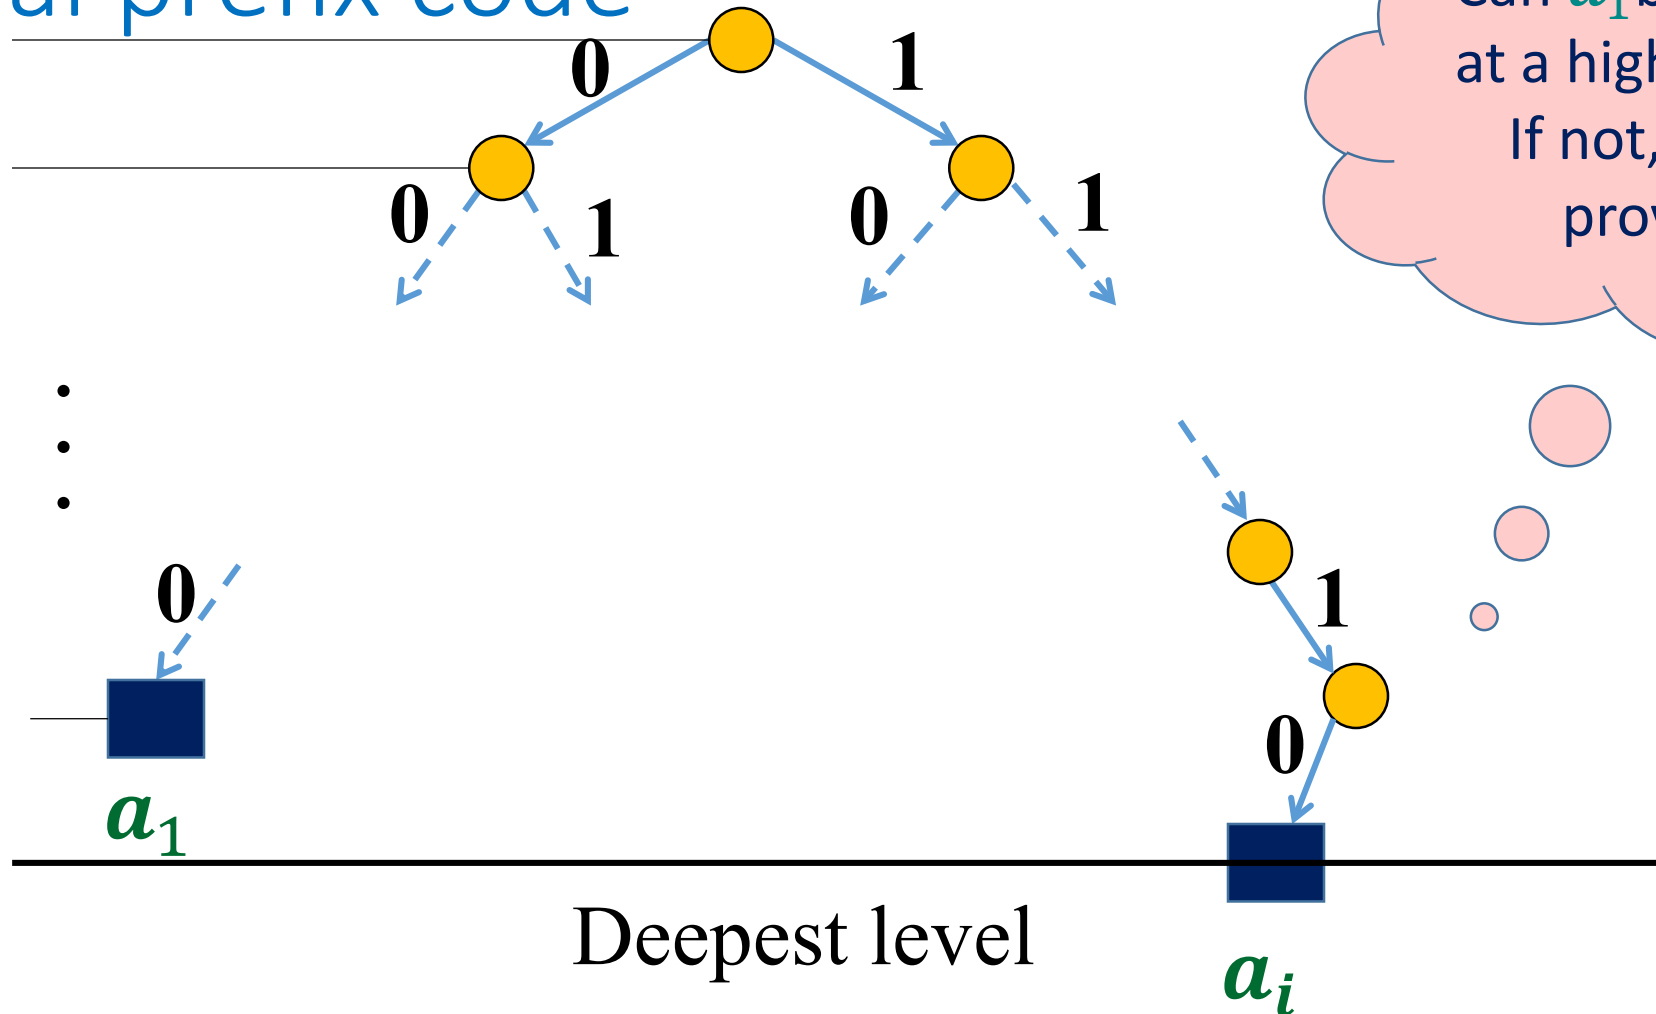
Intuitively, **more frequent** alphabets should be **closer to the root** and **less frequent** alphabets should be **farther from the root**.

But how to organize them to achieve optimal prefix code ?

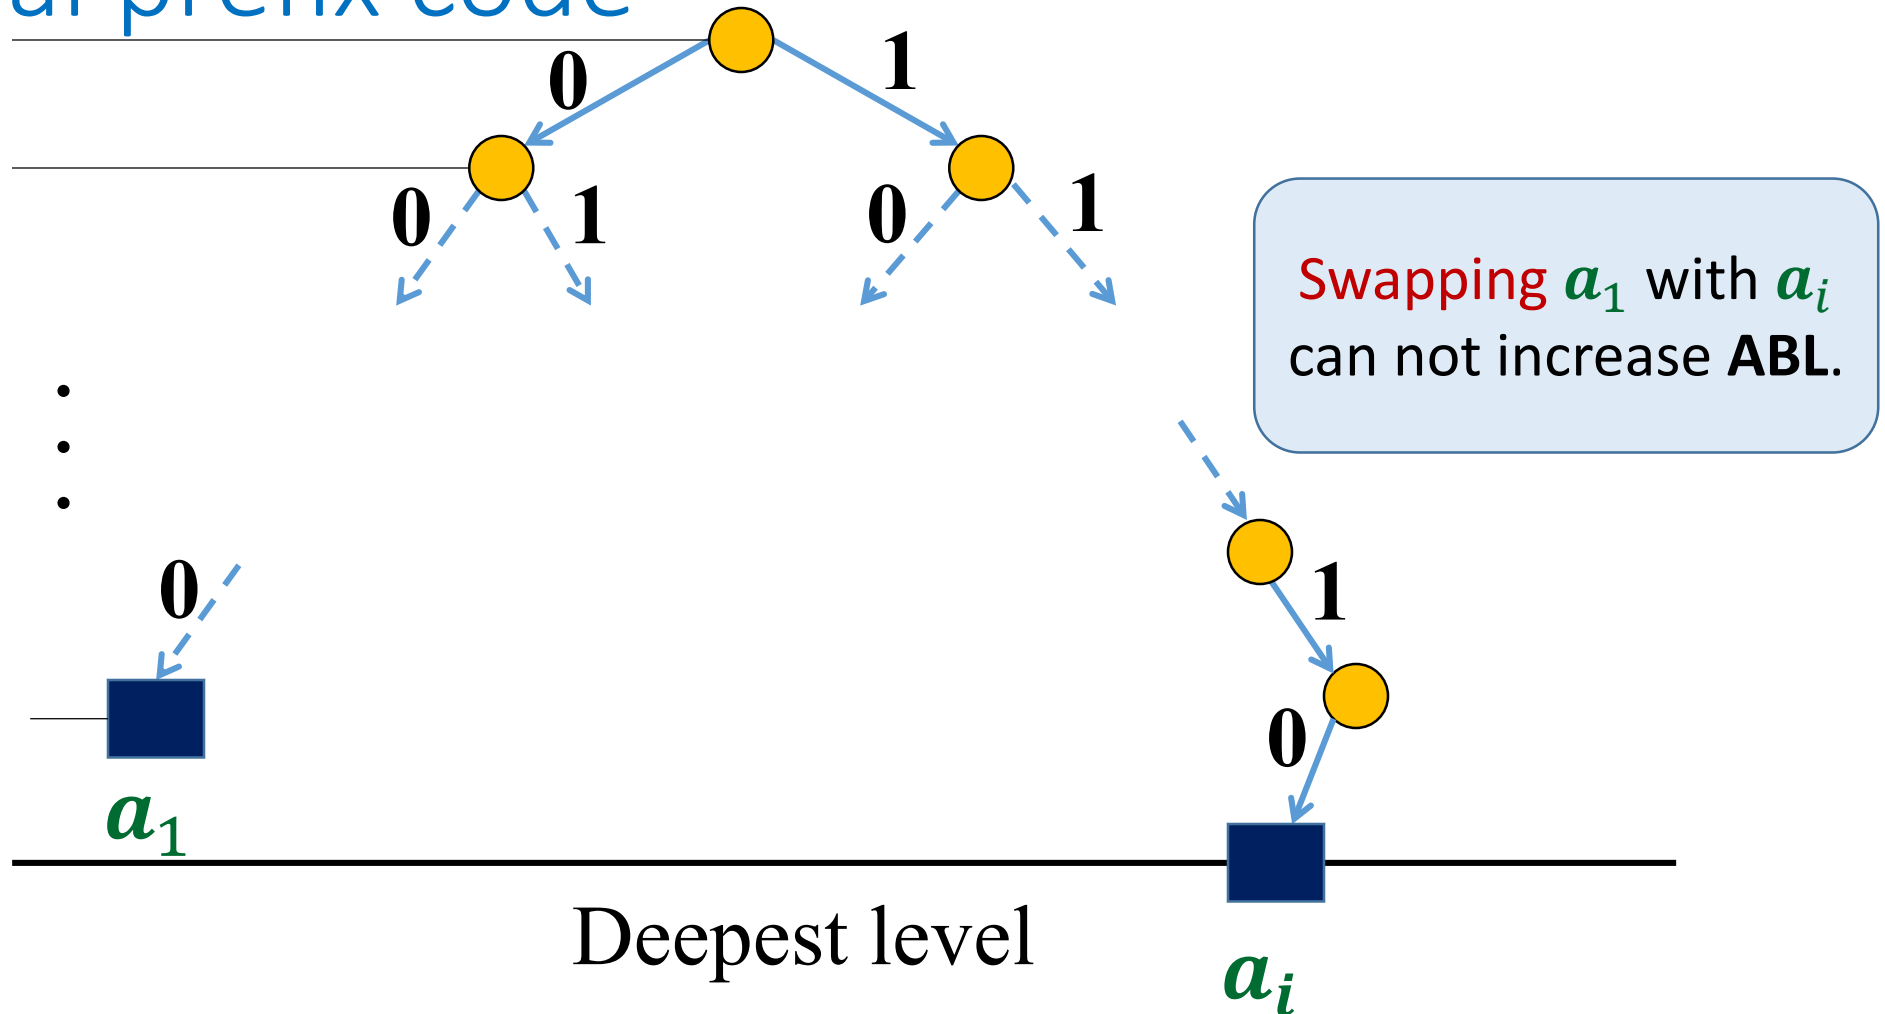
- We shall now make some simple observations about the structure of the binary tree corresponding to the optimal prefix codes.
- These observations will be about some local structure in the tree.
- Nevertheless, these observations will play a crucial role in the design of a binary tree with optimal prefix code for given A .

Please pay full attention on the next few slides.

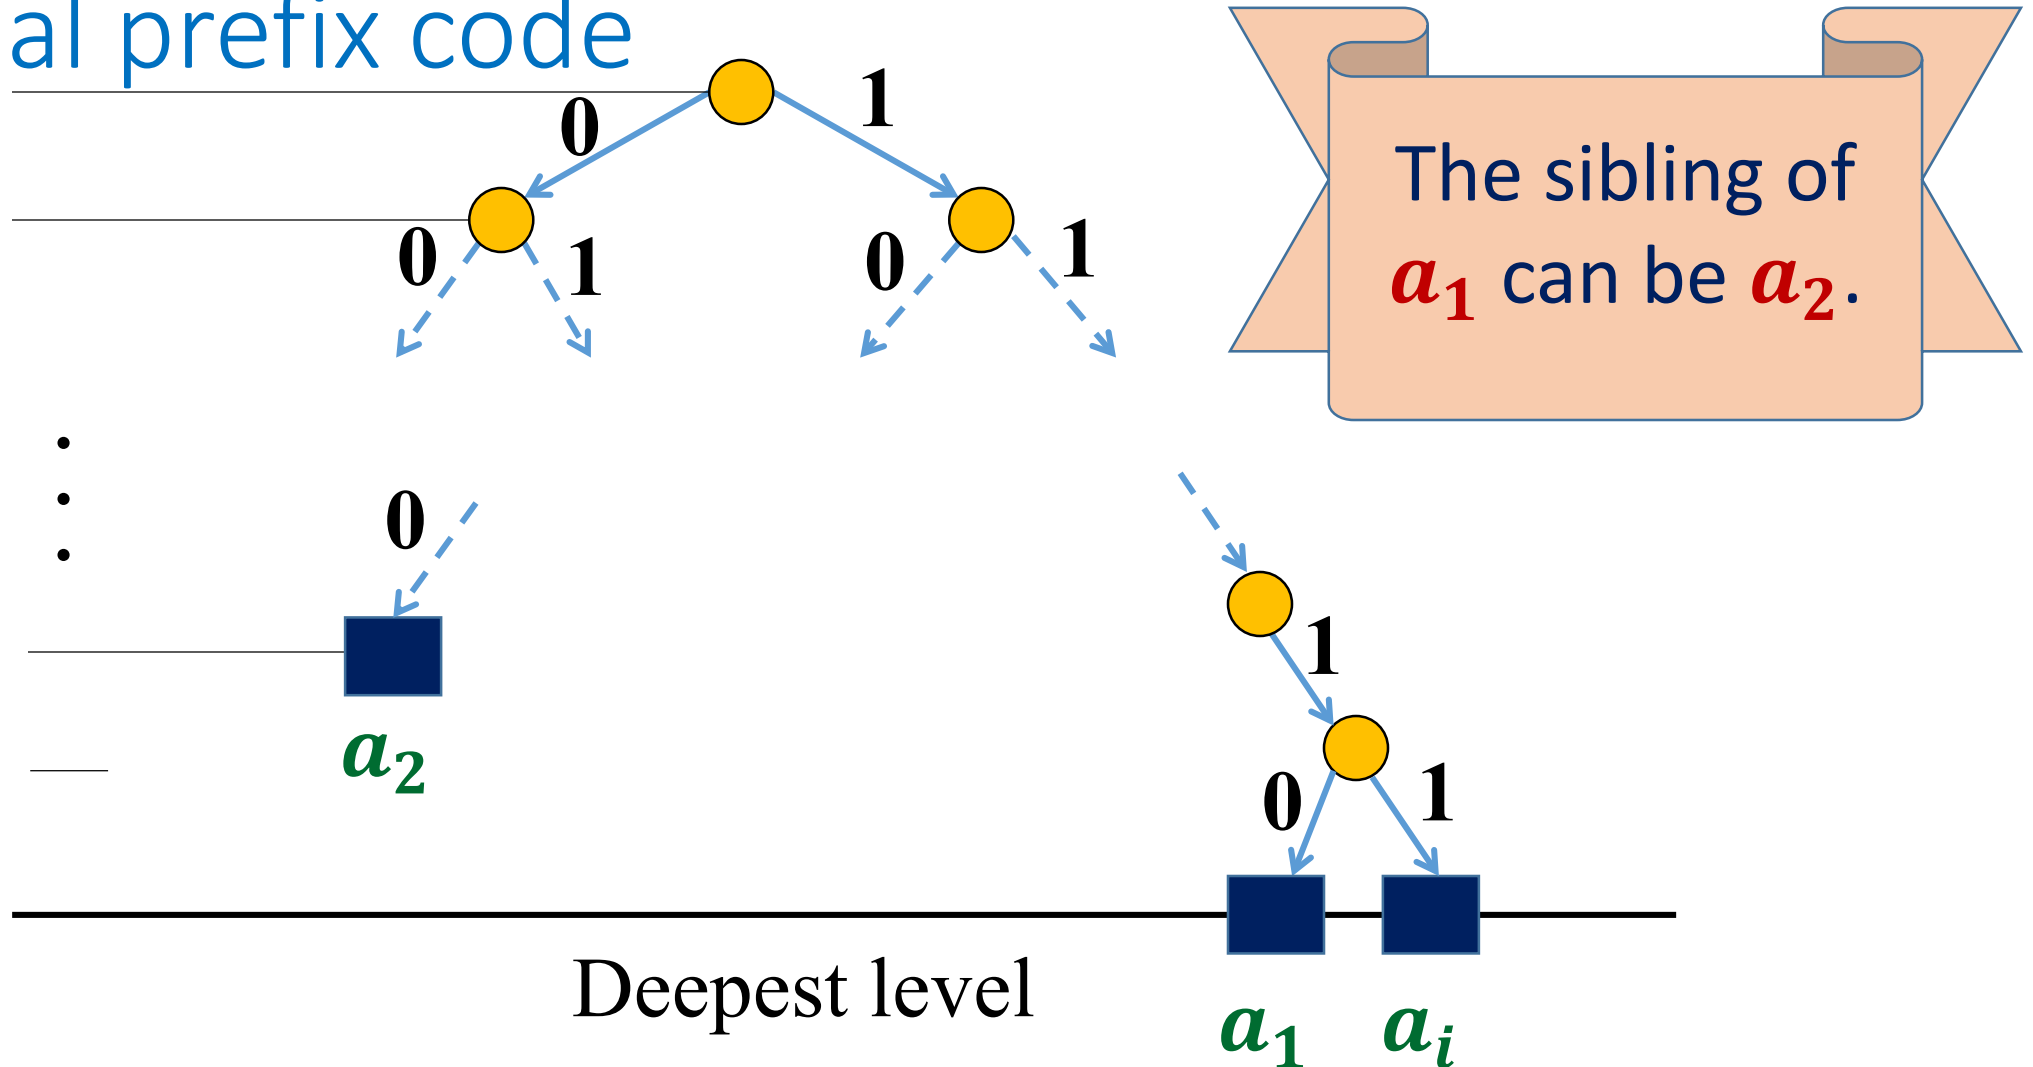
Observations on the binary tree of the optimal prefix code



Observations on the binary tree of the optimal prefix code



Observations on the binary tree of the optimal prefix code



An important observation

Lemma: There exists an optimal prefix coding in which a_1 and a_2 appear as siblings in the corresponding labeled binary tree.

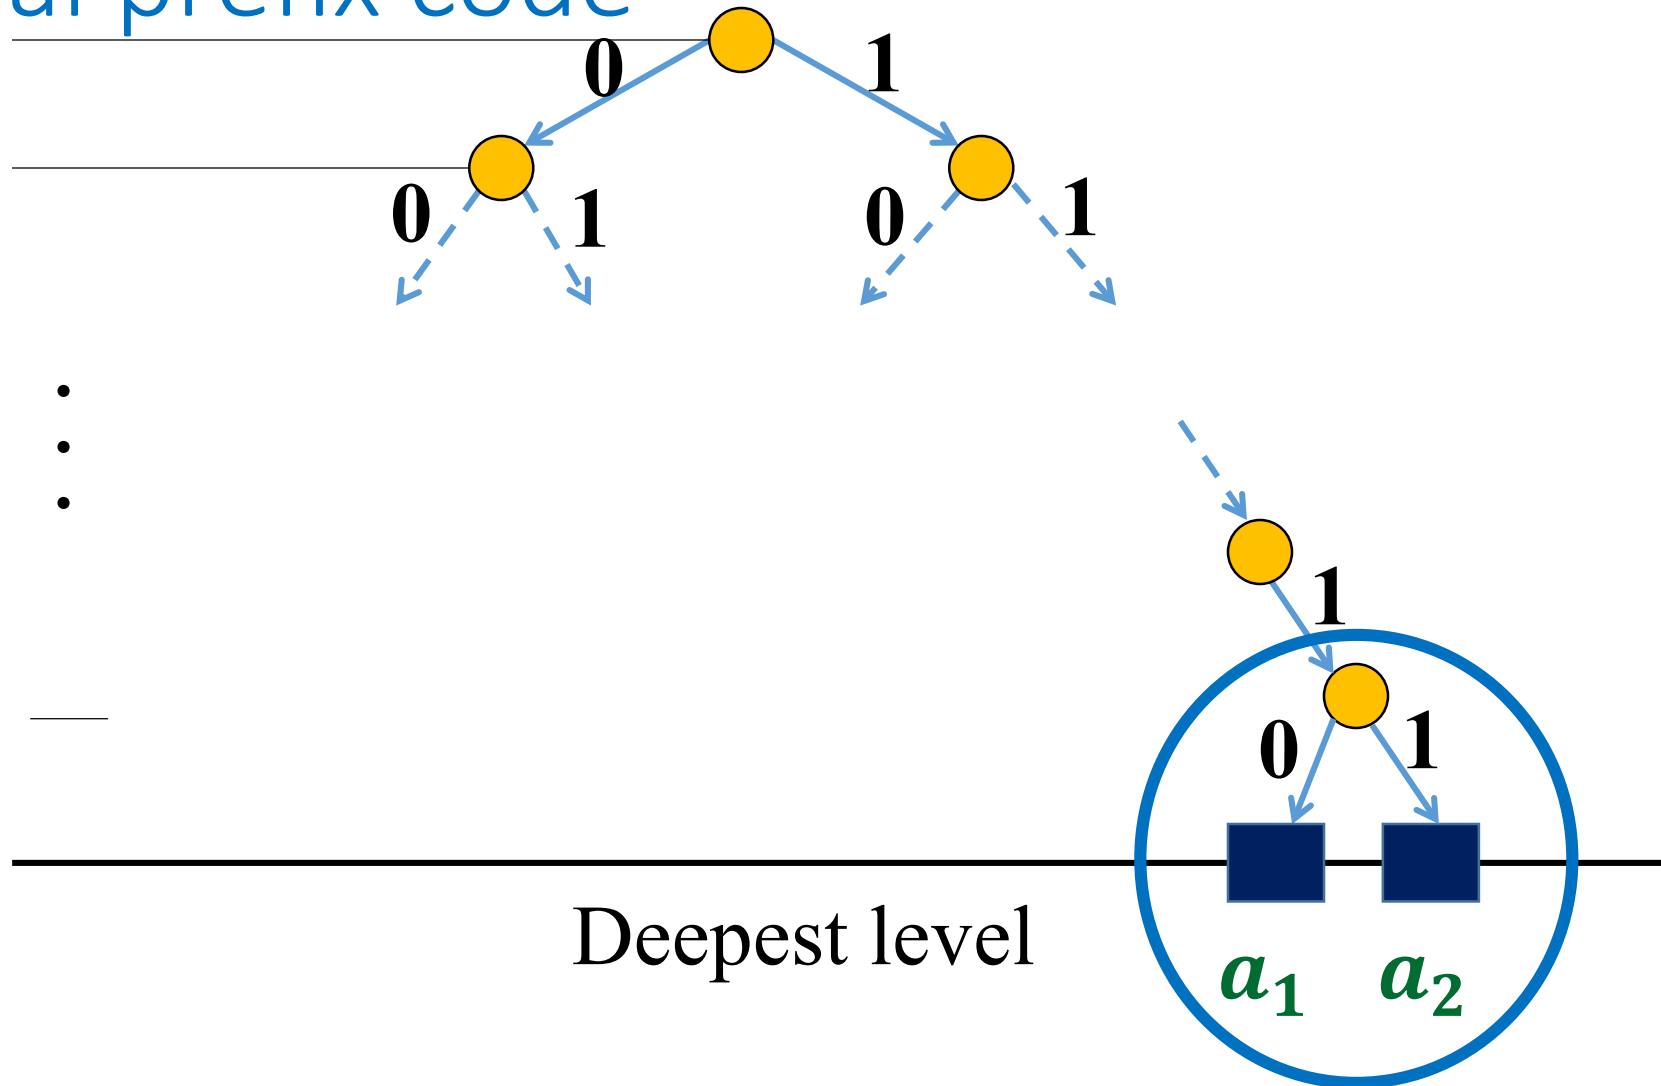
Important note: It is inaccurate to claim that “In every optimal prefix coding, a_1 and a_2 appear as siblings in the labeled binary string.”

But algorithmic implication of the Lemma mentioned above is quite important:


→ We just need to focus on that binary tree of optimal prefix coding in which a_1 and a_2 appear as siblings.

This lemma is a powerful hint to the design of optimal prefix code.

Observations on the binary tree of the optimal prefix code




Observations on the binary tree of the optimal prefix code



Key Idea to design an algorithm

$A = a_1, a_2, \dots, a_n$ be n alphabets in increasing order of frequencies




$A' = a_3, \dots, a', \dots, a_n$ be $n - 1$ alphabets in increasing order of frequencies with $f(a') = f(a_1) + f(a_2)$

Intuition (from the previous slide):

May be : The optimal prefix code of $A' \rightarrow$ optimal prefix code of A

Key Idea to design an algorithm

$A = a_1, a_2, \dots, a_n$ be n alphabets in increasing order of frequencies



$A' = a_3, \dots, a', \dots, a_n$ be $n - 1$ alphabets in increasing order of frequencies with $f(a') = f(a_1) + f(a_2)$

Question: What should be the relation between $\text{OPT}_{\text{ABL}}(A)$ and $\text{OPT}_{\text{ABL}}(A')$?

Answer: $\text{OPT}_{\text{ABL}}(A) = \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$

Observation: If this relation is true, we have an algorithm for optimal prefix codes.

The algorithm based on

$$\text{OPT}_{\text{ABL}}(A) = \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$$

OPT(A)

{ If $|A|=2$, return  ;

else

{ Let a_1 and a_2 be the two alphabets with **least frequencies**.

Remove a_1 and a_2 from A;

Create a new alphabet a' ;

$f(a') \leftarrow f(a_1) + f(a_2)$;

Insert a' into A;

$T \leftarrow \text{OPT}(A)$;

Replace node  in T by  ;

return T ;

}}

Sort the alphabets according to frequencies
Takes $O(n \log n)$ time

Do **binary search** to update the sorted list in $O(\log n)$ time

Overall time = $O(n \log n)$

How to prove


$$\text{OPT}_{\text{ABL}}(A) = \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2) \quad ?$$

Question 1: Can we derive a prefix coding for A from $\text{OPT}(A')$?

Question 2: Can we derive a prefix coding for A' from $\text{OPT}(A)$?


A prefix coding for A from $\text{OPT}(A')$

T' : the binary tree corresponding to $\text{OPT}_{\text{ABL}}(A')$



A prefix coding for A from $\mathbf{OPT}(A')$

T' : the binary tree corresponding to $\text{OPT}_{\text{ABL}}(A')$




This gives a prefix coding for A with $\mathbf{ABL} = ??$

Question 4

Express ABL in terms of $\text{OPT}_{\text{ABL}}(A')$, $f(a_1)$ and $f(a_2)$.

A prefix coding for A from $\text{OPT}(A')$

T' : the binary tree corresponding to $\text{OPT}_{\text{ABL}}(A')$




This gives a prefix coding for A with $\mathbf{ABL} = \mathbf{OPT}_{\mathbf{ABL}}(A') + f(a_1) + f(a_2)$

$$\Rightarrow \text{OPT}_{\text{ABL}}(A) \leq \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$$


A prefix coding for A' from $\text{OPT}(A)$

T : the binary tree corresponding to $\text{OPT}_{\text{ABL}}(A)$



A prefix coding for A' from $\text{OPT}(A)$

T : the binary tree corresponding to $\text{OPT}_{\text{ABL}}(A)$



This gives a prefix coding for A' with $\text{ABL} = \text{OPT}_{\text{ABL}}(A) - f(a_1) - f(a_2)$

$$\rightarrow \text{OPT}_{\text{ABL}}(A') \leq \text{OPT}_{\text{ABL}}(A) - f(a_1) - f(a_2)$$

How to prove

$$\text{OPT}_{\text{ABL}}(A) = \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2) \quad ?$$

We proved

- $\text{OPT}_{\text{ABL}}(A) \leq \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$

(1)

- $\text{OPT}_{\text{ABL}}(A') \leq \text{OPT}_{\text{ABL}}(A) - f(a_1) - f(a_2)$



$$\text{OPT}_{\text{ABL}}(A) \geq \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$$

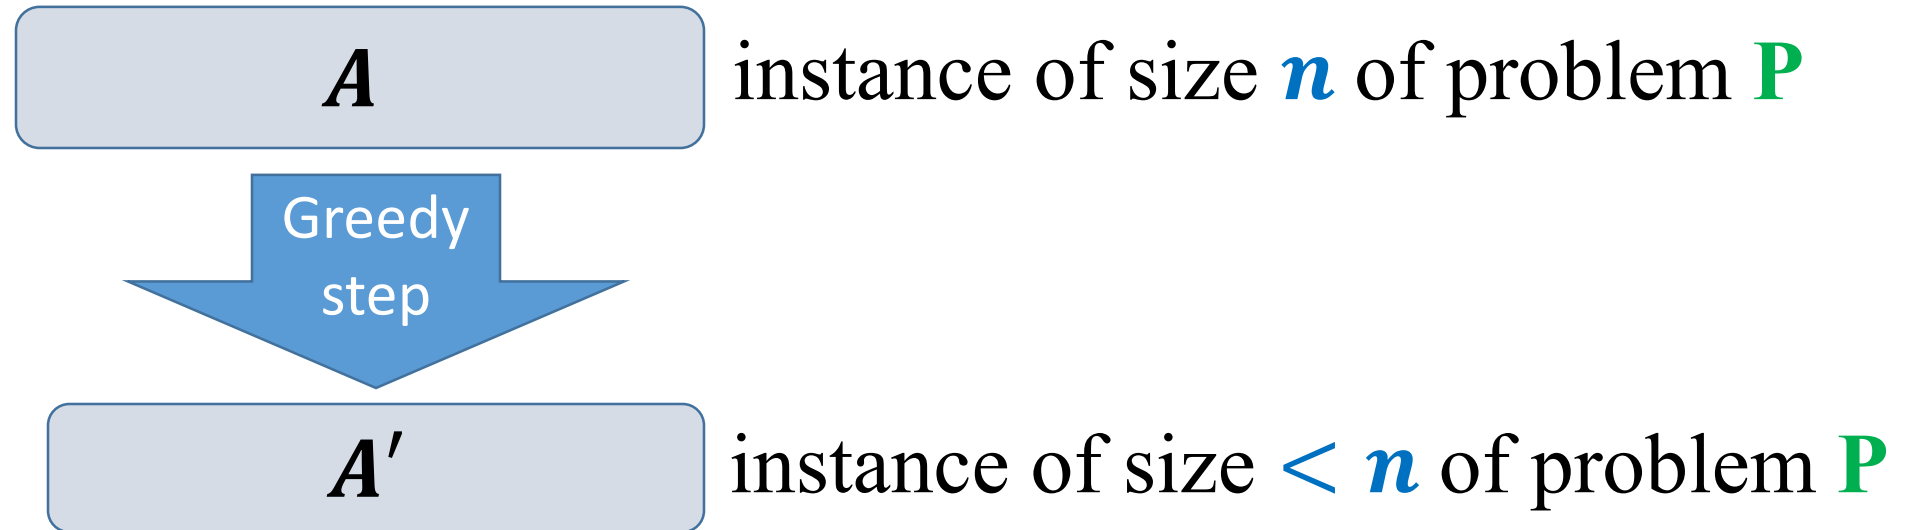
(2)

Using (1) and (2)

$$\text{OPT}_{\text{ABL}}(A) = \text{OPT}_{\text{ABL}}(A') + f(a_1) + f(a_2)$$

To prove that a greedy strategy works

P: A given optimization problem



1. Try to establish a relation between **OPT**(**A**) and **OPT**(**A'**);
2. Try to prove the relation formally by
 - ❑ deriving a (not necessary optimal) solution of **A** from **OPT**(**A'**)
 - ❑ deriving a (not necessary optimal) solution of **A'** from **OPT**(**A**)
3. If you succeed, this would give you an algorithm.

Fractional Knapsack

Input:

$(w_1, v_1), (w_2, v_2), \dots, (w_n, v_n)$ and W

Output:

Weights x_1, \dots, x_n that maximize $\sum_i v_i \cdot \frac{x_i}{w_i}$ subject to:

$$\sum_i x_i \leq W \text{ and } 0 \leq x_j \leq w_j \text{ for all } j \in [n].$$

Question 5

Prove the following: Let j^* be the item with the maximum value/kg, v_j/w_j . Then, there exists an optimal knapsack containing $\min(w_{j^*}, W)$ kgs of item j^* .

Use the above as greedy strategy to design an algorithm for the fractional knapsack problem.

Acknowledgement

- The slides are modified from
 - The slides from Prof. Kevin Wayne
 - The slides from Prof. Surender Baswana
 - The slides from Prof. Erik D. Demaine and Prof. Charles E. Leiserson
 - The slides from Prof. Arnab Bhattacharya and Prof. Wing-Kin Sung