



CS3243: Introduction to Artificial Intelligence

Semester 2, 2019/2020

Midterm – CS3243 2020

- Closed book examination;
- we provide a “cheat sheet” in the test body.
- **One** A4 size sheet with your own notes, written in reasonable font + NUS approved calculator.

Venue : To be Announced (several)

Time: 15:00 – 17:00 March 7th (Saturday) 2020.

Previously...

- Minimax algorithm
 - What if MIN plays sub-optimally and MAX plays MINIMAX strategy?
 - Generalize to non-zero-sum, multi-player game?
See section 5.2.2 of AIMA 3rd Ed.
- α - β pruning algorithm
 - Prune subtrees that won't affect minimax decision
- H -Minimax algorithm
 - Use heuristic functions and impose depth limit
- Expected Minimax algorithm
 - Solve stochastic games



Constraint Satisfaction Problems

AIMA Chapter 6.1 – 6.4

Outline

- Constraint satisfaction problems (CSPs)
- Backtracking search for CSPs
- Local consistency in constraint propagation
- Local search for CSPs

CSP: a Broad Modelling Framework

Linear programming

- Variables are all rational values
- Linear constraints

Combinatorial optimization

- Graph problems: vertex cover, bipartite matching, minimum spanning tree...
- Operations Research: scheduling, bin packing, planning, task allocation
- Game theory and economics: social choice, market pricing, goods allocation...


Effective CSP Solvers

- Variable dependencies reduce search space significantly
- Solve very large problem instances very quickly

Constraint Satisfaction Problems (CSPs)


- Standard search problem:
 - **state** is atomic – any data structure that supports transition function, heuristic function, and goal test
- CSP
 - Variables $\vec{X} = X_1, \dots, X_n$; each variable X_i has a domain D_i
 - Constraints \vec{C} : what variable combinations are allowed?
 - Constraint scope: which variables are involved
 - Constraint relation: what is the relation between them.
 - Objective: find a legal assignment (y_1, \dots, y_n) of values to variables
 - $y_i \in D_i$ for all $i \in [n]$
 - Constraints are all satisfied.

Example: Graph Coloring



Variables:	$\vec{X} = \langle WA, NT, Q, NSW, V, SA, T \rangle$
Domains:	$D_i = \{R, G, B\}$
Constraints:	If $(X_i, X_j) \in E$ then $color(X_i) \neq color(X_j)$


Example: Graph Coloring



Find a **complete, consistent** assignment

- **Complete:** all variables are set
- **Consistent:** no constraint is violated

Constraint Graph



- **Binary CSP:** each constraint relates two variables
- **Constraint graph:** nodes are variables, links are constraints

Example: Job-Shop Scheduling

- Car assembly consists of 15 tasks
- Variables: $Axle_F, Axle_B, Wheel_{LF}, Wheel_{RF}, Wheel_{LB}, Wheel_{RB}, Nuts_{LF}, Nuts_{RF}, Nuts_{LB}, Nuts_{RB}, Cap_{LF}, Cap_{RF}, Cap_{LB}, Cap_{RB}, Inspect$
- Domain: $D_i = \{1, 2, \dots, 27\}$
- **Precedence** constraints
 - e.g., $Axle_F + 10 \leq Wheel_{RF}$
- **Disjunctive** constraint
 - e.g., $(Axle_F + 10 \leq Axle_B) \text{ or } (Axle_B + 10 \leq Axle_F)$

CSP Variants

Discrete variables

- Finite domains:
 - n variables, domain size $d \rightarrow \mathcal{O}(d^n)$ complete assignments
 - e.g., 8-queens
- Infinite domains:
 - integers, strings, etc.
 - e.g., job-shop scheduling: variables are start/end days for each job
 - need a constraint language, e.g., $StartJob_3 + 5 \leq StartJob_4$

Continuous variables

- e.g., start/end times for Hubble Space Telescope observations
- Linear programming problems (i.e., with linear constraints) can be solved in polynomial time

Constraint Variants

Unary constraints

- Single variable: e.g., $SA \neq G$

Binary constraints


- e.g., $SA \neq WA$

Global (i.e., higher-order) constraints

- 3 or more variables e.g. $X_1 + X_2 - 4X_7 \leq 15$

Cryptarithmic Puzzles

$$\begin{array}{r} TWO \\ +TWO \\ \hline FO\ UR \end{array}$$



Variables:	$\vec{X} = \langle F, T, U, W, R, O, C_1, C_2, C_3 \rangle$
Domains:	$D_i = \{0, \dots, 9\}$
Constraints:	$AllDiff(F, T, U, W, R, O)$ $O + O = R + 10C_1$ $C_1 + W + W = U + 10C_2$ $C_2 + T + T = O + 10C_3$ $C_3 = F$ $T, F \neq 0$

Sudoku

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

(a)

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

(b)

Variables:	$A_1, \dots, A_9, \dots, I_1, \dots, I_9$ (81 variables)
Domains:	$D_i = \{1, \dots, 9\}$
Constraints:	$AllDiff(\dots) \times 27$ (9 columns, 9 rows, 9 boxes) e.g. $AllDiff(A_1, \dots, A_3; B_1, \dots, B_3; C_1, \dots, C_3)$ is the constraint for the top-left box.

Standard Search Formulation (Incremental)


- Each state is a partial variable assignment
 - **Initial state**: the empty assignment []
 - **Transition function**: assign a **valid** value to an unassigned variable → fail if no **valid** assignments
 - **Goal test**: all variables have been assigned.
1. Uniform model for all CSPs; not domain-specific.
 2. Every solution appears at depth n (n variables assigned)
 3. Search path is irrelevant, can also use complete-state formulation.

What's the best search technique?

CSP Search Tree Size

CSP

- Variables: A, B, C, D
- Domains: $\{1, 2, 3\}$
- Assume no constraints



b at depth 0 = 4 variables \times 3 values = 12 states

b at depth 1 = 3 variables \times 3 values = 9 states

b at depth 2 = 2 variables \times 3 values = 6 states

b at depth 3 = 1 variable \times 3 values = 3 states

At depth ℓ : $(n - \ell)d$

Total number of leaves = $n! d^n$

Backtracking Search

Order of variable assignment is irrelevant

- $[WA = R \text{ then } NT = G]$ equivalent to $[NT = G \text{ then } WA = R]$

At every level, consider assignments to a single variable.

- Fix an order of variable assignment $\rightarrow b = d$ and there are d^n leaves

DFS for CSPs with single-variable/level assignments is called backtracking search

- Is the basic uninformed search algorithm for CSPs
- Backtracks when no legal assignments
- Can solve n -queens for $n \approx 25$


Backtracking Search

function BACKTRACKING-SEARCH(*csp*) **returns** a solution, or failure
return BACKTRACK($\{ \}$, *csp*)

function BACKTRACK(*assignment*, *csp*) **returns** a solution, or failure
if *assignment* is complete **then return** *assignment*
var \leftarrow SELECT-UNASSIGNED-VARIABLE(*csp*)
for each *value* **in** ORDER-DOMAIN-VALUES(*var*, *assignment*, *csp*) **do**
 if *value* is consistent with *assignment* **then**
 add $\{var = value\}$ to *assignment*
 inferences \leftarrow INFERENCE(*csp*, *var*, *value*)
 if *inferences* \neq failure **then**
 add *inferences* to *assignment*
 result \leftarrow BACKTRACK(*assignment*, *csp*)
 if *result* \neq failure **then**
 return *result*
 remove $\{var = value\}$ and *inferences* from *assignment*
return failure

Backtracking Example

Assign variables in the order of *WA*, *NT*, *Q*, ...




Improving Backtracking Efficiency

- **General-purpose** heuristics can give huge time gains:
 - SELECT-UNASSIGNED-VARIABLE: Which variable should be assigned next?
 - ORDER-DOMAIN-VALUES: In what order should its values be tried?
 - INFERENCE: How can domain reductions on unassigned variables be inferred?
 - Can we detect inevitable failure early?

Most Constrained Variable


- Most constrained variable:
 - choose the variable with fewest legal values



- a.k.a. **minimum-remaining-values (MRV)** heuristic


Most Constraining Variable

- Most constraining variable:
 - choose the variable with most constraints on remaining unassigned variables
- Tie-breaker among most constrained variables
- a.k.a. **degree** heuristic



Least Constraining Value

- Given a variable, choose the least constraining value:
 - the one that rules out the fewest values for neighboring unassigned variables





Inference in CSPs

Forward Checking and Arc Consistency

Forward Checking

Idea:

- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values




WA	NT	Q	NSW	V	SA	T
  	  	  	  	  	  	  

Forward Checking

Idea:


- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Forward Checking

Idea:

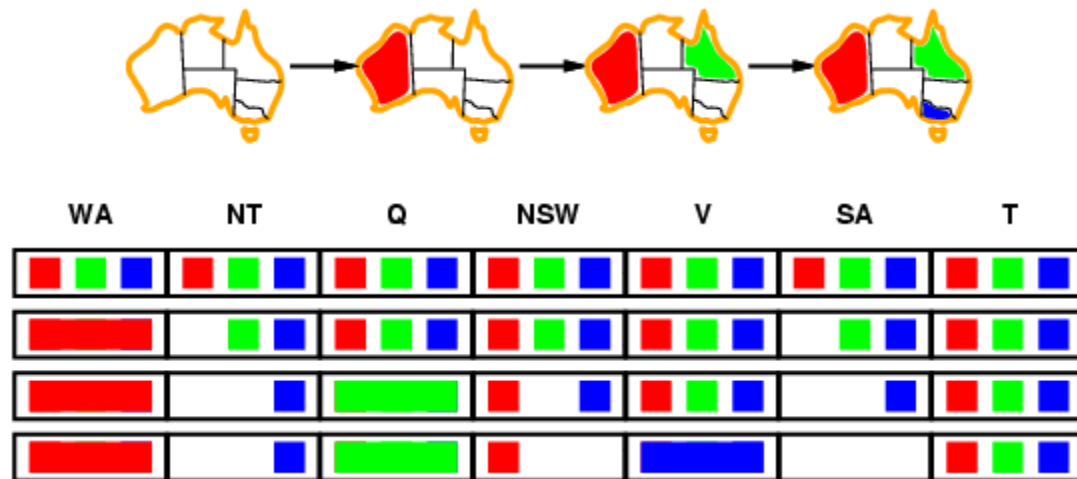
- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Forward Checking

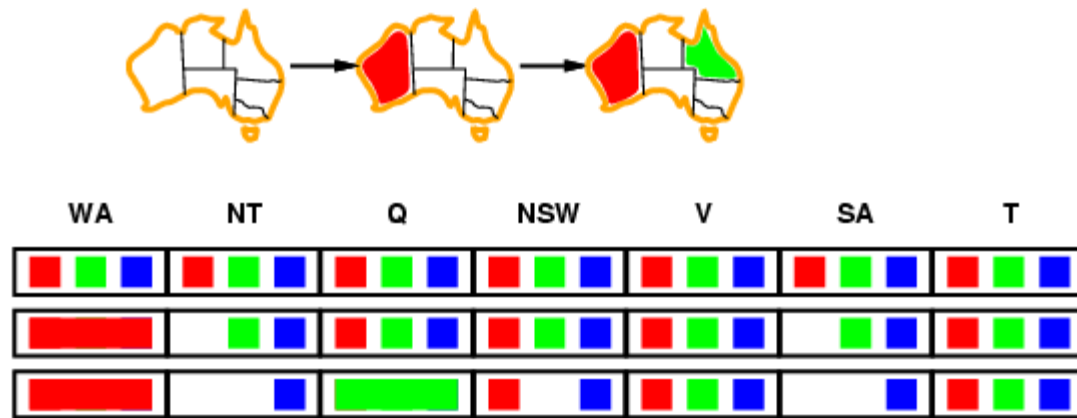
Idea:

- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Constraint Propagation

- Problem: Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:




- NT and SA cannot both be blue!
- Solution: **Constraint propagation** repeatedly locally enforces constraints.

Inference in CSPs


Try to infer illegal values for variables by performing constraint propagation

- Node consistency for unary constraints
- Arc consistency for binary constraints
- ...
- Can be interleaved with search, or run as a preprocessing step



Arc Consistency

X_i is arc-consistent wrt X_j (the arc (X_i, X_j) is consistent) iff for every value $x \in D_i$ there exists some value $y \in D_j$ that satisfies the binary constraint on the arc (X_i, X_j)




More on Arc Consistency

- If $X \leftarrow x$ makes a constraint impossible to satisfy, remove x from consideration!
- A value is impossible: the constraint provides ***no support*** for the value.
- E.g. $D_i = \{1,4,5\}$, $D_j = \{1,2,3\}$ and we have $X_i > X_j$; then $X_i \leftarrow 1$ is impossible.

Arc Consistency


- Simplest form of propagation makes each arc **consistent**
- X_i is arc-consistent wrt X_j (the arc (X_i, X_j) is consistent) iff for every value $x \in D_i$ there exists some value $y \in D_j$ that satisfies the binary constraint on the arc (X_i, X_j)



- Arcs are directed, a binary constraint becomes two arcs
- Arc $NSW \rightarrow SA$ is originally not consistent, but is consistent after deleting $NSW = \text{blue}$

Arc Consistency

- Simplest form of propagation makes each arc **consistent**
- X_i is arc-consistent wrt X_j (the arc (X_i, X_j) is consistent) iff for every value $x \in D_i$ there exists some value $y \in D_j$ that satisfies the binary constraint on the arc (X_i, X_j)



- If X_i loses a value, neighbors of X_i need to be (re)checked
- Arc $V \rightarrow NSW$ is originally not consistent, but is consistent after deleting $V = \text{red}$

Arc Consistency Propagation

- Reducing D_i may result in domain reductions for others.
- E.g. $D_1 = \{1, 4, 5\}$, $D_2 = \{1, 2, 3\}$, $D_3 = \{2, 3, 4, 5\}$
 - Constraints: $X_3 > X_1$, $X_1 > X_2$
 - Remove 1 from D_1 so $D_1 = \{4, 5\}$
 - No support for $X_3 = 2, 3, 4$
 - We can remove those values: $D_3 = \{5\} \Rightarrow X_3 = 5$
 - Before applying AC to arc (X_1, X_2) , we cannot reduce D_3
- A chain reaction of domain reductions.


Sudoku Chain Reaction

		3		2		6		
9			3		5			1
		1	8		6	4		
		8	1		2	9		
								8
		6	7		8	2		
		2	6		9	5		
8			2		3			9
		5		1		3		

- *Alldiff* constraint on middle box makes domain of red square $\{3, 4, 5, 6, 9\}$. Column constraint reduces domain to $\{4\}$.
- Consider orange square. Original column and box constraints yield domain of $\{4, 7\}$. Red square forces $\{7\}$.
- Blue box must be $\{1\}$ as column already has eight values.

Arc Consistency

- Simplest form of propagation makes each arc **consistent**
- X_i is arc-consistent wrt X_j (the arc (X_i, X_j) is consistent) iff for every value $x \in D_i$ there exists some value $y \in D_j$ that satisfies the binary constraint on the arc (X_i, X_j)



- Arc consistency propagation detects failure earlier than forward checking

Can be run as a preprocessing step or after each assignment

Arc Consistency Algorithm AC-3

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X , D , C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REVISE(*csp*, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to *queue*

return true

function REVISE(*csp*, X_i , X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow false

for each x **in** D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

revised \leftarrow true

return *revised*

Time complexity: $\mathcal{O}(n^2 d^3)$

Time Complexity of AC-3

- CSP has at most n^2 directed arcs
- Each arc (X_i, X_j) can be inserted at most d times because X_i has at most d values to delete.
- REVISE: Checking the consistency of an arc takes $\mathcal{O}(d^2)$ time
- $\mathcal{O}(n^2 \times d \times d^2) = \mathcal{O}(n^2 d^3)$

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X, D, C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REVISE(*csp*, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to *queue*

return true

function REVISE(*csp*, X_i , X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow false

for each x **in** D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

revised \leftarrow true

return *revised*

Maintaining AC (MAC)

- Like any other propagation, we can use AC in search
- i.e., search proceeds as follows:
 - establish AC at the root
 - when AC-3 terminates, choose a new variable/value
 - re-establish AC given the new variable choice (i.e. maintain AC)
 - repeat;
 - backtrack if AC gives **empty domain**
- Exception: initially, insert into queue only arcs of neighboring unassigned variables/nodes
- The hard part of implementation is undoing effects of AC

Special Types of Consistency

- Some types of constraints lend themselves to special types of arc-consistency
- Consider the *AllDiff* constraint
 - the named variables must all take different values
 - not a binary constraint
 - can be expressed as $\frac{n(n-1)}{2}$ not-equal binary constraints
- We can apply, e.g., AC-3 as usual
- But there is a much better option

Generalized Arc Consistency and k -Consistency

- Suppose $D_1 = D_2 = \{2,3\}$ and $D_3 = \{1,2,3\}$.
- The constraints $X_i \neq X_j$ are arc-consistent
 - e.g., $X_2 = 2$ supports the value $X_3 = 3$
- The single ternary constraint $Alldiff(X_1, X_2, X_3)$ is not!
 - We must reduce $D_3 = \{1\}$.
- A special-purpose algorithm exists for $Alldiff$ constraint to establish generalized arc consistency in efficient time
 - Special-purpose propagation algorithms are vital
- Arc Consistency (2-consistency) can be extended to k -consistency

Local Search for CSPs

- Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - allow states that violate constraints
 - operators **reassign** variable values
- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts** heuristic:
 - choose value that violates the fewest constraints
 - i.e., hill-climb with $h(n)$ = total # of violated constraints

Min-Conflicts

function MIN-CONFLICTS(*csp*, *max_steps*) **returns** a solution or failure

inputs: *csp*, a constraint satisfaction problem

max_steps, the number of steps allowed before giving up

current \leftarrow an initial complete assignment for *csp*

for *i* = 1 to *max_steps* **do**

if *current* is a solution for *csp* **then return** *current*

var \leftarrow a randomly chosen conflicted variable from *csp*.VARIABLES

value \leftarrow the value *v* for *var* that minimizes CONFLICTS(*var*, *v*, *current*, *csp*)


 set *var* = *value* in *current*

return *failure*

Figure 6.8 The MIN-CONFLICTS algorithm for solving CSPs by local search. The initial state may be chosen randomly or by a greedy assignment process that chooses a minimal-conflict value for each variable in turn. The CONFLICTS function counts the number of constraints violated by a particular value, given the rest of the current assignment.

Example: 4-Queens

- **States:** 4 queens in 4 columns ($4^4 = 256$ states)
- **Actions:** move queen in column
- **Goal test:** no attacks
- **Evaluation:** $h(n)$ = number of attacks



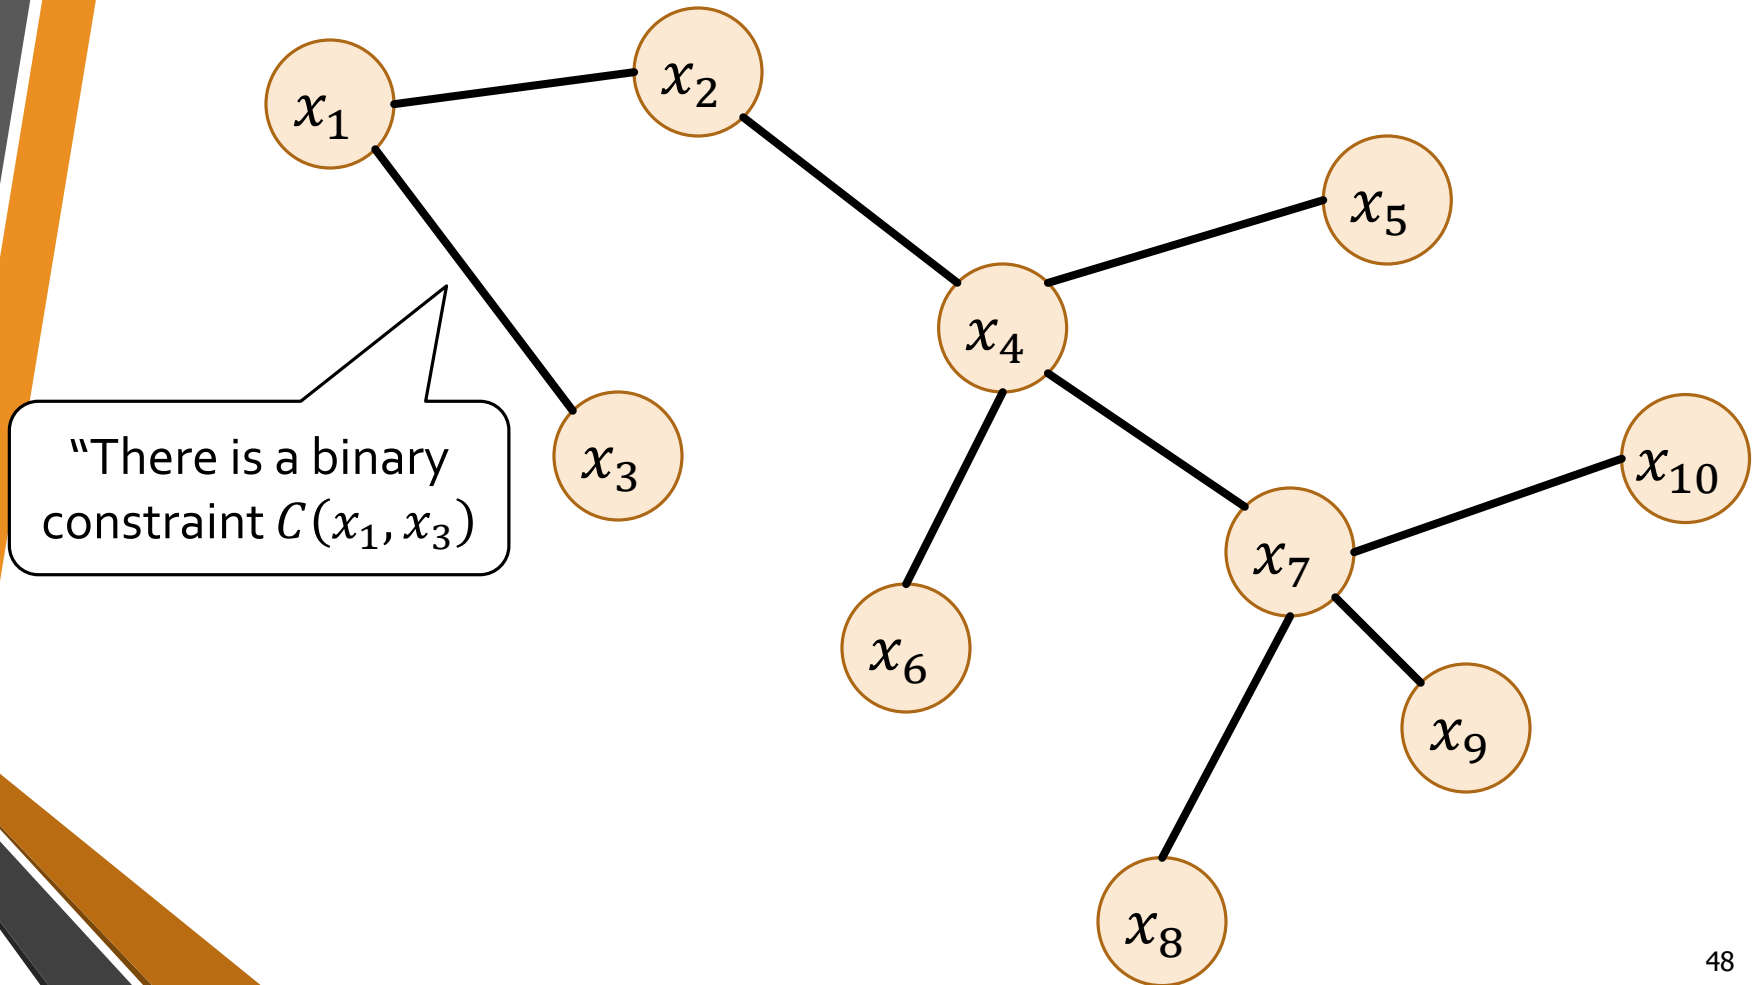
- Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10^6$)



Constraint Satisfaction Problems – Structured CSPs

AIMA Chapter 6.5


CSPs With Binary Constraints as Graphs



CSPs With Binary Constraints as Graphs


Theorem: if CSP constraint graph is a **tree**,
then we can compute a satisfying
assignment (or decide that one does not
exist) in $\mathcal{O}(nd^2)$ time

No cycles!



CSPs With Binary Constraints as Graphs

Theorem: if CSP constraint graph is a **tree**, then we can compute a satisfying assignment (or decide that one does not exist) in $\mathcal{O}(nd^2)$ time



Step 1: make parents arc-consistent with children


Step 2: if step 1 successful, assign values from root down.

CSPs With Binary Constraints as Graphs

Theorem: if CSP constraint graph is a **tree**, then we can compute a satisfying assignment (or decide that one does not exist) in $\mathcal{O}(nd^2)$ time

Step 1: make parents arc-consistent with children

Step 2: if step 1 successful, assign values from root down.



CSPs With Binary Constraints as Graphs

Theorem: if CSP constraint graph is a **tree**, then we can compute a satisfying assignment (or decide that one does not exist) in $\mathcal{O}(nd^2)$ time


Topological sort takes $\mathcal{O}(n)$ time on tree.

Step 1: make parents arc-consistent with children

n steps taking $\mathcal{O}(d^2)$ time each

Step 2: if step 1 successful, assign values from root down.

n steps taking $\mathcal{O}(d)$ time each



CSPs With Binary Constraints as Graphs

Theorem: if CSP constraint graph is a **tree**, then we can compute a satisfying assignment (or decide that one does not exist) in $\mathcal{O}(nd^2)$ time

Could it be that this step **removes options** for a **valid** CSP assignment? That would be very bad...

Step 1: make parents arc-consistent with children

Step 2: if step 1 successful, assign values from root down.

Because constraint graph is a tree, no need to backtrack when assigning! Assignment for node affects only its subtree.

