



CS3243: Introduction to Artificial Intelligence

Semester 2, 2019/2020



Reinforcement Learning

AIMA Chapter 21

Based in part on slides from Zemel, Urtason and Fidler (2016), as well as Li, Johnson and Yeung (2017)

Outline

- Introduction to Learning Agents
- Reinforcement Learning Formulation
- Agent policy and optimal policies
- Learning an optimal policy
- Q-learning

Supervised Learning

Agent



=dog



=dog



=cat



=cat



=cat



=dog




=?

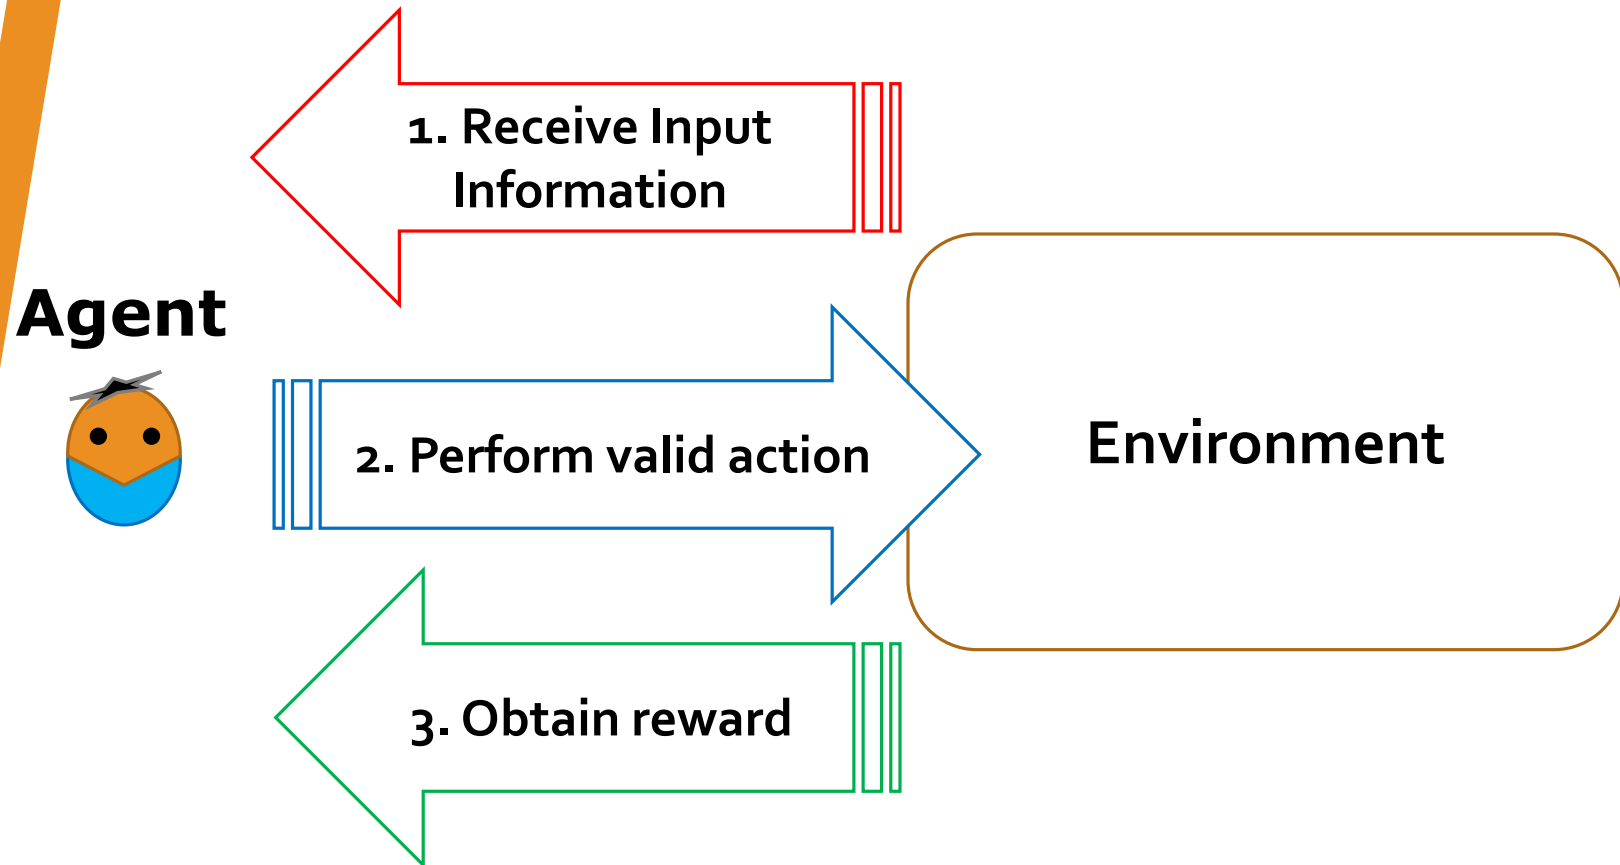
Unsupervised Learning



Agent



Reinforcement Learning



Reinforcement Learning

Agent



1. Receive Input Information

2. Perform valid action

3. Obtain reward

At time t :


State: s_t

State: s_{t+1}

Action: a_t
(state dependent)


Reward:
 r_t

Reinforcement Learning



Markov property: next state is determined only based on current action and state, not entire sequence

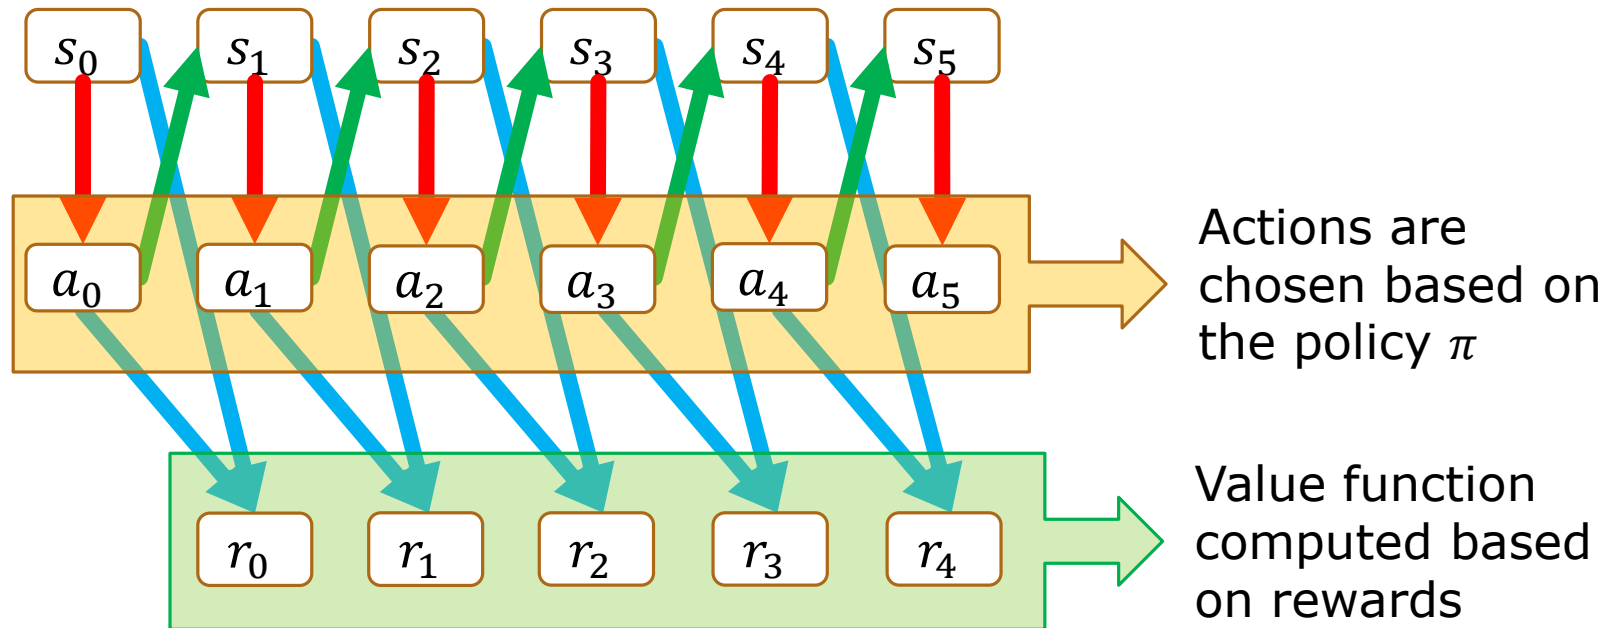
Policy



Policy π determines agent's behavior, i.e. actions

- Deterministic policy: $a_t = \pi(s_t)$
- Stochastic policy: $\pi(a \mid s_t) = \Pr[a_t = a \mid s_t]$

Value Function



Value function tries to predict how good a state is, given the rewards.

$$\begin{aligned} V^\pi(s_t) &= r_t(a_t, s_t) + \gamma r_{t+1}(a_{t+1}, s_{t+1}) + \gamma^2 r_{t+2}(a_{t+2}, s_{t+2}) + \dots \\ &= \sum_{\ell=0}^{\infty} \gamma^\ell r_{t+\ell}(a_{t+\ell}, s_{t+\ell}) \end{aligned}$$

Value Function

Value function tries to predict how good a state is, given the rewards.

$$\begin{aligned} V^\pi(s_t) &= r_t(a_t, s_t) + \gamma r_{t+1}(a_{t+1}, s_{t+1}) + \gamma^2 r_{t+2}(a_{t+2}, s_{t+2}) + \dots \\ &= \sum_{\ell=0}^{\infty} \gamma^\ell r_{t+\ell}(a_{t+\ell}, s_{t+\ell}) \end{aligned}$$

The value γ (a value between 0 and 1) is called the **discount rate**.

- High value of γ – long-sighted agent, cares about future rewards.
- Low value of γ – agent is greedy, who cares about the future?

Value Function – The Challenge

We want to choose a **value maximizing policy**.

However, we are missing two key bits of information:

The rewards of unobserved states.

What the next state will be when we take an action.

The challenge:

Infer rewards/state transitions as we go along

...while maximizing revenue.

State:

(x, y) position

Actions:

Up/Down/Left/Right

Reward:


-1 per time step
(+10 for reaching goal)

Policy:

Direction to go from
each position (can be
randomized).

Value Function:



Total reward of policy
execution from given
state.



State:

(1,1) position



Actions:



State:

(2,1) position



Actions:



State:

(2,2) position



Actions:



State:

(2,3) position



Actions:



State:

(3,3) position



Actions:



State:

(4,3) position



Actions:



State:

(5,3) position



Actions:



State:

(5,4) position



Actions:



State:

(5,5) position


Actions:



State:

(5,6) position

Actions:



spend as little
time in the
maze as
possible, get to
the goal.

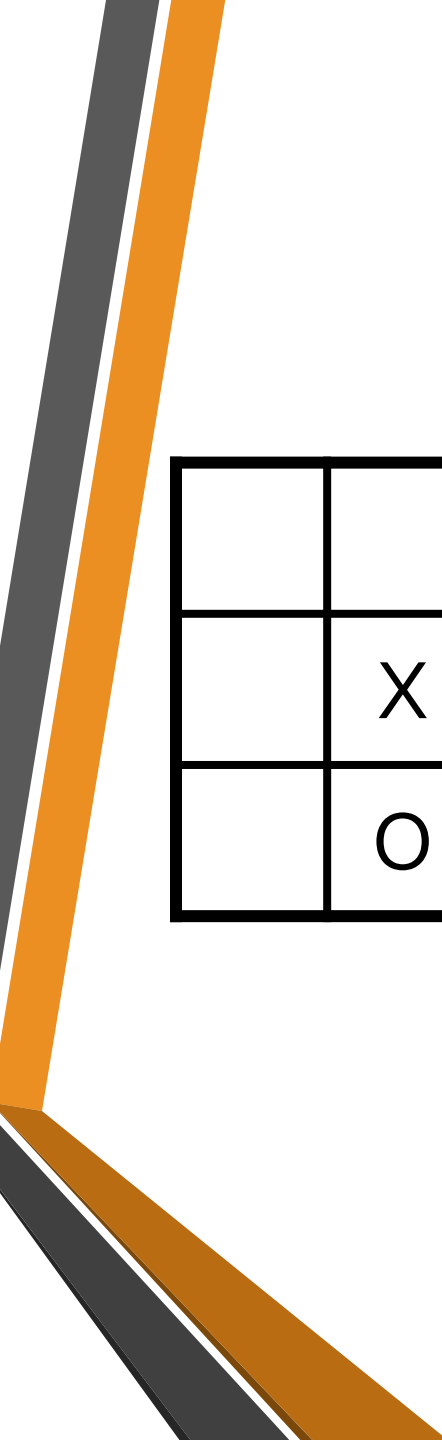


Value Function:

Discounted reward
if starting from this
state.

Numbers shown
for discount factor
 $\gamma = 1$

-8					-9	-8	-7	-8	-9
-7	-6	-5	-6				-6		
-8		-4			-3	-4	-5	-6	
		-3			-2			-7	
	-3	-2	-1	0	-1			-8	
				1					
	-1	0	1	2	1	0	-1		7
	0			3					8
	1	2	3	4	5	6	7	8	9
				3			6		10



		X
	X	
	O	O

Reward:

−1,0,+1 lose/tie/win (seen only on final move)

State:

Current positions of X's and O's on board

Policy:

What moves to make in given position?

Value function:

predict future reward given state.

Win Prob.

		X
	X	
	O	O

s_t

X		X
	X	
	O	O

= 0.5

	X	X
	X	
	O	O

= 0.5

		X
X	X	
	O	O

= 0.5

		X
	X	X
	O	O

= 0.5

		X
	X	
X	O	O

= 1

s_{t+1}

- All values are 0/0.5/1 initially
- At each turn choose move with highest win prob.
- Update table entries based on the game outcome.
- Value function will eventually represent true win probabilities

Alternatively:

- pick with probability proportional to win prob
- make a random choice.

Update strategy is critical:

- Necessary for convergence to optimal strategy.
- Some will work better than others.

Markov Decision Problem

Completely specified by a distribution:

$$\Pr[s_{t+1} = s; r_{t+1} = r \mid s_t, a_t]$$

"What is the next state and reward given current state and action?"

Planning: given an MDP, compute optimal policy

Learning: don't know the MDP, learn a strategy.

Markov Decision Problem

Given complete knowledge of MDP:

$$\Pr[s_{t+1} = s; r_{t+1} = r \mid s_t, a_t]$$

The optimal policy is deterministic - select optimal action in each state.

But... agent doesn't know the underlying MDP.

Needs to perform trial-and-error, interact with environment.

... and not lose too much reward along the way.

Learning Optimal Policies

We want to maximize (discounted) revenue.

Note that:

$$V^\pi(s_t) = r_{t+1} + \gamma V^\pi(s_{t+1})$$

Value
now

Reward
now

Value
later

Optimal policy: $\pi(s)$ is an action in

$$\operatorname{argmax}_a \{r(s, a) + \gamma V(\delta(s, a))\}$$

Reward
now

Value
later

$\delta(s, a)$: next
state given
current state
and action

Learning Optimal Policies

Optimal policy: $\pi^*(s)$ is an action in

$$\operatorname{argmax}_a \{r(s, a) + \gamma V^*(\delta(s, a))\}$$

Its value is:

$$V^*(s_t) = r_{t+1} + \gamma V^*(s_{t+1})$$

We could identify optimal policy $\pi^*(s)$ if we knew $r(s, a)$ and $\delta(s, a)$.

We don't. Cannot choose optimal actions.

Q-Learning

Define:

$$Q^{\pi}(s, a) = r(s, a) + \gamma V^{\pi}(\delta(s, a))$$

$Q^{\pi}(s, a)$ is the utility we obtain if we take action a at state s , and then follow the policy π from then on.

Define:

$$Q^*(s, a) = r(s, a) + \gamma V^*(\delta(s, a))$$

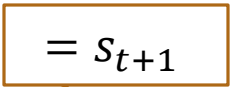
$Q^*(s, a)$ is the utility we obtain if we take action a at state s , and then follow the optimal policy from then on.

Q-Learning

Q^* and V^* are very similar:

$$V^*(s) = \max_a Q(s, a)$$

Therefore:

$$\begin{aligned} Q^*(s_t, a_t) &= r(s_t, a_t) + \gamma V^*(\delta(s_t, a_t)) \\ &= r(s_t, a_t) + \gamma \max_a Q(s_{t+1}, a) \end{aligned}$$


A diagram consisting of a rectangular box containing the text $= s_{t+1}$. An orange arrow points from the bottom-left corner of this box to the argument $\delta(s_t, a_t)$ in the equation above.

Q-Learning – Value Iteration Algorithm

1. Initialize $\hat{Q}(s, a) \leftarrow 0$ for all s, a .
2. Start at s_0
3. For $t = 0, \dots, \infty$:
 1. For every a : $\hat{Q}(s_t, a) \leftarrow r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a')$
 2. Pick action a_t maximizing $\hat{Q}(s_t, a)$
 3. Set $s_{t+1} \leftarrow \delta(s_t, a_t)$

Key observation: when $r(s, a) \geq 0$ and $\hat{Q} = 0$, $\hat{Q}(s, a) \leq Q^*(s, a)$ always.

In other words – we always underestimate the optimal Q values.

They always increase at every iteration, thus we converge to Q^* ... and in particular to an optimal policy!

Q-Learning – Value Iteration Algorithm

1. Initialize $\hat{Q}(s, a) \leftarrow 0$ for all s, a .
2. Start at s_0
3. For $t = 0, \dots, \infty$:
 1. For every a : $\hat{Q}(s_t, a) \leftarrow r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a')$
 2. Pick action a_t maximizing $\hat{Q}(s_t, a)$
 3. Set $s_{t+1} \leftarrow \delta(s_t, a_t)$

Problem:

We ignore current $\hat{Q}(s_t, a)$ value in computation.

$r(s, a)$ can be stochastic, as is $\delta(s, a)$.

One bad experience can result in bad underestimate of $Q^*(s, a)$.

Q-Learning – Value Iteration Algorithm

1. Initialize $\hat{Q}(s, a) \leftarrow 0$ for all s, a .
2. Start at s_0
3. For $t = 0, \dots, \infty$:
 1. For every a : $\hat{Q}(s_t, a) \leftarrow r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a')$
 2. Pick action a_t maximizing $\hat{Q}(s_t, a)$
 3. Set $s_{t+1} \leftarrow \delta(s_t, a_t)$

Solution:

We need to maintain value of \hat{Q}_i stable as we observe it more.

Change update rule:

$$\hat{Q}(s_t, a) \leftarrow (1 - \alpha_t) \hat{Q}(s_t, a) + \alpha_t \left(r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a') \right)$$

Where

$$\alpha_t = \frac{1}{1 + N[s_t, a]}$$

of times action a
taken at state s_t .

Q-Learning – Value Iteration Algorithm

1. Initialize $\hat{Q}(s, a) \leftarrow 0$ for all s, a .
2. Start at s_0
3. For $t = 0, \dots, \infty$:
 1. For every a : $\hat{Q}(s_t, a) \leftarrow r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a')$
 2. Pick action a_t maximizing $\hat{Q}(s_t, a)$
 3. Set $s_{t+1} \leftarrow \delta(s_t, a_t)$

Problem:

We greedily, deterministically, pick action a_t maximizing $\hat{Q}(s_t, a)$.

Deterministic algorithms can be 'fooled' by stochastic (or adversarial) inputs (more of this next lecture).

Q-Learning – Value Iteration Algorithm

1. Initialize $\hat{Q}(s, a) \leftarrow 0$ for all s, a .
2. Start at s_0
3. For $t = 0, \dots, \infty$:
 1. For every a : $\hat{Q}(s_t, a) \leftarrow r(s_t, a) + \gamma \max_{a'} \hat{Q}_i(\delta(s_t, a), a')$
 2. Pick action a_t maximizing $\hat{Q}(s_t, a)$
 3. Set $s_{t+1} \leftarrow \delta(s_t, a_t)$

Solution:

Pick action a_t **randomly**.

1. Totally randomly?
2. Randomly amongst current best actions?

3. Set $\Pr[\text{choosing action } a \mid s] = \frac{e^{\epsilon \hat{Q}(s, a)}}{\sum_{a'} e^{\epsilon \hat{Q}(s, a')}} \leftarrow \Pr[a \mid s] \sim e^{\epsilon \hat{Q}(s, a)}$