

# **CS4225/CS5425 Big Data Systems for Data Science**

## NoSQL Overview

Bryan Hooi  
School of Computing  
National University of Singapore  
[bhooi@comp.nus.edu.sg](mailto:bhooi@comp.nus.edu.sg)



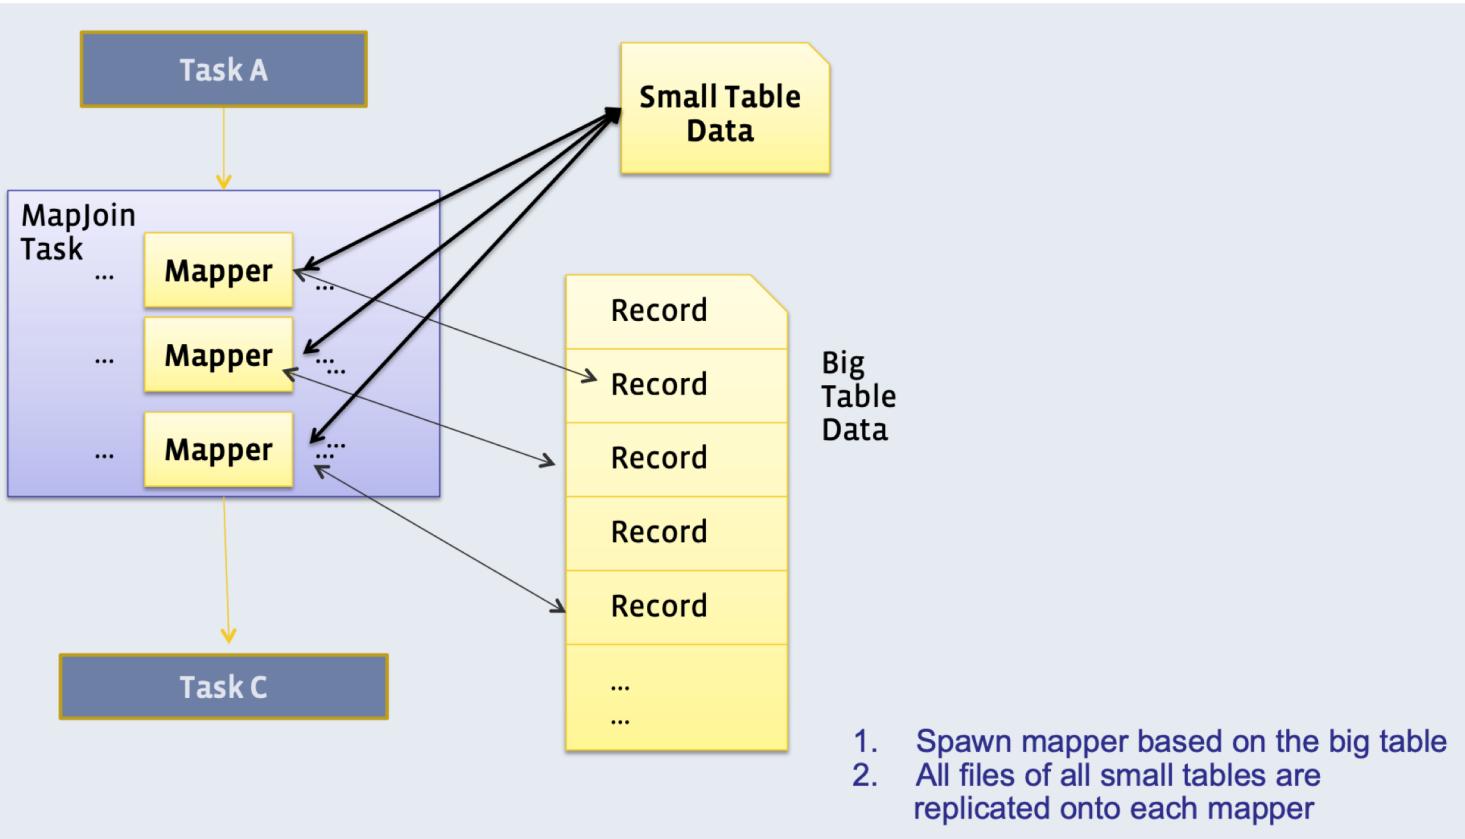
# Announcements

- Since we are quite ahead of schedule (compared to the syllabus / last semester), next week will be a shorter than usual lecture
- HWI will be released by this weekend (due 10 Oct 11.59pm)

| Week   | Date   | Topics                             | Tutorial                    | Due Dates                                       |
|--------|--------|------------------------------------|-----------------------------|---|
| 1      | 12 Aug | Overview and Introduction          |                             |   |
| 2      | 29 Aug | MapReduce - Introduction           |                             |   |
| 3      | 26 Aug | MapReduce and Relational Databases |                             |   |
| 4      | 2 Sep  | MapReduce and Data Mining          | Tutorial: Hadoop            | Assignment 1 released                           |
| 5      | 9 Sep  | NoSQL Overview 1                   |                             |   |
| 6      | 16 Sep | NoSQL Overview 2                   |                             |   |
| Recess |        |                                    |                             |   |
| 7      | 30 Sep | Apache Spark 1                     | Tutorial: NoSQL & Spark     | Assignment 1 due, Assignment 2 released (3 Oct) |
| 8      | 7 Oct  | Apache Spark 2                     |                             |   |
| 9      | 14 Oct | Large Graph Processing 1           | Tutorial: Graph Processing  |   |
| 10     | 21 Oct | Large Graph Processing 2           |                             |   |
| 11     | 28 Oct | Stream Processing                  | Tutorial: Stream Processing | Assignment 2 due (31 Oct)                       |
| 12     | 4 Nov  | Deepavali – No Class               |                             |   |
| 13     | 11 Nov | Test                               |                             |   |

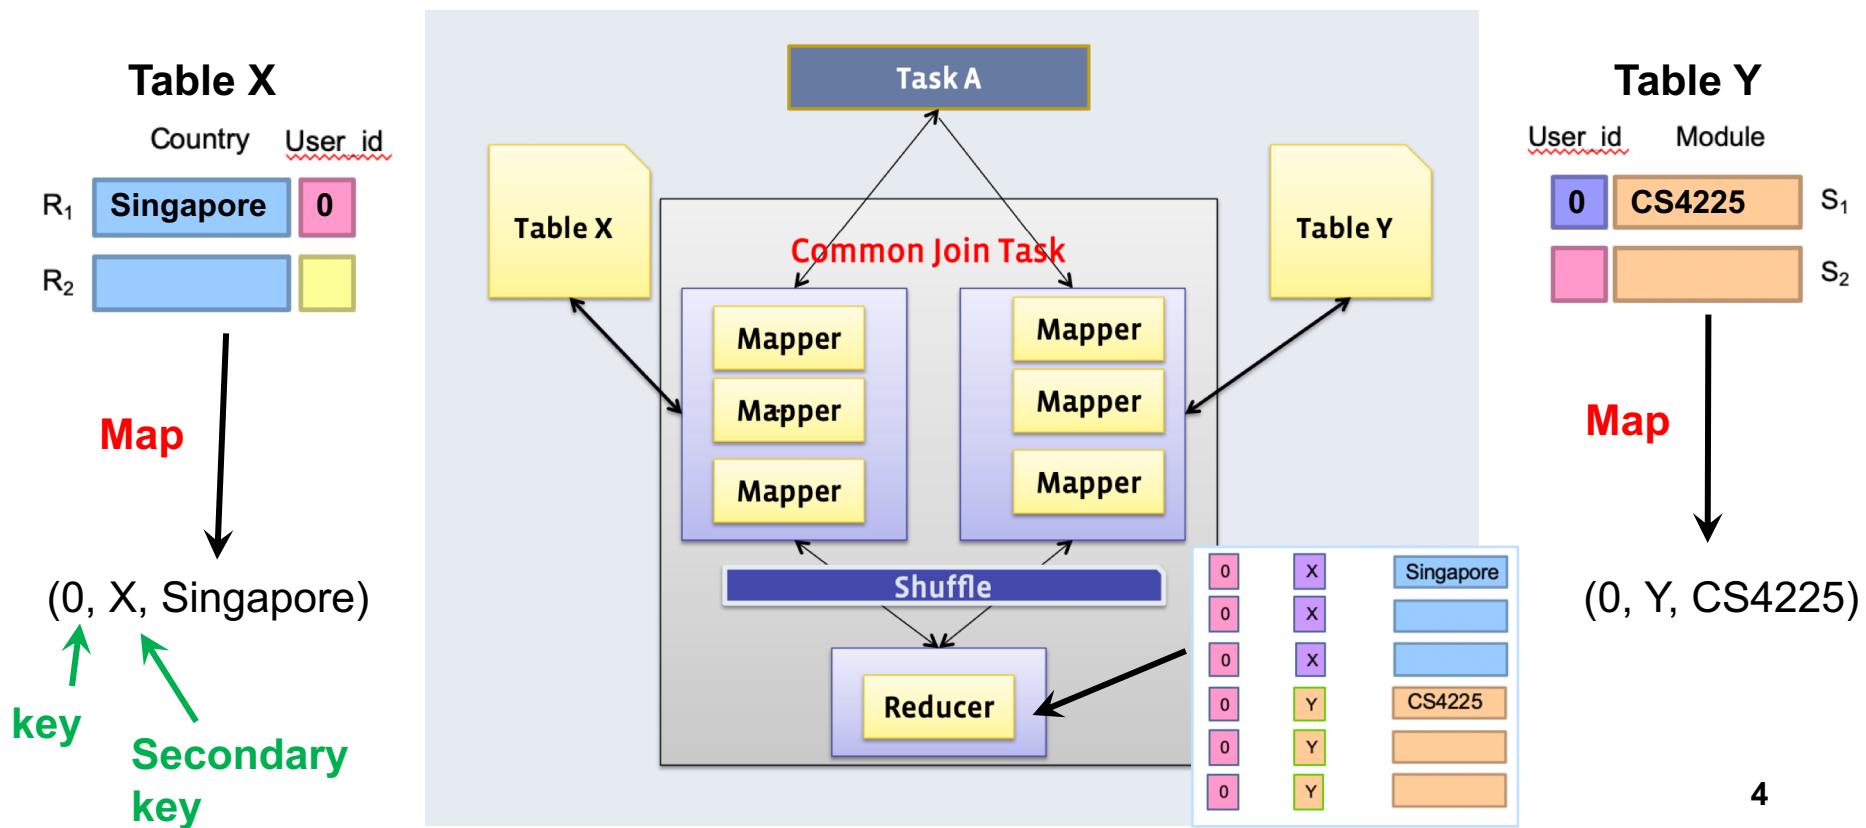
# Recap: Map-Side Join

- Requires one of the tables to fit in memory
  - All mappers store a copy of the small table
  - They iterate over the big table, and join the records with the small table



# Recap: Reduce-side (“Common”) Join

- Doesn't require a dataset to fit in memory, but slower than map-side join
  - Different mappers operate on each table, and emit records, with key as the variable to join by



# Recap: Jaccard similarity/distance

- **Jaccard Similarity**  
(between **sets** A and B)

$$s_{\text{Jaccard}}(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

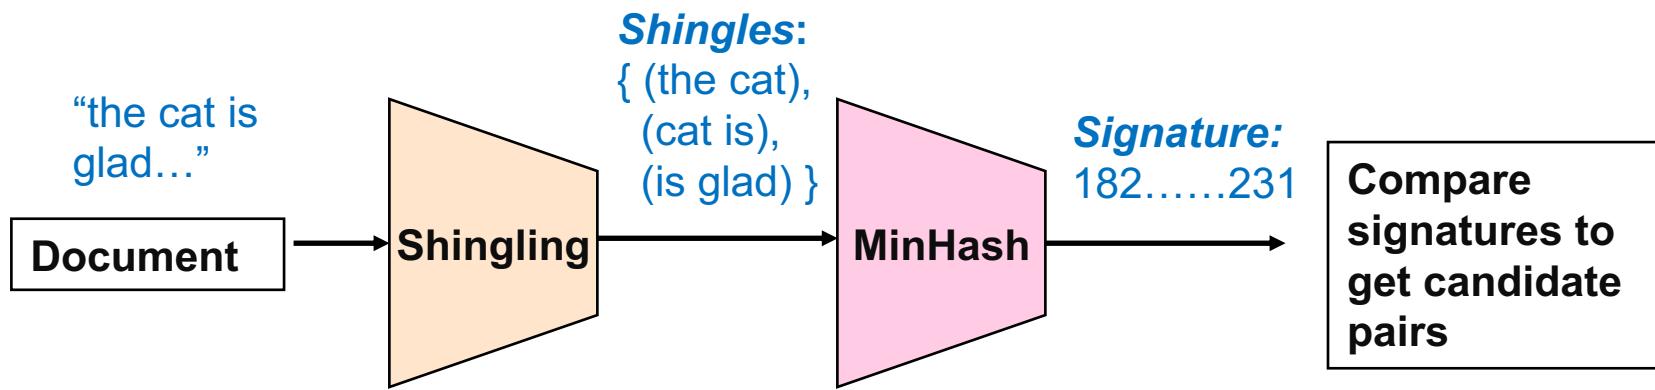
$$A = \{ \text{bread} , \text{milk} \} \quad B = \{ \text{cheese} , \text{milk} \}$$

$$s_{\text{Jaccard}} = \frac{\text{milk}}{\text{cheese}, \text{bread}, \text{milk}} = 1/3$$

- **Jaccard Distance**

$$d_{\text{Jaccard}}(A, B) = 1 - s_{\text{Jaccard}}(A, B)$$

# Recap: Near-Duplicate Document Search



# Recap: K-Means Algorithm

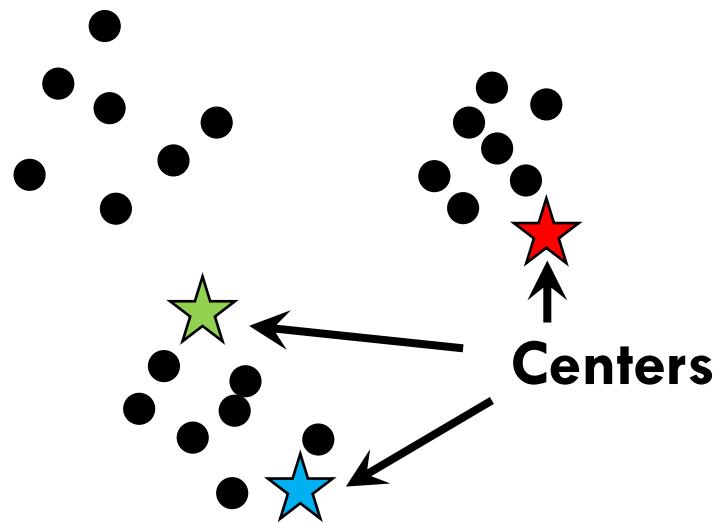
**1. Initialization:** Pick K random points as centers

**2. Repeat:**

a) **Assignment:** assign each point to nearest cluster

b) **Update:** move each cluster center to average of its assigned points

**Stop** if no assignments change



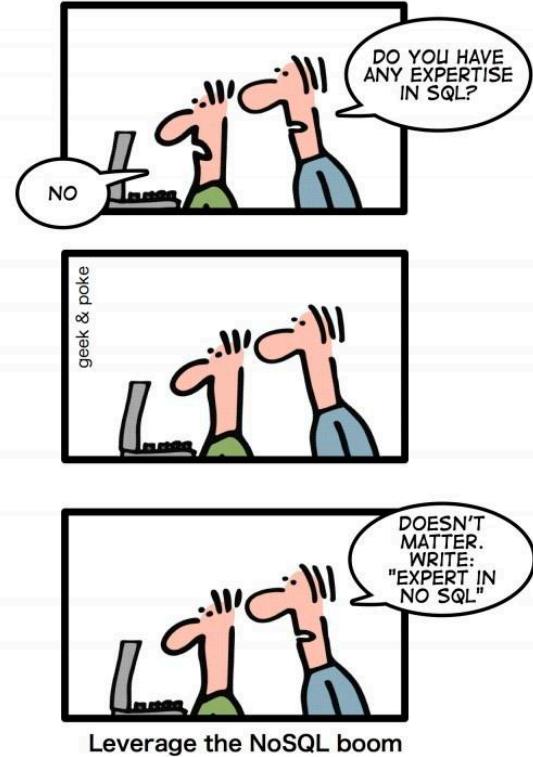
# Today's Plan

## I. What is NoSQL?

2. Major types of NoSQL systems
3. Key concepts

- We will see, but **no need to remember for test / HW:**
  - Historical details
  - Writing code in any NoSQL systems (e.g. MongoDB)
  - Details about any specific NoSQL systems

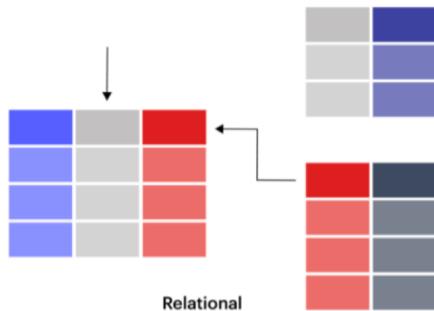
## HOW TO WRITE A CV



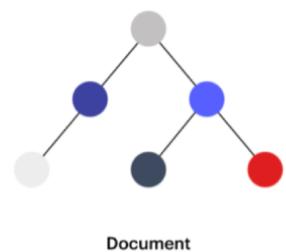
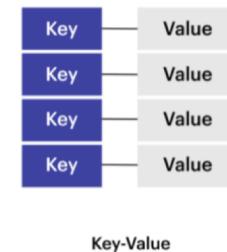
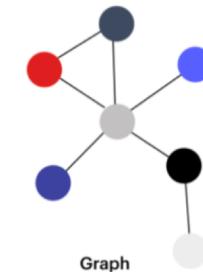
# What is NoSQL?

- NoSQL mainly refers to a **non-relational database**, i.e. it stores data in a format other than relational tables
- "SQL" here refers to traditional relational database management systems ("DBMS")
  - (This is a slight misnomer; here "SQL" is used to refer to relational databases, not the querying language)
- NoSQL has come to stand for "Not Only SQL", i.e. using relational and non-relational databases alongside one another, each for the tasks they are most suited for

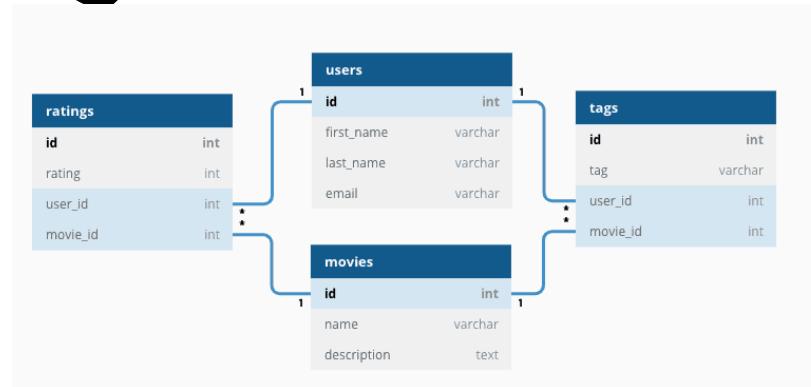
SQL database



NoSQL database



# What is NoSQL?



Traditional Relational Database Management System (DBMS)

```
>>> db.inventory.find()
[
  {
    _id: ObjectId("6138d8b1611d9dc433ad73dd"),
    item: 'canvas',
    qty: 100,
    tags: [ 'cotton' ],
    size: { h: 28, w: 35.5, uom: 'cm' }
  },
  {
    _id: ObjectId("6138d917611d9dc433ad73de"),
    item: 'table',
    qty: 5,
    price: 10
  }
]
```

Phone directory

| Key   | Value            |
|-------|------------------|
| Paul  | (091) 9786453778 |
| Greg  | (091) 9686154559 |
| Marco | (091) 9868564334 |



**mongoDB®** (Document Store)

**redis**

(Key-Value Store)

# Metaphor: NoSQL = No Rules?

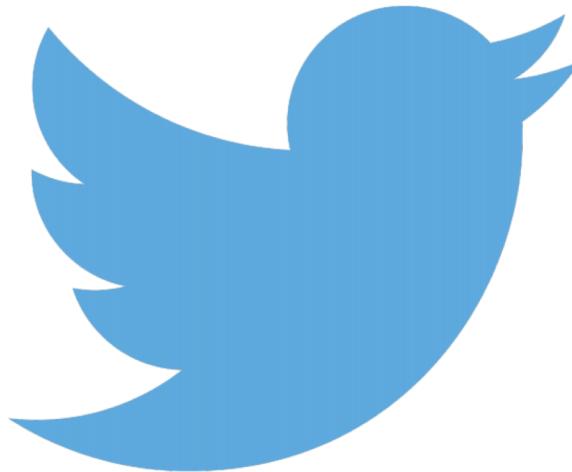


Relational Databases



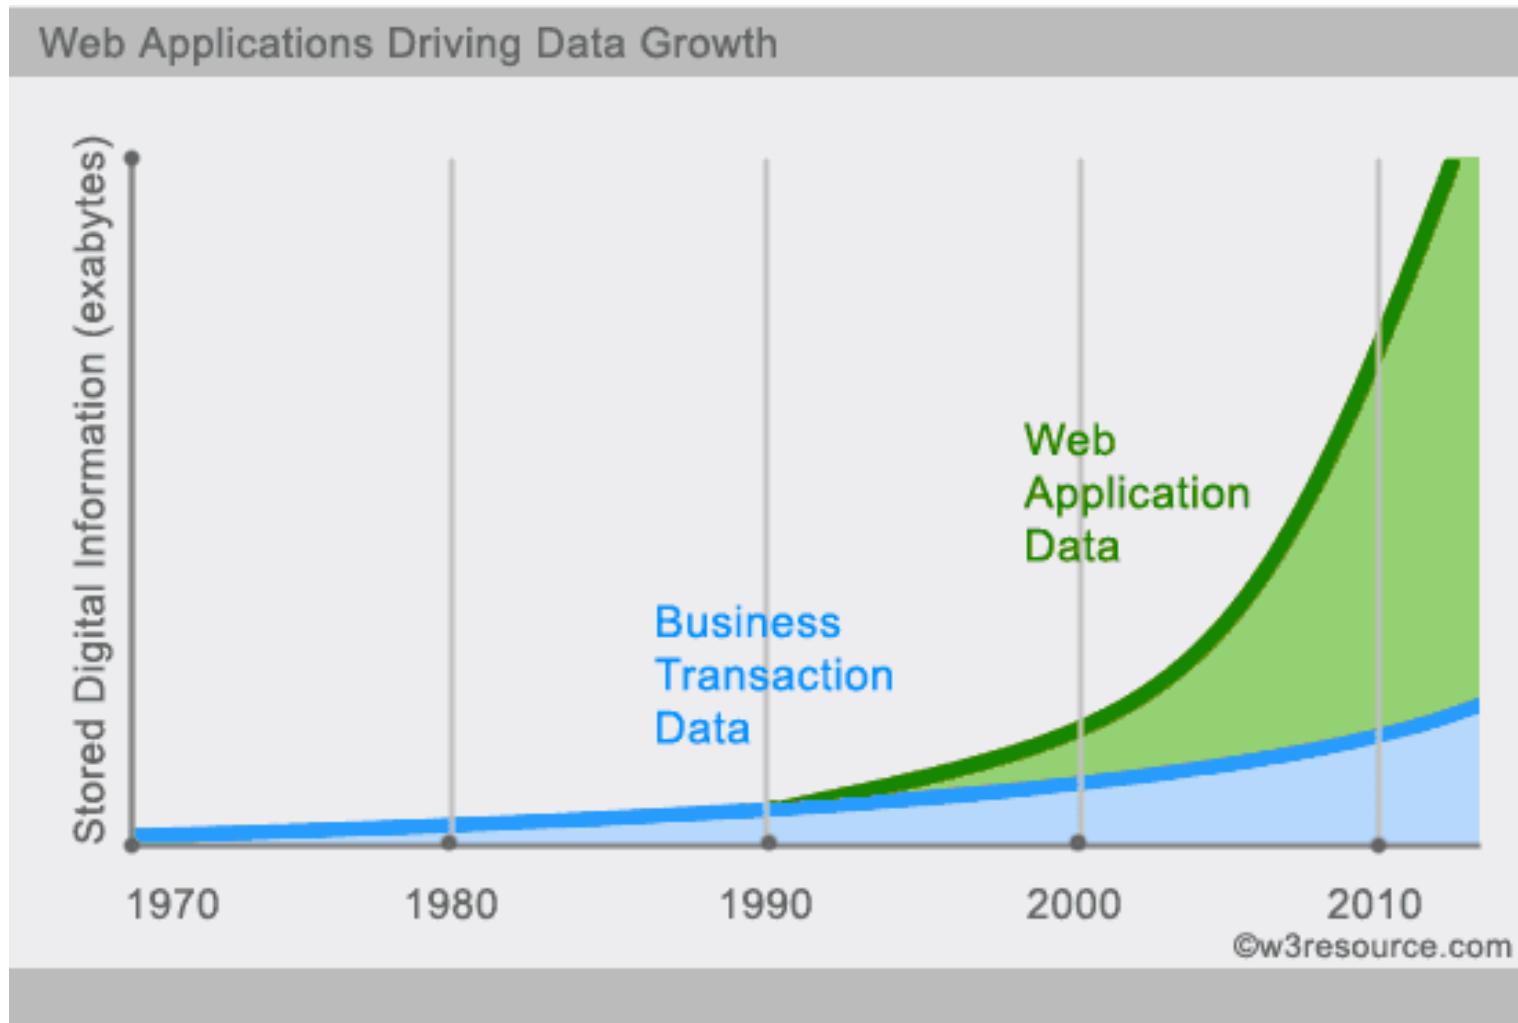
Non-Relational Databases (NoSQL)

# Who uses NoSQL?

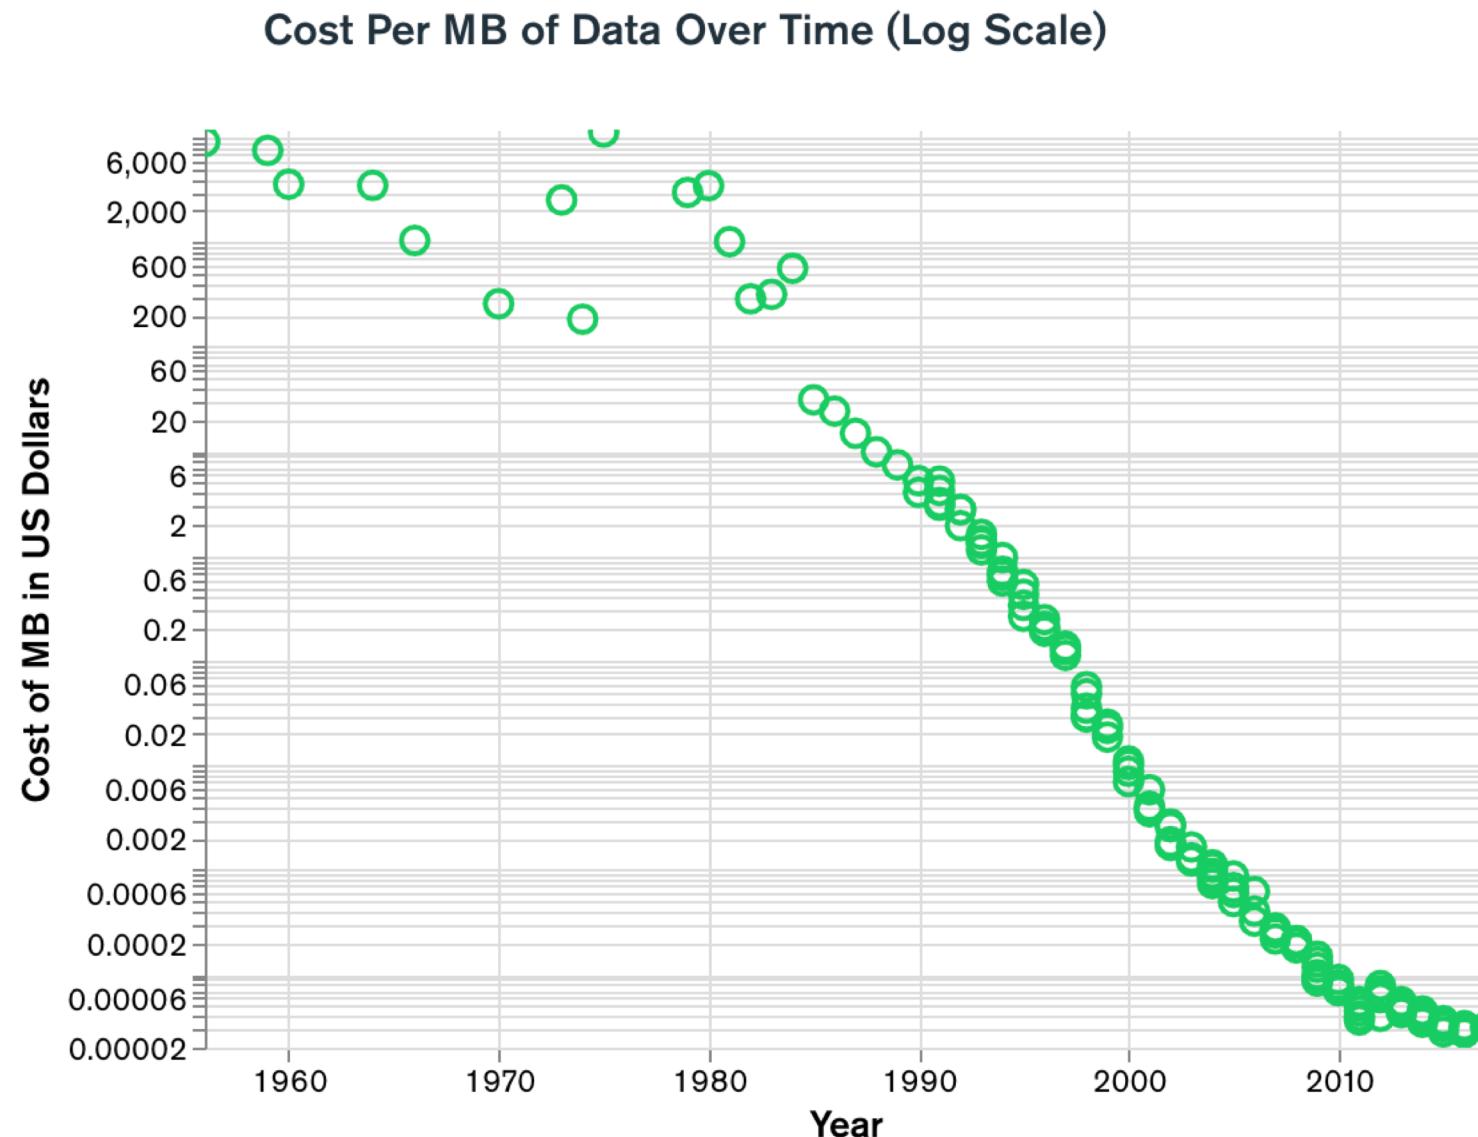


Adobe

# Why NoSQL: Growth in Web Applications



# Why NoSQL: Volume, Variety, Velocity



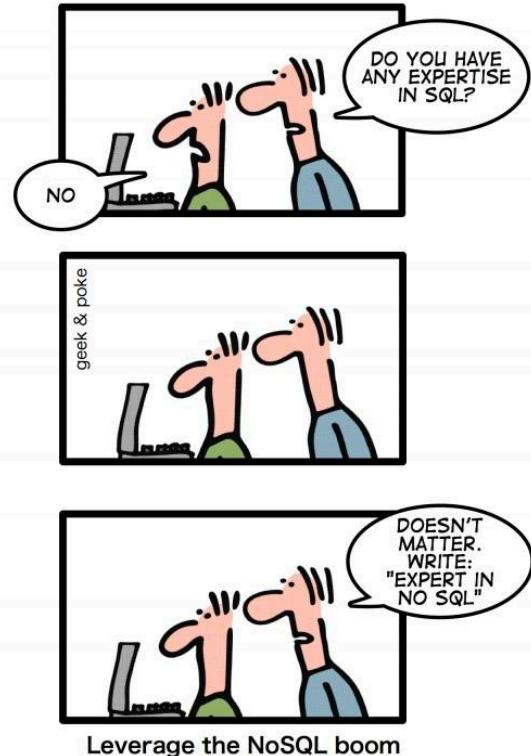
# A Brief Overview of NoSQL Systems

- 1. Horizontally scalability
  - 2. Replicate/distribute data over many servers
  - 3. Simple call interface
  - 4. Often weaker concurrency model than RDBMS
  - 5. Efficient use of distributed indexes and RAM
  - 6. Flexible schemas
- 
- Volume, Velocity
  - Variety

# Today's Plan

1. What is NoSQL?
2. Major types of NoSQL systems
3. Key concepts

*HOW TO WRITE A CV*



# (Major) Types of NoSQL databases

- Key-value stores



DynamoDB

- Wide column databases



- Document stores



- Graph databases



# Key-Value Stores



# Key-Value Stores: Data Model

Phone directory

| Key   | Value            |
|-------|------------------|
| Paul  | (091) 9786453778 |
| Greg  | (091) 9686154559 |
| Marco | (091) 9868564334 |

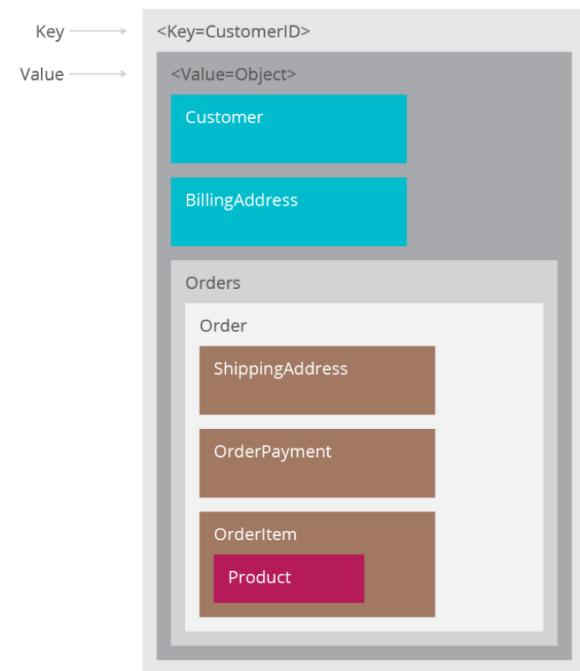
MAC table

| Key           | Value             |
|---------------|-------------------|
| 10.94.214.172 | 3c:22:fb:86:c1:b1 |
| 10.94.214.173 | 00:0a:95:9d:68:16 |
| 10.94.214.174 | 3c:1b:fb:45:c4:b1 |

- Stores associations between keys and values
  - Based on Amazon's Dynamo paper (2007)
- Keys are usually primitives and can be queried
  - For example, ints, strings, raw bytes, etc.
- Values can be primitive or complex; usually cannot be queried
  - Examples: ints, strings, lists, JSON, HTML fragments, BLOB (basic large object), etc.

# Key-Value Stores: Operations

- Very simple API:
  - Get – fetch value associated with key
  - Put – set value associated with key
- Optional operations:
  - Multi-get
  - Multi-put
  - Range queries
- Suitable for:
  - Small continuous read and writes
  - Storing ‘basic’ information (e.g. raw chunks of bytes), or no clear schema
  - When complex queries are not required / rarely required
- Example Applications
  - Storing user sessions
  - Caches
  - User data that is often processed individually: e.g. patient medical data?
    - Only if no queries like ‘How many patients who took X treatment recovered?’



# Key-Value Stores: Implementation

- Non-persistent:
  - Just a big in-memory hash table
  - Examples: Memcached, Redis
    - (But: these can also back up the data to disk periodically)
- Persistent
  - Data is stored persistently to disk
  - Examples: RocksDB, Dynamo, Riak



**Document Stores**

# Document Stores: Document Model



- A database can have multiple **collections**
- Collections have multiple **documents**
- A document is a JSON-like object: it has **fields and values**
  - Different documents can have different fields
  - Can be nested: i.e. JSON objects as values

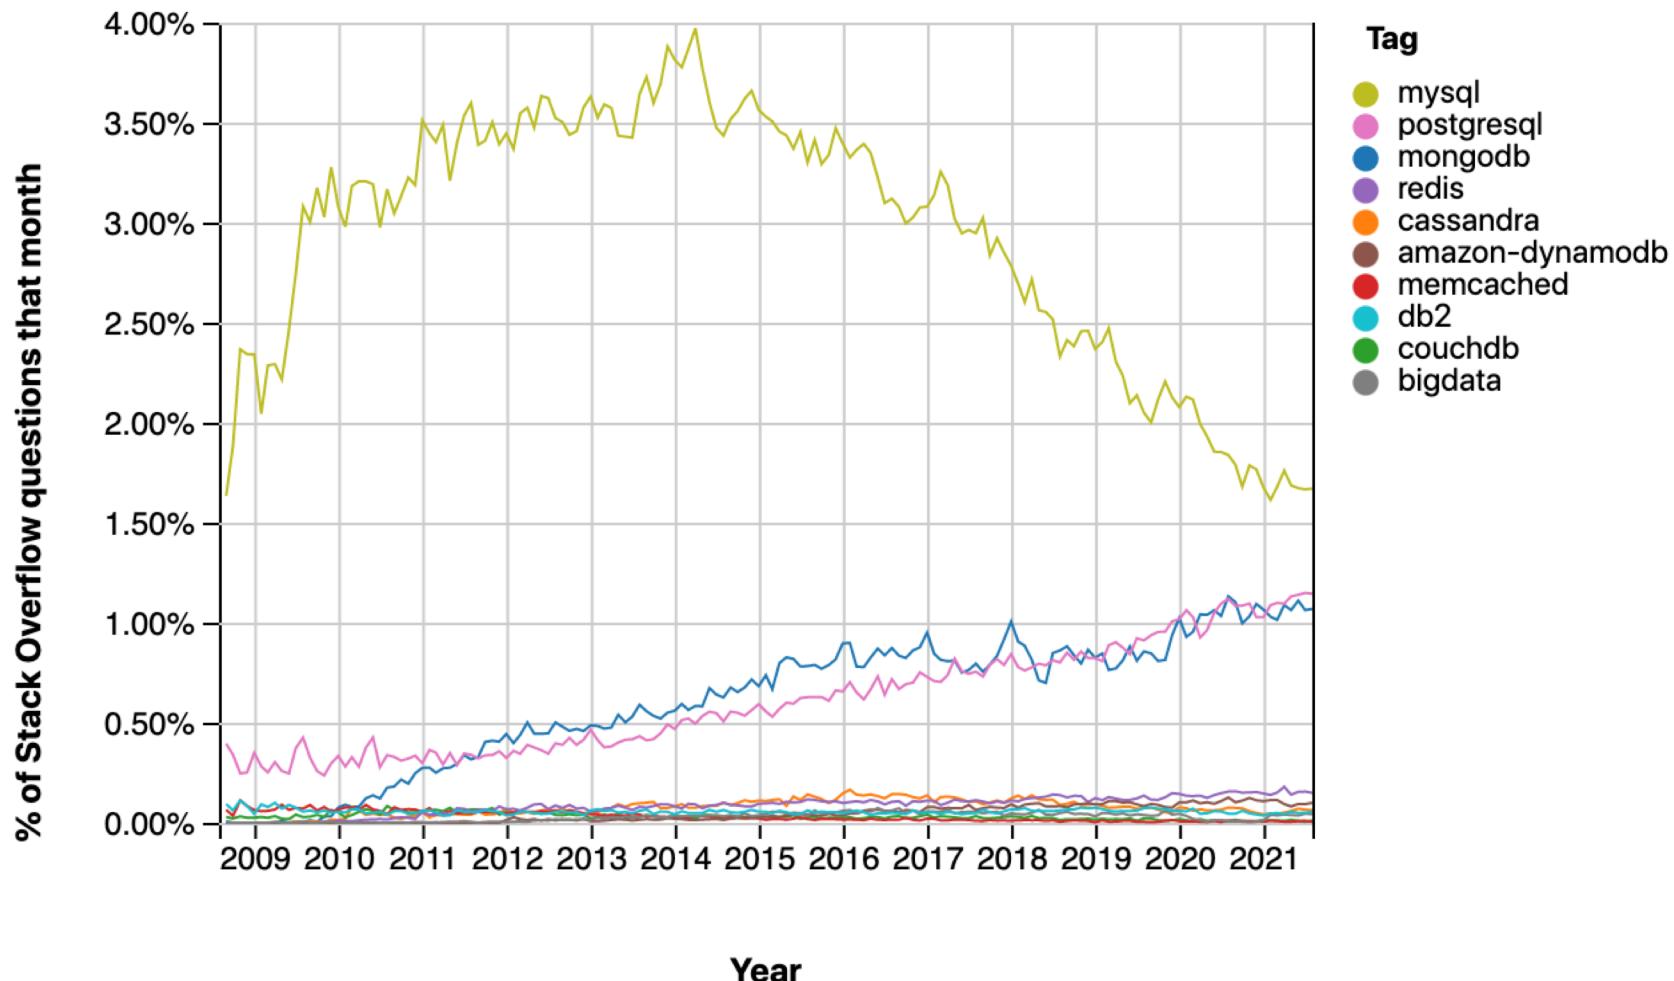
# Document Stores: Querying

- Unlike (basic) key value stores, document stores allow some querying based on the content of a document
- CRUD = Create, Read, Update, Delete

# Popularity of DBMS

| Rank     |          |          | DBMS   | Database Model   | Score    |          |          |
|----------|----------|----------|--|--|----------|----------|----------|
| Sep 2021 | Aug 2021 | Sep 2020 |  |  | Sep 2021 | Aug 2021 | Sep 2020 |
| 1.       | 1.       | 1.       | Oracle                | Relational, Multi-model     | 1271.55  | +2.29    | -97.82   |
| 2.       | 2.       | 2.       | MySQL                 | Relational, Multi-model     | 1212.52  | -25.69   | -51.72   |
| 3.       | 3.       | 3.       | Microsoft SQL Server  | Relational, Multi-model     | 970.85   | -2.50    | -91.91   |
| 4.       | 4.       | 4.       | PostgreSQL            | Relational, Multi-model     | 577.50   | +0.45    | +35.22   |
| 5.       | 5.       | 5.       | MongoDB               | Document, Multi-model       | 496.50   | -0.04    | +50.02   |
| 6.       | 6.       | ↑ 7.     | Redis                 | Key-value, Multi-model      | 171.94   | +2.05    | +20.08   |
| 7.       | 7.       | ↓ 6.     | IBM Db2  | Relational, Multi-model     | 166.56   | +1.09    | +5.32    |
| 8.       | 8.       | 8.       | Elasticsearch  | Search engine, Multi-model  | 160.24   | +3.16    | +9.74    |
| 9.       | 9.       | 9.       | SQLite                | Relational   | 128.65   | -1.16    | +1.98    |
| 10.      | ↑ 11.    | 10.      | Cassandra             | Wide column  | 118.99   | +5.33    | -0.18    |
| 11.      | ↓ 10.    | 11.      | Microsoft Access   | Relational   | 116.94   | +2.10    | -1.51    |
| 12.      | 12.      | 12.      | MariaDB               | Relational, Multi-model     | 100.70   | +1.72    | +9.09    |
| 13.      | 13.      | 13.      | Splunk   | Search engine  | 91.61    | +1.01    | +3.71    |
| 14.      | 14.      | ↑ 15.    | Hive   | Relational   | 85.58    | +1.64    | +14.41   |
| 15.      | 15.      | ↑ 17.    | Microsoft Azure SQL Database   | Relational, Multi-model   | 78.26    | +3.11    | +17.81   |
| 16.      | 16.      | 16.      | Amazon DynamoDB     | Multi-model               | 76.93    | +2.03    | +10.75   |
| 17.      | 17.      | ↓ 14.    | Teradata   | Relational, Multi-model   | 69.68    | +0.86    | -6.72    |
| 18.      | 18.      | ↑ 21.    | Neo4j               | Graph  | 57.63    | +0.68    | +7.00    |
| 19.      | 19.      | 19.      | SAP HANA            | Relational, Multi-model   | 56.24    | +0.66    | +3.38    |
| 20.      | ↑ 21.    | ↑ 23.    | FileMaker  | Relational   | 52.32    | +2.04    | +4.75    |

# Stack Overflow Trends



# CRUD: Create

In MongoDB

```
db.users.insert ( ← collection
  {
    name: "sue", ← field: value
    age: 26, ← field: value
    status: "A" ← field: value
  }
)
```

} document

In SQL

```
INSERT INTO users ← table
          ( name, age, status ) ← columns
VALUES      ( "sue", 26, "A" ) ← values/row
```

# CRUD: Read

## In MongoDB

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
).limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

## In SQL

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

← projection  
← table  
← select criteria  
← cursor modifier

# CRUD: Update

In MongoDB

```
db.users.update(  
    { age: { $gt: 18 } },           ← collection  
    { $set: { status: "A" } },      ← update criteria  
    { multi: true }               ← update action  
)  
                                ← update option
```

In SQL

```
UPDATE users                ← table  
SET status = 'A'             ← update action  
WHERE age > 18              ← update criteria
```

# CRUD: Delete

In MongoDB

```
db.users.remove(  
    { status: "D" }  
)
```



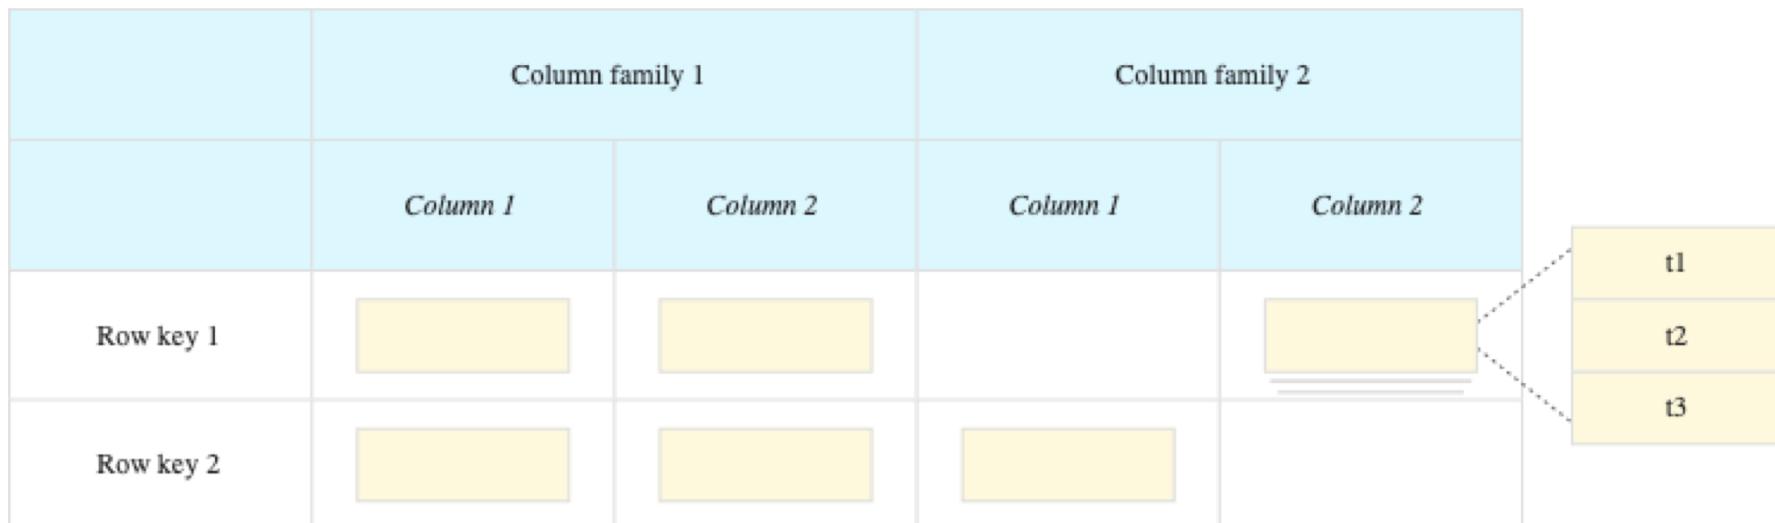
In SQL

```
DELETE FROM users  
WHERE status = 'D'
```

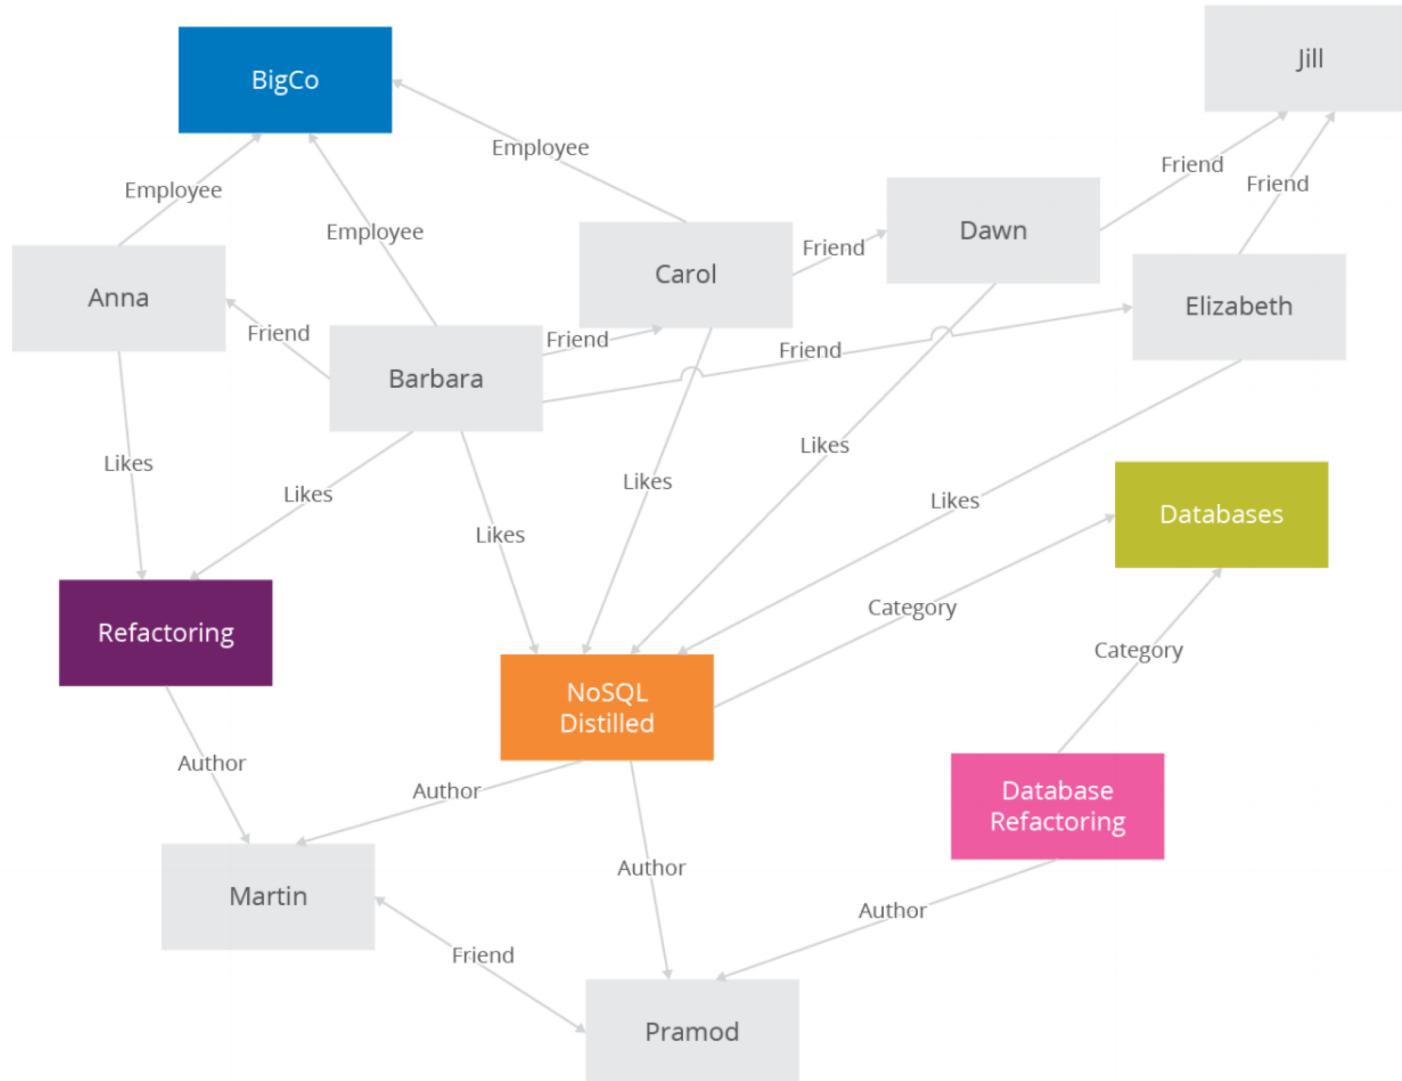


# Wide Column Stores

- Rows describe entities
- Related groups of columns are grouped as **column families**
- **Sparsity**: if a column is not used for a row, it doesn't use space
- Examples: BigTable, Cassandra, HBase



# Graph Databases



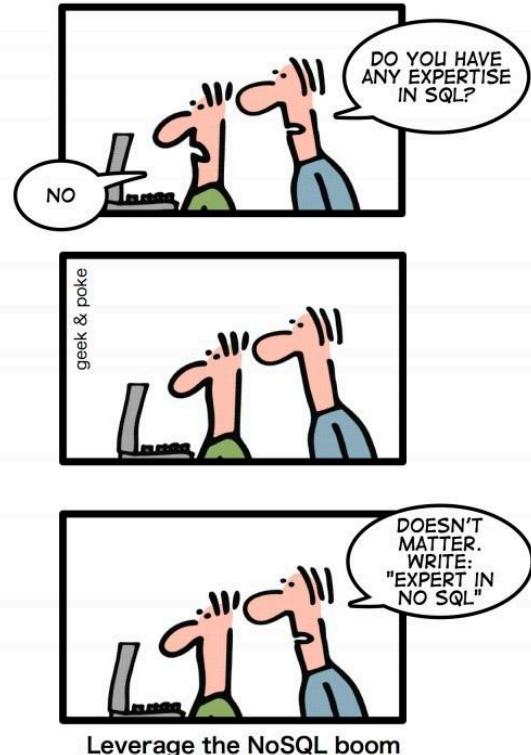
# Today's Plan

1. What is NoSQL?
2. Major types of NoSQL systems

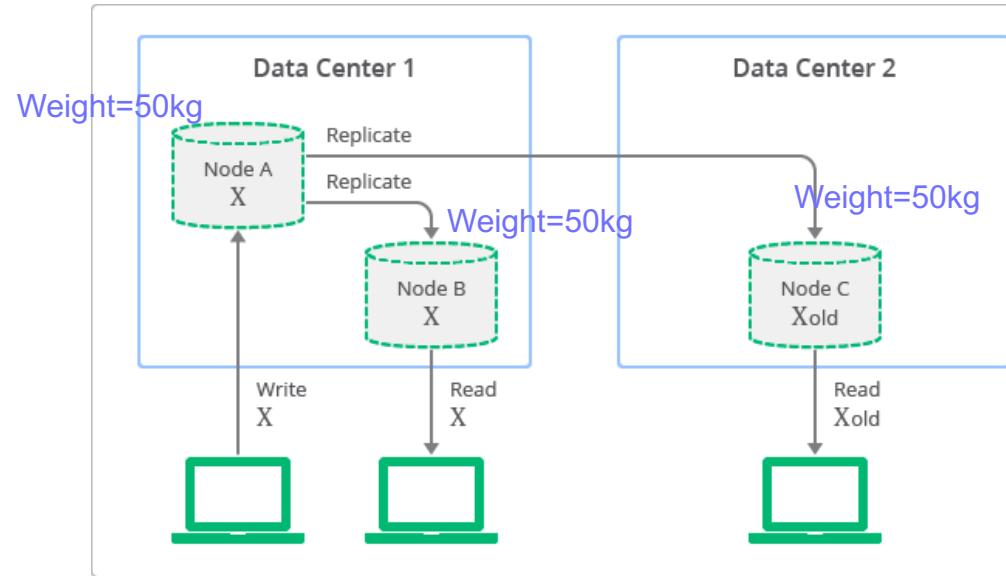
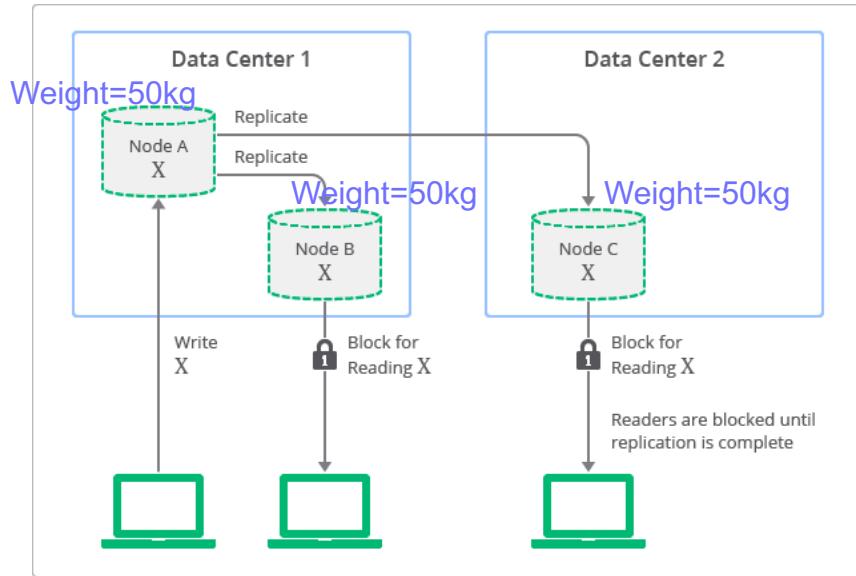
## 3. Key concepts

- Eventual Consistency
- Duplication

*HOW TO WRITE A CV*



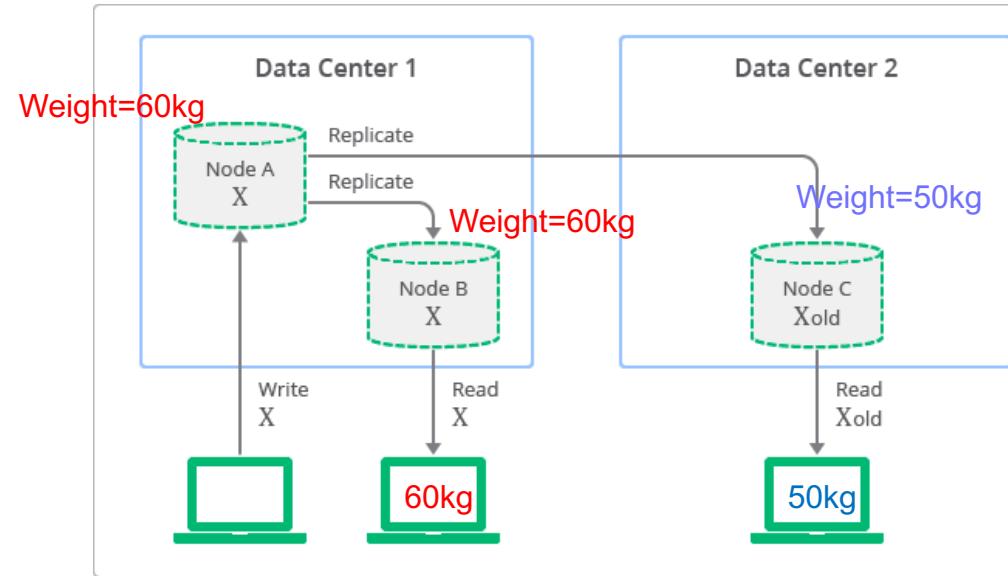
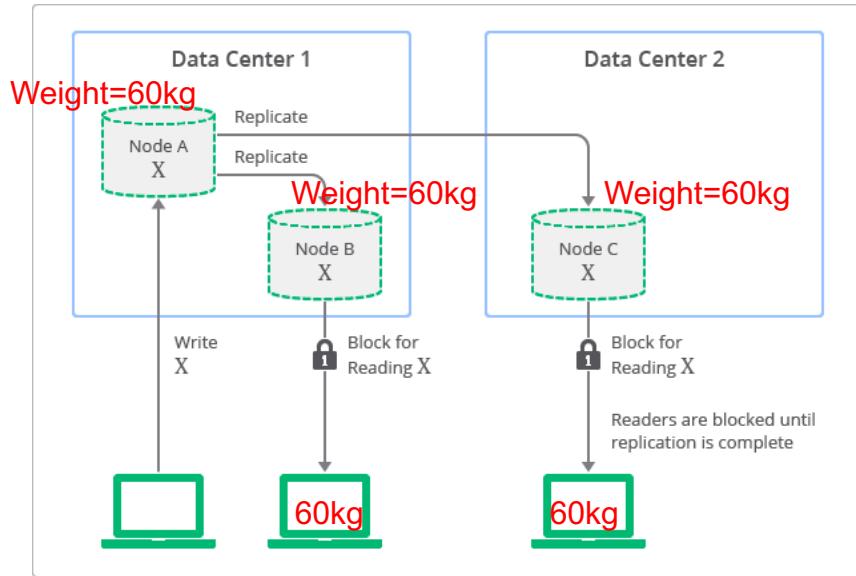
# Strong vs Eventual Consistency



**Strong consistency:** any reads immediately after an update must give the same result on all observers

**Eventual consistency:** if the system is functioning and we wait long enough, eventually all reads will return the last written value

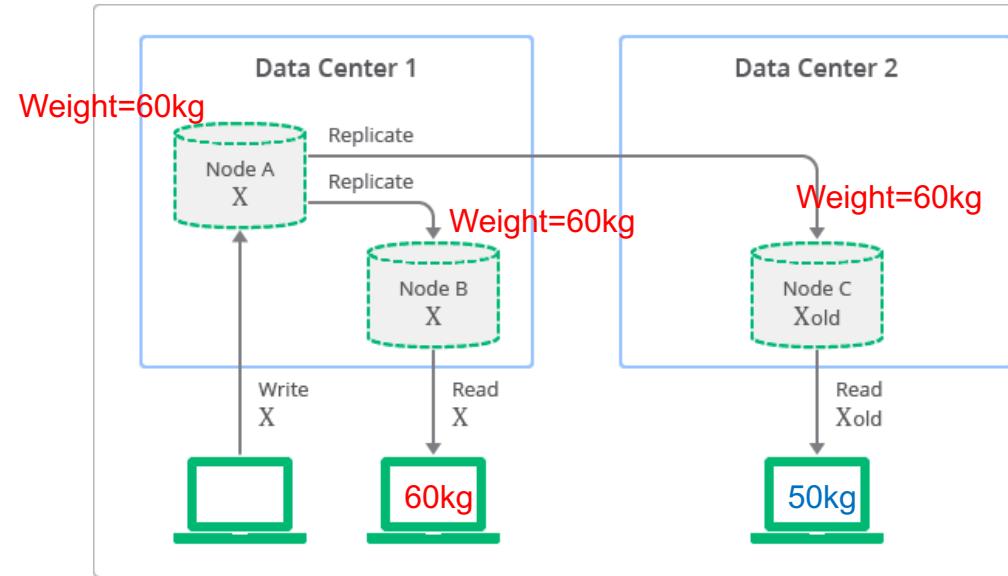
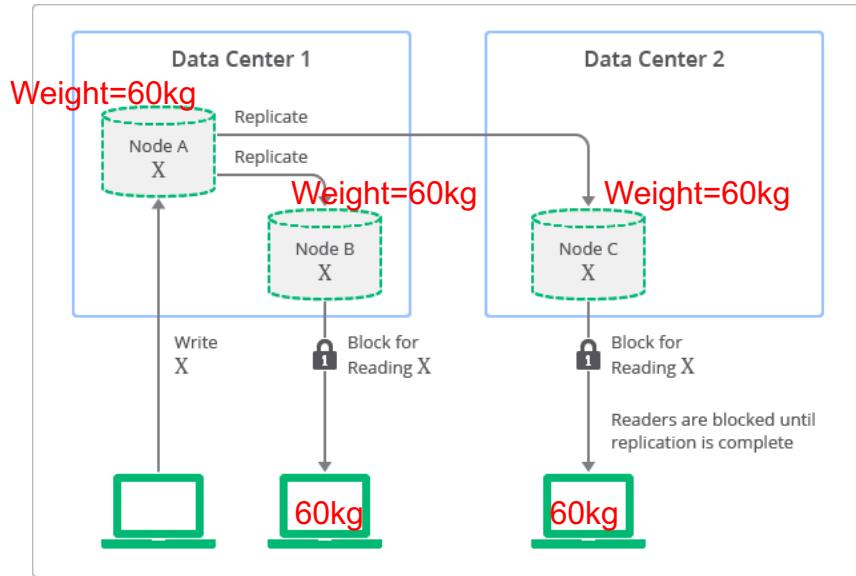
# Strong vs Eventual Consistency



**Strong consistency:** any reads immediately after an update must give the same result on all observers

**Eventual consistency:** if the system is functioning and we wait long enough, eventually all reads will return the last written value

# Strong vs Eventual Consistency



**Strong consistency:** any reads immediately after an update must give the same result on all observers

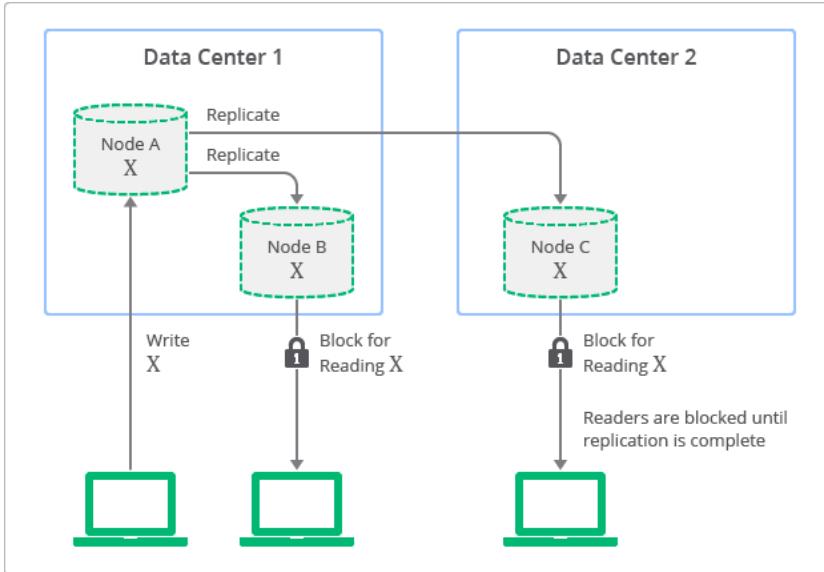
**Eventual consistency:** if the system is functioning and we wait long enough, eventually all reads will return the last written value

# Eventual Consistency

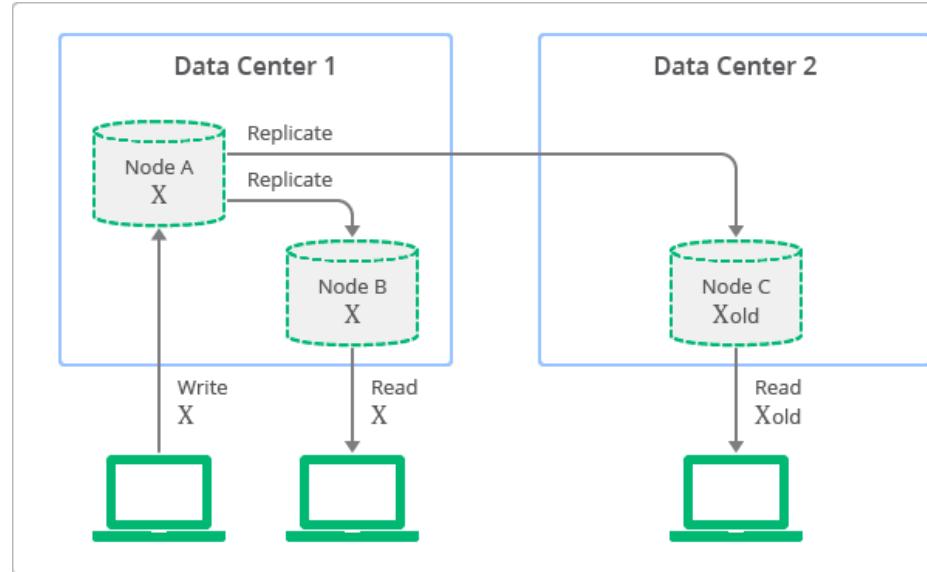
- **ACID vs BASE:** Relational DBMS provide stronger (ACID) guarantees, but many NoSQL system relax this to weaker “BASE” approach:
  - **Basically Available:** basic reading and writing operations are available most of the time
  - **Soft State:** without guarantees, we only have some probability of knowing the state at any time
  - **Eventually Consistent:** Contrast to “**strong consistency**”:

# Strong vs Eventual Consistency

## Strong consistency



## Eventual consistency



- **Implications:** eventual consistency offers better **availability** at the cost of a much weaker consistency guarantee. This may be acceptable for some applications (e.g. statistical queries, tweets, social network feed, ...) but not for others (e.g. financial transactions)
  - Note that while NoSQL systems allow for weaker consistency guarantees, many more recent systems / versions are often configurable, i.e. can be configured for multiple different consistency levels (including strong) – ‘tunable consistency’

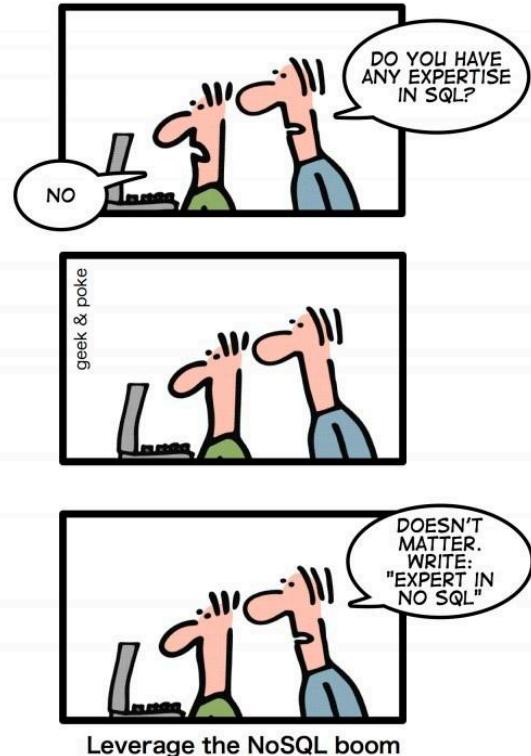
# Today's Plan

1. What is NoSQL?
2. Major types of NoSQL systems

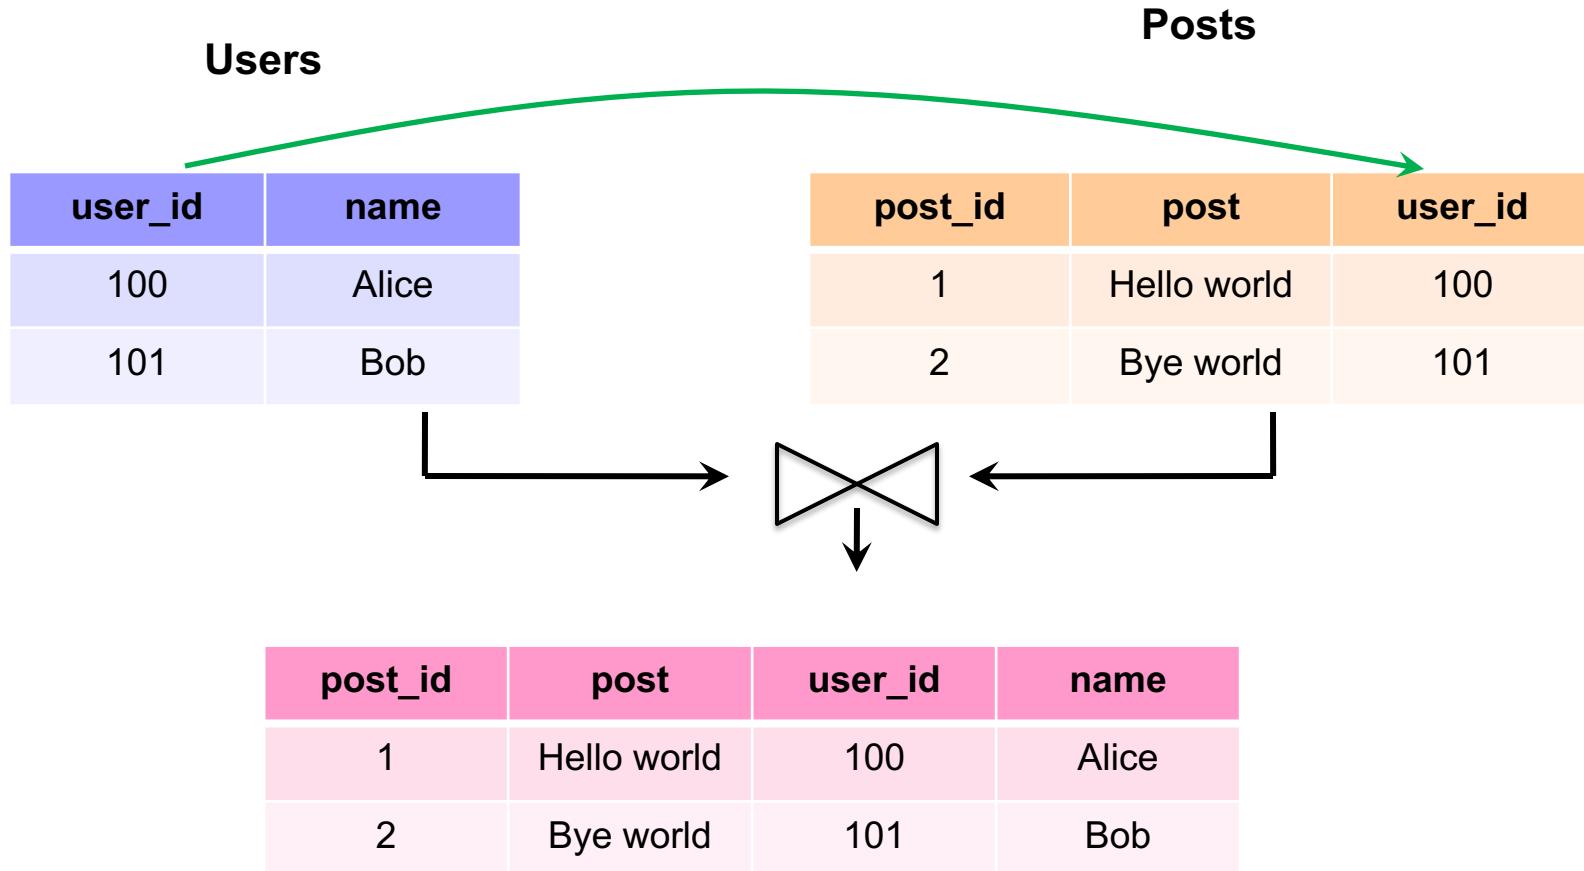
## 3. Key concepts

- Eventual Consistency
- Duplication

*HOW TO WRITE A CV*



# Recall: Joins in RDBMS



# What if we want to display posts with usernames?



- Answer 1: some NoSQL databases do support joins (e.g. later versions of MongoDB)
- Answer 2: Duplication (i.e. ‘denormalization’)
  - ‘Storage is cheap: why not just duplicate data to improve efficiency?’
  - Tables are designed around the queries we expect to receive
  - Leads to a new problem: what if user changes their name? (this needs to be propagated to multiple tables)

# Conclusion: Pros & Cons of NoSQL Systems

## Pros

- + **Flexible / Dynamic Schema:** suitable for less well-structured data
- + **Massive Scalability / High Performance**
- + **High Availability**
- + **Relaxed consistency:** higher performance and availability

## Cons

- **No declarative query language:** query logic (e.g. joins) may have to be handled on the application side, which can add additional programming
- **Relaxed consistency:** fewer guarantees; application may receive stale data that may need to be handled on the application side

**Conclusion:** no one size fits all. Depends on needs of application: complexity of queries (joins vs simple read/writes), importance of consistency (e.g. financial transactions vs tweets), data volume / need for availability.

- Binary view is oversimplified: NoSQL databases often still have SQL-like query languages (e.g. Cassandra) and tunable consistency levels.
- Other alternative solutions: “NewSQL”, distributed SQL (e.g. Presto), PostgreSQL JSON, ...
- Need to understand tradeoffs to make an informed decision!

# Acknowledgements

- CS4225 slides by He Bingsheng
- Penn State CMPSC 431W Database Management Systems (Fall 2015) – McGrawHill, Wang-Chien Lee, Yu-San Lin
- mongodb.com, redis.com
- Otter video from  
[https://www.reddit.com/r/singapore/comments/mezgg3/otters\\_on\\_the\\_grass\\_at\\_botanic\\_gardens\\_today/](https://www.reddit.com/r/singapore/comments/mezgg3/otters_on_the_grass_at_botanic_gardens_today/)