

# **CS4225/CS5425 Big Data Systems for Data Science**

## **PageRank and Graph Processing**

Bryan Hooi  
School of Computing  
National University of Singapore  
[bhooi@comp.nus.edu.sg](mailto:bhooi@comp.nus.edu.sg)



# Announcements

- Next week is last lecture + tutorial (on stream processing)
- Practice test will be posted next week. For those who want to do the test in school, you can come to COM1-02-06 Seminar Room I.
- HW2 is due 31 Oct 11.59pm

Week	Date	Topics	Tutorial	Due Dates
1	12 Aug	Overview and Introduction		
2	19 Aug	MapReduce - Introduction		
3	26 Aug	MapReduce and Relational Databases		
4	2 Sep	MapReduce and Data Mining	Tutorial: Hadoop	
5	9 Sep	NoSQL Overview 1		Assignment 1 released
6	16 Sep	NoSQL Overview 2		
Recess				
7	30 Sep	Apache Spark 1	Tutorial: NoSQL & Spark	Assignment 2 released
8	7 Oct	Apache Spark 2		Assignment 1 due (13 Oct)
9	14 Oct	Large Graph Processing 1	Tutorial: Graph Processing	
10	21 Oct	Large Graph Processing 2		
11	28 Oct	Stream Processing	Tutorial: Stream Processing	Assignment 2 due (31 Oct)
12	4 Nov	Deepavali – No Class		
13	11 Nov	Test		

# Recap: PageRank Flow Formulation

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

Importance of  $j$

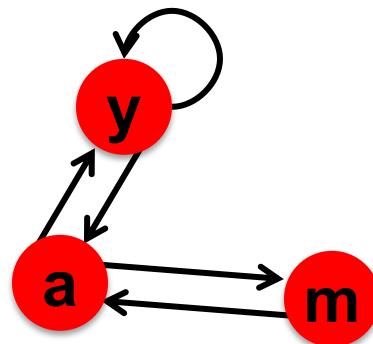
Sum of importances of pages linking to  $j$ , each divided by number of out-links

$d_i$  = number of out-links (or “out-degree”) of node  $i$

**Matrix Form:**

$$\mathbf{r} = \mathbf{M} \cdot \mathbf{r}$$

## Example:



$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} & 0 \\ \frac{1}{2} & 0 & 1 \\ 0 & \frac{1}{2} & 0 \end{bmatrix} \begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix}$$

# Recap: Power Iteration

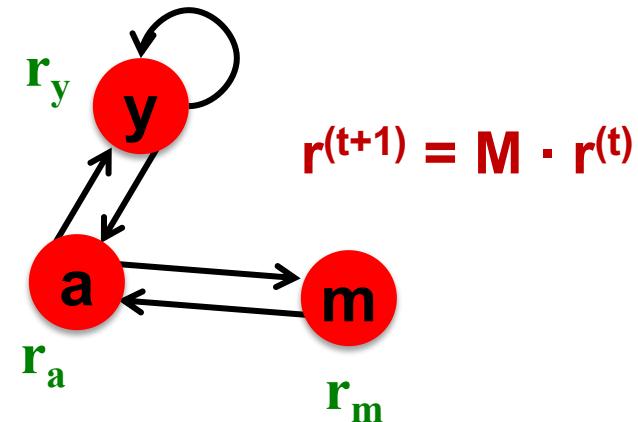
## Power Iteration:

- Suppose there are  $N$  web pages
- Initialize:  $\mathbf{r}^{(0)} = [1/N, \dots, 1/N]^T$
- Iterate:  $\mathbf{r}^{(t+1)} = \mathbf{M} \cdot \mathbf{r}^{(t)}$
- Stop when  $|\mathbf{r}^{(t+1)} - \mathbf{r}^{(t)}|_1 < \varepsilon$

### Example:

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{array}{cccccc} 1/3 & 1/3 & 5/12 & 9/24 & & 2/5 \\ 1/3 & 3/6 & 1/3 & 11/24 & \dots & 2/5 \\ 1/3 & 1/6 & 3/12 & 1/6 & & 1/5 \end{array}$$

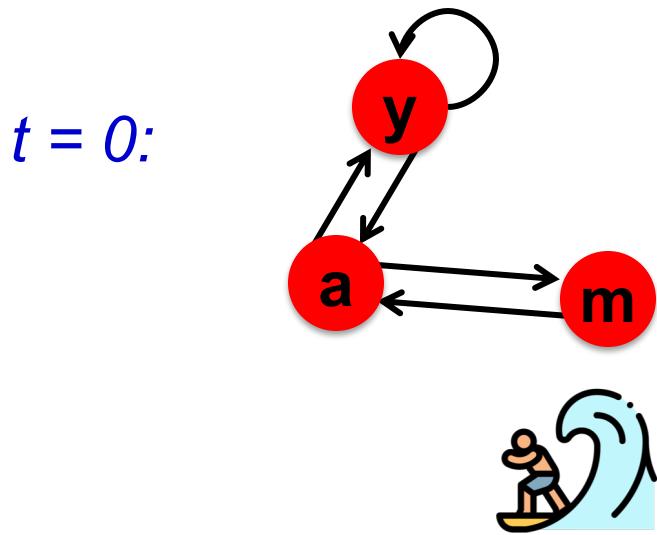
Iteration 0      Iteration 1      ...



# Recap: Random Walk Interpretation

- Imagine a random web surfer:

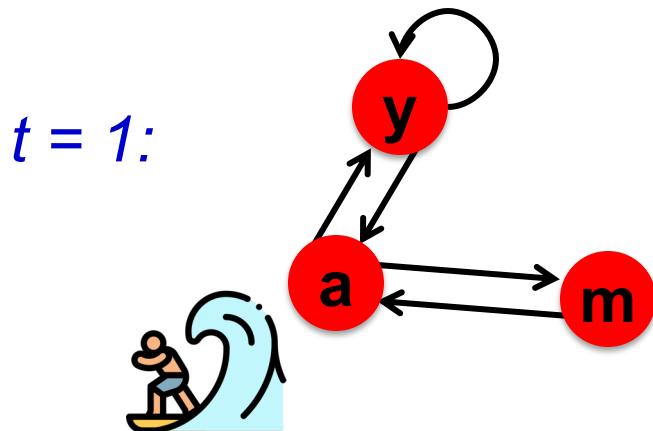
- At time  $t = 0$ , surfer starts on a random page
- At any time  $t$ , surfer is on some page  $i$
- At time  $t + 1$ , the surfer follows an out-link from  $i$  uniformly at random
- Process repeats indefinitely



# Recap: Random Walk Interpretation

- Imagine a random web surfer:

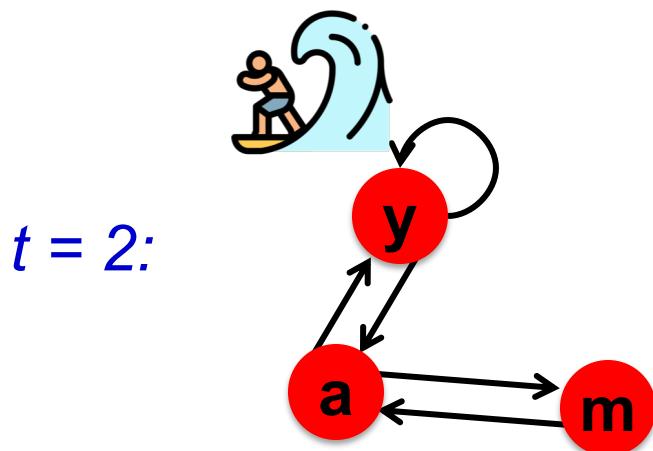
- At time  $t = 0$ , surfer starts on a random page
- At any time  $t$ , surfer is on some page  $i$
- At time  $t + 1$ , the surfer follows an out-link from  $i$  uniformly at random
- Process repeats indefinitely



# Recap: Random Walk Interpretation

- Imagine a random web surfer:

- At time  $t = 0$ , surfer starts on a random page
- At any time  $t$ , surfer is on some page  $i$
- At time  $t + 1$ , the surfer follows an out-link from  $i$  uniformly at random
- Process repeats indefinitely

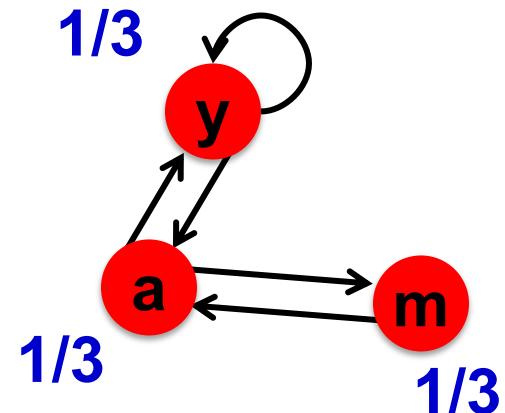


# Recap: Random Walk Interpretation

○ Let:

- $p(t)$  ... vector whose  $i$ th coordinate is the prob. that the surfer is at page  $i$  at time  $t$
- So,  $p(t)$  is a probability distribution over pages

$p(t)$  where  $t = 0$ :

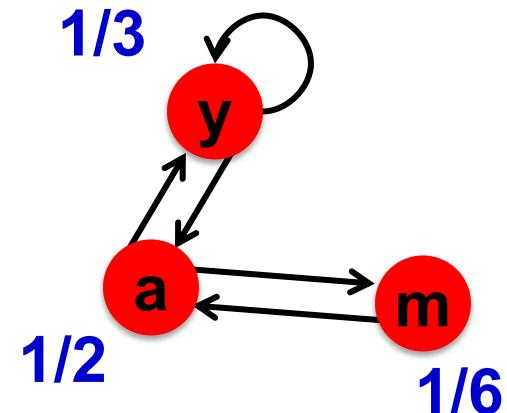


# Recap: Random Walk Interpretation

○ Let:

- $p(t)$  ... vector whose  $i$ th coordinate is the prob. that the surfer is at page  $i$  at time  $t$
- So,  $p(t)$  is a probability distribution over pages

$p(t)$  where  $t = 1$ :



# Recap: Random Walk Interpretation

○ Let:

- $p(t)$  ... vector whose  $i$ th coordinate is the prob. that the surfer is at page  $i$  at time  $t$
- So,  $p(t)$  is a probability distribution over pages

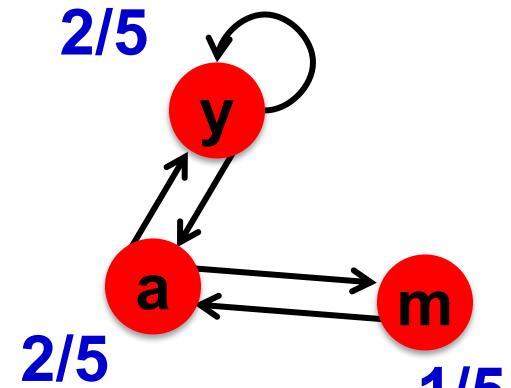
○ **Stationary Distribution:** as  $t \rightarrow \infty$ , the probability distribution gradually converges to a ‘steady state’

○ PageRank equations are the same, except that  $r_j$  represents the (steady state) probability of the random walker being at node  $j$ :

Probability  
of random  
walker being  
at  $j$

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

$p(t)$  where  $t \rightarrow \infty$ :



# Today's Plan

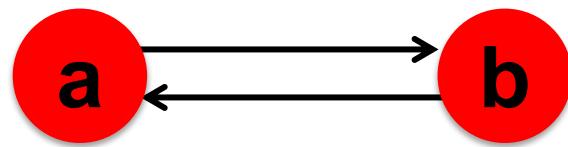
- **PageRank with Teleports**
- Topic Sensitive PageRank & TrustRank
- Graph Processing: Pregel and Giraph

# PageRank: Three Questions

$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i} \quad \text{or equivalently} \quad r = Mr$$

- o Does this converge?
- o Does it converge to what we want?
- o Are results reasonable?

# Does this converge?



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

- Answer: not always. Example:

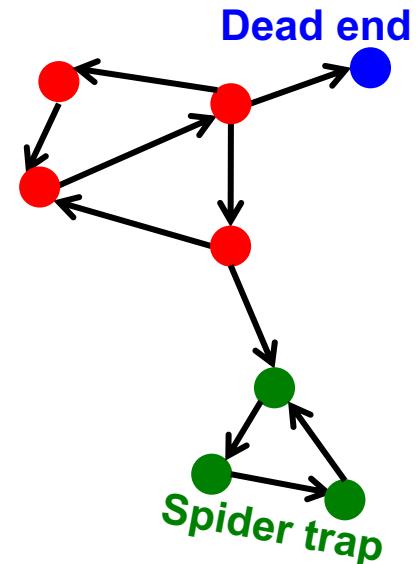
$$\begin{array}{lllll} r_a & = & 1 & 0 & 1 & 0 \\ & & \equiv & & & \\ r_b & & 0 & 1 & 0 & 1 \end{array}$$

Iteration 0, 1, 2, ...

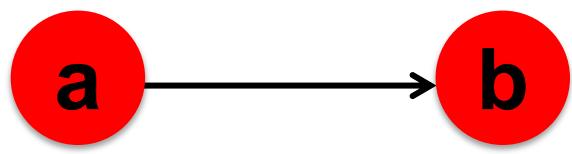
# Does it converge to what we want?

Answer: not always. 2 problems:

- (1) Some pages are **dead ends** (have no out-links)
  - Random walk has “nowhere” to go to
  - Such pages cause importance to “leak out”
- (2) **Spider traps:**  
(all out-links are within the group)
  - Random walk gets “stuck” in a trap
  - And eventually spider traps absorb all importance



# Problem I: Dead Ends



$$r_j^{(t+1)} = \sum_{i \rightarrow j} \frac{r_i^{(t)}}{d_i}$$

- Example:

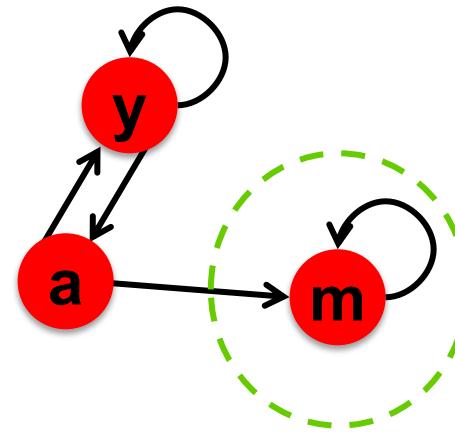
$$\begin{array}{lcl} r_a & = & 1 \quad 0 \quad 0 \quad 0 \\ & & \mid \\ r_b & = & 0 \quad 1 \quad 0 \quad 0 \end{array}$$

Iteration 0, 1, 2, ...

# Problem 2: Spider Traps

## ○ Power Iteration:

- Set  $r_j = 1/N$
- $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$ 
  - And iterate



$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2$$

$$r_m = r_a/2 + r_m$$

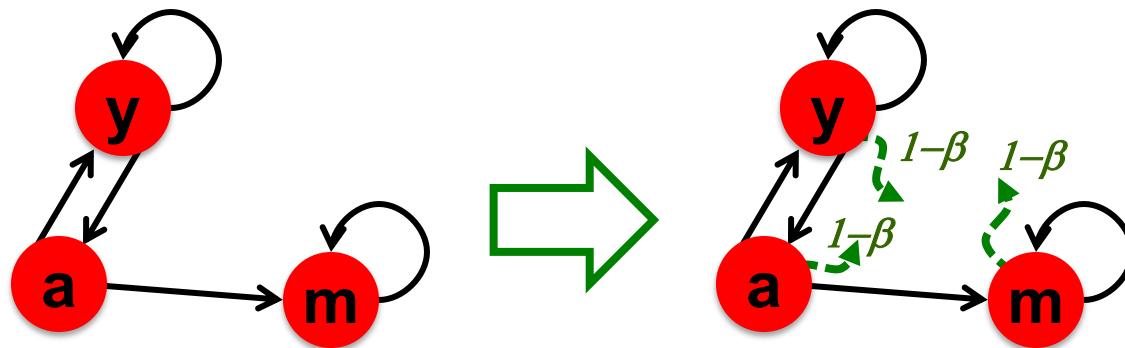
$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} 1/3 & 2/6 & 3/12 & 5/24 & \dots & 0 \\ 1/3 & 1/6 & 2/12 & 3/24 & \dots & 0 \\ 1/3 & 3/6 & 7/12 & 16/24 & \dots & 1 \end{bmatrix}$$

Iteration 0, 1, 2, ...

All the PageRank score gets “trapped” in node m.

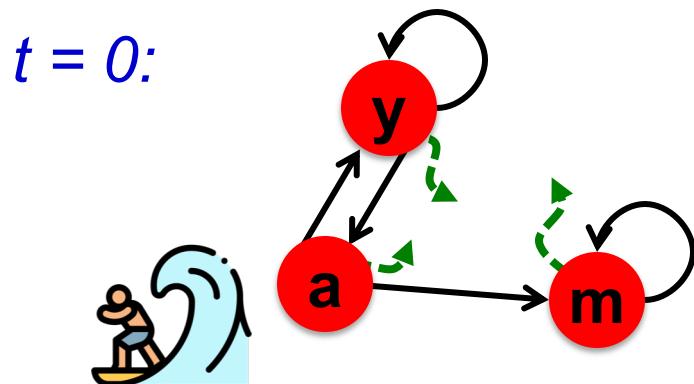
# Solution: Teleports!

- The Google solution for spider traps: At each time step, the random surfer has two options
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9



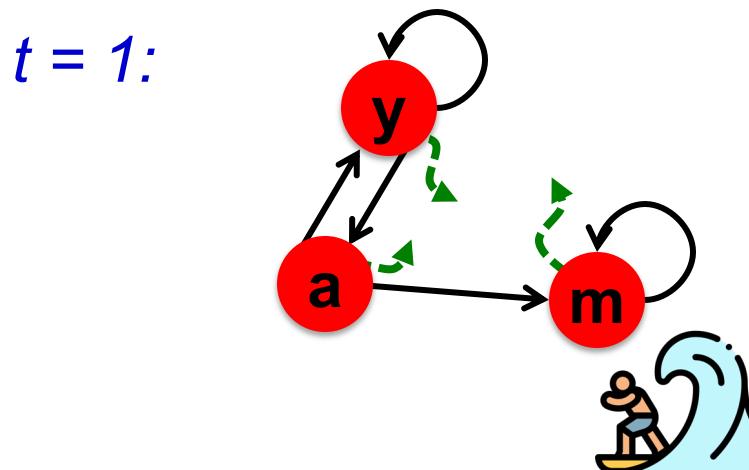
# Solution: Teleports!

- **The Google solution for spider traps: At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9



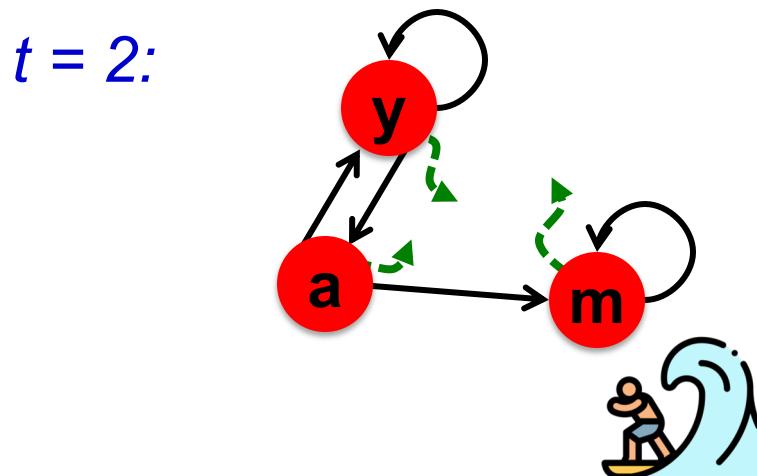
# Solution: Teleports!

- **The Google solution for spider traps: At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9



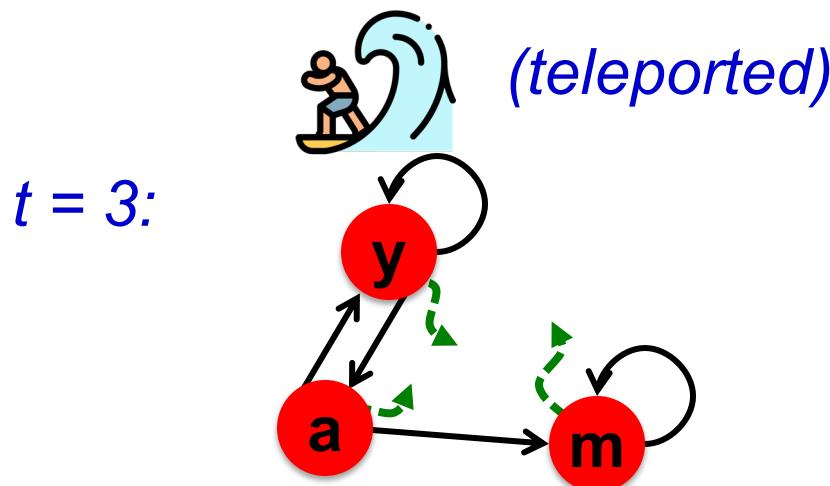
# Solution: Teleports!

- **The Google solution for spider traps: At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9



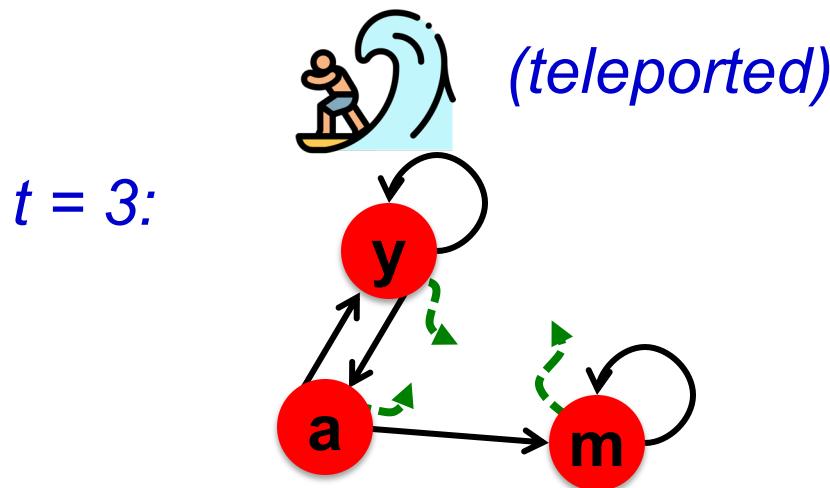
# Solution: Teleports!

- **The Google solution for spider traps: At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9



# Solution: Teleports!

- **The Google solution for spider traps: At each time step, the random surfer has two options**
  - With prob.  $\beta$ , follow a link at random
  - With prob.  $1-\beta$ , jump to some random page
  - Common values for  $\beta$  are in the range 0.8 to 0.9

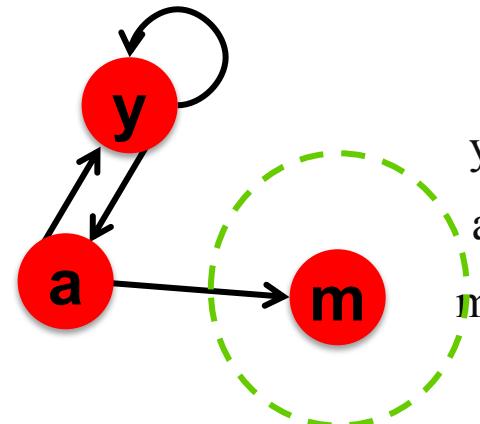


- **Conclusion:** surfer will quickly teleport out of any spider trap

# Problem: Dead Ends

## ○ Power Iteration:

- Set  $r_j = 1/N$
- $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
- And iterate



y	a	m
$\frac{1}{2}$	$\frac{1}{2}$	0
$\frac{1}{2}$	0	0
0	$\frac{1}{2}$	0

$$r_y = r_y/2 + r_a/2$$

$$r_a = r_y/2$$

$$r_m = r_a/2$$

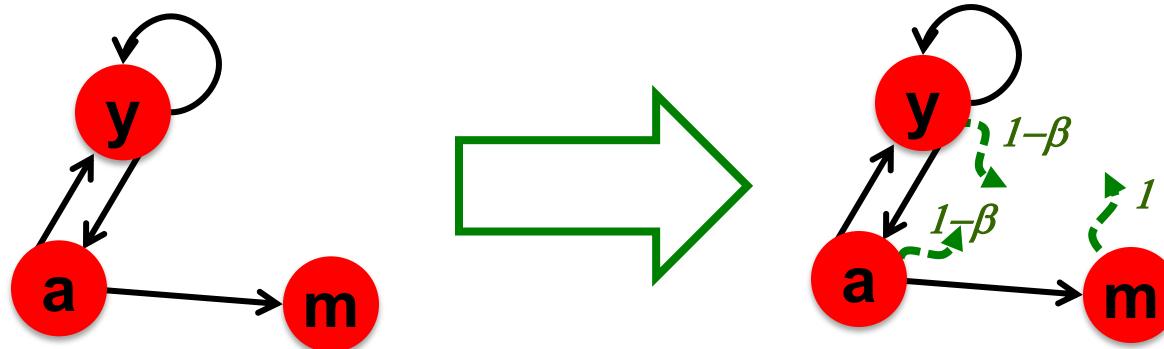
$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{matrix} 1/3 & 2/6 & 3/12 & 5/24 & \dots & 0 \\ 1/3 & 1/6 & 2/12 & 3/24 & \dots & 0 \\ 1/3 & 1/6 & 1/12 & 2/24 & \dots & 0 \end{matrix}$$

Iteration 0, 1, 2, ...

Here the PageRank “leaks” out since the matrix is not stochastic.

# Solution: If at a Dead End, Always Teleport

- **Teleports:** Follow random teleport links with probability 1.0 from dead-ends
  - Adjust matrix accordingly



# Why Teleports Solve the Problem?

**Why are dead-ends and spider traps a problem and why do teleports solve the problem?**

- **Spider-traps** cause random walker to get stuck in them, absorbing all importance
  - **Solution:** Never get stuck in a spider trap by teleporting out of it in a finite number of steps
- **Dead-ends** cause importance to “leak” out of the system
  - The matrix is not column stochastic
  - **Solution:** Make matrix column stochastic by always teleporting when at a dead end

# Random Teleport: Equations

- At each step, random surfer has two options:
  - With probability  $\beta$ , follow a link at random
  - With probability  $1-\beta$ , jump to some random page
- PageRank equation** [Brin-Page, 98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

Same as simplified PageRank      Teleport term

$d_i$  ... out-degree of node i

- Compare to Simplified PageRank:

$$r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$$

(For now, assume that there are no dead ends. If there are, we can handle them - see next slide)

# The Google Matrix

- **PageRank equation** [Brin-Page, '98]

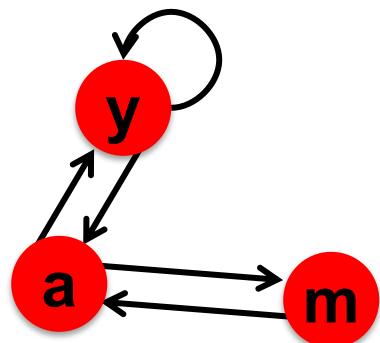
$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

- We can also write this as a matrix equation, by defining the **Google Matrix A**:

$$A = \beta M + (1 - \beta) \begin{bmatrix} 1 \\ \vdots \\ N \end{bmatrix}_{N \times N}$$

N by N matrix  
where all entries are 1/N

y	a	m
$\frac{1}{2}$	$\frac{1}{2}$	0
$\frac{1}{2}$	0	1
0	$\frac{1}{2}$	0



y	a	m
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$

# The Google Matrix

- **PageRank equation** [Brin-Page, '98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

- **The Google Matrix A:**

y	a	m
$\frac{1}{2}$	$\frac{1}{2}$	0
$\frac{1}{2}$	0	1
0	$\frac{1}{2}$	0

$$A = \beta M + (1 - \beta) \begin{bmatrix} \frac{1}{N} \\ \vdots \\ \frac{1}{N} \end{bmatrix}_{N \times N}$$

y	a	m
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$

- **We have a recursive problem:**  $r = A \cdot r$   
**And the Power Iteration method still works!**

- **What is  $\beta$ ?**

- In practice  $\beta = 0.8, 0.9$  (roughly, 5-10 steps on avg before a teleport)

# Some Problems with Page Rank

- **Measures generic popularity of a page**
  - Doesn't consider popularity based on specific topics
  - **Solution:** Topic-Specific PageRank
- **Uses a single measure of importance**
  - Other models of importance
  - **Solution:** Hubs-and-Authorities
- **Susceptible to Link spam**
  - Artificial link topographies created in order to boost page rank
  - **Solution:** TrustRank

# Today's Plan

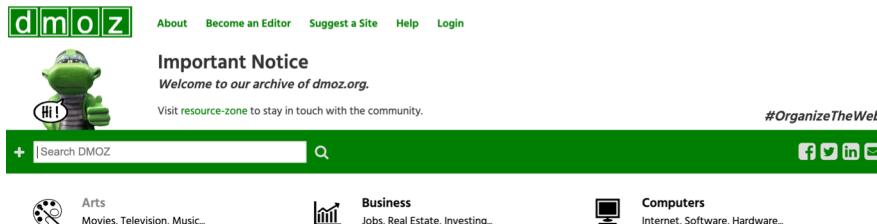
- PageRank with Teleports
- **Topic Sensitive PageRank & TrustRank**
- Graph Processing: Pregel and Giraph

# Topic-Specific PageRank

- Instead of generic popularity, can we measure popularity within a topic?
- Goal: Evaluate Web pages not just according to their popularity, but by how close they are to a particular topic, e.g. “sports” or “history”
- Allows search queries to be answered based on interests of the user
  - Example: Query “Trojan” wants different pages depending on whether you are interested in sports, history and computer security

# Topic-Specific PageRank

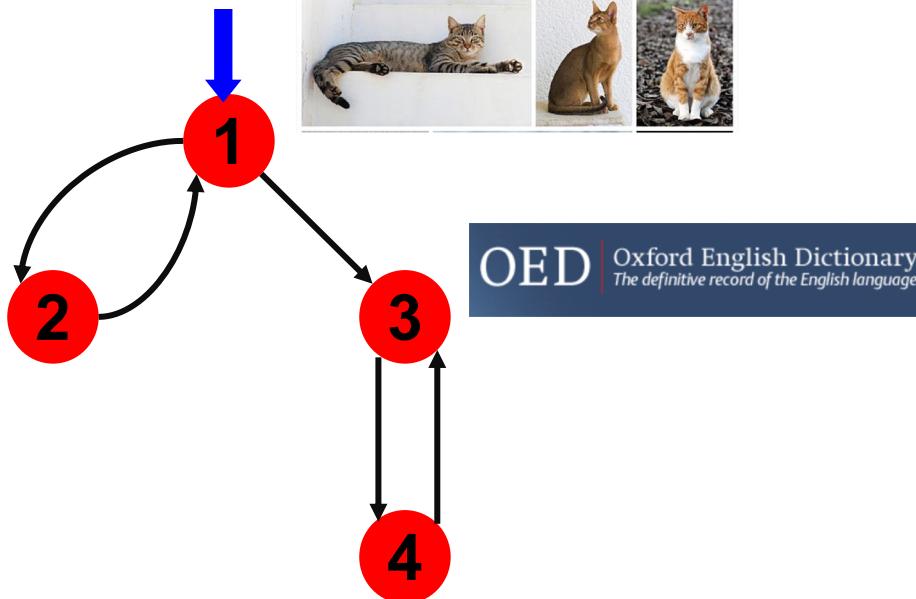
- Random walker has a small probability of teleporting at any step
- **Teleport can go to:**
  - **Standard PageRank:** Any page with equal probability
    - To avoid dead-end and spider-trap problems
  - **Topic Specific PageRank:** A topic-specific set of “relevant” pages (**teleport set**)
- **Idea: Bias the random walk**
  - When random walker teleports, it picks a page from a set **S**
  - **S** contains only pages that are relevant to the topic
    - E.g., Open Directory (DMOZ) pages for a given topic/query
  - For each teleport set **S**, we get a different vector **r<sub>s</sub>**



# Topic Specific PageRank

Example:  $S = \text{Wikipedia page on cats}$

Topic: cats



- **Idea:** By setting the teleport set  $S$  as a small set of relevant pages for a topic, higher importance is assigned to pages in  $S$  and the web pages closely linked to  $S$

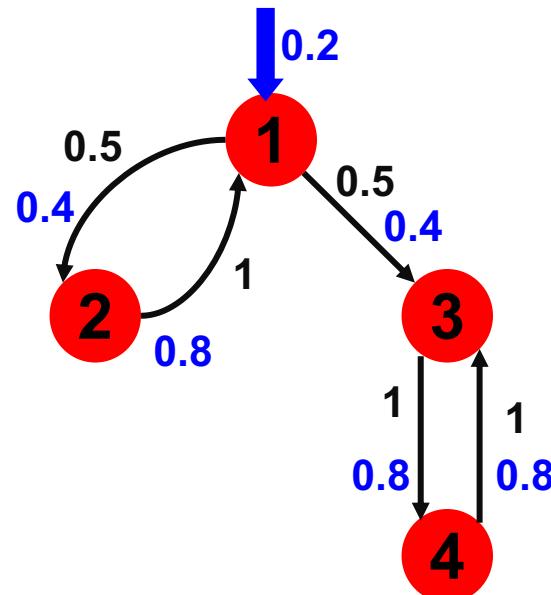
# Matrix Formulation

- To make this work all we need is to update the teleportation part of the PageRank formulation:

$$A_{ij} = \begin{cases} \beta M_{ij} + (1 - \beta)/|S| & \text{if } i \in S \\ \beta M_{ij} + 0 & \text{otherwise} \end{cases}$$

- $A$  is stochastic!
- We weighted all pages in the teleport set  $S$  equally
- Compute as for regular PageRank:
  - Multiply by  $M$ , then add a vector
  - Maintains sparseness

# Example: Topic-Specific PageRank



Probabilities (no teleport)

Probabilities (with teleport)

$$S = \{1\} \quad \beta = 0.8$$

Node	Iteration				
	0	1	2	...	stable
1	0.25	0.4	0.28		0.294
2	0.25	0.1	0.16		0.118
3	0.25	0.3	0.32		0.327
4	0.25	0.2	0.24		0.261

$S=\{1,2,3,4\}, \beta=0.8:$

$r=[0.13, 0.10, 0.39, 0.36]$

$S=\{1,2,3\}, \beta=0.8:$

$r=[0.17, 0.13, 0.38, 0.30]$

$S=\{1,2\}, \beta=0.8:$

$r=[0.26, 0.20, 0.29, 0.23]$

$S=\{1\}, \beta=0.8:$

$r=[0.29, 0.11, 0.32, 0.26]$

$S=\{1\}, \beta=0.90:$

$r=[0.17, 0.07, 0.40, 0.36]$

$S=\{1\}, \beta=0.8:$

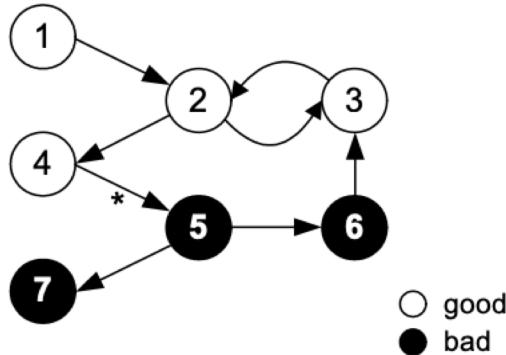
$r=[0.29, 0.11, 0.32, 0.26]$

$S=\{1\}, \beta=0.70:$

$r=[0.39, 0.14, 0.27, 0.19]$

# TrustRank

- **Goal:** identify trusted vs untrusted pages (e.g. spam pages). Start with a small set of trusted pages and propagate trust (assuming that “good” pages only rarely link to “bad” pages)



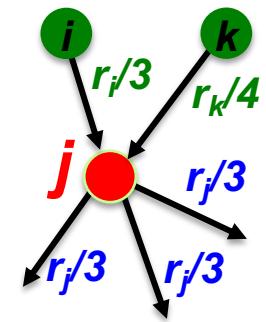
- **Approach:**
  - Sample a set of **seed pages** from the web
  - Have an **oracle (human)** to identify the good pages and the spam pages in the seed set
  - Perform a topic-sensitive PageRank with **teleport set = trusted pages**

# Today's Plan

- PageRank with Teleports
- Topic Sensitive PageRank & TrustRank
- **Graph Processing: Pregel and Giraph**

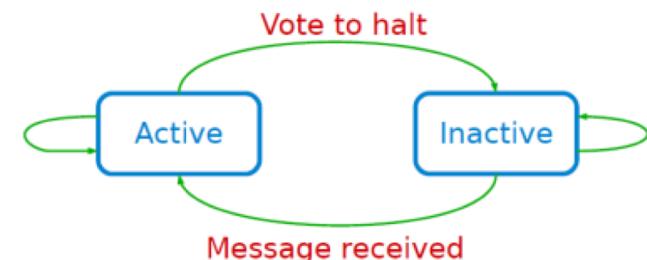
# Characteristics of Graph Algorithms

- What are some common features of graph algorithms?
  - Local computations at each node
  - Passing messages to other nodes
- **Think like a vertex:** algorithms are implemented from the view of a single vertex, performing one iteration based on messages from its neighbor
  - Similar to MapReduce, the user only implements a simple function, `compute()`, that describes the algorithm's behavior at one vertex, in one step.
  - The framework abstracts away the scheduling / implementation details.



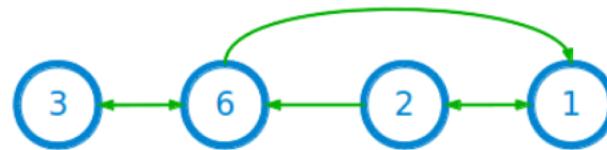
# Pregel: Computational Model

- Computation consists of a series of **supersteps**
- In each superstep, the framework **invokes a user-defined function**, `compute ()`, **for each vertex** (conceptually in parallel)
- `compute ()` specifies **behavior at a single vertex**  $v$  and a superstep  $s$ :
  - It can **read messages** sent to  $v$  in superstep  $s - 1$ ;
  - It can **send messages** to other vertices that will be read in superstep  $s + 1$ ;
  - It can **read or write the value** of  $v$  and the value of its outgoing edges (or even add or remove edges)
- Termination:
  - A vertex can choose to deactivate itself
  - Is “woken up” if new messages received
  - Computation halts when all vertices are inactive



# Example: Computing Max Value

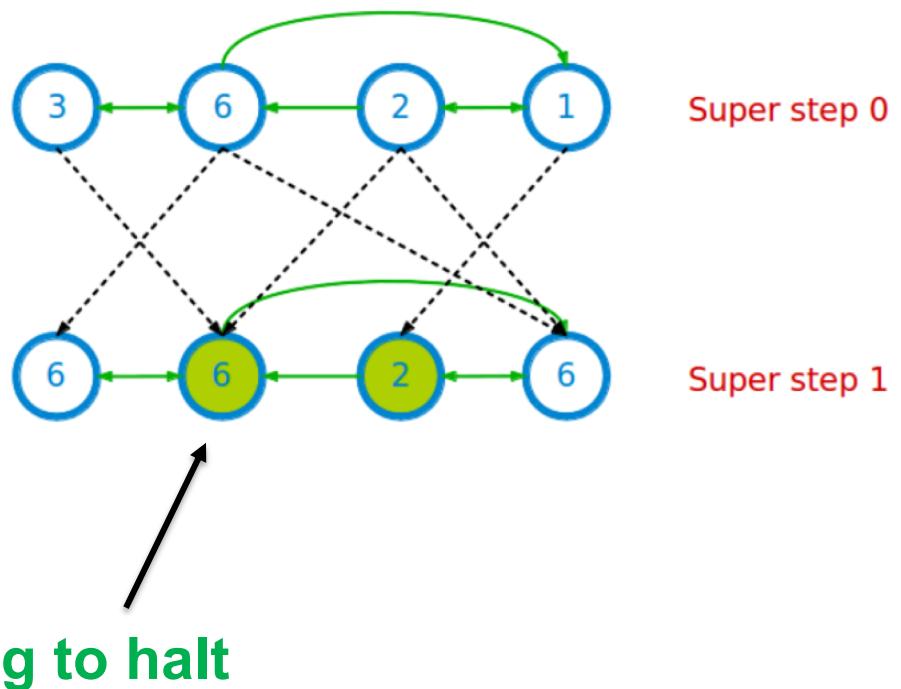
```
Compute(v, messages):
    changed = False
    for m in messages:
        if v.getValue() < m:
            v.setValue(m)
            changed = True
    if changed:
        for each outneighbor w:
            sendMessage(w, v.getValue())
    else:
        voteToHalt()
```



Super step 0

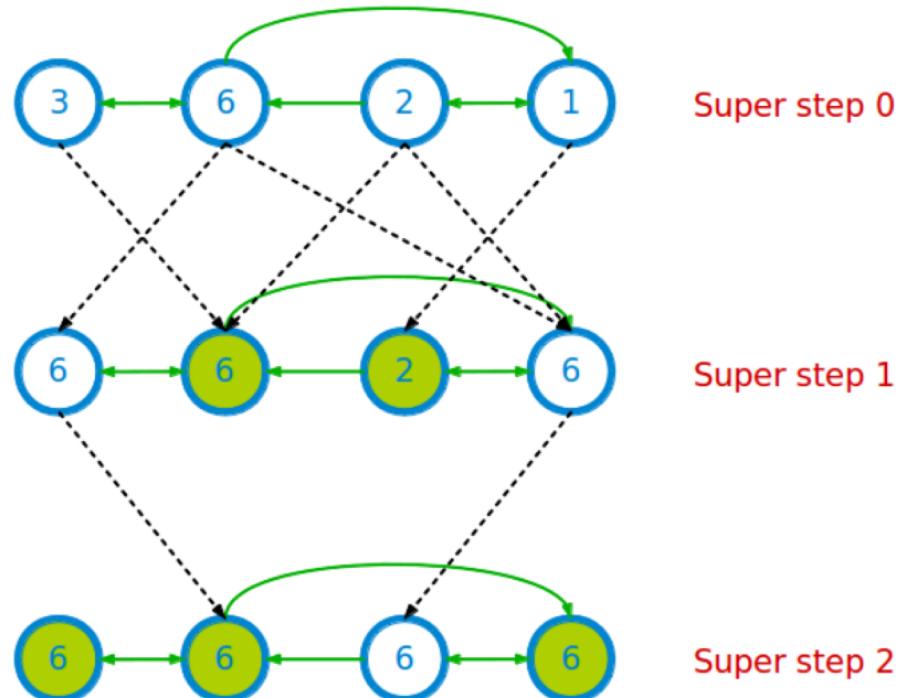
# Example: Computing Max Value

```
Compute(v, messages):  
    changed = False  
    for m in messages:  
        if v.getValue() < m:  
            v.setValue(m)  
            changed = True  
    if changed:  
        for each outneighbor w:  
            sendMessage(w, v.getValue())  
    else:  
        voteToHalt()
```



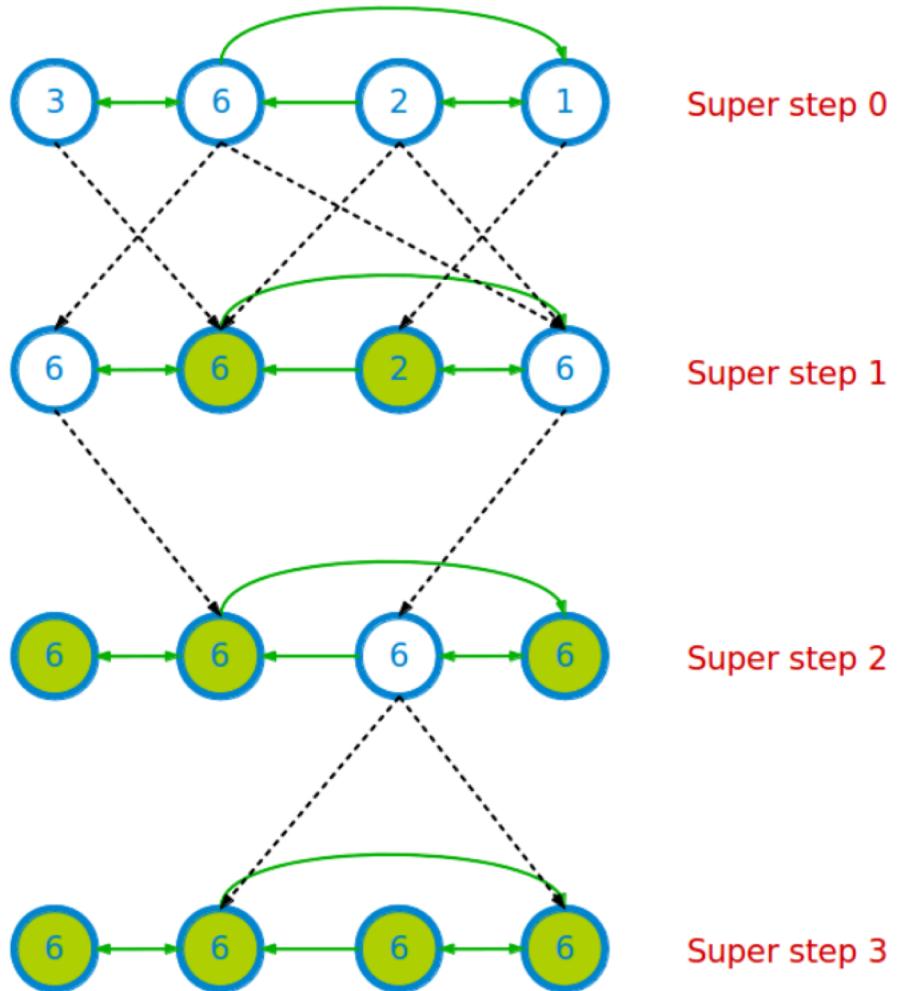
# Example: Computing Max Value

```
Compute(v, messages):  
    changed = False  
    for m in messages:  
        if v.getValue() < m:  
            v.setValue(m)  
            changed = True  
    if changed:  
        for each outneighbor w:  
            sendMessage(w, v.getValue())  
    else:  
        voteToHalt()
```



# Example: Computing Max Value

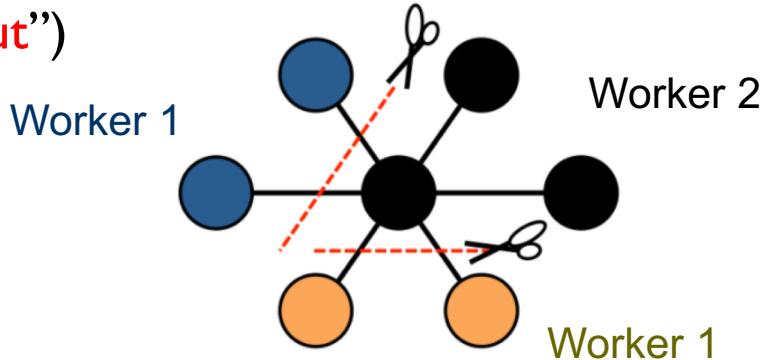
```
Compute(v, messages):  
    changed = False  
    for m in messages:  
        if v.getValue() < m:  
            v.setValue(m)  
            changed = True  
    if changed:  
        for each outneighbor w:  
            sendMessage(w, v.getValue())  
    else:  
        voteToHalt()
```



# Pregel: Implementation

- Master & workers architecture

- Vertices are hash partitioned (by default) and assigned to workers (“**edge cut**”)



- Each worker maintains the state of its portion of the graph **in memory**
- Computations happen in memory
- In each superstep, each worker loops through its vertices and executes `compute()`
- Messages from vertices are sent, either to vertices on the same worker, or to vertices on different workers (**buffered locally** and sent as a batch to reduce network traffic)

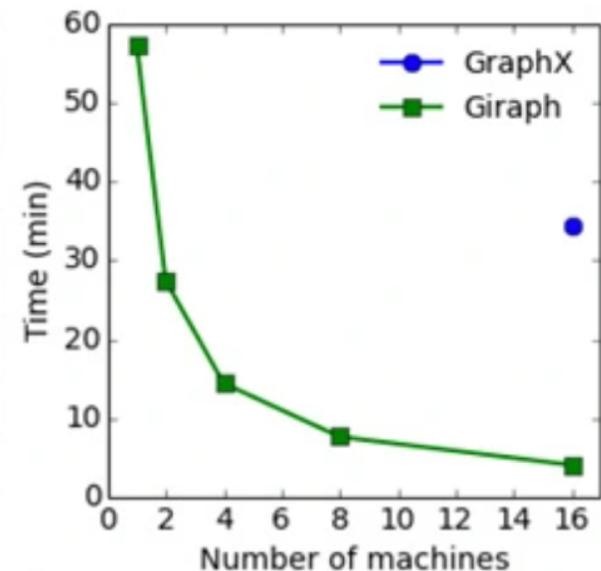
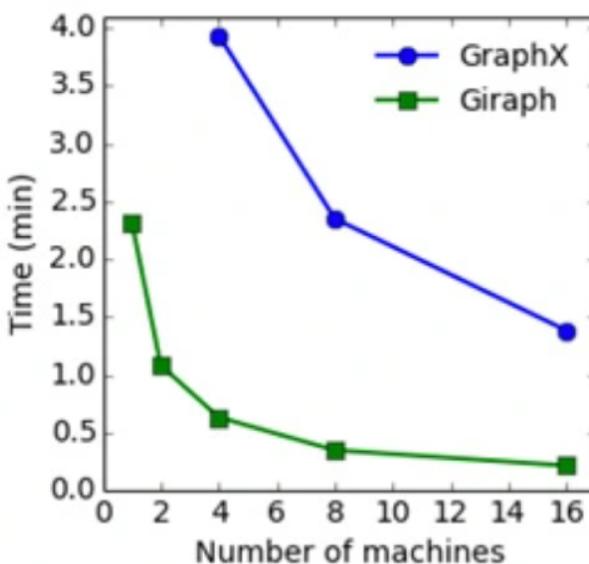
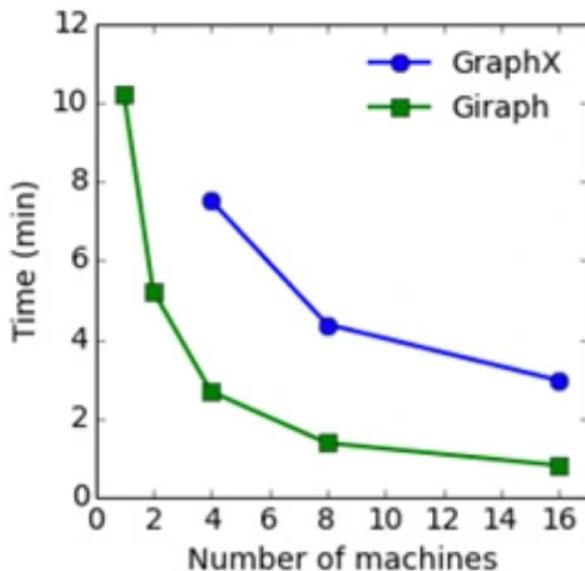
# Efficiency Comparison (on ~ trillion-scale Facebook social graph)

FACEBOOK Engineering



POSTED ON OCTOBER 19, 2016 TO CORE DATA, DATA INFRASTRUCTURE

## A comparison of state-of-the-art graph processing systems



# Giraph Architecture

- Giraph originated as the open-source counterpart to *Pregel*, with several features beyond the basic Pregel model.
- Employs **Hadoop's Map phase** to run computations
- Employs **ZooKeeper** (service that provides distributed synchronization) to enforce barrier waits

# Example (Giraph): Max Value

```
package org.apache.giraph.examples;

public class MaxComputation extends BasicComputation<IntWritable, IntWritable,
NullWritable, IntWritable> {

    @Override
    public void compute(Vertex<IntWritable, IntWritable, NullWritable, IntWritable> vertex,
                        Iterable<IntWritable> messages) throws IOException {

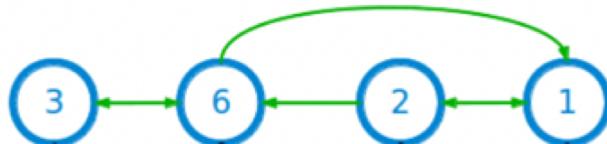
        boolean changed = false;
        for (IntWritable message : messages) {
            if (vertex.getValue().get() < message.get()) {
                vertex.setValue(message);
                changed = true;
            }
        }
        if (getSuperstep() == 0 || changed) {
            sendMessageToAllEdges(vertex, vertex.getValue());
        }
        vertex.voteToHalt(); // reactivation only
                            // after incoming message
    }
}
```

vertex id, vertex data  
edge data, message type

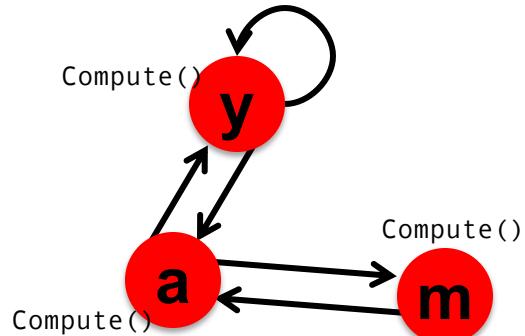
process messages  
from previous superstep

maximum changes

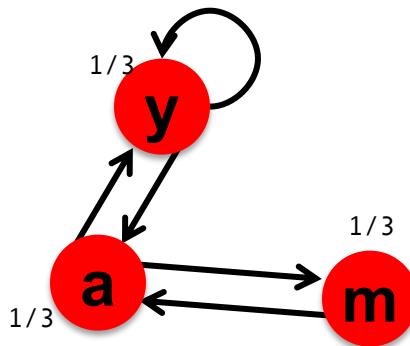
at start or after change,  
message connected vertices



# PageRank in Pregel

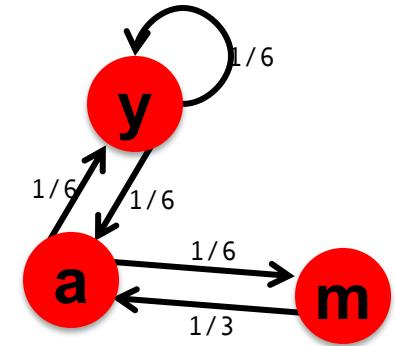


**Superstep 1:**  
Invoke Compute()



**Superstep 1:**  
Aggregate messages from  
neighbors to compute  
PageRank update

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$



**Superstep 1:**  
Send message to  
outgoing neighbors



Repeat this execution  
till Superstep reaches  
30

# PageRank in Pregel

```
class PageRankVertex : public Vertex<double, void, double> {  
public:  
    virtual void Compute(MessageIterator* msgs) {  
        if (superstep() >= 1) {  
            double sum = 0;  
            for (; !msgs->Done(); msgs->Next()) } Compute sum of incoming messages  
                sum += msgs->Value(); } PageRank update  
            *MutableValue() = 0.15 / NumVertices() + 0.85 * sum;  
        }  
  
        if (superstep() < 30) {  
            const int64 n = GetOutEdgeIterator().size();  
            SendMessageToAllNeighbors(GetValue() / n); } Send outgoing messages  
        } else {  
            VoteToHalt(); } Stop after fixed no. of iterations  
    }  
};
```

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

# Acknowledgements

- Amir H. Payberah - Large Scale Graph Processing - Pregel, GraphLab, and Xstream
- <https://engineering.fb.com/2016/10/19/core-data/a-comparison-of-state-of-the-art-graph-processing-systems/>
- <https://spark.apache.org/docs/latest/graphx-programming-guide.html>
- Claudia Hauff - Big Data Processing  
([https://chauff.github.io/documents/bdp/graph\\_giraph.pdf](https://chauff.github.io/documents/bdp/graph_giraph.pdf))