



Markov Decision Processes

Online Algorithms

CS4246/CS5446

AI Planning and Decision Making

Sem 1, AY2021-22



Topics

- Online algorithms for solving MDPs (RN17.2.4)
 - Monte Carlo Tree Search (RN5.4, SB 8.10, 8.11)

Markov Decision Process (MDP)

- Formally:

- An MDP $M \triangleq (S, A, T, R)$ consists of
- A set S of states
- A set A of actions
- A transition function $T: S \times A \times S \rightarrow [0,1]$ such that:

$$\forall s \in S, \forall a \in A: \sum_{s' \in S} T(s, a, s') = \sum_{s' \in S} P(s'|s, a) = 1$$

- A reward function $R: S \rightarrow \mathfrak{R}$
- Solution is a policy – a function to recommend an action in each state: $\pi: S \rightarrow A$
 - Solution involves careful balancing of risk and reward



Handling Curse of Dimensionality

- For large problems:
 - State space grows exponentially with number of variables
 - Value iteration and policy iteration iterate through all states
 - Exponential with number of variables
- New solution approaches:
 - To do **online search** with sampling – decision-time planning
 - Real-time dynamic programming
 - Monte Carlo Tree Search
 - To use **function approximation** of the utility function – compact representation
 - Linear function of features
 - Deep neural networks
 - Etc.



Online Search

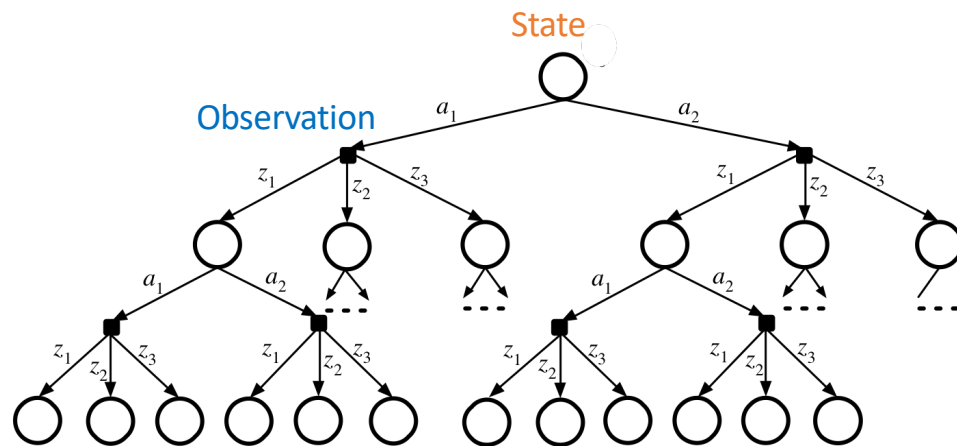
Decision-time planning



Online Algorithms


- Approach: Decision-time planning
 - Significant amount of computation at each decision point, rather than operating primarily with precomputed information
- Methods
 - Real-time dynamic programming
 - Good for mid-size problems.
 - State space with very few repeated states for any manageable set of explored states
 - Simple heuristic for frontier nodes may not be enough to guide well, if rewards are sparse
 - Apply reinforcement learning to generate more accurate heuristics
 - Monte Carlo Tree Search – To look further ahead in the MDP
 - UCT Algorithm for MDP – better suited for large domains, where payoffs go far enough into the future to assess risky potential move

Online Search



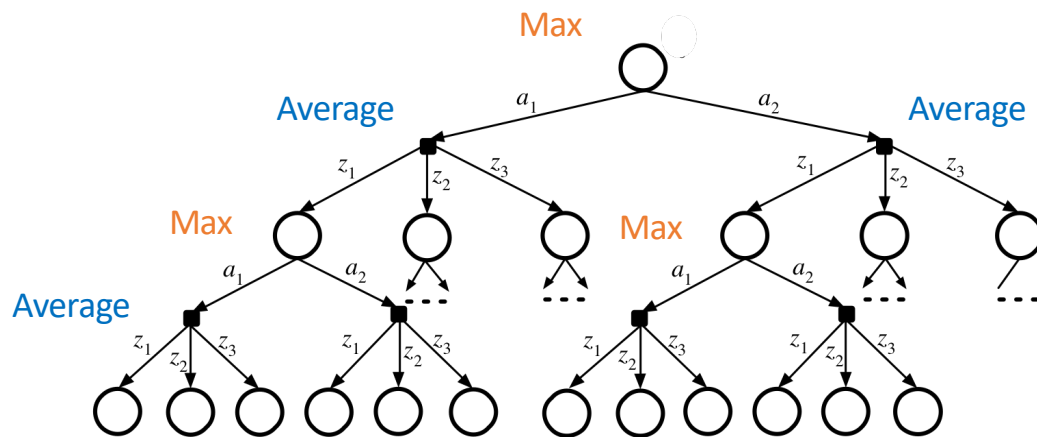
- At every step, construct a search tree.
 - Up to fixed depth D .
 - Root is current state.
 - $|A|$ actions – children of root (and other state nodes).
 - $|Z| = |S|$ children of observation nodes (next state nodes)

Online Search



- To compute value at the root:
 - Initialize leaf with utility estimates (or zeros).
 - At **observation** nodes, compute expected utilities of the children
 - At **state** nodes (where actions can be taken), compute max of the children.

Sparse Sampling



- Approach:

- Tree size is $|A^D| |S^D|$.
- With sparse sampling* [1], estimate by sampling k observations at observation nodes, instead of using all $|S|$ states as possible observations.
- Tree size $|A^D| |k^D|$
- **Question:** Have we solved the curse of dimensionality?

*Michael Kearns, Yishay Mansour, and Andrew Y Ng. "A sparse sampling algorithm for near-optimal planning in large Markov decision processes". In: Machine learning 49.2-3 (2002), pp. 193-208.



Rollout

- Assume:
 - You already have a policy π .
- In rollout, start at state s , try to obtain a policy better than π by:
 - Estimate Q-function $Q(s, a)$ at s by simulating many trajectories from s using each action a where the simulations are done using π .
 - Select action that has the highest average return.
- Improvement:
 - If estimates are accurate enough, the policy improvement theorem implies that rollout improves on π

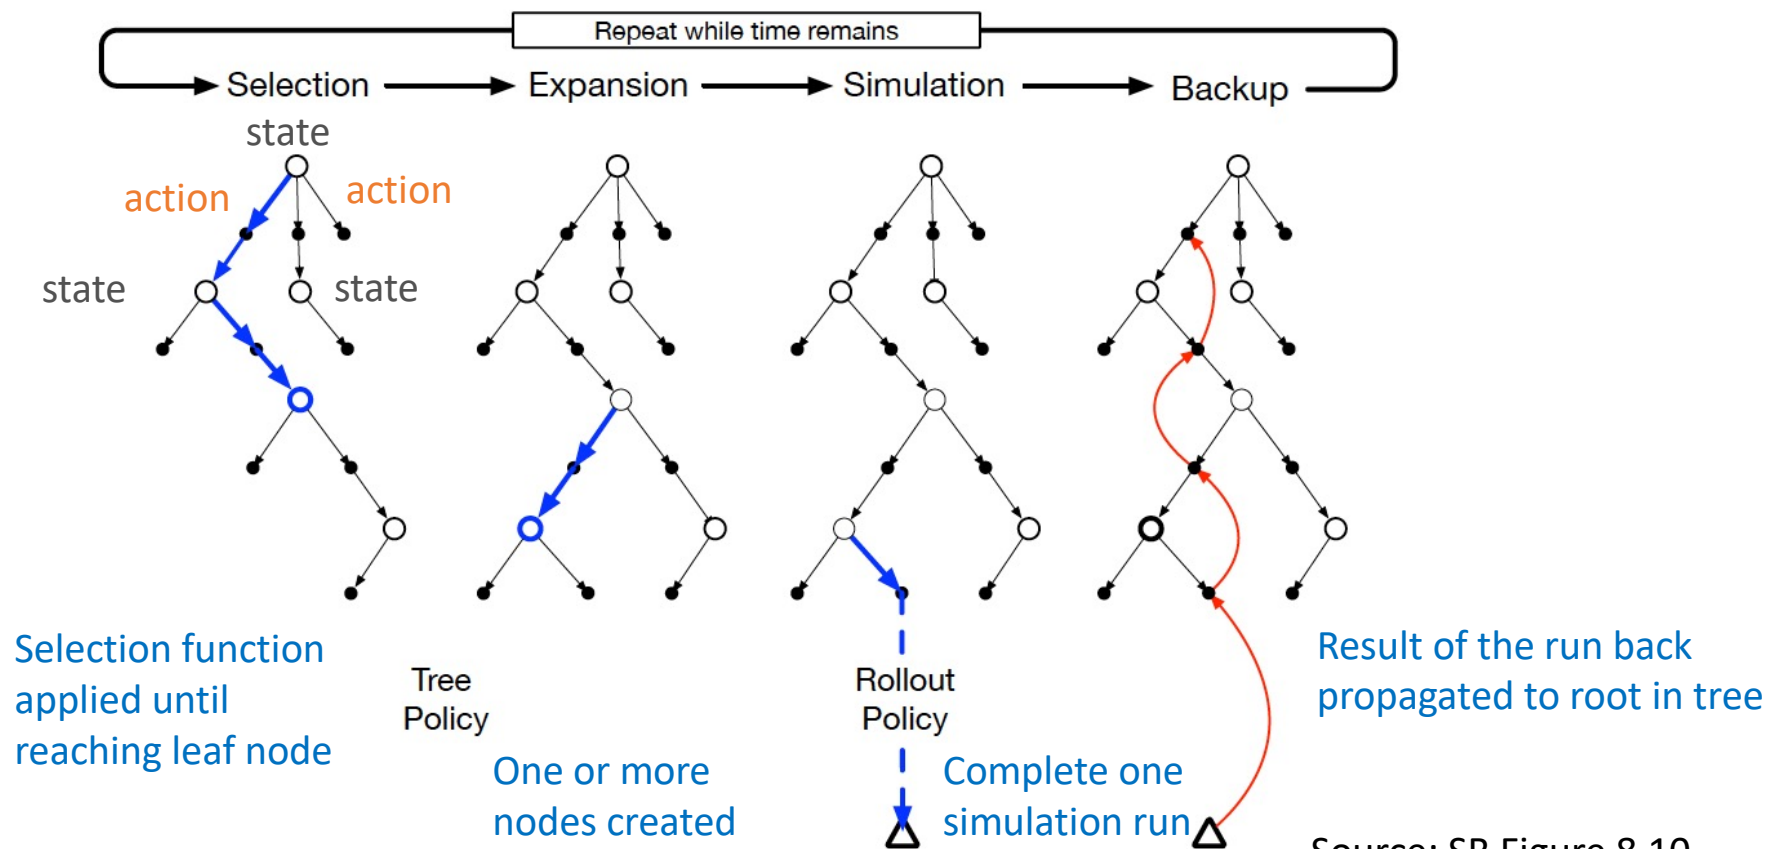


Monte Carlo Tree Search

Online search with simulation

Source: Lee WS, Lecture notes, 2020, MDP p39 -54

Monte Carlo Tree Search



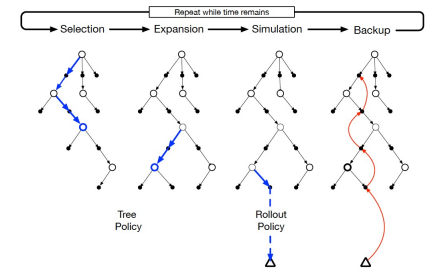
Source: SB Figure 8.10

Monte Carlo Tree Search Algorithm

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action  
  tree  $\leftarrow$  NODE(state)  
  while IS-TIME-REMAINING() do  
    leaf  $\leftarrow$  SELECT(tree)  
    child  $\leftarrow$  EXPAND(leaf)  
    result  $\leftarrow$  SIMULATE(child)  
    BACK-PROPAGATE(result, child)  
  return the move in ACTIONS(state) whose node has highest number of playouts
```

Figure 5.11 The Monte Carlo tree search algorithm. A game tree, *tree*, is initialized, and then we repeat a cycle of SELECT / EXPAND / SIMULATE / BACK-PROPAGATE until we run out of time, and return the move that led to the node with the highest number of playouts.


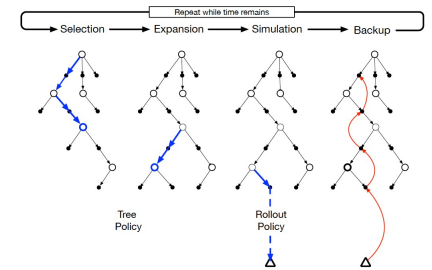
Monte Carlo Tree Search



- Commonly used in MDPs and games*
- Uses both tree search and rollout
 - Rollout at leaf of tree instead, selectively expand tree
 - Repeatedly run trials from the root (current state in online search)
- Trial
 - Repeatedly select node to go to at next level until target depth reached, or
 - Selected node has not been discovered
 - Expansion – create a new node, run a simulation using a rollout policy till required depth
 - Back up the outcomes all the way to the root.
- Anytime policy:
 - When time is up, use the action that looks best at the root at that time.

*Guillaume Chaslot et al. “Monte-Carlo Tree Search: A New Framework for Game AI”. In: AIED. 2008.

MCTS in MDP



- For an MDP:
 - A tree (actually DAG) node n is associated with a state s .
 - A node n' at the next level is selected by applying an action a to s , then sampling the next state s' (corresponding to n') according to $p(s'|s, a)$
- Action selection
 - The action a is selected by balancing exploitation with exploration
- Estimated utility:
 - $\hat{U}(n)$ at a node n is the average return of all the trials at n .
 - The return $r_t(n)$ of trial t starting from n with state s and next node n' is $R(s) + \gamma r_t(n')$.
- Estimated Q-function (action-value function)
 - Estimated Q-function at n , $\hat{Q}(n, a)$ is the average return of all trials at n that starts with action a .
 - $\hat{Q}(r, a)$ at root r used to select the action to take at the root.
- All these are updated in the back-up operation to the root.

Upper Confidence Bounds applied to Trees (UCT)

- Selection policy of UCT¹ [2] algorithm at node n :

$$\pi_{UCT}(n) = \operatorname{argmax}_a (\hat{Q}(n, a) + c \sqrt{\frac{\ln(N(n))}{N(n, a)}})$$

- Where:
 - $\hat{Q}(n, a)$ is the average return of all trials at n that starts with action a
 - $N(n, a)$ is the number of trials through node n that starts with action a
 - $N(n)$ is the number of trials through node n
- **Exploitation term:** $\hat{Q}(n, a)$ – average utility of n that starts with action a
- **Exploration term:** Square-rooted term with count $N(n, a)$ in the denominator
 - Will be high for nodes that have only been explored a few times
 - Will go to zero as the counts increase if (n, a) is selected some non-zero percentage of time
 - Eventually playouts given to the node with the highest utility
- c – Constant balancing exploitation and exploration – often tuned to do well on the problem

¹Levente Kocsis and Csaba Szepesvari. “Bandit based monte-carlo planning”. In: European conference on machine learning. Springer. 2006, pp. 282-293.
Markov Decision Process

Upper Confidence Bounds applied to Trees (UCT)

- Observations:

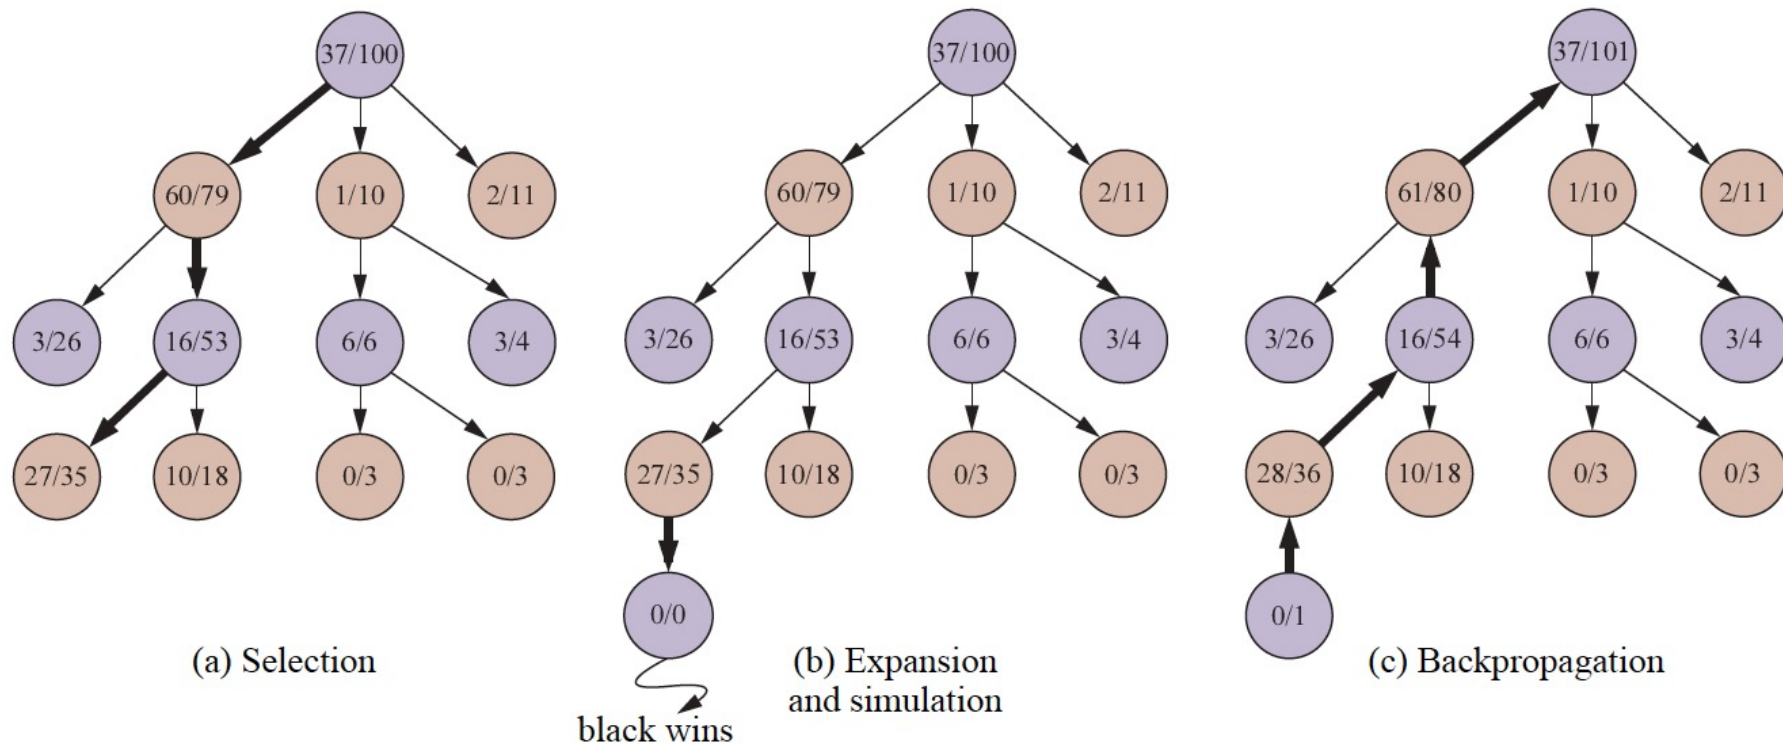
- UCT will eventually converge to the optimal policy with enough trials
- Worst case can be very bad¹ : $\Omega(\overbrace{\exp(\exp(\dots \exp(1) \dots)))^{D-1 \text{ times}})$
- Often works well in practice.
 - PROST Planner² won the ICAPS International Probabilistic Planning Competition for MDP in 2011 and 2014 uses UCT.

¹Coquelin, P.-A. and R. Munos, *Bandit algorithms for tree search*, in *Proceedings of the Twenty-Third Conference on Uncertainty in Artificial Intelligence*. 2007, AUAI Press: Vancouver, BC, Canada. p. 67–74.

²<https://github.com/prost-planner/prost>

Example: Monte Carlo Tree Search in Game

Using upper confidence bounds applied to trees (UCT) selection metric





MCTS in practice

- Visualizing MCTS: Player 2 uses MCTS
 - https://www.youtube.com/watch?v=FvRSxNLTg7U&ab_channel=DaveDyer
- Playing Super Mario:
 - https://www.youtube.com/watch?v=HRiEUUC9TUA&ab_channel=Emil
- AlphaGo Zero
 - <https://youtu.be/tXlM99xPQC8>
 - Toward breakthrough in AI



AlphaGo Zero

<https://youtu.be/tXlM99xPQC8>

David Silver et al. Mastering the game of Go without human knowledge". In: Nature 550.7676 (2017), p. 354.



AlphaGo Zero

- Go Game
 - Space size of about 10^{170} with branching factor that starts at 361
 - Difficult to define good evaluation function
 - Need function approximation to represent value and policy functions
- AlphaGo Zero* [3] uses combination of MCTS and policy iteration
- For playout policy:
 - AlphaGo Lee (which defeated Lee Sedol) used a combination of expert games as well as self play.
 - AlphaGo Zero uses only self-play (only provided with the rules of Go); defeated AlphaGo Lee 100-0.
- Use deep neural network with two “heads”
 - Value head – outputs real value estimate of the value function
 - Policy head – outputs vector of size 190×190
 - Each component represents the probability that the policy will play that board position.

*David Silver et al. Mastering the game of Go without human knowledge". In: Nature 550.7676 (2017), p. 354.



AlphaGo Zero

- For (action) selection policy:
 - Variant of UCT that exploits policy head output of neural network $P(s, a)$

$$\pi_{UCT}(s) = \operatorname{argmax}_a \hat{Q}(s, a) + cP(s, a) \sqrt{\frac{\sum_b N(s, b)}{1 + N(s, a)}}$$

- When a leaf node is reached, the value head of the neural network is used to evaluate the state instead of doing a roll-out (simulation).
- Go is a zero-sum turn taking game instead of an MDP:
 - Search alternates between:
 - Selecting action that maximizes when it is first player's turn
 - Selecting action that minimizes (multiply estimate by -1 then maximize) for second player's turn
 - At termination: reward $+1$ for first player win and -1 for second player win.



AlphaGo Zero

- Approximate policy iteration with self-play
 - Policy iteration has 2 stages: policy evaluation and policy improvement
 - AlphaGo Zero does both using supervised learning
 - With current value function, MCTS viewed as policy improvement operator gives improved policy values for evaluated states
 - Improve policy for a set of states
 - Self-play with search gives the policy evaluation for the evaluated states
 - Evaluate value of improved policy
 - Supervised learning is used to interpolate values and policy over whole domain using data from a set of states.



Homework

- Readings

- [RN] 17.2.4, 5.4 (Online algorithms, MCTS)
- [SB] 8.10, 8.11 (Online algorithms, MCTS)
- [SB] Sutton, R. S. and A. G. Barto. Reinforcement Learning: An introduction. 2nd ed. MIT Press, 2018, 2020
[Book website: <http://incompleteideas.net/book/the-book.html>]
[e-Book for personal use:
<http://incompleteideas.net/book/RLbook2020.pdf>]



References

- Sparse Sampling:

1. Michael Kearns, Yishay Mansour, and Andrew Y Ng. A sparse sampling algorithm for near-optimal planning in large Markov decision processes". In: Machine learning 49.2-3 (2002), pp. 193–208.

- The UCT algorithm:

2. Kocsis, L. and C. Szepesvári, Bandit based Monte-Carlo planning, in Proceedings of the 17th European conference on Machine Learning. 2006, Springer-Verlag: Berlin, Germany. p. 282–293.

- AlphaGo Zero:

3. David Silver et al. Mastering the game of Go without human knowledge". In: Nature 550.7676 (2017), p. 354.