

CS4246 / CS5446

Tutorial Week 10

Muhammad **Rizki** Maulana
rizki@u.nus.edu

First

Review


- Linear Regression
- Neural Networks
- Gradient Descent
- Common Issues

Linear Regression

Single variable

$$f(x) = wx + b$$

(linear) function approximator



Linear Regression


Single variable

$$f(x) = wx + b$$

Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$



Linear Regression


Single variable

$$f(x) = wx + b$$

Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \dots + w_F x_F + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$



Linear Regression

Single variable


$$f(x) = wx + b$$

Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \dots + w_F x_F + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$

Goal: minimize $\sum_{(x,y)} [f(x) - y]^2$



Linear Regression


Single variable

$$f(x) = wx + b$$

Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \dots + w_F x_F + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$



Goal: minimize $\sum_{(x,y)} [f(x) - y]^2$

Squared distance

Linear Regression

Single variable

$$f(x) = wx + b$$


Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \dots + w_F x_F + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$

Goal: minimize $\sum_{(x,y)} [f(x) - y]^2$

Squared distance



Linear Regression


Single variable

$$f(x) = wx + b$$

Multiple variables:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \dots + w_F x_F + b$$

$$\text{where } \mathbf{w} = [w_1, \dots, w_F]^T, \mathbf{x} = [x_1, \dots, x_F]^T$$



Goal: minimize $\sum_{(x,y)} [f(x) - y]^2 = L_{(x,y)}(f(x), y)$

Generally: loss function

Neural Networks (NN)




(another) function approximator

Neural Networks (NN)

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

Recall our linear function



where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$

Neural Networks (NN)

Simple NN (Neuron)

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

Now it goes **neural**!



where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$

Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function


Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

Features,
not data points!




Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function



Neural Networks (NN)


Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

“Deep” NN

$$f(\mathbf{x}) = f_{i1}(\mathbf{x})$$



Neural Networks (NN)


Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

“Deep” NN

$$f(\mathbf{x}) = f_{l1}(\mathbf{x}) \dots$$



Neural Networks (NN)


Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

“Deep” NN

$$f(\mathbf{x}) = f_{H1}([f_{H1}(\mathbf{x}) \dots]^T), \dots, f_{HM}([\dots]^T)$$



Neural Networks (NN)


Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

“Deep” NN

$$f(\mathbf{x}) = f_O([f_{H1}([f_{H1}(\mathbf{x}) \dots]^T), \dots, f_{HM}([\dots]^T)]^T)$$



Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$


where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

Convolutional NN

“building blocks”

“Deep” NN

$$f(\mathbf{x}) = f_O([f_{H1}([f_{I1}(\mathbf{x}) \dots]^T), \dots, f_{HM}([\dots]^T)]^T)$$



Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function


Convolutional NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T * \mathbf{x})$$

where $\mathbf{w} = [w_1, \dots, w_K]^T$, $K \leq F$

“Deep” NN

$$f(\mathbf{x}) = f_O([f_{H1}([f_{I1}(\mathbf{x}) \dots]^T), \dots, f_{HM}([\dots]^T)]^T)$$



Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$,
 φ non-linear activation function

Convolutional NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T * \mathbf{x})$$


where $\mathbf{w} = [w_1, \dots, w_K]^T$, $K \leq F$

\mathbf{x}

$f(\mathbf{x})$

“Deep” NN

$$f(\mathbf{x}) = f_O([f_{H1}([f_{I1}(\mathbf{x}) \dots]^T), \dots, f_{HM}([\dots]^T)]^T)$$



Neural Networks (NN)

Simple NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

where $\mathbf{w} = [w_1, \dots, w_F]^T$, $\mathbf{x} = [x_1, \dots, x_F]^T$

φ non-linear activation func


Convolutional NN

$$f(\mathbf{x}) = \varphi(\mathbf{w}^T * \mathbf{x})$$


where $\mathbf{w} = [w_1, \dots, w_K]^T$, $K \leq F$

“Deep” NN

$$= f_O([f_{H1}([f_{H1}(\mathbf{x})]^\top), \dots, f_{HM}([\dots]^\top)]^\top)$$



Goal: Minimize Loss



Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

Minimizing Loss

Consider:

- Loss = $\sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$



The only independent variable!

Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$


Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$




Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$



Minimizing Loss


Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"



Minimizing Loss


Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"



Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$


How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"

How do we know which one?

(should be general enough for any function)



Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?


$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"

How do we know which one?

$$dL_{(x,y)} / dw = 2[wx - y]x$$

"too small" \rightarrow gradient -ve
"too large" \rightarrow gradient +ve



Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"


How do we know which one?

$$dL_{(x,y)} / dw = 2[wx - y]x$$

"too small" \rightarrow gradient -ve
"too large" \rightarrow gradient +ve

Update:

$$w_t = w_{t-1} - \text{gradient}$$



Minimizing Loss

Consider:

- $\text{Loss} = \sum_{(x,y)} [f(x) - y]^2$
- $f(x) = wx$

How to find w such that we get minimum loss?

$$L_{(x,y)} = [f(x) - y]^2 = [wx - y]^2$$

$wx < y \rightarrow$ "too small"
 $wx > y \rightarrow$ "too large"

How do we know which one?


$$dL_{(x,y)} / dw = 2[wx - y]x$$

"too small" \rightarrow gradient -ve
"too large" \rightarrow gradient +ve


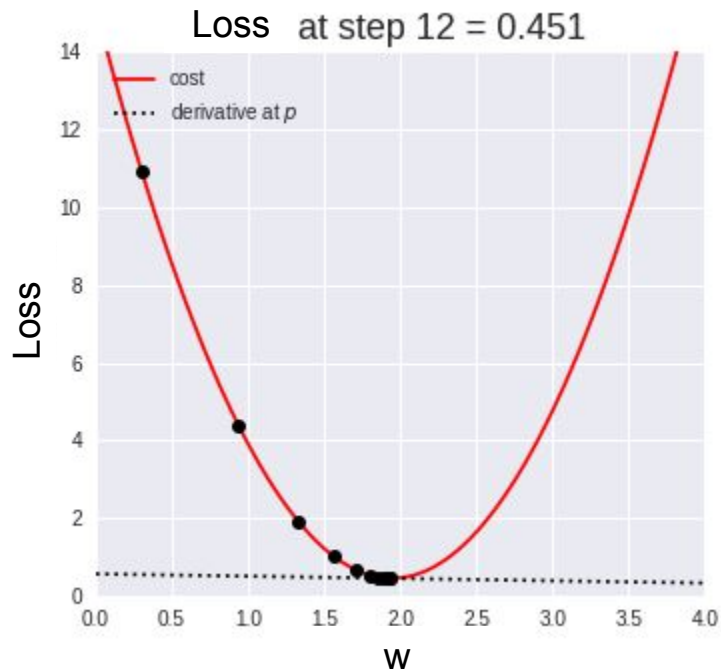
Update:

$$w_t = w_{t-1} - \text{gradient (do in a loop)}$$

\leftarrow Gradient Descent



Gradient Descent: Illustration



Gradient Descent: Backpropagation

What about complicated functions (e.g., neural networks)?

Gradient Descent: Backpropagation

What about complicated functions (e.g., neural networks)?

Backpropagation = Chain Rule of Calculus


$$\frac{da}{dx} a(b(c(d(e(f(g(x)))))))$$

$$\frac{da}{\partial x} = \frac{da}{db} \times \frac{db}{dc} \times \frac{dc}{dd} \times \frac{dd}{de} \times \frac{de}{df} \times \frac{df}{dg} \times \frac{dg}{dx}$$

Gradient Descent: Issues with Deep Neural Networks


Gradient Descent: Issues with Deep Neural Networks

Local Minima




Gradient Descent: Issues with Deep Neural Networks

Local Minima



Vanishing Gradient



Issue with Function Approximation in RL

The Deadly Triad in RL:

- **Function Approximation**
- Bootstrapping
- Off-policy Learning

Second

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g)^2 + (y - y_g)^2}$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g)^2 + (y - y_g)^2} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(R(s) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right)^2$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g)^2 + (y - y_g)^2} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s') + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g)^2 + (y - y_g)^2} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s) + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Question

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g) + (y - y_g)}$$

$$s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s)}_{\text{Target}} + \gamma \underbrace{\hat{U}_{\theta}(s') - \hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Target Prediction

$$\frac{\partial L}{\partial \theta_i} = \left(R(s) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-\frac{\partial \hat{U}_{\theta}}{\partial \theta_i} \right]$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g) + (y - y_g)} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s_t) + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Target Prediction

$$\frac{\partial L}{\partial \theta_i} = \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[- \frac{\partial \hat{U}_{\theta}}{\partial \theta_i} \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_0} = 1$$

$$\theta_0 \leftarrow \theta_0 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-1 \right]$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g) + (y - y_g)} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s_t) + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Target Prediction

$$\frac{\partial L}{\partial \theta_i} = \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[- \frac{\partial \hat{U}_{\theta}}{\partial \theta_i} \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_0} = 1$$

$$\theta_0 \leftarrow \theta_0 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-1 \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_1} = x$$

$$\theta_1 \leftarrow \theta_1 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-x \right]$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g) + (y - y_g)} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s_t) + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Target Prediction

$$\frac{\partial L}{\partial \theta_i} = \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[- \frac{\partial \hat{U}_{\theta}}{\partial \theta_i} \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_0} = 1$$

$$\theta_0 \leftarrow \theta_0 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-1 \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_1} = x$$

$$\theta_1 \leftarrow \theta_1 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-x \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_2} = y$$

$$\theta_2 \leftarrow \theta_2 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-y \right]$$

[RN 21.4] Write out the parameter update equations for TD learning with

$$\hat{U}_{\theta}(x, y) = \theta_0 + \theta_1 x + \theta_2 y + \theta_3 \sqrt{(x - x_g) + (y - y_g)} \quad s = (x, y)$$

$$L = \frac{1}{2} \left(\underbrace{R(s_t) + \gamma \hat{U}_{\theta}(s')}_{\text{Target}} - \underbrace{\hat{U}_{\theta}(s)}_{\text{Prediction}} \right)^2 \quad \theta_i \leftarrow \theta_i - \alpha \frac{\partial L}{\partial \theta_i}$$

Target Prediction

$$\frac{\partial L}{\partial \theta_i} = \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[- \frac{\partial \hat{U}_{\theta}}{\partial \theta_i} \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_0} = 1$$

$$\theta_0 \leftarrow \theta_0 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-1 \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_1} = x$$

$$\theta_1 \leftarrow \theta_1 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-x \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_2} = y$$

$$\theta_2 \leftarrow \theta_2 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-y \right]$$

$$\frac{\partial \hat{U}_{\theta}}{\partial \theta_3} = \sqrt{(x - x_g) + (y - y_g)}$$

$$\theta_3 \leftarrow \theta_3 - \alpha \left(R(s_t) + \gamma \hat{U}_{\theta}(s') - \hat{U}_{\theta}(s) \right) \left[-\sqrt{(x - x_g) + (y - y_g)} \right]$$

Third

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a)).$$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a)).$$

- i. First observed transition: initial value of $x = 0$, observed reward $r = 10$, action a_1 , next state $x = 1$.

Question

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a)).$$

- i. First observed transition: initial value of $x = 0$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(0, a_1) \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0) = 5$$

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 5 | 0 |
| a_2 | 0 | 0 |

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a)).$$

- i. First observed transition: initial value of $x = 0$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(0, a_1) \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0) = 5$$

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 5 | 0 |
| a_2 | 0 | 0 |

- ii. Second observed transition: from $x = 1$, observed reward $r = -5$, action a_2 , next state $x = 0$.

Question

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (a) Perform two steps of Q-learning with the observed transitions shown below in (i) and (ii) using a table representation of the Q-function. Use a learning rate $\alpha = 0.5$ starting from a table with all entries initialized to 0. Show the Q-function after each step.

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 0 | 0 |
| a_2 | 0 | 0 |

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a)).$$

- i. First observed transition: initial value of $x = 0$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(0, a_1) \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0) = 5$$

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 5 | 0 |
| a_2 | 0 | 0 |

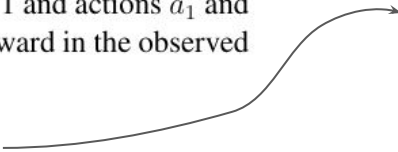
- ii. Second observed transition: from $x = 1$, observed reward $r = -5$, action a_2 , next state $x = 0$.

$$Q(1, a_2) \leftarrow 0 + 0.5(-5 + 0.9 \max(5, 0) - 0) = -0.25$$

| Q | $x = 0$ | $x = 1$ |
|-------|---------|---------|
| a_1 | 5 | 0 |
| a_2 | 0 | -0.25 |

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

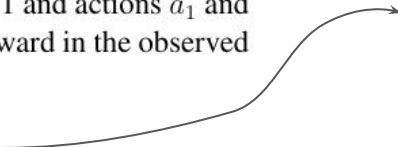
(b) Now perform Q-learning with function approximation using

- 
- $Q(x, a_1) = \beta_1 x$
 - $Q(x, a_2) = \beta_2 x$
 - $\alpha = 0.5$
 - $\beta_1 = 0$
 - $\beta_2 = 0$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

(b) Now perform Q-learning with function approximation using

$$\beta_i \leftarrow \beta_i + \alpha(R(x) + \gamma \max_{a'} Q(x', a') - Q(x, a)) \frac{\partial Q(x, a)}{\partial \beta_i}$$

- 
- $Q(x, a_1) = \beta_1 x$
 - $Q(x, a_2) = \beta_2 x$
 - $\alpha = 0.5$
 - $\beta_1 = 0$
 - $\beta_2 = 0$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

(b) Now perform Q-learning with function approximation using

$$\beta_i \leftarrow \beta_i + \alpha (R(x) + \gamma \max_{a'} Q(x', a') - Q(x, a)) \frac{\partial Q(x, a)}{\partial \beta_i}$$

- $Q(x, a_1) = \beta_1 x$
- $Q(x, a_2) = \beta_2 x$
- $\alpha = 0.5$
- $\beta_1 = 0$
- $\beta_2 = 0$

- i. First observed transition: initial value of $x = 1$, observed reward $r = 10$, action a_1 , next state $x = 1$.

Question

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

(b) Now perform Q-learning with function approximation using

$$\beta_i \leftarrow \beta_i + \alpha(R(x) + \gamma \max_{a'} Q(x', a') - Q(x, a)) \frac{\partial Q(x, a)}{\partial \beta_i}$$

- $Q(x, a_1) = \beta_1 x$

- $Q(x, a_2) = \beta_2 x$

- $\alpha = 0.5$

- $\beta_1 = 0$

- $\beta_2 = 0$

i. First observed transition: initial value of $x = 1$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(x, a_1) = \beta_1 x$$

$$\beta_1 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)1 = 5$$

$$\beta_2 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)0 = 0$$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

(b) Now perform Q-learning with function approximation using

$$\beta_i \leftarrow \beta_i + \alpha (R(x) + \gamma \max_{a'} Q(x', a') - Q(x, a)) \frac{\partial Q(x, a)}{\partial \beta_i}$$

- $Q(x, a_1) = \beta_1 x$
- $Q(x, a_2) = \beta_2 x$
- $\alpha = 0.5$
- $\beta_1 = 0$
- $\beta_2 = 0$

i. First observed transition: initial value of $x = 1$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(x, a_1) = \beta_1 x$$

$$\beta_1 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)1 = 5$$

$$\beta_2 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)0 = 0$$

ii. Second observed transition: from $x = 1$, observed reward $r = -5$, action a_2 , next state $x = 0$.

Question

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

(b) Now perform Q-learning with function approximation using

$$\beta_i \leftarrow \beta_i + \alpha(R(x) + \gamma \max_{a'} Q(x', a') - Q(x, a)) \frac{\partial Q(x, a)}{\partial \beta_i}$$

$$\bullet Q(x, a_1) = \beta_1 x$$

$$\bullet Q(x, a_2) = \beta_2 x$$

$$\bullet \alpha = 0.5$$

$$\bullet \beta_1 = 0$$

$$\bullet \beta_2 = 0$$

i. First observed transition: initial value of $x = 1$, observed reward $r = 10$, action a_1 , next state $x = 1$.

$$Q(x, a_1) = \beta_1 x$$

$$\beta_1 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)1 = 5$$

$$\beta_2 \leftarrow 0 + 0.5(10 + 0.9 \max(0, 0) - 0)0 = 0$$

ii. Second observed transition: from $x = 1$, observed reward $r = -5$, action a_2 , next state $x = 0$.

$$Q(x, a_2) = \beta_2 x$$

$$\beta_1 \leftarrow 5 + 0.5(-5 + 0.9 \max(0, 0) - 0)0 = 5$$

$$\beta_2 \leftarrow 0 + 0.5(-5 + 0.9 \max(0, 0) - 0)1 = -2.5$$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (c) After enough data is observed, which method would give better performance, the tabular method in (a) or the function approximation method in (b)? Why? Suggest how the poorer performing method can be improved.

Question

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (c) After enough data is observed, which method would give better performance, the tabular method in (a) or the function approximation method in (b)? Why? Suggest how the poorer performing method can be improved.

Tabular is better

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (c) After enough data is observed, which method would give better performance, the tabular method in (a) or the function approximation method in (b)? Why? Suggest how the poorer performing method can be improved.

Tabular is better

Failure case of the function approximation

$$Q(x, a_1) = \beta_1 x \text{ and } Q(x, a_2) = \beta_2 x$$

Can't have non-zero value for $x=0$

Consider a system with a single state variable x that can take value 0 or 1 and actions a_1 and a_2 . An agent can observe the value of the state variable as well as the reward in the observed state. Assume a discount factor $\gamma = 0.9$.

- (c) After enough data is observed, which method would give better performance, the tabular method in (a) or the function approximation method in (b)? Why? Suggest how the poorer performing method can be improved.

Tabular is better

Failure case of the function approximation

$$Q(x, a_1) = \beta_1 x \text{ and } Q(x, a_2) = \beta_2 x$$

Can't have non-zero value for $x=0$

Improvements

$$Q(x, a_1) = \beta_1 x + \delta_1 \text{ and } Q(x, a_2) = \beta_2 x + \delta_2$$

Question?

<EOF>