

# Intro to Connectionist Machine Learning

CS4248 Natural Language Processing

Week 05

Shreyas Kuthanoor PRAKESH and Min-Yen KAN

5

Slides from NUS CS3244 and Dan Jurafsky (Stanford)

$$\frac{d}{dx} \ln x = \frac{1}{x}, \quad \frac{d\sigma(z)}{dz} = \sigma(z)(1-\sigma(z)), \quad f(x) = \text{ulv}(x), \quad \frac{df(x)}{dx} = \frac{dy}{dw} \cdot \frac{dw}{dx}$$

## Week 04 Agenda

Text Classification

Case Study: Sentiment Analysis

TF-IDF

Vector Space Model

Naïve Bayes  
and a Runthrough (time permitting)

Evaluating Text Classification

$$\frac{\partial L_{CE}}{\partial w_{ij}}$$

$$= \frac{\partial}{\partial w_{ij}} \left[ - \left[ y \log \sigma(w \cdot x + b) + (1-y) (\log(1 - \sigma(w \cdot x + b))) \right] \right]$$

$$= - \left[ \frac{\partial}{\partial y} y \log \sigma(w \cdot x + b) + \frac{\partial}{\partial y} (1-y) (\log(1 - \sigma(w \cdot x + b))) \right]$$

$$\frac{\partial L_{CE}}{\partial w_{ij}} = \frac{-y}{\sigma(w \cdot x + b)} \frac{\partial}{\partial y} \sigma(w \cdot x + b) - \frac{1-y}{1 - \sigma(w \cdot x + b)} \frac{\partial}{\partial y} (1 - \sigma(w \cdot x + b))$$

$$\frac{\partial L_{CE}}{\partial w_{ij}} = - \left[ \frac{y}{\sigma(w \cdot x + b)} - \frac{1-y}{1 - \sigma(w \cdot x + b)} \right] \frac{\partial}{\partial y} \sigma(w \cdot x + b)$$



School of  
Computing

$$\begin{aligned} \frac{\partial L_{CE}}{\partial w_j} &= - \left[ \frac{y - \sigma(wx+b)}{\sigma(wx+b)[1-\sigma(wx+b)]} \right] \sigma(wx+b)[1-\sigma(wx+b)] \frac{\partial (wx+b)}{\partial w_j} \\ &= - \left[ \frac{y - \sigma(wx+b)}{\sigma(wx+b)[1-\sigma(wx+b)]} \right] \sigma(wx+b)[1-\sigma(wx+b)] x_j \\ &= - [y - \sigma(wx+b)] x_j \\ &= [\sigma(wx+b) - y] x_j \end{aligned}$$



School of  
Computing

## Week 05 Agenda

Generative vs. Discriminative Classifiers

Classification with Logistic Regression  
and a Runthrough

Cross Entropy

Stochastic Gradient Descent

LR as a Probabilistic ML Classifier

Regularization

XOR

Neural Networks

connectionless classification

# Generative vs. Discriminative Classifiers

*Slide Credits: Dan Jurafsky (Stanford)*

# Logistic Regression

Important analytic tool in natural and social sciences.

Baseline supervised machine learning tool for classification.

It's also the foundation of neural networks.

*Slide Credits: Dan Jurafsky (Stanford)*

# Generative vs. Discriminative Classifiers

**Naïve Bayes** is a generative classifier

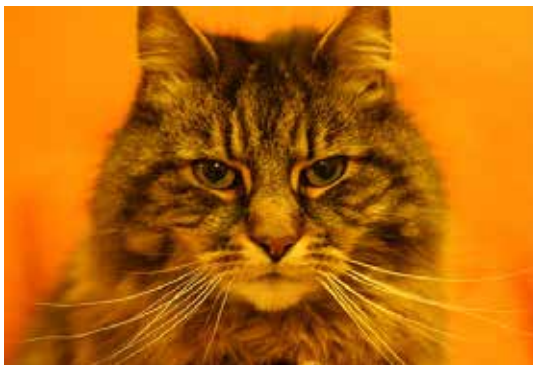
But in contrast:

**Logistic Regression** is a discriminative classifier

*Slide Credits: Dan Jurafsky (Stanford)*

# What are Generative and Discriminative Classifiers?

Suppose we're distinguishing cat from dog images:



*Photos sourced from ImageNet. Slide Credits: Dan Jurafsky (Stanford)*

# Generative Classifier

Build a model of what's in a cat image

- Knows about whiskers, ears, eyes
- Assigns a probability to any image:
  - How cat-y is this image?
  - Also build a model for dog images



Given a new image at test time:

Run both models and see which one fits better

*Photos sourced from ImageNet. Slide Credits: Dan Jurafsky (Stanford)*



# Discriminative Classifier

Just tries to distinguish dogs from cats.  $p(y|x)$

→ find the right class.



Oh look, dogs have collars! Don't need anything else.

*Photos sourced from ImageNet. Slide Credits: Dan Jurafsky (Stanford)*

# Classifying $y$ given document $x$ in Generative vs Discriminative Classifiers

Naïve Bayes - strong independence assumptions, multiply  $f_1$  and  $f_2$ .

$$\hat{y} = \operatorname{argmax} P(x|y)P(y)$$

Logistic Regression - weight distributed to two correlated features  $f_1$  and  $f_2$ . Works

$$\hat{y} = \operatorname{argmax} P(y|x)$$

Slide Credits: Dan Jurafsky (Stanford)

# Components of a probabilistic machine learning classifier

Given  $m$  input/output pairs  $(\mathbf{x}^{(j)}, y^{(j)})$ :

1. A feature representation of the input.  $\mathbf{x}^{(j)} = [x_1, x_2, \dots, x_i, \dots, x_n]$ .  
 The  $i$ th feature is denoted  $x_i$ , or more completely  $x_i^{(j)}$ , sometimes  $f_i(\mathbf{x})$ .
2. A classification function that computes  $\hat{y}$ , the estimated class, via  $p(y|\mathbf{x})$ , like the **sigmoid** or **softmax** functions.
3. An objective function for learning, like **cross-entropy** loss.
4. An algorithm for optimizing the objective function: **stochastic gradient descent**.

need  
all  
 $\phi$ .


Slide Credits: Dan Jurafsky (Stanford)

# Classification with Logistic Regression

want to find optimal weights-

# The Two Phases of Logistic Regression

**Training:** we learn weights  $\theta$  and bias  $b$  using **stochastic gradient descent** and **cross-entropy loss**.

 **Notation Varies:** Weights are also called *parameters*, sometime denoted as  $w$  (as used in the SLP3 textbook).

Why is  $\theta$  separate from  $b$ ? Mull on that.

**Test:** Given a test example  $x$ , we compute  $p(y|x)$  using learned weights  $\theta$  and bias  $b$ , and return whichever label ( $y = 1$  or  $y = 0$ ) has higher probability.

Slide Credits: Dan Jurafsky (Stanford)

# Logistic Regression: Weighted Features

For feature  $x_i$ , weight  $\theta_i$  tells is how important is  $x_i$ :

- $x_1$  = “review contains *awesome*”:  $\theta_1 = +10$
- $x_2$  = “review contains *abysmal*”:  $\theta_2 = -10$
- $x_3$  = “review contains *mediocre*”:  $\theta_3 = -2$

bias

weight

We'll sum up all the weighted features and the bias:

$$z = \left( \sum_{i=1}^n \theta_i x_i \right) + b$$

$$z = \theta \cdot x + b$$

If this sum is “high”, we say  $\hat{y} = 1$ ; if low, then  $\hat{y} = 0$ .

What about the bias  $b$ ?  
What does that correspond to?

Can also view  $b$  as a form of  $w_0$  for an  $x_0$  that is always observed.

Then what happens to these formulas?

linear regression. need to make into probability.

Slide Credits: Dan Jurafsky (Stanford)

# But we want a probabilistic classifier...

We need to formalize “sum is high”.

We’d want a principled classifier that gives us a probability, like Naïve Bayes

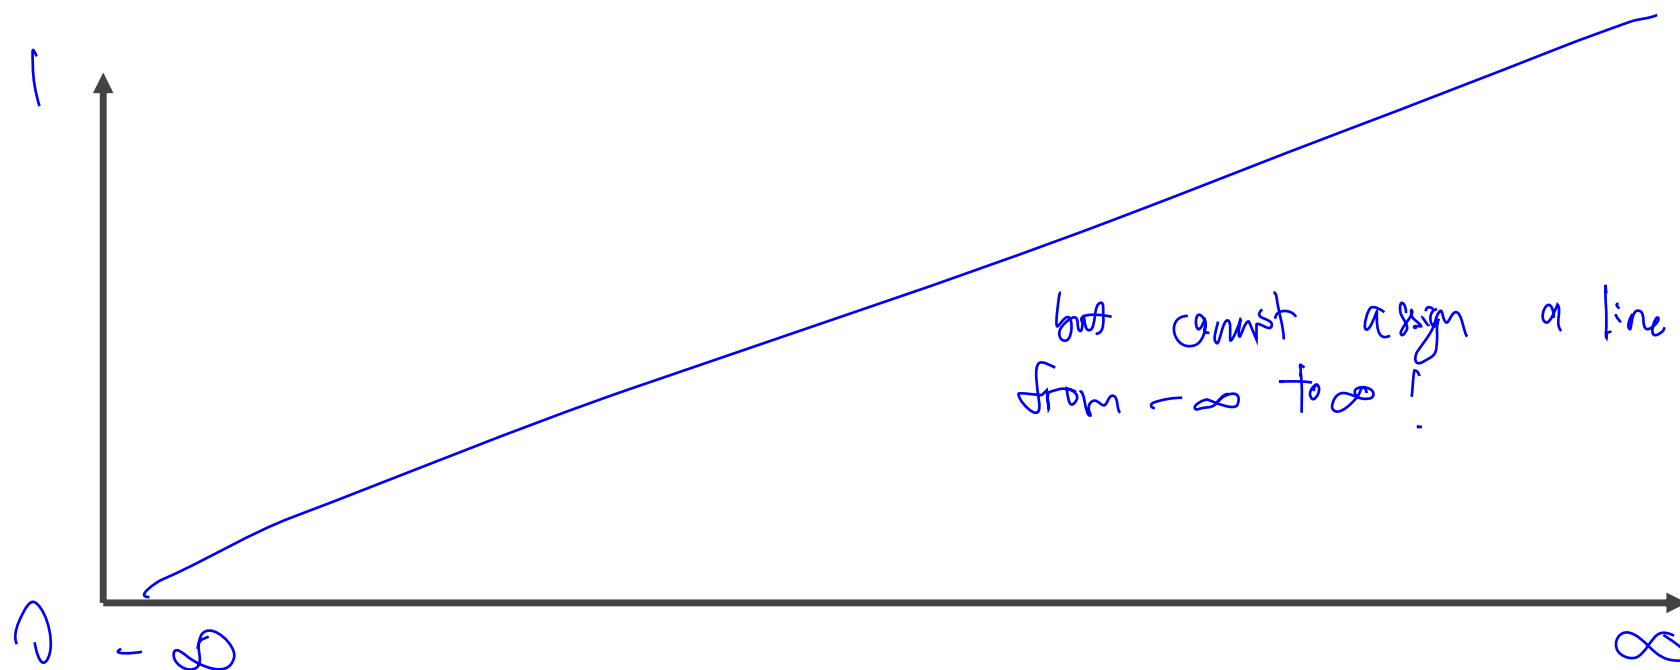
Want a model that tells us  $P(y = 1|x; \theta)$  and  $P(y = 0|x; \theta)$

But:  $z$  isn't a probability, but a number! How do we solve this?

*Slide Credits: Dan Jurafsky (Stanford)*

# Map to $\mathbb{R}$ interval $[0,1]$

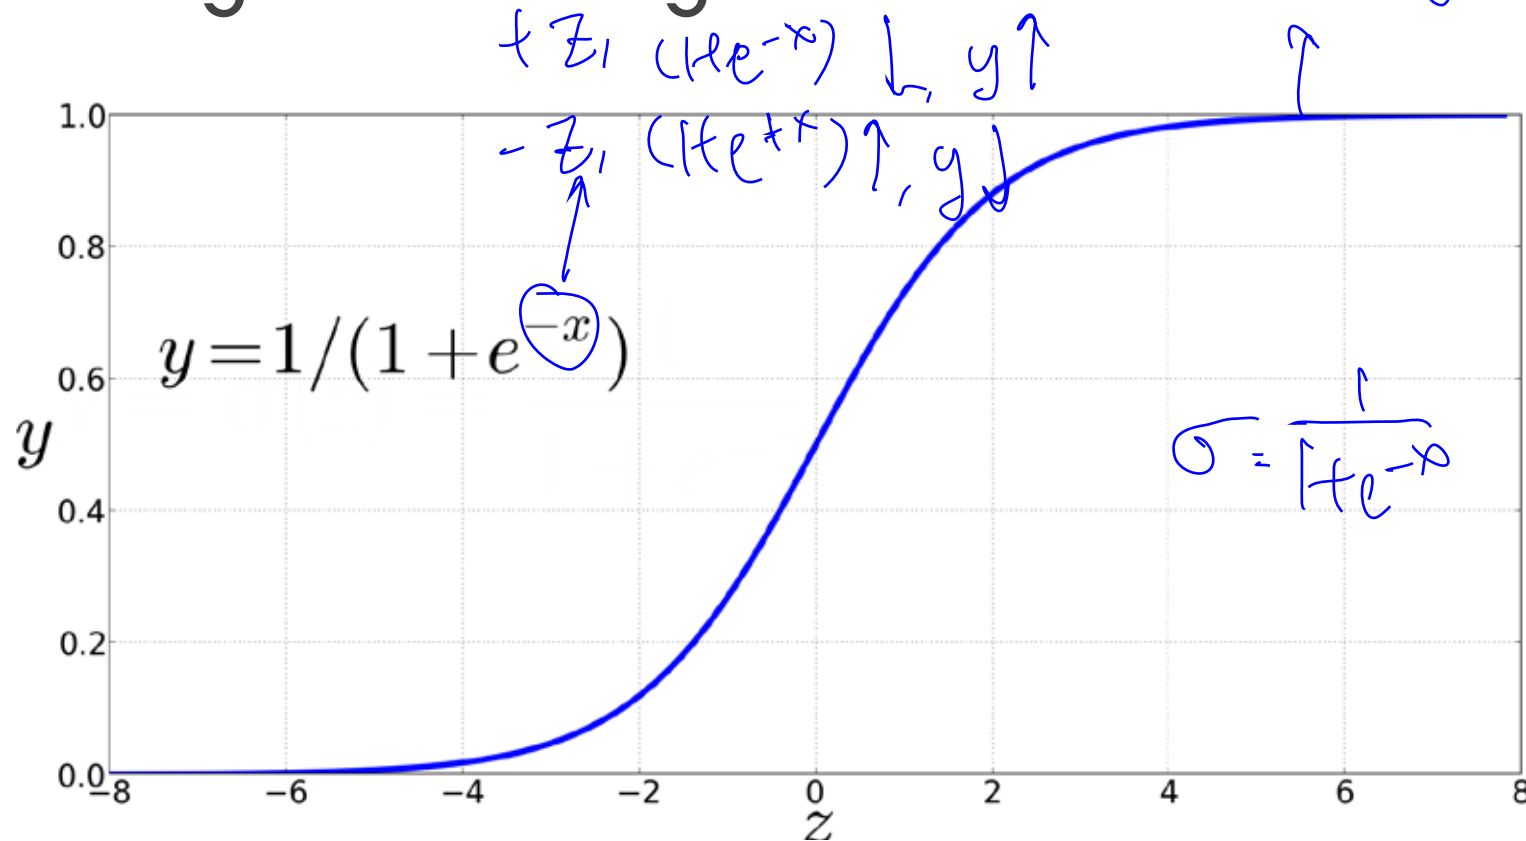
Need a function that maps real numbers to the unit interval.





# The sigmoid or logistic function

cannot be uniformly high



# Idea of logistic regression

We'll compute our signal  $z = \theta \cdot x$   
*conv. infinite range  $\rightarrow$   $[0,1]$  range*

We'll pass it through the sigmoid function  $\sigma(\theta \cdot x)$

And we'll just treat it as a probability.

*Slide Credits: Dan Jurafsky (Stanford)*

# Making probabilities with sigmoids

$$P(y = 1) = \boxed{\sigma(\theta \cdot x)} \rightarrow \text{This returns a probability} \in [0, 1].$$

$$= \frac{1}{1 + \exp(-(\theta \cdot x))}$$

$$P(y = 0) = 1 - \sigma(\theta \cdot x) \rightarrow 1 - P(y = 1)$$

$$= 1 - \frac{1}{1 + e^{-\theta x}}$$

$$= \frac{1 + e^{-\theta x} - 1}{1 + e^{-\theta x}} = \frac{e^{-\theta x}}{1 + e^{-\theta x}}$$

Slide Credits: Dan Jurafsky (Stanford)

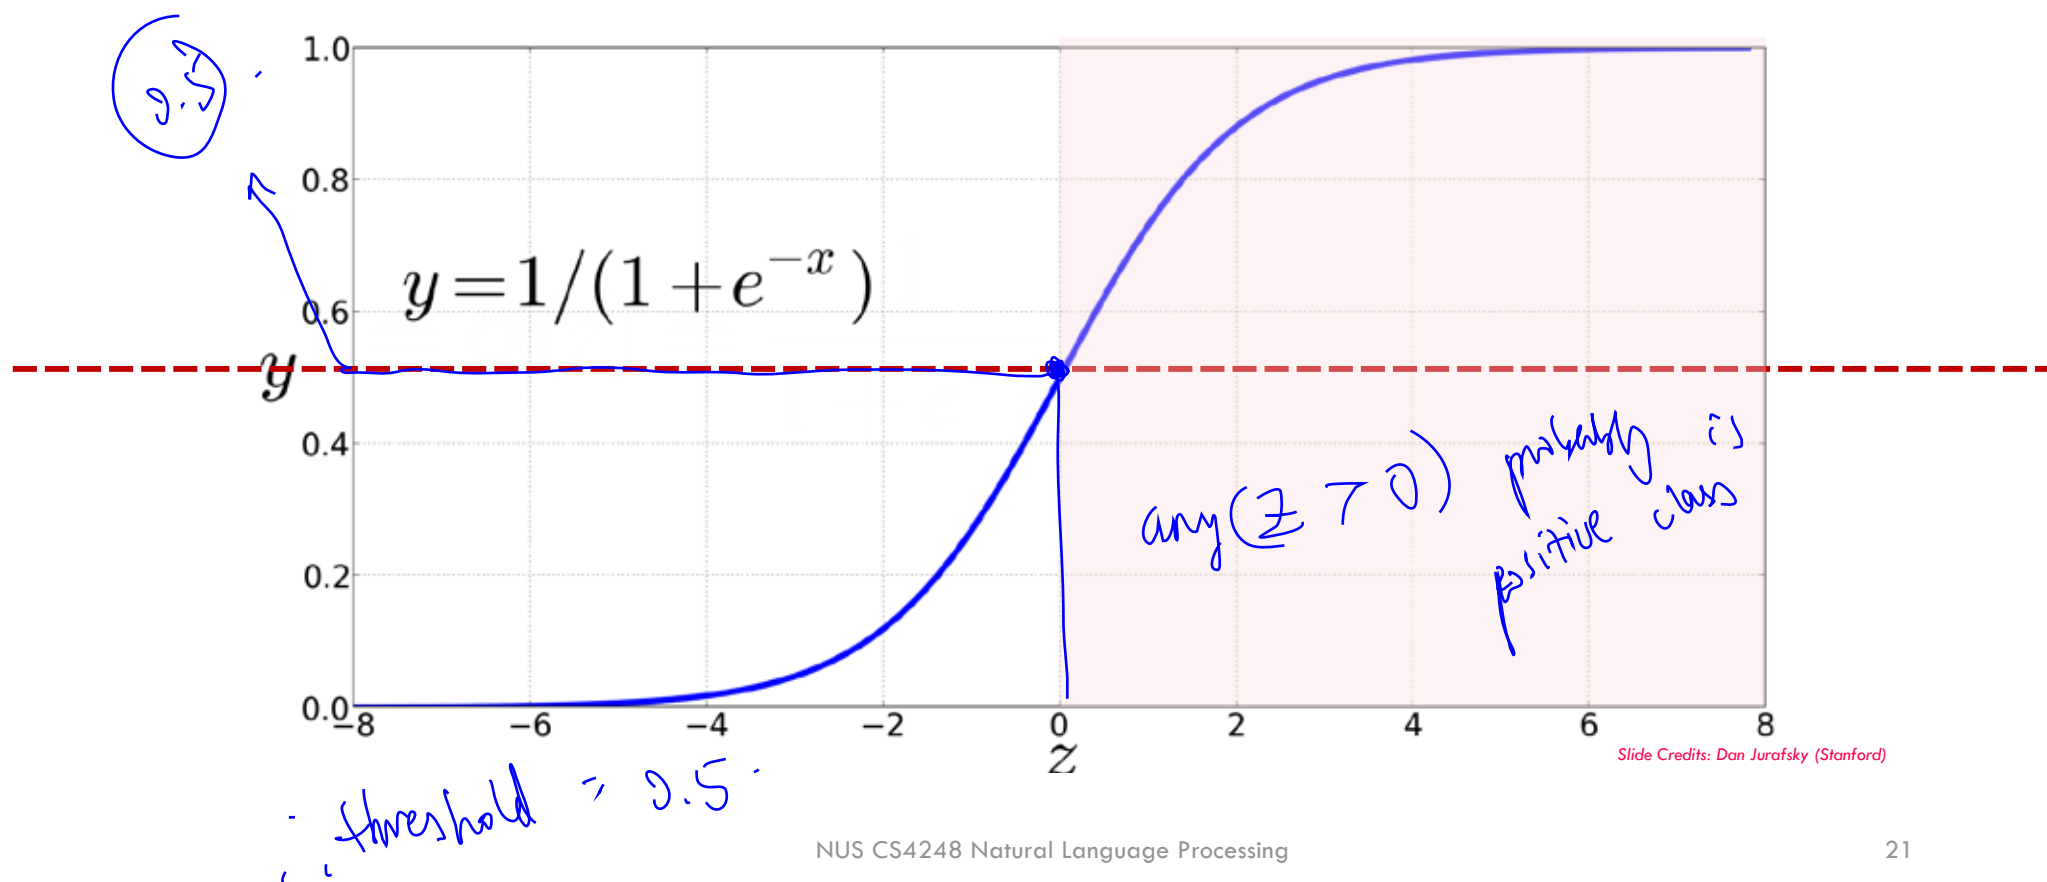
# Making probabilities with sigmoids

$$P(y = 1) = \sigma(\theta \cdot x)$$
$$= \frac{1}{1 + \exp(-(\theta \cdot x))}$$

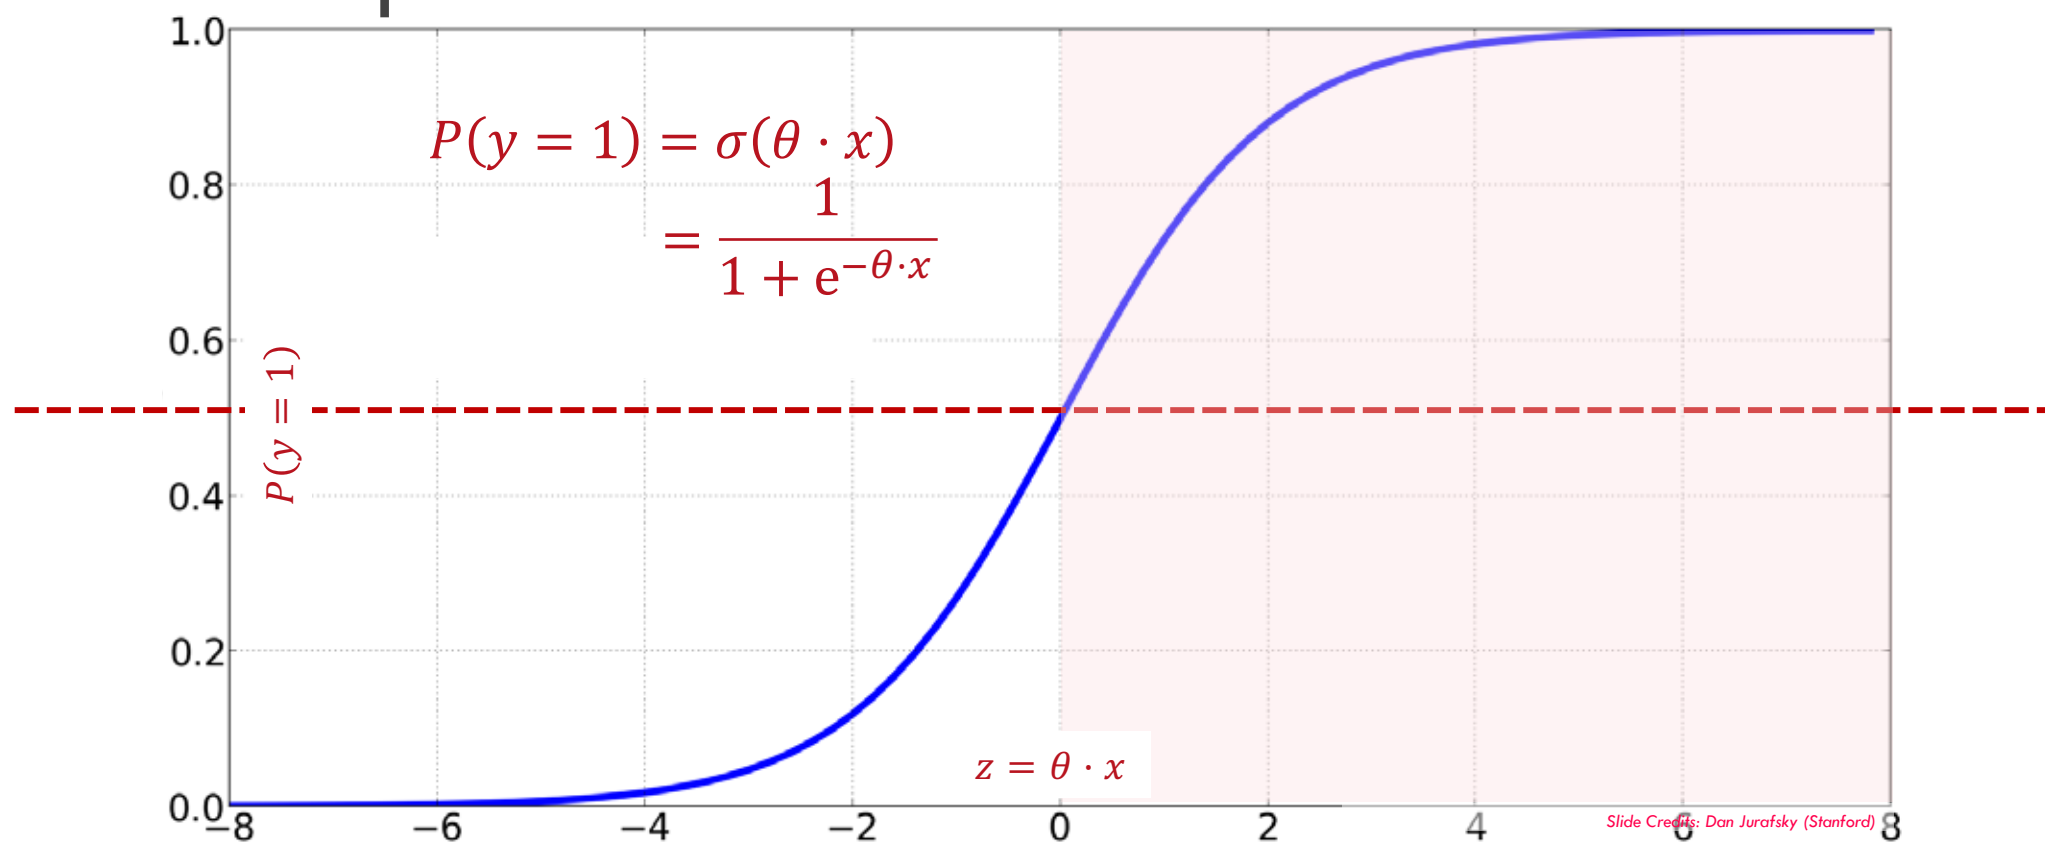
$$P(y = 0) = 1 - \sigma(\theta \cdot x)$$
$$= 1 - \frac{1}{1 + \exp(-(\theta \cdot x))}$$
$$= \frac{\exp(-(\theta \cdot x))}{1 + \exp(-(\theta \cdot x))}$$

*Slide Credits: Dan Jurafsky (Stanford)*

# The sigmoid or logistic function



# The probabilistic classifier



# LR Runthrough

## Sentiment Analysis Case Study

# Movie Review: does $\hat{y} = 1$ or 0?


It's hokey . There are virtually no surprises , and the writing is second-rate .  
So why was it so enjoyable ? For one thing , the cast is great . Another nice  
touch is the music . I was overcome with the urge to get off the couch and  
start dancing . It sucked me in , and it'll do the same to you .

Feature	Description	Value
$x_1$	Count of words in +ve lexicon	
$x_2$	Count of words in -ve lexicon	
$x_3$	1 if “no” in doc; 0 otherwise	
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	
$x_5$	1 if “I” in doc; 0 otherwise	
$x_6$	Log of the word count	



It's **hokey** . There are virtually no surprises , and the writing is **second-rate** . So why was it so **enjoyable** ? For one thing , the cast is **great** . Another **nice** touch is the music . *I* was overcome with the urge to get off the couch and start dancing . It sucked *me* in , and it'll do the same to *you* .

Feature	Description	Value
$x_1$	Count of words in +ve lexicon	<b>3</b>
$x_2$	Count of words in -ve lexicon	<b>2</b>
$x_3$	1 if “no” in doc; 0 otherwise	<u>1</u>
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	<b>3</b>
$x_5$	1 if “!” in doc; 0 otherwise	0
$x_6$	ln of the word count	$\ln(66) = 4.19$

 **Important:** LR and NB both require **feature engineering** as they do not combine primitive features together to make composite ones.

*Adapted from Dan Jurafsky (Stanford)*

# Now factor in the weights

Feature	Description	Value ( $x$ )	Weight ( $\theta$ ; Assumed)	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	
$x_1$	Count of words in +ve lexicon	3	2.5	
$x_2$	Count of words in -ve lexicon	2	-5.0	
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	
$x_5$	1 if "!" in doc; 0 otherwise	0	2.0	
$x_6$	ln of the word count	4.19	0.7	

*Adapted from Dan Jurafsky (Stanford)*

$$Z = \sum_{i=1}^n \theta_i x_i = 1 \times 0.1 + 3 \times 2.5 + 2 \times (-5) + 1 \times (-1.2) + 3 \times 0.5 + 4.19 \times 0.7 = 2.933$$

## LR Calculation

Feature	Description	Value ( $x$ )	Weight ( $\theta$ )	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "!" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	2.933

$$P(+|x) = P(y = 1|x) = \sigma(\theta \cdot x) =$$

$$P(+|x) = P(y = 0|x) = 1 - \sigma(\theta \cdot x) =$$

$$\frac{1}{1 + e^{-z}}$$

$$1.883$$

Adapted from Dan Jurafsky (Stanford)

# LR Calculation

Feature	Description	Value ( $x$ )	Weight ( $\theta$ ; Assumed)	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "!" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	+ 2.933

$$P(+|x) = P(y = 1|x) = \sigma(\theta \cdot x) =$$

$$P(-|x) = P(y = 0|x) = 1 - \sigma(\theta \cdot x) =$$

Adapted from Dan Jurafsky (Stanford)

# LR Calculation

Feature	Description	Value ( $x$ )	Weight ( $\theta$ ; Assumed)	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "!" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	+ 2.933
				0.833

$$P(+|x) = P(y = 1|x) = \sigma(\theta \cdot x) =$$

$$\sigma(0.833) = 0.7$$

$$P(-|x) = P(y = 0|x) = 1 - \sigma(\theta \cdot x) =$$

$$1 - \sigma(0.833) = 1 - 0.7 = 0.3$$

Adapted from Dan Jurafsky (Stanford)

# LR Calculation: *It's positive!*

Feature	Description	Value ( $x$ )	Weight ( $\theta$ ; Assumed)	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "!" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	+ 2.933
				0.833

$$P(+|x) = P(y = 1|x) = \sigma(\theta \cdot x) = \sigma(0.833) = 0.7$$

$$P(-|x) = P(y = 0|x) = 1 - \sigma(\theta \cdot x) = 1 - \sigma(0.833) = 1 - 0.7 = 0.3$$

$P(+|x) > 0.5, \hat{y} = +$  classified as positive

Adapted from Dan Jurafsky (Stanford)

# Cross Entropy

What's the goal of learning?  
To get the best weights.

# Wait, where did the $\theta$ 's come from?

Supervised classification: we know the correct label  $y$  (either 0 or 1) for each  $x$ .

What the system produces is an estimate,  $\hat{y}$ .

We want to set  $\theta$  to **minimize the difference** between our estimate  $\hat{y}^{(i)}$  and the true  $y^{(i)}$ .

1. We need a metric: a **loss function** (also termed **cost function**); and
2. We need an **optimization algorithm** to update  $\theta$  to minimize the loss.



Let's deal  
with #1 first.

*Adapted from Dan Jurafsky (Stanford)*



# Our Objective Function

We want to know how far is the classifier output:

$$\hat{y} = \sigma(\theta x)$$

from the true output:

$$y \text{ [= either 0 or 1]}$$

We'll call this difference:

$$L(\hat{y}, y) = \text{how much } \hat{y} \text{ differs from the true } y$$

Machine Learning likes to  
minimize loss, hence  $L$ .

maximize utility

Slide Credit: Dan Jurafsky (Stanford)

# Negative log likelihood loss

*(Also termed cross-entropy loss)*

A case of conditional maximum likelihood estimation:

We choose the parameters  $\theta$

That maximize the log probability of the true  $y$  labels in the training data

Given the observations  $x$ .

*Slide Credit: Dan Jurafsky (Stanford)*

# Cross-entropy Loss

**Goal:** maximize probability of the correct label  $p(y|x)$

Since there are 2 discrete outcomes (0 or 1), we can express the probability  $p(y|x)$  from our classifier (what we want to maximize) as:

$$\begin{array}{ll} 1 - \hat{y}, & \text{if } y = 0; \quad \text{— negative labels} \\ \hat{y}, & \text{if } y = 1. \quad \text{— positive labels} \end{array}$$

We can combine both cases into one formula this way:

$$P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

$\nearrow y=0, p(y|x) = 1 - \hat{y}$   
 $\nwarrow y=1, p(y|x) = \hat{y}$

Slide Credit: Dan Jurafsky (Stanford)

# Cross-entropy Loss

**Goal:** maximize  $P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$ .

Take logs of both sides (monotonically equivalent):

*Slide Credit: Dan Jurafsky (Stanford)*

# Cross-entropy Loss

**Goal:** maximize  $P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$ .

Take logs of both sides (monotonically equivalent):

$$\text{maximize } \log P(y|x) = \log[\hat{y}^y (1 - \hat{y})^{1-y}]$$

$$\text{maximize } \log P(y|x) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

*Slide Credit: Dan Jurafsky (Stanford)*

# Cross-entropy Loss

$$\text{maximize } \log P(y|x) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

We have a maximization, but ML terminology prefers losses to minimize (as in the title). Let's flip the sign.

The result is **cross-entropy loss** (cross entropy between  $\hat{y}$  and  $y$ ):

$$\text{minimize } L_{ce}(\hat{y}, y) = -\log p(y|x) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

*Slide Credit: Dan Jurafsky (Stanford)*

# Does it work?

Loss should be:

- **Smaller** if the model estimate is close to correct
- **Larger** when the model is confused

For our sentiment example, let's examine it, pretending if it was positive or negative.

*“ It's hokey . There are virtually no surprises , and the writing is second-rate . So why was it so enjoyable ? For one thing , the cast is great . Another nice touch is the music . I was overcome with the urge to get off the couch and start dancing . It sucked me in , and it'll do the same to you . ”*

*Slide Credit: Dan Jurafsky (Stanford)*

We calculated:

$$P(+|x) = \sigma(0.833) = 0.7$$

$$P(-|x) = 1 - \sigma(0.833) = 0.3$$

Feature	Description	Value ( $x$ )	Weight ( $\theta$ )	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "I" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	2.933

$$L_{ce}(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

If the model was right ( $y = 1$ )

$$-1 (\log 0.7)$$

2

If the model was wrong ( $y = 0$ )

$$-\log(0.3)$$

2

Slide Credit: Dan Jurafsky (Stanford)



# Loss is bigger when the model is wrong

We calculated:

$$P(+|x) = \sigma(0.833) = 0.7$$

$$P(-|x) = 1 - \sigma(0.833) = 0.3$$

Feature	Description	Value (x)	Weight ( $\theta$ )	Product ( $\theta x$ )
$x_0$	Bias $b$	1	0.1	0.1
$x_1$	Count of words in +ve lexicon	3	2.5	7.5
$x_2$	Count of words in -ve lexicon	2	-5.0	-10.0
$x_3$	1 if "no" in doc; 0 otherwise	1	-1.2	-1.2
$x_4$	Count of 1 <sup>st</sup> & 2 <sup>nd</sup> person pronouns	3	0.5	1.5
$x_5$	1 if "I" in doc; 0 otherwise	0	2.0	0
$x_6$	ln of the word count	4.19	0.7	2.933

$$L_{ce}(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

If the model was right ( $y = 1$ )

$$\begin{aligned}
 &= -[\log(\hat{y})] \\
 &= -[\log(0.7)] \\
 &= 0.36
 \end{aligned}$$

If the model was wrong ( $y = 0$ )

$$\begin{aligned}
 &= -[\log(1 - \hat{y})] \\
 &= -[\log(0.3)] \\
 &= 1.2
 \end{aligned}$$

Slide Credit: Dan Jurafsky (Stanford)

loss is bigger when model is wrong

# Multiclass with Logistic Regression

We can generalize logistic regression for 2 or more classes:

- Generalize to **Multinomial** Logistic Regression
- Features have separate weights for each of the  $k$  classes
- Upgrade Sigmoid to the **Softmax**, keeping the  $\mathbb{R} \rightarrow [0,1]$  idea:

$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}, 1 \leq i \leq k$$

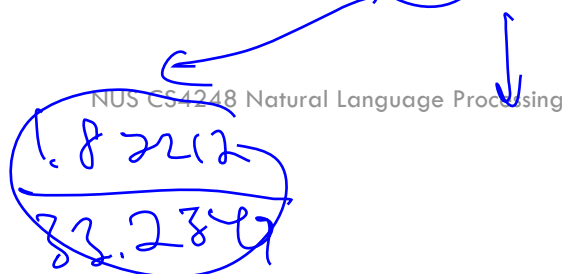
Also termed *Maximum Entropy Modeling (MaxEnt)*

Sums to unity

(e.g.,  $z = [0.6, 1.1, -1.5, 1.2, 3.2, -1.1] \rightarrow [0.055, 0.090, 0.0067, 0.1, 0.74, 0.01]$ )

Slide Credit: Dan Jurafsky (Stanford)

NUS CS4248 Natural Language Processing



$$\frac{1.8212}{32.2867}$$

# Stochastic Gradient Descent

Now that we know how we're doing,  
how can we do better?

Slide Credits: NUS CS3244 and Dan Jurafsky (Stanford)

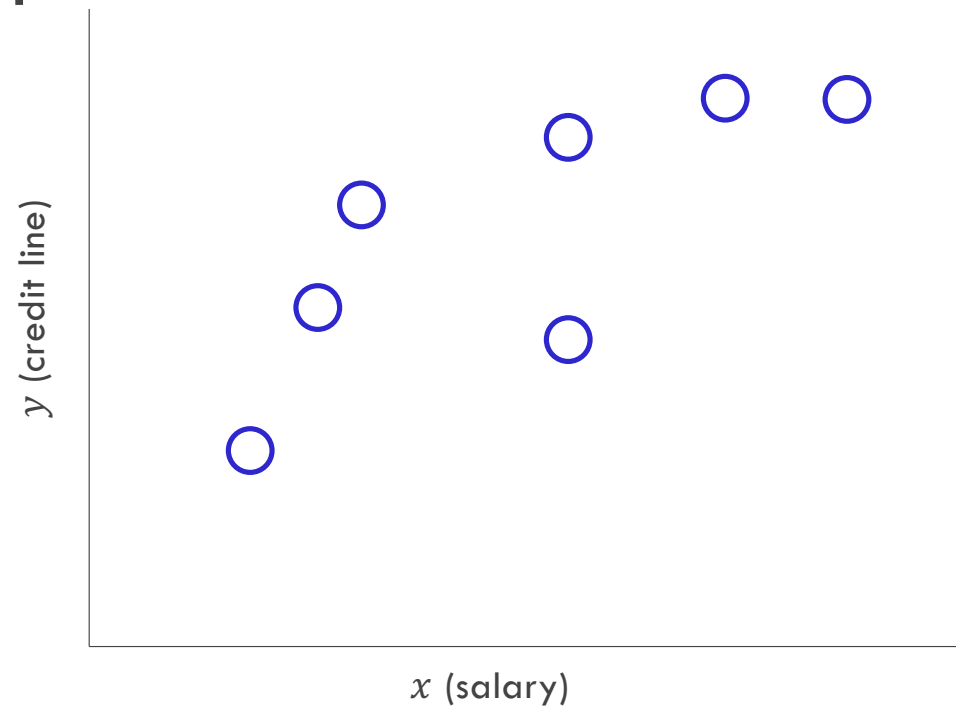
# Gradient Descent

Climbing up (down)  
one step at a time



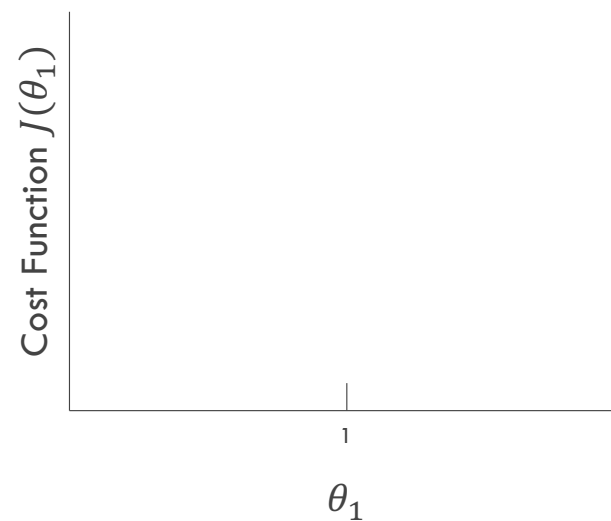
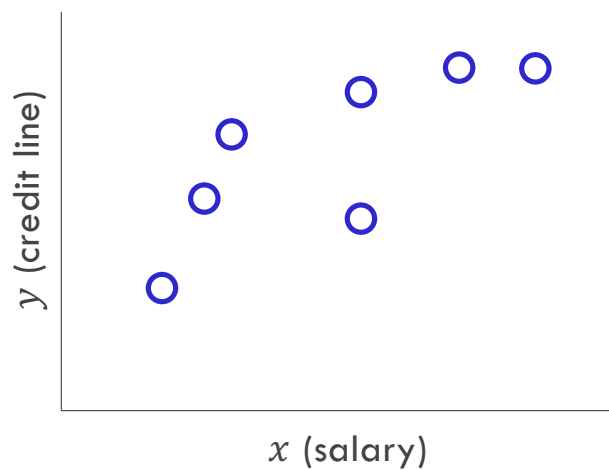
Slide Credits: NUS CS3244

# Univariate Linear Regression: Salary to predict Credit Line



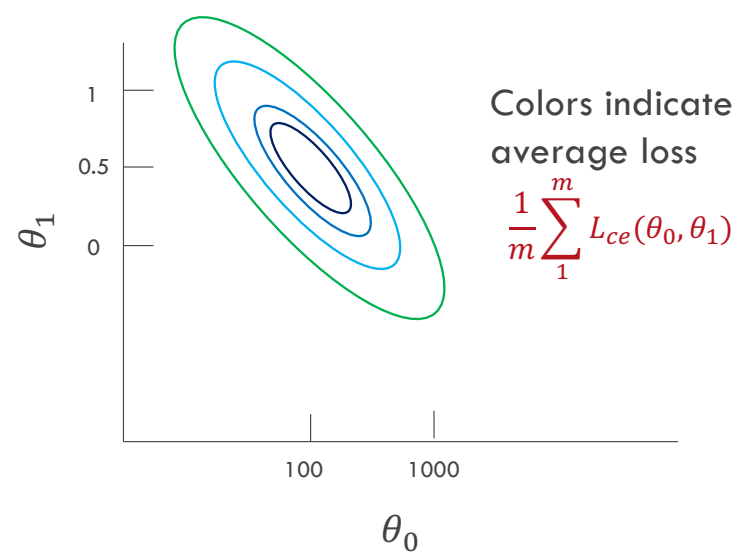
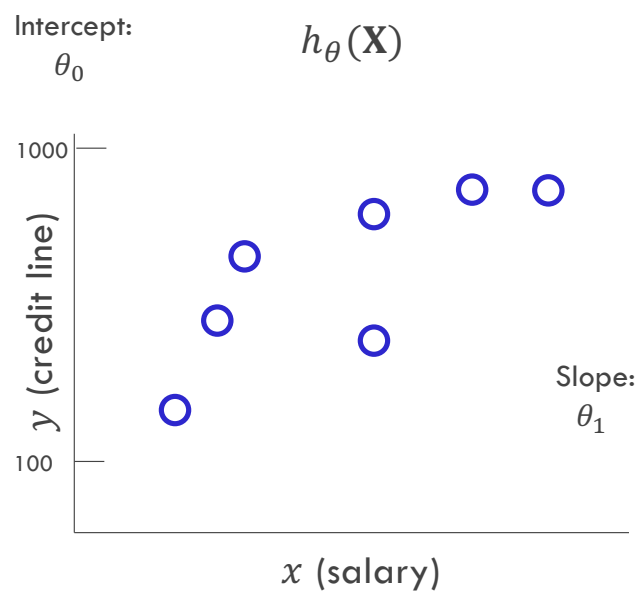
Slide Credits: NUS CS3244

# Ignoring the bias term



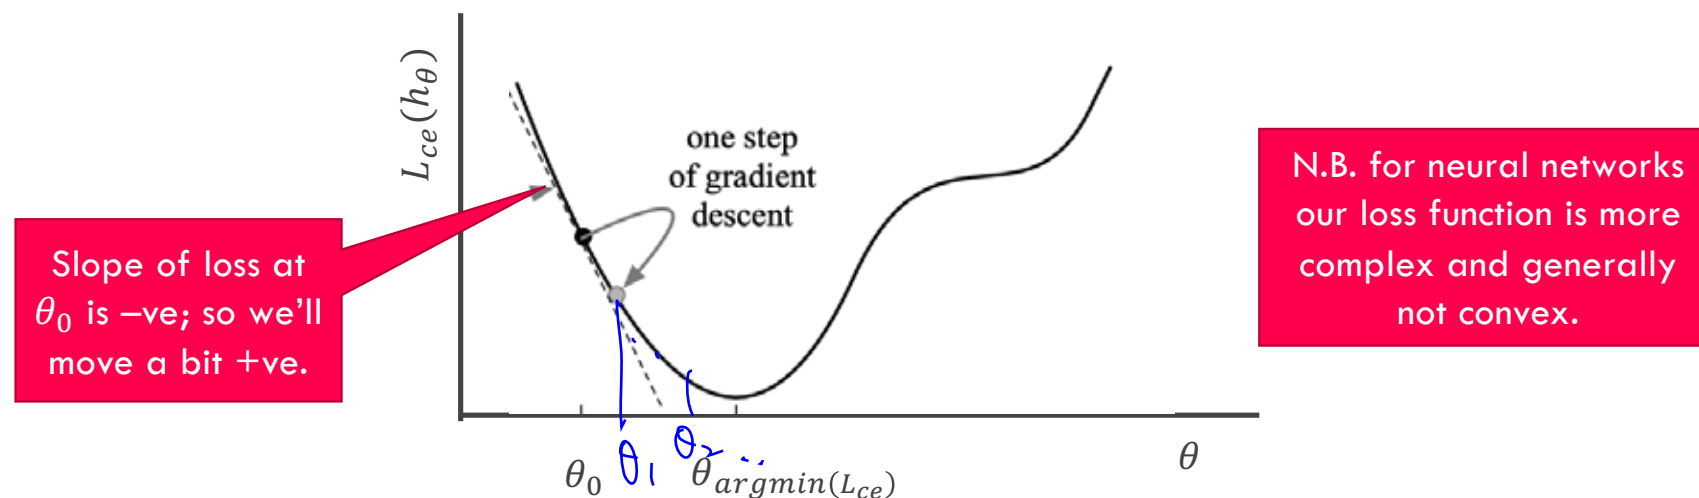
Slide Credits: NUS CS3244

# Loss Function $L_{ce}(\theta_0, \theta_1)$



Slide Credits: NUS CS3244

# Cross Entropy Loss Function Curve

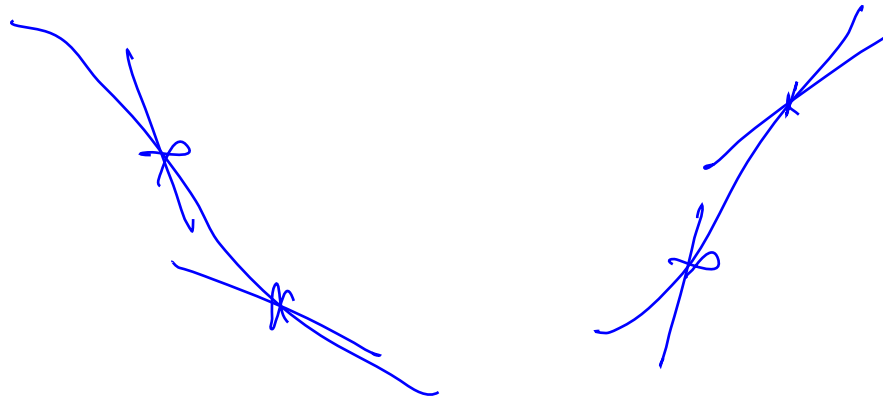


... Because  $L_{ce}$  for linear regression is a **convex function** of  $\theta$ .

Slide Credits: NUS CS3244 and Dan Jurafsky (Stanford)



# Gradients



The **gradient** of a function of many variables is a vector pointing in the direction of the greatest increase in a function.

**Gradient Descent:** Find the gradient of the loss function at the current point and move in the **opposite** direction.

*Slide Credit: Dan Jurafsky (Stanford)*

# Iterative method: gradient descent

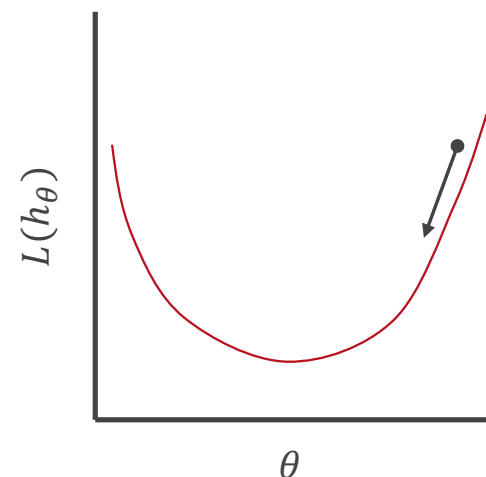
General method for nonlinear optimization

Start at  $\theta(t)$ ; take a step along the direction with the steepest gradient.

How big a step / How fast to learn?

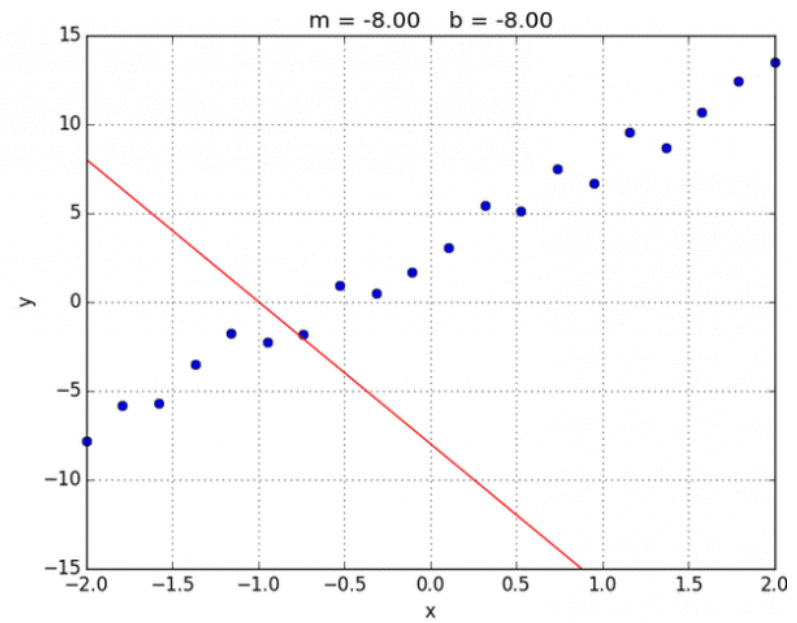
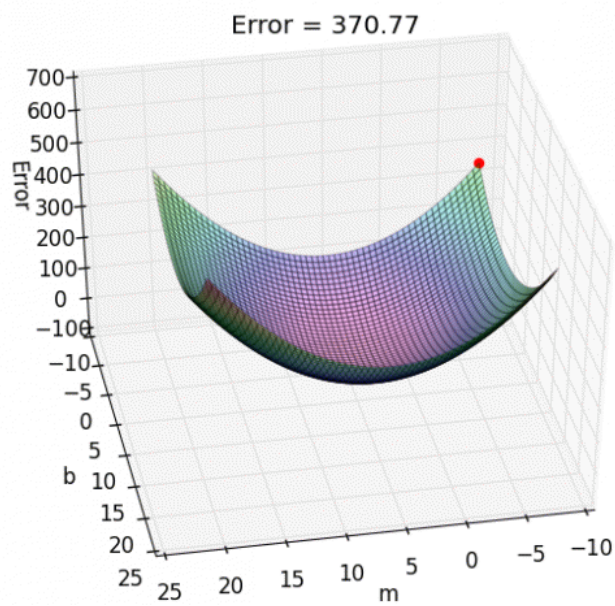
Dependent on a fixed step size  $\eta$  :

$$\theta(t+1) = \theta(t) - \eta \cdot \frac{d}{d\theta} h_{\theta}(x)$$



Gradient descent can minimize any smooth function (= needs a derivative).

Slide Credits: NUS CS3244



Credits: [Alykhan Tejani's Medium Post](#)

$$L_{ce}(\hat{y}, y) = -[y \log \sigma(w \cdot x + b) + (1-y) \log(1 - \sigma(w \cdot x + b))] \Rightarrow [\sigma(w \cdot x + b) - y] x_j$$

## The Gradient

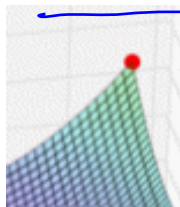
We'll represent  $y$  as  $h(x; \theta)$  to make the dependence on  $\theta$  more obvious:

gradient of loss fn.

$$\nabla_{\theta} L(h(x; \theta), y) = \begin{bmatrix} \frac{\partial}{\partial \theta_0} L(h(x; \theta), y) \\ \vdots \\ \frac{\partial}{\partial \theta_n} L(h(x; \theta), y) \end{bmatrix}$$

$\rightarrow [\sigma(w \cdot x + b) - y] x_j$

difference between true  $y$  and  $\hat{y} = \sigma(w \cdot x + b)$ , multiplied by  $x_j$ .



The final equation for updating  $\theta$  based on the gradient is thus:

$$\theta(t+1) = \theta(t) - \eta \nabla_{\theta} L(f(x; \theta), y)$$

Slide Credit: Dan Jurafsky (Stanford)

step size

# LR as a Probabilistic ML classifier

Putting it together

[Click to edit Master Attribution style.](#)

# Components of a probabilistic machine learning classifier

Given  $m$  input/output pairs  $(\mathbf{x}^{(j)}, y^{(j)})$ :

1. A feature representation of the input.  $\mathbf{x}^{(j)} = [x_1, x_2, \dots, x_i, \dots, x_n]$ .  
The  $i$ th feature is denoted  $x_i$ , or more completely  $x_i^{(j)}$ , sometimes  $f_i(\mathbf{x})$ .
2. A classification function that computes  $\hat{y}$ , the estimated class, via  $p(y|\mathbf{x})$ , like the **sigmoid** or **softmax** functions.
3. An objective function for learning, like **cross-entropy** loss.
4. An algorithm for optimizing the objective function: **stochastic gradient descent**.

Slide Credits: Dan Jurafsky (Stanford)

# Logistic regression algorithm

1. Initialize the weights at  $t = 0$  to  $\theta(0)$  → doesn't matter can start anywhere in the hypothesis, gradient descent later.
2. Do
3.     Compute the gradient gradient loss

$$\nabla(t) = \boxed{\nabla L_{ce}(\theta(t))} = -\frac{1}{m} \sum_{j=1}^m \frac{y^{(j)} x^{(j)}}{1 + e^{y^{(j)} \theta(t) \cdot x^{(j)}}}$$
sigmoid
4.     // Move in the direction  $v(t) = -\nabla(t)$   
       Update the weights  $\theta(t + 1) = \theta(t) - \alpha \nabla L_{ce}$  ← move in opposite direction of gradient
5.     Continue to next iteration, until it is time to stop.
6.     Return the final weights  $\theta^*$

Slide Credits: NUS CS3244

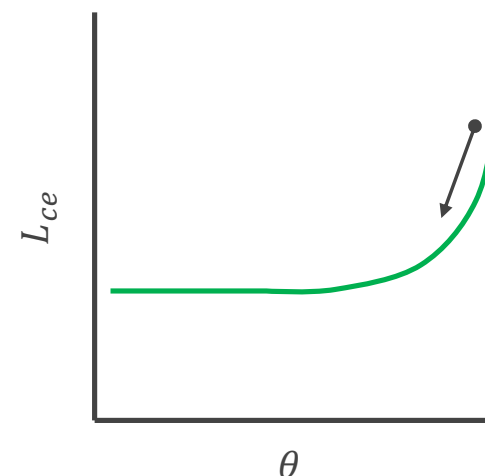
# Termination Condition

When to stop?

Natural choice:  $\text{gradient} < \text{threshold}$

But lots of flat regions in  
most spaces:

Instead, use criteria:



1. error **change** is small and/or;
2. error is small;
3. maximum number of iterations is reached.

Slide Credits: NUS CS3244



# Mini-batch training

**Stochastic gradient descent** chooses a single random example at a time. That can result in choppy movements

More common to compute gradient over batches of training instances.

*mini batches = easily vectorized  $\rightarrow$  computational efficiency.*

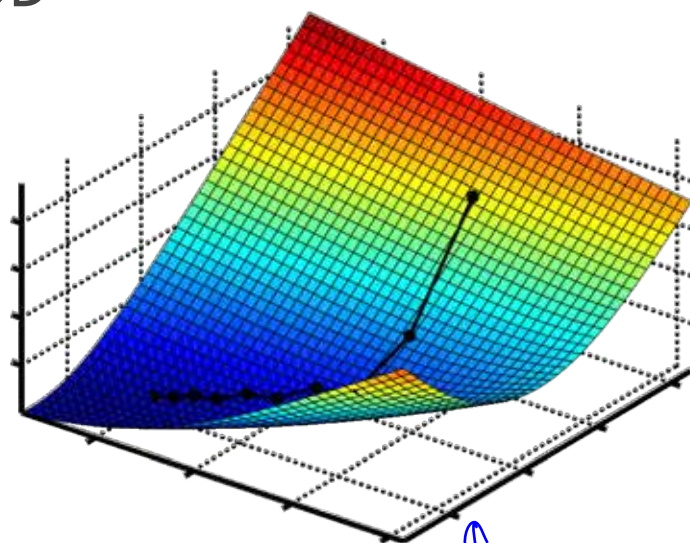
- **Batch training:** entire dataset
- **Mini-batch training:**  $m$  examples (512, or 1024)

$$\text{Mini batch gradient} = \frac{1}{m} \sum_{i=1}^m \text{LCE}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m [\sigma(w \cdot x^{(i)} + b) - y^{(i)}] x_j^{(i)}$$

Slide Credit: Dan Jurafsky (Stanford)

# GD vs. SGD on $m = 10$

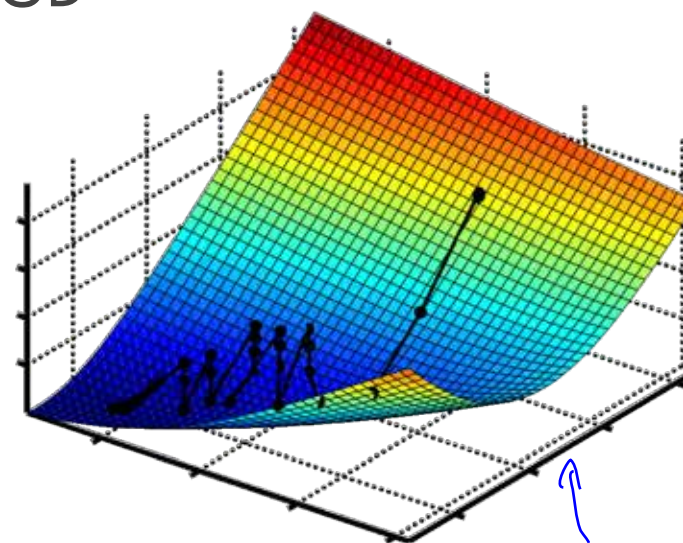
GD



10 steps

↑  
accurate but  
expensive

SGD



30 steps

↑  
cheap

Slide Credits: NUS CS3244

# Regularization

Sample Lecture Title

# Overfitting

A model that perfectly matches the training data often has a problem.

It may **overfit** to the data, modeling noise

- A random word that perfectly predicts  $y$  (it happens to only occur in one class) will get a very high weight.
- Failing to generalize to a test set without this word.

A good model should be able to **generalize**

# Overfitting

*What are some good  $n$ -gram features?*

Your answers:

+ This movie drew me in, and it'll do the same to you.

- I can't tell you how much I hated this movie. It sucked.

*Slide Credit: Dan Jurafsky (Stanford)*

# Overfitting

*What are some good  $n$ -gram features?*

+ This movie drew me in, and it'll do the same to you.

– I can't tell you how much I hated this movie. It sucked.

Your answers:

How do you feel about these?

$x_1$  = “the same to you”

$x_2$  = “tell you how much”

4-gram features that are very specific that just “memorize” training set might cause problems.

Slide Credit: Dan Jurafsky (Stanford)

# Overfitting

4-gram model on tiny data will just memorize the data

- 100% accuracy on the training set

But it will be surprised by the novel 4-grams in the test data

- Low accuracy on test set

Models that are too powerful can overfit the data

- Fitting the details of the training data so exactly that the model doesn't generalize well to the test set

*Slide Credit: Dan Jurafsky (Stanford)*

# Regularization

A solution for overfitting

Add an overfit penalty  $\Omega(\theta)$  to the loss function:

$$\hat{\theta} = \operatorname{argmax}_{\theta} \sum_{j=1}^m \log P(y^{(j)} | x^{(j)}) - \Omega(\theta)$$

Idea: choose a  $\Omega(\theta)$  that penalizes large weights.

*Fitting the data well with lots of big weights not as good as fitting the data a little less well, with small weights*



# Regularization Variants

L2 Regularization (= Ridge Regression)

$$\hat{\theta} = \operatorname{argmax} \sum_{j=1}^m \log P(y^{(j)} | x^{(j)}) - \sum_{i=1}^n \theta_i^2$$

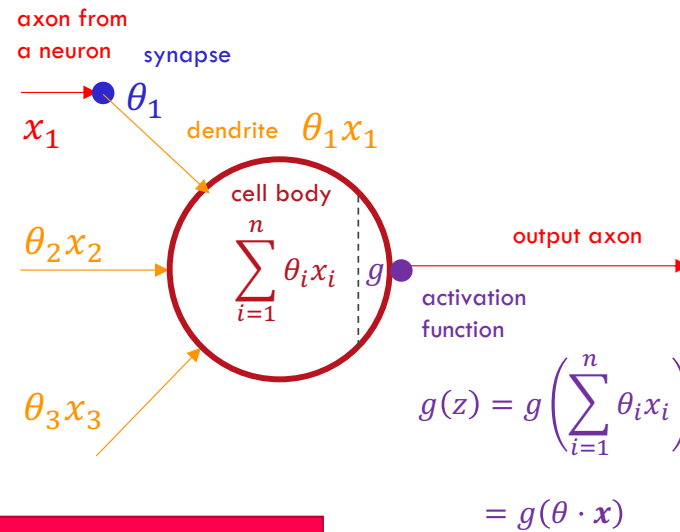
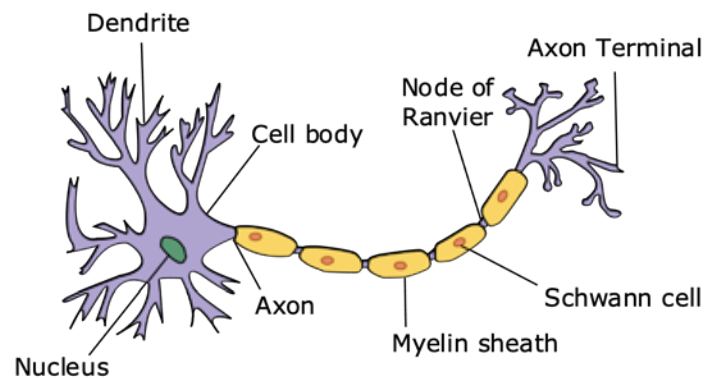
L1 Regularization (= Lasso Regression)

$$\hat{\theta} = \operatorname{argmax} \sum_{j=1}^m \log P(y^{(j)} | x^{(j)}) - \sum_{i=1}^n |\theta_i|$$

# XOR

## Motivating Neural Networks

# Biological Inspiration



*Looks familiar? A neuron is a LR unit!*

Diagram credits: Dhp1080 - Own work, CC BY-SA 3.0 via Wikimedia Commons.

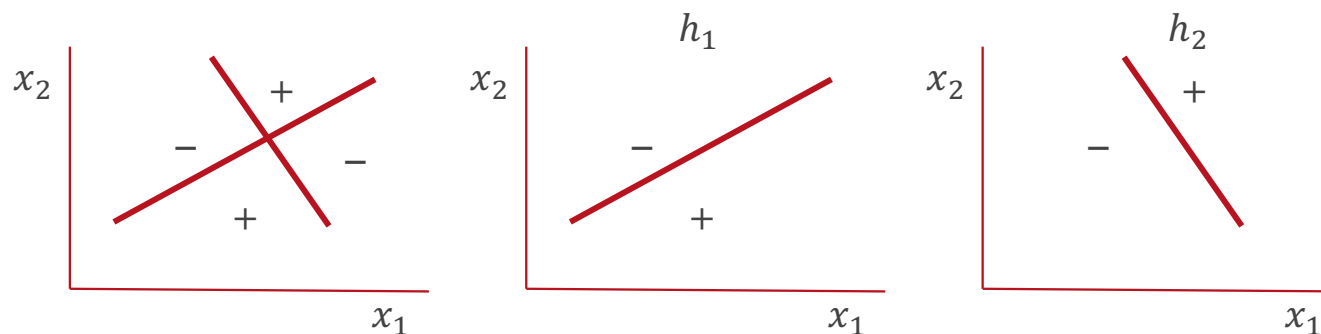
# Stacked Logistic Regression

Logistic Regression can represent any linear combination of the features, mapping them non-linearly to a probability.

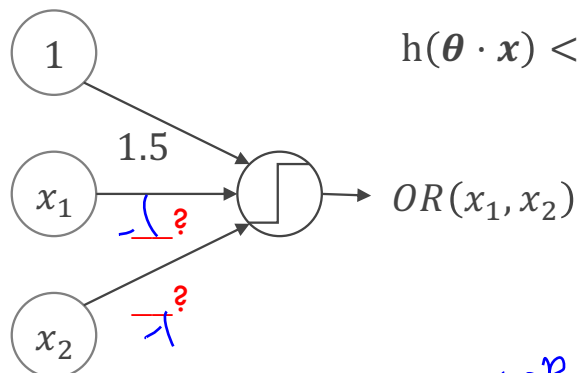
But what if we want to represent some non-linear relationship between features? Sorry, can't do it.

*Idea:* Use LR to create such non-linear features. Feed LR outputs from other LRs.

# Combining Linear Classifiers

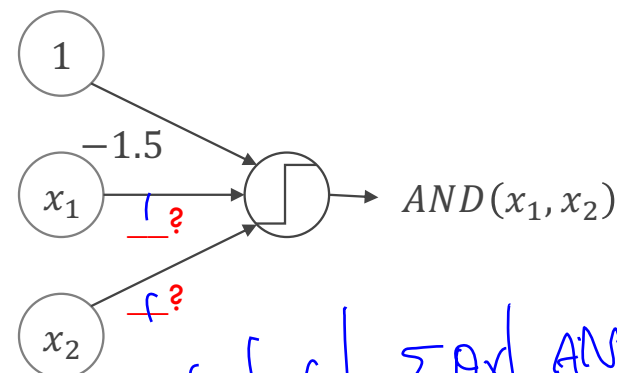


$h(\theta \cdot x) < 0.5 \quad y = \pm 1$



$x_1$	$x_2$	$\sum \theta x$	$OR(x_1, x_2)$
0	0	1.5	1
0	1	0.5	1
1	0	0.5	1
1	1	-0.5	1

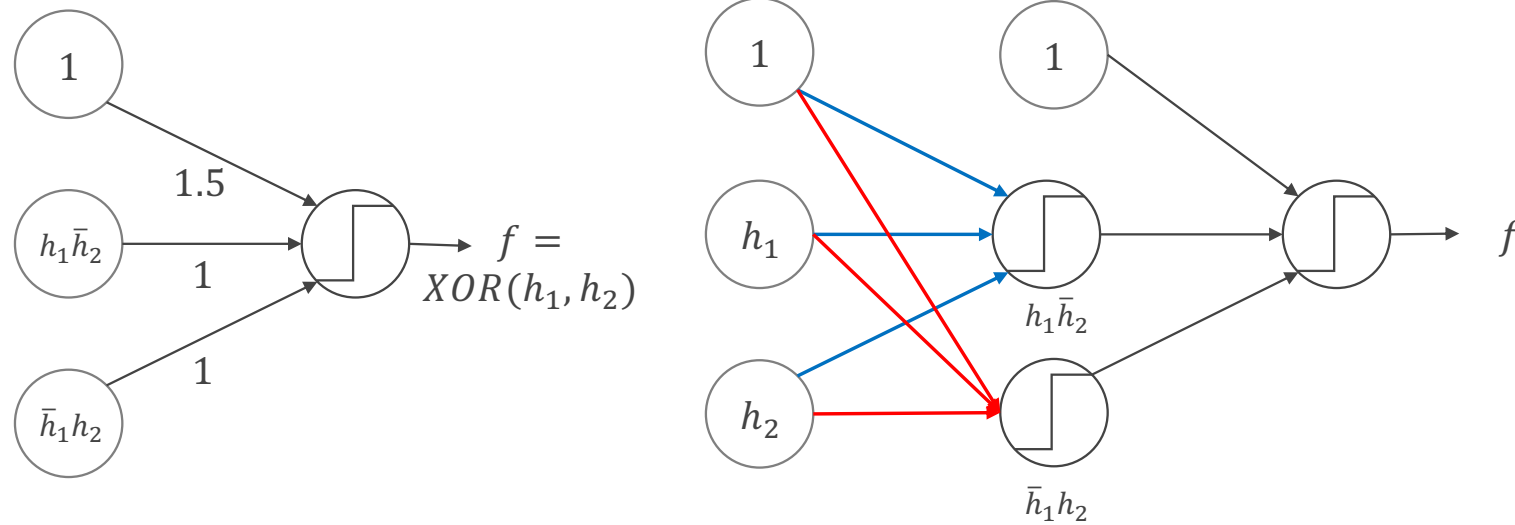
NUS CS4248 Natural Language Processing



$x_1$	$x_2$	$\sum \theta x$	$AND(x_1, x_2)$
0	0	-1.5	0
0	1	-0.5	0
1	0	-0.5	0
1	1	0.5	1

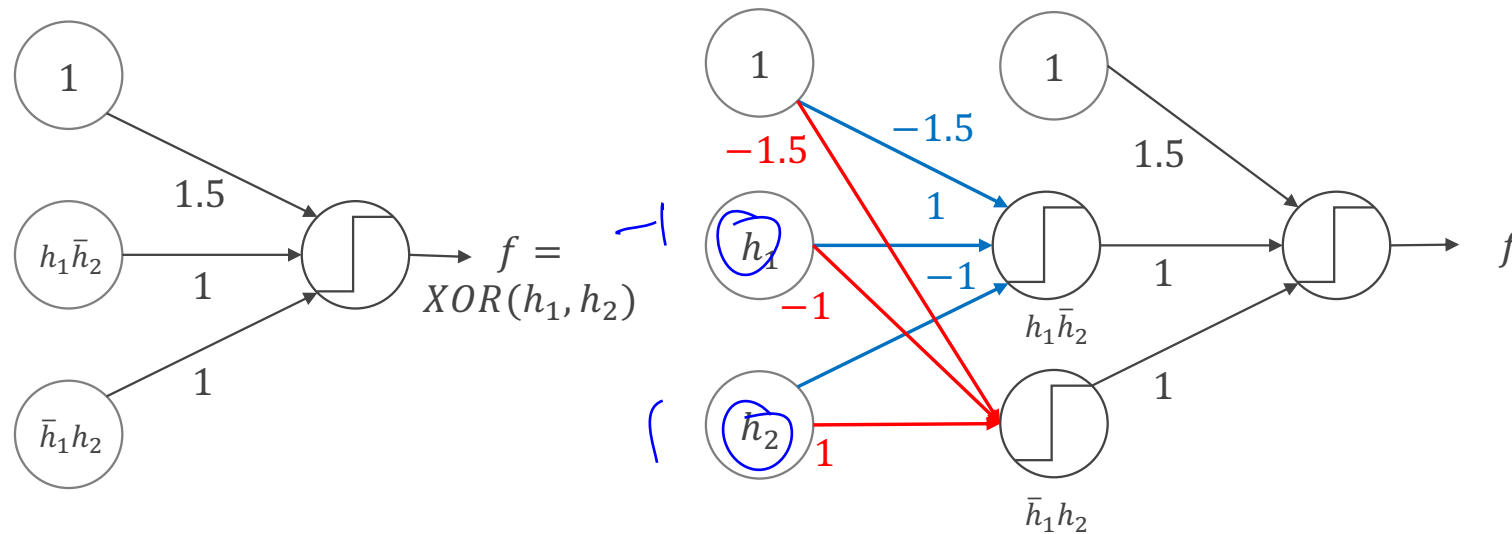
Slide Credits: CS3244

# Creating Layers



Slide Credits: CS3244

# Creating Layers

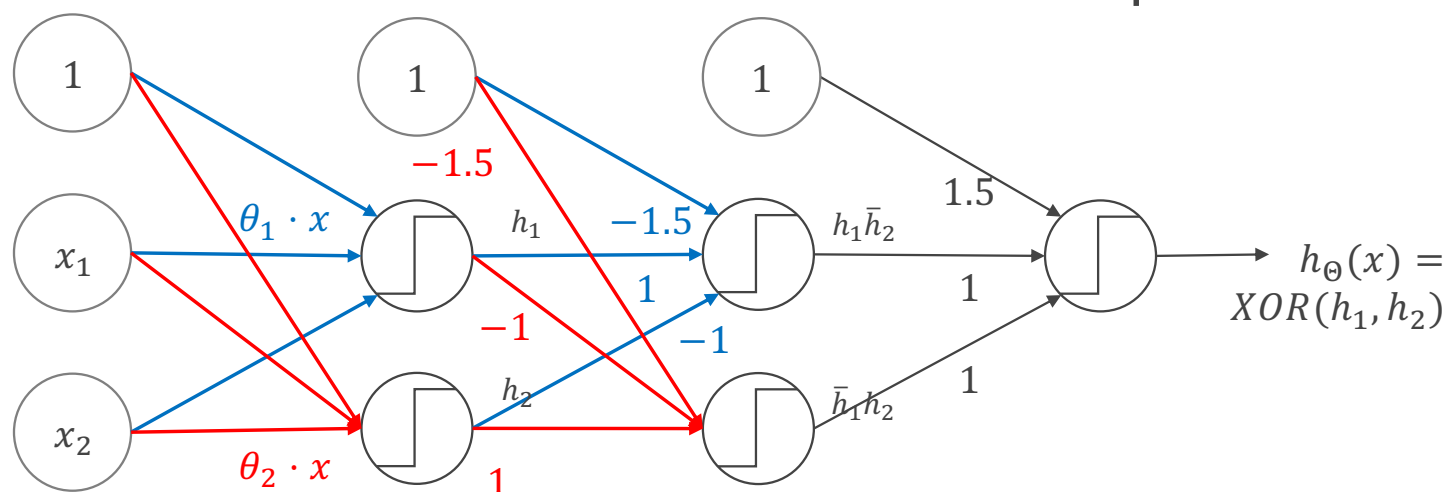


Slide Credits: CS3244

# NN = Stacked Linear Regression

3 layers

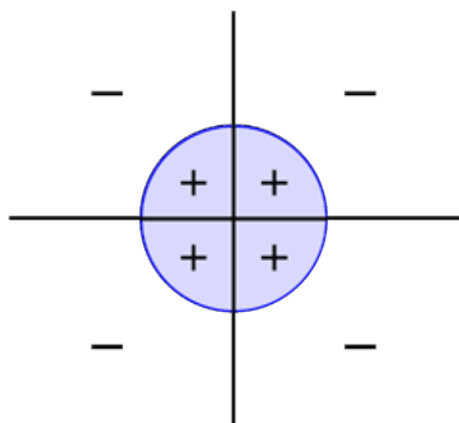
“feedforward”  $\equiv$  no loops



Slide Credits: CS3244

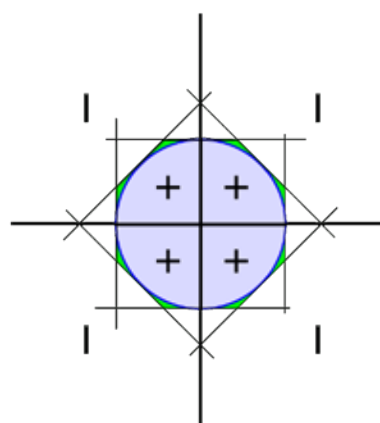


# A Powerful Model

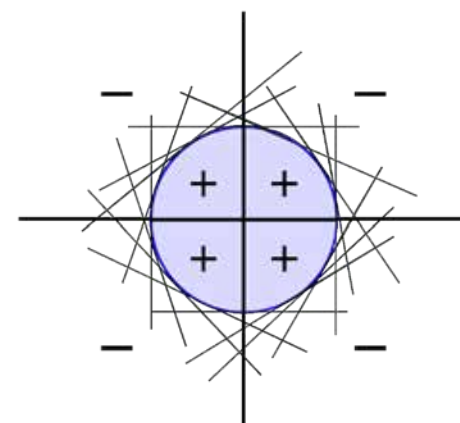


Target

2 red flags:



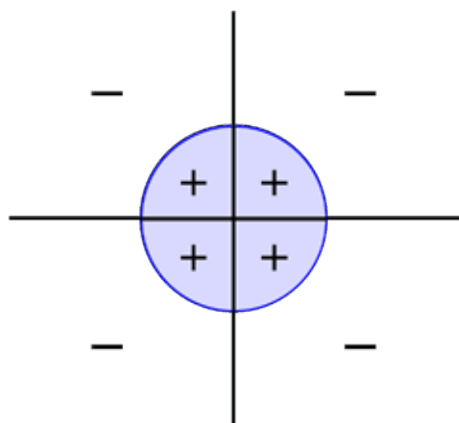
8 perceptrons



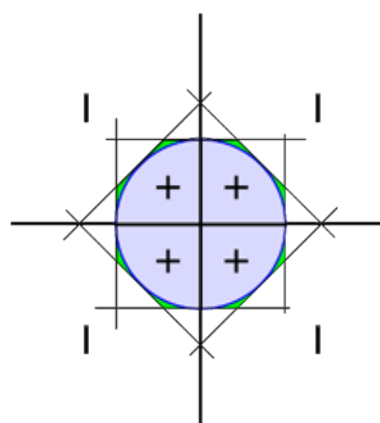
16 perceptrons

Slide Credits: CS3244

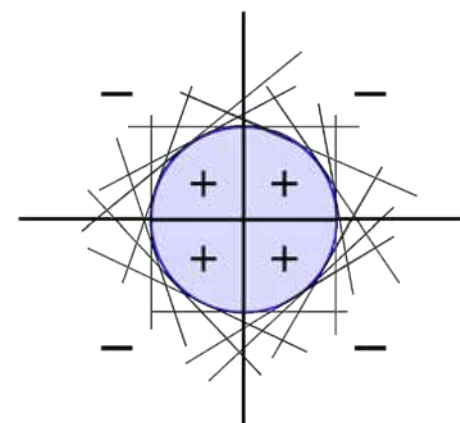
# A Powerful Model



Target



8 perceptrons



16 perceptrons

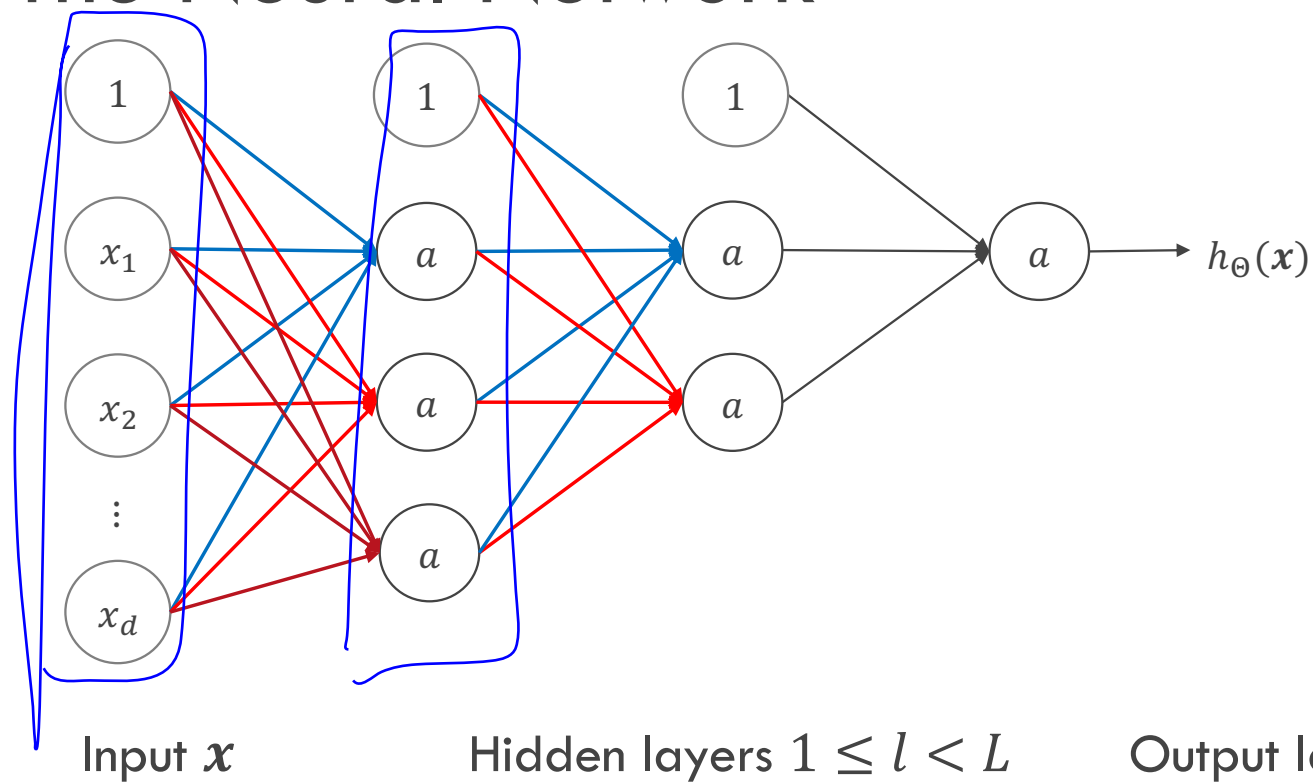
**2 red flags:**

for generalization and  
for optimization

Slide Credits: CS3244

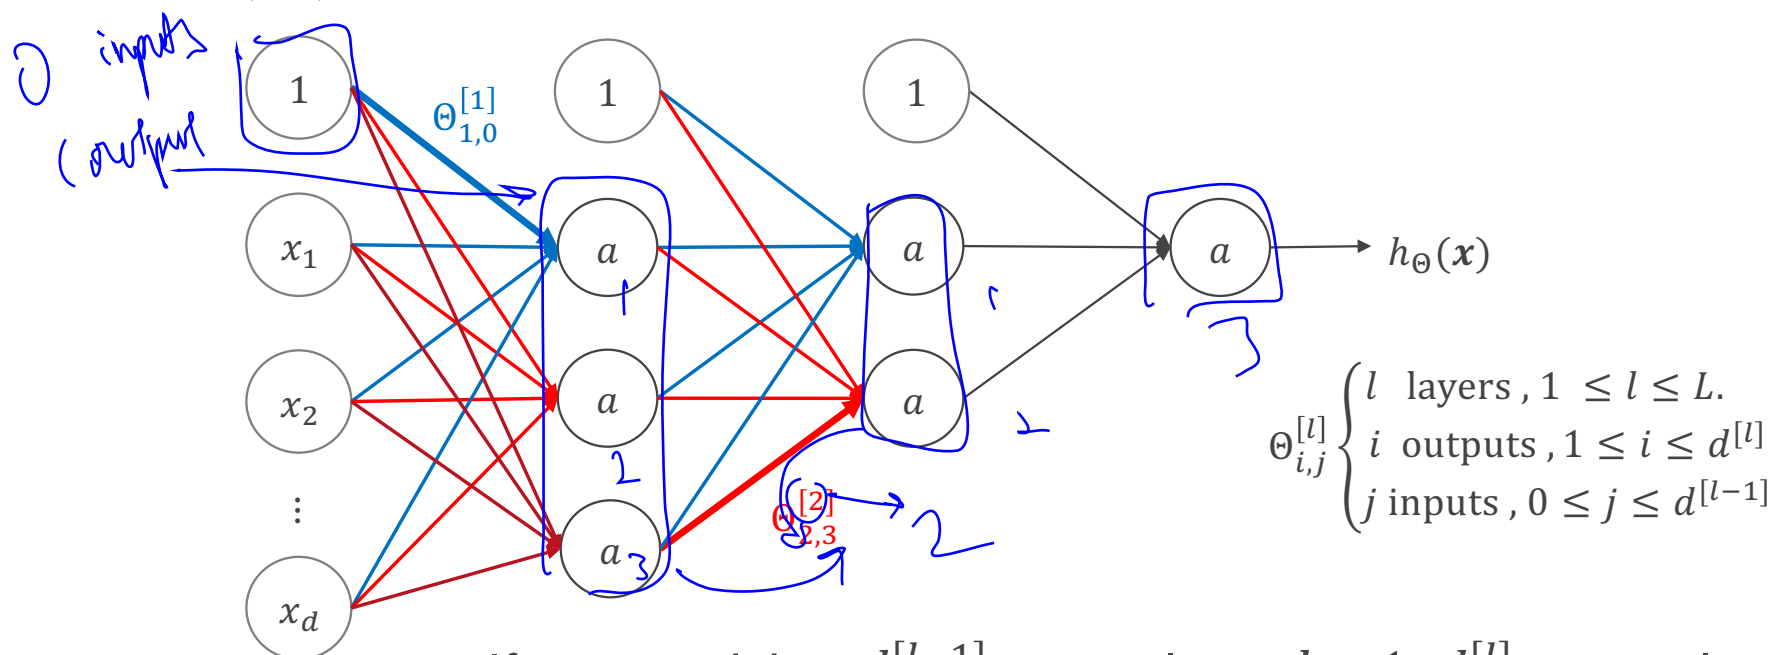
# Neural Networks

# The Neural Network



Slide Credits: CS3244

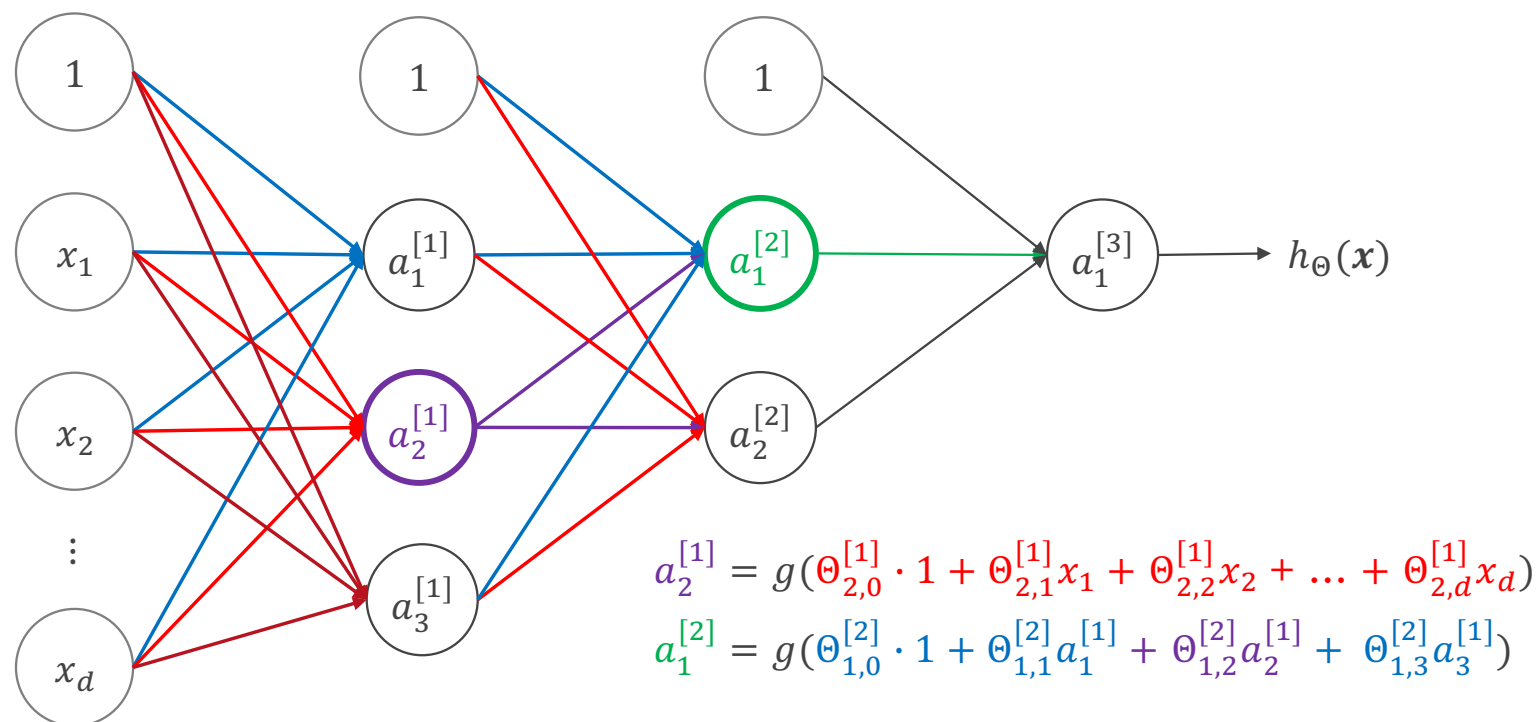
# NN Indices



If a network has  $d^{[l-1]}$  units in layer  $l - 1$ ,  $d^{[l]}$  units at layer  $l$ ,  
 then  $\Theta^{[l]}$  will be of dimension  $(d^{[l-1]} + 1) \times d^{[l]}$ .

Slide Credits: CS3244

# NN Activations



Slide Credits: CS3244

# How the Network Operates

$$\Theta_{i,j}^{[l]} \begin{cases} 1 \leq l \leq L & \text{layers} \\ 1 \leq i \leq d^{[l]} & \text{outputs} \\ 0 \leq j \leq d^{[l-1]} & \text{inputs} \end{cases}$$

$$\begin{aligned} x_i^{[l]} &= a_i^{[l]} = g(z_i^{[l]}) = g\left(\sum_{j=0}^{d^{[l-1]}} \Theta_{i,j}^{[l]} x_j^{[l-1]}\right) \\ &= g((\Theta_i^{[l]}) \cdot \mathbf{x}^{[l-1]}) \end{aligned}$$

Feedforward: Apply  $\mathbf{x}$  to  $x_1^{[0]} \dots x_{d^{[0]}}^{[0]} \rightarrow x_1^{[L]} = h(\mathbf{x})$

Slide Credits: CS3244

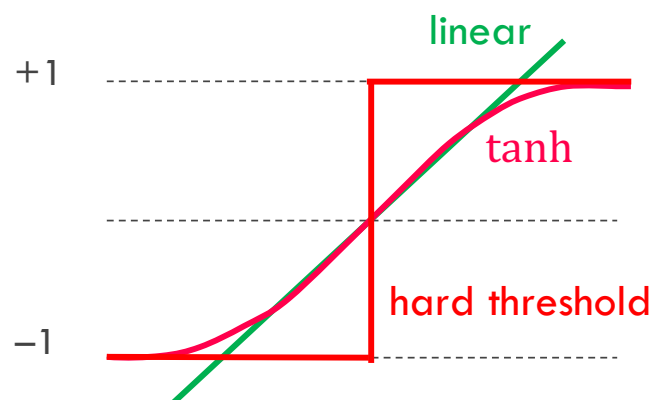
# Activation Functions

map sigmoid from  $[0,1]$  to  $[-1,1]$ .

## Final Layer

Replace sigmoid  $[0,1]$  with hyperbolic tangent  $[-1,1]$  if we want -ve and +ve classes.

$$g(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

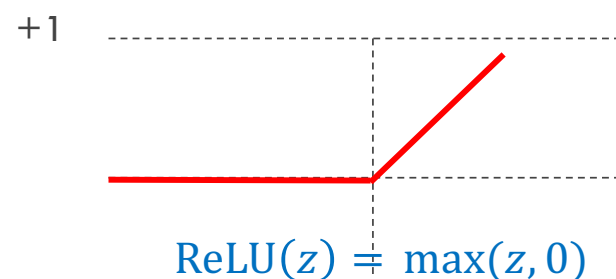


## Intermediate Layer(s):

Outputs are features for an upstream unit.

Don't need a probabilistic interpretation, since features can have any value; just need to be non-linear.

Use the Rectified Linear Unit (ReLU):



if result zero, neuron does not get activated.



# NN Summary

Intermediate NN layers create real valued features.

- Use a simpler activation function, the **ReLU**.

With NN being complex, our loss function is no longer convex with one minimum.

- Then SGD finds a local minimum, rather than a global.
- Good initialization is more important.

The overfitting issue is more extreme in NNs due to their complexity.

- Regularization more important
- One method you'll hear about: **dropout**.